

**Matrix computations
on the GPU**
CUBLAS, CUSOLVER and MAGMA by example

Andrzej Chrzęszczyk

Jan Kochanowski University, Kielce, Poland

Jacob Anders

CSIRO, Canberra, Australia

Version 2017

Foreword

Many scientific computer applications need high-performance matrix algebra. The major hardware developments always influenced new developments in linear algebra libraries. For example in the 80's the cache-based machines appeared and LAPACK based on Level 3 BLAS was developed. In the 90's new parallel platforms influenced ScaLAPACK developments.

To fully exploit the power of current heterogeneous systems of multi/many core CPUs and GPUs (Graphics Processing Units) new tools are needed. The main purpose of this document is to present three of them, CUBLAS, MAGMA and CUSOLVER linear algebra C/C++ libraries.

We propose a practical, hands-on approach. We show how to install and use these libraries. The detailed table of contents allows for easy navigation through over 200 code samples. We believe that the presented document can be an useful addition to the existing documentation for CUBLAS, CUSOLVER and MAGMA.

Remarks on using unified memory.

Unified memory is a single memory address space which allows applications to allocate data, that can be read or written from code running on either CPU or GPU. Unification of memory spaces means that there is no need for explicit memory transfers between host and device. This makes the CUDA programming easier. Allocating unified memory is as simple as replacing calls to `malloc` or `cudaMalloc` with calls to `cudaMallocManaged`. When

code running on CPU or GPU accesses data allocated this way, the CUDA system takes care of migrating memory pages to the memory of the accessing processor. Let us note however, that a carefully tuned CUDA program that uses streams and `cudaMemcpyAsync` to efficiently overlap execution with data transfer may perform better than a CUDA program that only uses unified memory. Users of unified memory are still free to use `cudaMemcpy` or `cudaMemcpyAsync` for performance optimization. Additionally, applications can guide the driver using `cudaMemAdvise` and explicitly migrate memory using `cudaMemPrefetchAsync`. Note also that unified memory examples, which do not call `cudaMemcpy`, require an explicit `cudaDeviceSynchronize` before the host program can safely use the output from the GPU. The memory allocated with `cudaMallocManaged` should be released with `cudaFree`.

Our main purpose is to show a set of examples containing matrix computations on GPUs which are *easy to understand*. On the other hand, the performance is the main reason for using GPUs in matrix computations. Therefore we have decided to present (almost) *all examples in two versions*. First we demonstrate a traditional version with explicit data copying between host and device. Next we give the same examples using the unified memory. The second approach is simpler and probably more appropriate for beginners. The users which need more control on the data will probably prefer the first one or will combine both.

Contents

Foreword	1
1 CUDA Toolkit	4
1.1 Installing CUDA Toolkit	4
1.2 Measuring GPUs performance	5
2 CUBLAS by example	8
2.1 General remarks on the examples	8
2.2 CUBLAS Level-1. Scalar and vector based operations	9
2.2.1 cublasIsamax, cublasIsamin - maximal, minimal elements	9
2.2.2 cublasIsamax, cublasIsamin - unified memory ver- sion	11
2.2.3 cublasSasum - sum of absolute values	11
2.2.4 cublasSasum - unified memory version	12
2.2.5 cublasSaxpy - compute $\alpha x + y$	13
2.2.6 cublasSaxpy - unified memory version	14
2.2.7 cublasScopy - copy vector into vector	15
2.2.8 cublasScopy - unified memory version	16
2.2.9 cublasSdot - dot product	17
2.2.10 cublasSdot - unified memory version	18
2.2.11 cublasSnrm2 - Euclidean norm	19
2.2.12 cublasSnrm2 - unified memory version	20
2.2.13 cublasSrot - apply the Givens rotation	21
2.2.14 cublasSrot - unified memory version	23
2.2.15 cublasSrotg - construct the Givens rotation matrix	24
2.2.16 cublasSrotm - apply the modified Givens rotation	25
2.2.17 cublasSrotm - unified memory version	27
2.2.18 cublasSrotmg - construct the modified Givens rota- tion matrix	28
2.2.19 cublasSscal - scale the vector	29
2.2.20 cublasSscal - unified memory version	30
2.2.21 cublasSswap - swap two vectors	31
2.2.22 cublasSswap - unified memory version	33

2.3	CUBLAS Level-2. Matrix-vector operations	34
2.3.1	cublasSgbmv – banded matrix-vector multiplication	34
2.3.2	cublasSgbmv – unified memory version	36
2.3.3	cublasSgemv – matrix-vector multiplication	37
2.3.4	cublasSgemv – unified memory version	39
2.3.5	cublasSger – rank one update	40
2.3.6	cublasSger – unified memory version	42
2.3.7	cublasSsbmv – symmetric banded matrix-vector multiplication	44
2.3.8	cublasSsbmv – unified memory version	46
2.3.9	cublasSspmv – symmetric packed matrix-vector multiplication	47
2.3.10	cublasSspmv – unified memory version	49
2.3.11	cublasSspr – symmetric packed rank-1 update	50
2.3.12	cublasSspr – unified memory version	52
2.3.13	cublasSspr2 – symmetric packed rank-2 update	53
2.3.14	cublasSspr2 – unified memory version	56
2.3.15	cublasSsymv – symmetric matrix-vector multiplication	57
2.3.16	cublasSsymv – unified memory version	59
2.3.17	cublasSsyr – symmetric rank-1 update	61
2.3.18	cublasSsyr – unified memory version	63
2.3.19	cublasSsyr2 – symmetric rank-2 update	64
2.3.20	cublasSsyr2 – unified memory version	66
2.3.21	cublasStbmv – triangular banded matrix-vector multiplication	68
2.3.22	cublasStbmv – unified memory version	70
2.3.23	cublasStbsv – solve the triangular banded linear system	71
2.3.24	cublasStbsv – unified memory version	72
2.3.25	cublasStpmv – triangular packed matrix-vector multiplication	73
2.3.26	cublasStpmv – unified memory version	75
2.3.27	cublasStpsv – solve the packed triangular linear system	76
2.3.28	cublasStpsv – unified memory version	77
2.3.29	cublasStrmv – triangular matrix-vector multiplication	78
2.3.30	cublasStrmv – unified memory version	80
2.3.31	cublasStrsv – solve the triangular linear system	81
2.3.32	cublasStrsv – unified memory version	83
2.3.33	cublasChemv – Hermitian matrix-vector multiplication	84
2.3.34	cublasChemv – unified memory version	86
2.3.35	cublasChbmv – Hermitian banded matrix-vector multiplication	88

2.3.36	cublasChbm - unified memory version	89
2.3.37	cublasChpm - Hermitian packed matrix-vector mul- tiplication	91
2.3.38	cublasChpm - unified memory version	93
2.3.39	cublasCher - Hermitian rank-1 update	94
2.3.40	cublasCher - unified memory version	96
2.3.41	cublasCher2 - Hermitian rank-2 update	98
2.3.42	cublasCher2 - unified memory version	100
2.3.43	cublasChpr - packed Hermitian rank-1 update	102
2.3.44	cublasChpr - unified memory version	104
2.3.45	cublasChpr2 - packed Hermitian rank-2 update . . .	106
2.3.46	cublasChpr2 - unified memory version	108
2.4	CUBLAS Level-3. Matrix-matrix operations	110
2.4.1	cublasSgem - matrix-matrix multiplication	110
2.4.2	cublasSgem - unified memory version	113
2.4.3	cublasSsym - symmetric matrix-matrix multiplication	115
2.4.4	cublasSsym - unified memory version	117
2.4.5	cublasSsy - symmetric rank-k update	119
2.4.6	cublasSsy - unified memory version	122
2.4.7	cublasSsy2k - symmetric rank-2k update	124
2.4.8	cublasSsy2k - unified memory version	126
2.4.9	cublasStrmm - triangular matrix-matrix multiplication	128
2.4.10	cublasStrmm - unified memory version	131
2.4.11	cublasStrsm - solving the triangular linear system .	133
2.4.12	cublasStrsm - unified memory version	135
2.4.13	cublasChem - Hermitian matrix-matrix multiplication	137
2.4.14	cublasChem - unified memory version	140
2.4.15	cublasCherk - Hermitian rank-k update	142
2.4.16	cublasCherk - unified memory version	145
2.4.17	cublasCher2k - Hermitian rank-2k update	147
2.4.18	cublasCher2k - unified memory version	149
3	CUSOLVER by example	152
3.1	General remarks on cuSolver	152
3.1.1	Remarks on installation and compilation	153
3.1.2	Remarks on hardware used in examples	153
3.2	LU decomposition and solving general linear systems . . .	154
3.2.1	cusolverDnSgetrf and cusolverDnSgetrs - solving general linear system using LU decomposition in sin- gle precision	154
3.2.2	cusolverDnSgetrf and cusolverDnSgetrs - unified memory version	156

3.2.3	<code>cusolverDnDgetrf</code> and <code>cusolverDnDgetrs</code> - solving general linear system using LU decomposition in double precision	157
3.2.4	<code>cusolverDnDgetrf</code> and <code>cusolverDnDgetrs</code> - unified memory version	160
3.3	QR decomposition and solving general linear systems	161
3.3.1	<code>cusolverDnSgeqrf</code> and <code>cusolverDnSorgqr</code> - QR decomposition and checking the orthogonality in single precision	161
3.3.2	<code>cusolverDnSgeqrf</code> and <code>cusolverDnSorgqr</code> - unified memory version	164
3.3.3	<code>cusolverDnDgeqrf</code> and <code>cusolverDnDorgqr</code> - QR decomposition and checking the orthogonality in double precision	166
3.3.4	<code>cusolverDnDgeqrf</code> and <code>cusolverDnDorgqr</code> - unified memory version	168
3.3.5	<code>cusolverDnSgeqrf</code> and <code>cusolverDnSormqr</code> , <code>cublasStrsm</code> - QR decomposition and solving a linear system in single precision	170
3.3.6	<code>cusolverDnSgeqrf</code> and <code>cusolverDnSormqr</code> , <code>cublasStrsm</code> - unified memory version	173
3.3.7	<code>cusolverDnDgeqrf</code> and <code>cusolverDnDormqr</code> , <code>cublasDtrsm</code> - QR decomposition and solving a linear system in double precision	175
3.3.8	<code>cusolverDnDgeqrf</code> and <code>cusolverDnDormqr</code> , <code>cublasDtrsm</code> - unified memory version	178
3.4	Cholesky decomposition and solving positive definite linear systems	180
3.4.1	<code>cusolverDnSpotrf</code> and <code>cusolverDnSpotrs</code> - Choleski decomposition and solving positive definite systems in single precision	180
3.4.2	<code>cusolverDnSpotrf</code> and <code>cusolverDnSpotrs</code> - unified memory version	182
3.4.3	<code>cusolverDnDpotrf</code> and <code>cusolverDnDpotrs</code> - Choleski decomposition and solving positive definite systems in double precision	184
3.4.4	<code>cusolverDnDpotrf</code> and <code>cusolverDnDpotrs</code> - unified memory version	186
3.5	Bunch-Kaufman decomposition and solving symmetric linear systems	188
3.5.1	<code>cusolverDnSsytrf</code> and <code>ssytrs</code> - Bunch-Kaufman decomposition and solving symmetric systems in single precision	188

3.5.2	<code>cusolverDnSsytrf</code> and <code>ssytrs</code> - unified memory version	190
3.5.3	<code>cusolverDnDsytrf</code> and <code>dsytrs</code> - Bunch-Kaufman decomposition and solving symmetric systems in double precision	191
3.5.4	<code>cusolverDnDsytrf</code> and <code>dsytrs</code> - unified memory version	194
3.6	SVD decomposition	195
3.6.1	<code>cusolverDnSgesvd</code> - SVD decomposition in single precision	195
3.6.2	<code>cusolverDnSgesvd</code> - unified memory version	198
3.6.3	<code>cusolverDnDgesvd</code> - SVD decomposition in double precision	200
3.6.4	<code>cusolverDnDgesvd</code> - unified memory version	202
3.7	Eigenvalues and eigenvectors for symmetric matrices	204
3.7.1	<code>cusolverDnSsyevd</code> - eigenvalues and eigenvectors for symmetric matrices in single precision	204
3.7.2	<code>cusolverDnSsyevd</code> - unified memory version	206
3.7.3	<code>cusolverDnDsyevd</code> - eigenvalues and eigenvectors for symmetric matrices in double precision	207
3.7.4	<code>cusolverDnDsyevd</code> - unified memory version	209
4	MAGMA by example	211
4.1	General remarks on Magma	211
4.1.1	Remarks on installation and compilation	213
4.1.2	Remarks on hardware used in examples	213
4.2	Magma BLAS	214
4.2.1	<code>magma_isamax</code> - find element with maximal absolute value	214
4.2.2	<code>magma_isamax</code> - unified memory version	215
4.2.3	<code>magma_sswap</code> - vectors swapping	215
4.2.4	<code>magma_sswap</code> - unified memory version	216
4.2.5	<code>magma_sgemv</code> - matrix-vector multiplication	217
4.2.6	<code>magma_sgemv</code> - unified memory version	219
4.2.7	<code>magma_ssylv</code> - symmetric matrix-vector multiplication	220
4.2.8	<code>magma_ssylv</code> - unified memory version	222
4.2.9	<code>magma_sgemm</code> - matrix-matrix multiplication	223
4.2.10	<code>magma_sgemm</code> - unified memory version	224
4.2.11	<code>magma_ssylv</code> - symmetric matrix-matrix multiplication	226
4.2.12	<code>magma_ssylv</code> - unified memory version	228
4.2.13	<code>magma_ssyrrk</code> - symmetric rank-k update	229
4.2.14	<code>magma_ssyrrk</code> - unified memory version	231
4.2.15	<code>magma_ssyrrk2k</code> - symmetric rank-2k update	232

4.2.16	<code>magma_ssyrr2k</code> - unified memory version	234
4.2.17	<code>magma_strmm</code> - triangular matrix-matrix multiplication	235
4.2.18	<code>magma_strmm</code> - unified memory version	237
4.2.19	<code>magmablas_sgeadd</code> - matrix-matrix addition	238
4.2.20	<code>magmablas_sgeadd</code> - unified memory version	240
4.3	LU decomposition and solving general linear systems	241
4.3.1	<code>magma_sgesv</code> - solve a general linear system in single precision, CPU interface	241
4.3.2	<code>magma_sgesv</code> - unified memory version	243
4.3.3	<code>magma_dgesv</code> - solve a general linear system in double precision, CPU interface	244
4.3.4	<code>magma_dgesv</code> - unified memory version	246
4.3.5	<code>magma_sgesv_gpu</code> - solve a general linear system in single precision, GPU interface	247
4.3.6	<code>magma_sgesv_gpu</code> - unified memory version	249
4.3.7	<code>magma_dgesv_gpu</code> - solve a general linear system in double precision, GPU interface	251
4.3.8	<code>magma_dgesv_gpu</code> - unified memory version	252
4.3.9	<code>magma_sgetrf</code> , <code>lapackf77_sgetrs</code> - LU factorization and solving factorized systems in single precision, CPU interface	254
4.3.10	<code>magma_sgetrf</code> , <code>lapackf77_sgetrs</code> - unified memory version	256
4.3.11	<code>magma_dgetrf</code> , <code>lapackf77_dgetrs</code> - LU factorization and solving factorized systems in double precision, CPU interface	257
4.3.12	<code>magma_dgetrf</code> , <code>lapackf77_dgetrs</code> - unified memory version	259
4.3.13	<code>magma_sgetrf_gpu</code> , <code>magma_sgetrs_gpu</code> - LU factorization and solving factorized systems in single precision, GPU interface	260
4.3.14	<code>magma_sgetrf_gpu</code> , <code>magma_sgetrs_gpu</code> - unified memory version	263
4.3.15	<code>magma_dgetrf_gpu</code> , <code>magma_dgetrs_gpu</code> - LU factorization and solving factorized systems in double precision, GPU interface	264
4.3.16	<code>magma_dgetrf_gpu</code> , <code>magma_dgetrs_gpu</code> - unified memory version	266
4.3.17	<code>magma_sgetrf_mgpu</code> - LU factorization in single precision on multiple GPUs	268
4.3.18	<code>magma_dgetrf_mgpu</code> - LU factorization in double precision on multiple GPUs	271

4.3.19	<code>magma_sgetri_gpu</code> - inverse matrix in single precision, GPU interface	273
4.3.20	<code>magma_sgetri_gpu</code> - unified memory version	275
4.3.21	<code>magma_dgetri_gpu</code> - inverse matrix in double precision, GPU interface	276
4.3.22	<code>magma_dgetri_gpu</code> - unified memory version	278
4.4	Cholesky decomposition and solving systems with positive definite matrices	280
4.4.1	<code>magma_sposv</code> - solve a system with a positive definite matrix in single precision, CPU interface	280
4.4.2	<code>magma_sposv</code> - unified memory version	281
4.4.3	<code>magma_dposv</code> - solve a system with a positive definite matrix in double precision, CPU interface	283
4.4.4	<code>magma_dposv</code> - unified memory version	284
4.4.5	<code>magma_sposv_gpu</code> - solve a system with a positive definite matrix in single precision, GPU interface	286
4.4.6	<code>magma_sposv_gpu</code> - unified memory version	288
4.4.7	<code>magma_dposv_gpu</code> - solve a system with a positive definite matrix in double precision, GPU interface	289
4.4.8	<code>magma_dposv_gpu</code> - unified memory version	291
4.4.9	<code>magma_spotrf</code> , <code>lapackf77_spotrs</code> - Cholesky decomposition and solving a system with a positive definite matrix in single precision, CPU interface	293
4.4.10	<code>magma_spotrf</code> , <code>lapackf77_spotrs</code> - unified memory version	294
4.4.11	<code>magma_dpotrf</code> , <code>lapackf77_dpotrs</code> - Cholesky decomposition and solving a system with a positive definite matrix in double precision, CPU interface	296
4.4.12	<code>magma_dpotrf</code> , <code>lapackf77_dpotrs</code> - unified memory version	298
4.4.13	<code>magma_spotrf_gpu</code> , <code>magma_spotrs_gpu</code> - Cholesky decomposition and solving a system with a positive definite matrix in single precision, GPU interface	299
4.4.14	<code>magma_spotrf_gpu</code> , <code>magma_spotrs_gpu</code> - unified memory version	301
4.4.15	<code>magma_dpotrf_gpu</code> , <code>magma_dpotrs_gpu</code> - Cholesky decomposition and solving a system with a positive definite matrix in double precision, GPU interface	303
4.4.16	<code>magma_dpotrf_gpu</code> , <code>magma_dpotrs_gpu</code> - unified memory version	305

4.4.17	<code>magma_spotrf_mgpu</code> , <code>lapackf77_spotrs</code> - Cholesky decomposition on multiple GPUs and solving a system with a positive definite matrix in single precision	307
4.4.18	<code>magma_dpotrf_mgpu</code> , <code>lapackf77_dpotrs</code> - Cholesky decomposition and solving a system with a positive definite matrix in double precision on multiple GPUs	309
4.4.19	<code>magma_spotri</code> - invert a symmetric positive definite matrix in single precision, CPU interface	312
4.4.20	<code>magma_spotri</code> - unified memory version	314
4.4.21	<code>magma_dpotri</code> - invert a positive definite matrix in double precision, CPU interface	315
4.4.22	<code>magma_dpotri</code> - unified memory version	316
4.4.23	<code>magma_spotri_gpu</code> - invert a positive definite matrix in single precision, GPU interface	317
4.4.24	<code>magma_spotri_gpu</code> - unified memory version	319
4.4.25	<code>magma_dpotri_gpu</code> - invert a positive definite matrix in double precision, GPU interface	320
4.4.26	<code>magma_dpotri_gpu</code> - unified memory version	322
4.5	QR decomposition and the least squares solution of general systems	324
4.5.1	<code>magma_sgels_gpu</code> - the least squares solution of a linear system using QR decomposition in single precision, GPU interface	324
4.5.2	<code>magma_sgels_gpu</code> - unified memory version	326
4.5.3	<code>magma_dgels_gpu</code> - the least squares solution of a linear system using QR decomposition in double precision, GPU interface	328
4.5.4	<code>magma_dgels_gpu</code> - unified memory version	331
4.5.5	<code>magma_sgels3_gpu</code> - the least squares solution of a linear system using QR decomposition in single precision, GPU interface	333
4.5.6	<code>magma_sgels3_gpu</code> - unified memory version	336
4.5.7	<code>magma_dgels3_gpu</code> - the least squares solution of a linear system using QR decomposition in double precision, GPU interface	338
4.5.8	<code>magma_dgels3_gpu</code> - unified memory version	341
4.5.9	<code>magma_sgeqrf</code> - QR decomposition in single precision, CPU interface	344
4.5.10	<code>magma_sgeqrf</code> - unified memory version	345
4.5.11	<code>magma_dgeqrf</code> - QR decomposition in double precision, CPU interface	347
4.5.12	<code>magma_dgeqrf</code> - unified memory version	348

4.5.13	<code>magma_sgeqrf_gpu</code> - QR decomposition in single precision, GPU interface	350
4.5.14	<code>magma_sgeqrf_gpu</code> - unified memory version	351
4.5.15	<code>magma_dgeqrf_gpu</code> - QR decomposition in double precision, GPU interface	353
4.5.16	<code>magma_dgeqrf_gpu</code> - unified memory version	355
4.5.17	<code>magma_sgeqrf_mgpu</code> - QR decomposition in single precision on multiple GPUs	356
4.5.18	<code>magma_dgeqrf_mgpu</code> - QR decomposition in double precision on multiple GPUs	358
4.5.19	<code>magma_sgelqf</code> - LQ decomposition in single precision, CPU interface	361
4.5.20	<code>magma_sgelqf</code> - unified memory version	362
4.5.21	<code>magma_dgelqf</code> - LQ decomposition in double precision, CPU interface	364
4.5.22	<code>magma_dgelqf</code> - unified memory version	365
4.5.23	<code>magma_sgelqf_gpu</code> - LQ decomposition in single precision, GPU interface	367
4.5.24	<code>magma_sgelqf_gpu</code> - unified memory version	369
4.5.25	<code>magma_dgelqf_gpu</code> - LQ decomposition in double precision, GPU interface	370
4.5.26	<code>magma_dgelqf_gpu</code> - unified memory version	372
4.5.27	<code>magma_sgeqp3</code> - QR decomposition with column pivoting in single precision, CPU interface	373
4.5.28	<code>magma_sgeqp3</code> - unified memory version	375
4.5.29	<code>magma_dgeqp3</code> - QR decomposition with column pivoting in double precision, CPU interface	376
4.5.30	<code>magma_dgeqp3</code> - unified memory version	378
4.6	Eigenvalues and eigenvectors for general matrices	379
4.6.1	<code>magma_sgeev</code> - compute the eigenvalues and optionally eigenvectors of a general real matrix in single precision, CPU interface, small matrix	379
4.6.2	<code>magma_sgeev</code> - unified memory version, small matrix	381
4.6.3	<code>magma_dgeev</code> - compute the eigenvalues and optionally eigenvectors of a general real matrix in double precision, CPU interface, small matrix	383
4.6.4	<code>magma_dgeev</code> - unified memory version, small matrix	385
4.6.5	<code>magma_sgeev</code> - compute the eigenvalues and optionally eigenvectors of a general real matrix in single precision, CPU interface, big matrix	387
4.6.6	<code>magma_sgeev</code> - unified memory version, big matrix	389

4.6.7	<code>magma_dgeev</code> - compute the eigenvalues and optionally eigenvectors of a general real matrix in double precision, CPU interface, big matrix	390
4.6.8	<code>magma_dgeev</code> - unified memory version, big matrix . .	392
4.6.9	<code>magma_sgehrd</code> - reduce a general matrix to the upper Hessenberg form in single precision, CPU interface .	393
4.6.10	<code>magma_sgehrd</code> - unified memory version	395
4.6.11	<code>magma_dgehrd</code> - reduce a general matrix to the upper Hessenberg form in double precision, CPU interface .	397
4.6.12	<code>magma_dgehrd</code> - unified memory version	399
4.7	Eigenvalues and eigenvectors for symmetric matrices	401
4.7.1	<code>magma_ssyevd</code> - compute the eigenvalues and optionally eigenvectors of a symmetric real matrix in single precision, CPU interface, small matrix	401
4.7.2	<code>magma_ssyevd</code> - unified memory version	403
4.7.3	<code>magma_dsyevd</code> - compute the eigenvalues and optionally eigenvectors of a symmetric real matrix in double precision, CPU interface, small matrix	405
4.7.4	<code>magma_dsyevd</code> - unified memory version	407
4.7.5	<code>magma_ssyevd</code> - compute the eigenvalues and optionally eigenvectors of a symmetric real matrix in single precision, CPU interface, big matrix	409
4.7.6	<code>magma_ssyevd</code> - unified memory version, big matrix .	410
4.7.7	<code>magma_dsyevd</code> - compute the eigenvalues and optionally eigenvectors of a symmetric real matrix in double precision, CPU interface, big matrix	412
4.7.8	<code>magma_dsyevd</code> - unified memory version, big matrix .	413
4.7.9	<code>magma_ssyevd_gpu</code> - compute the eigenvalues and optionally eigenvectors of a symmetric real matrix in single precision, GPU interface, small matrix	415
4.7.10	<code>magma_ssyevd_gpu</code> - unified memory version, small matrix	417
4.7.11	<code>magma_dsyevd_gpu</code> - compute the eigenvalues and optionally eigenvectors of a symmetric real matrix in double precision, GPU interface, small matrix	419
4.7.12	<code>magma_dsyevd_gpu</code> - unified memory version, small matrix	421
4.7.13	<code>magma_ssyevd_gpu</code> - compute the eigenvalues and optionally eigenvectors of a symmetric real matrix in single precision, GPU interface, big matrix	423
4.7.14	<code>magma_ssyevd_gpu</code> - unified memory version, big matrix	425

4.7.15	<code>magma_dsyevd_gpu</code> - compute the eigenvalues and optionally eigenvectors of a symmetric real matrix in double precision, GPU interface, big matrix	426
4.7.16	<code>magma_dsyevd_gpu</code> - unified memory version, big matrix	428
4.8	Singular value decomposition	429
4.8.1	<code>magma_sgesvd</code> - compute the singular value decomposition of a general real matrix in single precision, CPU interface	429
4.8.2	<code>magma_sgesvd</code> - unified memory version	432
4.8.3	<code>magma_dgesvd</code> - compute the singular value decomposition of a general real matrix in double precision, CPU interface	433
4.8.4	<code>magma_dgesvd</code> - unified memory version	435
4.8.5	<code>magma_sgebrd</code> - reduce a real matrix to bidiagonal form by orthogonal transformations in single precision, CPU interface	437
4.8.6	<code>magma_sgebrd</code> - unified memory version	438
4.8.7	<code>magma_dgebrd</code> - reduce a real matrix to bidiagonal form by orthogonal transformations in double precision, CPU interface	440
4.8.8	<code>magma_dgebrd</code> - unified memory version	441

Chapter 1

CUDA Toolkit

1.1 Installing CUDA Toolkit

All libraries described in our text: CUBLAS, CUSOLVER and MAGMA need CUDA (Compute Unified Device Architecture) environment. In fact CUBLAS and CUSOLVER are parts of CUDA. The environment can be downloaded from <https://developer.nvidia.com/cuda-downloads>. On the same page one can find complete documentation, including installation instructions for Windows, Linux and Mac OSX. At the time of writing, the current release was CUDA v8.0. In our examples we shall use Ubuntu 16.04. In this system one can install CUDA using the command

```
# apt-get install cuda
```

Let us remark however that `apt-get` changes the paths and makes the installation of other tools more difficult.

If the CUDA software is installed and configured correctly, and the CUDA code samples are copied to the `$HOME` directory, the executable:

```
$ ~/NVIDIA_CUDA-8.0_Samples/1_Uutilities/deviceQuery/deviceQuery
```

should display the properties of the detected CUDA devices.

The `nbody` executable:

```
$ ~/NVIDIA_CUDA-8.0_Samples/5_Simulations/nbody/  
nbody -benchmark -numbodies=256000 -numdevices=1  
# (in the case of one device)
```

gives the opportunity to check GPU performance. On GeForce GTX 1080 card, one can obtain for example

```
Compute 6.1 CUDA device: [GeForce GTX 1080]  
number of bodies = 256000  
256000 bodies, total time for 10 iterations: 2413.114 ms  
= 271.583 billion interactions per second  
= 5431.653 single-precision GFLOP/s at 20 flops per interaction
```

The state of devices can be checked using

```
$ nvidia-smi          # man nvidia-smi
```

```
Wed Jul 19 11:29:44 2017
```

NVIDIA-SMI 378.13										Driver Version: 378.13																																							
-----+-----+-----																																																	
GPU Name					Persistence-M					Bus-Id					Disp.A					Volatile					Uncorr. ECC																								
Fan Temp Perf					Pwr:Usage/Cap					Memory-Usage					GPU-Util					Compute M.																													
=====+=====+=====																																																	
0 GeForce GTX 1080					Off					0000:01:00.0					On										N/A																								
24% 36C P8					10W / 180W					222MiB /					8110MiB										0%					Default																			
-----+-----+-----																																																	
+-----+-----+-----																																																	
Processes:																														GPU Memory																			
GPU					PID Type					Process name										Usage																													
=====+=====+=====																																																	
0					1006 G					/usr/lib/xorg/Xorg										131MiB																													
0					1546 G					compiz										88MiB																													
-----+-----+-----																																																	

1.2 Measuring GPUs performance

It seems that one of the simplest ways of benchmarking systems with GPU devices is to use the Magma library. The library will be introduced in one of the next chapters but now let us remark, that as an by-product of Magma installation, one obtains the directory `testing` with ready to use testing binaries. Bellow we present the results of running four of them. We tested a system with Linux Ubuntu 16.04 and

- Intel i7 6700K CPU, 16GB RAM
- Nvidia GeForce GTX 1080 card,
- magma-2.2.0 compiled with OpenBLAS.

As we remarked such a system seems to be sufficient for training and single precision calculations on GPU, but GeForce GTX cards have strongly restricted double precision capabilities. As a consequence in the benchmarks we present, the double precision performance is not so impressive. Using professional cards, one can expect that the double precision functions are only two times slower than the corresponding single precision ones.

Solving the general NxN linear system in single precision.

```
./testing_sgesv --lapack
%   N  NRHS   CPU Gflop/s (sec)   GPU Gflop/s (sec)
%=====
1088    1    75.89 (  0.01)    74.72 (  0.01)
2112    1   181.11 (  0.03)   230.11 (  0.03)
3136    1   208.44 (  0.10)   438.46 (  0.05)
4160    1   227.21 (  0.21)   494.41 (  0.10)
5184    1   240.69 (  0.39)   632.86 (  0.15)
6208    1   250.69 (  0.64)   778.43 (  0.20)
7232    1   266.17 (  0.95)   920.08 (  0.27)
8256    1   273.73 (  1.37)  1072.90 (  0.35)
9280    1   285.11 (  1.87)  1220.60 (  0.44)
```

Let us repeat that the executables used in this section are contained in **testing** subdirectory of Magma installation directory.

Solving the general NxN linear system in double precision.

```
./testing_dgesv --lapack
%   N  NRHS   CPU Gflop/s (sec)   GPU Gflop/s (sec)
%=====
1088    1    86.97 (  0.01)    58.13 (  0.01)
2112    1    92.03 (  0.07)   117.56 (  0.05)
3136    1   102.69 (  0.20)   145.29 (  0.14)
4160    1   111.11 (  0.43)   179.07 (  0.27)
5184    1   120.00 (  0.77)   200.15 (  0.46)
6208    1   126.85 (  1.26)   211.75 (  0.75)
7232    1   132.66 (  1.90)   219.35 (  1.15)
8256    1   137.11 (  2.74)   225.68 (  1.66)
9280    1   142.41 (  3.74)   232.11 (  2.30)
10304   1   145.79 (  5.00)   236.96 (  3.08)
```

Matrix-matrix product in single precision.

```
./testing_sgemm --lapack
% transA = No transpose, transB = No transpose
%   M   N   K   MAGMA Gflop/s (ms)   cuBLAS Gflop/s (ms)   CPU Gflop/s (ms)
%=====
1088 1088 1088 2985.30 (  0.86)   5377.70 (  0.48)   344.92 (  7.47)
2112 2112 2112 4317.43 (  4.36)   6461.15 (  2.92)   317.30 ( 59.38)
3136 3136 3136 4507.27 ( 13.68)   6807.34 (  9.06)   347.10 (177.71)
4160 4160 4160 4531.90 ( 31.77)   6897.06 ( 20.88)   380.82 (378.08)
5184 5184 5184 5212.77 ( 53.45)   7501.88 ( 37.14)   390.93 (712.74)
6208 6208 6208 4896.58 ( 97.72)   7407.86 ( 64.59)   398.93 (1199.45)
7232 7232 7232 5021.23 (150.66)   7600.73 ( 99.53)   401.84 (1882.57)
8256 8256 8256 4974.09 (226.27)   7444.21 (151.19)   414.80 (2713.30)
9280 9280 9280 4931.62 (324.10)   7460.25 (214.25)   412.64 (3873.52)
10304 10304 10304 4820.21 (453.92)   7505.78 (291.51)   418.97 (5222.33)
```

Matrix-matrix product in double precision.

```

./testing_dgemm --lapack
% transA = No transpose, transB = No transpose
%   M      N      K   MAGMA Gflop/s (ms)   cuBLAS Gflop/s (ms)   CPU Gflop/s (ms)
%=====
1088 1088 1088   213.21 ( 12.08)    218.40 ( 11.79)    72.51 ( 35.52)
2112 2112 2112   220.69 ( 85.37)    231.89 ( 81.25)   165.29 (113.99)
3136 3136 3136   257.86 (239.21)    277.21 (222.51)   180.38 (341.96)
4160 4160 4160   261.32 (550.98)    278.40 (517.18)   183.05 (786.56)
5184 5184 5184   259.34 (1074.36)    278.91 (998.97)   190.42 (1463.25)
6208 6208 6208   258.53 (1850.89)    279.77 (1710.36)  190.80 (2507.84)
7232 7232 7232   260.83 (2900.31)    281.12 (2690.99)  194.81 (3883.22)
8256 8256 8256   260.32 (4323.45)    279.53 (4026.29)  196.29 (5733.71)
9280 9280 9280   258.77 (6176.75)    277.68 (5756.20)  197.17 (8106.52)
10304 10304 10304 257.44 (8499.15)    277.09 (7896.35)  196.48 (11135.99)

```

Chapter 2

CUBLAS by example

2.1 General remarks on the examples

CUBLAS is an abbreviation for CUDA Basic Linear Algebra Subprograms. In the file `/usr/local/cuda/doc/pdf/CUBLAS_Library.pdf` one can find a detailed description of the CUBLAS library syntax and we shall avoid to repeat the information contained there. Instead we present a series of examples how to use the library.

All subprograms have four versions corresponding to four data types

- `s,S` - `float` - real single-precision
- `d,D` - `double` - real double-precision,
- `c,C` - `cuComplex` - complex single-precision,
- `z,Z` - `cuDoubleComplex` -complex double-precision.

For example `cublasI<t>amax` is a template which can represent `cublasIsamax`, `cublasIdamax`, `cublasIcamax` or `cublasIzamax`.

- We shall restrict our examples in this chapter to single precision versions. The reason is that low-end devices have restricted double precision capabilities. On the other hand the changes needed in the double precision case are not significant. In most examples we use real data but the complex cases are also considered (see the subsections with the title of the form `cublasC*`).
- CUBLAS Library User Guide contains an example showing how to check for errors returned by API calls. Ideally we should check for errors on every API call. Unfortunately such an approach doubles the length of our sample codes (which are as short as possible by design). Since our set of CUBLAS sample code (without error checking) is rather long, we have decided to ignore the error checking and to focus on the explanations which cannot be found in User Guide. The reader

can add the error checking code from CUBLAS Library User Guide example with minor modifications.

- To obtain more compact explanations in our examples we restrict the full generality of CUBLAS to the special case where the leading dimension of matrices is equal to the number of rows and the stride between consecutive elements of vectors is equal to 1. CUBLAS allows for more flexible approach giving the user the access to submatrices and subvectors. The corresponding explanations can be found in CUBLAS Library User Guide and in BLAS manual.

Remarks on compilation. All examples in this chapter contain simple compilation instructions. Notice that the examples, in which we use the unified memory, have the names of the form `example.cu`, while the other ones have the form `example.c`. The simplest compilation method we know is respectively

```
nvcc example.c -lcublas
```

and

```
nvcc example.cu -lcublas}
```

If the extension `c` is preferred in the second case, then all occurrences of the function `cudaMallocManaged` should have the third argument (integer) 1.

If `g++` command is preferred, then the syntax of the form

```
cudaMallocManaged((void*)&x,n*sizeof(float),1);
```

should be used instead of `cudaMallocManaged(&x,n*sizeof(float),1);`

the constant `EXIT_SUCCESS` should be replaced by 0, and the header `cuda_runtime_api.h` should be included.

An example of compilation with `g++`:

```
g++ 001isamaxu.c -I/usr/local/cuda-8.0/include
-L/usr/local/cuda/lib64 -lcuda -lcublas -lcudart
```

2.2 CUBLAS Level-1. Scalar and vector based operations

2.2.1 `cublasIsamax`, `cublasIsamin` - maximal, minimal elements

This function finds the smallest index of the element of an array with the maximum /minimum magnitude.

```

//nvcc 001isamax.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"
#define n 6 // length of x
int main(void){
    cudaError_t cudaStat; // cudaMalloc status
    cublasStatus_t stat; // CUBLAS functions status
    cublasHandle_t handle; // CUBLAS context
    int j; // index of elements
    float* x; // n-vector on the host
    x=(float *)malloc (n*sizeof(*x)); // host memory alloc
    for(j=0;j<n;j++)
        x[j]=(float)j; // x={0,1,2,3,4,5}
    printf("x: ");
    for(j=0;j<n;j++)
        printf("%4.0f",x[j]); // print x
    printf("\n");
    // on the device
    float* d_x; // d_x - x on the device
    cudaStat=cudaMalloc((void**)&d_x,n*sizeof(*x)); // device
    // memory alloc for x
    stat = cublasCreate(&handle); // initialize CUBLAS context
    stat = cublasSetVector(n,sizeof(*x),x,1,d_x,1); //cp x ->d_x
    int result; // index of the maximal/minimal element
    // find the smallest index of the element of d_x with maximum
    // absolute value

    stat=cublasIsamax(handle,n,d_x,1,&result);

    printf("max |x[i]|:%4.0f\n",fabs(x[result-1])); // print
    // max{|x[0]|,...,|x[n-1]|}
    // find the smallest index of the element of d_x with minimum
    // absolute value

    stat=cublasIsamin(handle,n,d_x,1,&result);

    printf("min |x[i]|:%4.0f\n",fabs(x[result-1])); // print
    // min{|x[0]|,...,|x[n-1]|}
    cudaFree(d_x); // free device memory
    cublasDestroy(handle); // destroy CUBLAS context
    free(x); // free host memory
    return EXIT_SUCCESS;
}
// x: 0, 1, 2, 3, 4, 5,
// max |x[i]|: 5
// min |x[i]|: 0

```

2.2.2 cublasIsamax, cublasIsamin - unified memory version

```
// nvcc 001isamax.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define n 6 // length of x
int main(void){
    cublasHandle_t handle; // CUBLAS context
    int j; // index of elements
    float* x; // n-vector
    cudaMallocManaged(&x,n*sizeof(float)); // unified mem.for x
    for(j=0;j<n;j++){
        x[j]=(float)j; // x={0,1,2,3,4,5}
    }
    printf("x: ");
    for(j=0;j<n;j++){
        printf("%4.0f",x[j]); // print x
    }
    printf("\n");
    cublasCreate(&handle); // initialize CUBLAS context
    int result; // index of the maximal/minimal element
    // find the smallest index of the element of x with maximal
    // absolute value

    cublasIsamax(handle,n,x,1,&result);

    cudaDeviceSynchronize();
    printf("max |x[i]|:%4.0f\n",fabs(x[result-1]));
    // max{|x[0]|,...,|x[n-1]|}

    // find the smallest index of the element of x with minimal
    // absolute value

    cublasIsamin(handle,n,x,1,&result);

    cudaDeviceSynchronize();
    printf("min |x[i]|:%4.0f\n",fabs(x[result-1]));
    // min{|x[0]|,...,|x[n-1]|}

    cudaFree(x); // free memory
    cublasDestroy(handle); // destroy CUBLAS context
    return EXIT_SUCCESS;
}
// x: 0, 1, 2, 3, 4, 5,
// max |x[i]|: 5
// min |x[i]|: 0
```

2.2.3 cublasSasum - sum of absolute values

This function computes the sum of the absolute values of the elements of an array.

```
//nvcc 003sasumVec.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
```

```

#include "cublas_v2.h"
#define n 6 // length of x
int main(void){
    cudaError_t cudaStat; // cudaMalloc status
    cublasStatus_t stat; // CUBLAS functions status
    cublasHandle_t handle; // CUBLAS context
    int j; // index of elements
    float* x; // n-vector on the host
    x=(float *)malloc (n*sizeof(*x)); // host memory alloc
    for(j=0;j<n;j++)
        x[j]=(float)j; // x={0,1,2,3,4,5}
    printf("x: ");
    for(j=0;j<n;j++)
        printf("%2.0f",x[j]); // print x
    printf("\n");
    // on the device
    float* d_x; // d_x - x on the device
    cudaStat=cudaMalloc((void**)&d_x,n*sizeof(*x)); //device
    // memory alloc
    stat = cublasCreate(&handle); // initialize CUBLAS context
    stat = cublasSetVector(n,sizeof(*x),x,1,d_x,1); // cp x->d_x
    float result;
    // add absolute values of elements of the array d_x:
    // |d_x[0]|+...+|d_x[n-1]|

    stat=cublasSasum(handle,n,d_x,1,&result);

    //print the result
    printf("sum of the absolute values of elements of x:%4.0f\n",
        result);
    cudaFree(d_x); // free device memory
    cublasDestroy(handle); // destroy CUBLAS context
    free(x); // free host memory
    return EXIT_SUCCESS;
}
// x: 0, 1, 2, 3, 4, 5,
// sum of the absolute values of elements of x: 15
// |0|+|1|+|2|+|3|+|4|+|5|=15

```

2.2.4 cublasSasum - unified memory version

```

//nvcc 003sasum.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define n 6 // length of x
int main(void){
    cublasHandle_t handle; // CUBLAS context
    int j; // index of elements
    float* x; // n-vector on the host
    cudaMallocManaged(&x,n*sizeof(float)); // unified mem.for x

```

```

for(j=0;j<n;j++)
    x[j]=(float)j;                                // x={0,1,2,3,4,5}
printf("x: ");
for(j=0;j<n;j++)
    printf("%2.0f",x[j]);                          // print x
printf("\n");
cublasCreate(&handle);                            // initialize CUBLAS context
float result;
// add absolute values of elements of the array x:
// |x[0]|+...+|x[n-1]|

cublasSasum(handle,n,x,1,&result);

cudaDeviceSynchronize();
//print the result
printf("sum of the absolute values of elements of x:%4.0f\n",
      result);

cudaFree(x);                                     // free memory
cublasDestroy(handle);                          // destroy CUBLAS context
return EXIT_SUCCESS;
}
// x:  0, 1, 2, 3, 4, 5,
// sum of the absolute values of elements of x:  15
// |0|+|1|+|2|+|3|+|4|+|5|=15

```

2.2.5 cublasSaxpy - compute $\alpha x + y$

This function multiplies the vector x by the scalar α and adds it to the vector y

$$y = \alpha x + y.$$

```

//nvcc 004saxpy.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"
#define n 6                                     // length of x,y
int main(void){
    cudaError_t cudaStat;                      // cudaMalloc status
    cublasStatus_t stat;                      // CUBLAS functions status
    cublasHandle_t handle;                   // CUBLAS context
    int j;                                    // index of elements
    float* x;                                // n-vector on the host
    float* y;                                // n-vector on the host
    x=(float *)malloc (n*sizeof(*x)); // host memory alloc for x
    for(j=0;j<n;j++)
        x[j]=(float)j;                        // x={0,1,2,3,4,5}
    y=(float *)malloc (n*sizeof(*y)); // host memory alloc for y
    for(j=0;j<n;j++)

```



```

    y[j]=(float)j;                                     // y={0,1,2,3,4,5}
    printf("x,y:\n");
    for(j=0;j<n;j++)
        printf("%2.0f",x[j]);                         // print x,y
    printf("\n");
    // on the device
    float* d_x;                                       // d_x - x on the device
    float* d_y;                                       // d_y - y on the device
    cudaStat=cudaMalloc((void**)&d_x,n*sizeof(*x)); //device
                                                // memory alloc for x
    cudaStat=cudaMalloc((void**)&d_y,n*sizeof(*y)); //device
                                                // memory alloc for y
    stat = cublasCreate(&handle); // initialize CUBLAS context
    stat = cublasSetVector(n,sizeof(*x),x,1,d_x,1); //cp x->d_x
    stat = cublasSetVector(n,sizeof(*y),y,1,d_y,1); //cp y->d_y
    float al=2.0;                                     // al=2
    // multiply the vector d_x by the scalar al and add to d_y
    // d_y = al*d_x + d_y, d_x,d_y - n-vectors; al - scalar

    stat=cublasSaxpy(handle,n,&al,d_x,1,d_y,1);

    stat=cublasGetVector(n,sizeof(float),d_y,1,y,1); //cp d_y->y
    printf("y after Saxpy:\n");                      // print y after Saxpy
    for(j=0;j<n;j++)
        printf("%2.0f",y[j]);
    printf("\n");
    cudaFree(d_x);                                    // free device memory
    cudaFree(d_y);                                    // free device memory
    cublasDestroy(handle);                            // destroy CUBLAS context
    free(x);                                           // free host memory
    free(y);                                           // free host memory
    return EXIT_SUCCESS;
}
// x,y:
// 0, 1, 2, 3, 4, 5,

// y after Saxpy:
// 0, 3, 6, 9,12,15, // 2*x+y = 2*{0,1,2,3,4,5} + {0,1,2,3,4,5}

```

2.2.6 cublasSaxpy - unified memory version

```

//nvcc 004saxpy.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define n 6                                           // length of x,y
int main(void){
    cublasHandle_t handle;                            // CUBLAS context
    int j;                                           // index of elements
    float* x;                                       // n-vector
    float* y;                                       // n-vector

```

```

    cudaMallocManaged(&x,n*sizeof(float)); // unified mem.for x
    for(j=0;j<n;j++)
        x[j]=(float)j;                      // x={0,1,2,3,4,5}
    cudaMallocManaged(&y,n*sizeof(float)); // unified mem.for y
    for(j=0;j<n;j++)
        y[j]=(float)j;                      // y={0,1,2,3,4,5}
    printf("x,y:\n");
    for(j=0;j<n;j++)
        printf("%2.0f",x[j]);               // print x,y
    printf("\n");
    cublasCreate(&handle);                   // initialize CUBLAS context
    float al=2.0;                            // al=2
    // multiply the vector x by the scalar al and add to y
    // y = al*x + y,    x,y - n-vectors; al - scalar

    cublasSaxpy(handle,n,&al,x,1,y,1);

    cudaDeviceSynchronize();
    printf("y after Saxpy:\n");              // print y after Saxpy
    for(j=0;j<n;j++)
        printf("%2.0f",y[j]);
    printf("\n");
    cudaFree(x);                             // free memory
    cudaFree(y);                             // free memory
    cublasDestroy(handle);                   // destroy CUBLAS context
    return EXIT_SUCCESS;
}
// x,y:
// 0, 1, 2, 3, 4, 5,

// y after Saxpy:
// 0, 3, 6, 9,12,15, // 2*x+y = 2*{0,1,2,3,4,5} + {0,1,2,3,4,5}

```

2.2.7 cublasScopy - copy vector into vector

This function copies the vector x into the vector y.

```

//nvcc 005scopy.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"
#define n 6                                // length of x,y
int main(void){
    cudaError_t cudaStat;                   // cudaMalloc status
    cublasStatus_t stat;                    // CUBLAS functions status
    cublasHandle_t handle;                  // CUBLAS context
    int j;                                 // index of elements
    float* x;                             // n-vector on the host
    float* y;                             // n-vector on the host
    x=(float *)malloc (n*sizeof(*x)); // host memory alloc for x

```

```

for(j=0;j<n;j++)
    x[j]=(float)j;                                // x={0,1,2,3,4,5}
printf("x: ");
for(j=0;j<n;j++)
    printf("%2.0f",x[j]);                          // print x
printf("\n");
y=(float *)malloc (n*sizeof(*y)); // host memory alloc for y
// on the device
float* d_x;                                       // d_x - x on the device
float* d_y;                                       // d_y - y on the device
cudaStat=cudaMalloc((void**)&d_x,n*sizeof(*x)); // device
// memory alloc for x
cudaStat=cudaMalloc((void**)&d_y,n*sizeof(*y)); // device
// memory alloc for y
stat = cublasCreate(&handle); // initialize CUBLAS context
stat = cublasSetVector(n,sizeof(*x),x,1,d_x,1); //cp x->d_x
// copy the vector d_x into d_y: d_x -> d_y

stat=cublasScopy(handle,n,d_x,1,d_y,1);

stat=cublasGetVector(n,sizeof(float),d_y,1,y,1); //cp d_y->y
printf("y after copy:\n");
for(j=0;j<n;j++)
    printf("%2.0f",y[j]);                          // print y
printf("\n");
cudaFree(d_x);                                     // free device memory
cudaFree(d_y);                                     // free device memory
cublasDestroy(handle);                             // destroy CUBLAS context
free(x);                                           // free host memory
free(y);                                           // free host memory
return EXIT_SUCCESS;
}
// x:  0, 1, 2, 3, 4, 5,

// y after Scopy:                                // {0,1,2,3,4,5} -> {0,1,2,3,4,5}
// 0, 1, 2, 3, 4, 5,

```

2.2.8 cublasScopy - unified memory version

```

//nvcc 005scopy.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define n 6                                     // length of x,y
int main(void){
    cublasHandle_t handle;                     // CUBLAS context
    int j;                                     // index of elements
    float* x;                                  // n-vector
    float* y;                                  // n-vector
    cudaMallocManaged((void**)&x,n*sizeof(float)); //u.mem for x
    for(j=0;j<n;j++)

```

```

        x[j]=(float)j;                                // x={0,1,2,3,4,5}
printf("x: ");
for(j=0;j<n;j++)
    printf("%.2f",x[j]);                                // print x
printf("\n");
cudaMallocManaged((void**)&y,n*sizeof(float)); //u.mem for y
cublasCreate(&handle);                                // initialize CUBLAS context
// copy the vector x into y:  x -> y

cublasScopy(handle,n,x,1,y,1);

cudaDeviceSynchronize();
printf("y after copy:\n");
for(j=0;j<n;j++)
    printf("%.2f",y[j]);                                // print y
printf("\n");
cudaFree(x);                                           // free memory
cudaFree(y);                                           // free memory
cublasDestroy(handle);                                // destroy CUBLAS context
return 0;
}
// x:  0, 1, 2, 3, 4, 5,

// y after Scopy:                                // {0,1,2,3,4,5} -> {0,1,2,3,4,5}
// 0, 1, 2, 3, 4, 5,

```

2.2.9 cublasSdot - dot product

This function computes the dot product of vectors x and y

$$x.y = x_0y_0 + \dots + x_{n-1}y_{n-1},$$

for real vectors x, y and

$$x.y = x_0\bar{y}_0 + \dots + x_{n-1}\bar{y}_{n-1},$$

for complex x, y .

```
//nvcc 006sdot.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"
#define n 6 // length of x,y
int main(void){
    cudaError_t cudaStat; // cudaMalloc status
    cublasStatus_t stat; // CUBLAS functions status
    cublasHandle_t handle;
    int j; // index of elements
    float* x; // n-vector on the host
```

```

float* y; // n-vector on the host
x=(float *)malloc (n*sizeof(*x)); // host memory alloc for x
for(j=0;j<n;j++)
    x[j]=(float)j; // x={0,1,2,3,4,5}
y=(float *)malloc (n*sizeof(*y)); // host memory alloc for y
for(j=0;j<n;j++)
    y[j]=(float)j; // y={0,1,2,3,4,5}
printf("x,y:\n");
for(j=0;j<n;j++)
    printf("%2.0f",x[j]); // print x,y
printf("\n");
// on the device
float* d_x; // d_x - x on the device
float* d_y; // d_y - y on the device
cudaStat=cudaMalloc((void**)&d_x,n*sizeof(*x)); //device
// memory alloc for x
cudaStat=cudaMalloc((void**)&d_y,n*sizeof(*y)); //device
// memory alloc for y
stat = cublasCreate(&handle); // initialize CUBLAS context
stat = cublasSetVector(n,sizeof(*x),x,1,d_x,1); // cp x->d_x
stat = cublasSetVector(n,sizeof(*y),y,1,d_y,1); // cp y->d_y
float result;
// dot product of two vectors d_x,d_y:
// d_x[0]*d_y[0]+...+d_x[n-1]*d_y[n-1]

stat=cublasSdot(handle,n,d_x,1,d_y,1,&result);

printf("dot product x.y:\n");
printf("%7.0f\n",result); // print the result
cudaFree(d_x); // free device memory
cudaFree(d_y); // free device memory
cublasDestroy(handle); // destroy CUBLAS context
free(x); // free host memory
free(y); // free host memory
return EXIT_SUCCESS;
}
// x,y:
// 0, 1, 2, 3, 4, 5,

// dot product x.y: // x.y=
// 55 // 1*1+2*2+3*3+4*4+5*5

```

2.2.10 cublasSdot - unified memory version

```

//nvcc 006sdot.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define n 6 // length of x,y
int main(void){
    cublasHandle_t handle;

```

```

int j;                                // index of elements
float* x;                             // n-vector
float* y;                             // n-vector
cudaMallocManaged(&x,n*sizeof(float)); // unified mem.for x
for(j=0;j<n;j++)
    x[j]=(float)j;                    // x={0,1,2,3,4,5}
cudaMallocManaged(&y,n*sizeof(float)); // unified mem.for y
for(j=0;j<n;j++)
    y[j]=(float)j;                    // y={0,1,2,3,4,5}
printf("x,y:\n");
for(j=0;j<n;j++)
    printf("%2.0f",x[j]);              // print x,y
printf("\n");
cublasCreate(&handle);                 // initialize CUBLAS context
float result;
// dot prod. of two vectors x,y: x[0]*y[0]+...+x[n-1]*y[n-1]

cublasSdot(handle,n,x,1,y,1,&result);

cudaDeviceSynchronize();
printf("dot product x.y:\n");
printf("%7.0f\n",result);              // print the result
cudaFree(x);                          // free memory
cudaFree(y);                          // free memory
cublasDestroy(handle);                 // destroy CUBLAS context
return EXIT_SUCCESS;
}
// x,y:
// 0, 1, 2, 3, 4, 5,

// dot product x.y:                    // x.y=
//      55                             // 1*1+2*2+3*3+4*4+5*5

```

2.2.11 cublasSnrm2 - Euclidean norm

This function computes the Euclidean norm of the vector x

$$\|x\| = \sqrt{|x_0|^2 + \dots + |x_{n-1}|^2},$$

where $x = \{x_0, \dots, x_{n-1}\}$.

```

//nvcc 007snrm2.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"
#define n 6                                // length of x
int main(void){
    cudaError_t cudaStat;                  // cudaMalloc status
    cublasStatus_t stat;                  // CUBLAS functions status
    cublasHandle_t handle;                // CUBLAS context

```

```

int j;                                // index of elements
float* x;                             // n-vector on the host
x=(float *)malloc (n*sizeof(*x)); // host memory alloc for x
for(j=0;j<n;j++)
    x[j]=(float)j;                    // x={0,1,2,3,4,5}
printf("x: ");
for(j=0;j<n;j++)
    printf("%2.0f",x[j]);             // print x
printf("\n");
// on the device
float* d_x;                           // d_x - x on the device
cudaStat=cudaMalloc((void**)&d_x,n*sizeof(*x)); // device
// memory alloc for x
stat = cublasCreate(&handle); // initialize CUBLAS context
stat = cublasSetVector(n,sizeof(*x),x,1,d_x,1); // cp x->d_x
float result;
// Euclidean norm of the vector d_x:
// \sqrt{d_x[0]^2+...+d_x[n-1]^2}

stat=cublasSnrm2(handle,n,d_x,1,&result);

printf("Euclidean norm of  x: ");
printf("%7.3f\n",result);             // print the result
cudaFree(d_x);                       // free device memory
cublasDestroy(handle);               // destroy CUBLAS context
free(x);                             // free host memory
return EXIT_SUCCESS;
}
// x:  0, 1, 2, 3, 4, 5,
// ||x||=
//Euclidean norm of x: 7.416 //\sqrt{0^2+1^2+2^2+3^2+4^2+5^2}

```

2.2.12 cublasSnrm2 - unified memory version

```

//nvcc 007snrm2.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define n 6                                // length of x
int main(void){
    cublasHandle_t handle;                // CUBLAS context
    int j;                                // index of elements
    float* x;                             // n-vector
    cudaMallocManaged((void**)&x,n*sizeof(float)); // unified
    for(j=0;j<n;j++)                      // memory for x
        x[j]=(float)j;                    // x={0,1,2,3,4,5}
    printf("x: ");
    for(j=0;j<n;j++)
        printf("%2.0f",x[j]);             // print x
    printf("\n");
    cublasCreate(&handle);                 // initialize CUBLAS context

```

```

float result;
// Euclidean norm of the vector x: \sqrt{x[0]^2+...+x[n-1]^2}

cublasSnrm2(handle,n,x,1,&result);

cudaDeviceSynchronize();
printf("Euclidean norm of  x: ");
printf("%7.3f\n",result);           // print the result
cudaFree(x);                       // free memory
cublasDestroy(handle);              // destroy CUBLAS context
return 0;
}
// x:  0, 1, 2, 3, 4, 5,
// ||x||=
//Euclidean norm of x: 7.416 //\sqrt{0^2+1^2+2^2+3^2+4^2+5^2}

```

2.2.13 cublasSrot - apply the Givens rotation

This function multiplies 2×2 Givens rotation matrix $\begin{pmatrix} c & s \\ -s & c \end{pmatrix}$ with the $2 \times n$ matrix $\begin{pmatrix} x_0 & \dots & x_{n-1} \\ y_0 & \dots & y_{n-1} \end{pmatrix}$.

```

// nvcc 008srot.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"
#define n 6 // length of x,y
int main(void){
    cudaError_t cudaStat; // cudaMalloc status
    cublasStatus_t stat; // CUBLAS functions status
    cublasHandle_t handle; // CUBLAS context
    int j; // index of elements
    float* x; // n-vector on the host
    float* y; // n-vector on the host
    x=(float *)malloc (n*sizeof(*x)); // host memory alloc for x
    for(j=0;j<n;j++)
        x[j]=(float)j; // x={0,1,2,3,4,5}
    y=(float *)malloc (n*sizeof(*y)); // host memory alloc for y
    for(j=0;j<n;j++)
        y[j]=(float)j*j; // y={0,1,4,9,16,25}
    printf("x: ");
    for(j=0;j<n;j++)
        printf("%7.0f,",x[j]); // print x
    printf("\n");
    printf("y: ");
    for(j=0;j<n;j++)
        printf("%7.0f,",y[j]); // print y
    printf("\n");
}

```



```

// on the device
float* d_x; // d_x - x on the device
float* d_y; // d_y - y on the device
cudaStat=cudaMalloc((void**)&d_x,n*sizeof(*x)); //device
// memory alloc for x
cudaStat=cudaMalloc((void**)&d_y,n*sizeof(*y)); //device
// memory alloc for y
stat = cublasCreate(&handle); // initialize CUBLAS context
stat = cublasSetVector(n,sizeof(*x),x,1,d_x,1); //cp x->d_x
stat = cublasSetVector(n,sizeof(*y),y,1,d_y,1); //cp y->d_y
float c=0.5;
float s=0.8669254; // s=sqrt(3.0)/2.0
// Givens rotation
//
// multiplies 2x2 matrix [ c s ] with 2xn matrix [ row(x) ]
// [-s c ] [ row(y) ]
//
// [1/2 sqrt(3)/2] [0,1,2,3, 4, 5]
// [-sqrt(3)/2 1/2 ] [0,1,4,9,16,25]

stat=cublasSrot(handle,n,d_x,1,d_y,1,&c,&s);

stat=cublasGetVector(n,sizeof(float),d_x,1,x,1); //cp d_x->x

printf("x after Srot:\n"); // print x after Srot
for(j=0;j<n;j++)
    printf("%7.3f",x[j]);
printf("\n");
stat=cublasGetVector(n,sizeof(float),d_y,1,y,1); //cp d_y->y
printf("y after Srot:\n"); // print y after Srot
for(j=0;j<n;j++)
    printf("%7.3f",y[j]);
printf("\n");
cudaFree(d_x); // free device memory
cudaFree(d_y); // free device memory
cublasDestroy(handle); // destroy CUBLAS context
free(x); // free host memory
free(y); // free host memory
return EXIT_SUCCESS;
}
// x: 0, 1, 2, 3, 4, 5,
// y: 0, 1, 4, 9, 16, 25,

// x after Srot:
// 0.000, 1.367, 4.468, 9.302, 15.871, 24.173,
// y after Srot:
// 0.000, -0.367, 0.266, 1.899, 4.532, 8.165,
// // [x] [ 0.5 0.867] [0 1 2 3 4 5]
// // [ ]= [ ]*[ ]
// // [y] [-0.867 0.5 ] [0 1 4 9 16 25]

```

2.2.14 cublasSrot - unified memory version

```

// nvcc 008srot.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define n 6 // length of x,y
int main(void){
    cublasHandle_t handle; // CUBLAS context
    int j; // index of elements
    float* x; // n-vector
    float* y; // n-vector
    cudaMallocManaged(&x,n*sizeof(float)); // unified mem.for x
    for(j=0;j<n;j++){
        x[j]=(float)j; // x={0,1,2,3,4,5}
    }
    cudaMallocManaged(&y,n*sizeof(float)); // unified mem.for y
    for(j=0;j<n;j++){
        y[j]=(float)j*j; // y={0,1,4,9,16,25}
    }
    printf("x: ");
    for(j=0;j<n;j++){
        printf("%7.0f",x[j]); // print x
    }
    printf("\n");
    printf("y: ");
    for(j=0;j<n;j++){
        printf("%7.0f",y[j]); // print y
    }
    printf("\n");
    cublasCreate(&handle); // initialize CUBLAS context
    float c=0.5;
    float s=0.8669254; // s=sqrt(3.0)/2.0
    // Givens rotation
    //      [ c s ]      [ row(x) ]
    // multiplies 2x2 matrix [   ] with 2xn matrix [   ]
    //      [ -s c ]      [ row(y) ]
    //
    //      [1/2      sqrt(3)/2]   [0,1,2,3, 4, 5]
    //      [-sqrt(3)/2  1/2 ]   [0,1,4,9,16,25]

    cublasSrot(handle,n,x,1,y,1,&c,&s);

    cudaDeviceSynchronize();
    printf("x after Srot:\n"); // print x after Srot
    for(j=0;j<n;j++){
        printf("%7.3f",x[j]);
    }
    printf("\n");

    printf("y after Srot:\n"); // print y after Srot
    for(j=0;j<n;j++){
        printf("%7.3f",y[j]);
    }
    printf("\n");
    cudaFree(x); // free memory
    cudaFree(y); // free memory
    cublasDestroy(handle); // destroy CUBLAS context
}

```

```

// x:      0,      1,      2,      3,      4,      5,
// y:      0,      1,      4,      9,     16,     25,

// x after Srot:
// 0.000,  1.367,  4.468,  9.302, 15.871, 24.173,
// y after Srot:
// 0.000, -0.367,  0.266,  1.899,  4.532,  8.165,
//                                     // [x]  [ 0.5   0.867] [0 1 2 3 4 5]
//                                     // [ ]= [          ]*[          ]
//                                     // [y]  [-0.867  0.5 ] [0 1 4 9 16 25]

```

2.2.15 cublasSrotg - construct the Givens rotation matrix

This function constructs the Givens rotation matrix $G = \begin{pmatrix} c & s \\ -s & c \end{pmatrix}$ that zeros out the second entry of 2×1 vector $\begin{pmatrix} a \\ b \end{pmatrix}$ i.e. $\begin{pmatrix} c & s \\ -s & c \end{pmatrix} \begin{pmatrix} a \\ b \end{pmatrix} = \begin{pmatrix} r \\ 0 \end{pmatrix}$, where $c^2 + s^2 = 1$, $r^2 = a^2 + b^2$.

```

// nvcc 009srotg.c -lcublas
// This function is provided for completeness and runs
// exclusively on the host
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"
int main(void){
    cublasStatus_t stat;           // CUBLAS functions status
    cublasHandle_t handle;         // CUBLAS context
    int j;
    float a=1.0;
    float b=1.0;
    printf("a: %7.3f\n",a);        // print a
    printf("b: %7.3f\n",b);        // print b
    stat = cublasCreate(&handle);  // initialize CUBLAS context
    float c;
    float s;
    //                                     [ c  s ]
    // find the Givens rotation matrix G =[  ]
    //                                     [-s c ]
    //          [a] [r]
    // such that G*[ ]=[ ]
    //          [b] [0]
    //
    // c^2+s^2=1,  r=\sqrt{a^2+b^2}, a is replaced by r

    stat=cublasSrotg(handle,&a,&b,&c,&s);

    printf("After Srotg:\n");

```

```

printf("a: %7.5f\n",a);           // print a
printf("c: %7.5f\n",c);           // print c
printf("s: %7.5f\n",s);           // print s
cublasDestroy(handle);             // destroy CUBLAS context
return EXIT_SUCCESS;
}
// a:    1.000
// b:    1.000

// After Srotg:
// a: 1.41421                      // \sqrt{1^2+1^2}
// c: 0.70711                      // cos(pi/4)
// s: 0.70711                      // sin(pi/4)
//                                // [ 0.70711  0.70711] [1] [1.41422]
//                                // [          ]*[ ]=[          ]
//                                // [-0.70711  0.70711] [1] [  0   ]

```

2.2.16 cublasSrotm - apply the modified Givens rotation

This function multiplies the modified Givens 2×2 matrix $\begin{pmatrix} h_{11} & h_{12} \\ h_{21} & h_{22} \end{pmatrix}$ with $2 \times n$ matrix $\begin{pmatrix} x_0 & \dots & x_{n-1} \\ y_0 & \dots & y_{n-1} \end{pmatrix}$.

```

// nvcc 010srotmVec.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"
#define n 6                      // length of x,y
int main(void){
    cudaError_t cudaStat;         // cudaMalloc status
    cublasStatus_t stat;          // CUBLAS functions status
    cublasHandle_t handle;        // CUBLAS context
    int j;                       // index of elements
    float* x;                    // n-vector on the host
    float* y;                    // n-vector on the host
    float* param;
    x=(float *)malloc (n*sizeof(*x)); // host memory alloc for x
    for(j=0;j<n;j++)
        x[j]=(float)j;           // x={0,1,2,3,4,5}
    printf("x:\n");
    for(j=0;j<n;j++)
        printf("%3.0f,",x[j]);   // print x
    printf("\n");
    y=(float *)malloc (n*sizeof(*y)); // host memory alloc for y
    for(j=0;j<n;j++)
        y[j]=(float)j*j;         // y={0,1,4,9,16,25}
    printf("y:\n");

```

```

for(j=0;j<n;j++)
    printf("%3.0f",y[j]);                                // print y
printf("\n");
param=(float *)malloc (5*sizeof(*param));
param[0]=1.0f;                                           // flag
param[1]=0.5f;                                           // param[1],...,param[4]
param[2]=1.0f;                                           // -entries of the Givens matrix
param[3]=-1.0f;                                          // h11=param[1]  h12=param[2]
param[4]=0.5f;                                          // h21=param[3]  h22=param[4]
// on the device
float* d_x;                                             // d_x - x on the device
float* d_y;                                             // d_y - y on the device
cudaStat=cudaMalloc((void**)&d_x,n*sizeof(*x));        //device
// memory alloc for x
cudaStat=cudaMalloc((void**)&d_y,n*sizeof(*y));        //device
// memory alloc for y
stat =cublasCreate(&handle); // initialize CUBLAS context
stat =cublasSetVector(n,sizeof(*x),x,1,d_x,1); //copy x->d_x
stat =cublasSetVector(n,sizeof(*y),y,1,d_y,1); //copy y->d_y
//
// multiply the 2x2 modified Givens matrix      H=[ 0.5  1.0 ]
// by the 2xn matrix with two rows d_x and d_y  [-1.0 0.5 ]

stat=cublasSrotm(handle,n,d_x,1,d_y,1,param);

stat=cublasGetVector(n,sizeof(float),d_x,1,x,1); //cp d_x->x
printf("x after Srotm x:\n");                    // print x after Srotm
for(j=0;j<n;j++)
    printf("%7.3f",x[j]);
printf("\n");
stat=cublasGetVector(n,sizeof(float),d_y,1,y,1); //cp d_y->y
printf("y after Srotm y:\n");                    // print y after Srotm
for(j=0;j<n;j++)
    printf("%7.3f",y[j]);
printf("\n");
cudaFree(d_x);                                       // free device memory
cudaFree(d_y);                                       // free device memory
cublasDestroy(handle);                             // destroy CUBLAS context
free(x);                                           // free host memory
free(y);                                           // free host memory
free(param);                                       // free host memory
return EXIT_SUCCESS;
}
// x:
// 0,  1,  2,  3,  4,  5,
// y:
// 0,  1,  4,  9, 16, 25,

// x after Srotm:
// 0.000,  1.500,  5.000, 10.500, 18.000, 27.500,
// y after Srotm:
// 0.000, -0.500,  0.000,  1.500,  4.000,  7.500,

```

```

//                                // [x]  [ 0.5   1 ] [0 1 2 3  4  5]
//                                // [ ]= [          ]*[          ]
//                                // [y]  [ -1    0.5] [0 1 4 9 16 25]

```

2.2.17 cublasSrotm - unified memory version

```

// nvcc 010srotmVec.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define n 6                                // length of x,y
int main(void){
    cublasHandle_t handle;                // CUBLAS context
    int j;                                // index of elements
    float* x;                             // n-vector
    float* y;                             // n-vector
    float* param;
    cudaMallocManaged((void**)&x,n*sizeof(float)); // unified
    for(j=0;j<n;j++)                      // memory for x
        x[j]=(float)j;                   // x={0,1,2,3,4,5}
    printf("x:\n");
    for(j=0;j<n;j++)
        printf("%3.0f,",x[j]);           // print x
    printf("\n");
    cudaMallocManaged((void**)&y,n*sizeof(float)); // unified
    for(j=0;j<n;j++)                      // memory for y
        y[j]=(float)j*j;                 // y={0,1,4,9,16,25}
    printf("y:\n");
    for(j=0;j<n;j++)
        printf("%3.0f,",y[j]);           // print y
    printf("\n");
    param=(float *)malloc (5*sizeof(*param));
    param[0]=1.0f;                        // flag
    param[1]=0.5f;                        // param[1],...,param[4]
    param[2]=1.0f;                        // -entries of the Givens matrix
    param[3]=-1.0f;                       // h11=param[1] h12=param[2]
    param[4]=0.5f;                       // h21=param[3] h22=param[4]

    cublasCreate(&handle);                // initialize CUBLAS context
    //                                [0.5  1.0 ]
    // multiply the 2x2 modified Givens matrix      H=[
    // by the 2xn matrix with two rows x and y      [-1.0 0.5 ]

    cublasSrotm(handle,n,x,1,y,1,param);

    cudaDeviceSynchronize();
    printf("x after Srotm x:\n");          // print x after Srotm
    for(j=0;j<n;j++)
        printf("%7.3f,",x[j]);
    printf("\n");

```

```

printf("y after Srotm y:\n");          // print y after Srotm
for(j=0;j<n;j++)
    printf("%7.3f",y[j]);
printf("\n");
cudaFree(x);                          // free memory
cudaFree(y);                          // free memory
cublasDestroy(handle);                 // destroy CUBLAS context
free(param);                          // free memory
return 0;
}
// x:
// 0, 1, 2, 3, 4, 5,
// y:
// 0, 1, 4, 9, 16, 25,

// x after Srotm:
// 0.000, 1.500, 5.000, 10.500, 18.000, 27.500,
// y after Srotm:
// 0.000, -0.500, 0.000, 1.500, 4.000, 7.500,
//                                     // [x] [ 0.5  1 ] [0 1 2 3 4 5]
//                                     // [ ] = [      ]*[      ]
//                                     // [y] [ -1   0.5] [0 1 4 9 16 25]

```

2.2.18 cublasSrotmg - construct the modified Givens rotation matrix

This function constructs the modified Givens transformation $\begin{pmatrix} h_{11} & h_{12} \\ h_{21} & h_{22} \end{pmatrix}$ that zeros out the second entry of the vector $\begin{pmatrix} \sqrt{d_1} * x_1 \\ \sqrt{d_2} * y_1 \end{pmatrix}$.

```

// nvcc 011srotmg.c -lcublas
// this function is provided for completeness
// and runs exclusively on the Host
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"
int main(void){
    cublasStatus_t stat;          // CUBLAS functions status
    cublasHandle_t handle;        // CUBLAS context
    float d1=5.0f;                // d1=5.0
    float d2=5.0f;                // d2=5.0
    float param[5];              // [param[1] param[2]] [h11 h12]
                                // [      ] = [      ]
                                // [param[3] param[4]] [h21 h22]
    param[0]=1.0f;                // param[0] is a flag
    // if param[0]=1.0, then h12=1=param[2], h21=-1=param[3]
    printf("d1: %7.3f\n",d1);     // print d1
    printf("d2: %7.3f\n",d2);     // print d2
    stat = cublasCreate(&handle); // initialize CUBLAS context
}

```

```

float x1=1.0f;           // x1=1
float y1=2.0f;           // y1=2
printf("x1: %7.3f\n",x1); // print x1
printf("y1: %7.3f\n",y1); // print y1
//find modified Givens rotation matrix H={{h11,h12},{h21,h22}}
//such that the second entry of H*{\sqrt{d1}*x1,\sqrt{d2}*y1}^T
//is zero

    stat=cublasSrotmg(handle,&d1,&d2,&x1,&y1,param);

printf("After srotmg:\n");
printf("param[0]: %4.2f\n",param[0]);
printf("h11: %7.5f\n",param[1]);
printf("h22: %7.5f\n",param[4]);
//check if the second entry of H*{\sqrt{d1}*x1,\sqrt{d2}*y1}^T
//is zero; the values of d1,d2,x1 are overwritten so we use
//their initial values
printf("%7.5f\n",(-1.0)*sqrt(5.0)*1.0+
                                param[4]*sqrt(5.0)*2.0);
cublasDestroy(handle);          // destroy CUBLAS context
return EXIT_SUCCESS;
}
// d1:    5.000           // [d1] [5]      [x1] [1]          [0.5  1 ]
// d2:    5.000           // [ ]=[ ],      [ ]=[ ],      H=[          ]
// x1:    1.000           // [d2] [5]      [x2] [2]          [-1  0.5]
// y1:    2.000

// After srotmg:
// param[0]: 1.00
// h11: 0.50000
// h22: 0.50000

//   [sqrt(d1)*x1] [0.5  1 ] [sqrt(5)*1] [5.59]
// H*[              ]=[          ]*[          ]=[          ]
//   [sqrt(d2)*y1] [-1  0.5] [sqrt(5)*2] [  0 ]

// 0.00000 <= the second entry of
// H*{\sqrt{d1}*x1,\sqrt{d2}*y1}^T

```

2.2.19 cublasSscal - scale the vector

This function scales the vector x by the scalar α .

$$x = \alpha x.$$

```

// nvcc 012sscal.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"

```



```

#define n 6                                // length of x
int main(void){
    cudaError_t cudaStat;                  // cudaMalloc status
    cublasStatus_t stat;                   // CUBLAS functions status
    cublasHandle_t handle;                 // CUBLAS context
    int j;                                // index of elements
    float* x;                              // n-vector on the host
    x=(float *)malloc (n*sizeof(*x)); // host memory alloc for x
    for(j=0;j<n;j++)
        x[j]=(float)j;                    // x={0,1,2,3,4,5}
    printf("x:\n");
    for(j=0;j<n;j++)
        printf("%2.0f",x[j]);              // print x
    printf("\n");
    // on the device
    float* d_x;                            // d_x - x on the device
    cudaStat=cudaMalloc((void**)&d_x,n*sizeof(*x)); //device
                                                // memory alloc for x
    stat = cublasCreate(&handle); // initialize CUBLAS context
    stat = cublasSetVector(n,sizeof(*x),x,1,d_x,1); // cp x->d_x
    float al=2.0;                          // al=2
    // scale the vector d_x by the scalar al: d_x = al*d_x

    stat=cublasSscal(handle,n,&al,d_x,1);

    stat=cublasGetVector(n,sizeof(float),d_x,1,x,1); //cp d_x->x
    printf("x after Sscal:\n");                // print x after Sscal:
    for(j=0;j<n;j++)
        printf("%2.0f",x[j]);                  // x={0,2,4,6,8,10}
    printf("\n");
    cudaFree(d_x);                            // free device memory
    cublasDestroy(handle);                     // destroy CUBLAS context
    free(x);                                  // free host memory
    return EXIT_SUCCESS;
}
// x:
// 0, 1, 2, 3, 4, 5,

// x after Sscal:
// 0, 2, 4, 6, 8,10,                          // 2*{0,1,2,3,4,5}

```

2.2.20 cublasSscal - unified memory version

```

// nvcc 012sscal.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define n 6                                // length of x
int main(void){
    cublasHandle_t handle;                  // CUBLAS context
    int j;                                // index of elements

```

```

float* x; // n-vector
cudaMallocManaged(&x,n*sizeof(float)); // unified mem.for x
for(j=0;j<n;j++)
    x[j]=(float)j; // x={0,1,2,3,4,5}
printf("x:\n");
for(j=0;j<n;j++)
    printf("%2.0f",x[j]); // print x
printf("\n");
cublasCreate(&handle); // initialize CUBLAS context
float al=2.0; // al=2
// scale the vector x by the scalar al: x = al*x

cublasSscal(handle,n,&al,x,1);

cudaDeviceSynchronize();
printf("x after Sscal:\n"); // print x after Sscal:
for(j=0;j<n;j++)
    printf("%2.0f",x[j]); // x={0,2,4,6,8,10}
printf("\n");
cudaFree(x); // free memory
cublasDestroy(handle); // destroy CUBLAS context
return EXIT_SUCCESS;
}
// x:
// 0, 1, 2, 3, 4, 5,

// x after Sscal:
// 0, 2, 4, 6, 8,10, // 2*{0,1,2,3,4,5}

```

2.2.21 cublasSswap - swap two vectors

This function interchanges the elements of vector x and y

$$x \leftarrow y, \quad y \leftarrow x.$$

```

// nvcc 013sswap.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"
#define n 6 // length of x,y
int main(void){
    cudaError_t cudaStat; // cudaMalloc status
    cublasStatus_t stat; // CUBLAS functions status
    cublasHandle_t handle; // CUBLAS context
    int j; // index of elements
    float* x; // n-vector on the host
    float* y; // n-vector on the host
    x=(float *)malloc (n*sizeof(*x)); // host memory alloc for x
    for(j=0;j<n;j++)
        x[j]=(float)j; // x={0,1,2,3,4,5}

```

```

printf("x:\n");
for(j=0;j<n;j++)
    printf("%2.0f",x[j]); // print x
printf("\n");
y=(float *)malloc (n*sizeof(*y)); // host memory alloc for y
for(j=0;j<n;j++)
    y[j]=(float)2*j; // y={0,2,4,6,8,10}
printf("y:\n");
for(j=0;j<n;j++)
    printf("%2.0f",y[j]); // print y
printf("\n");
// on the device
float* d_x; // d_x - x on the device
float* d_y; // d_y - y on the device
cudaStat=cudaMalloc((void**)&d_x,n*sizeof(*x)); // device
// memory alloc for x
cudaStat=cudaMalloc((void**)&d_y,n*sizeof(*y)); // device
//memory alloc for y
stat = cublasCreate(&handle); // initialize CUBLAS context
stat = cublasSetVector(n,sizeof(*x),x,1,d_x,1); // cp x->d_x
stat = cublasSetVector(n,sizeof(*y),y,1,d_y,1); // cp y->d_y
// swap the vectors d_x,d_y: d_x<--d_y, d_y<--d_x

stat=cublasSswap(handle,n,d_x,1,d_y,1);

stat=cublasGetVector(n,sizeof(float),d_y,1,y,1); //cp d_y->y
stat=cublasGetVector(n,sizeof(float),d_x,1,x,1); //cp d_x->x
printf("x after Sswap:\n"); // print x after Sswap:
for(j=0;j<n;j++)
    printf("%2.0f",x[j]); // x={0,2,4,6,8,10}
printf("\n");
printf("y after Sswap:\n"); // print y after Sswap:
for(j=0;j<n;j++)
    printf("%2.0f",y[j]); // y={0,1,2,3,4,5}
printf("\n");
cudaFree(d_x); // free device memory
cudaFree(d_y); // free device memory
cublasDestroy(handle); // destroy CUBLAS context
free(x); // free host memory
free(y); // free host memory
return EXIT_SUCCESS;
}
// x:
// 0, 1, 2, 3, 4, 5,
// y:
// 0, 2, 4, 6, 8,10,

// x after Sswap:
// 0, 2, 4, 6, 8,10, // x <- y
// y after Sswap:
// 0, 1, 2, 3, 4, 5, // y <- x

```

2.2.22 cublasSswap - unified memory version

```

// nvcc 013sswap.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define n 6 // length of x,y
int main(void){
    cublasHandle_t handle; // CUBLAS context
    int j; // index of elements
    float* x; // n-vector
    float* y; // n-vector
    cudaMallocManaged(&x,n*sizeof(float)); // unified mem.for x
    for(j=0;j<n;j++){
        x[j]=(float)j; // x={0,1,2,3,4,5}
    }
    printf("x:\n");
    for(j=0;j<n;j++){
        printf("%2.0f",x[j]); // print x
    }
    printf("\n");
    cudaMallocManaged(&y,n*sizeof(float)); // unified mem.for y
    for(j=0;j<n;j++){
        y[j]=(float)2*j; // y={0,2,4,6,8,10}
    }
    printf("y:\n");
    for(j=0;j<n;j++){
        printf("%2.0f",y[j]); // print y
    }
    printf("\n");
    cublasCreate(&handle); // initialize CUBLAS context
    // swap x and y

    cublasSswap(handle,n,x,1,y,1);

    cudaDeviceSynchronize();
    printf("x after Sswap:\n"); // print x after Sswap:
    for(j=0;j<n;j++){
        printf("%2.0f",x[j]); // x={0,2,4,6,8,10}
    }
    printf("\n");
    printf("y after Sswap:\n"); // print y after Sswap:
    for(j=0;j<n;j++){
        printf("%2.0f",y[j]); // y={0,1,2,3,4,5}
    }
    printf("\n");
    cudaFree(x); // free memory
    cudaFree(y); // free memory
    cublasDestroy(handle); // destroy CUBLAS context
    return EXIT_SUCCESS;
}

// x:
// 0, 1, 2, 3, 4, 5,
// y:
// 0, 2, 4, 6, 8,10,
// x after Sswap:
// 0, 2, 4, 6, 8,10, // x <- y
// y after Sswap:
// 0, 1, 2, 3, 4, 5, // y <- x

```

2.3 CUBLAS Level-2. Matrix-vector operations

2.3.1 cublasSgbmv – banded matrix-vector multiplication

This function performs the banded matrix-vector multiplication

$$y = \alpha \operatorname{op}(A)x + \beta y,$$

where A is a banded matrix with ku superdiagonals and kl subdiagonals, x, y are vectors, α, β are scalars and $\operatorname{op}(A)$ can be equal to A (CUBLAS_OP_N case), A^T (transposition) in CUBLAS_OP_T case or A^H (conjugate transposition) in CUBLAS_OP_C case. The highest superdiagonal is stored in row 0, starting from position ku , the next superdiagonal is stored in row 1 starting from position $ku - 1, \dots$. The main diagonal is stored in row ku , starting from position 0, the first subdiagonal is stored in row $ku+1$, starting from position 0, the next subdiagonal is stored in row $ku + 2$ from position 0, \dots .

```
// nvcc 013sgbm.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))
#define m 5 // number of rows
#define n 6 // number of columns
#define ku 2 // number of superdiagonals
#define kl 1 // number of subdiagonals
int main(void){
    cudaError_t cudaStat; // cudaMalloc status
    cublasStatus_t stat; // CUBLAS functions status
    cublasHandle_t handle; // CUBLAS context
    int i,j; // row and column index
    // declaration and allocation of a,x,y on the host
    float* a; //mxn matrix on the host // a:
    float* x; //n-vector on the host // 20 15 11
    float* y; //m-vector on the host // 25 21 16 12
    a=(float*)malloc(m*n*sizeof(float)); // 26 22 17 13
    // host memory alloc for a // 27 23 18 14
    x=(float*)malloc(n*sizeof(float)); // 28 24 19
    // host memory alloc for x
    y=(float*)malloc(m*sizeof(float)); //host memory alloc for y
    int ind=11;
    // highest superdiagonal 11,12,13,14 in first row,
    // starting from i=ku
    for(i=ku;i<n;i++) a[IDX2C(0,i,m)]=(float)ind++;
    // next superdiagonal 15,16,17,18,19 in next row,
    // starting from i=ku-1
    for(i=ku-1;i<n;i++) a[IDX2C(1,i,m)]=(float)ind++;
```

```

// main diagonal 20,21,22,23,24 in row ku, starting from i=0
for(i=0;i<n-1;i++) a[IDX2C(ku,i,m)]=(float)ind++;
// subdiagonal 25,26,27,28 in ku+1 row, starting from i=0
for(i=0;i<n-2;i++) a[IDX2C(ku+1,i,m)]=(float)ind++;
for(i=0;i<n;i++) x[i]=1.0f;           // x={1,1,1,1,1,1}^T
for(i=0;i<m;i++) y[i]=0.0f;          // y={0,0,0,0,0}^T
// on the device
float* d_a;                          // d_a - a on the device
float* d_x;                          // d_x - x on the device
float* d_y;                          // d_y - y on the device
cudaStat=cudaMalloc((void**)&d_a,m*n*sizeof(*a)); // device
// memory alloc for a
cudaStat=cudaMalloc((void**)&d_x,n*sizeof(*x));   // device
// memory alloc for x
cudaStat=cudaMalloc((void**)&d_y,m*sizeof(*y));   // device
// memory alloc for y

stat = cublasCreate(&handle);
stat = cublasSetMatrix(m,n,sizeof(*a),a,m,d_a,m); //cp a->d_a
stat = cublasSetVector(n,sizeof(*x),x,1,d_x,1);   //cp x->d_x
stat = cublasSetVector(m,sizeof(*y),y,1,d_y,1);   //cp y->d_y
float al=1.0f;                                     // al=1
float bet=1.0f;                                    // bet=1

// banded matrix-vector multiplication:
// d_y = al*d_a*d_x + bet*d_y;          d_a - mxn banded matrix;
// d_x - n-vector, d_y - m-vector; al,bet - scalars

stat=cublasSgbmv(handle,CUBLAS_OP_N,m,n,kl,ku,&al,d_a,m,d_x,1,
                &bet,d_y,1);

stat=cublasGetVector(m,sizeof(*y),d_y,1,y,1); // copy d_y->y
printf("y after Sgbmv:\n");                  // print y after Sgbmv
for(j=0;j<m;j++){
    printf("%7.0f",y[j]);
    printf("\n");
}
cudaFree(d_a);                               // free device memory
cudaFree(d_x);                               // free device memory
cudaFree(d_y);                               // free device memory
cublasDestroy(handle);                       // destroy CUBLAS context
free(a);                                     // free host memory
free(x);                                     // free host memory
free(y);                                     // free host memory
return EXIT_SUCCESS;
}

// y after Sgbmv:
//      46
//      74
//      78
//      82
//      71
//      [ 20 15 11
//      [ 25 21 16 12
//      [      26 22 17 13
//      [      27 23 18 14
//      [      28 24 19
//      [1]
//      [1]
//      [1]
//      [1]
//      [1]
//      [1]

```

2.3.2 cublasSgbmv – unified memory version

```

// nvcc 013sgbmv.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))
#define m 5 // number of rows
#define n 6 // number of columns
#define ku 2 // number of superdiagonals
#define kl 1 // number of subdiagonals
int main(void){
    cublasHandle_t handle; // CUBLAS context
    int i,j; // row and column index
    // declaration and allocation of a,x,y in unified memory
    float* a; //mxn matrix // a:
    float* x; //n-vector // 20 15 11
    float* y; //m-vector // 25 21 16 12
    // 26 22 17 13
    // 27 23 18 14
    // 28 24 19
    cudaMallocManaged(&a,m*n*sizeof(float)); //unified mem.for a
    cudaMallocManaged(&x,n*sizeof(float)); //unified mem.for x
    cudaMallocManaged(&y,m*sizeof(float)); //unified mem.for y
    int ind=11;
    // highest superdiagonal 11,12,13,14 in first row,
    // starting from i=ku
    for(i=ku;i<n;i++) a[IDX2C(0,i,m)]=(float)ind++;
    // next superdiagonal 15,16,17,18,19 in next row,
    // starting from i=ku-1
    for(i=ku-1;i<n;i++) a[IDX2C(1,i,m)]=(float)ind++;
    // main diagonal 20,21,22,23,24 in row ku, starting from i=0
    for(i=0;i<n-1;i++) a[IDX2C(ku,i,m)]=(float)ind++;
    // subdiagonal 25,26,27,28 in ku+1 row, starting from i=0
    for(i=0;i<n-2;i++) a[IDX2C(ku+1,i,m)]=(float)ind++;
    for(i=0;i<n;i++) x[i]=1.0f; // x={1,1,1,1,1,1}^T
    for(i=0;i<m;i++) y[i]=0.0f; // y={0,0,0,0,0}^T
    cublasCreate(&handle);
    float al=1.0f; // al=1
    float bet=1.0f; // bet=1
    // banded matrix-vector multiplication:
    // y = al*a*x + bet*y; a - mxn banded matrix;
    // x - n-vector, y - m-vector; al,bet - scalars

    cublasSgbmv(handle,CUBLAS_OP_N,m,n,kl,ku,&al,a,m,x,1,&bet,y,1);

    cudaDeviceSynchronize();
    printf("y after Sgbmv:\n"); // print y after Sgbmv
    for(j=0;j<m;j++){
        printf("%7.0f",y[j]);
        printf("\n");
    }
    cudaFree(a); // free memory

```

```

    cudaFree(x);                                // free memory
    cudaFree(y);                                // free memory
    cublasDestroy(handle);                       // destroy CUBLAS context
    return EXIT_SUCCESS;
}
// y after Sgbmv:                               //
//      46                                     // [ 20 15 11          ] [1]
//      74                                     // [ 25 21 16 12          ] [1]
//      78                                     // [      26 22 17 13      ]*[ ]
//      82                                     // [          27 23 18 14 ] [1]
//      71                                     // [          28 24 19 ] [1]
//                                           // [1]

```

2.3.3 cublasSgemv – matrix-vector multiplication

This function performs matrix-vector multiplication

$$y = \alpha \operatorname{op}(A)x + \beta y,$$

where A is a matrix, x, y are vectors, α, β are scalars and $\operatorname{op}(A)$ can be equal to A (CUBLAS_OP_N case), A^T (transposition) in CUBLAS_OP_T case or A^H (conjugate transposition) in CUBLAS_OP_C case. A is stored column by column.

```

// nvcc 014sgemv.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))
#define m 6 // number of rows of a
#define n 5 // number of columns of a
int main(void){
    cudaError_t cudaStat; // cudaMalloc status
    cublasStatus_t stat; // CUBLAS functions status
    cublasHandle_t handle; // CUBLAS context
    int i,j; // i-row index, j-column index
    float* a; // a -mxn matrix on the host
    float* x; // x - n-vector on the host
    float* y; // y - m-vector on the host
    a=(float*)malloc(m*n*sizeof(float)); //host mem. alloc for a
    x=(float*)malloc(n*sizeof(float)); //host mem. alloc for x
    y=(float*)malloc(m*sizeof(float)); //host mem. alloc for y
    // define an mxn matrix a - column by column
    int ind=11; // a:
    for(j=0;j<n;j++){ // 11,17,23,29,35
        for(i=0;i<m;i++){ // 12,18,24,30,36
            a[IDX2C(i,j,m)]= (float) ind++; // 13,19,25,31,37
        } // 14,20,26,32,38
    } // 15,21,27,33,39
    } // 16,22,28,34,40

```



```

printf("a:\n");
for(i=0;i<m;i++){
    for(j=0;j<n;j++){
        printf("%4.0f",a[IDX2C(i,j,m)]); // print a row by row
    }
    printf("\n");
}
for(i=0;i<n;i++) x[i]=1.0f; // x={1,1,1,1,1}^T
for(i=0;i<m;i++) y[i]=0.0f; // y={0,0,0,0,0}^T
// on the device
float* d_a; // d_a - a on the device
float* d_x; // d_x - x on the device
float* d_y; // d_y - y on the device
cudaStat=cudaMalloc((void**)&d_a,m*n*sizeof(*a)); // device
// memory alloc for a
cudaStat=cudaMalloc((void**)&d_x,n*sizeof(*x)); // device
// memory alloc for x
cudaStat=cudaMalloc((void**)&d_y,m*sizeof(*y)); // device
// memory alloc for y

stat = cublasCreate(&handle);
stat = cublasSetMatrix(m,n,sizeof(*a),a,m,d_a,m); //cp a->d_a
stat = cublasSetVector(n,sizeof(*x),x,1,d_x,1); //cp x->d_x
stat = cublasSetVector(m,sizeof(*y),y,1,d_y,1); //cp y->d_y
float al=1.0f; // al=1
float bet=0.0f; // bet=1
// matrix-vector multiplication: d_y = al*d_a*d_x + bet*d_y
// d_a - mxn matrix; d_x - n-vector, d_y - m-vector;
// al,bet - scalars
stat=cublasSgemv(handle,CUBLAS_OP_N,m,n,&al,d_a,m,d_x,1,&bet,
                d_y,1);

stat=cublasGetVector(m,sizeof(*y),d_y,1,y,1); //copy d_y->y
printf("y after Sgemv::\n");
for(j=0;j<m;j++){
    printf("%5.0f",y[j]); // print y after Sgemv
    printf("\n");
}

cudaFree(d_a); // free device memory
cudaFree(d_x); // free device memory
cudaFree(d_y); // free device memory
cublasDestroy(handle); // destroy CUBLAS context
free(a); // free host memory
free(x); // free host memory
free(y); // free host memory
return EXIT_SUCCESS;
}
// a:
// 11 17 23 29 35
// 12 18 24 30 36
// 13 19 25 31 37
// 14 20 26 32 38
// 15 21 27 33 39
// 16 22 28 34 40

```

```

// y after Sgemv:
// 115           // [11 17 23 29 35] [1]
// 120           // [12 18 24 30 36] [1]
// 125           // [13 19 25 31 37]* [1]
// 130           // [14 20 26 32 38] [1]
// 135           // [15 21 27 33 39] [1]
// 140           // [16 22 28 34 40]

```

2.3.4 cublasSgemv – unified memory version

```

// nvcc 014sgemv.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))
#define m 6           // number of rows of a
#define n 5           // number of columns of a
int main(void){
    cublasHandle_t handle;           // CUBLAS context
    int i,j;           // i-row index, j-column index
    float* a;          // a -mxn matrix
    float* x;          // x - n-vector
    float* y;          // y - m-vector
    cudaMallocManaged(&a,m*n*sizeof(float)); //unified mem.for a
    cudaMallocManaged(&x,n*sizeof(float));   //unified mem.for x
    cudaMallocManaged(&y,m*sizeof(float));   //unified mem.for y
    // define an mxn matrix a - column by column
    int ind=11;         // a:
    for(j=0;j<n;j++){  // 11,17,23,29,35
        for(i=0;i<m;i++){ // 12,18,24,30,36
            a[IDX2C(i,j,m)]=(float)ind++; // 13,19,25,31,37
        } // 14,20,26,32,38
    } // 15,21,27,33,39
    // 16,22,28,34,40

    printf("a:\n");
    for(i=0;i<m;i++){
        for(j=0;j<n;j++){
            printf("%4.0f",a[IDX2C(i,j,m)]); // print a row by row
        }
        printf("\n");
    }
    for(i=0;i<n;i++) x[i]=1.0f;           // x={1,1,1,1,1}^T
    for(i=0;i<m;i++) y[i]=0.0f;          // y={0,0,0,0,0}^T
    cublasCreate(&handle);
    float al=1.0f;           // al=1
    float bet=0.0f;          // bet=1
    // matrix-vector multiplication: y = al*a*x + bet*y
    // a - mxn matrix; x - n-vector, y - m-vector;
    // al,bet - scalars

```

```

cublasSgemv(handle,CUBLAS_OP_N,m,n,&al,a,m,x,1,&bet,y,1);

cudaDeviceSynchronize();
printf("y after Sgemv::\n");
for(j=0;j<m;j++){
    printf("%5.0f",y[j]);           // print y after Sgemv
    printf("\n");
}
cudaFree(a);                       // free memory
cudaFree(x);                       // free memory
cudaFree(y);                       // free memory
cublasDestroy(handle);             // destroy CUBLAS context
return EXIT_SUCCESS;
}
// a:
//  11  17  23  29  35
//  12  18  24  30  36
//  13  19  25  31  37
//  14  20  26  32  38
//  15  21  27  33  39
//  16  22  28  34  40

// y after Sgemv:
//  115           //  [11  17  23  29  35]  [1]
//  120           //  [12  18  24  30  36]  [1]
//  125           //  [13  19  25  31  37]* [1]
//  130           //  [14  20  26  32  38]  [1]
//  135           //  [15  21  27  33  39]  [1]
//  140           //  [16  22  28  34  40]

```

2.3.5 cublasSger - rank one update

This function performs the rank-1 update

$$A = \alpha xy^T + A \quad \text{or} \quad A = \alpha xy^H + A,$$

where x, y are vectors, A is a $m \times n$ matrix and α is a scalar.

```

// nvcc 015sger.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))
#define m 6           // number of rows of a
#define n 5           // number of columns of a
int main(void){
    cudaError_t cudaStat;           // cudaMalloc status
    cublasStatus_t stat;           // CUBLAS functions status
    cublasHandle_t handle;         // CUBLAS context

```

```

int i,j;                                // i-row index, j-column index
float* a;                                // a -mxn matrix on the host
float* x;                                // x - n-vector on the host
float* y;                                // y - m-vector on the host
a=(float*)malloc(m*n*sizeof(float)); //host mem. alloc for a
x=(float*)malloc(n*sizeof(float));   //host mem. alloc for x
y=(float*)malloc(m*sizeof(float));   //host mem. alloc for y
// define an mxn matrix a column by column
int ind=11;                             // a:
for(j=0;j<n;j++){                       // 11,17,23,29,35
    for(i=0;i<m;i++){                   // 12,18,24,30,36
        a[IDX2C(i,j,m)]=(float)ind++;  // 13,19,25,31,37
    }                                   // 14,20,26,32,38
}                                       // 15,21,27,33,39
                                       // 16,22,28,34,40

printf("a:\n");
for(i=0;i<m;i++){
    for(j=0;j<n;j++){
        printf("%4.0f",a[IDX2C(i,j,m)]); // print a row by row
    }
    printf("\n");
}
for(i=0;i<m;i++) x[i]=1.0f;           // x={1,1,1,1,1,1}^T
for(i=0;i<n;i++) y[i]=1.0f;           // y={1,1,1,1,1}^T
// on the device
float* d_a;                             // d_a - a on the device
float* d_x;                             // d_x - x on the device
float* d_y;                             // d_y - y on the device
cudaStat=cudaMalloc((void**)&d_a,m*n*sizeof(*a)); // device
// memory alloc for a
cudaStat=cudaMalloc((void**)&d_x,m*sizeof(*x));   // device
// memory alloc for x
cudaStat=cudaMalloc((void**)&d_y,n*sizeof(*y));   // device
// memory alloc for y
stat = cublasCreate(&handle); // initialize CUBLAS context
stat = cublasSetMatrix(m,n,sizeof(*a),a,m,d_a,m); //cp a->d_a
stat = cublasSetVector(m,sizeof(*x),x,1,d_x,1); //cp x->d_x
stat = cublasSetVector(n,sizeof(*y),y,1,d_y,1); //cp y->d_y
float al=2.0f;                          // al=2
// rank-1 update of d_a: d_a = al*d_x*d_y^T + d_a
// d_a -mxn matrix; d_x -m-vector, d_y -n-vector; al -scalar

stat=cublasSger(handle,m,n,&al,d_x,1,d_y,1,d_a,m);

stat=cublasGetMatrix(m,n,sizeof(*a),d_a,m,a,m); //cp d_a->a
// print the updated a row by row
printf("a after Sger :\n");
for(i=0;i<m;i++){
    for(j=0;j<n;j++){
        printf("%4.0f",a[IDX2C(i,j,m)]); // print a after Sger
    }
    printf("\n");
}

```

```

    }
    cudaFree(d_a);                // free device memory
    cudaFree(d_x);                // free device memory
    cudaFree(d_y);                // free device memory
    cublasDestroy(handle);        // destroy CUBLAS context
    free(a);                      // free host memory
    free(x);                      // free host memory
    free(y);                      // free host memory
    return EXIT_SUCCESS;
}

// a:
//  11  17  23  29  35
//  12  18  24  30  36
//  13  19  25  31  37
//  14  20  26  32  38
//  15  21  27  33  39
//  16  22  28  34  40

// a after Sger :
//  13  19  25  31  37
//  14  20  26  32  38
//  15  21  27  33  39
//  16  22  28  34  40
//  17  23  29  35  41
//  18  24  30  36  42

//      [1]                [11  17  23  29  35]
//      [1]                [12  18  24  30  36]
//      [1]                [13  19  25  31  37]
// =  2*[ ]*[1,1,1,1,1] + [
//      [1]                [14  20  26  32  38]
//      [1]                [15  21  27  33  39]
//      [1]                [16  22  28  34  40]

```

2.3.6 cublasSger - unified memory version

```

// nvcc 015sger.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))
#define m 6                // number of rows of a
#define n 5                // number of columns of a
int main(void){
    cublasHandle_t handle;    // CUBLAS context
    int i,j;                 // i-row index, j-column index
    float* a;                // a -mxn matrix
    float* x;                // x - n-vector
    float* y;                // y - m-vector
    cudaMallocManaged(&a,m*n*sizeof(float)); //unified mem.for a
    cudaMallocManaged(&x,n*sizeof(float));    //unified mem.for x

```

```

    cudaMallocManaged(&y,m*sizeof(float)); //unified mem.for y
// define an mxn matrix a column by column
int ind=11;
for(j=0;j<n;j++){
    for(i=0;i<m;i++){
        a[IDX2C(i,j,m)]=(float)ind++;
    }
}

printf("a:\n");
for(i=0;i<m;i++){
    for(j=0;j<n;j++){
        printf("%4.0f",a[IDX2C(i,j,m)]); // print a row by row
    }
    printf("\n");
}
for(i=0;i<m;i++) x[i]=1.0f;
for(i=0;i<n;i++) y[i]=1.0f;
cublasCreate(&handle);
float al=2.0f;
// rank-1 update of a: a = al*x*y^T + a
// a - mxn matrix; x -m-vector, y - n-vector; al -scalar

cublasSger(handle,m,n,&al,x,1,y,1,a,m);

cudaDeviceSynchronize();

// print the updated a, row by row
printf("a after Sger :\n");
for(i=0;i<m;i++){
    for(j=0;j<n;j++){
        printf("%4.0f",a[IDX2C(i,j,m)]); // print a after Sger
    }
    printf("\n");
}

cudaFree(a);
cudaFree(x);
cudaFree(y);
cublasDestroy(handle);
return EXIT_SUCCESS;
}

// a:
// 11 17 23 29 35
// 12 18 24 30 36
// 13 19 25 31 37
// 14 20 26 32 38
// 15 21 27 33 39
// 16 22 28 34 40

// a after Sger :
```

```

// 13 19 25 31 37
// 14 20 26 32 38
// 15 21 27 33 39
// 16 22 28 34 40
// 17 23 29 35 41
// 18 24 30 36 42

//      [1]                [11 17 23 29 35]
//      [1]                [12 18 24 30 36]
//      [1]                [13 19 25 31 37]
// = 2*[ ]*[1,1,1,1,1] + [
//      [1]                [14 20 26 32 38]
//      [1]                [15 21 27 33 39]
//      [1]                [16 22 28 34 40]

```

2.3.7 cublasSsbmv - symmetric banded matrix-vector multiplication

This function performs the symmetric banded matrix-vector multiplication

$$y = \alpha Ax + \beta y,$$

where A is an $n \times n$ symmetric banded matrix with k subdiagonals and superdiagonals, x, y are vectors and α, β are scalars. The matrix A can be stored in lower (CUBLAS_FILL_MODE_LOWER) or upper mode (CUBLAS_FILL_MODE_UPPER). In both modes it is stored column by column. In lower mode the main diagonal is stored in row 0 (starting at position 0) the second subdiagonal in row 1 (starting at position 0) and so on.

```

// nvcc 016ssbmv.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))
#define n 6 // number of rows and columns of a
#define k 1 // number of subdiagonals and superdiagonals
int main(void){
    cudaError_t cudaStat; // cudaMalloc status
    cublasStatus_t stat; // CUBLAS functions status
    cublasHandle_t handle; // CUBLAS context
    int i,j; // row index, column index
    float *a; //nxn matrix a on the host //lower triangle of a:
    float *x; // n-vector x on the host //11
    float *y; // n-vector y on the host //17,12
    a=(float*)malloc(n*n*sizeof(float)); // 18,13
    // memory alloc for a on the host // 19,14
    x=(float*)malloc(n*sizeof(float)); // 20,15
    // memory alloc for x on the host // 21,16

```

```

    y=(float*)malloc(n*sizeof(float));          // mem. alloc for y
                                              //on the host
// main diagonal and subdiagonals of a in rows
int ind=11;
for(i=0;i<n;i++) a[i*n]=(float)ind++;          // main diagonal:
                                              // 11,12,13,14,15,16
for(i=0;i<n-1;i++) a[i*n+1]=(float)ind++;      // first subdiag.:
                                              // 17,18,19,20,21
for(i=0;i<n;i++){x[i]=1.0f;y[i]=0.0f;}        // x={1,1,1,1,1,1}^T
                                              // y={0,0,0,0,0,0}^T
// on the device
float* d_a;                                   // d_a - a on the device
float* d_x;                                   // d_x - x on the device
float* d_y;                                   // d_y - y on the device
cudaStat=cudaMalloc((void**)&d_a,n*n*sizeof(*a)); // device
                                              // memory alloc for a
cudaStat=cudaMalloc((void**)&d_x,n*sizeof(*x));   // device
                                              // memory alloc for x
cudaStat=cudaMalloc((void**)&d_y,n*sizeof(*y));   // device
                                              // memory alloc for y
stat = cublasCreate(&handle); // initialize CUBLAS context
stat = cublasSetMatrix(n,n,sizeof(*a),a,n,d_a,n); //cp a->d_a
stat = cublasSetVector(n,sizeof(*x),x,1,d_x,1); //cp x->d_x
stat = cublasSetVector(n,sizeof(*y),y,1,d_y,1); //cp y->d_y
float al=1.0f;                                // al=1
float bet=1.0f;                                // bet=1
// symmetric banded matrix-vector multiplication:
// d_y = al*d_a*d_x + bet*d_y,
// d_a - nxn symmetric banded matrix;
// d_x,d_y - n-vectors; al,bet - scalars

stat=cublasSsbmv(handle,CUBLAS_FILL_MODE_LOWER,n,k,&al,d_a,n,
                 d_x,1,&bet,d_y,1);

stat=cublasGetVector(n,sizeof(*y),d_y,1,y,1); //copy d_y->y
printf("y after Ssbmv:\n");
for(j=0;j<n;j++){
    printf("%7.0f",y[j]);                      // print y after Ssbmv
    printf("\n");
}
cudaFree(d_a);                                // free device memory
cudaFree(d_x);                                // free device memory
cudaFree(d_y);                                // free device memory
cublasDestroy(handle);                        // destroy CUBLAS context
free(a);                                       // free host memory
free(x);                                       // free host memory
free(y);                                       // free host memory
return EXIT_SUCCESS;
}

// y after Ssbmv:

```



```

//      28                      // [11 17          ] [1]    [28]
//      47                      // [17 12 18        ] [1]    [47]
//      50                      // [   18 13 19        ] [1] = [50]
//      53                      // [           19 14 20    ] [1]    [53]
//      56                      // [               20 15 21] [1]    [56]
//      37                      // [                   21 16] [1]    [37]

```

2.3.8 cublasSsbmv - unified memory version

```

// nvcc 016ssbmv.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))
#define n 6          // number of rows and columns of a
#define k 1          // number of subdiagonals and superdiagonals
int main(void){
    cublasHandle_t handle;                // CUBLAS context
    int i,j;                             // row index, column index
    float *a; // nxn matrix a             //lower triangle of a:
    float *x; // n-vector x                //11
    float *y; // n-vector y                //17,12
    cudaMallocManaged(&a,n*sizeof(float)); // 18,13
    // unified memory for a                // 19,14
    cudaMallocManaged(&x,n*sizeof(float)); // 20,15
    // unified memory for x                // 21,16
    cudaMallocManaged(&y,n*sizeof(float)); //unified mem.for y
    // main diagonal and subdiagonals of a in rows
    int ind=11;
    for(i=0;i<n;i++) a[i*n]=(float)ind++; // main diagonal:
                                           // 11,12,13,14,15,16
    for(i=0;i<n-1;i++) a[i*n+1]=(float)ind++; // first subdiag.:
                                           // 17,18,19,20,21
    for(i=0;i<n;i++){x[i]=1.0f;y[i]=0.0f;} // x={1,1,1,1,1,1}^T
                                           // y={0,0,0,0,0,0}^T
    cublasCreate(&handle);                // initialize CUBLAS context
    float al=1.0f;                        // al=1
    float bet=1.0f;                       // bet=1
    // symmetric banded matrix-vector multiplication:
    // y = al*a*x + bet*y,
    // a - nxn symmetric banded matrix;
    // x,y - n-vectors; al,bet - scalars

    cublasSsbmv(handle,CUBLAS_FILL_MODE_LOWER,n,k,&al,a,n,x,1,
                &bet,y,1);

    cudaDeviceSynchronize();
    printf("y after Ssbmv:\n");
    for(j=0;j<n;j++){
        printf("%7.0f",y[j]);                // print y after Ssbmv
        printf("\n");
    }
}

```

```

}
cudaFree(a);                // free memory
cudaFree(x);                // free memory
cudaFree(y);                // free memory
cublasDestroy(handle);      // destroy CUBLAS context
}
// y after Ssbmv:
//      28                // [11 17                ] [1]    [28]
//      47                // [17 12 18              ] [1]    [47]
//      50                // [   18 13 19            ] [1]    [50]
//      53                // [           19 14 20      ] [1]    [53]
//      56                // [                   20 15 21] [1]    [56]
//      37                // [                         21 16] [1]    [37]

```

2.3.9 cublasSspmv - symmetric packed matrix-vector multiplication

This function performs the symmetric packed matrix-vector multiplication

$$y = \alpha Ax + \beta y,$$

where A is a symmetric matrix in packed format, x, y are vectors and α, β - scalars. The symmetric $n \times n$ matrix A can be stored in lower (CUBLAS_FILL_MODE_LOWER) or upper mode (CUBLAS_FILL_MODE_UPPER). In lower mode the elements of the lower triangular part of A are packed together column by column without gaps.

```

// nvcc 017sspmv.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"
#define n 6                // number of rows and columns of a
int main(void){
    cudaError_t cudaStat;                // cudaMalloc status
    cublasStatus_t stat;                // CUBLAS functions status
    cublasHandle_t handle;              // CUBLAS context
    int i,j,l,m; // indices            // a:
    float *a; // lower triangle of nxn // 11
    // matrix a on the host            // 12,17
    float *x; // n-vector x on the host // 13,18,22
    float *y; // n-vector y on the host // 14,19,23,26
    a=(float*)malloc(n*(n+1)/2*sizeof(*a)); // 15,20,24,27,29
    //memory alloc for a on the host      // 16,21,25,28,30,31
    x=(float*)malloc(n*sizeof(float));    //memory alloc for x
    //on the host
    y=(float*)malloc(n*sizeof(float));    //memory alloc for y
    //on the host
    //define the lower triangle of a symmetric a in packed format
    //column by column without gaps

```

```

    for(i=0;i<n*(n+1)/2;i++) a[i]=(float)(11+i);
// print the upper triangle of a row by row
printf("upper triangle of a:\n");
l=n;j=0;m=0;
while(l>0){
    for(i=0;i<m;i++) printf("  ");
    for(i=j;i<j+1;i++) printf("%3.0f",a[i]);
    printf("\n");
    m++;j=j+1;l--;
}
for(i=0;i<n;i++){x[i]=1.0f;y[i]=0.0f;} // x={1,1,1,1,1,1}^T
// y={0,0,0,0,0,0}^T

// on the device
float* d_a; // d_a - a on the device
float* d_x; // d_x - x on the device
float* d_y; // d_y - y on the device
cudaStat=cudaMalloc((void**)&d_a,n*(n+1)/2*sizeof(*a));
// device memory alloc for a
cudaStat=cudaMalloc((void**)&d_x,n*sizeof(*x)); //device
// memory alloc for x
cudaStat=cudaMalloc((void**)&d_y,n*sizeof(*x)); //device
// memory alloc for y
stat = cublasCreate(&handle); // initialize CUBLAS context
stat = cublasSetVector(n*(n+1)/2,sizeof(*a),a,1,d_a,1);
// copy a->d_a
stat = cublasSetVector(n,sizeof(*x),x,1,d_x,1); //cp x->d_x
stat = cublasSetVector(n,sizeof(*y),y,1,d_y,1); //cp y->d_y
float al=1.0f; // al=1
float bet=1.0f; // bet=1
// symmetric packed matrix-vector multiplication:
// d_y = al*d_a*d_x + bet*d_y; d_a -symmetric nxn matrix
// in packed format; d_x,d_y - n-vectors; al,bet - scalars

stat=cublasSspmv(handle,CUBLAS_FILL_MODE_LOWER,n,&al,d_a,d_x,1,
                &bet,d_y,1);

stat=cublasGetVector(n,sizeof(*y),d_y,1,y,1); // copy d_y->y
printf("y after Sspmv:\n"); // print y after Sspmv
for(j=0;j<n;j++){
    printf("%7.0f",y[j]);
    printf("\n");
}
cudaFree(d_a); // free device memory
cudaFree(d_x); // free device memory
cudaFree(d_y); // free device memory
cublasDestroy(handle); // destroy CUBLAS context
free(a); // free host memory
free(x); // free host memory
free(y); // free host memory
return EXIT_SUCCESS;
}
// upper triangle of a:

```

```

// 11 12 13 14 15 16
//    17 18 19 20 21
//      22 23 24 25
//        26 27 28
//          29 30
//            31

// y after Sspmv:
//      81                // [11 12 13 14 15 16] [1]    [ 81]
//     107                // [12 17 18 19 20 21] [1]    [107]
//     125                // [13 18 22 23 24 25] [1]    [125]
//     137                // [14 19 23 26 27 28] [1]    [137]
//     145                // [15 20 24 27 29 30] [1]    [145]
//     151                // [16 21 25 28 30 31] [1]    [151]

```

2.3.10 cublasSspmv - unified memory version

```

// nvcc 017sspmv.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define n 6                // number of rows and columns of a
int main(void){
    cublasHandle_t handle;    // CUBLAS context
    int i,j,l,m; // indices    // a:
    float *a; //lower triangle of nxn    // 11
    // matrix a    // 12,17
    float *x; // n-vector x    // 13,18,22
    float *y; // n-vector y    // 14,19,23,26
    // 15,20,24,27,29
    // 16,21,25,28,30,31
    // unified memory for a
    cudaMallocManaged(&a,n*(n+1)/2*sizeof(float));
    cudaMallocManaged(&x,n*sizeof(float)); // unified mem.for x
    cudaMallocManaged(&y,n*sizeof(float)); // unified mem.for y
    //define the lower triangle of a symmetric a in packed format
    //column by column without gaps
    for(i=0;i<n*(n+1)/2;i++) a[i]=(float)(11+i);
    // print the upper triangle of a row by row
    printf("upper triangle of a:\n");
    l=n;j=0;m=0;
    while(l>0){
        for(i=0;i<m;i++) printf(" ");
        for(i=j;i<j+1;i++) printf("%3.0f",a[i]);
        printf("\n");
        m++;j=j+1;l--;
    }
    for(i=0;i<n;i++){x[i]=1.0f;y[i]=0.0f;} // x={1,1,1,1,1,1}^T
    // y={0,0,0,0,0,0}^T
    cublasCreate(&handle); // initialize CUBLAS context
    float al=1.0f; // al=1
    float bet=1.0f; // bet=1

```

```

// symmetric packed matrix-vector multiplication:
// y = al*a*x + bet*y; a -symmetric nxn matrix
// in packed format; x,y - n-vectors;  al,bet - scalars

cublasSspmv(handle,CUBLAS_FILL_MODE_LOWER,n,&al,a,x,1,&bet,y,1);

cudaDeviceSynchronize();
printf("y after Sspmv:\n");           // print y after Sspmv
for(j=0;j<n;j++){
    printf("%7.0f",y[j]);
    printf("\n");
}
cudaFree(a);                          // free memory
cudaFree(x);                          // free memory
cudaFree(y);                          // free memory
cublasDestroy(handle);                // destroy CUBLAS context
return EXIT_SUCCESS;
}
// upper triangle of a:
// 11 12 13 14 15 16
//    17 18 19 20 21
//      22 23 24 25
//        26 27 28
//          29 30
//            31

// y after Sspmv:
//      81                // [11 12 13 14 15 16] [1]    [ 81]
//     107                // [12 17 18 19 20 21] [1]    [107]
//     125                // [13 18 22 23 24 25] [1] = [125]
//     137                // [14 19 23 26 27 28] [1]    [137]
//     145                // [15 20 24 27 29 30] [1]    [145]
//     151                // [16 21 25 28 30 31] [1]    [151]

```

2.3.11 cublasSspr - symmetric packed rank-1 update

This function performs the symmetric packed rank-1 update

$$A = \alpha x x^T + A,$$

where A is a symmetric matrix in packed format, x is a vector and α is a scalar. The symmetric $n \times n$ matrix A can be stored in lower (CUBLAS_FILL_MODE_LOWER) or upper mode (CUBLAS_FILL_MODE_UPPER). In lower mode the elements of the lower triangular part of A are packed together column by column without gaps.

```

// nvcc 018sspr.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>

```

```

#include "cublas_v2.h"
#define n 6 // number of rows and columns of a
int main(void){
    cudaError_t cudaStat; // cudaMalloc status
    cublasStatus_t stat; // CUBLAS functions status
    cublasHandle_t handle; // CUBLAS context
    int i,j,l,m; //a:
    float *a; // lower triangle of a //11
    float *x; // n-vector x //12,17
    a=(float*)malloc(n*(n+1)/2*sizeof(*a)); //13,18,22
    // memory alloc for a on the host //14,19,23,26
    x=(float*)malloc(n*sizeof(float)); //15,20,24,27,29
    // memory alloc for x on the host //16,21,25,28,30,31
    //define the lower triangle of a symmetric a in packed format
    //column by column without gaps
    for(i=0;i<n*(n+1)/2;i++) a[i]=(float)(11+i);
    // print the upper triangle of a row by row
    printf("upper triangle of a:\n");
    l=n;j=0;m=0;
    while(l>0){
        for(i=0;i<m;i++) printf(" ");
        for(i=j;i<j+1;i++) printf("%3.0f",a[i]);
        printf("\n");
        m++;j=j+1;l--;
    }
    for(i=0;i<n;i++){x[i]=1.0f;} // x={1,1,1,1,1,1}^T
    // on the device
    float* d_a; // d_a - a on the device
    float* d_x; // d_x - x on the device
    cudaStat=cudaMalloc((void**)&d_a,n*(n+1)/2*sizeof(*a));
    // device memory alloc for a
    cudaStat=cudaMalloc((void**)&d_x,n*sizeof(*x)); // device
    // memory alloc for x
    stat = cublasCreate(&handle); // initialize CUBLAS context
    stat = cublasSetVector(n*(n+1)/2,sizeof(*a),a,1,d_a,1);
    // copy a -> d_a
    stat = cublasSetVector(n,sizeof(*x),x,1,d_x,1); // cp x->d_x
    float al=1.0f; // al=1
    // rank-1 update of a symmetric matrix d_a :
    // d_a = al*d_x*d_x^T + d_a
    // d_a - symmetric nxn matrix in packed format; d_x n-vector;
    // al - scalar

    stat=cublasSspr(handle,CUBLAS_FILL_MODE_LOWER,n,&al,d_x,1,d_a);

    stat=cublasGetVector(n*(n+1)/2,sizeof(*a),d_a,1,a,1);
    // copy d_a -> a
    printf("upper triangle of updated a after Sspr:\n");
    l=n;j=0;m=0;
    while(l>0){
        for(i=0;i<m;i++) printf(" "); // upper triangle
        for(i=j;i<j+1;i++) printf("%3.0f",a[i]); //of a after Sspr
    }
}

```

```

    printf("\n");
    m++; j=j+1; l--;
}
cudaFree(d_a);           // free device memory
cudaFree(d_x);           // free device memory
cublasDestroy(handle);   // destroy CUBLAS context
free(a);                 // free host memory
free(x);                 // free host memory
return EXIT_SUCCESS;
}
// upper triangle of a:
// 11 12 13 14 15 16
//    17 18 19 20 21
//      22 23 24 25
//        26 27 28
//          29 30
//            31

// upper triangle of a after Sspr://    [1]
// 12 13 14 15 16 17           //    [1]
//    18 19 20 21 22           //    [1]
//      23 24 25 26           // 1*[ ]*[1,1,1,1,1,1] + a
//        27 28 29           //    [1]
//          30 31           //    [1]
//            32           //    [1]

```

2.3.12 cublasSspr - unified memory version

```

// nvcc 018sspr.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define n 6           // number of rows and columns of a
int main(void){
    cublasHandle_t handle;           // CUBLAS context
    int i,j,l,m;                     //a:
    float *a; // lower triangle of a //11
    float *x; // n-vector x          //12,17
                                     //13,18,22
                                     //14,19,23,26
                                     //15,20,24,27,29
                                     //16,21,25,28,30,31
    // unified memory for a
    cudaMallocManaged(&a,n*(n+1)/2*sizeof(float));
    cudaMallocManaged(&x,n*sizeof(float)); //unified mem. for x
    //define the lower triangle of a symmetric a in packed format
    //column by column without gaps
    for(i=0;i<n*(n+1)/2;i++) a[i]=(float)(11+i);
    // print the upper triangle of a row by row
    printf("upper triangle of a:\n");
    l=n;j=0;m=0;
    while(l>0){

```

```

    for(i=0;i<m;i++) printf(" ");
    for(i=j;i<j+1;i++) printf("%3.0f",a[i]);
    printf("\n");
    m++;j=j+1;l--;
}
for(i=0;i<n;i++){x[i]=1.0f;} // x={1,1,1,1,1,1}^T
cublasCreate(&handle); // initialize CUBLAS context
float al=1.0f; // al=1
// rank-1 update of a symmetric matrix a :
// a = al*x*x^T + a
// a - symmetric nxn matrix in packed format; x n-vector;
// al - scalar

cublasSspr(handle,CUBLAS_FILL_MODE_LOWER,n,&al,x,1,a);

cudaDeviceSynchronize();
printf("upper triangle of updated a after Sspr:\n");
l=n;j=0;m=0;
while(l>0){
    for(i=0;i<m;i++) printf(" "); // upper triangle
    for(i=j;i<j+1;i++) printf("%3.0f",a[i]); //of a after Sspr
    printf("\n");
    m++;j=j+1;l--;
}
cudaFree(a); // free memory
cudaFree(x); // free memory
cublasDestroy(handle); // destroy CUBLAS context
return EXIT_SUCCESS;
}
// upper triangle of a:
// 11 12 13 14 15 16
//    17 18 19 20 21
//      22 23 24 25
//        26 27 28
//          29 30
//            31

// upper triangle of a after Sspr:// [1]
// 12 13 14 15 16 17 // [1]
//    18 19 20 21 22 // [1]
//      23 24 25 26 // 1*[ ]*[1,1,1,1,1,1] + a
//        27 28 29 // [1]
//          30 31 // [1]
//            32 // [1]

```

2.3.13 cublasSspr2 - symmetric packed rank-2 update

This function performs the symmetric packed rank-2 update

$$A = \alpha(xy^T + yx^T) + A,$$

where A is a symmetric matrix in packed format, x, y are vectors and α is a scalar. The symmetric $n \times n$ matrix A can be stored in lower (CUBLAS_FILL_MODE_LOWER) or upper mode (CUBLAS_FILL_MODE_UPPER). In lower mode the elements of the lower triangular part of A are packed together column by column without gaps.

```
// nvcc 019ssspr2.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"
#define n 6 // number of rows and columns of a
int main(void){
    cudaError_t cudaStat; // cudaMalloc status
    cublasStatus_t stat; // CUBLAS functions status
    cublasHandle_t handle; // CUBLAS context
    int i,j,l,m; // indices
    float *a; // lower triangle of an nxn matrix on the host
    float *x; // n-vector x on the host // a:
    float *y; // n-vector x on the host // 11
    a=(float*)malloc(n*(n+1)/2*sizeof(*a)); // 12,17
    // memory alloc for a on the host // 13,18,22
    x=(float*)malloc(n*sizeof(float)); // 14,19,23,26
    // memory alloc for x on the host // 15,20,24,27,29
    y=(float*)malloc(n*sizeof(float)); // 16,21,25,28,30,31
    // memory alloc for y on the host
    //define the lower triangle of a symmetric matrix a in packed
    // format column by column without gaps
    for(i=0;i<n*(n+1)/2;i++) a[i]=(float)(11+i);
    // print the upper triangle of a row by row
    printf("upper triangle of a:\n");
    l=n;j=0;m=0;
    while(l>0){
        for(i=0;i<m;i++) printf(" ");
        for(i=j;i<j+1;i++) printf("%3.0f",a[i]);
        printf("\n");
        m++;j=j+1;l--;
    }
    for(i=0;i<n;i++){x[i]=1.0f;y[i]=2.0;} // x={1,1,1,1,1,1}^T
                                           // y={2,2,2,2,2,2}^T
    // on the device
    float* d_a; // d_a - a on the device
    float* d_x; // d_x - x on the device
    float* d_y; // d_y - y on the device
    cudaStat=cudaMalloc((void**)&d_a,n*(n+1)/2*sizeof(*a));
    // device memory alloc for a
    cudaStat=cudaMalloc((void**)&d_x,n*sizeof(*x)); // device
    // memory alloc for x
    cudaStat=cudaMalloc((void**)&d_y,n*sizeof(*y)); // device
    // memory alloc for y
    stat = cublasCreate(&handle); // initialize CUBLAS context
    stat = cublasSetVector(n*(n+1)/2,sizeof(*a),a,1,d_a,1);
```

```

// copy a -> d_a
stat = cublasSetVector(n,sizeof(*x),x,1,d_x,1); //cp x->d_x
stat = cublasSetVector(n,sizeof(*y),y,1,d_y,1); //cp y->d_y
float al=1.0f; // al=1.0
// rank-2 update of symmetric matrix d_a :
// d_a = al*(d_x*d_y^T + d_y*d_x^T) + d_a
// d_a - symmetric nxn matrix in packed form; d_x,d_y -n-vect.
// al - scalar
stat=cublasSspr2(handle,CUBLAS_FILL_MODE_LOWER,n,&al,d_x,1,d_y,
1,d_a);

stat=cublasGetVector(n*(n+1)/2,sizeof(*a),d_a,1,a,1);
// copy d_a -> a
// print the updated upper triangle of a row by row
printf("upper triangle of updated a after Sspr2:\n");
l=n;j=0;m=0;
while(l>0){
    for(i=0;i<m;i++) printf(" ");
    for(i=j;i<j+1;i++) printf("%3.0f",a[i]);
    printf("\n");
    m++;j=j+1;l--;
}
cudaFree(d_a); // free device memory
cudaFree(d_x); // free device memory
cudaFree(d_y); // free device memory
cublasDestroy(handle); // destroy CUBLAS context
free(a); // free host memory
free(x); // free host memory
free(y); // free host memory
return EXIT_SUCCESS;
}

// upper triangle of a:
// 11 12 13 14 15 16
//    17 18 19 20 21
//      22 23 24 25
//        26 27 28
//          29 30
//            31

// upper triangle of a after Sspr2:
// 15 16 17 18 19 20
//    21 22 23 24 25
//      26 27 28 29
//        30 31 32
//          33 34
//            35
// [15 16 17 18 19 20] [1] [2]
// [16 21 22 23 24 25] [1] [2]
// [17 22 26 27 28 29] [1] [2]
// [    ]=1*([ ]*[2,2,2,2,2,2]+[ ]*[1,1,1,1,1,1])+a
// [18 23 27 30 31 32] [1] [2]
// [19 24 28 31 33 34] [1] [2]
// [20 25 29 33 34 35] [1] [2]

```

2.3.14 cublasSspr2 - unified memory version

```

// nvcc 019ssspr2.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define n 6 // number of rows and columns of a
int main(void){
    cublasHandle_t handle; // CUBLAS context
    int i,j,l,m; // indices
    float *a; // lower triangle of an nxn matrix
    float *x; // n-vector // a:
    float *y; // n-vector // 11
                                // 12,17
                                // 13,18,22
                                // 14,19,23,26
                                // 15,20,24,27,29
                                // 16,21,25,28,30,31

    // unified memory for a
    cudaMallocManaged(&a,n*(n+1)/2*sizeof(float));
    cudaMallocManaged(&x,n*sizeof(float)); //unified mem. for x
    cudaMallocManaged(&y,n*sizeof(float)); //unified mem. for y
    //define the lower triangle of a symmetric matrix a in packed
    // format column by column without gaps
    for(i=0;i<n*(n+1)/2;i++) a[i]=(float)(11+i);
    // print the upper triangle of a row by row
    printf("upper triangle of a:\n");
    l=n;j=0;m=0;
    while(l>0){
        for(i=0;i<m;i++) printf(" ");
        for(i=j;i<j+1;i++) printf("%3.0f",a[i]);
        printf("\n");
        m++;j=j+1;l--;
    }
    for(i=0;i<n;i++){x[i]=1.0f;y[i]=2.0f;} // x={1,1,1,1,1,1}^T
                                // y={2,2,2,2,2,2}^T

    cublasCreate(&handle); // initialize CUBLAS context
    float al=1.0f; // al=1.0

    // rank-2 update of symmetric matrix a :
    // a = al*(x*y^T + y*x^T) + a
    // a - symmetric nxn matrix in packed form; x,y - n-vect.;
    // al - scalar

    cublasSspr2(handle,CUBLAS_FILL_MODE_LOWER,n,&al,x,1,y,1,a);

    cudaDeviceSynchronize();
    // print the updated upper triangle of a row by row
    printf("upper triangle of updated a after Sspr2:\n");
    l=n;j=0;m=0;
    while(l>0){
        for(i=0;i<m;i++) printf(" ");
        for(i=j;i<j+1;i++) printf("%3.0f",a[i]);
        printf("\n");
    }

```

```

    m++;j=j+1;l--;
}
cudaFree(a);                                // free memory
cudaFree(x);                                // free memory
cudaFree(y);                                // free memory
cublasDestroy(handle);                       // destroy CUBLAS context
return EXIT_SUCCESS;
}
// upper triangle of a:
// 11 12 13 14 15 16
//    17 18 19 20 21
//      22 23 24 25
//        26 27 28
//          29 30
//            31

// upper triangle of a after Sspr2:
// 15 16 17 18 19 20
//    21 22 23 24 25
//      26 27 28 29
//        30 31 32
//          33 34
//            35
// [15 16 17 18 19 20]    [1]                [2]
// [16 21 22 23 24 25]    [1]                [2]
// [17 22 26 27 28 29]    [1]                [2]
// [      ]=1*([      ]*[2,2,2,2,2,2]+[      ]*[1,1,1,1,1,1])+a
// [18 23 27 30 31 32]    [1]                [2]
// [19 24 28 31 33 34]    [1]                [2]
// [20 25 29 33 34 35]    [1]                [2]

```

2.3.15 cublasSsymv - symmetric matrix-vector multiplication

This function performs the symmetric matrix-vector multiplication.

$$y = \alpha Ax + \beta y,$$

where A is an $n \times n$ symmetric matrix, x, y are vectors and α, β are scalars. The matrix A can be stored in lower (CUBLAS_FILL_MODE_LOWER) or upper (CUBLAS_FILL_MODE_UPPER) mode.

```

// nvcc 020ssymv.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))
#define n 6                                // number of rows and columns of a
int main(void){
    cudaError_t cudaStat;                    // cudaMalloc status
    cublasStatus_t stat;                     // CUBLAS functions status

```

```

cublasHandle_t handle;                                // CUBLAS context
int i,j;                                                // i - row index, j - column index
float* a;                                              // nxn matrix on the host
float* x;                                              // n-vector on the host
float* y;                                              // n-vector on the host
a=(float*)malloc(n*n*sizeof(float));                // host memory for a
x=(float*)malloc(n*sizeof(float));                  // host memory for x
y=(float*)malloc(n*sizeof(float));                  // host memory for y
// define the lower triangle of an nxn symmetric matrix a
// in lower mode column by column
int ind=11;
for(j=0;j<n;j++){
    for(i=0;i<n;i++){
        if(i>=j){
            a[IDX2C(i,j,n)]=(float)ind++;
        }
    }
}
// print the lower triangle of a row by row
printf("lower triangle of a:\n");
for(i=0;i<n;i++){
    for(j=0;j<n;j++){
        if(i>=j)
            printf("%5.0f",a[IDX2C(i,j,n)]);
    }
    printf("\n");
}
for(i=0;i<n;i++){x[i]=1.0f;y[i]=1.0;} // x={1,1,1,1,1,1}^T
// y={1,1,1,1,1,1}^T

// on the device
float* d_a; // d_a - a on the device
float* d_x; // d_x - x on the device
float* d_y; // d_y - y on the device
cudaStat=cudaMalloc((void**)&d_a,n*n*sizeof(*a)); // device
// memory alloc for a
cudaStat=cudaMalloc((void**)&d_x,n*sizeof(*x)); // device
// memory alloc for x
cudaStat=cudaMalloc((void**)&d_y,n*sizeof(*y)); // device
// memory alloc for y
stat = cublasCreate(&handle); // initialize CUBLAS context
stat = cublasSetMatrix(n,n,sizeof(*a),a,n,d_a,n);
// copy a -> d_a
stat = cublasSetVector(n,sizeof(*x),x,1,d_x,1); // cp x->d_x
stat = cublasSetVector(n,sizeof(*y),y,1,d_y,1); // cp y->d_y
float al=1.0f; // al=1.0
float bet=1.0f; // bet=1.0
// symmetric matrix-vector multiplication:
// d_y = al*d_a*d_x + bet*d_y
// d_a - nxn symmetric matrix; d_x,d_y - n-vectors;
// al,bet - scalars

```

```

stat=cublasSsymv(handle,CUBLAS_FILL_MODE_LOWER,n,&al,d_a,n,
                d_x,1,&bet,d_y,1);

stat=cublasGetVector(n,sizeof(float),d_y,1,y,1); // d_y->y
printf("y after Ssymv:\n");                      // print y after Ssymv
for(j=0;j<n;j++)
printf("%7.0f\n",y[j]);
cudaFree(d_a);                                     // free device memory
cudaFree(d_x);                                     // free device memory
cudaFree(d_y);                                     // free device memory
cublasDestroy(handle);                             // destroy CUBLAS context
free(a);                                           // free host memory
free(x);                                           // free host memory
free(y);                                           // free host memory
return EXIT_SUCCESS;
}
// lower triangle of a:
//   11
//   12  17
//   13  18  22
//   14  19  23  26
//   15  20  24  27  29
//   16  21  25  28  30  31

// y after Ssymv:
//   82
//  108
//  126
//  138
//  146
//  152
//
//   [11  12  13  14  15  16] [1]      [1]    [ 82]
//   [12  17  18  19  20  21] [1]      [1]    [108]
//  1*[13  18  22  23  24  25]*[1] + 1*[1] = [126]
//   [14  19  23  26  27  28] [1]      [1]    [138]
//   [15  20  24  27  29  30] [1]      [1]    [146]
//   [16  21  25  28  30  31] [1]      [1]    [152]

```

2.3.16 cublasSsymv - unified memory version

```

// nvcc 020ssymv.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))
#define n 6                                     // number of rows and columns of a
int main(void){
    cublasHandle_t handle;                     // CUBLAS context
    int i,j;                                  // i - row index, j - column index
    float* a;                                  // nxn matrix
    float* x;                                  // n-vector

```

```

float* y; // n-vector
cudaMallocManaged(&a,n*n*sizeof(float)); //unif.memory for a
cudaMallocManaged(&x,n*sizeof(float)); //unif.memory for x
cudaMallocManaged(&y,n*sizeof(float)); //unif.memory for y
// define the lower triangle of an nxn symmetric matrix a
// in lower mode column by column
int ind=11; // a:
for(j=0;j<n;j++){ // 11
    for(i=0;i<n;i++){ // 12,17
        if(i>=j){ // 13,18,22
            a[IDX2C(i,j,n)]=(float)ind++; // 14,19,23,26
        } // 15,20,24,27,29
    } // 16,21,25,28,30,31
}
// print the lower triangle of a row by row
printf("lower triangle of a:\n");
for(i=0;i<n;i++){
    for(j=0;j<n;j++){
        if(i>=j)
            printf("%5.0f",a[IDX2C(i,j,n)]);
    }
    printf("\n");
}
for(i=0;i<n;i++){x[i]=1.0f;y[i]=1.0;} // x={1,1,1,1,1,1}^T
// y={1,1,1,1,1,1}^T

cublasCreate(&handle); // initialize CUBLAS context
float al=1.0f; // al=1.0
float bet=1.0f; // bet=1.0
// symmetric matrix-vector multiplication:
// y = al*a*x + bet*y
// a - nxn symmetric matrix; x,y - n-vectors;
// al,bet - scalars

cublasSsymv(handle,CUBLAS_FILL_MODE_LOWER,n,&al,a,n,x,1,&bet,
            y,1);

cudaDeviceSynchronize();
printf("y after Ssymv:\n"); // print y after Ssymv
for(j=0;j<n;j++)
    printf("%7.0f\n",y[j]);
cudaFree(a); // free memory
cudaFree(x); // free memory
cudaFree(y); // free memory
cublasDestroy(handle); // destroy CUBLAS context
return EXIT_SUCCESS;
}
// lower triangle of a:
// 11
// 12 17
// 13 18 22
// 14 19 23 26
// 15 20 24 27 29
// 16 21 25 28 30 31

```

```

// y after Ssymv:
//      82
//     108
//     126
//     138
//     146
//     152
//
//      [11  12  13  14  15  16] [1]      [1]      [ 82]
//      [12  17  18  19  20  21] [1]      [1]      [108]
//  1*[13  18  22  23  24  25]*[1] + 1*[1] = [126]
//      [14  19  23  26  27  28] [1]      [1]      [138]
//      [15  20  24  27  29  30] [1]      [1]      [146]
//      [16  21  25  28  30  31] [1]      [1]      [152]

```

2.3.17 cublasSsyr - symmetric rank-1 update

This function performs the symmetric rank-1 update

$$A = \alpha x x^T + A,$$

where A is an $n \times n$ symmetric matrix, x is a vector and α is a scalar. A is stored in column-major format in lower (CUBLAS_FILL_MODE_LOWER) or upper (CUBLAS_FILL_MODE_UPPER) mode.

```

// nvcc 021ssyr.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))
#define n 6 // number of rows and columns of a
int main(void){
    cudaError_t cudaStat; // cudaMalloc status
    cublasStatus_t stat; // CUBLAS functions status
    cublasHandle_t handle; // CUBLAS context
    int i,j; // i - row index, j - column index
    float* a; // nxn matrix on the host
    float* x; // n-vector on the host
    a=(float*)malloc(n*n*sizeof(float)); // host memory for a
    x=(float*)malloc(n*sizeof(float)); // host memory for x
    // define the lower triangle of an nxn symmetric matrix a
    // in lower mode column by column
    int ind=11; // a:
    for(j=0;j<n;j++){ // 11
        for(i=0;i<n;i++){ // 12,17
            if(i>=j){ // 13,18,22
                a[IDX2C(i,j,n)]=(float)ind++; // 14,19,23,26
            } // 15,20,24,27,29
        } // 16,21,25,28,30,31
    }
}

```



```

    }
    // print the lower triangle of a row by row
    printf("lower triangle of a:\n");
    for(i=0;i<n;i++){
        for(j=0;j<n;j++){
            if(i>=j)
                printf("%5.0f",a[IDX2C(i,j,n)]);
        }
        printf("\n");
    }
    for(i=0;i<n;i++){x[i]=1.0f;} // x={1,1,1,1,1,1}^T

    // on the device
    float* d_a; // d_a - a on the device
    float* d_x; // d_x - x on the device
    cudaStat=cudaMalloc((void**)&d_a,n*n*sizeof(*a)); // device
    // memory alloc for a
    cudaStat=cudaMalloc((void**)&d_x,n*sizeof(*x)); // device
    // memory alloc for x
    stat = cublasCreate(&handle); // initialize CUBLAS context
    stat = cublasSetMatrix(n,n,sizeof(*a),a,n,d_a,n); //a -> d_a
    stat = cublasSetVector(n,sizeof(*x),x,1,d_x,1); //cp x->d_x
    float al=1.0f; // al=1.0
    // symmetric rank-1 update of d_a: d_a = al*d_x*d_x^T + d_a
    // d_a - nxn symmetric matrix; d_x - n-vector; al - scalar

    stat=cublasSsyr(handle,CUBLAS_FILL_MODE_LOWER,n,&al,d_x,1,
        d_a,n);
    stat=cublasGetMatrix(n,n,sizeof(*a),d_a,n,a,n); //cp d_a->a
    // print the lower triangle of the updated a after Ssyr
    printf("lower triangle of updated a after Ssyr :\n");
    for(i=0;i<n;i++){
        for(j=0;j<n;j++){
            if(i>=j)
                printf("%5.0f",a[IDX2C(i,j,n)]);
        }
        printf("\n");
    }
    cudaFree(d_a); // free device memory
    cudaFree(d_x); // free device memory
    cublasDestroy(handle); // destroy CUBLAS context
    free(a); // free host memory
    free(x); // free host memory
    return EXIT_SUCCESS;
}

// lower triangle of a:
// 11
// 12 17
// 13 18 22
// 14 19 23 26
// 15 20 24 27 29
// 16 21 25 28 30 31

```

```

// lower triangle of a after Ssyr://      [1]
//   12                          //      [1]
//   13   18                      //      [1]
//   14   19   23                  // a=1*[ ]*[1,1,1,1,1,1]+ a
//   15   20   24   27              //      [1]
//   16   21   25   28   30         //      [1]
//   17   22   26   29   31   32    //      [1]

```

2.3.18 cublasSsyr - unified memory version

```

// nvcc 021ssyr.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))
#define n 6 // number of rows and columns of a
int main(void){
    cublasHandle_t handle; // CUBLAS context
    int i,j; // i - row index, j - column index
    float* a; // nxn matrix
    float* x; // n-vector
    cudaMallocManaged(&a,n*n*sizeof(float)); //unif.memory for a
    cudaMallocManaged(&x,n*sizeof(float)); //unif.memory for x
    // define the lower triangle of an nxn symmetric matrix a
    // in lower mode column by column
    int ind=11; // a:
    for(j=0;j<n;j++){ // 11
        for(i=0;i<n;i++){ // 12,17
            if(i>=j){ // 13,18,22
                a[IDX2C(i,j,n)]=(float)ind++; // 14,19,23,26
            } // 15,20,24,27,29
        } // 16,21,25,28,30,31
    }
    // print the lower triangle of a row by row
    printf("lower triangle of a:\n");
    for(i=0;i<n;i++){
        for(j=0;j<n;j++){
            if(i>=j)
                printf("%5.0f",a[IDX2C(i,j,n)]);
        }
        printf("\n");
    }
    for(i=0;i<n;i++){x[i]=1.0f;} // x={1,1,1,1,1,1}^T
    cublasCreate(&handle); // initialize CUBLAS context
    float al=1.0f; // al=1.0
    // symmetric rank-1 update of a: a = al*x*x^T + a
    // a - nxn symmetric matrix; x - n-vector; al - scalar

    cublasSsyr(handle,CUBLAS_FILL_MODE_LOWER,n,&al,x,1,a,n);

```

```

    cudaDeviceSynchronize();
// print the lower triangle of the updated a after Ssyr
printf("lower triangle of updated a after Ssyr :\n");
for(i=0;i<n;i++){
    for(j=0;j<n;j++){
        if(i>=j)
            printf("%5.0f",a[IDX2C(i,j,n)]);
    }
    printf("\n");
}
cudaFree(a); // free memory
cudaFree(x); // free memory
cublasDestroy(handle); // destroy CUBLAS context
return EXIT_SUCCESS;
}
// lower triangle of a:
//  11
//  12  17
//  13  18  22
//  14  19  23  26
//  15  20  24  27  29
//  16  21  25  28  30  31

// lower triangle of a after Ssyr://      [1]
//  12          //      [1]
//  13  18          //      [1]
//  14  19  23          // a=1*[ ]*[1,1,1,1,1,1]+ a
//  15  20  24  27          //      [1]
//  16  21  25  28  30          //      [1]
//  17  22  26  29  31  32 //      [1]

```

2.3.19 cublasSsyr2 - symmetric rank-2 update

This function performs the symmetric rank-2 update

$$A = \alpha(xy^T + yx^T) + A,$$

where A is an $n \times n$ symmetric matrix, x, y are vectors and α is a scalar. A is stored in column-major format in lower (CUBLAS_FILL_MODE_LOWER) or upper (CUBLAS_FILL_MODE_UPPER) mode.

```

// nvcc 022ssyr2.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))
#define n 6 // number of rows and columns of a
int main(void){
    cudaError_t cudaStat; // cudaMalloc status

```

```

cublasStatus_t stat;           // CUBLAS functions status
cublasHandle_t handle;         // CUBLAS context
int i,j;                       // i - row index, j - column index
float* a;                      // nxn matrix on the host
float* x;                      // n-vector on the host
float* y;                      // n-vector on the host
a=(float*)malloc(n*n*sizeof(float)); // host memory for a
x=(float*)malloc(n*sizeof(float));   // host memory for x
y=(float*)malloc(n*sizeof(float));   // host memory for y
// define the lower triangle of an nxn symmetric matrix a
// in lower mode column by column
int ind=11;                    // a:
for(j=0;j<n;j++){             // 11
    for(i=0;i<n;i++){          // 12,17
        if(i>=j){              // 13,18,22
            a[IDX2C(i,j,n)]=(float)ind++; // 14,19,23,26
        }                      // 15,20,24,27,29
    }                          // 16,21,25,28,30,31
}
// print the lower triangle of a row by row
printf("lower triangle of a:\n");
for(i=0;i<n;i++){
    for(j=0;j<n;j++){
        if(i>=j)
            printf("%5.0f",a[IDX2C(i,j,n)]);
    }
    printf("\n");
}
for(i=0;i<n;i++){x[i]=1.0f;y[i]=2.0f;} // x={1,1,1,1,1,1}^T
// y={2,2,2,2,2,2}^T

// on the device
float* d_a;                    // d_a - a on the device
float* d_x;                    // d_x - x on the device
float* d_y;                    // d_y - y on the device
cudaStat=cudaMalloc((void**)&d_a,n*n*sizeof(*a)); // device
// memory alloc for a
cudaStat=cudaMalloc((void**)&d_x,n*sizeof(*x));   // device
// memory alloc for x
cudaStat=cudaMalloc((void**)&d_y,n*sizeof(*y));   // device
// memory alloc for y
stat = cublasCreate(&handle); // initialize CUBLAS context
stat = cublasSetMatrix(n,n,sizeof(*a),a,n,d_a,n); //a -> d_a
stat = cublasSetVector(n,sizeof(*x),x,1,d_x,1); //cp x->d_x
stat = cublasSetVector(n,sizeof(*y),y,1,d_y,1); //cp y->d_y
float al=1.0f;                // al=1
// symmetric rank-2 update of d_a:
// d_a = al*(d_x*d_y^T + d_y*d_x^T) + d_a
// d_a - nxn symmetric matrix; d_x,d_y -n-vectors; al -scalar

stat=cublasSsyr2(handle,CUBLAS_FILL_MODE_LOWER,n,&al,d_x,1,
                d_y,1,d_a,n);

```

```

    stat=cublasGetMatrix(n,n,sizeof(*a),d_a,n,a,n); //cp d_a->a
// print the lower triangle of the updated a
printf("lower triangle of a after Ssyr2 :\n");
for(i=0;i<n;i++){
    for(j=0;j<n;j++){
        if(i>=j)
            printf("%5.0f",a[IDX2C(i,j,n)]);
    }
    printf("\n");
}
cudaFree(d_a);                // free device memory
cudaFree(d_x);                // free device memory
cudaFree(d_y);                // free device memory
cublasDestroy(handle);        // destroy CUBLAS context
free(a);                      // free host memory
free(x);                      // free host memory
free(y);                      // free host memory
return EXIT_SUCCESS;
}

// lower triangle of a:
//   11
//   12   17
//   13   18   22
//   14   19   23   26
//   15   20   24   27   29
//   16   21   25   28   30   31

// lower triangle of a after Ssyr2 :
//   15
//   16   21
//   17   22   26
//   18   23   27   30
//   19   24   28   31   33
//   20   25   29   32   34   35

//[15 16 17 18 19 20]    [1]                [2]
//[16 21 22 23 24 25]    [1]                [2]
//[17 22 26 27 28 29]    [1]                [2]
//[
                        ]=1*([ ]*[2,2,2,2,2,2]+[ ]*[1,1,1,1,1,1])+a
//[18 23 27 30 31 32]    [1]                [2]
//[19 24 28 31 33 34]    [1]                [2]
//[20 25 29 33 34 35]    [1]                [2]

```

2.3.20 cublasSsyr2 - unified memory version

```

// nvcc 022ssyr2.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))

```

```

#define n 6                                // number of rows and columns of a
int main(void){
    cublasHandle_t handle;                  // CUBLAS context
    int i,j;                               // i - row index, j - column index
    float* a;                              // nxn matrix
    float* x;                              // n-vector
    float* y;                              // n-vector
    cudaMallocManaged(&a,n*n*sizeof(float)); //unif.memory for a
    cudaMallocManaged(&x,n*sizeof(float));   //unif.memory for x
    cudaMallocManaged(&y,n*sizeof(float));   //unif.memory for y
    // define the lower triangle of an nxn symmetric matrix a
    // in lower mode column by column
    int ind=11;                            // a:
    for(j=0;j<n;j++){                      // 11
        for(i=0;i<n;i++){                  // 12,17
            if(i>=j){                      // 13,18,22
                a[IDX2C(i,j,n)]=(float)ind++; // 14,19,23,26
            }                               // 15,20,24,27,29
        }                                  // 16,21,25,28,30,31
    }
    // print the lower triangle of a row by row
    printf("lower triangle of a:\n");
    for(i=0;i<n;i++){
        for(j=0;j<n;j++){
            if(i>=j)
                printf("%5.0f",a[IDX2C(i,j,n)]);
        }
        printf("\n");
    }
    for(i=0;i<n;i++){x[i]=1.0f;y[i]=2.0;} // x={1,1,1,1,1,1}^T
                                           // y={2,2,2,2,2,2}^T

    cublasCreate(&handle);                  // initialize CUBLAS context
    float al=1.0f;                          // al=1
    // symmetric rank-2 update of a:
    // a = al*(x*y^T + y*x^T) + a
    // a - nxn symmetric matrix; x,y - n-vectors; al -scalar

    cublasSsyr2(handle,CUBLAS_FILL_MODE_LOWER,n,&al,x,1,y,1,a,n);

    cudaDeviceSynchronize();
    // print the lower triangle of the updated a
    printf("lower triangle of a after Ssyr2 :\n");
    for(i=0;i<n;i++){
        for(j=0;j<n;j++){
            if(i>=j)
                printf("%5.0f",a[IDX2C(i,j,n)]);
        }
        printf("\n");
    }
    cudaFree(a);                            // free memory
    cudaFree(x);                            // free memory

```

```

    cudaFree(y);                                     // free memory
    cublasDestroy(handle);                           // destroy CUBLAS context
    return EXIT_SUCCESS;
}

// lower triangle of a:
//   11
//   12   17
//   13   18   22
//   14   19   23   26
//   15   20   24   27   29
//   16   21   25   28   30   31

// lower triangle of a after Ssyr2 :
//   15
//   16   21
//   17   22   26
//   18   23   27   30
//   19   24   28   31   33
//   20   25   29   32   34   35

//[15 16 17 18 19 20]      [1]                [2]
//[16 21 22 23 24 25]      [1]                [2]
//[17 22 26 27 28 29]      [1]                [2]
//[
                        ]=1*([ ]*[2,2,2,2,2,2]+[ ]*[1,1,1,1,1,1])+a
//[18 23 27 30 31 32]      [1]                [2]
//[19 24 28 31 33 34]      [1]                [2]
//[20 25 29 33 34 35]      [1]                [2]

```

2.3.21 cublasStbmv - triangular banded matrix-vector multiplication

This function performs the triangular banded matrix-vector multiplication

$$x = op(A)x,$$

where A is a triangular banded matrix, x is a vector and $op(A)$ can be equal to A (CUBLAS_OP_N case), A^T (transposition) in CUBLAS_OP_T case or A^H (Hermitian transposition) in CUBLAS_OP_C case. The matrix A is stored in lower (CUBLAS_FILL_MODE_LOWER) or upper (CUBLAS_FILL_MODE_UPPER) mode. In the lower mode the main diagonal is stored in row 0, the first subdiagonal in row 1 and so on. If the diagonal of the matrix A has non-unit elements, then the parameter CUBLAS_DIAG_NON_UNIT should be used (in the opposite case - CUBLAS_DIAG_UNIT).

```

// nvcc 023stbmv.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"

```

```

#define n 6                                // number of rows and columns of a
#define k 1                                // number of subdiagonals
int main(void){
    cudaError_t cudaStat;                  // cudaMalloc status
    cublasStatus_t stat;                   // CUBLAS functions status
    cublasHandle_t handle;                 // CUBLAS context
    int i,j;                              // lower triangle of a:
    float *a; //nxn matrix a on the host // 11
    float *x; //n-vector x on the host // 17,12
    a=(float*)malloc(n*n*sizeof(float)); // 18,13
    x=(float*)malloc(n*sizeof(float)); // 19,14
    // main diagonal and subdiagonals // 20,15
    // of a in rows // 21,16
    int ind=11;
    // main diagonal: 11,12,13,14,15,16 in row 0:
    for(i=0;i<n;i++) a[i*n]=(float)ind++;
    // first subdiagonal: 17,18,19,20,21 in row 1:
    for(i=0;i<n-1;i++) a[i*n+1]=(float)ind++;
    for(i=0;i<n;i++){x[i]=1.0f;} // x={1,1,1,1,1,1}^T
    // on the device
    float* d_a; // d_a - a on the device
    float* d_x; // d_x - x on the device
    cudaStat=cudaMalloc((void**)&d_a,n*n*sizeof(*a)); // device
    // memory alloc for a
    cudaStat=cudaMalloc((void**)&d_x,n*sizeof(*x)); // device
    // memory alloc for x
    stat = cublasCreate(&handle); // initialize CUBLAS context
    stat = cublasSetMatrix(n,n,sizeof(*a),a,n,d_a,n); // a->d_a
    stat = cublasSetVector(n,sizeof(*x),x,1,d_x,1); // cp x->d_x
    // triangular banded matrix-vector multiplication:
    // d_x = d_a*d_x;
    // d_a - nxn lower triangular banded matrix; d_x - n-vector

    stat=cublasStbmv(handle,CUBLAS_FILL_MODE_LOWER,CUBLAS_OP_N,
        CUBLAS_DIAG_NON_UNIT,n,k,d_a,n,d_x,1);

    stat=cublasGetVector(n,sizeof(*x),d_x,1,x,1); //copy d_x->x
    printf("x after Stbmv :\n"); // print x after Stbmv
    for(j=0;j<n;j++){
        printf("%7.0f",x[j]);
        printf("\n");
    }
    cudaFree(d_a); // free device memory
    cudaFree(d_x); // free device memory
    cublasDestroy(handle); // destroy CUBLAS context
    free(a); // free host memory
    free(x); // free host memory
    return EXIT_SUCCESS;
}

// x after Stbmv :

```



```

//      11      //      [11  0  0  0  0  0] [1]
//      29      //      [17 12  0  0  0  0] [1]
//      31      //      = [ 0 18 13  0  0  0]*[1]
//      33      //      [ 0  0 19 14  0  0] [1]
//      35      //      [ 0  0  0 20 15  0] [1]
//      37      //      [ 0  0  0  0 21 16] [1]

```

2.3.22 cublasStbmv - unified memory version

```

// nvcc 023stbmv.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define n 6 // number of rows and columns of a
#define k 1 // number of subdiagonals
int main(void){
    cublasHandle_t handle; // CUBLAS context
    int i,j; // lower triangle of a:
    float *a; //nxn matrix a // 11
    float *x; //n-vector x // 17,12
                                // 18,13
                                // 19,14
                                // 20,15
                                // 21,16
    cudaMallocManaged(&a,n*n*sizeof(float)); //unif.memory for a
    cudaMallocManaged(&x,n*sizeof(float)); //unif.memory for x
    // main diagonal and subdiagonals of a in rows
    int ind=11;
    // main diagonal: 11,12,13,14,15,16 in row 0:
    for(i=0;i<n;i++) a[i*n]=(float)ind++;
    // first subdiagonal: 17,18,19,20,21 in row 1:
    for(i=0;i<n-1;i++) a[i*n+1]=(float)ind++;
    for(i=0;i<n;i++){x[i]=1.0f;} // x={1,1,1,1,1,1}^T
    cublasCreate(&handle); // initialize CUBLAS context
    // triangular banded matrix-vector multiplication: x = a*x;
    // a - nxn lower triangular banded matrix; x - n-vector

    cublasStbmv(handle,CUBLAS_FILL_MODE_LOWER,CUBLAS_OP_N,
                CUBLAS_DIAG_NON_UNIT,n,k,a,n,x,1);

    cudaDeviceSynchronize();
    printf("x after Stbmv :\n"); // print x after Stbmv
    for(j=0;j<n;j++){
        printf("%7.0f",x[j]);
        printf("\n");
    }
    cudaFree(a); // free memory
    cudaFree(x); // free memory
    cublasDestroy(handle); // destroy CUBLAS context
    return EXIT_SUCCESS;
}

```

```
// x after Stbmv :
//      11          //      [11   0   0   0   0   0] [1]
//      29          //      [17  12   0   0   0   0] [1]
//      31          // = [ 0   18  13   0   0   0]*[1]
//      33          //      [ 0   0  19  14   0   0] [1]
//      35          //      [ 0   0   0  20  15   0] [1]
//      37          //      [ 0   0   0   0  21  16] [1]
```

2.3.23 cublasStbsv - solve the triangular banded linear system

This function solves the triangular banded linear system with a single right-hand-side

$$op(A)x = b,$$

where A is a triangular banded matrix, x, b are vectors and $op(A)$ can be equal to A (CUBLAS_OP_N case), A^T (transposition) in CUBLAS_OP_T case, or A^H (Hermitian transposition) in CUBLAS_OP_C case. The matrix A is stored in lower (CUBLAS_FILL_MODE_LOWER) or upper (CUBLAS_FILL_MODE_UPPER) mode. In the lower mode the main diagonal is stored in row 0, the first subdiagonal in row 1 and so on. If the diagonal of the matrix A has non-unit elements, then the parameter CUBLAS_DIAG_NON_UNIT should be used (in the opposite case - CUBLAS_DIAG_UNIT).

```
// nvcc 024stbsv.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"
#define n 6          // number of rows and columns of a
#define k 1          // number of subdiagonals
int main(void){
    cudaError_t cudaStat;          // cudaMalloc status
    cublasStatus_t stat;          // CUBLAS functions status
    cublasHandle_t handle;        // CUBLAS context
    int i,j;                      // lower triangle of a:
    float *a; //nxn matrix a on the host // 11
    float *x; //n-vector x on the host // 17,12
    a=(float*)malloc(n*n*sizeof(float)); // 18,13
    // memory allocation for a on the host // 19,14
    x=(float*)malloc(n*sizeof(float)); // 20,15
    // memory allocation for x on the host // 21,16
    //main diagonal and subdiagonals of a in rows:
    int ind=11;
    // main diagonal: 11,12,13,14,15,16 in row 0
    for(i=0;i<n;i++) a[i*n]=(float)ind++;
    // first subdiagonal: 17,18,19,20,21 in row 1
```

```

    for(i=0;i<n-1;i++) a[i*n+1]=(float)ind++;
    for(i=0;i<n;i++){x[i]=1.0f;}           // x={1,1,1,1,1,1}^T
// on the device
float* d_a;                               // d_a - a on the device
float* d_x;                               // d_x - x on the device
cudaStat=cudaMalloc((void**)&d_a,n*n*sizeof(*a)); // device
                                           // memory alloc for a
cudaStat=cudaMalloc((void**)&d_x,n*sizeof(*x));   // device
                                           // memory alloc for x
stat = cublasCreate(&handle); // initialize CUBLAS context
stat = cublasSetMatrix(n,n,sizeof(*a),a,n,d_a,n); //a -> d_a
stat = cublasSetVector(n,sizeof(*x),x,1,d_x,1); //cp x->d_x
// solve a triangular banded linear system: d_a*X=d_x;
// the solution X overwrites the right hand side d_x;
// d_a - nxn banded lower triangular matrix; d_x - n-vector

    stat=cublasStbsv(handle,CUBLAS_FILL_MODE_LOWER,CUBLAS_OP_N,
                     CUBLAS_DIAG_NON_UNIT,n,k,d_a,n,d_x,1);

    stat=cublasGetVector(n,sizeof(*x),d_x,1,x,1); //copy d_x->x
// print the solution
    printf("solution : \n");                // print x after Stbsv
    for(j=0;j<n;j++){
        printf("%9.6f",x[j]);
        printf("\n");
    }
    cudaFree(d_a);                          // free device memory
    cudaFree(d_x);                          // free device memory
    cublasDestroy(handle);                  // destroy CUBLAS context
    free(a);                                // free host memory
    free(x);                                // free host memory
    return EXIT_SUCCESS;
}
// solution :
// 0.090909 // [11    0    0    0    0    0] [ 0.090909] [1]
// -0.045455 // [17   12    0    0    0    0] [-0.045455] [1]
// 0.139860 // [ 0   18   13    0    0    0] [ 0.139860] = [1]
// -0.118382 // [ 0    0   19   14    0    0] [-0.118382] [1]
// 0.224509 // [ 0    0    0   20   15    0] [ 0.224509] [1]
// -0.232168 // [ 0    0    0    0   21   16] [-0.232168] [1]

```

2.3.24 cublasStbsv - unified memory version

```

// nvcc 024stbsv.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define n 6 // number of rows and columns of a
#define k 1 // number of subdiagonals
int main(void){
    cublasHandle_t handle; // CUBLAS context

```

```

int i,j;                                // lower triangle of a:
float *a; //nxn matrix a                // 11
float *x; //n-vector x                  // 17,12
                                         // 18,13
                                         // 19,14
                                         // 20,15
                                         // 21,16
    cudaMallocManaged(&a,n*n*sizeof(float)); //unif.memory for a
    cudaMallocManaged(&x,n*sizeof(float));   //unif.memory for x
//main diagonal and subdiagonals of a in rows:
    int ind=11;
// main diagonal: 11,12,13,14,15,16 in row 0
    for(i=0;i<n;i++) a[i*n]=(float)ind++;
// first subdiagonal: 17,18,19,20,21 in row 1
    for(i=0;i<n-1;i++) a[i*n+1]=(float)ind++;
    for(i=0;i<n;i++){x[i]=1.0f;}          // x={1,1,1,1,1,1}^T
    cublasCreate(&handle);                // initialize CUBLAS context
// solve a triangular banded linear system: a*X=x;
// the solution X overwrites the right hand side x;
// a - nxn banded lower triangular matrix; x - n-vector

    cublasStbsv(handle,CUBLAS_FILL_MODE_LOWER,CUBLAS_OP_N,
                CUBLAS_DIAG_NON_UNIT,n,k,a,n,x,1);

    cudaDeviceSynchronize();
// print the solution
    printf("solution :\n");                // print x after Stbsv
    for(j=0;j<n;j++){
        printf("%9.6f",x[j]);
        printf("\n");
    }
    cudaFree(a);                          // free memory
    cudaFree(x);                          // free memory
    cublasDestroy(handle);                // destroy CUBLAS context
    return EXIT_SUCCESS;
}
// solution :
// 0.090909 // [11    0    0    0    0    0] [ 0.090909] [1]
// -0.045455 // [17   12    0    0    0    0] [-0.045455] [1]
// 0.139860 // [ 0   18   13    0    0    0] [ 0.139860] = [1]
// -0.118382 // [ 0    0   19   14    0    0] [-0.118382] [1]
// 0.224509 // [ 0    0    0   20   15    0] [ 0.224509] [1]
// -0.232168 // [ 0    0    0    0   21   16] [-0.232168] [1]

```

2.3.25 cublasStpmv - triangular packed matrix-vector multiplication

This function performs the triangular packed matrix-vector multiplication

$$x = op(A)x,$$

where A is a triangular packed matrix, x is a vector and $op(A)$ can be equal to A (CUBLAS_OP_N case), A^T (CUBLAS_OP_T case - transposition) or A^H (CUBLAS_OP_C case - conjugate transposition). A can be stored in lower (CUBLAS_FILL_MODE_LOWER) or upper (CUBLAS_FILL_MODE_UPPER) mode. In lower mode the elements of the lower triangular part of A are packed together column by column without gaps. If the diagonal of the matrix A has non-unit elements, then the parameter CUBLAS_DIAG_NON_UNIT should be used (in the opposite case - CUBLAS_DIAG_UNIT).

```
// nvcc 025stpmv.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"
#define n 6 // number of rows and columns of a
int main(void){
    cudaError_t cudaStat; // cudaMalloc status
    cublasStatus_t stat; // CUBLAS functions status
    cublasHandle_t handle; // CUBLAS context
    int i,j;
    float* a; // lower triangle of an nxn matrix on the host
    float* x; // n-vector on the host
    a=(float*)malloc(n*(n+1)/2*sizeof(float)); // host memory
    // alloc for a
    x=(float*)malloc(n*sizeof(float)); // host memory alloc for x
    //define a triangular matrix a in packed format column
    // by column without gaps //a:
    for(i=0;i<n*(n+1)/2;i++) //11
        a[i]=(float)(11+i); //12,17
    for(i=0;i<n;i++){x[i]=1.0f;} //13,18,22
    // x={1,1,1,1,1,1}^T //14,19,23,2
    // on the device //15,20,24,27,29
    float* d_a; // d_a - a on the device //16,21,25,28,30,31
    float* d_x; // d_x - x on the device
    cudaStat=cudaMalloc((void**)&d_a,n*(n+1)/2*sizeof(*a));
    // device memory alloc for a
    cudaStat=cudaMalloc((void**)&d_x,n*sizeof(*x)); // device
    // memory alloc for x
    stat = cublasCreate(&handle); // initialize CUBLAS context
    stat = cublasSetVector(n*(n+1)/2,sizeof(*a),a,1,d_a,1);
    // copy a -> d_a
    stat = cublasSetVector(n,sizeof(*x),x,1,d_x,1); // cp x->d_x
    // triangular packed matrix-vector multiplication:
    // d_x = d_a*d_x; d_a - nxn lower triangular matrix
    // in packed format; d_x - n-vector

    stat=cublasStpmv(handle,CUBLAS_FILL_MODE_LOWER,CUBLAS_OP_N,
        CUBLAS_DIAG_NON_UNIT,n,d_a,d_x,1);

    stat=cublasGetVector(n,sizeof(*x),d_x,1,x,1); //copy d_x->x
```

```

printf("x after Stpmv :\n");          // print x after Stpmv
for(j=0;j<n;j++){
    printf("%7.0f",x[j]);
    printf("\n");
}
cudaFree(d_a);                        // free device memory
cudaFree(d_x);                        // free device memory
cublasDestroy(handle);                // destroy CUBLAS context
free(a);                             // free host memory
free(x);                             // free host memory
return EXIT_SUCCESS;
}
// x after Stpmv :
//      11      //      [11   0   0   0   0   0] [1]
//      29      //      [12  17   0   0   0   0] [1]
//      53      //      = [13  18  22   0   0   0]*[1]
//      82      //      [14  19  23  26   0   0] [1]
//     115      //      [15  20  24  27  29   0] [1]
//     151      //      [16  21  25  28  30  31] [1]

```

2.3.26 cublasStpmv - unified memory version

```

// nvcc 025stpmv.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define n 6                      // number of rows and columns of a
int main(void){
    cublasHandle_t handle;        // CUBLAS context
    int i,j;
    float* a;                    // lower triangle of an nxn matrix
    float* x;                    // n-vector
    //unified memory for a
    cudaMallocManaged(&a,n*(n+1)/2*sizeof(float));
    cudaMallocManaged(&x,n*sizeof(float)); //unif. memory for x
    //define a triangular matrix a in packed format column
    // by column without gaps      //a:
    for(i=0;i<n*(n+1)/2;i++){      //11
        a[i]=(float)(11+i);        //12,17
        for(i=0;i<n;i++){x[i]=1.0f;} //13,18,22
    }
    // x={1,1,1,1,1,1}^T          //14,19,23,2
                                    //15,20,24,27,29
                                    //16,21,25,28,30,31
    cublasCreate(&handle);          // initialize CUBLAS context
    // triangular packed matrix-vector multiplication:
    // x = a*x;  a - nxn lower triangular matrix
    // in packed format;  x - n-vector

    cublasStpmv(handle,CUBLAS_FILL_MODE_LOWER,CUBLAS_OP_N,
                CUBLAS_DIAG_NON_UNIT,n,a,x,1);

```



```

    x=(float*)malloc(n*sizeof(float)); //host memory alloc for x
    // define a triangular a in packed format
    // column by column without gaps          //a:
    for(i=0;i<n*(n+1)/2;i++)                //11
        a[i]=(float)(11+i);                  //12,17
    for(i=0;i<n;i++){x[i]=1.0f;}             //13,18,22
    // x={1,1,1,1,1,1}^T                    //14,19,23,26
    // on the device                          //15,20,24,27,29
    float* d_a; // d_a - a on the device      //16,21,25,28,30,31
    float* d_x; // d_x - x on the device
    cudaStat=cudaMalloc((void**)&d_a,n*(n+1)/2*sizeof(*a));
                                // device memory alloc for a
    cudaStat=cudaMalloc((void**)&d_x,n*sizeof(*x)); // device
                                // memory alloc for x
    stat = cublasCreate(&handle); // initialize CUBLAS context
    stat = cublasSetVector(n*(n+1)/2,sizeof(*a),a,1,d_a,1);
                                // copy a -> d_a
    stat = cublasSetVector(n,sizeof(*x),x,1,d_x,1); // cp x->d_x
    // solve the packed triangular linear system: d_a*X=d_x,
    // the solution X overwrites the right hand side d_x
    // d_a -nxn lower triang. matrix in packed form; d_x -n-vect.

    stat=cublasStpsv(handle,CUBLAS_FILL_MODE_LOWER,CUBLAS_OP_N,
                      CUBLAS_DIAG_NON_UNIT,n,d_a,d_x,1);

    stat=cublasGetVector(n,sizeof(*x),d_x,1,x,1); //copy d_x->x
    printf("solution : \n");                      // print x after Stpsv
    for(j=0;j<n;j++){
        printf("%9.6f",x[j]);
        printf("\n");
    }
    cudaFree(d_a); // free device memory
    cudaFree(d_x); // free device memory
    cublasDestroy(handle); // destroy CUBLAS context
    free(a); // free host memory
    free(x); // free host memory
    return EXIT_SUCCESS;
}
// solution :
// 0.090909 //      [11  0  0  0  0  0] [ 0.090909] [1]
// -0.005348 //      [12 17  0  0  0  0] [-0.005348] [1]
// -0.003889 //      [13 18 22  0  0  0]*[-0.003889]=[1]
// -0.003141 //      [14 19 23 26  0  0] [-0.003141] [1]
// -0.002708 //      [15 20 24 27 29  0] [-0.002708] [1]
// -0.002446 //      [16 21 25 28 30 31] [-0.002446] [1]

```

2.3.28 cublasStpsv - unified memory version

```

// nvcc 026stpsv.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"

```



```

#define n 6                                // number of rows and columns of a
int main(void){
    cublasHandle_t handle;                  // CUBLAS context
    int i,j;                               // i-row index, j-column index
    float* a;                              // nxn matrix a
    float* x;                              // n-vector x
    // unified memory for a
    cudaMallocManaged(&a,n*(n+1)/2*sizeof(float));
    cudaMallocManaged(&x,n*sizeof(float)); //unif. memory for x
    // define a triangular a in packed format
    // column by column without gaps        //a:
    for(i=0;i<n*(n+1)/2;i++){              //11
        a[i]=(float)(11+i);                //12,17
    }                                       //13,18,22
    for(i=0;i<n;i++){x[i]=1.0f;}          //14,19,23,26
    // x={1,1,1,1,1,1}^T                 //15,20,24,27,29
                                         //16,21,25,28,30,31

    cublasCreate(&handle);                 // initialize CUBLAS context
    // solve the packed triangular linear system: a*X=x,
    // the solution X overwrites the right hand side x
    // a - nxn lower triang. matrix in packed form; x - n-vector

    cublasStpsv(handle,CUBLAS_FILL_MODE_LOWER,CUBLAS_OP_N,
                CUBLAS_DIAG_NON_UNIT,n,a,x,1);

    cudaDeviceSynchronize();
    printf("solution : \n");               // print x after Stpsv
    for(j=0;j<n;j++){
        printf("%9.6f",x[j]);
        printf("\n");
    }
    cudaFree(a);                          // free memory
    cudaFree(x);                          // free memory
    cublasDestroy(handle);                 // destroy CUBLAS context
    return EXIT_SUCCESS;
}

// solution :
// 0.090909 //      [11  0  0  0  0  0] [ 0.090909] [1]
// -0.005348 //      [12 17  0  0  0  0] [-0.005348] [1]
// -0.003889 //      [13 18 22  0  0  0]*[-0.003889]=[1]
// -0.003141 //      [14 19 23 26  0  0] [-0.003141] [1]
// -0.002708 //      [15 20 24 27 29  0] [-0.002708] [1]
// -0.002446 //      [16 21 25 28 30 31] [-0.002446] [1]

```

2.3.29 cublasStrmv - triangular matrix-vector multiplication

This function performs the triangular matrix-vector multiplication

$$x = op(A)x,$$

where A is a triangular $n \times n$ matrix, x is an n -vector and $op(A)$ can be equal to A (CUBLAS_OP_N case), A^T (transposition) in CUBLAS_OP_T case or A^H

(conjugate transposition) in CUBLAS_OP_C case. The matrix A can be stored in lower (CUBLAS_FILL_MODE_LOWER) or upper (CUBLAS_FILL_MODE_UPPER) mode. If the diagonal of the matrix A has non-unit elements, then the parameter CUBLAS_DIAG_NON_UNIT should be used (in the opposite case - CUBLAS_DIAG_UNIT).

```
// nvcc 027strmv.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))
#define n 6 // number of rows and columns of a
int main(void){
    cudaError_t cudaStat; // cudaMalloc status
    cublasStatus_t stat; // CUBLAS functions status
    cublasHandle_t handle; // CUBLAS context
    int i,j; // i-row index, j-column index
    float* a; // nxn matrix a on the host
    float* x; // n-vector x on the host
    a=(float*)malloc(n*n*sizeof(*a)); //host memory alloc for a
    x=(float*)malloc(n*sizeof(*x)); //host memory alloc for x
    // define an nxn triangular matrix a in lower mode
    // column by column
    int ind=11; // a:
    for(j=0;j<n;j++){ // 11
        for(i=0;i<n;i++){ // 12,17
            if(i>=j){ // 13,18,22
                a[IDX2C(i,j,n)]=(float)ind++; // 14,19,23,26
            } // 15,20,24,27,29
        } // 16,21,25,28,30,31
    }
    for(i=0;i<n;i++) x[i]=1.0f; // x={1,1,1,1,1,1}^T
    // on the device
    float* d_a; // d_a - a on the device
    float* d_x; // d_x - x on the device
    cudaStat=cudaMalloc((void**)&d_a,n*n*sizeof(*a)); // device
    // memory alloc for a
    cudaStat=cudaMalloc((void**)&d_x,n*sizeof(*x)); //device
    // memory alloc for x
    stat = cublasCreate(&handle); // initialize CUBLAS context
    stat = cublasSetMatrix(n,n,sizeof(*a),a,n,d_a,n); // a->d_a
    stat = cublasSetVector(n,sizeof(*x),x,1,d_x,1); //cp x->d_x
    // triangular matrix-vector multiplication: d_x = d_a*d_x
    // d_a - triangular nxn matrix in lower mode; d_x - n-vector

    stat=cublasStrmv(handle,CUBLAS_FILL_MODE_LOWER,CUBLAS_OP_N,
        CUBLAS_DIAG_NON_UNIT,n,d_a,n,d_x,1);

    stat=cublasGetVector(n,sizeof(*x),d_x,1,x,1); //copy d_x->x
    printf("multiplication result :\n"); // print x after Strmv
```

```

    for(j=0;j<n;j++){
        printf("%7.0f",x[j]);
        printf("\n");
    }
    cudaFree(d_a);                // free device memory
    cudaFree(d_x);                // free device memory
    cublasDestroy(handle);        // destroy CUBLAS context
    free(a);                      // free host memory
    free(x);                      // free host memory
    return EXIT_SUCCESS;
}
// multiplication result :
//      11      //      [11   0   0   0   0   0] [1]
//      29      //      [12  17   0   0   0   0] [1]
//      53      //      = [13  18  22   0   0   0]*[1]
//      82      //      [14  19  23  26   0   0] [1]
//      115     //      [15  20  24  27  29   0] [1]
//      151     //      [16  21  25  28  30  31] [1]

```

2.3.30 cublasStrmv - unified memory version

```

// nvcc 027strmv.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))
#define n 6                // number of rows and columns of a
int main(void){
    cublasHandle_t handle;   // CUBLAS context
    int i,j;                // i-row index, j-column index
    float* a;               // nxn matrix a
    float* x;               // n-vector x
    cudaMallocManaged(&a,n*n*sizeof(float)); //unif.memory for a
    cudaMallocManaged(&x,n*sizeof(float));    //unif.memory for x
    // define an nxn triangular matrix a in lower mode
    // column by column
    int ind=11;              // a:
    for(j=0;j<n;j++){        // 11
        for(i=0;i<n;i++){    // 12,17
            if(i>=j){        // 13,18,22
                a[IDX2C(i,j,n)]=(float)ind++; // 14,19,23,26
            }                // 15,20,24,27,29
        }                   // 16,21,25,28,30,31
    }
    for(i=0;i<n;i++) x[i]=1.0f; // x={1,1,1,1,1,1}^T
    cublasCreate(&handle);      // initialize CUBLAS context
    // triangular matrix-vector multiplication: x = a*x
    // a - triangular nxn matrix in lower mode; x - n-vector

    cublasStrmv(handle,CUBLAS_FILL_MODE_LOWER,CUBLAS_OP_N,
                CUBLAS_DIAG_NON_UNIT,n,a,n,x,1);

```

```

    cudaDeviceSynchronize();
    printf("multiplication result :\n"); // print x after Strmv
    for(j=0;j<n;j++){
        printf("%7.0f",x[j]);
        printf("\n");
    }
    cudaFree(a); // free memory
    cudaFree(x); // free memory
    cublasDestroy(handle); // destroy CUBLAS context
    return EXIT_SUCCESS;
}
// multiplication result :
//      11      //      [11   0   0   0   0   0] [1]
//      29      //      [12  17   0   0   0   0] [1]
//      53      //      = [13  18  22   0   0   0]*[1]
//      82      //      [14  19  23  26   0   0] [1]
//      115     //      [15  20  24  27  29   0] [1]
//      151     //      [16  21  25  28  30  31] [1]

```

2.3.31 cublasStrsv - solve the triangular linear system

This function solves the triangular linear system

$$op(A)x = b,$$

where A is a triangular $n \times n$ matrix, x, b are n -vectors and $op(A)$ can be equal to A (CUBLAS_OP_N case), A^T (transposition) in CUBLAS_OP_T case or A^H (conjugate transposition) in CUBLAS_OP_C case. A can be stored in lower (CUBLAS_FILL_MODE_LOWER) or upper (CUBLAS_FILL_MODE_UPPER) mode. If the diagonal of the matrix A has non-unit elements, then the parameter CUBLAS_DIAG_NON_UNIT should be used (in the opposite case - CUBLAS_DIAG_UNIT).

```

// nvcc 028strsv.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))
#define n 6 // number of rows and columns of a
int main(void){
    cudaError_t cudaStat; // cudaMalloc status
    cublasStatus_t stat; // CUBLAS functions status
    cublasHandle_t handle; // CUBLAS context
    int i,j; // i-row index, j-column index
    float* a; // nxn matrix a on the host
    float* x; // n-vector x on the host
    a=(float*)malloc(n*n*sizeof(*a)); //host memory alloc for a

```

```

    x=(float*)malloc(n*sizeof(*x));    //host memory alloc for x
    // define an nxn triangular matrix a in lower mode
    // column by column
    int ind=11;                        // a:
    for(j=0;j<n;j++){                  // 11
        for(i=0;i<n;i++){              // 12,17
            if(i>=j){                  // 13,18,22
                a[IDX2C(i,j,n)]=(float)ind++;    // 14,19,23,26
            }                          // 15,20,24,27,29
        }                             // 16,21,25,28,30,31
    }
    for(i=0;i<n;i++) x[i]=1.0f;        // x={1,1,1,1,1,1}^T
    // on the device
    float* d_a;                        // d_a - a on the device
    float* d_x;                        // d_x - x on the device
    cudaStat=cudaMalloc((void**)&d_a,n*n*sizeof(*a)); // device
                                                // memory alloc for a
    cudaStat=cudaMalloc((void**)&d_x,n*sizeof(*x));    //device
                                                // memory alloc for x
    stat = cublasCreate(&handle);    // initialize CUBLAS context
    stat = cublasSetMatrix(n,n,sizeof(*a),a,n,d_a,n); // a->d_a
    stat = cublasSetVector(n,sizeof(*x),x,1,d_x,1); //cp x->d_x
    // solve the triangular linear system: d_a*X=d_x,
    // the solution X overwrites the right hand side d_x,
    // d_a - nxn triangular matrix in lower mode; d_x - n-vector

    stat=cublasStrsv(handle,CUBLAS_FILL_MODE_LOWER,CUBLAS_OP_N,
                      CUBLAS_DIAG_NON_UNIT,n,d_a,n,d_x,1);

    stat=cublasGetVector(n,sizeof(*x),d_x,1,x,1); //copy x->d_x
    printf("solution : \n");                    // print x after Strsv
    for(j=0;j<n;j++){
        printf("%9.6f",x[j]);
        printf("\n");
    }
    cudaFree(d_a);                            // free device memory
    cudaFree(d_x);                            // free device memory
    cublasDestroy(handle);                    // destroy CUBLAS context
    free(a);                                  // free host memory
    free(x);                                  // free host memory
    return EXIT_SUCCESS;
}
// solution :

// 0.090909 // [11  0  0  0  0  0] [ 0.090909] [1]
// -0.005348 // [12 17  0  0  0  0] [-0.005348] [1]
// -0.003889 // [13 18 22  0  0  0]*[-0.003889]=[1]
// -0.003141 // [14 19 23 26  0  0] [-0.003141] [1]
// -0.002708 // [15 20 24 27 29  0] [-0.002708] [1]
// -0.002446 // [16 21 25 28 30 31] [-0.002446] [1]

```

2.3.32 cublasStrsv - unified memory version

```

// nvcc 028strsv.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))
#define n 6 // number of rows and columns of a
int main(void){
    cublasHandle_t handle; // CUBLAS context
    int i,j; // i-row index, j-column index
    float* a; // nxn matrix
    float* x; // n-vector
    cudaMallocManaged(&a,n*n*sizeof(float)); //unif.memory for a
    cudaMallocManaged(&x,n*sizeof(float)); //unif.memory for x
    // define an nxn triangular matrix a in lower mode
    // column by column
    int ind=11; // a:
    for(j=0;j<n;j++){ // 11
        for(i=0;i<n;i++){ // 12,17
            if(i>=j){ // 13,18,22
                a[IDX2C(i,j,n)]=(float)ind++; // 14,19,23,26
            } // 15,20,24,27,29
        } // 16,21,25,28,30,31
    }
    for(i=0;i<n;i++) x[i]=1.0f; // x={1,1,1,1,1,1}^T
    cublasCreate(&handle); // initialize CUBLAS context
    // solve the triangular linear system: a*X=x,
    // the solution X overwrites the right hand side x,
    // a - nxn triangular matrix in lower mode; x - n-vector

    cublasStrsv(handle,CUBLAS_FILL_MODE_LOWER,CUBLAS_OP_N,
                CUBLAS_DIAG_NON_UNIT,n,a,n,x,1);

    cudaDeviceSynchronize();
    printf("solution :\n"); // print x after Strsv
    for(j=0;j<n;j++){
        printf("%9.6f",x[j]);
        printf("\n");
    }
    cudaFree(a); // free memory
    cudaFree(x); // free memory
    cublasDestroy(handle); // destroy CUBLAS context
    return EXIT_SUCCESS;
}
// solution :
// 0.090909 // [11 0 0 0 0 0] [ 0.090909] [1]
// -0.005348 // [12 17 0 0 0 0] [-0.005348] [1]
// -0.003889 // [13 18 22 0 0 0]*[-0.003889]=[1]
// -0.003141 // [14 19 23 26 0 0] [-0.003141] [1]
// -0.002708 // [15 20 24 27 29 0] [-0.002708] [1]
// -0.002446 // [16 21 25 28 30 31] [-0.002446] [1]

```



```

    printf("\n");
}
for(i=0;i<n;i++){x[i].x=1.0f;y[i].x=1.0;}
//x={1,1,1,1,1,1}^T  y={1,1,1,1,1,1}^T
// on the device
cuComplex* d_a; // d_a - a on the device
cuComplex* d_x; // d_x - x on the device
cuComplex* d_y; // d_y - y on the device
cudaStat=cudaMalloc((void**)&d_a,n*n*sizeof(cuComplex));
//device memory alloc for a
cudaStat=cudaMalloc((void**)&d_x,n*sizeof(cuComplex));
//device memory alloc for x
cudaStat=cudaMalloc((void**)&d_y,n*sizeof(cuComplex));
//device memory alloc for y
stat = cublasCreate(&handle); // initialize CUBLAS context
stat = cublasSetMatrix(n,n,sizeof(*a),a,n,d_a,n);
// copy a -> d_a
stat = cublasSetVector(n,sizeof(*x),x,1,d_x,1); //cp x->d_x
stat = cublasSetVector(n,sizeof(*y),y,1,d_y,1); //cp y->d_y
cuComplex al={1.0f,0.0f}; // al=1
cuComplex bet={1.0f,0.0f}; // bet=1
// Hermitian matrix-vector multiplication:
// d_y=al*d_a*d_x + bet*d_y
// d_a - nxn Hermitian matrix; d_x,d_y - n-vectors;
// al,bet -scalars

stat=cublasChemv(handle,CUBLAS_FILL_MODE_LOWER,n,&al,d_a,n,
                d_x,1,&bet,d_y,1);

stat=cublasGetVector(n,sizeof(*y),d_y,1,y,1); //copy d_y->y
printf("y after Chemv:\n"); // print y after Chemv
for(j=0;j<n;j++){
    printf("%4.0f+%1.0f*I",y[j].x,y[j].y);
    printf("\n");
}
cudaFree(d_a); // free device memory
cudaFree(d_x); // free device memory
cudaFree(d_y); // free device memory
cublasDestroy(handle); // destroy CUBLAS context
free(a); // free host memory
free(y); // free host memory
return EXIT_SUCCESS;
}
// lower triangle of a:
// 11+0*I
// 12+0*I 17+0*I
// 13+0*I 18+0*I 22+0*I
// 14+0*I 19+0*I 23+0*I 26+0*I
// 15+0*I 20+0*I 24+0*I 27+0*I 29+0*I
// 16+0*I 21+0*I 25+0*I 28+0*I 30+0*I 31+0*I

// y after Chemv:

```



```

// 82+0*I
// 108+0*I
// 126+0*I
// 138+0*I
// 146+0*I
// 152+0*I
//
//      [11  12  13  14  15  16] [1]      [1]      [ 82]
//      [12  17  18  19  20  21] [1]      [1]      [108]
// 1*[13  18  22  23  24  25]*[1] + 1*[1] = [126]
//      [14  19  23  26  27  28] [1]      [1]      [138]
//      [15  20  24  27  29  30] [1]      [1]      [146]
//      [16  21  25  28  30  31] [1]      [1]      [152]

```

2.3.34 cublasChemv - unified memory version

```

// nvcc 029Chemv.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#include "cuComplex.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))
#define n 6 // number of rows and columns of a
int main(void){
    cublasHandle_t handle; // CUBLAS context
    int i,j; // i-row index, j-column index
    cuComplex *a; // complex nxn matrix
    cuComplex *x; // complex n-vector
    cuComplex *y; // complex n-vector
    cudaMallocManaged(&a,n*n*sizeof(cuComplex)); //unif.memory a
    cudaMallocManaged(&x,n*sizeof(cuComplex)); //unif.memory x
    cudaMallocManaged(&y,n*sizeof(cuComplex)); //unif.memory y
    // define the lower triangle of an nxn Hermitian matrix a in
    // lower mode column by column
    int ind=11; // c:
    for(j=0;j<n;j++){ // 11
        for(i=0;i<n;i++){ // 12,17
            if(i>=j){ // 13,18,22
                a[IDX2C(i,j,n)].x=(float)ind++; // 14,19,23,26
                a[IDX2C(i,j,n)].y=0.0f; // 15,20,24,27,29
            } // 16,21,25,28,30,31
        }
    }
    // print the lower triangle of a row by row
    printf("lower triangle of a:\n");
    for(i=0;i<n;i++){
        for(j=0;j<n;j++){
            if(i>=j)
                printf("%5.0f+%.10f*I",a[IDX2C(i,j,n)].x,
                    a[IDX2C(i,j,n)].y);
        }
    }
}

```

```

    printf("\n");
}
for(i=0;i<n;i++){x[i].x=1.0f;y[i].x=1.0;}
                //x={1,1,1,1,1,1}^T ;y={1,1,1,1,1,1}^T
cublasCreate(&handle);           // initialize CUBLAS context
cuComplex al={1.0f,0.0f};        // al=1
cuComplex bet={1.0f,0.0f};       // bet=1
// Hermitian matrix-vector multiplication:
// y=al*a*x + bet*y
// a - nxn Hermitian matrix; x,y - n-vectors;
// al,bet -scalars

    cublasChemv(handle,CUBLAS_FILL_MODE_LOWER,n,&al,a,n,x,1,&bet,
                y,1);

    cudaDeviceSynchronize();
    printf("y after Chemv:\n");           // print y after Chemv
    for(j=0;j<n;j++){
        printf("%4.0f+%.10f*I",y[j].x,y[j].y);
        printf("\n");
    }
    cudaFree(a);                        // free memory
    cudaFree(x);                        // free memory
    cudaFree(y);                        // free memory
    cublasDestroy(handle);              // destroy CUBLAS context
    return EXIT_SUCCESS;
}
// lower triangle of a:
//   11+0*I
//   12+0*I   17+0*I
//   13+0*I   18+0*I   22+0*I
//   14+0*I   19+0*I   23+0*I   26+0*I
//   15+0*I   20+0*I   24+0*I   27+0*I   29+0*I
//   16+0*I   21+0*I   25+0*I   28+0*I   30+0*I   31+0*I

// y after Chemv:
//   82+0*I
//  108+0*I
//  126+0*I
//  138+0*I
//  146+0*I
//  152+0*I
//
//   [11  12  13  14  15  16] [1]      [1]      [ 82]
//   [12  17  18  19  20  21] [1]      [1]      [108]
//  1*[13  18  22  23  24  25]*[1] + 1*[1] = [126]
//   [14  19  23  26  27  28] [1]      [1]      [138]
//   [15  20  24  27  29  30] [1]      [1]      [146]
//   [16  21  25  28  30  31] [1]      [1]      [152]

```

2.3.35 cublasChbmv - Hermitian banded matrix-vector multiplication

This function performs the Hermitian banded matrix-vector multiplication

$$y = \alpha Ax + \beta y,$$

where A is an $n \times n$ complex Hermitian banded matrix with k subdiagonals and superdiagonals, x, y are complex n -vectors and α, β are complex scalars. A can be stored in lower (CUBLAS_FILL_MODE_LOWER) or upper (CUBLAS_FILL_MODE_UPPER) mode. If A is stored in lower mode, then the main diagonal is stored in row 0, the first subdiagonal in row 1, the second subdiagonal in row 2, etc.

```
// nvcc 030Chbmv.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"
#include "cuComplex.h"
#define n 6 // number of rows and columns of a
#define k 1 // number of subdiagonals and superdiagonals
int main(void){
    cudaError_t cudaStat; // cudaMalloc status
    cublasStatus_t stat; // CUBLAS functions status
    cublasHandle_t handle; // CUBLAS context
    int i,j; // i-row index, j-column index
                // lower triangle of a:
    cuComplex *a;//nxn matrix a on the host //11
    cuComplex *x; //n-vector x on the host //17,12
    cuComplex *y; //n-vector y on the host // 18,13
    a=(cuComplex*)malloc(n*n*sizeof(*a)); // 19,14
    // host memory alloc for a // 20,15
    x=(cuComplex*)malloc(n*sizeof(*x)); // 21,16
    // host memory alloc for x
    y=(cuComplex*)malloc(n*sizeof(*y));
    // host memory alloc for y
    // main diagonal and subdiagonals of a in rows
    int ind=11;
    for(i=0;i<n;i++) a[i*n].x=(float)ind++;
    // main diagonal: 11,12,13,14,15,16 in row 0
    for(i=0;i<n-1;i++) a[i*n+1].x=(float)ind++;
    // first subdiagonal: 17,18,19,20,21 in row 1
    for(i=0;i<n;i++){x[i].x=1.0f;y[i].x=0.0f;}
                //x={1,1,1,1,1,1}^T; y={0,0,0,0,0,0}^T
    // on the device
    cuComplex* d_a; // d_a - a on the device
    cuComplex* d_x; // d_x - x on the device
    cuComplex* d_y; // d_y - y on the device
    cudaStat=cudaMalloc((void**)&d_a,n*n*sizeof(*a)); //device
```

```

// memory alloc for a
cudaStat=cudaMalloc((void**)&d_x,n*sizeof(*x)); //device
// memory alloc for x
cudaStat=cudaMalloc((void**)&d_y,n*sizeof(*y)); //device
// memory alloc for y
stat = cublasCreate(&handle); // initialize CUBLAS context
// copy matrix and vectors from host to device
stat = cublasSetMatrix(n,n,sizeof(*a),a,n,d_a,n); // a-> d_a
stat = cublasSetVector(n,sizeof(*x),x,1,d_x,1); // x-> d_x
stat = cublasSetVector(n,sizeof(*y),y,1,d_y,1); // y-> d_y
cuComplex al={1.0f,0.0f}; // al=1
cuComplex bet={1.0f,0.0f}; // bet=1
// Hermitian banded matrix-vector multiplication:
// d_y = al*d_a*d_x + bet*d_y
// d_a - complex Hermitian banded nxn matrix;
// d_x,d_y -complex n-vectors; al,bet - complex scalars

stat=cublasChbm(v(handle,CUBLAS_FILL_MODE_LOWER,n,k,&al,d_a,n,
d_x,1,&bet,d_y,1);
stat=cublasGetVector(n,sizeof(*y),d_y,1,y,1); //copy d_y->y
printf("y after Chbm(v:\n"); // print y after Chbm(v
for(j=0;j<n;j++){
printf("%3.0f+%1.0f*I",y[j].x,y[j].y);
printf("\n");
}
cudaFree(d_a); // free device memory
cudaFree(d_x); // free device memory
cudaFree(d_y); // free device memory
cublasDestroy(handle); // destroy CUBLAS context
free(a); // free host memory
free(x); // free host memory
free(y); // free host memory
return EXIT_SUCCESS;
}
// y after Chbm(v:
// 28+0*I // [11 17 ] [1] [28]
// 47+0*I // [17 12 18 ] [1] [47]
// 50+0*I // [ 18 13 19 ] [1] = [50]
// 53+0*I // [ 19 14 20 ] [1] [53]
// 56+0*I // [ 20 15 21] [1] [56]
// 37+0*I // [ 21 16] [1] [37]

```

2.3.36 cublasChbm(v - unified memory version

```

// nvcc 030Chbm(v.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#include "cuComplex.h"
#define n 6 // number of rows and columns of a
#define k 1 // number of subdiagonals and superdiagonals

```

```

int main(void){
    cublasHandle_t handle;                                // CUBLAS context
    int i,j;                                              // i-row index, j-column index
                                                    // lower triangle of a:

    cuComplex *a; //nxn matrix a                        //11
    cuComplex *x; //n-vector x                          //17,12
    cuComplex *y; //n-vector y                          // 18,13
                                                    //      19,14
                                                    //      20,15
                                                    //      21,16

    cudaMallocManaged(&a,n*n*sizeof(cuComplex)); //unif.memory a
    cudaMallocManaged(&x,n*sizeof(cuComplex)); //unif.memory x
    cudaMallocManaged(&y,n*sizeof(cuComplex)); //unif.memory y
    // main diagonal and subdiagonals of a in rows
    int ind=11;
    for(i=0;i<n;i++) a[i*n].x=(float)ind++;
    // main diagonal: 11,12,13,14,15,16 in row 0
    for(i=0;i<n-1;i++) a[i*n+1].x=(float)ind++;
    // first subdiagonal: 17,18,19,20,21 in row 1
    for(i=0;i<n;i++){x[i].x=1.0f;y[i].x=0.0f;}
                                //x={1,1,1,1,1,1}^T; y={0,0,0,0,0,0}^T
    cublasCreate(&handle); // initialize CUBLAS context
    cuComplex al={1.0f,0.0f}; // al=1
    cuComplex bet={1.0f,0.0f}; // bet=1
    // Hermitian banded matrix-vector multiplication:
    // y = al*a*x + bet*y
    // a - complex Hermitian banded nxn matrix;
    // x,y -complex n-vectors; al,bet - complex scalars

    cublasChbm(v(handle,CUBLAS_FILL_MODE_LOWER,n,k,&al,a,n,x,1,
                                                    &bet,y,1);

    cudaDeviceSynchronize();
    printf("y after Chbm(v:\n"); // print y after Chbm(v
    for(j=0;j<n;j++){
        printf("%3.0f+%1.0f*I",y[j].x,y[j].y);
        printf("\n");
    }
    cudaFree(a); // free memory
    cudaFree(x); // free memory
    cudaFree(y); // free memory
    cublasDestroy(handle); // destroy CUBLAS context
    return EXIT_SUCCESS;
}

// y after Chbm(v:
// 28+0*I // [11 17 ] [1] [28]
// 47+0*I // [17 12 18 ] [1] [47]
// 50+0*I // [ 18 13 19 ] [1] = [50]
// 53+0*I // [ 19 14 20 ] [1] [53]
// 56+0*I // [ 20 15 21] [1] [56]
// 37+0*I // [ 21 16] [1] [37]

```

2.3.37 cublasChpmv - Hermitian packed matrix-vector multiplication

This function performs the Hermitian packed matrix-vector multiplication

$$y = \alpha Ax + \beta y,$$

where A is an $n \times n$ complex Hermitian packed matrix, x, y are complex n -vectors and α, β are complex scalars. A can be stored in lower (CUBLAS_FILL_MODE_LOWER) or upper (CUBLAS_FILL_MODE_UPPER) mode. If A is stored in lower mode, then the elements of the lower triangular part of A are packed together column by column without gaps.

```
// nvcc 031Chpmv.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"
#define n 6 // number of rows and columns of a
int main(void){
    cudaError_t cudaStat; // cudaMalloc status
    cublasStatus_t stat; // CUBLAS functions status
    cublasHandle_t handle; // CUBLAS context
    int i,j,l,m; // i-row index, j-column index
    // data preparation on the host
    cuComplex *a; // lower triangle of a complex
    // nxn matrix on the host
    cuComplex *x; // complex n-vector x on the host
    cuComplex *y; // complex n-vector y on the host
    a=(cuComplex*)malloc(n*(n+1)/2*sizeof(cuComplex)); // host
    // memory alloc for a
    x=(cuComplex*)malloc(n*sizeof(cuComplex)); // host memory
    // alloc for x
    y=(cuComplex*)malloc(n*sizeof(cuComplex)); // host memory
    // alloc for y
    // define the lower triangle of a Hermitian matrix a:
    // in packed format, column by column // 11
    // without gaps // 12,17
    for(i=0;i<n*(n+1)/2;i++) // 13,18,22
        a[i].x=(float)(11+i); // 14,19,23,26
    // print the upp.triang.of a row by row // 15,20,24,27,29
    printf("upper triangle of a:\n"); // 16,21,25,28,30,31
    l=n;j=0;m=0;
    while(l>0){ // print the upper
        for(i=0;i<m;i++) printf(" "); // triangle of a
        for(i=j;i<j+l;i++) printf("%3.0f+%1.0f*I",a[i].x,a[i].y);
        printf("\n");
        m++;j=j+l;l--;
    }
    for(i=0;i<n;i++){x[i].x=1.0f;y[i].x=0.0f;}
```

```

//x={1,1,1,1,1,1}^T;    y={0,0,0,0,0,0}^T
// on the device
cuComplex* d_a;           // d_a - a on the device
cuComplex* d_x;           // d_x - x on the device
cuComplex* d_y;           // d_y - y on the device
cudaStat=cudaMalloc((void**)&d_a,n*(n+1)/2*sizeof(*a));
                        //device memory alloc for a
cudaStat=cudaMalloc((void**)&d_x,n*sizeof(cuComplex));
                        //device memory alloc for x
cudaStat=cudaMalloc((void**)&d_y,n*sizeof(cuComplex));
                        // device memory alloc for y
stat = cublasCreate(&handle); // initialize CUBLAS context
// copy matrix and vectors from the host to the device
stat = cublasSetVector(n*(n+1)/2, sizeof(*a), a, 1, d_a, 1);
                        //copy a-> d_a
stat = cublasSetVector(n, sizeof(cuComplex), x, 1, d_x, 1);
                        //copy x-> d_x
stat = cublasSetVector(n, sizeof(cuComplex), y, 1, d_y, 1);
                        //copy y-> d_y
cuComplex al={1.0f,0.0f}; // al=1
cuComplex bet={1.0f,0.0f}; // bet=1
// Hermitian packed matrix-vector multiplication:
// d_y = al*d_a*d_x + bet*d_y;  d_a - nxn Hermitian matrix
// in packed format; d_x,d_y - complex n-vectors;
// al,bet - complex scalars

stat=cublasChpmv(handle,CUBLAS_FILL_MODE_LOWER,n,&al,d_a,d_x,1,
                &bet,d_y,1);
stat=cublasGetVector(n, sizeof(cuComplex), d_y, 1, y, 1);
                        // copy d_y->y
printf("y after Chpmv :\n"); // print y after Chpmv
for(j=0;j<n;j++){
    printf("%3.0f+%1.0f*I",y[j].x,y[j].y);
    printf("\n");
}
cudaFree(d_a);           // free device memory
cudaFree(d_x);           // free device memory
cudaFree(d_y);           // free device memory
cublasDestroy(handle);   // destroy CUBLAS context
free(a);                 // free host memory
free(x);                 // free host memory
free(y);                 // free host memory
return EXIT_SUCCESS;
}
// upper triangle of a:
// 11+0*I 12+0*I 13+0*I 14+0*I 15+0*I 16+0*I
//      17+0*I 18+0*I 19+0*I 20+0*I 21+0*I
//      22+0*I 23+0*I 24+0*I 25+0*I
//      26+0*I 27+0*I 28+0*I
//      29+0*I 30+0*I
//      31+0*I

```

```
// y after Chpmv :
// 81+0*I // [11 12 13 14 15 16] [1] [0] [ 81]
// 107+0*I // [12 17 18 19 20 21] [1] [0] [107]
// 125+0*I // 1*[13 18 22 23 24 25]*[1] + 1*[0] = [125]
// 137+0*I // [14 19 23 26 27 28] [1] [0] [137]
// 145+0*I // [15 20 24 27 29 30] [1] [0] [145]
// 151+0*I // [16 21 25 28 30 31] [1] [0] [151]
```

2.3.38 cublasChpmv - unified memory version

```
// nvcc 031Chpmv.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define n 6 // number of rows and columns of a
int main(void){
    cublasHandle_t handle; // CUBLAS context
    int i,j,l,m; // i-row index, j-column index
    // data preparation
    cuComplex *a; // lower triangle of a complex
    // nxn matrix
    cuComplex *x; // complex n-vector
    cuComplex *y; // complex n-vector
    // unified memory for a,x,y
    cudaMallocManaged(&a,n*(n+1)/2*sizeof(cuComplex));
    cudaMallocManaged(&x,n*sizeof(cuComplex));
    cudaMallocManaged(&y,n*sizeof(cuComplex));
    // define the lower triangle of a Hermitian matrix a:
    // in packed format, column by column // 11
    // without gaps // 12,17
    for(i=0;i<n*(n+1)/2;i++) // 13,18,22
        a[i].x=(float)(11+i); // 14,19,23,26
    // print the upp.triang.of a row by row // 15,20,24,27,29
    printf("upper triangle of a:\n"); // 16,21,25,28,30,31
    l=n;j=0;m=0;
    while(l>0){ // print the upper
        for(i=0;i<m;i++) printf(" "); // triangle of a
        for(i=j;i<j+1;i++) printf("%3.0f+%1.0f*I",a[i].x,a[i].y);
        printf("\n");
        m++;j=j+1;l--;
    }
    for(i=0;i<n;i++){x[i].x=1.0f;y[i].x=0.0f;}
    //x={1,1,1,1,1,1}^T; y={0,0,0,0,0,0}^T
    cublasCreate(&handle); // initialize CUBLAS context
    cuComplex al={1.0f,0.0f}; // al=1
    cuComplex bet={1.0f,0.0f}; // bet=1
    // Hermitian packed matrix-vector multiplication:
    // y = al*a*x + bet*y; a - nxn Hermitian matrix
    // in packed format; x,y - complex n-vectors;
    // al,bet - complex scalars
```



```

cublasChpmv(handle,CUBLAS_FILL_MODE_LOWER,n,&al,a,x,1,&bet,y,1);

cudaDeviceSynchronize();
printf("y after Chpmv :\n");          // print y after Chpmv
for(j=0;j<n;j++){
    printf("%3.0f+%.10f*I",y[j].x,y[j].y);
    printf("\n");
}
cudaFree(a);                          // free memory
cudaFree(x);                          // free memory
cudaFree(y);                          // free memory
cublasDestroy(handle);                // destroy CUBLAS context
return EXIT_SUCCESS;
}
// upper triangle of a:
// 11+0*I 12+0*I 13+0*I 14+0*I 15+0*I 16+0*I
//      17+0*I 18+0*I 19+0*I 20+0*I 21+0*I
//      22+0*I 23+0*I 24+0*I 25+0*I
//      26+0*I 27+0*I 28+0*I
//      29+0*I 30+0*I
//      31+0*I

// y after Chpmv:
// 81+0*I // [11 12 13 14 15 16] [1] [0] [ 81]
// 107+0*I // [12 17 18 19 20 21] [1] [0] [107]
// 125+0*I // 1*[13 18 22 23 24 25]*[1] + 1*[0] = [125]
// 137+0*I // [14 19 23 26 27 28] [1] [0] [137]
// 145+0*I // [15 20 24 27 29 30] [1] [0] [145]
// 151+0*I // [16 21 25 28 30 31] [1] [0] [151]

```

2.3.39 cublasCher - Hermitian rank-1 update

This function performs the Hermitian rank-1 update

$$A = \alpha x x^H + A,$$

where A is an $n \times n$ Hermitian complex matrix, x is a complex n -vector and α is a scalar. A is stored in column-major format. A can be stored in lower (CUBLAS_FILL_MODE_LOWER) or upper (CUBLAS_FILL_MODE_UPPER) mode.

```

// nvcc 032cher.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))
#define n 6 // number of rows and columns of a
int main(void){
    cudaError_t cudaStat; // cudaMalloc status
    cublasStatus_t stat; // CUBLAS functions status

```

```

    cublasHandle_t handle;                                // CUBLAS context
    int i,j;                                              // i-row index, j-column index
// data preparation on the host
    cuComplex *a;                                        //nxn complex matrix a on the host
    cuComplex *x;                                        //complex n-vector x on the host
    a=(cuComplex*)malloc(n*n*sizeof(cuComplex)); // host memory
                                                    // alloc for a
    x=(cuComplex*)malloc(n*sizeof(cuComplex)); // host memory
                                                    // alloc for x
// define the lower triangle of an nxn Hermitian matrix a
// column by column
    int ind=11;                                          // a:
    for(j=0;j<n;j++){                                  // 11
        for(i=0;i<n;i++){                              // 12,17
            if(i>=j){                                  // 13,18,22
                a[IDX2C(i,j,n)].x=(float)ind++;        // 14,19,23,26
                a[IDX2C(i,j,n)].y=0.0f;                // 15,20,24,27,29
            }                                           // 16,21,25,28,30,31
        }
    }
// print the lower triangle of a row by row
    printf("lower triangle of a:\n");
    for(i=0;i<n;i++){
        for(j=0;j<n;j++){
            if(i>=j)
                printf("%5.0f+%1.0f*I",a[IDX2C(i,j,n)].x,a[IDX2C(i,j,n)].y);
        }
        printf("\n");
    }
    for(i=0;i<n;i++) x[i].x=1.0f;                      // x={1,1,1,1,1,1}^T

// on the device
    cuComplex* d_a;                                     // d_a - a on the device
    cuComplex* d_x;                                     // d_x - x on the device
    cudaStat=cudaMalloc((void**)&d_a,n*n*sizeof(cuComplex));
                                                    //device memory alloc for a
    cudaStat=cudaMalloc((void**)&d_x,n*sizeof(cuComplex));
                                                    //device memory alloc for x
    stat = cublasCreate(&handle); // initialize CUBLAS context
// copy the matrix and vector from the host to the device
    stat = cublasSetMatrix(n,n,sizeof(*a),a,n,d_a,n); //a -> d_a
    stat = cublasSetVector(n,sizeof(*x),x,1,d_x,1); //x -> d_x
    float al=1.0f;                                     // al=1
// rank-1 update of the Hermitian matrix d_a:
// d_a = al*d_x*d_x^H + d_a
// d_a - nxn Hermitian matrix; d_x - n-vector; al - scalar

    stat=cublasCher(handle,CUBLAS_FILL_MODE_LOWER,n,&al,d_x,1,d_a,n);

    stat=cublasGetMatrix(n,n,sizeof(cuComplex),d_a,n,a,n);
                                                    // copy d_a-> a
// print the lower triangle of updated a

```

```

printf("lower triangle of updated a after Cher:\n");
for(i=0;i<n;i++){
    for(j=0;j<n;j++){
        if(i>=j)
            printf("%5.0f+%1.0f*I",a[IDX2C(i,j,n)].x,a[IDX2C(i,j,n)].y);
        }
    printf("\n");
}
cudaFree(d_a);                // free device memory
cudaFree(d_x);                // free device memory
cublasDestroy(handle);        // destroy CUBLAS context
free(a);                      // free host memory
free(x);                      // free host memory
return EXIT_SUCCESS;
}
// lower triangle of a:
//   11+0*I
//   12+0*I   17+0*I
//   13+0*I   18+0*I   22+0*I
//   14+0*I   19+0*I   23+0*I   26+0*I
//   15+0*I   20+0*I   24+0*I   27+0*I   29+0*I
//   16+0*I   21+0*I   25+0*I   28+0*I   30+0*I   31+0*I

// lower triangle of updated a after Cher:
//   12+0*I
//   13+0*I   18+0*I
//   14+0*I   19+0*I   23+0*I
//   15+0*I   20+0*I   24+0*I   27+0*I
//   16+0*I   21+0*I   25+0*I   28+0*I   30+0*I
//   17+0*I   22+0*I   26+0*I   29+0*I   31+0*I   32+0*I
//           [1]
//           [1]
//           [1]
//   a = 1*[ ]*[1,1,1,1,1,1]+ a
//           [1]
//           [1]
//           [1]

```

2.3.40 cublasCher - unified memory version

```

// nvcc 032cher.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))
#define n 6                // number of rows and columns of a
int main(void){
    cublasHandle_t handle;    // CUBLAS context

```

```

int i,j;                                // i-row index, j-column index
cuComplex *a;                          //nxn complex matrix a
cuComplex *x;                          //complex n-vector x
cudaMallocManaged(&a,n*n*sizeof(cuComplex)); //unif.memory a
cudaMallocManaged(&x,n*sizeof(cuComplex)); //unif.memory x
// define the lower triangle of an nxn Hermitian matrix a
// column by column
int ind=11;                             // a:
for(j=0;j<n;j++){                       // 11
    for(i=0;i<n;i++){                   // 12,17
        if(i>=j){                      // 13,18,22
            a[IDX2C(i,j,n)].x=(float)ind++; // 14,19,23,26
            a[IDX2C(i,j,n)].y=0.0f;      // 15,20,24,27,29
        }                               // 16,21,25,28,30,31
    }
}
// print the lower triangle of a row by row
printf("lower triangle of a:\n");
for(i=0;i<n;i++){
    for(j=0;j<n;j++){
        if(i>=j)
            printf("%5.0f+%1.0f*I",a[IDX2C(i,j,n)].x,a[IDX2C(i,j,n)].y);
    }
    printf("\n");
}
for(i=0;i<n;i++) x[i].x=1.0f;           // x={1,1,1,1,1,1}^T
cublasCreate(&handle);                  // initialize CUBLAS context
float al=1.0f;                          // al=1
// rank-1 update of the Hermitian matrix d_a:
// a = al*x*x^H + a
// a - nxn Hermitian matrix; x - n-vector; al - scalar

cublasCher(handle,CUBLAS_FILL_MODE_LOWER,n,&al,x,1,a,n);

cudaDeviceSynchronize();
// print the lower triangle of updated a
printf("lower triangle of updated a after Cher :\n");
for(i=0;i<n;i++){
    for(j=0;j<n;j++){
        if(i>=j)
            printf("%5.0f+%1.0f*I",a[IDX2C(i,j,n)].x,a[IDX2C(i,j,n)].y);
    }
    printf("\n");
}
cudaFree(a);                           // free memory
cudaFree(x);                           // free memory
cublasDestroy(handle);                  // destroy CUBLAS context
return EXIT_SUCCESS;
}

// lower triangle of a:

```



```

x=(cuComplex*)malloc(n*sizeof(cuComplex)); // host memory
// alloc for x
y=(cuComplex*)malloc(n*sizeof(cuComplex)); // host memory
// alloc for y
// define the lower triangle of an nxn Hermitian matrix a
// column by column
int ind=11; // a:
for(j=0;j<n;j++){ // 11
    for(i=0;i<n;i++){ // 12,17
        if(i>=j){ // 13,18,22
            a[IDX2C(i,j,n)].x=(float)ind++; // 14,19,23,26
            a[IDX2C(i,j,n)].y=0.0f; // 15,20,24,27,29
        } // 16,21,25,28,30,31
    }
}
// print the lower triangle of a row by row
printf("lower triangle of a:\n");
for(i=0;i<n;i++){
    for(j=0;j<n;j++){
        if(i>=j)
            printf("%5.0f+%1.0f*I",a[IDX2C(i,j,n)].x,a[IDX2C(i,j,n)].y);
    }
    printf("\n");
}
for(i=0;i<n;i++){x[i].x=1.0f;y[i].x=2.0;}
//x={1,1,1,1,1,1}^T, y={2,2,2,2,2,2}^T
// on the device
cuComplex* d_a; // d_a - a on the device
cuComplex* d_x; // d_x - x on the device
cuComplex* d_y; // d_y - y on the device
cudaStat=cudaMalloc((void**)&d_a,n*n*sizeof(cuComplex));
//device memory alloc for a
cudaStat=cudaMalloc((void**)&d_x,n*sizeof(cuComplex));
//device memory alloc for x
cudaStat=cudaMalloc((void**)&d_y,n*sizeof(cuComplex));
//device memory alloc for y
stat = cublasCreate(&handle); // initialize CUBLAS context
// copy the matrix and vectors from the host to the device
stat = cublasSetMatrix(n,n,sizeof(*a),a,n,d_a,n); //a -> d_a
stat = cublasSetVector(n,sizeof(*x),x,1,d_x,1); //x -> d_x
stat = cublasSetVector(n,sizeof(*y),y,1,d_y,1); //y -> d_y
cuComplex al={1.0f,0.0f}; // al=1
// rank-2 update of the Hermitian matrix d_a:
// d_a = al*d_x*d_y^H + \bar{al}*d_y*d_x^H + d_a
// d_a - nxn Hermitian matrix; d_x,d_y -n-vectors; al -scalar

stat=cublasCher2(handle,CUBLAS_FILL_MODE_LOWER,n,&al,d_x,1,d_y,
1,d_a,n);
stat=cublasGetMatrix(n,n,sizeof(*a),d_a,n,a,n); //cp d_a->a
// print the lower triangle of updated a
printf("lower triangle of updated a after Cher2:\n");
for(i=0;i<n;i++){

```

```

    for(j=0;j<n;j++){
        if(i>=j)
            printf("%5.0f+%1.0f*I",a[IDX2C(i,j,n)].x,a[IDX2C(i,j,n)].y);
        }
        printf("\n");
    }
    cudaFree(d_a);                // free device memory
    cudaFree(d_x);                // free device memory
    cudaFree(d_y);                // free device memory
    cublasDestroy(handle);        // destroy CUBLAS context
    free(a);                      // free host memory
    free(x);                      // free host memory
    free(y);                      // free host memory
    return EXIT_SUCCESS;
}
// lower triangle of a:
//   11+0*I
//   12+0*I   17+0*I
//   13+0*I   18+0*I   22+0*I
//   14+0*I   19+0*I   23+0*I   26+0*I
//   15+0*I   20+0*I   24+0*I   27+0*I   29+0*I
//   16+0*I   21+0*I   25+0*I   28+0*I   30+0*I   31+0*I

// lower triangle of updated a after Cher2:
//   15+0*I
//   16+0*I   21+0*I
//   17+0*I   22+0*I   26+0*I
//   18+0*I   23+0*I   27+0*I   30+0*I
//   19+0*I   24+0*I   28+0*I   31+0*I   33+0*I
//   20+0*I   25+0*I   29+0*I   32+0*I   34+0*I   35+0*I

//[15 16 17 18 19 20]   [1]                [2]
//[16 21 22 23 24 25]   [1]                [2]
//[17 22 26 27 28 29]   [1]                [2]
//[
                        ]=1*[ ]*[2,2,2,2,2,2]+1*[ ]*[1,1,1,1,1,1])+a
//[18 23 27 30 31 32]   [1]                [2]
//[19 24 28 31 33 34]   [1]                [2]
//[20 25 29 33 34 35]   [1]                [2]

```

2.3.42 cublasCher2 - unified memory version

```

// nvcc 033cher2.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))
#define n 6                // number of rows and columns of a
int main(void){
    cublasHandle_t handle;    // CUBLAS context
    int i,j;                 // i-row index, j-column index
    cuComplex *a;            // nxn complex matrix

```

```

    cuComplex *x;                                     //complex n-vector
    cuComplex *y;                                     //complex n-vector
    cudaMallocManaged(&a,n*n*sizeof(cuComplex)); //unif.memory a
    cudaMallocManaged(&x,n*sizeof(cuComplex));    //unif.memory x
    cudaMallocManaged(&y,n*sizeof(cuComplex));    //unif.memory y
// define the lower triangle of an nxn Hermitian matrix a
// column by column
int ind=11;                                           // a:
for(j=0;j<n;j++){                                   // 11
    for(i=0;i<n;i++){                               // 12,17
        if(i>=j){                                   // 13,18,22
            a[IDX2C(i,j,n)].x=(float)ind++;         // 14,19,23,26
            a[IDX2C(i,j,n)].y=0.0f;                 // 15,20,24,27,29
        }                                           // 16,21,25,28,30,31
    }
}
// print the lower triangle of a row by row
printf("lower triangle of a:\n");
for(i=0;i<n;i++){
    for(j=0;j<n;j++){
        if(i>=j)
            printf("%5.0f+%1.0f*I",a[IDX2C(i,j,n)].x,a[IDX2C(i,j,n)].y);
    }
    printf("\n");
}
for(i=0;i<n;i++){x[i].x=1.0f;y[i].x=2.0f;}
//x={1,1,1,1,1,1}^T, y={2,2,2,2,2,2}^T
cublasCreate(&handle);                             // initialize CUBLAS context
cuComplex al={1.0f,0.0f};                          // al=1
// rank-2 update of the Hermitian matrix a:
// a = al*x*y^H + \bar{al}*y*x^H + a
// a - nxn Hermitian matrix; x,y - n-vectors; al -scalar

    cublasCher2(handle,CUBLAS_FILL_MODE_LOWER,n,&al,x,1,y,1,a,n);

    cudaDeviceSynchronize();
// print the lower triangle of updated a
printf("lower triangle of updated a after Cher2 :\n");
for(i=0;i<n;i++){
    for(j=0;j<n;j++){
        if(i>=j)
            printf("%5.0f+%1.0f*I",a[IDX2C(i,j,n)].x,a[IDX2C(i,j,n)].y);
    }
    printf("\n");
}
cudaFree(a);                                         // free memory
cudaFree(x);                                         // free memory
cudaFree(y);                                         // free memory
cublasDestroy(handle);                             // destroy CUBLAS context
return EXIT_SUCCESS;
}
// lower triangle of a:

```



```

//      11+0*I
//      12+0*I      17+0*I
//      13+0*I      18+0*I      22+0*I
//      14+0*I      19+0*I      23+0*I      26+0*I
//      15+0*I      20+0*I      24+0*I      27+0*I      29+0*I
//      16+0*I      21+0*I      25+0*I      28+0*I      30+0*I      31+0*I

// lower triangle of updated a after Cher2 :
//      15+0*I
//      16+0*I      21+0*I
//      17+0*I      22+0*I      26+0*I
//      18+0*I      23+0*I      27+0*I      30+0*I
//      19+0*I      24+0*I      28+0*I      31+0*I      33+0*I
//      20+0*I      25+0*I      29+0*I      32+0*I      34+0*I      35+0*I

//[15 16 17 18 19 20]      [1]
//[16 21 22 23 24 25]      [1]
//[17 22 26 27 28 29]      [1]
//[
                        ]=1*[
                        ]*[2,2,2,2,2,2]+1*[
                        ]*[1,1,1,1,1,1])+a
//[18 23 27 30 31 32]      [1]
//[19 24 28 31 33 34]      [1]
//[20 25 29 33 34 35]      [1]

```

2.3.43 cublasChpr - packed Hermitian rank-1 update

This function performs the Hermitian rank-1 update

$$A = \alpha x x^H + A,$$

where A is an $n \times n$ complex Hermitian matrix in packed format, x is a complex n -vector and α is a scalar. A can be stored in lower (CUBLAS_FILL_MODE_LOWER) or upper (CUBLAS_FILL_MODE_UPPER) mode. If A is stored in lower mode, then the elements of the lower triangular part of A are packed together column by column without gaps.

```

// nvcc 034chpr.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"
#define n 6 // number of rows and columns of a
int main(void){
    cudaError_t cudaStat; // cudaMalloc status
    cublasStatus_t stat; // CUBLAS functions status
    cublasHandle_t handle; // CUBLAS context
    int i,j,l,m; // i-row index, j-column index
    // data preparation on the host
    cuComplex *a; // lower triangle of a complex
    // nxn matrix a on the host
    cuComplex *x; // complex n-vector x on the host

```

```

a=(cuComplex*)malloc(n*(n+1)/2*sizeof(*a)); // host memory
// alloc for a
x=(cuComplex*)malloc(n*sizeof(cuComplex)); // host memory
// alloc for x

// define the lower triangle of a Hermi- //11
// tian a in packed format column by //12,17
// column without gaps //13,18,22
for(i=0;i<n*(n+1)/2;i++) //14,19,23,26
    a[i].x=(float)(11+i); //15,20,24,27,29
// print upper triangle of a row by row //16,21,25,28,30,31
printf("upper triangle of a:\n");
l=n;j=0;m=0;
while(l>0){ // print the lower
    for(i=0;i<m;i++) printf(" "); // triangle of a
    for(i=j;i<j+1;i++) printf("%3.0f+%1.0f*I",a[i].x,a[i].y);
    printf("\n");
    m++;j=j+1;l--;
}
for(i=0;i<n;i++){x[i].x=1.0f;} //x={1,1,1,1,1,1}^T
// on the device
cuComplex* d_a; // d_a - a on the device
cuComplex* d_x; // d_x - x on the device
cudaStat=cudaMalloc((void**)&d_a,n*(n+1)/2*sizeof(*a));
//device memory alloc for a
cudaStat=cudaMalloc((void**)&d_x,n*sizeof(cuComplex));
//device memory alloc for x
stat = cublasCreate(&handle); // initialize CUBLAS context
// copy the matrix and vector from the host to the device
stat = cublasSetVector(n*(n+1)/2,sizeof(*a),a,1,d_a,1);
// copy a-> d_a
stat = cublasSetVector(n,sizeof(*x),x,1,d_x,1); // x-> d_x
float al=1.0f; // al=1
// rank-1 update of a Hermitian packed complex matrix d_a:
// d_a = al*d_x*d_x^H + d_a; d_a - Hermitian nxn complex
// matrix in packed format; d_x - n-vector; al - scalar

stat=cublasChpr(handle,CUBLAS_FILL_MODE_LOWER,n,&al,d_x,1,d_a);

stat=cublasGetVector(n*(n+1)/2,sizeof(*a),d_a,1,a,1);
// copy d_a-> a
// print the updated upper triangle of a row by row
printf("updated upper triangle of a after Chpr:\n");
l=n;j=0;m=0;
while(l>0){
    for(i=0;i<m;i++) printf(" ");
    for(i=j;i<j+1;i++)
        printf("%3.0f+%1.0f*I",a[i].x,a[i].y);
    printf("\n");
    m++;j=j+1;l--;
}
cudaFree(d_a); // free device memory
cudaFree(d_x); // free device memory

```

```

    cublasDestroy(handle);           // destroy CUBLAS context
    free(a);                         // free host memory
    free(x);                         // free host memory
return EXIT_SUCCESS;
}
// upper triangle of a:
// 11+0*I 12+0*I 13+0*I 14+0*I 15+0*I 16+0*I
//      17+0*I 18+0*I 19+0*I 20+0*I 21+0*I
//      22+0*I 23+0*I 24+0*I 25+0*I
//      26+0*I 27+0*I 28+0*I
//      29+0*I 30+0*I
//      31+0*I

// updated upper triangle of a after Chpr:
// 12+0*I 13+0*I 14+0*I 15+0*I 16+0*I 17+0*I
//      18+0*I 19+0*I 20+0*I 21+0*I 22+0*I
//      23+0*I 24+0*I 25+0*I 26+0*I
//      27+0*I 28+0*I 29+0*I
//      30+0*I 31+0*I
//      32+0*I
//      [1]
//      [1]
//      [1]
//  a = 1*[ ]*[1,1,1,1,1,1]+ a
//      [1]
//      [1]
//      [1]

```

2.3.44 cublasChpr - unified memory version

```

// nvcc 034chpr.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define n 6           // number of rows and columns of a
int main(void){
    cublasHandle_t handle;           // CUBLAS context
    int i,j,l,m;                   // i-row index, j-column index
    cuComplex *a; // lower triangle of a complex nxn matrix a
    cuComplex *x; // complex n-vector x
// unified memory for a,x
    cudaMallocManaged(&a,n*(n+1)/2*sizeof(cuComplex));
    cudaMallocManaged(&x,n*sizeof(cuComplex));
// define the lower triangle of a Hermi- //11
// tian a in packed column by //12,17
// column without gaps //13,18,22
    for(i=0;i<n*(n+1)/2;i++) //14,19,23,26
        a[i].x=(float)(11+i); //15,20,24,27,29
// print upper triangle of a row by row //16,21,25,28,30,31
printf("upper triangle of a:\n");
    l=n;j=0;m=0;

```

```

while(l>0){
    for(i=0;i<m;i++){
        for(j=i;j<m;j++){
            printf("%3.0f+%1.0f*I",a[i].x,a[i].y);
            printf("\n");
            m++;j=j+1;l--;
        }
    }
    for(i=0;i<n;i++){x[i].x=1.0f;} //x={1,1,1,1,1,1}^T
    cublasCreate(&handle); // initialize CUBLAS context
    float al=1.0f; // al=1
    // rank-1 update of a Hermitian packed complex matrix a:
    // a = al*x*x^H + a; a - Hermitian nxn complex
    // matrix in packed format; x - n-vector; al - scalar

    cublasChpr(handle,CUBLAS_FILL_MODE_LOWER,n,&al,x,1,a);

    cudaDeviceSynchronize();
    // print the updated upper triangle of a row by row
    printf("updated upper triangle of a after Chpr:\n");
    l=n;j=0;m=0;
    while(l>0){
        for(i=0;i<m;i++){
            for(j=i;j<m;j++){
                printf("%3.0f+%1.0f*I",a[i].x,a[i].y);
                printf("\n");
                m++;j=j+1;l--;
            }
        }
        cudaFree(a); // free memory
        cudaFree(x); // free memory
        cublasDestroy(handle); // destroy CUBLAS context
    }
    return EXIT_SUCCESS;
}

// upper triangle of a:
// 11+0*I 12+0*I 13+0*I 14+0*I 15+0*I 16+0*I
//      17+0*I 18+0*I 19+0*I 20+0*I 21+0*I
//          22+0*I 23+0*I 24+0*I 25+0*I
//              26+0*I 27+0*I 28+0*I
//                  29+0*I 30+0*I
//                      31+0*I
// updated upper triangle of a after Chpr:
// 12+0*I 13+0*I 14+0*I 15+0*I 16+0*I 17+0*I
//      18+0*I 19+0*I 20+0*I 21+0*I 22+0*I
//          23+0*I 24+0*I 25+0*I 26+0*I
//              27+0*I 28+0*I 29+0*I
//                  30+0*I 31+0*I
//                      32+0*I
//
//      [1]
//      [1]
//      [1]
// a = 1*[ ]*[1,1,1,1,1,1]+ a
//      [1]
//      [1]
//      [1]

```

2.3.45 cublasChpr2 - packed Hermitian rank-2 update

This function performs the Hermitian rank-2 update

$$A = \alpha xy^H + \bar{\alpha}yx^H + A,$$

where A is an $n \times n$ Hermitian complex matrix in packed format, x, y are complex n -vectors and α is a complex scalar. A can be stored in lower (CUBLAS_FILL_MODE_LOWER) or upper (CUBLAS_FILL_MODE_UPPER) mode. If A is stored in lower mode, then the elements of the lower triangular part of A are packed together column by column without gaps.

```
// nvcc 035chpr2.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"
#define n 6 // number of rows and columns of a
int main(void){
    cudaError_t cudaStat; // cudaMalloc status
    cublasStatus_t stat; // CUBLAS functions status
    cublasHandle_t handle; // CUBLAS context
    int i,j,l,m; // i-row index, j-column index
    // data preparation on the host
    cuComplex *a; // lower triangle of a complex
    // nxn matrix a on the host
    cuComplex *x; // complex n-vector x on the host
    cuComplex *y; // complex n-vector y on the host
    a=(cuComplex*)malloc(n*(n+1)/2*sizeof(*a)); // host memory
    // alloc for a
    x=(cuComplex*)malloc(n*sizeof(cuComplex)); // host memory
    // alloc for x
    y=(cuComplex*)malloc(n*sizeof(cuComplex)); // host memory
    // alloc for y
    // define the lower triangle of a Hermitian //11
    // tian a in packed format column by //12,17
    // column without gaps //13,18,22
    for(i=0;i<n*(n+1)/2;i++) //14,19,23,26
        a[i].x=(float)(11+i); //15,20,24,27,29
    // print upper triangle of a row by row //16,21,25,28,30,31
    printf("upper triangle of a:\n");
    l=n;j=0;m=0;
    while(l>0){ // print the upper
        for(i=0;i<m;i++) printf(" "); // triangle of a
        for(i=j;i<j+1;i++) printf("%3.0f+%1.0f*I",a[i].x,a[i].y);
        printf("\n");
        m++;j=j+1;l--;
    }
    for(i=0;i<n;i++){x[i].x=1.0f;y[i].x=2.0f;
        //x={1,1,1,1,1,1}^T; y={2,2,2,2,2,2}^T
    }
    // on the device
```

```

cuComplex* d_a; // d_a - a on the device
cuComplex* d_x; // d_x - x on the device
cuComplex* d_y; // d_y - y on the device
cudaStat=cudaMalloc((void**)&d_a,n*(n+1)/2*sizeof(*a));
//device memory alloc for a
cudaStat=cudaMalloc((void**)&d_x,n*sizeof(cuComplex));
//device memory alloc for x
cudaStat=cudaMalloc((void**)&d_y,n*sizeof(cuComplex));
//device memory alloc for y
stat = cublasCreate(&handle); // initialize CUBLAS context
// copy matrix and vectors from the host to the device
stat = cublasSetVector(n*(n+1)/2, sizeof(*a), a, 1, d_a, 1);
// copy a-> d_a
stat = cublasSetVector(n, sizeof(*x), x, 1, d_x, 1); // x-> d_x
stat = cublasSetVector(n, sizeof(*y), y, 1, d_y, 1); // y-> d_y
cuComplex al={1.0f,0.0f}; //al=1
// rank-2 update of a Hermitian matrix d_a:
// d_a = al*d_x*d_y^H + \bar{al}*d_y*d_x^H + d_a; d_a -Herm.
// nxn matrix in packed format; d_x,d_y - n-vectors; al -scal.

stat=cublasChpr2(handle,CUBLAS_FILL_MODE_LOWER,n,&al,d_x,1,
                d_y,1,d_a);

stat=cublasGetVector(n*(n+1)/2, sizeof(cuComplex), d_a, 1, a, 1);
// copy d_a -> a
// print the updated upper triangle of a row by row
printf("updated upper triangle of a after Chpr2:\n");
l=n;j=0;m=0;
while(l>0){
    for(i=0;i<m;i++) printf(" ");
    for(i=j;i<j+1;i++) printf("%3.0f+%1.0f*I",a[i].x,a[i].y);
    printf("\n");
    m++;j=j+1;l--;
}
cudaFree(d_a); // free device memory
cudaFree(d_x); // free device memory
cudaFree(d_y); // free device memory
cublasDestroy(handle); // destroy CUBLAS context
free(a); // free host memory
free(x); // free host memory
free(y); // free host memory
return EXIT_SUCCESS;
}
// upper triangle of a:
// 11+0*I 12+0*I 13+0*I 14+0*I 15+0*I 16+0*I
//      17+0*I 18+0*I 19+0*I 20+0*I 21+0*I
//      22+0*I 23+0*I 24+0*I 25+0*I
//      26+0*I 27+0*I 28+0*I
//      29+0*I 30+0*I
//      31+0*I

```

```

// updated upper triangle of a after Chpr2:
// 15+0*I 16+0*I 17+0*I 18+0*I 19+0*I 20+0*I
//      21+0*I 22+0*I 23+0*I 24+0*I 25+0*I
//      26+0*I 27+0*I 28+0*I 29+0*I
//      30+0*I 31+0*I 32+0*I
//      33+0*I 34+0*I
//      35+0*I

//[15 16 17 18 19 20] [1] [2]
//[16 21 22 23 24 25] [1] [2]
//[17 22 26 27 28 29] [1] [2]
//[      ]=1*[ ]*[2,2,2,2,2,2]+1*[ ]*[1,1,1,1,1,1])+a
//[18 23 27 30 31 32] [1] [2]
//[19 24 28 31 33 34] [1] [2]
//[20 25 29 33 34 35] [1] [2]

```

2.3.46 cublasChpr2 - unified memory version

```

// nvcc 035chpr2.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define n 6 // number of rows and columns of a
int main(void){
    cublasHandle_t handle; // CUBLAS context
    int i,j,l,m; // i-row index, j-column index
    cuComplex *a; // lower triangle of a complex
                  // nxn matrix
    cuComplex *x; // complex n-vector
    cuComplex *y; // complex n-vector
    // unified memory for a,x,y
    cudaMallocManaged(&a,n*(n+1)/2*sizeof(cuComplex));
    cudaMallocManaged(&x,n*sizeof(cuComplex));
    cudaMallocManaged(&y,n*sizeof(cuComplex));
    // define the lower triangle of a Hermi- //11
    // tian a in packed format column by //12,17
    // column without gaps //13,18,22
    for(i=0;i<n*(n+1)/2;i++) //14,19,23,26
        a[i].x=(float)(11+i); //15,20,24,27,29
    // print upper triangle of a row by row //16,21,25,28,30,31
    printf("upper triangle of a:\n");
    l=n;j=0;m=0;
    while(l>0){ // print the upper
        for(i=0;i<m;i++) printf(" "); // triangle of a
        for(i=j;i<j+l;i++) printf("%3.0f+%1.0f*I",a[i].x,a[i].y);
        printf("\n");
        m++;j=j+l;l--;
    }
    for(i=0;i<n;i++){x[i].x=1.0f;y[i].x=2.0f;}
    //x={1,1,1,1,1,1}^T; y={2,2,2,2,2,2}^T
    cublasCreate(&handle); // initialize CUBLAS context
}

```

```

    cuComplex al={1.0f,0.0f};                                     //al=1
// rank-2 update of a Hermitian matrix a:
// a = al*x*y^H + \bar{al}*y*x^H + a; a -Hermitian
// nxn matrix in packed format; x,y - n-vectors; al -scalar

    cublasChpr2(handle,CUBLAS_FILL_MODE_LOWER,n,&al,x,1,y,1,a);

    cudaDeviceSynchronize();
// print the updated upper triangle of a row by row
printf("updated upper triangle of a after Chpr2:\n");
l=n;j=0;m=0;
while(l>0){
    for(i=0;i<m;i++) printf("          ");
    for(i=j;i<j+1;i++) printf("%3.0f+%1.0f*I",a[i].x,a[i].y);
    printf("\n");
    m++;j=j+1;l--;
}
cudaFree(a);                                                     // free memory
cudaFree(x);                                                     // free memory
cudaFree(y);                                                     // free memory
cublasDestroy(handle);                                           // destroy CUBLAS context
return EXIT_SUCCESS;
}
// upper triangle of a:
// 11+0*I 12+0*I 13+0*I 14+0*I 15+0*I 16+0*I
//          17+0*I 18+0*I 19+0*I 20+0*I 21+0*I
//                22+0*I 23+0*I 24+0*I 25+0*I
//                        26+0*I 27+0*I 28+0*I
//                                29+0*I 30+0*I
//                                    31+0*I

// updated upper triangle of a after Chpr2:
// 15+0*I 16+0*I 17+0*I 18+0*I 19+0*I 20+0*I
//          21+0*I 22+0*I 23+0*I 24+0*I 25+0*I
//                26+0*I 27+0*I 28+0*I 29+0*I
//                        30+0*I 31+0*I 32+0*I
//                                33+0*I 34+0*I
//                                    35+0*I

// [15 16 17 18 19 20]    [1]                                [2]
// [16 21 22 23 24 25]    [1]                                [2]
// [17 22 26 27 28 29]    [1]                                [2]
// [          ]=1*[ ]*[2,2,2,2,2,2]+1*[ ]*[1,1,1,1,1,1])+a
// [18 23 27 30 31 32]    [1]                                [2]
// [19 24 28 31 33 34]    [1]                                [2]
// [20 25 29 33 34 35]    [1]                                [2]

```


2.4 CUBLAS Level-3. Matrix-matrix operations

2.4.1 cublasSgemm - matrix-matrix multiplication

This function performs the matrix-matrix multiplication

$$C = \alpha op(A)op(B) + \beta C,$$

where A, B are matrices in column-major format and α, β are scalars. The value of $op(A)$ can be equal to A (CUBLAS_OP_N case), A^T (transposition) in CUBLAS_OP_T case, or A^H (conjugate transposition) in CUBLAS_OP_C case and similarly for $op(B)$.

```
// nvcc 036sgemm.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))
#define m 6 // a - mxk matrix
#define n 4 // b - kxn matrix
#define k 5 // c - mxn matrix
int main(void){
    cudaError_t cudaStat; // cudaMalloc status
    cublasStatus_t stat; // CUBLAS functions status
    cublasHandle_t handle; // CUBLAS context
    int i,j; // i-row index,j-column index
    float* a; // mxk matrix a on the host
    float* b; // kxn matrix b on the host
    float* c; // mxn matrix c on the host
    a=(float*)malloc(m*k*sizeof(float)); // host memory for a
    b=(float*)malloc(k*n*sizeof(float)); // host memory for b
    c=(float*)malloc(m*n*sizeof(float)); // host memory for c
    // define an mxk matrix a column by column
    int ind=11; // a:
    for(j=0;j<k;j++){ // 11,17,23,29,35
        for(i=0;i<m;i++){ // 12,18,24,30,36
            a[IDX2C(i,j,m)]=(float)ind++; // 13,19,25,31,37
        } // 14,20,26,32,38
    } // 15,21,27,33,39
    // 16,22,28,34,40

    // print a row by row
    printf("a:\n");
    for(i=0;i<m;i++){
        for(j=0;j<k;j++){
            printf("%5.0f",a[IDX2C(i,j,m)]);
        }
        printf("\n");
    }
    // define a kxn matrix b column by column
```

```

ind=11;                                // b:
for(j=0;j<n;j++){                       // 11,16,21,26
    for(i=0;i<k;i++){                   // 12,17,22,27
        b[IDX2C(i,j,k)]=(float)ind++;  // 13,18,23,28
    }                                   // 14,19,24,29
}                                       // 15,20,25,30
// print b row by row
printf("b:\n");
for(i=0;i<k;i++){
    for(j=0;j<n;j++){
        printf("%5.0f",b[IDX2C(i,j,k)]);
    }
    printf("\n");
}
// define an mxn matrix c column by column
ind=11;                                // c:
for(j=0;j<n;j++){                       // 11,17,23,29
    for(i=0;i<m;i++){                   // 12,18,24,30
        c[IDX2C(i,j,m)]=(float)ind++;  // 13,19,25,31
    }                                   // 14,20,26,32
}                                       // 15,21,27,33
// 16,22,28,34
// print c row by row
printf("c:\n");
for(i=0;i<m;i++){
    for(j=0;j<n;j++){
        printf("%5.0f",c[IDX2C(i,j,m)]);
    }
    printf("\n");
}
// on the device
float* d_a;                             // d_a - a on the device
float* d_b;                             // d_b - b on the device
float* d_c;                             // d_c - c on the device
cudaStat=cudaMalloc((void**)&d_a,m*k*sizeof(*a)); //device
// memory alloc for a
cudaStat=cudaMalloc((void**)&d_b,k*n*sizeof(*b)); //device
// memory alloc for b
cudaStat=cudaMalloc((void**)&d_c,m*n*sizeof(*c)); //device
// memory alloc for c
stat = cublasCreate(&handle); // initialize CUBLAS context
// copy matrices from the host to the device
stat = cublasSetMatrix(m,k,sizeof(*a),a,m,d_a,m); //a -> d_a
stat = cublasSetMatrix(k,n,sizeof(*b),b,k,d_b,k); //b -> d_b
stat = cublasSetMatrix(m,n,sizeof(*c),c,m,d_c,m); //c -> d_c
float al=1.0f;                          // al=1
float bet=1.0f;                         // bet=1
// matrix-matrix multiplication: d_c = al*d_a*d_b + bet*d_c
// d_a -mxk matrix, d_b -kxn matrix, d_c -mxn matrix;
// al,bet -scalars

```

```

stat=cublasSgemm(handle,CUBLAS_OP_N,CUBLAS_OP_N,m,n,k,&al,d_a,
                m,d_b,k,&bet,d_c,m);

stat=cublasGetMatrix(m,n,sizeof(*c),d_c,m,c,m); //cp d_c->c
printf("c after Sgemm :\n");
for(i=0;i<m;i++){
    for(j=0;j<n;j++){
        printf("%7.0f",c[IDX2C(i,j,m)]); //print c after Sgemm
    }
    printf("\n");
}
cudaFree(d_a); // free device memory
cudaFree(d_b); // free device memory
cudaFree(d_c); // free device memory
cublasDestroy(handle); // destroy CUBLAS context
free(a); // free host memory
free(b); // free host memory
free(c); // free host memory
return EXIT_SUCCESS;
}

// a:
// 11 17 23 29 35
// 12 18 24 30 36
// 13 19 25 31 37
// 14 20 26 32 38
// 15 21 27 33 39
// 16 22 28 34 40
// b:
// 11 16 21 26
// 12 17 22 27
// 13 18 23 28
// 14 19 24 29
// 15 20 25 30
// c:
// 11 17 23 29
// 12 18 24 30
// 13 19 25 31
// 14 20 26 32
// 15 21 27 33
// 16 22 28 34

// c after Sgemm :
// 1566 2147 2728 3309
// 1632 2238 2844 3450
// 1698 2329 2960 3591 // c=al*a*b+bet*c
// 1764 2420 3076 3732
// 1830 2511 3192 3873
// 1896 2602 3308 4014

```

2.4.2 cublasSgemm - unified memory version

```

// nvcc 036sgemm.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))
#define m 6 // a - mxk matrix
#define n 4 // b - kxn matrix
#define k 5 // c - mxn matrix
int main(void){
    cublasHandle_t handle; // CUBLAS context
    int i,j; // i-row index,j-column index
    float* a; // mxk matrix
    float* b; // kxn matrix
    float* c; // mxn matrix
    // unified memory for a,b,c
    cudaMallocManaged(&a,m*k*sizeof(cuComplex));
    cudaMallocManaged(&b,k*n*sizeof(cuComplex));
    cudaMallocManaged(&c,m*n*sizeof(cuComplex));
    // define an mxk matrix a column by column
    int ind=11; // a:
    for(j=0;j<k;j++){ // 11,17,23,29,35
        for(i=0;i<m;i++){ // 12,18,24,30,36
            a[IDX2C(i,j,m)]=(float)ind++; // 13,19,25,31,37
        } // 14,20,26,32,38
    } // 15,21,27,33,39
    // 16,22,28,34,40

    // print a row by row
    printf("a:\n");
    for(i=0;i<m;i++){
        for(j=0;j<k;j++){
            printf("%5.0f",a[IDX2C(i,j,m)]);
        }
        printf("\n");
    }
    // define a kxn matrix b column by column
    ind=11; // b:
    for(j=0;j<n;j++){ // 11,16,21,26
        for(i=0;i<k;i++){ // 12,17,22,27
            b[IDX2C(i,j,k)]=(float)ind++; // 13,18,23,28
        } // 14,19,24,29
    } // 15,20,25,30

    // print b row by row
    printf("b:\n");
    for(i=0;i<k;i++){
        for(j=0;j<n;j++){
            printf("%5.0f",b[IDX2C(i,j,k)]);
        }
        printf("\n");
    }
    // define an mxn matrix c column by column
    ind=11; // c:

```

```

    for(j=0;j<n;j++){
        for(i=0;i<m;i++){
            c[IDX2C(i,j,m)]=(float)ind++;
        }
    }
// 11,17,23,29
// 12,18,24,30
// 13,19,25,31
// 14,20,26,32
// 15,21,27,33
// 16,22,28,34

// print c row by row
printf("c:\n");
for(i=0;i<m;i++){
    for(j=0;j<n;j++){
        printf("%5.0f",c[IDX2C(i,j,m)]);
    }
    printf("\n");
}

cublasCreate(&handle);           // initialize CUBLAS context
float al=1.0f;                  // al=1
float bet=1.0f;                 // bet=1
// matrix-matrix multiplication: c = al*a*b + bet*c
// a -mxk matrix, b -kxn matrix, c -mxn matrix;
// al,bet -scalars

cublasSgemv(handle,CUBLAS_OP_N,CUBLAS_OP_N,m,n,k,&al,a,m,b,k,
              &bet,c,m);

cudaDeviceSynchronize();
printf("c after Sgemv :\n");
for(i=0;i<m;i++){
    for(j=0;j<n;j++){
        printf("%7.0f",c[IDX2C(i,j,m)]); //print c after Sgemv
    }
    printf("\n");
}

cudaFree(a);                   // free memory
cudaFree(b);                   // free memory
cudaFree(c);                   // free memory
cublasDestroy(handle);         // destroy CUBLAS context
return EXIT_SUCCESS;
}

// a:
// 11 17 23 29 35
// 12 18 24 30 36
// 13 19 25 31 37
// 14 20 26 32 38
// 15 21 27 33 39
// 16 22 28 34 40
// b:
// 11 16 21 26
// 12 17 22 27
// 13 18 23 28
// 14 19 24 29
// 15 20 25 30

```

```

// c:
//   11   17   23   29
//   12   18   24   30
//   13   19   25   31
//   14   20   26   32
//   15   21   27   33
//   16   22   28   34

// c after Sgemm :
//   1566   2147   2728   3309
//   1632   2238   2844   3450
//   1698   2329   2960   3591      // c=al*a*b+bet*c
//   1764   2420   3076   3732
//   1830   2511   3192   3873
//   1896   2602   3308   4014

```

2.4.3 cublasSsymm - symmetric matrix-matrix multiplication

This function performs the left or right symmetric matrix-matrix multiplications

$$C = \alpha AB + \beta C \quad \text{in CUBLAS_SIDE_LEFT case,}$$

$$C = \alpha BA + \beta C \quad \text{in CUBLAS_SIDE_RIGHT case.}$$

The symmetric matrix A has dimension $m \times m$ in the first case and $n \times n$ in the second one. The general matrices B, C have dimensions $m \times n$ and α, β are scalars. The matrix A can be stored in lower (CUBLAS_FILL_MODE_LOWER) or upper (CUBLAS_FILL_MODE_UPPER) mode.

```

// nvcc 037ssymm.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))
#define m 6 // a - mxm matrix
#define n 4 // b,c - mxn matrices
int main(void){
    cudaError_t cudaStat; // cudaMalloc status
    cublasStatus_t stat; // CUBLAS functions status
    cublasHandle_t handle; // CUBLAS context
    int i,j; // i-row ind., j-column ind.
    float* a; // mxm matrix a on the host
    float* b; // mxn matrix b on the host
    float* c; // mxn matrix c on the host
    a=(float*)malloc(m*m*sizeof(float)); // host memory for a
    b=(float*)malloc(m*n*sizeof(float)); // host memory for b
    c=(float*)malloc(m*n*sizeof(float)); // host memory for c
    // define the lower triangle of an mxm symmetric matrix a in

```

```

// lower mode column by column
int ind=11;
for(j=0;j<m;j++){
    for(i=0;i<m;i++){
        if(i>=j){
            a[IDX2C(i,j,m)]=(float)ind++;
        }
    }
}
// print the lower triangle of a row by row
printf("lower triangle of a:\n");
for(i=0;i<m;i++){
    for(j=0;j<m;j++){
        if(i>=j)
            printf("%5.0f",a[IDX2C(i,j,m)]);
    }
    printf("\n");
}
// define mxn matrices b,c column by column
ind=11;
for(j=0;j<n;j++){
    for(i=0;i<m;i++){
        b[IDX2C(i,j,m)]=(float)ind;
        c[IDX2C(i,j,m)]=(float)ind;
        ind++;
    }
}
// print b(=c) row by row
printf("b(=c):\n");
for(i=0;i<m;i++){
    for(j=0;j<n;j++){
        printf("%5.0f",b[IDX2C(i,j,m)]);
    }
    printf("\n");
}
// on the device
float* d_a;
float* d_b;
float* d_c;
cudaStat=cudaMalloc((void**)&d_a,m*m*sizeof(*a)); //device
cudaStat=cudaMalloc((void**)&d_b,m*n*sizeof(*b)); //device
cudaStat=cudaMalloc((void**)&d_c,m*n*sizeof(*c)); //device
stat = cublasCreate(&handle); // initialize CUBLAS context
// copy matrices from the host to the device
stat = cublasSetMatrix(m,m,sizeof(*a),a,m,d_a,m); //a -> d_a
stat = cublasSetMatrix(m,n,sizeof(*b),b,m,d_b,m); //b -> d_b
stat = cublasSetMatrix(m,n,sizeof(*c),c,m,d_c,m); //c -> d_c
float al=1.0f;
float bet=1.0f;

```

```

// symmetric matrix-matrix multiplication:
// d_c = al*d_a*d_b + bet*d_c; d_a - mxm symmetric matrix;
// d_b,d_c - mxn general matrices; al,bet - scalars

    stat=cublasSsymm(handle,CUBLAS_SIDE_LEFT,CUBLAS_FILL_MODE_LOWER,
                    m,n,&al,d_a,m,d_b,m,&bet,d_c,m);
    stat=cublasGetMatrix(m,n,sizeof(*c),d_c,m,c,m); //d_c -> c
    printf("c after Ssymm :\n");                    //print c after Ssymm
    for(i=0;i<m;i++){
        for(j=0;j<n;j++){
            printf("%7.0f",c[IDX2C(i,j,m)]);
        }
        printf("\n");
    }
    cudaFree(d_a);                                // free device memory
    cudaFree(d_b);                                // free device memory
    cudaFree(d_c);                                // free device memory
    cublasDestroy(handle);                        // destroy CUBLAS context
    free(a);                                       // free host memory
    free(b);                                       // free host memory
    free(c);                                       // free host memory
    return EXIT_SUCCESS;
}

// lower triangle of a:
//   11
//   12   17
//   13   18   22
//   14   19   23   26
//   15   20   24   27   29
//   16   21   25   28   30   31
// b(=c):
//   11   17   23   29
//   12   18   24   30
//   13   19   25   31
//   14   20   26   32
//   15   21   27   33
//   16   22   28   34
// c after Ssymm :
//   1122   1614   2106   2598
//   1484   2132   2780   3428
//   1740   2496   3252   4008           // c=al*a*b+bet*c
//   1912   2740   3568   4396
//   2025   2901   3777   4653
//   2107   3019   3931   4843

```

2.4.4 cublasSsymm - unified memory version

```

// nvcc 037ssymm.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))

```



```

#define m 6                                // a - mxm matrix
#define n 4                                // b,c - mxn matrices
int main(void){
    cublasHandle_t handle;                  // CUBLAS context
    int i,j;                               // i-row ind., j-column ind.
    float* a;                              // mxm matrix a
    float* b;                              // mxn matrix b
    float* c;                              // mxn matrix c
    // unified memory for a,b,c
    cudaMallocManaged(&a,m*m*sizeof(cuComplex));
    cudaMallocManaged(&b,m*n*sizeof(cuComplex));
    cudaMallocManaged(&c,m*n*sizeof(cuComplex));
    // define the lower triangle of an mxm symmetric matrix a in
    // lower mode column by column
    int ind=11;                             // a:
    for(j=0;j<m;j++){                       // 11
        for(i=0;i<m;i++){                   // 12,17
            if(i>=j){                       // 13,18,22
                a[IDX2C(i,j,m)]=(float)ind++; // 14,19,23,26
            }                               // 15,20,24,27,29
        }                                   // 16,21,25,28,30,31
    }
    // print the lower triangle of a row by row
    printf("lower triangle of a:\n");
    for(i=0;i<m;i++){
        for(j=0;j<m;j++){
            if(i>=j)
                printf("%5.0f",a[IDX2C(i,j,m)]);
        }
        printf("\n");
    }
    // define mxn matrices b,c column by column
    ind=11;                                 // b,c:
    for(j=0;j<n;j++){                       // 11,17,23,29
        for(i=0;i<m;i++){                   // 12,18,24,30
            b[IDX2C(i,j,m)]=(float)ind;      // 13,19,25,31
            c[IDX2C(i,j,m)]=(float)ind;      // 14,20,26,32
            ind++;                           // 15,21,27,33
        }                                   // 16,22,28,34
    }
    // print b(=c) row by row
    printf("b(=c):\n");
    for(i=0;i<m;i++){
        for(j=0;j<n;j++){
            printf("%5.0f",b[IDX2C(i,j,m)]);
        }
        printf("\n");
    }
    cublasCreate(&handle);                  // initialize CUBLAS context
    float al=1.0f;                          // al=1
    float bet=1.0f;                         // bet=1
    // symmetric matrix-matrix multiplication:

```

```

// c = al*a*b + bet*c; a - mxm symmetric matrix;
// b,c - mxn general matrices; al,bet - scalars

cublasSsymm(handle,CUBLAS_SIDE_LEFT,CUBLAS_FILL_MODE_LOWER,
            m,n,&al,a,m,b,m,&bet,c,m);

cudaDeviceSynchronize();
printf("c after Ssymm :\n");           //print c after Ssymm
for(i=0;i<m;i++){
    for(j=0;j<n;j++){
        printf("%7.0f",c[IDX2C(i,j,m)]);
    }
    printf("\n");
}
cudaFree(a);                          // free memory
cudaFree(b);                          // free memory
cudaFree(c);                          // free memory
cublasDestroy(handle);                 // destroy CUBLAS context
return EXIT_SUCCESS;
}

// lower triangle of a:
//   11
//   12   17
//   13   18   22
//   14   19   23   26
//   15   20   24   27   29
//   16   21   25   28   30   31
// b(=c):
//   11   17   23   29
//   12   18   24   30
//   13   19   25   31
//   14   20   26   32
//   15   21   27   33
//   16   22   28   34

// c after Ssymm :
//   1122   1614   2106   2598
//   1484   2132   2780   3428
//   1740   2496   3252   4008           // c=al*a*b+bet*c
//   1912   2740   3568   4396
//   2025   2901   3777   4653
//   2107   3019   3931   4843

```

2.4.5 cublasSsyrrk - symmetric rank-k update

This function performs the symmetric rank-k update

$$C = \alpha \text{op}(A)\text{op}(A)^T + \beta C,$$

where $\text{op}(A)$ is an $n \times k$ matrix, C is a symmetric $n \times n$ matrix stored in lower (CUBLAS_FILL_MODE_LOWER) or upper (CUBLAS_FILL_MODE_UPPER) mode and

α, β are scalars. The value of $op(A)$ can be equal to A in CUBLAS_OP_N case or A^T (transposition) in CUBLAS_OP_T case.

```
// nvcc 038ssyrk.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))
#define n 6 // a - nxk matrix
#define k 4 // c - nxn matrix
int main(void){
    cudaError_t cudaStat; // cudaMalloc status
    cublasStatus_t stat; // CUBLAS functions status
    cublasHandle_t handle; // CUBLAS context
    int i,j; // i-row index, j-column index
    float* a; // nxk matrix a on the host
    float* c; // nxn matrix c on the host
    a=(float*)malloc(n*k*sizeof(float)); // host memory for a
    c=(float*)malloc(n*n*sizeof(float)); // host memory for c
    // define the lower triangle of an nxn symmetric matrix c
    // column by column
    int ind=11; // c:
    for(j=0;j<n;j++){ // 11
        for(i=0;i<n;i++){ // 12,17
            if(i>=j){ // 13,18,22
                c[IDX2C(i,j,n)]=(float)ind++; // 14,19,23,26
            } // 15,20,24,27,29
        } // 16,21,25,28,30,31
    }
    // print the lower triangle of c row by row
    printf("lower triangle of c:\n");
    for(i=0;i<n;i++){
        for(j=0;j<n;j++){
            if(i>=j)
                printf("%5.0f",c[IDX2C(i,j,n)]);
        }
        printf("\n");
    }
    // define an nxk matrix a column by column
    ind=11; // a:
    for(j=0;j<k;j++){ // 11,17,23,29
        for(i=0;i<n;i++){ // 12,18,24,30
            a[IDX2C(i,j,n)]=(float)ind; // 13,19,25,31
            ind++; // 14,20,26,32
        } // 15,21,27,33
    } // 16,22,28,34
    printf("a:\n");
    for(i=0;i<n;i++){
        for(j=0;j<k;j++){
            printf("%5.0f",a[IDX2C(i,j,n)]); // print a row by row
        }
    }
}
```

```

    }
    printf("\n");
}

// on the device
float* d_a; // d_a - a on the device
float* d_c; // d_c - c on the device
cudaStat=cudaMalloc((void**)&d_a,n*k*sizeof(*a)); //device
// memory alloc for a
cudaStat=cudaMalloc((void**)&d_c,n*n*sizeof(*c)); //device
// memory alloc for c
stat = cublasCreate(&handle); // initialize CUBLAS context
// copy matrices from the host to the device
stat = cublasSetMatrix(n,k,sizeof(*a),a,n,d_a,n); //a -> d_a
stat = cublasSetMatrix(n,n,sizeof(*c),c,n,d_c,n); //c -> d_c
float al=1.0f; // al=1
float bet=1.0f; //bet=1
// symmetric rank-k update: d_c = al*d_a*d_a^T + bet*d_c;
// d_c - symmetric nxn matrix, d_a - general nxk matrix;
// al,bet - scalars

stat=cublasSsyrrk(handle,CUBLAS_FILL_MODE_LOWER,CUBLAS_OP_N,
                  n,k,&al,d_a,n,&bet,d_c,n);

stat=cublasGetMatrix(n,n,sizeof(*c),d_c,n,c,n); // d_c -> c
printf("lower triangle of updated c after Ssyrrk :\n");
for(i=0;i<n;i++){
    for(j=0;j<n;j++){
        if(i>=j) //print the lower triangle
            printf("%7.0f",c[IDX2C(i,j,n)]); //of c after Ssyrrk
    }
    printf("\n");
}
cudaFree(d_a); // free device memory
cudaFree(d_c); // free device memory
cublasDestroy(handle); // destroy CUBLAS context
free(a); // free host memory
free(c); // free host memory
return EXIT_SUCCESS;
}

// lower triangle of c:
// 11
// 12 17
// 13 18 22
// 14 19 23 26
// 15 20 24 27 29
// 16 21 25 28 30 31

// a:
```

```

//   11   17   23   29
//   12   18   24   30
//   13   19   25   31
//   14   20   26   32
//   15   21   27   33
//   16   22   28   34

// lower triangle of updated c after Ssyrrk: c=al*a*a^T+bet*c
//   1791
//   1872   1961
//   1953   2046   2138
//   2034   2131   2227   2322
//   2115   2216   2316   2415   2513
//   2196   2301   2405   2508   2610   2711

```

2.4.6 cublasSsyrrk - unified memory version

```

// nvcc 038ssyrrk.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))
#define n 6 // a - nxk matrix
#define k 4 // c - nxn matrix
int main(void){
    cublasHandle_t handle; // CUBLAS context
    int i,j; // i-row index, j-column index
    float* a; // nxk matrix
    float* c; // nxn matrix

    // unified memory for a,c
    cudaMallocManaged(&a,n*k*sizeof(cuComplex));
    cudaMallocManaged(&c,n*n*sizeof(cuComplex));
    // define the lower triangle of an nxn symmetric matrix c
    // column by column
    int ind=11; // c:
    for(j=0;j<n;j++){ // 11
        for(i=0;i<n;i++){ // 12,17
            if(i>=j){ // 13,18,22
                c[IDX2C(i,j,n)]=(float)ind++; // 14,19,23,26
            } // 15,20,24,27,29
        } // 16,21,25,28,30,31
    }

    // print the lower triangle of c row by row
    printf("lower triangle of c:\n");
    for(i=0;i<n;i++){
        for(j=0;j<n;j++){
            if(i>=j)
                printf("%5.0f",c[IDX2C(i,j,n)]);
        }
        printf("\n");
    }
}

```

```

// define an nxk matrix a column by column
ind=11;
for(j=0;j<k;j++){
    for(i=0;i<n;i++){
        a[IDX2C(i,j,n)]=(float)ind;
        ind++;
    }
}
printf("a:\n");
for(i=0;i<n;i++){
    for(j=0;j<k;j++){
        printf("%5.0f",a[IDX2C(i,j,n)]); // print a row by row
    }
    printf("\n");
}
cublasCreate(&handle); // initialize CUBLAS context
float al=1.0f; // al=1
float bet=1.0f; //bet=1
// symmetric rank-k update: c = al*a*a^T + bet*c;
// c - symmetric nxn matrix, a - general nxk matrix;
// al,bet - scalars

cublasSsyrc(handle,CUBLAS_FILL_MODE_LOWER, CUBLAS_OP_N,
            n,k,&al,a,n,&bet,c,n);

cudaDeviceSynchronize();
printf("lower triangle of updated c after Ssyrc :\n");
for(i=0;i<n;i++){
    for(j=0;j<n;j++){
        if(i>=j) //print the lower triangle
            printf("%7.0f",c[IDX2C(i,j,n)]); //of c after Ssyrc
    }
    printf("\n");
}
cudaFree(a); // free memory
cudaFree(c); // free memory
cublasDestroy(handle); // destroy CUBLAS context
return EXIT_SUCCESS;
}

// lower triangle of c:
// 11
// 12 17
// 13 18 22
// 14 19 23 26
// 15 20 24 27 29
// 16 21 25 28 30 31
// a:
// 11 17 23 29
// 12 18 24 30
// 13 19 25 31
// 14 20 26 32
// 15 21 27 33
// 16 22 28 34

```

```
// lower triangle of updated c after Ssyrrk: c=al*a*a^T+bet*c
//   1791
//   1872   1961
//   1953   2046   2138
//   2034   2131   2227   2322
//   2115   2216   2316   2415   2513
//   2196   2301   2405   2508   2610   2711
```

2.4.7 cublasSsyrr2k - symmetric rank-2k update

This function performs the symmetric rank-2k update

$$C = \alpha(op(A)op(B)^T + op(B)op(A)^T) + \beta C,$$

where $op(A)$, $op(B)$ are $n \times k$ matrices, C is a symmetric $n \times n$ matrix stored in lower (CUBLAS_FILL_MODE_LOWER) or upper (CUBLAS_FILL_MODE_UPPER) mode and α, β are scalars. The value of $op(A)$ can be equal to A in CUBLAS_OP_N case or A^T (transposition) in CUBLAS_OP_T case and similarly for $op(B)$.

```
// nvcc 039ssyrrk.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))
#define n 6 // c - nxn matrix
#define k 4 // a,b - nxk matrices
int main(void){
    cudaError_t cudaStat; // cudaMalloc status
    cublasStatus_t stat; // CUBLAS functions status
    cublasHandle_t handle; // CUBLAS context
    int i,j; // i-row index, j-col. index
    float* a; // nxk matrix on the host
    float* b; // nxk matrix on the host
    float* c; // nxn matrix on the host
    a=(float*)malloc(n*k*sizeof(float)); // host memory for a
    b=(float*)malloc(n*k*sizeof(float)); // host memory for b
    c=(float*)malloc(n*n*sizeof(float)); // host memory for c
    // define the lower triangle of an nxn symmetric matrix c in
    // lower mode column by column
    int ind=11; // c:
    for(j=0;j<n;j++){ // 11
        for(i=0;i<n;i++){ // 12,17
            if(i>=j){ // 13,18,22
                c[IDX2C(i,j,n)]=(float)ind++; // 14,19,23,26,
            } // 15,20,24,27,29
        } // 16,21,25,28,30,31
    }
}
```

```

// print the lower triangle of c row by row
printf("lower triangle of c:\n");
for(i=0;i<n;i++){
    if(i>=j)
        for(j=0;j<n;j++){
            printf("%5.0f",c[IDX2C(i,j,n)]);
        }
    printf("\n");
}
// define nxk matrices a,b column by column
ind=11;
for(j=0;j<k;j++){
    for(i=0;i<n;i++){
        a[IDX2C(i,j,n)]=(float)ind;
        b[IDX2C(i,j,n)]=(float)ind;
        ind++;
    }
}
printf("a(=b):\n");
for(i=0;i<n;i++){
    for(j=0;j<k;j++){
        printf("%5.0f",a[IDX2C(i,j,n)]); // print a row by row
    }
    printf("\n");
}
// on the device
float* d_a; // d_a - a on the device
float* d_b; // d_b - b on the device
float* d_c; // d_c - c on the device
cudaStat=cudaMalloc((void**)&d_a,n*k*sizeof(*a)); //device
// memory alloc for a
cudaStat=cudaMalloc((void**)&d_b,n*k*sizeof(*b)); //device
// memory alloc for b
cudaStat=cudaMalloc((void**)&d_c,n*n*sizeof(*c)); //device
// memory alloc for c
stat = cublasCreate(&handle); // initialize CUBLAS context
// copy matrices from the host to the device
stat = cublasSetMatrix(n,k,sizeof(*a),a,n,d_a,n); //a -> d_a
stat = cublasSetMatrix(n,k,sizeof(*b),b,n,d_b,n); //b -> d_b
stat = cublasSetMatrix(n,n,sizeof(*c),c,n,d_c,n); //c -> d_c
float al=1.0f; // al=1
float bet=1.0f; //bet=1
// symmetric rank-2k update:
// d_c=al*(d_a*d_b^T+d_b*d_a^T)+bet*d_c
// d_c - symmetric nxn matrix, d_a,d_b - general nxk matrices
// al,bet - scalars

stat=cublasSsyr2k(handle,CUBLAS_FILL_MODE_LOWER,CUBLAS_OP_N,
                  n,k,&al,d_a,n,d_b,n,&bet,d_c,n);
stat=cublasGetMatrix(n,n,sizeof(*c),d_c,n,c,n); //d_c -> c
printf("lower triangle of updated c after Ssyr2k :\n");
for(i=0;i<n;i++){

```



```

    for(j=0;j<n;j++){
        if(i>=j)                //print the lower triangle
            printf("%7.0f",c[IDX2C(i,j,n)]);    //of c after Ssyr2k
    }
    printf("\n");
}
cudaFree(d_a);                // free device memory
cudaFree(d_b);                // free device memory
cudaFree(d_c);                // free device memory
cublasDestroy(handle);        // destroy CUBLAS context
free(a);                      // free host memory
free(b);                      // free host memory
free(c);                      // free host memory
return EXIT_SUCCESS;
}
// lower triangle of c:
//   11
//   12   17
//   13   18   22
//   14   19   23   26
//   15   20   24   27   29
//   16   21   25   28   30   31
// a(=b):
//   11   17   23   29
//   12   18   24   30
//   13   19   25   31
//   14   20   26   32
//   15   21   27   33
//   16   22   28   34

// lower triangle of updated c after Ssyr2k :
//   3571
//   3732   3905
//   3893   4074   4254
//   4054   4243   4431   4618
//   4215   4412   4608   4803   4997
//   4376   4581   4785   4988   5190   5391

// c = al(a*b^T + b*a^T) + bet*c

```

2.4.8 cublasSsyr2k - unified memory version

```

// nvcc 039ssyr2k.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))
#define n 6                // c - nxn matrix
#define k 4                // a,b - nxk matrices
int main(void){
    cublasHandle_t handle;    // CUBLAS context

```

```

int i,j;                                // i-row index, j-col. index
float* a;                                // nxk matrix
float* b;                                // nxk matrix
float* c;                                // nxn matrix
// unified memory for a,b,c
cudaMallocManaged((void**)&a,n*k*sizeof(float));
cudaMallocManaged((void**)&b,n*k*sizeof(float));
cudaMallocManaged((void**)&c,n*n*sizeof(float));
// define the lower triangle of an nxn symmetric matrix c in
// lower mode column by column
int ind=11;                               // c:
for(j=0;j<n;j++){                         // 11
    for(i=0;i<n;i++){                     // 12,17
        if(i>=j){                         // 13,18,22
            c[IDX2C(i,j,n)]=(float)ind++; // 14,19,23,26,
        }                                 // 15,20,24,27,29
    }                                     // 16,21,25,28,30,31
}
// print the lower triangle of c row by row
printf("lower triangle of c:\n");
for(i=0;i<n;i++){
    for(j=0;j<n;j++){
        if(i>=j)
            printf("%5.0f",c[IDX2C(i,j,n)]);
    }
    printf("\n");
}
// define nxk matrices a,b column by column
ind=11;
for(j=0;j<k;j++){
    for(i=0;i<n;i++){
        a[IDX2C(i,j,n)]=(float)ind;
        b[IDX2C(i,j,n)]=(float)ind;
        ind++;
    }
}
printf("a(=b):\n");
for(i=0;i<n;i++){
    for(j=0;j<k;j++){
        printf("%5.0f",a[IDX2C(i,j,n)]); // print a row by row
    }
    printf("\n");
}
cublasCreate(&handle);                    // initialize CUBLAS context
float al=1.0f;                            // al=1
float bet=1.0f;                           // bet=1
// symmetric rank-2k update: c=al*(a*b^T+b*a^T)+bet*c
// c - symmetric nxn matrix, a,b - general nxk matrices
// al,bet - scalars

cublasSsyr2k(handle,CUBLAS_FILL_MODE_LOWER,CUBLAS_OP_N,n,k,
              &al,a,n,b,n,&bet,c,n);

```

```

    cudaDeviceSynchronize();
    printf("lower triangle of updated c after Ssyr2k :\n");
    for(i=0;i<n;i++){
        for(j=0;j<n;j++){
            if(i>=j)                //print the lower triangle
                printf("%7.0f",c[IDX2C(i,j,n)]);    //of c after Ssyr2k
        }
        printf("\n");
    }
    cudaFree(a);                    // free memory
    cudaFree(b);                    // free memory
    cudaFree(c);                    // free memory
    cublasDestroy(handle);          // destroy CUBLAS context
    return 0;
}
// lower triangle of c:
//   11
//   12  17
//   13  18  22
//   14  19  23  26
//   15  20  24  27  29
//   16  21  25  28  30  31
// a(=b):
//   11  17  23  29
//   12  18  24  30
//   13  19  25  31
//   14  20  26  32
//   15  21  27  33
//   16  22  28  34
// lower triangle of updated c after Ssyr2k :
//   3571
//   3732  3905
//   3893  4074  4254
//   4054  4243  4431  4618
//   4215  4412  4608  4803  4997
//   4376  4581  4785  4988  5190  5391
// c = al(a*b^T + b*a^T) + bet*c

```

2.4.9 cublasStrmm - triangular matrix-matrix multiplication

This function performs the left or right triangular matrix-matrix multiplications

$$\begin{aligned}
 C &= \alpha \operatorname{op}(A) B && \text{in CUBLAS_SIDE_LEFT case,} \\
 C &= \alpha B \operatorname{op}(A) && \text{in CUBLAS_SIDE_RIGHT case,}
 \end{aligned}$$

where A is a triangular matrix, C, B are $m \times n$ matrices and α is a scalar. The value of $\operatorname{op}(A)$ can be equal to A in CUBLAS_OP_N case, A^T (transposition) in CUBLAS_OP_T case or A^H (conjugate transposition) in CUBLAS_OP_C

case. A has dimension $m \times m$ in the first case and $n \times n$ in the second case. A can be stored in lower (CUBLAS_FILL_MODE_LOWER) or upper (CUBLAS_FILL_MODE_UPPER) mode. If the diagonal of the matrix A has non-unit elements, then the parameter CUBLAS_DIAG_NON_UNIT should be used (in the opposite case - CUBLAS_DIAG_UNIT).

```
// nvcc 040strmm.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))
#define m 6 // a - mxm matrix
#define n 5 // b,c - mxn matrices
int main(void){
    cudaError_t cudaStat; // cudaMalloc status
    cublasStatus_t stat; // CUBLAS functions status
    cublasHandle_t handle; // CUBLAS context
    int i,j; // i-row index, j-col. index
    float* a; // mxm matrix a on the host
    float* b; // mxn matrix b on the host
    float* c; // mxn matrix c on the host
    a=(float*)malloc(m*m*sizeof(float)); // host memory for a
    b=(float*)malloc(m*n*sizeof(float)); // host memory for b
    c=(float*)malloc(m*n*sizeof(float)); // host memory for c
    // define the lower triangle of an mxm triangular matrix a in
    // lower mode column by column
    int ind=11; // a:
    for(j=0;j<m;j++){ // 11
        for(i=0;i<m;i++){ // 12,17
            if(i>=j){ // 13,18,22
                a[IDX2C(i,j,m)]=(float)ind++; // 14,19,23,26
            } // 15,20,24,27,29
        } // 16,21,25,28,30,31
    }
    // print the lower triangle of a row by row
    printf("lower triangle of a:\n");
    for(i=0;i<m;i++){
        for(j=0;j<m;j++){
            if(i>=j)
                printf("%5.0f",a[IDX2C(i,j,m)]);
        }
        printf("\n");
    }
    // define an mxn matrix b column by column
    ind=11; // b:
    for(j=0;j<n;j++){ // 11,17,23,29,35
        for(i=0;i<m;i++){ // 12,18,24,30,36
            b[IDX2C(i,j,m)]=(float)ind++; // 13,19,25,31,37
        } // 14,20,26,32,38
    } // 15,21,27,33,39
    } // 16,22,28,34,40
```

```

printf("b:\n");
for(i=0;i<m;i++){
    for(j=0;j<n;j++){
        printf("%5.0f",b[IDX2C(i,j,m)]); // print b row by row
    }
    printf("\n");
}
// on the device
float* d_a; // d_a - a on the device
float* d_b; // d_b - b on the device
float* d_c; // d_c - c on the device
cudaStat=cudaMalloc((void**)&d_a,m*m*sizeof(*a)); //device
// memory alloc for a
cudaStat=cudaMalloc((void**)&d_b,m*n*sizeof(*b)); //device
// memory alloc for b
cudaStat=cudaMalloc((void**)&d_c,m*n*sizeof(*c)); //device
// memory alloc for c
stat = cublasCreate(&handle); // initialize CUBLAS context
// copy matrices from the host to the device
stat = cublasSetMatrix(m,m,sizeof(*a),a,m,d_a,m); //a -> d_a
stat = cublasSetMatrix(m,n,sizeof(*b),b,m,d_b,m); //b -> d_b
float al=1.0f;
// triangular matrix-matrix multiplication: d_c = al*d_a*d_b;
// d_a - mxm triangular matrix in lower mode,
// d_b,d_c -mxn general matrices; al- scalar

stat=cublasStrmm(handle,CUBLAS_SIDE_LEFT,CUBLAS_FILL_MODE_LOWER,
    CUBLAS_OP_N,CUBLAS_DIAG_NON_UNIT,m,n,&al,d_a,m,d_b,m,d_c,m);
stat=cublasGetMatrix(m,n,sizeof(*c),d_c,m,c,m); //d_c -> c
printf("c after Strmm :\n");
for(i=0;i<m;i++){
    for(j=0;j<n;j++){
        printf("%7.0f",c[IDX2C(i,j,m)]); //print c after Strmm
    }
    printf("\n");
}
cudaFree(d_a); // free device memory
cudaFree(d_b); // free device memory
cudaFree(d_c); // free device memory
cublasDestroy(handle); // destroy CUBLAS context
free(a); // free host memory
free(b); // free host memory
free(c); // free host memory
return EXIT_SUCCESS;
}
// lower triangle of a:
// 11
// 12 17
// 13 18 22
// 14 19 23 26
// 15 20 24 27 29
// 16 21 25 28 30 31

```

```

// b:
//   11   17   23   29   35
//   12   18   24   30   36
//   13   19   25   31   37
//   14   20   26   32   38
//   15   21   27   33   39
//   16   22   28   34   40
// c after Strmm :
//   121   187   253   319   385
//   336   510   684   858   1032
//   645   963  1281  1599  1917   // c = al*a*b
//  1045  1537  2029  2521  3013
//  1530  2220  2910  3600  4290
//  2091  2997  3903  4809  5715

```

2.4.10 cublasStrmm - unified memory version

```

// nvcc 040strmm.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))
#define m 6 // a - mxm matrix
#define n 5 // b,c - mxn matrices
int main(void){
    cublasHandle_t handle; // CUBLAS context
    int i,j; // i-row index, j-col. index
    float* a; // mxm matrix
    float* b; // mxn matrix
    float* c; // mxn matrix
    // unified memory for a,b,c
    cudaMallocManaged(&a,m*m*sizeof(float));
    cudaMallocManaged(&b,m*n*sizeof(float));
    cudaMallocManaged(&c,m*n*sizeof(float));
    // define the lower triangle of an mxm triangular matrix a in
    // lower mode column by column
    int ind=11; // a:
    for(j=0;j<m;j++){ // 11
        for(i=0;i<m;i++){ // 12,17
            if(i>=j){ // 13,18,22
                a[IDX2C(i,j,m)]=(float)ind++; // 14,19,23,26
            } // 15,20,24,27,29
        } // 16,21,25,28,30,31
    }
    // print the lower triangle of a row by row
    printf("lower triangle of a:\n");
    for(i=0;i<m;i++){
        for(j=0;j<m;j++){
            if(i>=j)
                printf("%5.0f",a[IDX2C(i,j,m)]);
        }
        printf("\n");
    }
}

```

```

    }
    // define an mxn matrix b column by column
    ind=11;
    for(j=0;j<n;j++){
        for(i=0;i<m;i++){
            b[IDX2C(i,j,m)]=(float)ind++;
        }
    }

    printf("b:\n");
    for(i=0;i<m;i++){
        for(j=0;j<n;j++){
            printf("%5.0f",b[IDX2C(i,j,m)]); // print b row by row
        }
        printf("\n");
    }
    cublasCreate(&handle); // initialize CUBLAS context
    float al=1.0f;
    // triangular matrix-matrix multiplication: c = al*a*b;
    // a - mxm triangular matrix in lower mode,
    // b,c -mxn general matrices; al- scalar

    cublasStrmm(handle,CUBLAS_SIDE_LEFT,CUBLAS_FILL_MODE_LOWER,
        CUBLAS_OP_N,CUBLAS_DIAG_NON_UNIT,m,n,&al,a,m,b,m,c,m);
    cudaDeviceSynchronize();
    printf("c after Strmm :\n");
    for(i=0;i<m;i++){
        for(j=0;j<n;j++){
            printf("%7.0f",c[IDX2C(i,j,m)]); //print c after Strmm
        }
        printf("\n");
    }
    cudaFree(a); // free memory
    cudaFree(b); // free memory
    cudaFree(c); // free memory
    cublasDestroy(handle); // destroy CUBLAS context
    return EXIT_SUCCESS;
}

// lower triangle of a:
//   11
//   12  17
//   13  18  22
//   14  19  23  26
//   15  20  24  27  29
//   16  21  25  28  30  31
// b:
//   11  17  23  29  35
//   12  18  24  30  36
//   13  19  25  31  37
//   14  20  26  32  38
//   15  21  27  33  39
//   16  22  28  34  40

```

```
// c after Strmm :
//   121   187   253   319   385
//   336   510   684   858  1032
//   645   963  1281  1599  1917 // c = al*a*b
//  1045  1537  2029  2521  3013
//  1530  2220  2910  3600  4290
//  2091  2997  3903  4809  5715
```

2.4.11 cublasStrsm - solving the triangular linear system

This function solves the triangular system

$$\begin{aligned} op(A) X &= \alpha B && \text{in CUBLAS_SIDE_LEFT case,} \\ X op(A) &= \alpha B && \text{in CUBLAS_SIDE_RIGHT case,} \end{aligned}$$

where A is a triangular matrix, X, B are $m \times n$ matrices and α is a scalar. The value of $op(A)$ can be equal to A in CUBLAS_OP_N case, A^T (transposition) in CUBLAS_OP_T case or A^H (conjugate transposition) in CUBLAS_OP_C case. A has dimension $m \times m$ in the first case and $n \times n$ in the second and third case. A can be stored in lower (CUBLAS_FILL_MODE_LOWER) or upper (CUBLAS_FILL_MODE_UPPER) mode. If the diagonal of the matrix A has non-unit elements, then the parameter CUBLAS_DIAG_NON_UNIT should be used (in the opposite case - CUBLAS_DIAG_UNIT).

```
// nvcc 041strsm.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))
#define m 6 // a - mxm matrix
#define n 5 // b,x - mxn matrices
int main(void){
    cudaError_t cudaStat; // cudaMalloc status
    cublasStatus_t stat; // CUBLAS functions status
    cublasHandle_t handle; // CUBLAS context
    int i,j; // i-row index, j-col. index
    float* a; // mxm matrix a on the host
    float* b; // mxn matrix b on the host
    a=(float*)malloc(m*m*sizeof(float)); // host memory for a
    b=(float*)malloc(m*n*sizeof(float)); // host memory for b
    // define the lower triangle of an mxm triangular matrix a in
    // lower mode column by column
    int ind=11; // a:
    for(j=0;j<m;j++){ // 11
        for(i=0;i<m;i++){ // 12,17
            if(i>=j){ // 13,18,22
                a[IDX2C(i,j,m)]=(float)ind++; // 14,19,23,26
            } // 15,20,24,27,29
        } // 16,21,25,28,30,31
    }
```



```

    }
    // print the lower triangle of a row by row
    printf("lower triangle of a:\n");
    for(i=0;i<m;i++){
        for(j=0;j<m;j++){
            if(i>=j)
                printf("%5.0f",a[IDX2C(i,j,m)]);
        }
        printf("\n");
    }
    // define an mxn matrix b column by column
    ind=11;
    for(j=0;j<n;j++){
        for(i=0;i<m;i++){
            b[IDX2C(i,j,m)]=(float)ind;
            ind++;
        }
    }
    printf("b:\n");
    for(i=0;i<m;i++){
        for(j=0;j<n;j++){
            printf("%5.0f",b[IDX2C(i,j,m)]); // print b row by row
        }
        printf("\n");
    }
    // on the device
    float* d_a; // d_a - a on the device
    float* d_b; // d_b - b on the device
    cudaStat=cudaMalloc((void**)&d_a,m*m*sizeof(*a)); //device
    // memory alloc for a
    cudaStat=cudaMalloc((void**)&d_b,m*n*sizeof(*b)); //device
    // memory alloc for b
    stat = cublasCreate(&handle); // initialize CUBLAS context
    // copy matrices from the host to the device
    stat = cublasSetMatrix(m,m,sizeof(*a),a,m,d_a,m); //a -> d_a
    stat = cublasSetMatrix(m,n,sizeof(*b),b,m,d_b,m); //b -> d_b
    float al=1.0f; // al=1
    // solve d_a*x=al*d_b; the solution x overwrites rhs d_b;
    // d_a - mxm triangular matrix in lower mode;
    // d_b,x - mxn general matrices; al - scalar

    stat=cublasStrsm(handle,CUBLAS_SIDE_LEFT,CUBLAS_FILL_MODE_LOWER,
        CUBLAS_OP_N,CUBLAS_DIAG_NON_UNIT,m,n,&al,d_a,m,d_b,m);

    stat=cublasGetMatrix(m,n,sizeof(*b),d_b,m,b,m); // d_b -> b
    printf("solution x from Strsm :\n");
    for(i=0;i<m;i++){
        for(j=0;j<n;j++){
            printf("%11.5f",b[IDX2C(i,j,m)]); //print b after Strsm
        }
        printf("\n");
    }
}

```

```

    cudaFree(d_a);                // free device memory
    cudaFree(d_b);                // free device memory
    cublasDestroy(handle);        // destroy CUBLAS context
    free(a);                      // free host memory
    free(b);                      // free host memory
    return EXIT_SUCCESS;
}
// lower triangle of a:
//   11
//   12   17
//   13   18   22
//   14   19   23   26
//   15   20   24   27   29
//   16   21   25   28   30   31
// b:
//   11   17   23   29   35
//   12   18   24   30   36
//   13   19   25   31   37
//   14   20   26   32   38
//   15   21   27   33   39
//   16   22   28   34   40

// solution x from Strsm : a*x=b
//   1.00000   1.54545   2.09091   2.63636   3.18182
//   0.00000   -0.03209   -0.06417   -0.09626   -0.12834
//   0.00000   -0.02333   -0.04667   -0.07000   -0.09334
//   0.00000   -0.01885   -0.03769   -0.05654   -0.07539
//   0.00000   -0.01625   -0.03250   -0.04874   -0.06499
//   0.00000   -0.01468   -0.02935   -0.04403   -0.05870

```

2.4.12 cublasStrsm - unified memory version

```

// nvcc 041strsm.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))
#define m 6 // a - mxm matrix
#define n 5 // b,x - mxn matrices
int main(void){
    cublasHandle_t handle; // CUBLAS context
    int i,j; // i-row index, j-col. index
    float* a; // mxm matrix
    float* b; // mxn matrix
    // unified memory for a,b
    cudaMallocManaged(&a,m*m*sizeof(float));
    cudaMallocManaged(&b,m*n*sizeof(float));
    // define the lower triangle of an mxm triangular matrix a in
    // lower mode column by column
    int ind=11; // a:

```

```

    for(j=0;j<m;j++){
        for(i=0;i<m;i++){
            if(i>=j){
                a[IDX2C(i,j,m)]=(float)ind++;
            }
        }
    }
// print the lower triangle of a row by row
printf("lower triangle of a:\n");
for(i=0;i<m;i++){
    for(j=0;j<m;j++){
        if(i>=j)
            printf("%5.0f",a[IDX2C(i,j,m)]);
    }
    printf("\n");
}
// define an mxn matrix b column by column
ind=11;
for(j=0;j<n;j++){
    for(i=0;i<m;i++){
        b[IDX2C(i,j,m)]=(float)ind;
        ind++;
    }
}
printf("b:\n");
for(i=0;i<m;i++){
    for(j=0;j<n;j++){
        printf("%5.0f",b[IDX2C(i,j,m)]); // print b row by row
    }
    printf("\n");
}
cublasCreate(&handle); // initialize CUBLAS context
float al=1.0f; // al=1
// solve a*x=al*b; the solution x overwrites rhs b;
// a - mxm triangular matrix in lower mode;
// b,x - mxn general matrices; al - scalar

cublasStrsm(handle,CUBLAS_SIDE_LEFT,CUBLAS_FILL_MODE_LOWER,
            CUBLAS_OP_N,CUBLAS_DIAG_NON_UNIT,m,n,&al,a,m,b,m);

cudaDeviceSynchronize();
printf("solution x from Strsm :\n");
for(i=0;i<m;i++){
    for(j=0;j<n;j++){
        printf("%11.5f",b[IDX2C(i,j,m)]); //print b after Strsm
    }
    printf("\n");
}
cudaFree(a); // free memory
cudaFree(b); // free memory
cublasDestroy(handle); // destroy CUBLAS context
return EXIT_SUCCESS;

```

```

}
// lower triangle of a:
//   11
//   12   17
//   13   18   22
//   14   19   23   26
//   15   20   24   27   29
//   16   21   25   28   30   31
// b:
//   11   17   23   29   35
//   12   18   24   30   36
//   13   19   25   31   37
//   14   20   26   32   38
//   15   21   27   33   39
//   16   22   28   34   40

// solution x from Strsm : a*x=b
//   1.00000   1.54545   2.09091   2.63636   3.18182
//   0.00000   -0.03209   -0.06417   -0.09626   -0.12834
//   0.00000   -0.02333   -0.04667   -0.07000   -0.09334
//   0.00000   -0.01885   -0.03769   -0.05654   -0.07539
//   0.00000   -0.01625   -0.03250   -0.04874   -0.06499
//   0.00000   -0.01468   -0.02935   -0.04403   -0.05870

```

2.4.13 cublasChemmm - Hermitian matrix-matrix multiplication

This function performs the Hermitian matrix-matrix multiplication

$$C = \alpha AB + \beta C \quad \text{in CUBLAS_SIDE_LEFT case,}$$

$$C = \alpha BA + \beta C \quad \text{in CUBLAS_SIDE_RIGHT case,}$$

where A is a Hermitian $m \times m$ matrix in the first case and $n \times n$ Hermitian matrix in the second case, B, C are general $m \times n$ matrices and α, β are scalars. A can be stored in lower (CUBLAS_FILL_MODE_LOWER) or upper (CUBLAS_FILL_MODE_UPPER) mode.

```

// nvcc 042chemm.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))
#define m 6 // a - mxm matrix
#define n 5 // b,c - mxn matrices
int main(void){
    cudaError_t cudaStat; // cudaMalloc status
    cublasStatus_t stat; // CUBLAS functions status
    cublasHandle_t handle; // CUBLAS context
    int i,j; // i-row index, j-col. ind.
    // data preparation on the host

```

```

cuComplex *a;           // mxm complex matrix a on the host
cuComplex *b;           // mxn complex matrix b on the host
cuComplex *c;           // mxn complex matrix c on the host
a=(cuComplex*)malloc(m*m*sizeof(cuComplex)); // host memory
// alloc for a
b=(cuComplex*)malloc(m*n*sizeof(cuComplex)); // host memory
// alloc for b
c=(cuComplex*)malloc(m*n*sizeof(cuComplex)); // host memory
// alloc for c

// define the lower triangle of an mxm Hermitian matrix a in
// lower mode column by column
int ind=11;              // a:
for(j=0;j<m;j++){        // 11
    for(i=0;i<m;i++){    // 12,17
        if(i>=j){        // 13,18,22
            a[IDX2C(i,j,m)].x=(float)ind++; // 14,19,23,26
            a[IDX2C(i,j,m)].y=0.0f;         // 15,20,24,27,29
        }                  // 16,21,25,28,30,31
    }
}

//print the lower triangle of a row by row
printf("lower triangle of a:\n");
for(i=0;i<m;i++){
    for(j=0;j<m;j++){
        if(i>=j)
            printf("%5.0f+%2.0f*I",a[IDX2C(i,j,m)].x,
                    a[IDX2C(i,j,m)].y);
    }
    printf("\n");
}

// define mxn matrices b,c column by column
ind=11;                  // b,c:
for(j=0;j<n;j++){        // 11,17,23,29,35
    for(i=0;i<m;i++){    // 12,18,24,30,36
        b[IDX2C(i,j,m)].x=(float)ind;      // 13,19,25,31,37
        b[IDX2C(i,j,m)].y=0.0f;            // 14,20,26,32,38
        c[IDX2C(i,j,m)].x=(float)ind;      // 15,21,27,33,39
        c[IDX2C(i,j,m)].y=0.0f;            // 16,22,28,34,40
        ind++;
    }
}

// print b(=c) row by row
printf("b,c:\n");
for(i=0;i<m;i++){
    for(j=0;j<n;j++){
        printf("%5.0f+%2.0f*I",b[IDX2C(i,j,m)].x,
                b[IDX2C(i,j,m)].y);
    }
    printf("\n");
}

// on the device

```

```

cuComplex* d_a;                // d_a - a on the device
cuComplex* d_b;                // d_b - b on the device
cuComplex* d_c;                // d_c - c on the device
cudaStat=cudaMalloc((void**)&d_a,m*m*sizeof(cuComplex));
                                //device memory alloc for a
cudaStat=cudaMalloc((void**)&d_b,n*m*sizeof(cuComplex));
                                //device memory alloc for b
cudaStat=cudaMalloc((void**)&d_c,n*m*sizeof(cuComplex));
                                //device memory alloc for c
stat = cublasCreate(&handle); // initialize CUBLAS context
// copy matrices from the host to the device
stat = cublasSetMatrix(m,m,sizeof(*a),a,m,d_a,m); //a -> d_a
stat = cublasSetMatrix(m,n,sizeof(*b),b,m,d_b,m); //b -> d_b
stat = cublasSetMatrix(m,n,sizeof(*c),c,m,d_c,m); //c -> d_c
cuComplex al={1.0f,0.0f};      // al=1
cuComplex bet={1.0f,0.0f};     // bet=1
// Hermitian matrix-matrix multiplication:
// d_c=al*d_a*d_b+bet*d_c;
// d_a - mxm hermitian matrix; d_b,d_c - mxn-general matrices;
// al,bet - scalars

stat=cublasChemmm(handle,CUBLAS_SIDE_LEFT,CUBLAS_FILL_MODE_LOWER,
                  m,n,&al,d_a,m,d_b,m,&bet,d_c,m);

stat=cublasGetMatrix(m,n,sizeof(*c),d_c,m,c,m); // d_c -> c
printf("c after Chemm :\n");
for(i=0;i<m;i++){
    for(j=0;j<n;j++){
        //print c after Chemm
        printf("%5.0f+%1.0f*I",c[IDX2C(i,j,m)].x,
                c[IDX2C(i,j,m)].y);
    }
    printf("\n");
}
cudaFree(d_a);                // free device memory
cudaFree(d_b);                // free device memory
cudaFree(d_c);                // free device memory
cublasDestroy(handle);        // destroy CUBLAS context
free(a);                      // free host memory
free(b);                      // free host memory
free(c);                      // free host memory
return EXIT_SUCCESS;
}

//lower triangle of a:
// 11+ 0*I
// 12+ 0*I 17+ 0*I
// 13+ 0*I 18+ 0*I 22+ 0*I
// 14+ 0*I 19+ 0*I 23+ 0*I 26+ 0*I
// 15+ 0*I 20+ 0*I 24+ 0*I 27+ 0*I 29+ 0*I
// 16+ 0*I 21+ 0*I 25+ 0*I 28+ 0*I 30+ 0*I 31+ 0*I

// b,c:

```

```

// 11+ 0*I 17+ 0*I 23+ 0*I 29+ 0*I 35+ 0*I
// 12+ 0*I 18+ 0*I 24+ 0*I 30+ 0*I 36+ 0*I
// 13+ 0*I 19+ 0*I 25+ 0*I 31+ 0*I 37+ 0*I
// 14+ 0*I 20+ 0*I 26+ 0*I 32+ 0*I 38+ 0*I
// 15+ 0*I 21+ 0*I 27+ 0*I 33+ 0*I 39+ 0*I
// 16+ 0*I 22+ 0*I 28+ 0*I 34+ 0*I 40+ 0*I

// c after Chemm :
// 1122+0*I 1614+0*I 2106+0*I 2598+0*I 3090+0*I //
// 1484+0*I 2132+0*I 2780+0*I 3428+0*I 4076+0*I //
// 1740+0*I 2496+0*I 3252+0*I 4008+0*I 4764+0*I //      c=a*b+c
// 1912+0*I 2740+0*I 3568+0*I 4396+0*I 5224+0*I //
// 2025+0*I 2901+0*I 3777+0*I 4653+0*I 5529+0*I //
// 2107+0*I 3019+0*I 3931+0*I 4843+0*I 5755+0*I //

```

2.4.14 cublasChemmm - unified memory version

```

// nvcc 042chemm.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))
#define m 6 // a - mxm matrix
#define n 5 // b,c - mxn matrices
int main(void){
    cublasHandle_t handle; // CUBLAS context
    int i,j; // i-row index, j-col. ind.
    // data preparation
    cuComplex *a; // mxm complex matrix a
    cuComplex *b; // mxn complex matrix b
    cuComplex *c; // mxn complex matrix c
    cudaMallocManaged(&a,m*m*sizeof(cuComplex)); //unif.memory a
    cudaMallocManaged(&b,m*n*sizeof(cuComplex)); //unif.memory b
    cudaMallocManaged(&c,m*n*sizeof(cuComplex)); //unif.memory c
    // define the lower triangle of an mxm Hermitian matrix a in
    // lower mode column by column
    int ind=11; // a:
    for(j=0;j<m;j++){ // 11
        for(i=0;i<m;i++){ // 12,17
            if(i>=j){ // 13,18,22
                a[IDX2C(i,j,m)].x=(float)ind++; // 14,19,23,26
                a[IDX2C(i,j,m)].y=0.0f; // 15,20,24,27,29
            } // 16,21,25,28,30,31
        }
    }
}

//print the lower triangle of a row by row
printf("lower triangle of a:\n");
for(i=0;i<m;i++){
    for(j=0;j<m;j++){
        if(i>=j)
            printf("%5.0f+%2.0f*I",a[IDX2C(i,j,m)].x,

```

```

        a[IDX2C(i,j,m)].y);
    }
    printf("\n");
}
// define mxn matrices b,c column by column
ind=11; // b,c:
for(j=0;j<n;j++){ // 11,17,23,29,35
    for(i=0;i<m;i++){ // 12,18,24,30,36
        b[IDX2C(i,j,m)].x=(float)ind; // 13,19,25,31,37
        b[IDX2C(i,j,m)].y=0.0f; // 14,20,26,32,38
        c[IDX2C(i,j,m)].x=(float)ind; // 15,21,27,33,39
        c[IDX2C(i,j,m)].y=0.0f; // 16,22,28,34,40
        ind++;
    }
}
// print b(=c) row by row
printf("b,c:\n");
for(i=0;i<m;i++){
    for(j=0;j<n;j++){
        printf("%5.0f+%2.0f*I",b[IDX2C(i,j,m)].x,
            b[IDX2C(i,j,m)].y);
    }
    printf("\n");
}
cublasCreate(&handle); // initialize CUBLAS context
cuComplex al={1.0f,0.0f}; // al=1
cuComplex bet={1.0f,0.0f}; // bet=1
// Hermitian matrix-matrix multiplication:
// c=al*a*b+bet*c;
// a - mxm hermitian matrix; b,c - mxn-general matrices;
// al,bet - scalars

cublasChemv(handle,CUBLAS_SIDE_LEFT,CUBLAS_FILL_MODE_LOWER,
            m,n,&al,a,m,b,m,&bet,c,m);

cudaDeviceSynchronize();
printf("c after Chemv :\n");
for(i=0;i<m;i++){
    for(j=0;j<n;j++){ //print c after Chemv
        printf("%5.0f+%1.0f*I",c[IDX2C(i,j,m)].x,
            c[IDX2C(i,j,m)].y);
    }
    printf("\n");
}
cudaFree(a); // free memory
cudaFree(b); // free memory
cudaFree(c); // free memory
cublasDestroy(handle); // destroy CUBLAS context
return EXIT_SUCCESS;
}

//lower triangle of a:

```



```

//      11+ 0*I
//      12+ 0*I      17+ 0*I
//      13+ 0*I      18+ 0*I      22+ 0*I
//      14+ 0*I      19+ 0*I      23+ 0*I      26+ 0*I
//      15+ 0*I      20+ 0*I      24+ 0*I      27+ 0*I      29+ 0*I
//      16+ 0*I      21+ 0*I      25+ 0*I      28+ 0*I      30+ 0*I      31+ 0*I

// b,c:
//      11+ 0*I      17+ 0*I      23+ 0*I      29+ 0*I      35+ 0*I
//      12+ 0*I      18+ 0*I      24+ 0*I      30+ 0*I      36+ 0*I
//      13+ 0*I      19+ 0*I      25+ 0*I      31+ 0*I      37+ 0*I
//      14+ 0*I      20+ 0*I      26+ 0*I      32+ 0*I      38+ 0*I
//      15+ 0*I      21+ 0*I      27+ 0*I      33+ 0*I      39+ 0*I
//      16+ 0*I      22+ 0*I      28+ 0*I      34+ 0*I      40+ 0*I

// c after Chemm :
// 1122+0*I 1614+0*I 2106+0*I 2598+0*I 3090+0*I //
// 1484+0*I 2132+0*I 2780+0*I 3428+0*I 4076+0*I //
// 1740+0*I 2496+0*I 3252+0*I 4008+0*I 4764+0*I //      c=a*b+c
// 1912+0*I 2740+0*I 3568+0*I 4396+0*I 5224+0*I //
// 2025+0*I 2901+0*I 3777+0*I 4653+0*I 5529+0*I //
// 2107+0*I 3019+0*I 3931+0*I 4843+0*I 5755+0*I //

```

2.4.15 cublasCherk - Hermitian rank-k update

This function performs the Hermitian rank-k update

$$C = \alpha op(A)op(A)^H + \beta C,$$

where C is a Hermitian $n \times n$ matrix, $op(A)$ is an $n \times k$ matrix and α, β are scalars. The value of $op(A)$ can be equal to A in CUBLAS_OP_N case or A^H in CUBLAS_OP_C case (conjugate transposition). C can be stored in lower (CUBLAS_FILL_MODE_LOWER) or upper (CUBLAS_FILL_MODE_UPPER) mode.

```

// nvcc 043cherk.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))
#define n 6 // c - nxn matrix
#define k 5 // a - nxk matrix
int main(void){
    cudaError_t cudaStat; // cudaMalloc status
    cublasStatus_t stat; // CUBLAS functions status
    cublasHandle_t handle; // CUBLAS context
    int i,j; // i-row index, j-col. index
// data preparation on the host
    cuComplex *a; // nxk complex matrix a on the host
    cuComplex *c; // nxn complex matrix c on the host

```

```

a=(cuComplex*)malloc(n*k*sizeof(cuComplex)); // host memory
// alloc for a
c=(cuComplex*)malloc(n*n*sizeof(cuComplex)); // host memory
// alloc for c
// define the lower triangle of an nxn Hermitian matrix c in
// lower mode column by column;
int ind=11; // c:
for(j=0;j<n;j++){ // 11
    for(i=0;i<n;i++){ // 12,17
        if(i>=j){ // 13,18,22
            c[IDX2C(i,j,n)].x=(float)ind++; // 14,19,23,26
            c[IDX2C(i,j,n)].y=0.0f; // 15,20,24,27,29
        } // 16,21,25,28,30,31
    }
}
// print the lower triangle of c row by row
printf("lower triangle of c:\n");
for(i=0;i<n;i++){
    for(j=0;j<n;j++){
        if(i>=j)
            printf("%5.0f+%2.0f*I",c[IDX2C(i,j,n)].x,
                    c[IDX2C(i,j,n)].y);
    }
    printf("\n");
}
// define an nxk matrix a column by column
ind=11; // a:
for(j=0;j<k;j++){ // 11,17,23,29,35
    for(i=0;i<n;i++){ // 12,18,24,30,36
        a[IDX2C(i,j,n)].x=(float)ind; // 13,19,25,31,37
        a[IDX2C(i,j,n)].y=0.0f; // 14,20,26,32,38
        ind++; // 15,21,27,33,39
    } // 16,22,28,34,40
}
// print a row by row
printf("a:\n");
for(i=0;i<n;i++){
    for(j=0;j<k;j++){
        printf("%5.0f+%2.0f*I",a[IDX2C(i,j,n)].x,
                a[IDX2C(i,j,n)].y);
    }
    printf("\n");
}

// on the device
cuComplex* d_a; // d_a - a on the device
cuComplex* d_c; // d_c - c on the device
cudaStat=cudaMalloc((void**)&d_a,n*k*sizeof(cuComplex));
//device memory alloc for a
cudaStat=cudaMalloc((void**)&d_c,n*n*sizeof(cuComplex));
//device memory alloc for c
stat = cublasCreate(&handle); // initialize CUBLAS context

```

```

// copy matrices from the host to the device
stat = cublasSetMatrix(n,k,sizeof(*a),a,n,d_a,n); //a -> d_a
stat = cublasSetMatrix(n,n,sizeof(*c),c,n,d_c,n); //c -> d_c
float al=1.0f; // al=1
float bet=1.0f; //bet=1
// rank-k update of a Hermitian matrix:
// d_c=al*d_a*d_a^H +bet*d_c
// d_c - nxn, Hermitian matrix; d_a - nxk general matrix;
// al,bet - scalars

stat=cublasCherk(handle,CUBLAS_FILL_MODE_LOWER,CUBLAS_OP_N,
                 n,k,&al,d_a,n,&bet,d_c,n);

stat=cublasGetMatrix(n,n,sizeof(*c),d_c,n,c,n); // d_c -> c
printf("lower triangle of c after Cherk :\n");
for(i=0;i<n;i++){
    for(j=0;j<n;j++){ // print c after Cherk
        if(i>=j)
            printf("%5.0f+%1.0f*I",c[IDX2C(i,j,n)].x,
                    c[IDX2C(i,j,n)].y);
    }
    printf("\n");
}
cudaFree(d_a); // free device memory
cudaFree(d_c); // free device memory
cublasDestroy(handle); // destroy CUBLAS context
free(a); // free host memory
free(c); // free host memory
return EXIT_SUCCESS;
}

// lower triangle of c:
// 11+ 0*I
// 12+ 0*I 17+ 0*I
// 13+ 0*I 18+ 0*I 22+ 0*I
// 14+ 0*I 19+ 0*I 23+ 0*I 26+ 0*I
// 15+ 0*I 20+ 0*I 24+ 0*I 27+ 0*I 29+ 0*I
// 16+ 0*I 21+ 0*I 25+ 0*I 28+ 0*I 30+ 0*I 31+ 0*I
// a:
// 11+ 0*I 17+ 0*I 23+ 0*I 29+ 0*I 35+ 0*I
// 12+ 0*I 18+ 0*I 24+ 0*I 30+ 0*I 36+ 0*I
// 13+ 0*I 19+ 0*I 25+ 0*I 31+ 0*I 37+ 0*I
// 14+ 0*I 20+ 0*I 26+ 0*I 32+ 0*I 38+ 0*I
// 15+ 0*I 21+ 0*I 27+ 0*I 33+ 0*I 39+ 0*I
// 16+ 0*I 22+ 0*I 28+ 0*I 34+ 0*I 40+ 0*I

// lower triangle of c after Cherk :
// 3016+0*I
// 3132+0*I 3257+0*I
// 3248+0*I 3378+0*I 3507+0*I // c=a*a^H +c
// 3364+0*I 3499+0*I 3633+0*I 3766+0*I
// 3480+0*I 3620+0*I 3759+0*I 3897+0*I 4034+0*I
// 3596+0*I 3741+0*I 3885+0*I 4028+0*I 4170+0*I 4311+0*I

```

2.4.16 cublasCherk - unified memory version

```

// nvcc 043cherk.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))
#define n 6 // c - nxn matrix
#define k 5 // a - nxk matrix
int main(void){
    cublasHandle_t handle; // CUBLAS context
    int i,j; // i-row index, j-col. index
// data preparation
    cuComplex *a; // nxk complex matrix a
    cuComplex *c; // nxn complex matrix c
    cudaMallocManaged(&a,n*k*sizeof(cuComplex)); //unif.memory a
    cudaMallocManaged(&c,n*n*sizeof(cuComplex)); //unif.memory c
// define the lower triangle of an nxn Hermitian matrix c in
// lower mode column by column;
    int ind=11; // c:
    for(j=0;j<n;j++){ // 11
        for(i=0;i<n;i++){ // 12,17
            if(i>=j){ // 13,18,22
                c[IDX2C(i,j,n)].x=(float)ind++; // 14,19,23,26
                c[IDX2C(i,j,n)].y=0.0f; // 15,20,24,27,29
            } // 16,21,25,28,30,31
        }
    }
// print the lower triangle of c row by row
    printf("lower triangle of c:\n");
    for(i=0;i<n;i++){
        for(j=0;j<n;j++){
            if(i>=j)
                printf("%5.0f+%2.0f*I",c[IDX2C(i,j,n)].x,
                    c[IDX2C(i,j,n)].y);
        }
        printf("\n");
    }
// define an nxk matrix a column by column
    ind=11; // a:
    for(j=0;j<k;j++){ // 11,17,23,29,35
        for(i=0;i<n;i++){ // 12,18,24,30,36
            a[IDX2C(i,j,n)].x=(float)ind; // 13,19,25,31,37
            a[IDX2C(i,j,n)].y=0.0f; // 14,20,26,32,38
            ind++; // 15,21,27,33,39
        } // 16,22,28,34,40
    }
// print a row by row
    printf("a:\n");
    for(i=0;i<n;i++){
        for(j=0;j<k;j++){
            printf("%5.0f+%2.0f*I",a[IDX2C(i,j,n)].x,
                a[IDX2C(i,j,n)].y);
        }
    }
}

```

```

    }
    printf("\n");
}

cublasCreate(&handle);           // initialize CUBLAS context
float al=1.0f;                   // al=1
float bet=1.0f;                  // bet=1
// rank-k update of a Hermitian matrix: c=al*a*a^H +bet*c
// c - nxn, Hermitian matrix; a - nxk general matrix;
// al,bet - scalars

cublasCherk(handle,CUBLAS_FILL_MODE_LOWER,CUBLAS_OP_N,n,k,
             &al,a,n,&bet,c,n);

cudaDeviceSynchronize();
printf("lower triangle of c after Cherk :\n");
for(i=0;i<n;i++){
    for(j=0;j<n;j++){
        // print c after Cherk
        if(i>=j)
            printf("%5.0f+%1.0f*I",c[IDX2C(i,j,n)].x,
                    c[IDX2C(i,j,n)].y);
    }
    printf("\n");
}
cudaFree(a);                    // free memory
cudaFree(c);                    // free memory
cublasDestroy(handle);          // destroy CUBLAS context
return EXIT_SUCCESS;
}

// lower triangle of c:
//   11+ 0*I
//   12+ 0*I   17+ 0*I
//   13+ 0*I   18+ 0*I   22+ 0*I
//   14+ 0*I   19+ 0*I   23+ 0*I   26+ 0*I
//   15+ 0*I   20+ 0*I   24+ 0*I   27+ 0*I   29+ 0*I
//   16+ 0*I   21+ 0*I   25+ 0*I   28+ 0*I   30+ 0*I   31+ 0*I
// a:
//   11+ 0*I   17+ 0*I   23+ 0*I   29+ 0*I   35+ 0*I
//   12+ 0*I   18+ 0*I   24+ 0*I   30+ 0*I   36+ 0*I
//   13+ 0*I   19+ 0*I   25+ 0*I   31+ 0*I   37+ 0*I
//   14+ 0*I   20+ 0*I   26+ 0*I   32+ 0*I   38+ 0*I
//   15+ 0*I   21+ 0*I   27+ 0*I   33+ 0*I   39+ 0*I
//   16+ 0*I   22+ 0*I   28+ 0*I   34+ 0*I   40+ 0*I

// lower triangle of c after Cherk :
// 3016+0*I
// 3132+0*I 3257+0*I
// 3248+0*I 3378+0*I 3507+0*I           // c=a*a^H +c
// 3364+0*I 3499+0*I 3633+0*I 3766+0*I
// 3480+0*I 3620+0*I 3759+0*I 3897+0*I 4034+0*I
// 3596+0*I 3741+0*I 3885+0*I 4028+0*I 4170+0*I 4311+0*I

```

2.4.17 cublasCher2k - Hermitian rank-2k update

This function performs the Hermitian rank-2k update

$$C = \alpha \operatorname{op}(A) \operatorname{op}(B)^H + \bar{\alpha} \operatorname{op}(B) \operatorname{op}(A)^H + \beta C,$$

where C is a Hermitian $n \times n$ matrix, $\operatorname{op}(A), \operatorname{op}(B)$ are $n \times k$ matrices and α, β are scalars. The value of $\operatorname{op}(A)$ can be equal to A in CUBLAS_OP_N case or A^H (conjugate transposition) in CUBLAS_OP_C case and similarly for $\operatorname{op}(B)$. C can be stored in lower (CUBLAS_FILL_MODE_LOWER) or upper (CUBLAS_FILL_MODE_UPPER) mode.

```
// nvcc 044cher2k.c -lcublas
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))
#define n 6 // c - nxn matrix
#define k 5 // a,b - nxk matrices
int main(void){
    cudaError_t cudaStat; // cudaMalloc status
    cublasStatus_t stat; // CUBLAS functions status
    cublasHandle_t handle; // CUBLAS context
    int i,j; // i-row index, j-col. ind.
    // data preparation on the host
    cuComplex *a; // nxk complex matrix a on the host
    cuComplex *b; // nxk complex matrix b on the host
    cuComplex *c; // nxn complex matrix c on the host
    a=(cuComplex*)malloc(n*k*sizeof(cuComplex)) // host memory
    // alloc for a
    b=(cuComplex*)malloc(n*k*sizeof(cuComplex)); // host memory
    // alloc for b
    c=(cuComplex*)malloc(n*n*sizeof(cuComplex)); // host memory
    // alloc for c
    // define the lower triangle of an nxn Hermitian matrix c in
    // lower mode column by column
    int ind=11; // c:
    for(j=0;j<n;j++){ // 11
        for(i=0;i<n;i++){ // 12 17
            if(i>=j){ // 13,18,22
                c[IDX2C(i,j,n)].x=(float)ind; // 14,19,23,26
                c[IDX2C(i,j,n)].y=0.0f; // 15,20,24,27,29
                ind++; // 16,21,25,28,30,31
            }
        }
    }
    // print the lower triangle of c row by row
    printf("lower triangle of c:\n");
    for(i=0;i<n;i++){
        for(j=0;j<n;j++){
```

```

        if(i>=j)
            printf("%5.0f+%2.0f*I",c[IDX2C(i,j,n)].x,
                    c[IDX2C(i,j,n)].y);
    }
    printf("\n");
}
// define nxk matrices a,b column by column
ind=11;
for(j=0;j<k;j++){
    for(i=0;i<n;i++){
        a[IDX2C(i,j,n)].x=(float)ind;
        a[IDX2C(i,j,n)].y=0.0f;
        b[IDX2C(i,j,n)].x=(float)ind++;
        b[IDX2C(i,j,n)].y=0.0f;
    }
}
// print a(=b) row by row
printf("a,b:\n");
for(i=0;i<n;i++){
    for(j=0;j<k;j++){
        printf("%5.0f+%2.0f*I",a[IDX2C(i,j,n)].x,
                a[IDX2C(i,j,n)].y);
    }
    printf("\n");
}
// on the device
cuComplex* d_a; // d_a - a on the device
cuComplex* d_b; // d_b - b on the device
cuComplex* d_c; // d_c - c on the device
cudaStat=cudaMalloc((void**)&d_a,n*k*sizeof(cuComplex));
//device memory alloc for a
cudaStat=cudaMalloc((void**)&d_b,n*k*sizeof(cuComplex));
//device memory alloc for b
cudaStat=cudaMalloc((void**)&d_c,n*n*sizeof(cuComplex));
//device memory alloc for c
stat = cublasCreate(&handle); // initialize CUBLAS context
stat = cublasSetMatrix(n,k,sizeof(*a),a,n,d_a,n); //a -> d_a
stat = cublasSetMatrix(n,k,sizeof(*a),b,n,d_b,n); //b -> d_b
stat = cublasSetMatrix(n,n,sizeof(*c),c,n,d_c,n); //c -> d_c
cuComplex al={1.0f,0.0f}; // al=1
float bet=1.0f; //bet=1
// Hermitian rank-2k update:
// d_c=al*d_a*d_b^H+\bar{al}*d_b*d_a^H + bet*d_c
// d_c - nxn, hermitian matrix; d_a,d_b -nxk general matrices;
// al,bet - scalars

stat=cublasCher2k(handle,CUBLAS_FILL_MODE_LOWER,CUBLAS_OP_N,
                  n,k,&al,d_a,n,d_b,n,&bet,d_c,n);

stat=cublasGetMatrix(n,n,sizeof(*c),d_c,n,c,n); //d_c -> c
// print the updated lower triangle of c row by row
printf("lower triangle of c after Cher2k:\n");

```

```

    for(i=0;i<n;i++){
        for(j=0;j<n;j++){
            if(i>=j)
                printf("%6.0f+%2.0f*I",c[IDX2C(i,j,n)].x,
                    c[IDX2C(i,j,n)].y);
        }
        printf("\n");
    }
    cudaFree(d_a);
    cudaFree(d_b);
    cudaFree(d_c);
    cublasDestroy(handle);
    free(a);
    free(b);
    free(c);
    return EXIT_SUCCESS;
}

// lower triangle of c:
//   11+ 0*I
//   12+ 0*I   17+ 0*I
//   13+ 0*I   18+ 0*I   22+ 0*I
//   14+ 0*I   19+ 0*I   23+ 0*I   26+ 0*I
//   15+ 0*I   20+ 0*I   24+ 0*I   27+ 0*I   29+ 0*I
//   16+ 0*I   21+ 0*I   25+ 0*I   28+ 0*I   30+ 0*I   31+ 0*I
// a,b:
//   11+ 0*I   17+ 0*I   23+ 0*I   29+ 0*I   35+ 0*I
//   12+ 0*I   18+ 0*I   24+ 0*I   30+ 0*I   36+ 0*I
//   13+ 0*I   19+ 0*I   25+ 0*I   31+ 0*I   37+ 0*I
//   14+ 0*I   20+ 0*I   26+ 0*I   32+ 0*I   38+ 0*I
//   15+ 0*I   21+ 0*I   27+ 0*I   33+ 0*I   39+ 0*I
//   16+ 0*I   22+ 0*I   28+ 0*I   34+ 0*I   40+ 0*I

// lower triangle of c after Cher2k: c = a*b^H + b*a^H + c
//   6021+0*I
//   6252+0*I   6497+0*I
//   6483+0*I   6738+0*I   6992+0*I
//   6714+0*I   6979+0*I   7243+0*I   7506+0*I
//   6945+0*I   7220+0*I   7494+0*I   7767+0*I   8039+0*I
//   7176+0*I   7461+0*I   7745+0*I   8028+0*I   8310+0*I   8591+0*I

```

2.4.18 cublasCher2k - unified memory version

```

// nvcc 044cher2k.cu -lcublas
#include <stdio.h>
#include "cublas_v2.h"
#define IDX2C(i,j,ld) (((j)*(ld))+( i ))
#define n 6
#define k 5
int main(void){
    cublasHandle_t handle;

```



```

int i,j;                                // i-row index, j-col. ind.
cuComplex *a;                           // nxk complex matrix
cuComplex *b;                           // nxk complex matrix
cuComplex *c;                           // nxn complex matrix
cudaMallocManaged(&a,n*k*sizeof(cuComplex)); //unif.memory a
cudaMallocManaged(&b,n*k*sizeof(cuComplex)); //unif.memory b
cudaMallocManaged(&c,n*n*sizeof(cuComplex)); //unif.memory c
// define the lower triangle of an nxn Hermitian matrix c in
// lower mode column by column
int ind=11;                             // c:
for(j=0;j<n;j++){                       // 11
    for(i=0;i<n;i++){                   // 12 17
        if(i>=j){                       // 13,18,22
            c[IDX2C(i,j,n)].x=(float)ind; // 14,19,23,26
            c[IDX2C(i,j,n)].y=0.0f;      // 15,20,24,27,29
            ind++;                       // 16,21,25,28,30,31
        }
    }
}
// print the lower triangle of c row by row
printf("lower triangle of c:\n");
for(i=0;i<n;i++){
    for(j=0;j<n;j++){
        if(i>=j)
            printf("%5.0f+%2.0f*I",c[IDX2C(i,j,n)].x,
                    c[IDX2C(i,j,n)].y);
    }
    printf("\n");
}
// define nxk matrices a,b column by column
ind=11;                                 // a,b:
for(j=0;j<k;j++){                       // 11,17,23,29,35
    for(i=0;i<n;i++){                   // 12,18,24,30,36
        a[IDX2C(i,j,n)].x=(float)ind;   // 13,19,25,31,37
        a[IDX2C(i,j,n)].y=0.0f;         // 14,20,26,32,38
        b[IDX2C(i,j,n)].x=(float)ind++; // 15,21,27,33,39
        b[IDX2C(i,j,n)].y=0.0f;         // 16,22,28,34,40
    }
}
// print a(=b) row by row
printf("a,b:\n");
for(i=0;i<n;i++){
    for(j=0;j<k;j++){
        printf("%5.0f+%2.0f*I",a[IDX2C(i,j,n)].x,
                a[IDX2C(i,j,n)].y);
    }
    printf("\n");
}
cublasCreate(&handle);                  // initialize CUBLAS context
cuComplex al={1.0f,0.0f};               // al=1
float bet=1.0f;                         // bet=1
// Hermitian rank-2k update: c=al*a*b^H+\bar{al}*b*a^H +bet*c

```

```

// c - nxn, hermitian matrix; a,b - nxk general matrices;
// al,bet - scalars

    cublasCher2k(handle,CUBLAS_FILL_MODE_LOWER,CUBLAS_OP_N,n,k,
                  &al,a,n,b,n,&bet,c,n);

    cudaDeviceSynchronize();
// print the updated lower triangle of c row by row
printf("lower triangle of c after Cher2k :\n");
    for(i=0;i<n;i++){
        for(j=0;j<n;j++){
            //print c after Cher2k
            if(i>=j)
                printf("%6.0f+%2.0f*I",c[IDX2C(i,j,n)].x,
                        c[IDX2C(i,j,n)].y);
        }
        printf("\n");
    }
    cudaFree(a);
    cudaFree(b);
    cudaFree(c);
    cublasDestroy(handle);
    return EXIT_SUCCESS;
}

// lower triangle of c:
//   11+ 0*I
//   12+ 0*I   17+ 0*I
//   13+ 0*I   18+ 0*I   22+ 0*I
//   14+ 0*I   19+ 0*I   23+ 0*I   26+ 0*I
//   15+ 0*I   20+ 0*I   24+ 0*I   27+ 0*I   29+ 0*I
//   16+ 0*I   21+ 0*I   25+ 0*I   28+ 0*I   30+ 0*I   31+ 0*I

// a,b:
//   11+ 0*I   17+ 0*I   23+ 0*I   29+ 0*I   35+ 0*I
//   12+ 0*I   18+ 0*I   24+ 0*I   30+ 0*I   36+ 0*I
//   13+ 0*I   19+ 0*I   25+ 0*I   31+ 0*I   37+ 0*I
//   14+ 0*I   20+ 0*I   26+ 0*I   32+ 0*I   38+ 0*I
//   15+ 0*I   21+ 0*I   27+ 0*I   33+ 0*I   39+ 0*I
//   16+ 0*I   22+ 0*I   28+ 0*I   34+ 0*I   40+ 0*I

// lower triangle of c after Cher2k : c = a*b^H + b*a^H + c
//   6021+0*I
//   6252+0*I   6497+0*I
//   6483+0*I   6738+0*I   6992+0*I
//   6714+0*I   6979+0*I   7243+0*I   7506+0*I
//   6945+0*I   7220+0*I   7494+0*I   7767+0*I   8039+0*I
//   7176+0*I   7461+0*I   7745+0*I   8028+0*I   8310+0*I   8591+0*I

```

Chapter 3

CUSOLVER by example

3.1 General remarks on cuSolver

CUSOLVER is a part of CUDA environment and consists of three subsets of routines:

- cuSolverDn - dense matrix routines,
- cuSolverSp - sparse matrix routines,
- cuSolverRf - sparse re-factorization (useful when solving a sequence of matrices where only the coefficients are changed but the sparsity pattern remains the same).

The most complete source of information on cuSolver is http://docs.nvidia.com/cuda/pdf/CUSOLVER_Library.pdf. In this chapter we restrict ourselves to presentation of cuSolverDn, which includes the following topics.

- LU, QR and Cholesky factorization.
- Linear solvers based on LU, QR and Cholesky decompositions.
- Bunch-Kaufman factorization of symmetric indefinite matrices.
- Symmetric Eigenvalue and singular value problem solvers.
- Singular value decomposition.

All subprograms have four versions corresponding to four data types:

- **S** - **float** – real single-precision,
- **D** - **double** – real double-precision,
- **C** - **complex** single-precision,
- **Z** - **complex** double-precision.

Note on error checking. It is obvious that we should check for errors on every function call. Unfortunately, such an approach doubles the length of sample codes. Therefore we decided to compute the error codes only in the case of functions contained in subsections titles. A more careful error checking one can find in examples contained in [\[CUSOLVER\]](#).

3.1.1 Remarks on installation and compilation

As we already said, cuSolver is a part of CUDA, so after CUDA installation, cuSolver is ready to use. Compilation of examples needs two steps. Below we show how to compile our first example with Openblas library.

```
nvcc -Wno-deprecated-gpu-targets -c -std=c++11 -I/usr/local/cuda/include -I/usr/include 001dgetrf_exampleBigOnes.cpp

g++ -fopenmp 001dgetrf_exampleBigOnes.o -L/usr/local/cuda/lib64 -L/usr/lib -lcudart -lcublas -lcusolver -lopenblas
```

In the case of the Bunch-Kaufman decomposition example, the routines `ssytrs/dsytrs` are needed. Since we had problems with finding them in Openblas, we used MKL:

```
nvcc -I/opt/intel/mkl/include -Wno-deprecated-gpu-targets -c -std=c++11 -I/usr/local/cuda/include 005dsytrf_dsytrs_exampleBigOnes.cpp

icpc -mkl -fopenmp 005dsytrf_dsytrs_exampleBigOnes.o -L/usr/local/cuda/lib64 -lcudart -lcublas -lcusolver
```

3.1.2 Remarks on hardware used in examples

In all examples we have measured the computations times. The times were obtained on the machine with Ubuntu 16.04, Cuda 8.0 and

- Intel i7 6700 CPU, 4GHz, 16 GB RAM,
- Nvidia GeForce GTX 1080 GPU.

3.2 LU decomposition and solving general linear systems

3.2.1 cusolverDnSgetrf and cusolverDnSgetrs - solving general linear system using LU decomposition in single precision

The function `cusolverDnSgetrf` computes in single precision an LU factorization of a general $m \times n$ matrix A using partial pivoting with row interchanges:

$$P A = L U,$$

where P is a permutation matrix, L is lower triangular with unit diagonal, and U is upper diagonal. The information on the interchanged rows is contained in `d_pivot`. Using the obtained factorization one can replace the problem of solving a general linear system by solving two triangular systems with matrices L and U respectively. The function `cusolverDnSgetrs` uses the L, U factors defined by `cusolverDnSgetrf` to solve in single precision a general linear system

$$A X = B.$$

```
#include <cbblas.h>
#include <time.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <cuda_runtime.h>
#include <cusolver_v2.h>
#include <cusolverDn.h>
#define N 8192
#define BILLION 1000000000LL;
using namespace std;
int main(int argc, char*argv[]){
    struct timespec start,stop;           // variables for timing
    double accum;                         // elapsed time variable
    cublasStatus_t stat;
    cudaError_t cudaStatus;
    cusolverStatus_t cusolverStatus;
    cusolverDnHandle_t handle;
    // declare arrays on the host
    float *A, *B1, *B; // A - NxN matrix, B1 - auxiliary N-vect.
                        // B=A*B1 - N-vector of rhs, all on the host
    // declare arrays on the device
    float *d_A, *d_B, *d_Work; // coeff.matrix, rhs, workspace
    int *d_pivot, *d_info, Lwork; // pivots, info, worksp. size
    int info_gpu = 0;
    // prepare memory on the host
    A = (float*)malloc(N*N*sizeof(float));
    B = (float*)malloc(N*sizeof(float));
```

```

B1 = (float*)malloc(N*sizeof(float));
for(int i=0;i<N*N;i++) A[i]=rand()/(float)RAND_MAX;// A=rand
for(int i=0;i<N;i++) B[i] = 0.0;           // initialize B
for(int i=0;i<N;i++) B1[i] = 1.0;         // B1 - N-vector of ones
float al=1.0, bet=0.0;                     // coefficients for sgemv
int incx=1, incy=1;
cblas_sgemv(CblasColMajor,CblasNoTrans,N,N,al,A,N,B1,incx,
            bet,B,incy);                   // multiply B=A*B1
cudaStatus = cudaGetDevice(0);
cusolverStatus = cusolverDnCreate(&handle);
// prepare memory on the device
cudaStatus = cudaMalloc((void**)&d_A,N*N*sizeof(float));
cudaStatus = cudaMalloc((void**)&d_B, N*sizeof(float));
cudaStatus = cudaMalloc((void**)&d_pivot, N*sizeof(int));
cudaStatus = cudaMalloc((void**)&d_info, sizeof(int));
cudaStatus = cudaMemcpy(d_A, A, N*N*sizeof(float),
                        cudaMemcpyHostToDevice); // copy d_A<-A
cudaStatus = cudaMemcpy(d_B, B, N*sizeof(float),
                        cudaMemcpyHostToDevice); // copy d_B<-B
cusolverStatus = cusolverDnSgetrf_bufferSize(handle, N, N,
        d_A, N, &Lwork); // compute buffer size and prep.memory
cudaStatus=cudaMalloc((void**)&d_Work,Lwork*sizeof(float));
clock_gettime(CLOCK_REALTIME,&start); // timer start
// LU factorization of d_A, with partial pivoting and row
// interchanges; row i is interchanged with row d_pivot(i);

cusolverStatus = cusolverDnSgetrf(handle,N,N,d_A,N,d_Work,
        d_pivot, d_info);

// use the LU factorization to solve the system d_A*x=d_B;
// the solution overwrites d_B

cusolverStatus = cusolverDnSgetrs(handle, CUBLAS_OP_N, N, 1,
        d_A, N, d_pivot, d_B,N, d_info);

cudaStatus = cudaDeviceSynchronize();
clock_gettime(CLOCK_REALTIME,&stop); // timer stop
accum=(stop.tv_sec-start.tv_sec)+ // elapsed time
        (stop.tv_nsec-start.tv_nsec)/(double)BILLION;
printf("getrf+getrs time: %lf sec.\n",accum);//print el.time
cudaStatus = cudaMemcpy(&info_gpu, d_info, sizeof(int),
        cudaMemcpyDeviceToHost); //d_info -> info_gpu
printf("after getrf+getrs: info_gpu = %d\n", info_gpu);
cudaStatus = cudaMemcpy(B1, d_B, N*sizeof(float),
        cudaMemcpyDeviceToHost); // copy d_B->B1
printf("solution: ");
for (int i = 0; i < 5; i++) printf("%g, ", B1[i]);
printf(" ..."); // print first components of the solution
printf("\n");
// free memory
cudaStatus = cudaFree(d_A);
cudaStatus = cudaFree(d_B);

```

```

    cudaStatus = cudaFree(d_pivot);
    cudaStatus = cudaFree(d_info);
    cudaStatus = cudaFree(d_Work);
    free(A); free(B); free(B1);
    cusolverStatus = cusolverDnDestroy(handle);
    cudaStatus = cudaDeviceReset();
    return 0;
}
//getrf+getrs time: 0.267574 sec.
//after getrf+getrs: info_gpu = 0
//solution: 1.04225, 0.873826, 1.05703, 1.03822, 0.883831, ...

```

3.2.2 cusolverDnSgetrf and cusolverDnSgetrs - unified memory version

```

#include <cblas.h>
#include <time.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <cuda_runtime.h>
#include <cublas_v2.h>
#include <cusolverDn.h>
#define N 8192
#define BILLION 1000000000L;
using namespace std;
int main(int argc, char*argv[]){
    struct timespec start,stop;           // variables for timing
    double accum;                         // elapsed time variable
    cublasStatus_t stat;
    cudaError_t cudaStatus;
    cusolverStatus_t cusolverStatus;
    cusolverDnHandle_t handle;
    // declare arrays
    float *A, *B1, *B; // A - NxN matrix, B1 - auxiliary N-vect.
                                // B=A*B1 - N-vector of rhs
    float *Work; // workspace
    int *pivot, *info, Lwork; // pivots, info, worksp. size
    cudaMallocManaged(&A,N*N*sizeof(float)); //unified mem.for A
    cudaMallocManaged(&B,N*sizeof(float)); //unified mem.for B
    cudaMallocManaged(&B1,N*sizeof(float)); //unified mem.for B1
    for(int i=0;i<N*N;i++) A[i]=rand()/(float)RAND_MAX; //A=rand
    for(int i=0;i<N;i++) B[i] = 0.0f; // initialize B
    for(int i=0;i<N;i++) B1[i] = 1.0f; // B1 - N-vector of ones
    float al=1.0, bet=0.0; // coefficients for sgemv
    int incx=1, incy=1;
    cblas_sgemv(CblasColMajor,CblasNoTrans,N,N,al,A,N,B1,incx,
                bet,B,incy); // multiply B=A*B1
    cudaStatus = cudaGetDevice(0);
    cusolverStatus = cusolverDnCreate(&handle);

```

```

    cudaMallocManaged(&pivot,N*sizeof(int)); //unif.mem.for pivot
    cudaMallocManaged(&info,sizeof(int));    //unif.mem.for info
    cusolverStatus = cusolverDnSgetrf_bufferSize(handle, N, N,
        A, N, &Lwork);    // compute buffer size and prep.memory
    cudaMallocManaged(&Work,Lwork*sizeof(float));
    clock_gettime(CLOCK_REALTIME,&start);    // timer start
    // LU factorization of A, with partial pivoting and row
    // interchanges; row i is interchanged with row pivot(i);

    cusolverStatus = cusolverDnSgetrf(handle,N,N,A,N,Work,
        pivot, info);

    // use the LU factorization to solve the system A*x=B;
    // the solution overwrites B

    cusolverStatus = cusolverDnSgetrs(handle, CUBLAS_OP_N, N, 1,
        A, N, pivot, B,N, info);

    cudaStatus = cudaDeviceSynchronize();
    clock_gettime(CLOCK_REALTIME,&stop);    // timer stop
    accum=(stop.tv_sec-start.tv_sec)+    // elapsed time
        (stop.tv_nsec-start.tv_nsec)/((double)BILLION;
    printf("getrf+getrs time: %lf sec.\n",accum); //pr.elaps.time
    printf("after getrf+getrs: info = %d\n", *info);
    printf("solution: ");
    for (int i = 0; i < 5; i++) printf("%g, ", B[i]);
    printf(" ...");    // print first components of the solution
    printf("\n");
    // free memory
    cudaStatus = cudaFree(A);
    cudaStatus = cudaFree(B);
    cudaStatus = cudaFree(pivot);
    cudaStatus = cudaFree(info);
    cudaStatus = cudaFree(Work);
    cusolverStatus = cusolverDnDestroy(handle);
    cudaStatus = cudaDeviceReset();
    return 0;
}
//getrf+getrs time: 0.295533 sec.
//after getrf+getrs: info = 0
//solution: 1.04225 0.873826, 1.05703, 1.03822, 0.883831, ...

```

3.2.3 cusolverDnDgetrf and cusolverDnDgetrs - solving general linear system using LU decomposition in double precision

The function `cusolverDnDgetrf` computes in double precision an LU factorization of a general $m \times n$ matrix A using partial pivoting with row interchanges:

$$P A = L U,$$

where P is a permutation matrix, L is lower triangular with unit diagonal, and U is upper triangular. The information on the interchanged rows is contained in `d_pivot`. Using the obtained factorization one can replace the problem of solving a general linear system by solving two triangular systems with matrices L and U respectively. The function `cusolverDnDgetrs` uses the L, U factors defined by `cusolverDnDgetrf` to solve in double precision a general linear system

$$A X = B.$$

```
#include <cbblas.h>
#include <time.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <cuda_runtime.h>
#include <cublas_v2.h>
#include <cusolverDn.h>
#define N 8192
#define BILLION 1000000000L;
using namespace std;
int main(int argc, char*argv[]){
    struct timespec start, stop;           // variables for timing
    double accum;                          // elapsed time variable
    cublasStatus_t stat;
    cudaError_t cudaStatus;
    cusolverStatus_t cusolverStatus;
    cusolverDnHandle_t handle;
    // declare arrays on the host
    double *A, *B1, *B; // A - NxN matrix, B1 - auxiliary N-vect.
                        // B=A*B1 - N-vector of rhs, all on the host
    // declare arrays on the device
    double *d_A, *d_B, *d_Work; // coeff.matrix, rhs, workspace
    int *d_pivot, *d_info, Lwork; // pivots, info, worksp. size
    int info_gpu = 0;
    // prepare memory on the host
    A = (double*)malloc(N*N*sizeof(double));
    B = (double*)malloc(N*sizeof(double));
    B1 = (double*)malloc(N*sizeof(double));
    for(int i=0; i<N*N; i++) A[i]=rand()/(double)RAND_MAX; //A=rand
    for(int i=0; i<N; i++) B[i] = 0.0; // initialize B
    for(int i=0; i<N; i++) B1[i] = 1.0; // B1 - N-vector of ones
    double al=1.0, bet=0.0; // coefficients for dgemv
    int incx=1, incy=1;
    cbblas_dgemv(CblasColMajor, CblasNoTrans, N, N, al, A, N, B1, incx,
                bet, B, incy); // multiply B=A*B1
    cudaStatus = cudaGetDevice(0);
    cusolverStatus = cusolverDnCreate(&handle);
    // prepare memory on the device
    cudaStatus = cudaMalloc((void**)&d_A, N*N*sizeof(double));
    cudaStatus = cudaMalloc((void**)&d_B, N*sizeof(double));
```

```

cudaStatus = cudaMalloc((void**)&d_pivot, N*sizeof(int));
cudaStatus = cudaMalloc((void**)&d_info, sizeof(int));
cudaStatus = cudaMemcpy(d_A, A, N*N*sizeof(double),
                        cudaMemcpyHostToDevice); // copy d_A<-A
cudaStatus = cudaMemcpy(d_B, B, N*sizeof(double),
                        cudaMemcpyHostToDevice); // copy d_B<-B
cusolverStatus = cusolverDnDgetrf_bufferSize(handle, N, N,
        d_A, N, &Lwork); // compute buffer size and prep.memory
cudaStatus=cudaMalloc((void**)&d_Work,Lwork*sizeof(double));
clock_gettime(CLOCK_REALTIME,&start); // timer start
// LU factorization of d_A, with partial pivoting and row
// interchanges; row i is interchanged with row d_pivot(i);

cusolverStatus = cusolverDnDgetrf(handle,N,N,d_A,N,d_Work,
                                d_pivot, d_info);

// use the LU factorization to solve the system d_A*x=d_B;
// the solution overwrites d_B

cusolverStatus = cusolverDnDgetrs(handle, CUBLAS_OP_N, N, 1,
                                d_A, N, d_pivot, d_B,N, d_info);

cudaStatus = cudaDeviceSynchronize();
clock_gettime(CLOCK_REALTIME,&stop); // timer stop
accum=(stop.tv_sec-start.tv_sec)+ // elapsed time
        (stop.tv_nsec-start.tv_nsec)/((double)BILLION;
printf("getrf+getrs time: %lf sec.\n",accum); //pr.elapsed.time
cudaStatus = cudaMemcpy(&info_gpu, d_info, sizeof(int),
                        cudaMemcpyDeviceToHost); //d_info -> info_gpu
printf("after getrf+getrs: info_gpu = %d\n", info_gpu);
cudaStatus = cudaMemcpy(B1, d_B, N*sizeof(double),
                        cudaMemcpyDeviceToHost); // copy d_B->B1
printf("solution: ");
for (int i = 0; i < 5; i++) printf("%g, ", B1[i]);
printf(" ..."); // print first components of the solution
printf("\n");
// free memory
cudaStatus = cudaFree(d_A);
cudaStatus = cudaFree(d_B);
cudaStatus = cudaFree(d_pivot);
cudaStatus = cudaFree(d_info);
cudaStatus = cudaFree(d_Work);
free(A); free(B); free(B1);
cusolverStatus = cusolverDnDestroy(handle);
cudaStatus = cudaDeviceReset();
return 0;
}
//getrf+getrs time: 1.511761 sec.
//after getrf+getrs: info_gpu = 0
//solution: 1, 1, 1, 1, 1, ...

```

3.2.4 cusolverDnDgetrf and cusolverDnDgetrs - unified memory version

```

#include <cblas.h>
#include <time.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <cuda_runtime.h>
#include <cublas_v2.h>
#include <cusolverDn.h>
#define N 8192
#define BILLION 1000000000L;
using namespace std;
int main(int argc, char*argv[]){
    struct timespec start,stop;           // variables for timing
    double accum;                         // elapsed time variable
    cublasStatus_t stat;
    cudaError_t cudaStatus;
    cusolverStatus_t cusolverStatus;
    cusolverDnHandle_t handle;
// declare arrays
    double *A, *B1, *B; // A - NxN matrix, B1 - auxiliary N-vect.
                        // B=A*B1 - N-vector of rhs
    double *Work;           // workspace
    int *pivot, *info, Lwork; // pivots, info, worksp. size
    cudaMallocManaged(&A,N*N*sizeof(double)); //unified mem.for A
    cudaMallocManaged(&B,N*sizeof(double)); //unified mem.for B
    cudaMallocManaged(&B1,N*sizeof(double)); //unified mem.for B1
    for(int i=0;i<N*N;i++) A[i]=rand()/((double)RAND_MAX); //A=rand
    for(int i=0;i<N;i++) B[i] = 0.0; // initialize B
    for(int i=0;i<N;i++) B1[i] = 1.0; // B1 - N-vector of ones
    double al=1.0, bet=0.0; // coefficients for dgemv
    int incx=1, incy=1;
    cblas_dgemv(CblasColMajor,CblasNoTrans,N,N,al,A,N,B1,incx,
                bet,B,incy); // multiply B=A*B1
    cudaStatus = cudaGetDevice(0);
    cusolverStatus = cusolverDnCreate(&handle);
    cudaMallocManaged(&pivot,N*sizeof(int)); //unif.mem.for pivot
    cudaMallocManaged(&info,sizeof(int)); //unif.mem.for info
    cusolverStatus = cusolverDnDgetrf_bufferSize(handle, N, N,
        A, N, &Lwork); // compute buffer size and prep.memory
    cudaMallocManaged(&Work,Lwork*sizeof(double)); //mem.for Work
    clock_gettime(CLOCK_REALTIME,&start); // timer start
// LU factorization of A, with partial pivoting and row
// interchanges; row i is interchanged with row pivot(i);

    cusolverStatus = cusolverDnDgetrf(handle,N,N,A,N,Work,
        pivot, info);

// use the LU factorization to solve the system A*x=B;

```

```

// the solution overwrites B

cusolverStatus = cusolverDnDgetrs(handle, CUBLAS_OP_N, N, 1,
                                   A, N, pivot, B, N, info);

cudaStatus = cudaDeviceSynchronize();
clock_gettime(CLOCK_REALTIME, &stop);           // timer stop
accum=(stop.tv_sec-start.tv_sec)+                // elapsed time
      (stop.tv_nsec-start.tv_nsec)/((double)BILLION;
printf("getrf+getrs time: %lf sec.\n", accum); //pr.elaps.time
printf("after getrf+getrs: info = %d\n", *info);
printf("solution: ");
for (int i = 0; i < 5; i++) printf("%g, ", B[i]);
printf(" ..."); // print first components of the solution
printf("\n");
// free memory
cudaStatus = cudaFree(A);
cudaStatus = cudaFree(B);
cudaStatus = cudaFree(pivot);
cudaStatus = cudaFree(info);
cudaStatus = cudaFree(Work);
cusolverStatus = cusolverDnDestroy(handle);
cudaStatus = cudaDeviceReset();
return 0;
}
//getrf+getrs time: 1.595864 sec.
//after getrf+getrs: info_gpu = 0
//solution: 1, 1, 1, 1, 1, ...

```

3.3 QR decomposition and solving general linear systems

3.3.1 cusolverDnSgeqrf and cusolverDnSorgqr - QR decomposition and checking the orthogonality in single precision

The function `cusolverDnSgeqrf` computes in single precision the QR factorization:

$$A = Q R,$$

where A is an $m \times n$ general matrix, R is upper triangular and Q is orthogonal. On exit the upper triangle of A contains R . The orthogonal matrix Q is represented as a product of elementary reflectors $H(1) \dots H(\min(m, n))$, where $H(k) = I - \tau_k v_k v_k^T$. The real scalars τ_k are stored in an array `d_tau` and the nonzero components of vectors v_k are stored on exit in parts of columns of A corresponding to the lower triangular part of A . The function `cusolverDnSorgqr` computes the orthogonal matrix Q using elementary reflectors vectors stored in A and elementary reflectors scalars stored in `d_tau`.

```

#include <time.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <cuda_runtime.h>
#include <cublas_v2.h>
#include <cusolverDn.h>
#define BILLION 1000000000L;
int main(int argc, char*argv[])
{
    struct timespec start, stop;           // variables for timing
    double accum;                         // elapsed time variable
    cusolverDnHandle_t cusolverH;
    cublasHandle_t cublasH;
    cublasStatus_t cublas_status = CUBLAS_STATUS_SUCCESS;
    cusolverStatus_t cusolver_status = CUSOLVER_STATUS_SUCCESS;
    cudaError_t cudaStat = cudaSuccess;
    const int m = 8192;                   // number of rows of A
    const int n = 8192;                   // number of columns of A
    const int lda = m;                   // leading dimension of A
    // declare matrices A and Q,R on the host
    float *A, *Q, *R;
    // prepare host memory
    A=(float*)malloc(lda*n*sizeof(float)); // matr. A on the host
    Q=(float*)malloc(lda*n*sizeof(float)); // orthogonal factor Q
    R=(float*)malloc(n*n*sizeof(float));    // R=I-Q^T*Q
    for(int i=0;i<lda*n;i++) A[i]=rand()/(float)RAND_MAX; //rand
    float *d_A, *d_R ;                   // matrices A, R on the device
    float *d_tau ; // scalars defining the elementary reflectors
    int *devInfo ;                        // info on the device
    float *d_work;                       // workspace on the device
    // workspace sizes
    int lwork_geqrf = 0;
    int lwork_orgqr = 0;
    int lwork = 0;
    int info_gpu = 0;                     // info copied from the device
    const float h_one = 1;                // constants used in
    const float h_minus_one = -1;         // computations of I-Q^T*Q
    // create cusolver and cublas handles
    cusolver_status = cusolverDnCreate(&cusolverH);
    cublas_status = cublasCreate(&cublasH);
    // prepare device memory
    cudaStat = cudaMalloc ((void**)&d_A , sizeof(float)*lda*n);
    cudaStat = cudaMalloc ((void**)&d_tau , sizeof(float)*n);
    cudaStat = cudaMalloc ((void**)&devInfo , sizeof(int));
    cudaStat = cudaMalloc ((void**)&d_R , sizeof(float)*n*n);
    cudaStat = cudaMemcpy(d_A, A, sizeof(float)*lda*n,
                          cudaMemcpyHostToDevice); // copy d_A <- A
    // compute working space for geqrf and orgqr
    cusolver_status = cusolverDnSgeqrf_bufferSize(cusolverH,
    m, n, d_A, lda,&lwork_geqrf); // compute Sgeqrf buffer size
    cusolver_status = cusolverDnSorgqr_bufferSize(cusolverH,

```

```

    m, n, n, d_A, lda, d_tau, &lwork_orgqr); //and Sorgqr b.size
    lwork=(lwork_geqrf > lwork_orgqr)? lwork_geqrf: lwork_orgqr;
// device memory for workspace
    cudaStat = cudaMalloc((void**)&d_work, sizeof(float)*lwork);
// QR factorization for d_A
    clock_gettime(CLOCK_REALTIME,&start);           // start timer

    cusolver_status = cusolverDnSgeqrf(cusolverH,m, n, d_A, lda,
                                       d_tau, d_work, lwork, devInfo);

    cudaStat = cudaDeviceSynchronize();
    clock_gettime(CLOCK_REALTIME,&stop);           // stop timer
    accum=(stop.tv_sec-start.tv_sec)+             // elapsed time
          (stop.tv_nsec-start.tv_nsec)/((double)BILLION;
    printf("Sgeqrf time: %lf sec.\n",accum); //print elapsed time

    cudaStat = cudaMemcpy(&info_gpu, devInfo, sizeof(int),
                          cudaMemcpyDeviceToHost); // copy devInfo->info_gpu
// check geqrf error code
    printf("after geqrf: info_gpu = %d\n", info_gpu);
// apply orgqr function to compute the orthogonal matrix Q
// using elementary reflectors vectors stored in d_A and
// elementary reflectors scalars stored in d_tau,

    cusolver_status= cusolverDnSorgqr(cusolverH, m, n, n, d_A,
                                       lda, d_tau, d_work, lwork, devInfo);

    cudaStat = cudaDeviceSynchronize();
    cudaStat = cudaMemcpy(&info_gpu, devInfo, sizeof(int),
                          cudaMemcpyDeviceToHost); // copy devInfo->info_gpu
// check orgqr error code
    printf("after orgqr: info_gpu = %d\n", info_gpu);
    cudaStat = cudaMemcpy(Q, d_A, sizeof(float)*lda*n,
                          cudaMemcpyDeviceToHost); // copy d_A->Q
    memset(R, 0, sizeof(double)*n*n);           // nxn matrix of zeros
    for(int j = 0 ; j < n ; j++){
        R[j + n*j] = 1.0f;                      // ones on the diagonal
    }
    cudaStat = cudaMemcpy(d_R, R, sizeof(float)*n*n,
                          cudaMemcpyHostToDevice); // copy R-> d_R
// compute R = -Q**T*Q + I
    cublas_status=cublasSgemm_v2(cublasH,CUBLAS_OP_T,CUBLAS_OP_N,
    n, n, m, &h_minus_one, d_A, lda, d_A, lda, &h_one, d_R,n);
    float dR_nrm2 = 0.0;                        // norm value
// compute the norm of R = -Q**T*Q + I
    cublas_status = cublasSnrm2_v2(cublasH,n*n,d_R,1,&dR_nrm2);
    printf("||I - Q^T*Q|| = %E\n", dR_nrm2); // print the norm
// free memory
    cudaFree(d_A);
    cudaFree(d_tau);
    cudaFree(devInfo);
    cudaFree(d_work);

```

```

    cudaFree(d_R);
    cublasDestroy(cublasH);
    cusolverDnDestroy(cusolverH);
    cudaDeviceReset();
    return 0;
}
//Sgeqrf time: 0.434779 sec.
//after geqrf: info_gpu = 0
//after orgqr: info_gpu = 0
//|I - Q**T*Q| = 2.515004E-04

```

3.3.2 cusolverDnSgeqrf and cusolverDnSorgqr - unified memory version

```

#include <time.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <cuda_runtime.h>
#include <cublas_v2.h>
#include <cusolverDn.h>
#define BILLION 1000000000L;
int main(int argc, char*argv[])
{
    struct timespec start, stop;           // variables for timing
    double accum;                         // elapsed time variable
    cusolverDnHandle_t cusolverH;
    cublasHandle_t cublasH;
    cublasStatus_t cublas_status = CUBLAS_STATUS_SUCCESS;
    cusolverStatus_t cusolver_status = CUSOLVER_STATUS_SUCCESS;
    cudaError_t cudaStat = cudaSuccess;
    const int m = 8192;                   // number of rows of A
    const int n = 8192;                   // number of columns of A
    const int lda = m;                    // leading dimension of A
    // declare matrices A and Q,R
    float *A, *Q, *R;
    cudaMallocManaged(&A, lda*n*sizeof(float)); //unif. mem. for A
    cudaMallocManaged(&Q, lda*n*sizeof(float)); //unif. mem. for Q
    cudaMallocManaged(&R, n*n*sizeof(float)); //mem. for R=I-Q^T*Q
    for(int i=0; i<lda*n; i++) A[i]=rand()/(float)RAND_MAX; //rand
    float *tau; // scalars defining the elementary reflectors
    int *Info; // info
    float *work; // workspace
    // workspace sizes
    int lwork_geqrf = 0;
    int lwork_orgqr = 0;
    int lwork = 0;
    const float h_one = 1; // constants used in
    const float h_minus_one = -1; // computations of I-Q^T*Q
    // create cusolver and cublas handles

```

```

    cusolver_status = cusolverDnCreate(&cusolverH);
    cublas_status = cublasCreate(&cublasH);
// prepare memory
    cudaMallocManaged(&tau,n*sizeof(float)); //unif.mem.for tau
    cudaMallocManaged(&Info,sizeof(int)); //unif.mem.for Info
// compute working space for geqrf and orgqr
    cusolver_status = cusolverDnSgeqrf_bufferSize(cusolverH,
    m, n, A, lda, &lwork_geqrf); // compute Sgeqrf buffer size
    cusolver_status = cusolverDnSorgqr_bufferSize(cusolverH,
    m, n, n, A, lda, tau, &lwork_orgqr); //and Sorgqr b.size
    lwork=(lwork_geqrf > lwork_orgqr)? lwork_geqrf: lwork_orgqr;
// memory for workspace
    cudaMallocManaged(&work,lwork*sizeof(float)); //mem.for work
// QR factorization for A
    clock_gettime(CLOCK_REALTIME,&start); // start timer

    cusolver_status = cusolverDnSgeqrf(cusolverH, m, n, A, lda,
    tau, work, lwork, Info);

    cudaStat = cudaDeviceSynchronize();
    clock_gettime(CLOCK_REALTIME,&stop); // stop timer
    accum=(stop.tv_sec-start.tv_sec)+ // elapsed time
    (stop.tv_nsec-start.tv_nsec)/(double)BILLION;
    printf("Sgeqrf time :%lf\n",accum); // print elapsed time
// check geqrf error code
    printf("after geqrf: info = %d\n", *Info);
// apply orgqr function to compute the orthogonal matrix Q
// using elementary reflectors vectors stored in A and
// elementary reflectors scalars stored in tau,

    cusolver_status= cusolverDnSorgqr(cusolverH, m, n, n, A,
    lda, tau, work, lwork, Info);

    cudaStat = cudaDeviceSynchronize();
// check orgqr error code
    printf("after orgqr: info = %d\n", *Info);
    memset(R, 0, sizeof(float)*n*n); // nxn matrix of zeros
    for(int j = 0 ; j < n ; j++){
        R[j + n*j] = 1.0; // ones on the diagonal
    }
// compute R = -Q**T*Q + I
    cublas_status=cublasSgemm_v2(cublasH,CUBLAS_OP_T,CUBLAS_OP_N,
    n, n, m, &h_minus_one, A, lda, A, lda, &h_one, R, n);
    float nrm2 = 0.0; // norm value
// compute the norm of R = -Q**T*Q + I
    cublas_status = cublasSnrm2_v2(cublasH,n*n,R,1,&nrm2);
    printf("||I - Q^T*Q|| = %E\n", nrm2); // print the norm
// free memory
    cudaFree(A);
    cudaFree(Q);
    cudaFree(R);
    cudaFree(tau);

```



```

    cudaFree(Info);
    cudaFree(work);
    cublasDestroy(cublasH);
    cusolverDnDestroy(cusolverH);
    cudaDeviceReset();
    return 0;
}
//Sgeqrf time :0.470025
//after geqrf: info = 0
//after orgqr: info = 0
//||I - Q^T*Q|| = 2.515004E-04

```

3.3.3 cusolverDnDgeqrf and cusolverDnDorgqr - QR decomposition and checking the orthogonality in double precision

The function `cusolverDnDgeqrf` computes in double precision the QR factorization:

$$A = Q R,$$

where A is an $m \times n$ general matrix, R is upper triangular (upper trapezoidal in general case) and Q is orthogonal. On exit the upper triangle (trapezoid) of A contains R . The orthogonal matrix Q is represented as a product of elementary reflectors $H(1) \dots H(\min(m, n))$, where $H(k) = I - \tau_k v_k v_k^T$. The real scalars τ_k are stored in an array `d_tau` and the nonzero components of vectors v_k are stored on exit in parts of columns of A corresponding to the lower triangular (trapezoidal) part of A . The function `cusolverDnDorgqr` computes the orthogonal matrix Q using elementary reflectors vectors stored in A and elementary reflectors scalars stored in `d_tau`.

```

#include <time.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <cuda_runtime.h>
#include <cublas_v2.h>
#include <cusolverDn.h>
#define BILLION 1000000000L;
int main(int argc, char*argv[])
{
    struct timespec start, stop;           // variables for timing
    double accum;                          // elapsed time variable
    cusolverDnHandle_t cusolverH;
    cublasHandle_t cublasH;
    cublasStatus_t cublas_status = CUBLAS_STATUS_SUCCESS;
    cusolverStatus_t cusolver_status = CUSOLVER_STATUS_SUCCESS;
    cudaError_t cudaStat = cudaSuccess;
    const int m = 8192;                   // number of rows of A
    const int n = 8192;                   // number of columns of A
    const int lda = m;                    // leading dimension of A
    // declare matrices A and Q,R on the host
    double *A, *Q, *R;

```

```

// prepare host memory
A=(double*)malloc(lda*n*sizeof(double)); //matr.A on the host
Q=(double*)malloc(lda*n*sizeof(double)); //orthogonal matr. Q
R=(double*)malloc(n*n*sizeof(double)); // R=I-Q^T*Q
for(int i=0;i<lda*n;i++) A[i]=rand()/(double)RAND_MAX; //rand
double *d_A, *d_R ; // matrices A, R on the device
double *d_tau ; // scalars defining the elementary reflectors
int *devInfo ; // info on the device
double *d_work; // workspace on the device
// workspace sizes
int lwork_geqrf = 0;
int lwork_orgqr = 0;
int lwork = 0;
int info_gpu = 0; // info copied from the device
const double h_one = 1; // constants used in
const double h_minus_one = -1; // computations of I-Q^T*Q
// create cusolver and cublas handles
cusolver_status = cusolverDnCreate(&cusolverH);
cublas_status = cublasCreate(&cublasH);
// prepare device memory
cudaStat = cudaMalloc ((void**)&d_A , sizeof(double)*lda*n);
cudaStat = cudaMalloc ((void**)&d_tau, sizeof(double)*n);
cudaStat = cudaMalloc ((void**)&devInfo, sizeof(int));
cudaStat = cudaMalloc ((void**)&d_R , sizeof(double)*n*n);
cudaStat = cudaMemcpy(d_A, A, sizeof(double)*lda*n,
                      cudaMemcpyHostToDevice); // copy d_A <- A
// compute working space for geqrf and orgqr
cusolver_status = cusolverDnDgeqrf_bufferSize(cusolverH,
m, n, d_A, lda,&lwork_geqrf); // compute Sgeqrf buffer size
cusolver_status = cusolverDnDorgqr_bufferSize(cusolverH,
m, n, n, d_A, lda, d_tau, &lwork_orgqr); //and Sorgqr b.size
lwork=(lwork_geqrf > lwork_orgqr)? lwork_geqrf: lwork_orgqr;
// device memory for workspace
cudaStat = cudaMalloc((void**)&d_work, sizeof(double)*lwork);
// QR factorization for d_A
clock_gettime(CLOCK_REALTIME,&start); // start timer

cusolver_status = cusolverDnDgeqrf(cusolverH, m, n, d_A, lda,
d_tau, d_work, lwork, devInfo);

cudaStat = cudaDeviceSynchronize();
clock_gettime(CLOCK_REALTIME,&stop); // stop timer
accum=(stop.tv_sec-start.tv_sec)+ // elapsed time
(stop.tv_nsec-start.tv_nsec)/(double)BILLION;
printf("Dgeqrf time :%lf sec.\n",accum); //print elapsed time

cudaStat = cudaMemcpy(&info_gpu, devInfo, sizeof(int),
                      cudaMemcpyDeviceToHost); // copy devInfo->info_gpu
// check geqrf error code
printf("after geqrf: info_gpu = %d\n", info_gpu);
// apply orgqr function to compute the orthogonal matrix Q
// using elementary reflectors vectors stored in d_A and

```

```

// elementary reflectors scalars stored in d_tau,

cusolver_status= cusolverDnDorgqr(cusolverH, m, n, n, d_A,
                                   lda, d_tau, d_work, lwork, devInfo);

cudaStat = cudaDeviceSynchronize();
cudaStat = cudaMemcpy(&info_gpu, devInfo, sizeof(int),
                      cudaMemcpyDeviceToHost); // copy devInfo->info_gpu
// check orgqr error code
printf("after orgqr: info_gpu = %d\n", info_gpu);
cudaStat = cudaMemcpy(Q, d_A, sizeof(double)*lda*n,
                      cudaMemcpyDeviceToHost); // copy d_A->Q
memset(R, 0, sizeof(double)*n*n); // nxn matrix of zeros
for(int j = 0 ; j < n ; j++){
    R[j + n*j] = 1.0; // ones on the diagonal
}
cudaStat = cudaMemcpy(d_R, R, sizeof(double)*n*n,
                      cudaMemcpyHostToDevice); // copy R-> d_R
// compute R = -Q**T*Q + I
cublas_status=cublasDgemm_v2(cublasH,CUBLAS_OP_T,CUBLAS_OP_N,
                             n, n, m, &h_minus_one, d_A, lda, d_A, lda, &h_one, d_R,n);
double dR_nrm2 = 0.0; // norm value
// compute the norm of R = -Q**T*Q + I
cublas_status = cublasDnrm2_v2(cublasH,n*n,d_R,1,&dR_nrm2);
printf("||I - Q^T*Q|| = %E\n", dR_nrm2); // print the norm
// free memory
cudaFree(d_A);
cudaFree(d_tau);
cudaFree(devInfo);
cudaFree(d_work);
cudaFree(d_R);
cublasDestroy(cublasH);
cusolverDnDestroy(cusolverH);
cudaDeviceReset();
return 0;
}
//Dgeqrf time: 3.324072 sec.
//after geqrf: info_gpu = 0
//after orgqr: info_gpu = 0
//||I - Q**T*Q| = 4.646390E-13

```

3.3.4 cusolverDnDgeqrf and cusolverDnDorgqr - unified memory version

```

#include <time.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <cuda_runtime.h>
#include <cublas_v2.h>

```

```

#include <cusolverDn.h>
#define BILLION 1000000000L;
int main(int argc, char*argv[])
{
    struct timespec start,stop;           // variables for timing
    double accum;                         // elapsed time variable
    cusolverDnHandle_t cusolverH;
    cublasHandle_t cublasH;
    cublasStatus_t cublas_status = CUBLAS_STATUS_SUCCESS;
    cusolverStatus_t cusolver_status = CUSOLVER_STATUS_SUCCESS;
    cudaError_t cudaStat = cudaSuccess;
    const int m = 8192;                   // number of rows of A
    const int n = 8192;                   // number of columns of A
    const int lda = m;                    // leading dimension of A
    // declare matrices A and Q,R
    double *A, *Q, *R;
    cudaMallocManaged(&A,lda*n*sizeof(double)); //unif. mem.for A
    cudaMallocManaged(&Q,lda*n*sizeof(double)); //unif. mem.for Q
    cudaMallocManaged(&R,n*n*sizeof(double)); //mem.for R=I-Q^T*A
    for(int i=0;i<lda*n;i++) A[i]=rand()/(double)RAND_MAX; //rand
    double *tau ; // scalars defining the elementary reflectors
    int *Info ; // info
    double *work; // workspace
    // workspace sizes
    int lwork_geqrf = 0;
    int lwork_orgqr = 0;
    int lwork = 0;
    const double h_one = 1; // constants used in
    const double h_minus_one = -1; // computations of I-Q^T*A
    // create cusolver and cublas handles
    cusolver_status = cusolverDnCreate(&cusolverH);
    cublas_status = cublasCreate(&cublasH);
    // prepare memory
    cudaMallocManaged(&tau,n*sizeof(double)); //unif.mem.for tau
    cudaMallocManaged(&Info,sizeof(int)); //unif.mem.for Info
    // compute working space for geqrf and orgqr
    cusolver_status = cusolverDnDgeqrf_bufferSize(cusolverH,
    m, n, A, lda, &lwork_geqrf); // compute Sgeqrf buffer size
    cusolver_status = cusolverDnDorgqr_bufferSize(cusolverH,
    m, n, n, A, lda, tau, &lwork_orgqr); //and Sorgqr buff.size
    lwork=(lwork_geqrf > lwork_orgqr)? lwork_geqrf: lwork_orgqr;
    // memory for workspace
    cudaMallocManaged(&work,lwork*sizeof(double)); //mem.for work
    // QR factorization for A
    clock_gettime(CLOCK_REALTIME,&start); // start timer

    cusolver_status = cusolverDnDgeqrf(cusolverH,m, n, A, lda,
    tau, work, lwork, Info);

    cudaStat = cudaDeviceSynchronize();
    clock_gettime(CLOCK_REALTIME,&stop); // stop timer
    accum=(stop.tv_sec-start.tv_sec)+ // elapsed time

```

```

        (stop.tv_nsec-start.tv_nsec)/((double)BILLION;
    printf("Dgeqrf time :%lf sec.\n",accum); //print elapsed time
    // check geqrf error code
    printf("after geqrf: info = %d\n", *Info);
    // apply orgqr function to compute the orthogonal matrix Q
    // using elementary reflectors vectors stored in A and
    // elementary reflectors scalars stored in tau,

    cusolver_status= cusolverDnDorgqr(cusolverH, m, n, n, A,
                                      lda, tau, work, lwork, Info);

    cudaStat = cudaDeviceSynchronize();
    // check orgqr error code
    printf("after orgqr: info = %d\n", *Info);
    memset(R, 0, sizeof(double)*n*n); // nxn matrix of zeros
    for(int j = 0 ; j < n ; j++){
        R[j + n*j] = 1.0; // ones on the diagonal
    }
    // compute R = -Q**T*Q + I
    cublas_status=cublasDgemm_v2(cublasH,CUBLAS_OP_T,CUBLAS_OP_N,
        n, n, m, &h_minus_one, A, lda, A, lda, &h_one, R, n);
    double nrm2 = 0.0; // norm value
    // compute the norm of R = -Q**T*Q + I
    cublas_status = cublasDnrm2_v2(cublasH,n*n,R,1,&nrm2);
    printf("||I - Q^T*Q|| = %E\n", nrm2); // print the norm
    // free memory
    cudaFree(A);
    cudaFree(Q);
    cudaFree(R);
    cudaFree(tau);
    cudaFree(Info);
    cudaFree(work);
    cublasDestroy(cublasH);
    cusolverDnDestroy(cusolverH);
    cudaDeviceReset();
    return 0;
}
//Dgeqrf time :3.398122 sec.
//after geqrf: info = 0
//after orgqr: info = 0
//||I - Q^T*Q|| = 4.646390E-13

```

3.3.5 cusolverDnSgeqrf and cusolverDnSormqr, cublasStrsm - QR decomposition and solving a linear system in single precision

The function `cusolverDnSgeqrf` computes in single precision the QR factorization:

$$A = Q R,$$

where A is an $m \times n$ general matrix, R is upper triangular and Q is orthogonal. On exit the upper triangle of A contains R . The orthogonal matrix Q is represented as a product of elementary reflectors $H(1) \dots H(\min(m, n))$, where $H(k) = I - \tau_k v_k v_k^T$. The real scalars τ_k are stored in an array `d_tau` and the nonzero components of vectors v_k are stored on exit in parts of columns of A corresponding to the lower triangular part of A . The function `cusolverDnSormqr` computes $Q^T * B$, the original system $A * X = (Q * R) * X = B$ can be written in the form $R * X = Q^T * B$ and `cublasStrsm` solves the obtained triangular system.

```
#include <cblas.h>
#include <time.h>
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include <cublas_v2.h>
#include <cusolverDn.h>
#define BILLION 1000000000L;
int main(int argc, char*argv[])
{
    struct timespec start, stop;           // variables for timing
    double accum;                         // elapsed time variable
    cusolverDnHandle_t cusolverH;         // cusolver handle
    cublasHandle_t cublasH;               // cublas handle
    cublasStatus_t cublas_status = CUBLAS_STATUS_SUCCESS;
    cusolverStatus_t cusolver_status = CUSOLVER_STATUS_SUCCESS;
    // variables for error checking in cudaMalloc
    cudaError_t cudaStat1 = cudaSuccess;
    cudaError_t cudaStat2 = cudaSuccess;
    cudaError_t cudaStat3 = cudaSuccess;
    cudaError_t cudaStat4 = cudaSuccess;
    const int m = 8192;                   // number of rows of A
    const int lda = m;                    // leading dimension of A
    const int ldb = m;                    // leading dimension of B
    const int nrhs = 1;                   // number of right hand sides
    // A - mxm coeff. matr., B=A*B1 -right hand side, B1 - mxnrhs
    float *A, *B, *B1, *X;               // - auxil.matrix, X - solution
    // prepare memory on the host
    A=(float*)malloc(lda*m*sizeof(float));
    B=(float*)malloc(ldb*nrhs*sizeof(float));
    B1=(float*)malloc(ldb*nrhs*sizeof(float));
    X=(float*)malloc(ldb*nrhs*sizeof(float));
    for(int i=0; i<lda*m; i++) A[i]=rand()/(float)RAND_MAX;
    for(int i=0; i<ldb*nrhs; i++) B[i]=0.0;
    for(int i=0; i<ldb*nrhs; i++) B1[i]=1.0;
    float al=1.0, bet=0.0;                // constants for sgemv
    int incx=1, incy=1;
    cblas_sgemv(CblasColMajor, CblasNoTrans, m, m, al, A, m, B1, incx,
               bet, B, incy);             // B=A*B1
    // declare arrays on the device
    float *d_A, *d_B, *d_tau, *d_work ;
```

```

int *devInfo ;                               // device version of info
int lwork = 0;                               // workspace size
int info_gpu = 0;                            // device info copied to host
const float one = 1;

// create cusolver and cublas handles
cusolver_status = cusolverDnCreate(&cusolverH);
cublas_status = cublasCreate(&cublasH);
// prepare memory on the device
cudaStat1 = cudaMalloc((void**)&d_A, sizeof(float)*lda*m);
cudaStat2 = cudaMalloc((void**)&d_tau, sizeof(float) * m);
cudaStat3 = cudaMalloc((void**)&d_B, sizeof(float)*ldb*nrhs);
cudaStat4 = cudaMalloc((void**)&devInfo, sizeof(int));
// copy A,B from host to device
cudaStat1 = cudaMemcpy(d_A,A,sizeof(float)*lda*m,
                      cudaMemcpyHostToDevice); // A->d_A
cudaStat2 = cudaMemcpy(d_B,B,sizeof(float)*ldb*nrhs,
                      cudaMemcpyHostToDevice); // B->d_B
// compute buffer size for geqrf and prepare worksp. on device
cusolver_status = cusolverDnSgeqrf_bufferSize(cusolverH, m,
                                              m, d_A, lda, &lwork);
cudaStat1 = cudaMalloc((void**)&d_work, sizeof(float)*lwork);
clock_gettime(CLOCK_REALTIME,&start);        // start timer
// QR factorization for d_A; R stored in upper triangle of
// d_A, elementary reflectors vectors stored in lower triangle
// of d_A, elementary reflectors scalars stored in d_tau

cusolver_status = cusolverDnSgeqrf(cusolverH, m, m, d_A, lda,
                                   d_tau, d_work, lwork, devInfo);

cudaStat1 = cudaDeviceSynchronize();          // stop timer
clock_gettime(CLOCK_REALTIME,&stop);
accum=(stop.tv_sec-start.tv_sec)+             // elapsed time
      (stop.tv_nsec-start.tv_nsec)/(double)BILLION;
printf("Sgeqrf time: %lf sec.\n",accum); //print elapsed time

cudaStat1 = cudaMemcpy(&info_gpu,devInfo,sizeof(int),
                      cudaMemcpyDeviceToHost); // devInfo -> info_gpu
// check error code of geqrf function
printf("after geqrf: info_gpu = %d\n", info_gpu);
// compute d_B=Q^T*B using ormqr function

cusolver_status=cusolverDnSormqr(cusolverH,CUBLAS_SIDE_LEFT,
CUBLAS_OP_T, m, nrhs, m, d_A, lda, d_tau, d_B, ldb, d_work,
                                lwork, devInfo);

cudaStat1 = cudaDeviceSynchronize();
cudaStat1 = cudaMemcpy(&info_gpu, devInfo, sizeof(int),
                      cudaMemcpyDeviceToHost); // devInfo -> info_gpu
// check error code of ormqr function
printf("after ormqr: info_gpu = %d\n", info_gpu);
// write the original system A*X=(Q*R)*X=B in the form
// R*X=Q^T*B and solve the obtained triangular system

```

```

cublas_status = cublasStrsm(cublasH,CUBLAS_SIDE_LEFT,
CUBLAS_FILL_MODE_UPPER,CUBLAS_OP_N, CUBLAS_DIAG_NON_UNIT,
    m, nrhs, &one, d_A, lda, d_B, ldb);

cudaStat1 = cudaDeviceSynchronize();
cudaStat1 = cudaMemcpy(X,d_B,sizeof(float)*ldb*nrhs,
    cudaMemcpyDeviceToHost);           // copy d_B->X
printf("solution: ");//show first components of the solution
for (int i = 0; i < 5; i++) printf("%g, ", X[i]);
printf(" ...");
printf("\n");
// free memory
cudaFree(d_A);
cudaFree(d_tau);
cudaFree(d_B);
cudaFree(devInfo);
cudaFree(d_work);
cublasDestroy(cublasH);
cusolverDnDestroy(cusolverH);
cudaDeviceReset();
return 0;
}
//Sgeqrf time: 0.435715 sec.
//after geqrf: info_gpu = 0
//after ormqr: info_gpu = 0
//solution: 1.00008, 1.02025, 1.00586, 0.999749, 1.00595, ...

```

3.3.6 cusolverDnSgeqrf and cusolverDnSormqr, cublasStrsm - unified memory version

```

#include <cblas.h>
#include <time.h>
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include <cublas_v2.h>
#include <cusolverDn.h>
#define BILLION 1000000000L;
int main(int argc, char*argv[])
{
    struct timespec start,stop;           // variables for timing
    double accum;                         // elapsed time variable
    cusolverDnHandle_t cusolverH;         // cusolver handle
    cublasHandle_t cublasH;               // cublas handle
    cublasStatus_t cublas_status = CUBLAS_STATUS_SUCCESS;
    cusolverStatus_t cusolver_status = CUSOLVER_STATUS_SUCCESS;
    cudaError_t cudaStat1 = cudaSuccess;
    const int m = 8192;                   // number of rows of A
    const int lda = m;                    // leading dimension of A

```



```

const int ldb = m; // leading dimension of B
const int nrhs = 1; // number of right hand sides
// A - mxm coeff. matr., B=A*B1 -right hand side, B1 - mxnrhs
float *A, *B, *B1, *X; // - auxil.matrix, X - solution
// prepare unified memory
cudaMallocManaged(&A, lda*m*sizeof(float)); //unif. mem.for A
cudaMallocManaged(&B, ldb*nrhs*sizeof(float)); //uni.mem.for A
cudaMallocManaged(&B1, ldb*nrhs*sizeof(float)); //u.mem.for B1
for(int i=0; i<lda*m; i++) A[i]=rand()/(float)RAND_MAX;
for(int i=0; i<ldb*nrhs; i++) B[i]=0.0f;
for(int i=0; i<ldb*nrhs; i++) B1[i]=1.0f;
float al=1.0, bet=0.0; // constants for sgemv
int incx=1, incy=1;
cblas_sgemv(CblasColMajor, CblasNoTrans, m, m, al, A, m, B1, incx,
            bet, B, incy); //B=A*B1
float *tau, *work; //elem. reflectors scalars, workspace
int *Info; // info
int lwork = 0; // workspace size
const float one = 1;
// create cusolver and cublas handles
cusolver_status = cusolverDnCreate(&cusolverH);
cublas_status = cublasCreate(&cublasH);
cudaMallocManaged(&tau, m*sizeof(float)); //unif. mem. for tau
cudaMallocManaged(&Info, sizeof(int)); //unif. mem. for Info
// compute buffer size for geqrf and prepare workspace
cusolver_status=cusolverDnSgeqrf_bufferSize(cusolverH, m, m,
            A, lda, &lwork);
cudaMallocManaged(&work, lwork*sizeof(float)); //mem.for work
clock_gettime(CLOCK_REALTIME, &start); // start timer
// QR factorization for A; R stored in upper triangle of A
// elementary reflectors vectors stored in lower triangle of A
// elementary reflectors scalars stored in tau

cusolver_status = cusolverDnSgeqrf(cusolverH, m, m, A, lda,
            tau, work, lwork, Info);

cudaStat1 = cudaDeviceSynchronize(); // stop timer
clock_gettime(CLOCK_REALTIME, &stop);
accum=(stop.tv_sec-start.tv_sec)+ // elapsed time
        (stop.tv_nsec-start.tv_nsec)/((double)BILLION;
printf("Sgeqrf time: %lf sec.\n", accum); //print elapsed time
// check error code of geqrf function
printf("after geqrf: info = %d\n", *Info);
// compute B=Q^T*B using ormqr function

cusolver_status=cusolverDnSormqr(cusolverH, CUBLAS_SIDE_LEFT,
            CUBLAS_OP_T, m, nrhs, m, A, lda, tau, B, ldb, work, lwork, Info);

cudaStat1 = cudaDeviceSynchronize();
// check error code of ormqr function
printf("after ormqr: info = %d\n", *Info);
// write the original system A*X=(Q*R)*X=B in the form

```

```

// R*X=Q^T*B and solve the obtained triangular system

cublas_status = cublasStrsm(cublasH,CUBLAS_SIDE_LEFT,
CUBLAS_FILL_MODE_UPPER,CUBLAS_OP_N, CUBLAS_DIAG_NON_UNIT,
                           m, nrhs, &one, A, lda, B, ldb);

cudaStat1 = cudaDeviceSynchronize();
printf("solution: ");//show first components of the solution
for (int i = 0; i < 5; i++) printf("%g, ", B[i]);
printf(" ...");
printf("\n");
// free memory
cudaFree(A);
cudaFree(tau);
cudaFree(B);
cudaFree(B1);
cudaFree(Info);
cudaFree(work);
cublasDestroy(cublasH);
cusolverDnDestroy(cusolverH);
cudaDeviceReset();
return 0;
}
//Sgeqrf time: 0.465168 sec.
//after geqrf: info = 0
//after ormqr: info = 0
//solution: 1.00008, 1.02025, 1.00586, 0.999749, 1.00595, ...

```

3.3.7 cusolverDnDgeqrf and cusolverDnDormqr, cublasDtrsm - QR decomposition and solving a linear system in double precision

The function `cusolverDnDgeqrf` computes in double precision the QR factorization:

$$A = Q R,$$

where A is an $m \times n$ general matrix, R is upper triangular and Q is orthogonal. On exit the upper triangle of A contains R . The orthogonal matrix Q is represented as a product of elementary reflectors $H(1) \dots H(\min(m, n))$, where $H(k) = I - \tau_k v_k v_k^T$. The real scalars τ_k are stored in an array `d_tau` and the nonzero components of vectors v_k are stored on exit in parts of columns of A corresponding to the lower triangular part of A . The function `cusolverDnDormqr` computes $Q^T * B$, the original system $A * X = (Q * R) * X = B$ can be written in the form $R * X = Q^T * B$ and `cublasDtrsm` solves the obtained triangular system.

```

#include <cbblas.h>
#include <time.h>
#include <stdio.h>

```

```

#include <stdlib.h>
#include <cuda_runtime.h>
#include <cublas_v2.h>
#include <cusolverDn.h>
#define BILLION 1000000000L;
int main(int argc, char*argv[])
{
    struct timespec start,stop;           // variables for timing
    double accum;                         // elapsed time variable
    cusolverDnHandle_t cusolverH;        // cusolver handle
    cublasHandle_t cublasH;              // cublas handle
    cublasStatus_t cublas_status = CUBLAS_STATUS_SUCCESS;
    cusolverStatus_t cusolver_status = CUSOLVER_STATUS_SUCCESS;
    // variables for error checking in cudaMalloc
    cudaError_t cudaStat1 = cudaSuccess;
    cudaError_t cudaStat2 = cudaSuccess;
    cudaError_t cudaStat3 = cudaSuccess;
    cudaError_t cudaStat4 = cudaSuccess;
    const int m = 8192;                  // number of rows of A
    const int lda = m;                   // leading dimension of A
    const int ldb = m;                   // leading dimension of B
    const int nrhs = 1;                  // number of right hand sides
    // A - mxm coeff. matr., B=A*B1 -right hand side, B1 - mxnrhs
    double *A, *B, *B1, *X;             // - auxil.matrix, X - solution
    // prepare memory on the host
    A=(double*)malloc(lda*m*sizeof(double));
    B=(double*)malloc(ldb*nrhs*sizeof(double));
    B1=(double*)malloc(ldb*nrhs*sizeof(double));
    X=(double*)malloc(ldb*nrhs*sizeof(double));
    for(int i=0;i<lda*m;i++) A[i]=rand()/(double)RAND_MAX;
    for(int i=0;i<ldb*nrhs;i++) B[i]=0.0;;
    for(int i=0;i<ldb*nrhs;i++) B1[i]=1.0;
    double al=1.0,bet=0.0;               // constants for dgemv
    int incx=1, incy=1;
    cblas_dgemv(CblasColMajor,CblasNoTrans,m,m,al,A,m,B1,incx,
                bet,B,incy);             //B=A*B1
    // declare arrays on the device
    double *d_A, *d_B, *d_tau, *d_work ;
    int *devInfo ;                       // device version of info
    int lwork = 0;                       // workspace size
    int info_gpu = 0;                    // device info copied to host
    const double one = 1;
    // create cusolver and cublas handles
    cusolver_status = cusolverDnCreate(&cusolverH);
    cublas_status = cublasCreate(&cublasH);
    // prepare memory on the device
    cudaStat1 = cudaMalloc((void**)&d_A, sizeof(double)*lda*m);
    cudaStat2 = cudaMalloc((void**)&d_tau, sizeof(double) * m);
    cudaStat3 = cudaMalloc((void**)&d_B, sizeof(double)*ldb*nrhs);
    cudaStat4 = cudaMalloc((void**)&devInfo, sizeof(int));
    // copy A,B from host to device
    cudaStat1 = cudaMemcpy(d_A,A,sizeof(double)*lda*m,

```

```

        cudaMemcpyHostToDevice); // A->d_A
    cudaStat2 = cudaMemcpy(d_B,B,sizeof(double)*ldb*nrhs,
        cudaMemcpyHostToDevice); // B->d_B
// compute buffer size for geqrf and prepare worksp. on device
    cusolver_status=cusolverDnDgeqrf_bufferSize(cusolverH, m, m,
        d_A, lda, &lwork);
    cudaStat1=cudaMalloc((void**)&d_work,sizeof(double)*lwork);
    clock_gettime(CLOCK_REALTIME,&start); // start timer
// QR factorization for d_A; R stored in upper triangle of
// d_A, elementary reflectors vectors stored in lower triangle
// of d_A, elementary reflectors scalars stored in d_tau

    cusolver_status = cusolverDnDgeqrf(cusolverH, m, m, d_A, lda,
        d_tau, d_work, lwork, devInfo);

    cudaStat1 = cudaDeviceSynchronize(); // stop timer
    clock_gettime(CLOCK_REALTIME,&stop);
    accum=(stop.tv_sec-start.tv_sec)+ // elapsed time
        (stop.tv_nsec-start.tv_nsec)/((double)BILLION;
    printf("Dgeqrf time: %lf\n",accum); // print elapsed time

    cudaStat1 = cudaMemcpy(&info_gpu,devInfo,sizeof(int),
        cudaMemcpyDeviceToHost); // devInfo -> info_gpu
// check error code of geqrf function
    printf("after geqrf: info_gpu = %d\n", info_gpu);
// compute d_B=Q^T*B using ormqr function

    cusolver_status=cusolverDnDormqr(cusolverH,CUBLAS_SIDE_LEFT,
    CUBLAS_OP_T, m, nrhs, m, d_A, lda, d_tau, d_B, ldb, d_work,
        lwork, devInfo);

    cudaStat1 = cudaDeviceSynchronize();
    cudaStat1 = cudaMemcpy(&info_gpu, devInfo, sizeof(int),
        cudaMemcpyDeviceToHost); // devInfo -> info_gpu
// check error code of ormqr function
    printf("after ormqr: info_gpu = %d\n", info_gpu);
// write the original system A*X=(Q*R)*X=B in the form
// R*X=Q^T*B and solve the obtained triangular system

    cublas_status = cublasDtrsm(cublasH,CUBLAS_SIDE_LEFT,
    CUBLAS_FILL_MODE_UPPER,CUBLAS_OP_N, CUBLAS_DIAG_NON_UNIT,
        m, nrhs, &one, d_A, lda, d_B, ldb);

    cudaStat1 = cudaDeviceSynchronize();
    cudaStat1 = cudaMemcpy(X,d_B,sizeof(double)*ldb*nrhs,
        cudaMemcpyDeviceToHost); // copy d_B->X
    printf("solution: ");//show first components of the solution
    for (int i = 0; i < 5; i++) printf("%g, ", X[i]);
    printf(" ...");
    printf("\n");
// free memory
    cudaFree(d_A);

```

```

    cudaFree(d_tau);
    cudaFree(d_B);
    cudaFree(devInfo);
    cudaFree(d_work);
    cublasDestroy(cublasH);
    cusolverDnDestroy(cusolverH);
    cudaDeviceReset();
    return 0;
}
//Dgeqrf time: 3.333913 sec.
//after geqrf: info_gpu = 0
//after ormqr: info_gpu = 0
//solution: 1, 1, 1, 1, 1, ...

```

3.3.8 cusolverDnDgeqrf and cusolverDnDormqr, cublasDtrsm - unified memory version

```

#include <cblas.h>
#include <time.h>
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include <cublas_v2.h>
#include <cusolverDn.h>
#define BILLION 1000000000L;
int main(int argc, char*argv[])
{
    struct timespec start,stop;           // variables for timing
    double accum;                         // elapsed time variable
    cusolverDnHandle_t cusolverH;         // cusolver handle
    cublasHandle_t cublasH;               // cublas handle
    cublasStatus_t cublas_status = CUBLAS_STATUS_SUCCESS;
    cusolverStatus_t cusolver_status = CUSOLVER_STATUS_SUCCESS;
    cudaError_t cudaStat1 = cudaSuccess;

    const int m = 8192;                  // number of rows of A
    const int lda = m;                   // leading dimension of A
    const int ldb = m;                   // leading dimension of B
    const int nrhs = 1;                  // number of right hand sides
    // A - mxm coeff. matr., B=A*B1 -right hand side, B1 - mxnrhs
    double *A, *B, *B1, *X;             // - auxil.matrix, X - solution
    // prepare unified memory
    cudaMallocManaged(&A,lda*m*sizeof(double)); //unif. mem.for A
    cudaMallocManaged(&B,ldb*nrhs*sizeof(double)); //un.mem.for A
    cudaMallocManaged(&B1,ldb*nrhs*sizeof(double)); //mem.for B1
    for(int i=0;i<lda*m;i++) A[i]=rand()/(double)RAND_MAX;
    for(int i=0;i<ldb*nrhs;i++) B[i]=0.0;;
    for(int i=0;i<ldb*nrhs;i++) B1[i]=1.0;
    double al=1.0,bet=0.0;                // constants for dgemv
    int incx=1, incy=1;
    cblas_dgemv(CblasColMajor,CblasNoTrans,m,m,al,A,m,B1,incx,

```

```

                                bet,B,incy);           //B=A*B1
double *tau, *work ; //elem. reflectors scalars, workspace
int *Info ; // info
int lwork = 0; // workspace size
const double one = 1;
// create cusolver and cublas handles
cusolver_status = cusolverDnCreate(&cusolverH);
cublas_status = cublasCreate(&cublasH);
cudaMallocManaged(&tau,m*sizeof(double)); //unif.mem.for tau
cudaMallocManaged(&Info,sizeof(int)); //unif.mem.for Info
// compute buffer size for geqrf and prepare workspace
cusolver_status=cusolverDnDgeqrf_bufferSize(cusolverH, m, m,
                                A, lda, &lwork);
cudaMallocManaged(&work,lwork*sizeof(double)); //mem.for work
clock_gettime(CLOCK_REALTIME,&start); // start timer
// QR factorization for A; R stored in upper triangle of A
// elementary reflectors vectors stored in lower triangle of A
// elementary reflectors scalars stored in tau

cusolver_status = cusolverDnDgeqrf(cusolverH, m, m, A, lda,
                                tau, work, lwork, Info);

cudaStat1 = cudaDeviceSynchronize(); // stop timer
clock_gettime(CLOCK_REALTIME,&stop);
accum=(stop.tv_sec-start.tv_sec)+ // elapsed time
      (stop.tv_nsec-start.tv_nsec)/(double)BILLION;
printf("Dgeqrf time: %lf sec.\n",accum); //print elapsed time
// check error code of geqrf function
printf("after geqrf: info = %d\n", *Info);
// compute B=Q^T*B using ormqr function

cusolver_status=cusolverDnDormqr(cusolverH,CUBLAS_SIDE_LEFT,
                                CUBLAS_OP_T,m,nrhs,m,A,lda,tau,B,ldb,work,lwork,Info);

cudaStat1 = cudaDeviceSynchronize();
// check error code of ormqr function
printf("after ormqr: info = %d\n", *Info);
// write the original system A*X=(Q*R)*X=B in the form
// R*X=Q^T*B and solve the obtained triangular system

cublas_status = cublasDtrsm(cublasH,CUBLAS_SIDE_LEFT,
                                CUBLAS_FILL_MODE_UPPER,CUBLAS_OP_N, CUBLAS_DIAG_NON_UNIT,
                                m, nrhs, &one, A, lda, B, ldb);

cudaStat1 = cudaDeviceSynchronize();
printf("solution: "); //show first components of the solution
for (int i = 0; i < 5; i++) printf("%g, ", B[i]);
printf(" ...");
printf("\n");
// free memory
cudaFree(A);
cudaFree(tau);

```

```

    cudaFree(B);
    cudaFree(B1);
    cudaFree(Info);
    cudaFree(work);
    cublasDestroy(cublasH);
    cusolverDnDestroy(cusolverH);
    cudaDeviceReset();
    return 0;
}
//Dgeqrf time: 3.386223 sec.
//after geqrf: info = 0
//after ormqr: info = 0
//solution: 1, 1, 1, 1, 1, ...

```

3.4 Cholesky decomposition and solving positive definite linear systems

3.4.1 cusolverDnSpotrf and cusolverDnSpotrs - Choleski decomposition and solving positive definite systems in single precision

The function `cusolverDnSpotrf` computes in single precision the Cholesky factorization for a symmetric, positive definite $m \times m$ matrix A :

$$A = \begin{cases} U^T U & \text{in CUBLAS_FILL_MODE_UPPER case,} \\ L L^T & \text{in CUBLAS_FILL_MODE_LOWER case,} \end{cases}$$

where U is an upper triangular matrix and L is a lower triangular matrix. Using the obtained factorization the function `cusolverDnSpotrs` computes in single precision the solution of the linear system

$$A X = B,$$

where B, X are general $m \times n$ matrices. The solution X overwrites B .

```

#include <time.h>
#include <cbblas.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <cuda_runtime.h>
#include <cusolverDn.h>
#define N 8192
#define BILLION 1000000000L;
using namespace std;

int main(int argc, char*argv[]){
    struct timespec start,stop;                // variables for timing

```

```

double accum; // elapsed time variable
float *A, *B, *B1; // declare arrays on the host
// prepare memory on the host
A = (float*)malloc(N*N*sizeof(float)); // NxN coeff. matrix
B = (float*)malloc(N*sizeof(float)); // N-vector rhs B=A*B1
B1 = (float*)malloc(N*sizeof(float)); // auxiliary N-vect.
for(int i=0;i<N*N;i++) A[i] = rand()/(float)RAND_MAX;
for(int i=0;i<N;i++) B[i] = 0.0;
for(int i=0;i<N;i++) B1[i] = 1.0; // N-vector of ones
for(int i=0;i<N;i++){
    A[i*N+i]=A[i*N+i]+(float)N; // make A positive definite
    for(int j=0;j<i;j++) A[i*N+j]=A[j*N+i]; //and symmetric
}
float al=1.0,bet=0.0; // constants for sgemv
int incx=1, incy=1;
cblas_sgemv(CblasColMajor,CblasNoTrans,N,N,al,A,N,B1,incx,
            bet,B,incy); // B=A*B1

cudaError_t cudaStatus;
cusolverStatus_t cusolverStatus;
cusolverDnHandle_t handle; // device versions of
float *d_A, *d_B, *Work; // matrix A, rhs B and worksp.
int *d_info, Lwork; // device version of info, worksp.size
int info_gpu = 0; // device info copied to host
cudaStatus = cudaGetDevice(0);
cusolverStatus = cusolverDnCreate(&handle); // create handle
cublasFillMode_t uplo = CUBLAS_FILL_MODE_LOWER;
// prepare memory on the device
cudaStatus = cudaMalloc((void**)&d_A, N*N*sizeof(float));
cudaStatus = cudaMalloc((void**)&d_B, N*sizeof(float));
cudaStatus = cudaMalloc((void**)&d_info, sizeof(int));
cudaStatus = cudaMemcpy(d_A, A, N*N*sizeof(float),
                        cudaMemcpyHostToDevice); // copy A->d_A
cudaStatus = cudaMemcpy(d_B, B, N*sizeof(float),
                        cudaMemcpyHostToDevice); // copy B->d_B
// compute workspace size and prepare workspace
cusolverStatus = cusolverDnSpotrf_bufferSize(handle,
        uplo,N,d_A,N,&Lwork );
cudaStatus = cudaMalloc((void**)&Work,Lwork*sizeof(float));
clock_gettime(CLOCK_REALTIME,&start); // start timer
// Cholesky decomposition d_A=L*L^T, lower triangle of d_A is
// replaced by the factor L

cusolverStatus = cusolverDnSpotrf(handle,uplo,N,d_A,N,Work,
        Lwork,d_info);

// solve d_A*X=d_B, where d_A is factorized by potrf function
// d_B is overwritten by the solution

cusolverStatus = cusolverDnSpotrs(handle,uplo,N, 1,d_A,N,
        d_B,N,d_info);

cudaStatus = cudaDeviceSynchronize();

```



```

    clock_gettime(CLOCK_REALTIME, &stop);           // stop timer
    accum=(stop.tv_sec-start.tv_sec)+              // elapsed time
        (stop.tv_nsec-start.tv_nsec)/(double)BILLION;
    printf("Spotrf+Spotrs time: %lf sec.\n", accum); //pr.el.time
    cudaStatus = cudaMemcpy(&info_gpu, d_info, sizeof(int),
        cudaMemcpyDeviceToHost); // copy d_info -> info_gpu
    printf("after Spotrf+Spotrs: info_gpu = %d\n", info_gpu);
    cudaStatus = cudaMemcpy(B, d_B, N*sizeof(float),
        cudaMemcpyDeviceToHost); // copy solution to host d_B->B
    printf("solution: ");
    for (int i = 0; i < 5; i++) printf("%g, ", B[i]); // print
    printf(" ..."); // first components of the solution
    printf("\n");
// free memory
    cudaStatus = cudaFree(d_A);
    cudaStatus = cudaFree(d_B);
    cudaStatus = cudaFree(d_info);
    cudaStatus = cudaFree(Work);
    cusolverStatus = cusolverDnDestroy(handle);
    cudaStatus = cudaDeviceReset();
    return 0;
}
//Spotrf+Spotrs time: 0.057328
//after Spotrf+Spotrs: info_gpu = 0
//solution: 1, 1, 1, 0.999999, 1, ...

```

3.4.2 cusolverDnSpotrf and cusolverDnSpotrs - unified memory version

```

#include <time.h>
#include <cbblas.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <cuda_runtime.h>
#include <cusolverDn.h>
#define N 8192
#define BILLION 1000000000L;
using namespace std;

int main(int argc, char*argv[]){
    struct timespec start, stop;           // variables for timing
    double accum;                          // elapsed time variable
    float *A, *B, *B1;                    // declare arrays
// prepare unified memory
    cudaMallocManaged(&A, N*N*sizeof(float)); // unified mem.for A
    cudaMallocManaged(&B, N*sizeof(float)); // unified mem.for B
    cudaMallocManaged(&B1, N*sizeof(float)); // unified mem.for B1
    for(int i=0; i<N*N; i++) A[i] = rand()/(float)RAND_MAX;
    for(int i=0; i<N; i++) B[i] = 0.0;

```

```

for(int i=0;i<N;i++) B1[i] = 1.0;           // N-vector of ones
for(int i=0;i<N;i++){
    A[i*N+i]=A[i*N+i]+(float)N;           // make A positive definite
    for(int j=0;j<i;j++) A[i*N+j]=A[j*N+i]; //and symmetric
}
float al=1.0,bet=0.0;                       // constants for sgemv
int incx=1, incy=1;
cblas_sgemv(CblasColMajor,CblasNoTrans,N,N,al,A,N,B1,incx,
            bet,B,incy);                     // B=A*B1

cudaError_t cudaStatus;
cusolverStatus_t cusolverStatus;
cusolverDnHandle_t handle;
float *Work;                                // workspace
int *info, Lwork;                           // info, workspace size
cudaStatus = cudaGetDevice(0);
cusolverStatus = cusolverDnCreate(&handle); // create handle
cublasFillMode_t uplo = CUBLAS_FILL_MODE_LOWER;
cudaMallocManaged(&info,sizeof(int)); //unified mem. for info
// compute workspace size and prepare workspace
cusolverStatus = cusolverDnSpotrf_bufferSize(handle,
        uplo,N,A,N,&Lwork );
cudaMallocManaged(&Work,Lwork*sizeof(float)); //mem.for Work
clock_gettime(CLOCK_REALTIME,&start);         // start timer
// Cholesky decomposition d_A=L*L^T, lower triangle of d_A is
// replaced by the factor L

    cusolverStatus = cusolverDnSpotrf(handle,uplo,N,A,N,Work,
        Lwork,info);

    cudaStatus = cudaDeviceSynchronize();
// solve A*X=B, where A is factorized by potrf function
// B is overwritten by the solution

    cusolverStatus = cusolverDnSpotrs(handle,uplo,N,1,A,N,
        B,N,info);

    cudaStatus = cudaDeviceSynchronize();
    clock_gettime(CLOCK_REALTIME,&stop);           // stop timer
    accum=(stop.tv_sec-start.tv_sec)+              // elapsed time
        (stop.tv_nsec-start.tv_nsec)/((double)BILLION;
    printf("Spotrf+Spotrs time: %lf sec.\n",accum); //pr.el.time
    printf("after Spotrf+Spotrs: info = %d\n", *info);
    printf("solution: ");
    for (int i = 0; i < 5; i++) printf("%g, ", B[i]); // print
    printf(" ..."); // first components of the solution
    printf("\n");
// free memory
    cudaStatus = cudaFree(A);
    cudaStatus = cudaFree(B);
    cudaStatus = cudaFree(B1);
    cudaStatus = cudaFree(info);
    cudaStatus = cudaFree(Work);

```

```

    cusolverStatus = cusolverDnDestroy(handle);
    cudaStatus = cudaDeviceReset();
    return 0;
}
//Spotrf+Spotrs time: 0.094803 sec.
//after Spotrf+Spotrs: info = 0
//solution: 1, 1, 1, 0.999999, 1, ...

```

3.4.3 cusolverDnDpotrf and cusolverDnDpotrs - Choleski decomposition and solving positive definite systems in double precision

The function `cusolverDnDpotrf` computes in double precision the Cholesky factorization for a symmetric, positive definite $m \times m$ matrix A :

$$A = \begin{cases} U^T U & \text{in CUBLAS_FILL_MODE_UPPER case,} \\ L L^T & \text{in CUBLAS_FILL_MODE_LOWER case,} \end{cases}$$

where U is an upper triangular matrix and L is a lower triangular matrix. Using the obtained factorization the function `cusolverDnDpotrs` computes in double precision the solution of the linear system

$$A X = B,$$

where B, X are general $m \times n$ matrices. The solution X overwrites B .

```

#include <time.h>
#include <cblas.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <cuda_runtime.h>
#include <cusolverDn.h>
#define N 8192
#define BILLION 1000000000L;
using namespace std;

int main(int argc, char*argv[]){
    struct timespec start, stop;           // variables for timing
    double accum;                          // elapsed time variable
    double *A, *B, *B1;                  // declare arrays on the host
    // prepare memory on the host
    A = (double*)malloc(N*N*sizeof(double)); // NxN coeff. matrix
    B = (double*)malloc(N*sizeof(double)); // N-vector rhs B=A*B1
    B1 = (double*)malloc(N*sizeof(double)); // auxiliary N-vect.
    for(int i=0; i<N*N; i++) A[i] = rand()/(double)RAND_MAX;
    for(int i=0; i<N; i++) B[i] = 0.0;
    for(int i=0; i<N; i++) B1[i] = 1.0;    // N-vector of ones
    for(int i=0; i<N; i++){

```

```

    A[i*N+i]=A[i*N+i]+(double)N; // make A positive definite
    for(int j=0;j<i;j++) A[i*N+j]=A[j*N+i]; //and symmetric
}
double al=1.0,bet=0.0; // constants for dgemv
int incx=1, incy=1;
cblas_dgemv(CblasColMajor,CblasNoTrans,N,N,al,A,N,B1,incx,
            bet,B,incy); // B=A*B1
cudaError_t cudaStatus;
cusolverStatus_t cusolverStatus;
cusolverDnHandle_t handle; // device versions of
double *d_A, *d_B, *Work; // matrix A, rhs B and worksp.
int *d_info, Lwork; // device version of info, worksp.size
int info_gpu = 0; // device info copied to host
cudaStatus = cudaGetDevice(0);
cusolverStatus = cusolverDnCreate(&handle); // create handle
cublasFillMode_t uplo = CUBLAS_FILL_MODE_LOWER;
// prepare memory on the device
cudaStatus = cudaMalloc((void**)&d_A, N*N*sizeof(double));
cudaStatus = cudaMalloc((void**)&d_B, N*sizeof(double));
cudaStatus = cudaMalloc((void**)&d_info, sizeof(int));
cudaStatus = cudaMemcpy(d_A, A, N*N*sizeof(double),
                        cudaMemcpyHostToDevice); // copy A->d_A
cudaStatus = cudaMemcpy(d_B, B, N*sizeof(double),
                        cudaMemcpyHostToDevice); // copy B->d_B
// compute workspace size and prepare workspace
cusolverStatus = cusolverDnDpotrf_bufferSize(handle,
        uplo,N,d_A,N,&Lwork );
cudaStatus = cudaMalloc((void**)&Work, Lwork*sizeof(double));
clock_gettime(CLOCK_REALTIME,&start); // start timer
// Cholesky decomposition d_A=L*L^T, lower triangle of d_A is
// replaced by the factor L

    cusolverStatus = cusolverDnDpotrf(handle,uplo,N,d_A,N,Work,
        Lwork,d_info);

// solve d_A*X=d_B, where d_A is factorized by potrf function
// d_B is overwritten by the solution

    cusolverStatus = cusolverDnDpotrs(handle, uplo,N, 1,d_A, N,
        d_B,N,d_info);

    cudaStatus = cudaDeviceSynchronize();
    clock_gettime(CLOCK_REALTIME,&stop); // stop timer
    accum=(stop.tv_sec-start.tv_sec)+ // elapsed time
        (stop.tv_nsec-start.tv_nsec)/((double)BILLION;
    printf("solution: ");
    printf("Dpotrf+Dpotrs time: %lf sec.\n",accum); //pr.el.time
    cudaStatus = cudaMemcpy(&info_gpu, d_info, sizeof(int),
        cudaMemcpyDeviceToHost); // copy d_info -> info_gpu
    printf("after Dpotrf+Dpotrs: info_gpu = %d\n", info_gpu);
    cudaStatus = cudaMemcpy(B, d_B, N*sizeof(double),
        cudaMemcpyDeviceToHost); // copy solution to host d_B->B

```

```

    for (int i = 0; i < 5; i++) printf("%g, ", B[i]); // print
    printf(" ..."); // first components of the solution
    printf("\n");
// free memory
    cudaStatus = cudaFree(d_A);
    cudaStatus = cudaFree(d_B);
    cudaStatus = cudaFree(d_info);
    cudaStatus = cudaFree(Work);
    cusolverStatus = cusolverDnDestroy(handle);
    cudaStatus = cudaDeviceReset();
    return 0;
}

//Dpotrf+Dpotrs time: 0.754091 sec.
//after potrf: info_gpu = 0
//solution: 1, 1, 1, 1, 1, ...

```

3.4.4 cusolverDnDpotrf and cusolverDnDpotrs - unified memory version

```

#include <time.h>
#include <cbblas.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <cuda_runtime.h>
#include <cusolverDn.h>
#define N 8192
#define BILLION 1000000000L;
using namespace std;

int main(int argc, char*argv[]){
    struct timespec start,stop; // variables for timing
    double accum; // elapsed time variable
    double *A, *B, *B1; // declare arrays
// prepare unified memory
    cudaMallocManaged(&A,N*N*sizeof(double)); //unified mem.for A
    cudaMallocManaged(&B,N*sizeof(double)); //unified mem.for B
    cudaMallocManaged(&B1,N*sizeof(double)); //unified mem.for B1
    for(int i=0;i<N*N;i++) A[i] = rand()/(double)RAND_MAX;
    for(int i=0;i<N;i++) B[i] = 0.0;
    for(int i=0;i<N;i++) B1[i] = 1.0; // N-vector of ones
    for(int i=0;i<N;i++){
        A[i*N+i]=A[i*N+i]+(double)N; // make A positive definite
        for(int j=0;j<i;j++) A[i*N+j]=A[j*N+i]; //and symmetric
    }
    double al=1.0,bet=0.0; // constants for dgemv
    int incx=1, incy=1;
    cbblas_dgemv(CblasColMajor,CblasNoTrans,N,N,al,A,N,B1,incx,
                bet,B,incy); // B=A*B1

```

```

    cudaError_t cudaStatus;
    cusolverStatus_t cusolverStatus;
    cusolverDnHandle_t handle;
    double *Work; // workspace
    int *info, Lwork; // info, workspace size
    cudaStatus = cudaGetDevice(0);
    cusolverStatus = cusolverDnCreate(&handle); // create handle
    cublasFillMode_t uplo = CUBLAS_FILL_MODE_LOWER;
    cudaMallocManaged(&info, sizeof(int)); // unified mem. for info
    // compute workspace size and prepare workspace
    cusolverStatus = cusolverDnDpotrf_bufferSize(handle,
        uplo, N, A, N, &Lwork );
    cudaMallocManaged(&Work, Lwork * sizeof(double)); // mem. for Work
    clock_gettime(CLOCK_REALTIME, &start); // start timer
    // Cholesky decomposition d_A = L * L^T, lower triangle of d_A is
    // replaced by the factor L

    cusolverStatus = cusolverDnDpotrf(handle, uplo, N, A, N, Work,
        Lwork, info);

    cudaStatus = cudaDeviceSynchronize();
    // solve A * X = B, where A is factorized by potrf function
    // B is overwritten by the solution

    cusolverStatus = cusolverDnDpotrs(handle, uplo, N, 1, A, N,
        B, N, info);

    cudaStatus = cudaDeviceSynchronize();
    clock_gettime(CLOCK_REALTIME, &stop); // stop timer
    accum = (stop.tv_sec - start.tv_sec) + // elapsed time
        (stop.tv_nsec - start.tv_nsec) / (double) BILLION;
    printf("Dpotrf+Dpotrs time: %lf sec.\n", accum); // pr.el.time
    printf("after Dpotrf+Dpotrs: info = %d\n", *info);
    printf("solution: ");
    for (int i = 0; i < 5; i++) printf("%g, ", B[i]); // print
    printf(" ..."); // first components of the solution
    printf("\n");
    // free memory
    cudaStatus = cudaFree(A);
    cudaStatus = cudaFree(B);
    cudaStatus = cudaFree(B1);
    cudaStatus = cudaFree(info);
    cudaStatus = cudaFree(Work);
    cusolverStatus = cusolverDnDestroy(handle);
    cudaStatus = cudaDeviceReset();
    return 0;
}

//Dpotrf+Dpotrs time: 0.807432 sec.
//after Dpotrf+Dpotrs: info = 0
//solution: 1, 1, 1, 1, 1, ...

```

3.5 Bunch-Kaufman decomposition and solving symmetric linear systems

3.5.1 cusolverDnSsytrf and ssytrs - Bunch-Kaufman decomposition and solving symmetric systems in single precision

The function `cusolverDnSsytrf` computes Bunch-Kaufman factorization of a symmetric indefinite matrix

$$A = L * D * L^T,$$

where D is symmetric, block-diagonal, with 1×1 or 2×2 blocks, L is a product of permutation and triangular matrices. The function `ssytrs` solves the system $A * X = B$, where $A = L * D * L^T$ is the matrix factored using Bunch-Kaufman method, B is overwritten by the solution.

```
#include <time.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <cuda_runtime.h>
#include <cusolverDn.h>
#include <mkl.h>
#define BILLION 1000000000L;
using namespace std;
int main(int argc, char*argv[]){
    int N=8192;
    int nrhs=1;
    struct timespec start,stop;           // variables for timing
    double accum;                         // elapsed time variable
    float *A;                             // NxN coefficient matrix
    float *B, *B1;                        // N-vectors, rhs B=A*B1
    // prepare memory on the host
    A = (float*)malloc(N*N*sizeof(float));
    B1 = (float*)malloc(N*sizeof(float));
    B = (float*)malloc(N*sizeof(float));
    for(int i=0;i<N*N;i++) A[i] = rand()/(float)RAND_MAX;
    for(int i=0;i<N;i++) B[i] = 0.0;
    for(int i=0;i<N;i++) B1[i] = 1.0;      // N-vector of ones
    for(int i=0;i<N;i++){                 // make A symmetric
        for(int j=0;j<i;j++) A[i*N+j]=A[j*N+i];
    }
    float al=1.0,bet=0.0;                 // constants for sgemv
    int incx=1, incy=1;
    const char tr='N';
    sgemv(&tr,&N,&N,&al,A,&N,B1,&incx,&bet,B,&incy); // B=A*B1
    cudaError_t cudaStatus;
    cusolverStatus_t cusolverStatus;
    cusolverDnHandle_t handle;
```

```

float *d_A, *Work; // coeff. matrix and workspace on device
int *d_pivot, *d_info, Lwork; // pivots and info on device
int *piv, info; // pivots and info on the host
int info_gpu = 0; // device info copied to host
cudaStatus = cudaGetDevice(0);
cusolverStatus = cusolverDnCreate(&handle); // create handle
cublasFillMode_t uplo=CUBLAS_FILL_MODE_LOWER;
const char upl='L'; //use lower triangular part of A
// prepare memory on the device
cudaStatus = cudaMalloc((void**)&d_A, N*N*sizeof(float));
cudaStatus = cudaMalloc((void**)&d_pivot, N*sizeof(int));
cudaStatus = cudaMalloc((void**)&d_info, sizeof(int));
cudaStatus = cudaMemcpy(d_A, A, N*N*sizeof(float),
                        cudaMemcpyHostToDevice); // copy A->d_A
piv=(int*)malloc(N*sizeof(int));
cusolverStatus=cusolverDnSsytrf_bufferSize(handle,N,d_A,N,
      &Lwork ); // compute buffer size and prepare memory
cudaStatus = cudaMalloc((void**)&Work, Lwork*sizeof(float));
clock_gettime(CLOCK_REALTIME,&start); // start timer
// Bunch-Kaufman factorization of an NxN symmetric indefinite
// matrix d_A=L*D*L^T, where D is symmetric, block-diagonal,
// with 1x1 or 2x2 blocks, L is a product of permutation and
// triangular matrices

cusolverStatus = cusolverDnSsytrf(handle,uplo,N,d_A,N,d_pivot,
      Work,Lwork,d_info );

cudaStatus = cudaDeviceSynchronize();
cudaStatus = cudaMemcpy(A, d_A, N*N*sizeof(float),
                        cudaMemcpyDeviceToHost); // copy d_A->A
cudaStatus = cudaMemcpy(piv,d_pivot , N*sizeof(int),
                        cudaMemcpyDeviceToHost); // copy d_pivot->piv
// solve the system A*X=B , where A=L*D*L^T - symmetric
// coefficient matrix factored using Bunch-Kaufman method,
// B is overwritten by the solution

ssytrs(&upl,&N,&nrhs,A,&N,piv,B,&N,&info);

clock_gettime(CLOCK_REALTIME,&stop); // stop timer
accum=(stop.tv_sec-start.tv_sec)+ // elapsed time
      (stop.tv_nsec-start.tv_nsec)/((double)BILLION;
printf("Ssytrf+ssytrs time: %lf sec.\n",accum); //pr.el.time
cudaStatus = cudaMemcpy(&info_gpu, d_info, sizeof(int),
                        cudaMemcpyDeviceToHost); // copy d_info->info_gpu
printf("after Sytrf: info_gpu = %d\n", info_gpu);
printf("solution: ");
for (int i = 0; i < 5; i++) printf("%g, ", B[i]);
printf("...\n"); // first components of the solution
// free memory
cudaStatus = cudaFree(d_A);
cudaStatus = cudaFree(d_pivot);
cudaStatus = cudaFree(d_info);

```



```

    cudaStatus = cudaFree(Work);
    cusolverStatus = cusolverDnDestroy(handle);
    cudaStatus = cudaDeviceReset();
    return 0;
}
//Ssytrf+Ssytrs time: 0.397637 sec.
//after Sytrf: info_gpu = 0
//solution: 1.01025, 1.0031, 0.994385, 1.00684, 0.986153, ...

```

3.5.2 cusolverDnSsytrf and ssytrs - unified memory version

```
#include <time.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <cuda_runtime.h>
#include <cusolverDn.h>
#include <mkl.h>
#define BILLION 1000000000L;
using namespace std;
int main(int argc, char*argv[]){
    int N=8192;
    int nrhs=1;
    struct timespec start,stop;           // variables for timing
    double accum;                         // elapsed time variable
    float *A;                             // NxN coefficient matrix
    float *B, *B1;                        // N-vectors, rhs B=A*B1
    // prepare unified memory
    cudaMallocManaged(&A,N*N*sizeof(float)); //unified mem.for A
    cudaMallocManaged(&B,N*sizeof(float));    //unified mem.for B
    cudaMallocManaged(&B1,N*sizeof(float));   //unified mem.for B1
    for(int i=0;i<N*N;i++) A[i] = rand()/(float)RAND_MAX;
    for(int i=0;i<N;i++) B[i] = 0.0;
    for(int i=0;i<N;i++) B1[i] = 1.0;        // N-vector of ones
    for(int i=0;i<N;i++){                   // make A symmetric
        for(int j=0;j<i;j++) A[i*N+j]=A[j*N+i];
    }
    float al=1.0,bet=0.0;                  // constants for sgemv
    int incx=1, incy=1;
    const char tr='N';
    sgemv(&tr,&N,&N,&al,A,&N,B1,&incx,&bet,B,&incy); // B=A*B1
    cudaError_t cudaStatus;
    cusolverStatus_t cusolverStatus;
    cusolverDnHandle_t handle;
    float *Work;                           // workspace
    int *pivot,*info,Lwork;                // pivots, info, workspace size
    cudaStatus = cudaGetDevice(0);
    cusolverStatus = cusolverDnCreate(&handle); // create handle
    cublasFillMode_t uplo=CUBLAS_FILL_MODE_LOWER;
    const char upl='L';                    //use lower triangular part of A
```

```

    cudaMallocManaged(&pivot,N*sizeof(int)); //unif.mem.for pivot
    cudaMallocManaged(&info,sizeof(int));    //unif.mem.for info
    cusolverStatus=cusolverDnSsytrf_bufferSize(handle,N,A,N,
        &Lwork );    // compute buffer size and prepare memory
    cudaMallocManaged(&Work,Lwork*sizeof(float)); //mem.for Work
    clock_gettime(CLOCK_REALTIME,&start);    // start timer
// Bunch-Kaufman factorization of an NxN symmetric indefinite
// matrix A=L*D*L^T, where D is symmetric, block-diagonal,
// with 1x1 or 2x2 blocks, L is a product of permutation and
// triangular matrices

    cusolverStatus = cusolverDnSsytrf(handle,uplo,N,A,N, pivot,
        Work,Lwork, info );

    cudaStatus = cudaDeviceSynchronize();
// solve the system A*X=B , where A=L*D*L^T - symmetric
// coefficient matrix factored using Bunch-Kaufman method,
// B is overwritten by the solution

    ssytrs(&upl,&N,&nrhs,A,&N,pivot,B,&N,info);

    clock_gettime(CLOCK_REALTIME,&stop);    // stop timer
    accum=(stop.tv_sec-start.tv_sec)+    // elapsed time
        (stop.tv_nsec-start.tv_nsec)/(double)BILLION;
    printf("Ssytrf+ssytrs time: %lf sec.\n",accum); //pr.el.time
    printf("after Ssytrf: info = %d\n", *info);
    printf("solution: ");
    for (int i = 0; i < 5; i++) printf("%g, ", B[i]);
    printf("...\n");    // first components of the solution
// free memory
    cudaStatus = cudaFree(A);
    cudaStatus = cudaFree(B);
    cudaStatus = cudaFree(B1);
    cudaStatus = cudaFree(pivot);
    cudaStatus = cudaFree(info);
    cudaStatus = cudaFree(Work);
    cusolverStatus = cusolverDnDestroy(handle);
    cudaStatus = cudaDeviceReset();
    return 0;
}

//Ssytrf+ssytrs time: 0.544929 sec.
//after Ssytrf: info = 0
//solution: 1.01025 1.0031, 0.994385, 1.00684, 0.986153, ...

```

3.5.3 cusolverDnDsytrf and dsytrs - Bunch-Kaufman decomposition and solving symmetric systems in double precision

The function `cusolverDnDsytrf` computes Bunch-Kaufman factorization of a symmetric indefinite matrix

$$A = L * D * L^T,$$

where D is symmetric, block-diagonal, with 1×1 or 2×2 blocks, L is a product of permutation and triangular matrices. The function `dsytrs` solves the system $A * X = B$, where $A = L * D * L^T$ is the matrix factored using Bunch-Kaufman method, B is overwritten by the solution.

```
#include <time.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <cuda_runtime.h>
#include <cusolverDn.h>
#include <mkl.h>
#define BILLION 1000000000L;
using namespace std;
int main(int argc, char*argv[]){
    int N=8192;
    int nrhs=1;
    struct timespec start,stop;           // variables for timing
    double accum;                         // elapsed time variable
    double *A;                           // NxN coefficient matrix
    double *B, *B1;                       // N-vectors, rhs B=A*B1
    // prepare memory on the host
    A = (double*)malloc(N*N*sizeof(double));
    B1 = (double*)malloc(N*sizeof(double));
    B = (double*)malloc(N*sizeof(double));
    for(int i=0;i<N*N;i++) A[i] = rand()/(double)RAND_MAX;
    for(int i=0;i<N;i++) B[i] = 0.0;
    for(int i=0;i<N;i++) B1[i] = 1.0;      // N-vector of ones
    for(int i=0;i<N;i++){                 // make A symmetric
        for(int j=0;j<i;j++) A[i*N+j]=A[j*N+i];
    }
    double al=1.0,bet=0.0;                // constants for dgemv
    int incx=1, incy=1;
    const char tr='N';
    dgemv(&tr,&N,&N,&al,A,&N,B1,&incx,&bet,B,&incy); // B=A*B1
    cudaError_t cudaStatus;
    cusolverStatus_t cusolverStatus;
    cusolverDnHandle_t handle;
    double *d_A, *Work; // coeff. matrix and workspace on device
    int *d_pivot, *d_info, Lwork; // pivots and info on device
    int *piv, info; // pivots and info on the host
    int info_gpu = 0; // device info copied to host
    cudaStatus = cudaGetDevice(0);
    cusolverStatus = cusolverDnCreate(&handle); // create handle
    cublasFillMode_t uplo=CUBLAS_FILL_MODE_LOWER;
    const char upl='L'; //use lower triangular part of A
    // prepare memory on the device
    cudaStatus = cudaMalloc((void**)&d_A, N*N*sizeof(double));
    cudaStatus = cudaMalloc((void**)&d_pivot, N*sizeof(int));
    cudaStatus = cudaMalloc((void**)&d_info, sizeof(int));
    cudaStatus = cudaMemcpy(d_A, A, N*N*sizeof(double),
```

```

        cudaMemcpyHostToDevice);           // copy A->d_A
piv=(int*)malloc(N*sizeof(int));
cusolverStatus=cusolverDnDsytrf_bufferSize(handle,N,d_A,N,
        &Lwork );           // compute buffer size and prepare memory
cudaStatus = cudaMalloc((void**)&Work,Lwork*sizeof(double));
clock_gettime(CLOCK_REALTIME,&start);      // start timer
// Bunch-Kaufman factorization of an NxN symmetric indefinite
// matrix d_A=L*D*L^T, where D is symmetric, block-diagonal,
// with 1x1 or 2x2 blocks, L is a product of permutation and
// triangular matrices

cusolverStatus = cusolverDnDsytrf(handle,uplo,N,d_A,N,d_pivot,
        Work,Lwork,d_info );

cudaStatus = cudaDeviceSynchronize();
cudaStatus = cudaMemcpy(A, d_A, N*N*sizeof(double),
        cudaMemcpyDeviceToHost);           // copy d_A->A
cudaStatus = cudaMemcpy(piv,d_pivot, N*sizeof(int),
        cudaMemcpyDeviceToHost);           // copy d_pivot->piv
// solve the system A*X=B, where A=L*D*L^T - symmetric
// coefficient matrix factored using Bunch-Kaufman method,
// B is overwritten by the solution

dsytrs(&upl,&N,&nrhs,A,&N,piv,B,&N,&info);

clock_gettime(CLOCK_REALTIME,&stop);           // stop timer
accum=(stop.tv_sec-start.tv_sec)+           // elapsed time
        (stop.tv_nsec-start.tv_nsec)/((double)BILLION);
printf("Dsytrf+dsytrs time: %lf sec.\n",accum); //pr.el.time
cudaStatus = cudaMemcpy(&info_gpu, d_info, sizeof(int),
        cudaMemcpyDeviceToHost); // copy d_info->info_gpu
printf("after Dsytrf: info_gpu = %d\n", info_gpu);
printf("solution: ");
for (int i = 0; i < 5; i++) printf("%g, ", B[i]);
printf("...\n");           // first components of the solution
// free memory
cudaStatus = cudaFree(d_A);
cudaStatus = cudaFree(d_pivot);
cudaStatus = cudaFree(d_info);
cudaStatus = cudaFree(Work);
cusolverStatus = cusolverDnDestroy(handle);
cudaStatus = cudaDeviceReset();
return 0;
}
//Dsytrf+dsytrs time: 1.173202 sec.
//after Dsytrf: info_gpu = 0
//solution: 1, 1, 1, 1, 1, ...

```

3.5.4 cusolverDnDsytrf and dsytrs - unified memory version

```

#include <time.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <cuda_runtime.h>
#include <cusolverDn.h>
#include <mkl.h>
#define BILLION 1000000000L;
using namespace std;
int main(int argc, char*argv[]){
    int N=8192;
    int nrhs=1;
    struct timespec start,stop;           // variables for timing
    double accum;                         // elapsed time variable
    double *A;                            // NxN coefficient matrix
    double *B, *B1;                       // N-vectors, rhs B=A*B1
    // prepare unified memory
    cudaMallocManaged(&A,N*N*sizeof(double)); //unified mem.for A
    cudaMallocManaged(&B,N*sizeof(double));    //unified mem.for B
    cudaMallocManaged(&B1,N*sizeof(double));   //unified mem.for B1
    for(int i=0;i<N*N;i++) A[i] = rand()/(double)RAND_MAX;
    for(int i=0;i<N;i++) B[i] = 0.0;
    for(int i=0;i<N;i++) B1[i] = 1.0;         // N-vector of ones
    for(int i=0;i<N;i++){                    // make A symmetric
        for(int j=0;j<i;j++) A[i*N+j]=A[j*N+i];
    }
    double al=1.0,bet=0.0;                  // constants for dgemv
    int incx=1, incy=1;
    const char tr='N';
    dgemv(&tr,&N,&N,&al,A,&N,B1,&incx,&bet,B,&incy); // B=A*B1
    cudaError_t cudaStatus;
    cusolverStatus_t cusolverStatus;
    cusolverDnHandle_t handle;
    double *Work;                           // workspace
    int *pivot, *info, Lwork;               // pivots, info, workspace size
    cudaStatus = cudaGetDevice(0);
    cusolverStatus = cusolverDnCreate(&handle); // create handle
    cublasFillMode_t uplo=CUBLAS_FILL_MODE_LOWER;
    const char upl='L';                     //use lower triangular part of A
    cudaMallocManaged(&pivot,N*sizeof(int)); //unif.mem.for pivot
    cudaMallocManaged(&info,sizeof(int));    //unif.mem.for info
    cusolverStatus=cusolverDnDsytrf_bufferSize(handle,N,A,N,
        &Lwork );                          // compute buffer size and prepare memory
    cudaMallocManaged(&Work,Lwork*sizeof(double)); //mem.for Work
    clock_gettime(CLOCK_REALTIME,&start);    // start timer
    // Bunch-Kaufman factorization of an NxN symmetric indefinite
    // matrix A=L*D*L^T, where D is symmetric, block-diagonal,
    // with 1x1 or 2x2 blocks, L is a product of permutation and
    // triangular matrices

```

```

    cusolverStatus = cusolverDnDsytrf(handle,uplo,N,A,N,pivot,
                                     Work,Lwork, info );

    cudaStatus = cudaDeviceSynchronize();
    // solve the system A*X=B , where A=L*D*L^T - symmetric
    // coefficient matrix factored using Bunch-Kaufman method,
    // B is overwritten by the solution

    dsytrs(&upl,&N,&nrhs,A,&N,pivot,B,&N,info);

    clock_gettime(CLOCK_REALTIME,&stop);           // stop timer
    accum=(stop.tv_sec-start.tv_sec)+              // elapsed time
          (stop.tv_nsec-start.tv_nsec)/((double)BILLION);
    printf("Dsytrf+dsytrs time: %lf sec.\n",accum); //pr.el.time
    printf("after Dsytrf: info = %d\n", *info);
    printf("solution: ");
    for (int i = 0; i < 5; i++) printf("%g, ", B[i]);
    printf("...\n");           // first components of the solution
    // free memory
    cudaStatus = cudaFree(A);
    cudaStatus = cudaFree(B);
    cudaStatus = cudaFree(B1);
    cudaStatus = cudaFree(pivot);
    cudaStatus = cudaFree(info);
    cudaStatus = cudaFree(Work);
    cusolverStatus = cusolverDnDestroy(handle);
    cudaStatus = cudaDeviceReset();
    return 0;
}
//Dsytrf+dsytrs time: 1.279214 sec.
//after Dsytrf: info = 0
//solution: 1, 1, 1, 1, 1, ...

```

3.6 SVD decomposition

3.6.1 cusolverDnSgesvd - SVD decomposition in single precision

This function computes in single precision the singular value decomposition of an $m \times n$ matrix:

$$A = u \sigma v^T,$$

where σ is an $m \times n$ matrix which is zero except for its $\min(m, n)$ diagonal elements (singular values), u is an $m \times m$ orthogonal matrix and v is an $n \times n$ orthogonal matrix. The first $\min(m, n)$ columns of u and v are the left and right singular vectors of A respectively.

```

#include <time.h>
#include <stdio.h>
#include <stdlib.h>

```

```

#include <string.h>
#include <cuda_runtime.h>
#include <cublas_v2.h>
#include <cusolverDn.h>
#define BILLION 1000000000L;
int main(int argc, char*argv[])
{
    struct timespec start, stop;           // variables for timing
    double accum;                         // elapsed time variable
    cusolverDnHandle_t cusolverH;         // cusolver handle
    cublasHandle_t cublasH;               // cublas handle
    cublasStatus_t cublas_status = CUBLAS_STATUS_SUCCESS;
    cusolverStatus_t cusolver_status = CUSOLVER_STATUS_SUCCESS;
    cudaError_t cudaStat = cudaSuccess;

    const int m = 2048;                   // number of rows of A
    const int n = 2048;                   // number of columns of A
    const int lda = m;                    // leading dimension of A
    // declare the factorized matrix A, orthogonal matrices U, VT
    float *A, *U, *VT, *S; // and sing.val. matrix S on the host
    A=(float*)malloc(lda*n*sizeof(float));
    U=(float*)malloc(lda*m*sizeof(float));
    VT=(float*)malloc(lda*n*sizeof(float));
    S= (float*)malloc(n*sizeof(float));
    for(int i=0;i<lda*n;i++) A[i]=rand()/(float)RAND_MAX;
    // the factorized matrix d_A, orthogonal matrices d_U, d_VT
    float *d_A, *d_U, *d_VT, *d_S; // and sing.val. matrix d_S
    int *devInfo; // on the device
    float *d_work, *d_rwork; // workspace on the device
    float *d_W; // auxiliary device array (d_W = d_S*d_VT)
    int lwork = 0;
    int info_gpu = 0; // info copied from device to host
    const float h_one = 1;
    const float h_minus_one = -1;
    // create cusolver and cublas handle
    cusolver_status = cusolverDnCreate(&cusolverH);
    cublas_status = cublasCreate(&cublasH);
    // prepare memory on the device
    cudaStat = cudaMalloc((void**)&d_A, sizeof(float)*lda*n);
    cudaStat = cudaMalloc((void**)&d_S, sizeof(float)*n);
    cudaStat = cudaMalloc((void**)&d_U, sizeof(float)*lda*m);
    cudaStat = cudaMalloc((void**)&d_VT, sizeof(float)*lda*n);
    cudaStat = cudaMalloc((void**)&devInfo, sizeof(int));
    cudaStat = cudaMalloc((void**)&d_W, sizeof(float)*lda*n);
    cudaStat = cudaMemcpy(d_A, A, sizeof(float)*lda*n,
                          cudaMemcpyHostToDevice); // copy A->d_A
    // compute buffer size and prepare workspace
    cusolver_status = cusolverDnSgesvd_bufferSize(cusolverH, m, n,
                                                  &lwork );
    cudaStat = cudaMalloc((void**)&d_work, sizeof(float)*lwork);
    // compute the singular value decomposition of d_A
    // and optionally the left and right singular vectors:
    // d_A = d_U*d_S*d_VT; the diagonal elements of d_S

```

```

// are the singular values of d_A in descending order
// the first min(m,n) columns of d_U contain the left sing.vec.
// the first min(m,n) cols of d_VT contain the right sing.vec.
signed char jobu = 'A';    // all m columns of d_U returned
signed char jobvt = 'A';   // all n columns of d_VT returned
clock_gettime(CLOCK_REALTIME,&start);    // start timer

cusolver_status = cusolverDnSgesvd (cusolverH, jobu, jobvt,
m, n, d_A, lda, d_S, dU, lda, d_VT, lda, d_work, lwork,
                                d_rwork, devInfo);

cudaStat = cudaDeviceSynchronize();
clock_gettime(CLOCK_REALTIME,&stop);    // stop timer
accum=(stop.tv_sec-start.tv_sec)+    // elapsed time
      (stop.tv_nsec-start.tv_nsec)/(double)BILLION;
printf("SVD time: %lf sec.\n",accum); // print elapsed time
cudaStat = cudaMemcpy(U,d_U,sizeof(float)*lda*m,
                      cudaMemcpyDeviceToHost);    // copy d_U->U
cudaStat = cudaMemcpy(VT,d_VT,sizeof(float)*lda*n,
                      cudaMemcpyDeviceToHost);    // copy d_VT->VT
cudaStat = cudaMemcpy(S,d_S,sizeof(float)*n,
                      cudaMemcpyDeviceToHost);    // copy d_S->S
cudaStat = cudaMemcpy(&info_gpu,devInfo,sizeof(int),
                      cudaMemcpyDeviceToHost);    // devInfo->info_gpu
printf("after gesvd: info_gpu = %d\n", info_gpu);
// multiply d_VT by the diagonal matrix corresponding to d_S
cublas_status = cublasSdggm(cublasH,CUBLAS_SIDE_LEFT,n,n,
d_VT, lda, d_S, 1, d_W, lda);    // d_W=d_S*d_VT
cudaStat = cudaMemcpy(d_A,A,sizeof(float)*lda*n,
                      cudaMemcpyHostToDevice);    // copy A->d_A
// compute the difference d_A-d_U*d_S*d_VT
cublas_status=cublasSgemm_v2(cublasH,CUBLAS_OP_N,CUBLAS_OP_N,
m, n, n, &h_minus_one,d_U, lda, d_W, lda, &h_one, d_A, lda);
float dR_fro = 0.0;    // variable for the norm
// compute the norm of the difference d_A-d_U*d_S*d_VT
cublas_status = cublasSnrm2_v2(cublasH,lda*n,d_A,1,&dR_fro);
printf("|A - U*S*VT| = %E \n", dR_fro);    // print the norm
// free memory
cudaFree(d_A);
cudaFree(d_S);
cudaFree(d_U);
cudaFree(d_VT);
cudaFree(devInfo);
cudaFree(d_work);
cudaFree(d_rwork);
cudaFree(d_W);
cublasDestroy(cublasH);
cusolverDnDestroy(cusolverH);
cudaDeviceReset();
return 0;
}

```



```
//SVD time: 18.911288 sec.
//after gesvd: info_gpu = 0
//|A - U*S*VT| = 5.613920E-03
```

3.6.2 cusolverDnSgesvd - unified memory version

```
#include <time.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <cuda_runtime.h>
#include <cublas_v2.h>
#include <cusolverDn.h>
#define BILLION 1000000000L;
int main(int argc, char*argv[])
{
    struct timespec start, stop;           // variables for timing
    double accum;                         // elapsed time variable
    cusolverDnHandle_t cusolverH;         // cusolver handle
    cublasHandle_t cublasH;               // cublas handle
    const int m = 2048;                   // number of rows of A
    const int n = 2048;                   // number of columns of A
    const int lda = m;                    // leading dimension of A
    // declare the factorized matrix A, orthogonal matrices U, VT,
    // sing.val. vector S and auxiliary matr. W = S*VT
    float *A, *A1, *U, *VT, *S, *W;
    cudaMallocManaged(&A, lda*n*sizeof(float)); //unif. mem.for A
    cudaMallocManaged(&A1, lda*n*sizeof(float)); //unif. mem.for A1
    cudaMallocManaged(&U, lda*m*sizeof(float)); //unif. mem.for U
    cudaMallocManaged(&VT, lda*n*sizeof(float)); //unif. mem.for VT
    cudaMallocManaged(&S, n*sizeof(float)); //unified mem.for S
    cudaMallocManaged(&W, lda*n*sizeof(float)); //unif. mem.for W
    for(int i=0; i<lda*n; i++) A[i]=rand()/(float)RAND_MAX;
    int *Info;                             // info for gesvd fun.
    float *work, *rwork;                   // workspace
    int lwork = 0;                         // workspace size
    const float h_one = 1;                 // constants used in SVD checking
    const float h_minus_one = -1;
    // create cusolver and cublas handle
    cusolverDnCreate(&cusolverH);
    cublasCreate(&cublasH);
    cudaMallocManaged(&Info, sizeof(int)); //unified mem.for Info
    cudaMemcpy(A1, A, sizeof(float)*lda*n,
               cudaMemcpyHostToDevice);    // copy A->A1
    // compute buffer size and prepare workspace
    cusolverDnSgesvd_bufferSize(cusolverH, m, n, &lwork);
    cudaMallocManaged(&work, lwork*sizeof(float)); //mem.for work
    // compute the singular value decomposition of A
    // and optionally the left and right singular vectors:
    // A = U*S*VT; the diagonal elements of S
```

```

// are the singular values of A in descending order
// the first min(m,n) columns of U contain the left sing.vec.
// the first min(m,n) cols of VT contain the right sing.vec.
signed char jobu = 'A';          // all m columns of U returned
signed char jobvt = 'A';        // all n columns of d_VT returned
clock_gettime(CLOCK_REALTIME,&start);          // start timer

cusolverDnSgesvd (cusolverH, jobu, jobvt,
m, n, A1, lda, S, U, lda, VT, lda, work,lwork, rwork, Info);

cudaDeviceSynchronize();
clock_gettime(CLOCK_REALTIME,&stop);           // stop timer
accum=(stop.tv_sec-start.tv_sec)+             // elapsed time
      (stop.tv_nsec-start.tv_nsec)/(double)BILLION;
printf("SVD time: %lf sec.\n",accum); // print elapsed time
printf("after gesvd: info = %d\n", *Info);
// multiply VT by the diagonal matrix corresponding to S
cublasSdggmm(cublasH,CUBLAS_SIDE_LEFT,n,n,
             VT, lda, S, 1, W, lda);          // W=S*VT
cudaMemcpy(A1,A,sizeof(float)*lda*n,
           cudaMemcpyHostToDevice);          // copy A->A1
// compute the difference A1-U*S*VT
cublasSgemm_v2(cublasH,CUBLAS_OP_N,CUBLAS_OP_N,
m, n, n, &h_minus_one, U, lda, W, lda, &h_one, A1, lda);
float nrm = 0.0;                          // variable for the norm
// compute the norm of the difference A1-U*S*VT
cublasSnrm2_v2(cublasH,lda*n,A1,1,&nrm);
printf("|A - U*S*VT| = %E \n", nrm);        // print the norm
// free memory
cudaFree(A);
cudaFree(A1);
cudaFree(U);
cudaFree(VT);
cudaFree(S);
cudaFree(W);
cudaFree(Info);
cudaFree(work);
cudaFree(rwork);
cublasDestroy(cublasH);
cusolverDnDestroy(cusolverH);
cudaDeviceReset();
return 0;
}
//SVD time: 19.200704 sec.
//after gesvd: info = 0
//|A - U*S*VT| = 5.613920E-03

```

3.6.3 cusolverDnDgesvd - SVD decomposition in double precision

This function computes in double precision the singular value decomposition of an $m \times n$ matrix:

$$A = u \sigma v^T,$$

where σ is an $m \times n$ matrix which is zero except for its $\min(m, n)$ diagonal elements (singular values), u is an $m \times m$ orthogonal matrix and v is an $n \times n$ orthogonal matrix. The first $\min(m, n)$ columns of u and v are the left and right singular vectors of A respectively.

```
#include <time.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <cuda_runtime.h>
#include <cublas_v2.h>
#include <cusolverDn.h>
#define BILLION 1000000000L;
int main(int argc, char*argv[])
{
    struct timespec start, stop;           // variables for timing
    double accum;                         // elapsed time variable
    cusolverDnHandle_t cusolverH;         // cusolver handle
    cublasHandle_t cublasH;               // cublas handle
    cublasStatus_t cublas_status = CUBLAS_STATUS_SUCCESS;
    cusolverStatus_t cusolver_status = CUSOLVER_STATUS_SUCCESS;
    cudaError_t cudaStat = cudaSuccess;
    const int m = 2048;                   // number of rows of A
    const int n = 2048;                   // number of columns of A
    const int lda = m;                    // leading dimension of A
    // declare the factorized matrix A, orthogonal matrices U, VT
    double *A, *U, *VT, *S; // and sing.val. matrix S on the host
    A=(double*)malloc(lda*n*sizeof(double));
    U=(double*)malloc(lda*m*sizeof(double));
    VT=(double*)malloc(lda*n*sizeof(double));
    S= (double*)malloc(n*sizeof(double));
    for(int i=0; i<lda*n; i++) A[i]=rand()/(double)RAND_MAX;
    // the factorized matrix d_A, orthogonal matrices d_U, d_VT
    double *d_A, *d_U, *d_VT, *d_S; // and sing.val. matrix d_S
    int *devInfo;                     // on the device
    double *d_work, *d_rwork;        // workspace on the device
    double *d_W;                      // auxiliary device array (d_W = d_S*d_VT)
    int lwork = 0;
    int info_gpu = 0;                 // info copied from device to host
    const double h_one = 1;
    const double h_minus_one = -1;
    // create cusolver and cublas handle
    cusolver_status = cusolverDnCreate(&cusolverH);
    cublas_status = cublasCreate(&cublasH);
    // prepare memory on the device
    cudaStat = cudaMalloc((void**)&d_A, sizeof(double)*lda*n);
```

```

    cudaStat = cudaMalloc((void**)&d_S, sizeof(double)*n);
    cudaStat = cudaMalloc((void**)&d_U, sizeof(double)*lda*m);
    cudaStat = cudaMalloc((void**)&d_VT, sizeof(double)*lda*n);
    cudaStat = cudaMalloc((void**)&devInfo, sizeof(int));
    cudaStat = cudaMalloc((void**)&d_W, sizeof(double)*lda*n);
    cudaStat = cudaMemcpy(d_A, A, sizeof(double)*lda*n,
                          cudaMemcpyHostToDevice); // copy A->d_A
// compute buffer size and prepare workspace
    cusolver_status = cusolverDnDgesvd_bufferSize(cusolverH, m, n,
                                                  &lwork );
    cudaStat = cudaMalloc((void**)&d_work, sizeof(double)*lwork);
// compute the singular value decomposition of d_A
// and optionally the left and right singular vectors:
// d_A = d_U*d_S*d_VT; the diagonal elements of d_S
// are the singular values of d_A in descending order
// the first min(m,n) columns of d_U contain the left sing.vec.
// the first min(m,n) cols of d_VT contain the right sing.vec.
    signed char jobu = 'A'; // all m columns of d_U returned
    signed char jobvt = 'A'; // all n columns of d_VT returned
    clock_gettime(CLOCK_REALTIME, &start); // start timer

    cusolver_status = cusolverDnDgesvd (cusolverH, jobu, jobvt,
    m, n, d_A, lda, d_S, d_U, lda, d_VT, lda, d_work, lwork,
    d_rwork, devInfo);

    cudaStat = cudaDeviceSynchronize();
    clock_gettime(CLOCK_REALTIME, &stop); // stop timer
    accum=(stop.tv_sec-start.tv_sec)+ // elapsed time
           (stop.tv_nsec-start.tv_nsec)/(double)BILLION;
    printf("SVD time: %lf sec.\n", accum); // print elapsed time
    cudaStat = cudaMemcpy(U, d_U, sizeof(double)*lda*m,
                          cudaMemcpyDeviceToHost); // copy d_U->U
    cudaStat = cudaMemcpy(VT, d_VT, sizeof(double)*lda*n,
                          cudaMemcpyDeviceToHost); // copy d_VT->VT
    cudaStat = cudaMemcpy(S, d_S, sizeof(double)*n,
                          cudaMemcpyDeviceToHost); // copy d_S->S
    cudaStat = cudaMemcpy(&info_gpu, devInfo, sizeof(int),
                          cudaMemcpyDeviceToHost); // devInfo->info_gpu
    printf("after gesvd: info_gpu = %d\n", info_gpu);
// multiply d_VT by the diagonal matrix corresponding to d_S
    cublas_status = cublasDdggmm(cublasH, CUBLAS_SIDE_LEFT, n, n,
    d_VT, lda, d_S, 1, d_W, lda); // d_W=d_S*d_VT
    cudaStat = cudaMemcpy(d_A, A, sizeof(double)*lda*n,
                          cudaMemcpyHostToDevice); // copy A->d_A
// compute the difference d_A-d_U*d_S*d_VT
    cublas_status=cublasDgemm_v2(cublasH, CUBLAS_OP_N, CUBLAS_OP_N,
    m, n, n, &h_minus_one, d_U, lda, d_W, lda, &h_one, d_A, lda);
    double dR_fro = 0.0; // variable for the norm
// compute the norm of the difference d_A-d_U*d_S*d_VT
    cublas_status = cublasDnrm2_v2(cublasH, lda*n, d_A, 1, &dR_fro);
    printf("|A - U*S*VT| = %E \n", dR_fro); // print the norm
// free memory

```

```

    cudaFree(d_A);
    cudaFree(d_S);
    cudaFree(d_U);
    cudaFree(d_VT);
    cudaFree(devInfo);
    cudaFree(d_work);
    cudaFree(d_rwork);
    cudaFree(d_W);
    cublasDestroy(cublasH);
    cusolverDnDestroy(cusolverH);
    cudaDeviceReset();
    return 0;
}
//SVD time: 22.178122 sec.
//after gesvd: info_gpu = 0
//|A - U*S*VT| = 8.710823E-12

```

3.6.4 cusolverDnDgesvd - unified memory version

```

#include <time.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <cuda_runtime.h>
#include <cublas_v2.h>
#include <cusolverDn.h>
#define BILLION 1000000000L;
int main(int argc, char*argv[])
{
    struct timespec start, stop;           // variables for timing
    double accum;                          // elapsed time variable
    cusolverDnHandle_t cusolverH;         // cusolver handle
    cublasHandle_t cublasH;               // cublas handle
    const int m = 2048;                   // number of rows of A
    const int n = 2048;                   // number of columns of A
    const int lda = m;                    // leading dimension of A
    // declare the factorized matrix A, orthogonal matrices U, VT,
    // sing.val. vector S and auxiliary matr. W = S*VT
    double *A, *A1, *U, *VT, *S, *W;
    cudaMallocManaged(&A, lda*n*sizeof(double)); //unif. mem. for A
    cudaMallocManaged(&A1, lda*n*sizeof(double)); //uni. mem. for A1
    cudaMallocManaged(&U, lda*m*sizeof(double)); //unif. mem. for U
    cudaMallocManaged(&VT, lda*n*sizeof(double)); //uni. mem. for VT
    cudaMallocManaged(&S, n*sizeof(double)); //unified mem. for S
    cudaMallocManaged(&W, lda*n*sizeof(double)); //unif. mem. for W
    for(int i=0; i<lda*n; i++) A[i]=rand()/(double)RAND_MAX;
    int *Info;                             // info for gesvd fun.
    double *work, *rwork;                  // workspace
    int lwork = 0;                          // workspace size
    const double h_one = 1;                // constants used in SVD checking

```

```

    const double h_minus_one = -1;
// create cusolver and cublas handle
    cusolverDnCreate(&cusolverH);
    cublasCreate(&cublasH);
    cudaMallocManaged(&Info, sizeof(int)); //unified mem.for Info
    cudaMemcpy(A1, A, sizeof(double)*lda*n,
               cudaMemcpyHostToDevice); // copy A->A1
// compute buffer size and prepare workspace
    cusolverDnDgesvd_bufferSize(cusolverH, m, n, &lwork );
    cudaMallocManaged(&work, lwork*sizeof(double)); //mem.for work
// compute the singular value decomposition of A
// and optionally the left and right singular vectors:
// A = U*S*VT; the diagonal elements of S
// are the singular values of A in descending order
// the first min(m,n) columns of U contain the left sing.vec.
// the first min(m,n) cols of VT contain the right sing.vec.
    signed char jobu = 'A'; // all m columns of U returned
    signed char jobvt = 'A'; // all n columns of d_VT returned
    clock_gettime(CLOCK_REALTIME, &start); // start timer

    cusolverDnDgesvd (cusolverH, jobu, jobvt, m, n, A1, lda, S, U,
                      lda, VT, lda, work, lwork, rwork, Info);

    cudaDeviceSynchronize();
    clock_gettime(CLOCK_REALTIME, &stop); // stop timer
    accum=(stop.tv_sec-start.tv_sec)+ // elapsed time
           (stop.tv_nsec-start.tv_nsec)/(double)BILLION;
    printf("SVD time: %lf sec.\n", accum); // print elapsed time
    printf("after gesvd: info = %d\n", *Info);
// multiply VT by the diagonal matrix corresponding to S
    cublasDdggmm(cublasH, CUBLAS_SIDE_LEFT, n, n,
                 VT, lda, S, 1, W, lda); // W=S*VT
    cudaMemcpy(A1, A, sizeof(double)*lda*n,
               cudaMemcpyHostToDevice); // copy A->A1
// compute the difference A1-U*S*VT
    cublasDgemm_v2(cublasH, CUBLAS_OP_N, CUBLAS_OP_N,
                  m, n, n, &h_minus_one, U, lda, W, lda, &h_one, A1, lda);
    double nrm = 0.0; // variable for the norm
// compute the norm of the difference A1-U*S*VT
    cublasDnrm2_v2(cublasH, lda*n, A1, 1, &nrm);
    printf("|A - U*S*VT| = %E \n", nrm); // print the norm
// free memory
    cudaFree(A);
    cudaFree(A1);
    cudaFree(U);
    cudaFree(VT);
    cudaFree(S);
    cudaFree(W);
    cudaFree(Info);
    cudaFree(work);
    cudaFree(rwork);
    cublasDestroy(cublasH);

```

```

    cusolverDnDestroy(cusolverH);
    cudaDeviceReset();
    return 0;
}
//SVD time: 22.325408 sec.
//after gesvd: info = 0
//|A - U*S*VT| = 8.710823E-12

```

3.7 Eigenvalues and eigenvectors for symmetric matrices

3.7.1 cusolverDnSsyevd - eigenvalues and eigenvectors for symmetric matrices in single precision

This function computes in single precision all eigenvalues and, optionally, eigenvectors of a real symmetric matrix A . The second parameter can take the values `CUSOLVER_EIG_MODE_VECTOR` or `CUSOLVER_EIG_MODE_NOVECTOR` and answers the question whether the eigenvectors are desired. The symmetric matrix A can be stored in lower (`CUBLAS_FILL_MODE_LOWER`) or upper (`CUBLAS_FILL_MODE_UPPER`) mode. If the eigenvectors are desired, then on exit A contains orthonormal eigenvectors. The eigenvalues are stored in an array W .

```

#include <time.h>
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include <cusolverDn.h>
#define BILLION 1000000000L;
int main(int argc, char*argv[])
{
    struct timespec start,stop;           // variables for timing
    double accum;                         // elapsed time variable
    cusolverDnHandle_t cusolverH;
    cusolverStatus_t cusolver_status = CUSOLVER_STATUS_SUCCESS;
    cudaError_t cudaStat = cudaSuccess;
    const int m = 2048;                  // number of rows and columns of A
    const int lda = m;                   // leading dimension of A
    float *A;                            // mxm matrix
    float *V;                             // mxm matrix of eigenvectors
    float *W;                             // m-vector of eigenvalues
    // prepare memory on the host
    A = (float*)malloc(lda*m*sizeof(float));
    V = (float*)malloc(lda*m*sizeof(float));
    W = (float*)malloc(m*sizeof(float));
    // define random A
    for(int i=0;i<lda*m;i++) A[i] = rand()/(float)RAND_MAX;
    // declare arrays on the device
    float *d_A;                          // mxm matrix A on the device
    float *d_W;                          // m-vector of eigenvalues on the device

```

```

    int *devInfo;                                // info on the device
    float *d_work;                               // workspace on the device
    int lwork = 0;                               // workspace size
    int info_gpu = 0;                            // info copied from device to host
// create cusolver handle
    cusolver_status = cusolverDnCreate(&cusolverH);
// prepare memory on the device
    cudaStat = cudaMalloc ((void**)&d_A, sizeof(float)*lda*m);
    cudaStat = cudaMalloc ((void**)&d_W, sizeof(float)*m);
    cudaStat = cudaMalloc ((void**)&devInfo, sizeof(int));
    cudaStat = cudaMemcpy(d_A, A, sizeof(float)*lda*m,
                          cudaMemcpyHostToDevice); // copy A->d_A
// compute eigenvalues and eigenvectors
    cusolverEigMode_t jobz = CUSOLVER_EIG_MODE_VECTOR;
// use lower left triangle of the matrix
    cublasFillMode_t uplo = CUBLAS_FILL_MODE_LOWER;
// compute buffer size and prepare workspace
    cusolver_status = cusolverDnSsyevd_bufferSize(cusolverH,
                                                  jobz, uplo, m, d_A, lda, d_W, &lwork);
    cudaStat = cudaMalloc((void**)&d_work, sizeof(float)*lwork);
    clock_gettime(CLOCK_REALTIME, &start); // start timer
// compute the eigenvalues and eigenvectors for a symmetric,
// real mxm matrix (only the lower left triangle of A is used)

    cusolver_status = cusolverDnSsyevd(cusolverH, jobz, uplo, m,
                                       d_A, lda, d_W, d_work, lwork, devInfo);

    cudaStat = cudaDeviceSynchronize();
    clock_gettime(CLOCK_REALTIME, &stop); // stop timer
    accum=(stop.tv_sec-start.tv_sec)+ // elapsed time
          (stop.tv_nsec-start.tv_nsec)/((double)BILLION);
    printf("Ssyevd time: %lf sec.\n", accum); // print elapsed time
    cudaStat = cudaMemcpy(W, d_W, sizeof(float)*m,
                          cudaMemcpyDeviceToHost); // copy d_W->W
    cudaStat = cudaMemcpy(V, d_A, sizeof(float)*lda*m,
                          cudaMemcpyDeviceToHost); // copy d_A->V
    cudaStat = cudaMemcpy(&info_gpu, devInfo, sizeof(int),
                          cudaMemcpyDeviceToHost); // copy devInfo->info_gpu
    printf("after syevd: info_gpu = %d\n", info_gpu);
    printf("eigenvalues:\n"); // print first eigenvalues
    for(int i = 0 ; i < 3 ; i++){
        printf("W[%d] = %E\n", i+1, W[i]);
    }
// free memory
    cudaFree(d_A);
    cudaFree(d_W);
    cudaFree(devInfo);
    cudaFree(d_work);
    cusolverDnDestroy(cusolverH);
    cudaDeviceReset();
    return 0;
}

```



```
//Ssyevd time: 2.110875 sec.
//after syevd: info_gpu = 0
//eigenvalues:
//W[1] = -2.582270E+01
//W[2] = -2.566824E+01
//W[3] = -2.563596E+01
```

3.7.2 cusolverDnSsyevd - unified memory version

```
#include <time.h>
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include <cusolverDn.h>
#define BILLION 1000000000L;
int main(int argc, char*argv[])
{
    struct timespec start, stop;           // variables for timing
    double accum;                         // elapsed time variable
    cusolverDnHandle_t cusolverH;
    const int m = 2048;                   // number of rows and columns of A
    const int lda = m;                    // leading dimension of A
    float *A;                             // mxm matrix
    float *W;                             // m-vector of eigenvalues
    // prepare memory
    cudaMallocManaged(&A, lda*m*sizeof(float)); //unif. mem. for A
    cudaMallocManaged(&W, m*sizeof(float));      //unif. mem. for W
    // define random A
    for(int i=0; i<lda*m; i++) A[i] = rand()/(float)RAND_MAX;
    int *Info;                             // info
    float *work;                           // workspace
    int lwork = 0;                         // workspace size
    // create cusolver handle
    cusolverDnCreate(&cusolverH);
    cudaMallocManaged(&Info, sizeof(int)); // unif.mem. for Info
    // compute eigenvalues and eigenvectors
    cusolverEigMode_t jobz = CUSOLVER_EIG_MODE_VECTOR;
    // use lower left triangle of the matrix
    cublasFillMode_t uplo = CUBLAS_FILL_MODE_LOWER;
    // compute buffer size and prepare workspace
    cusolverDnSsyevd_bufferSize(cusolverH,
                                jobz, uplo, m, A, lda, W, &lwork);
    cudaMallocManaged(&work, lwork*sizeof(float)); //mem. for work
    clock_gettime(CLOCK_REALTIME, &start); // start timer
    // compute the eigenvalues and eigenvectors for a symmetric,
    // real mxm matrix (only the lower left triangle of A is used)

    cusolverDnSsyevd(cusolverH, jobz, uplo, m, A, lda, W, work,
                     lwork, Info);
```

```

    cudaDeviceSynchronize();
    clock_gettime(CLOCK_REALTIME,&stop);           // stop timer
    accum=(stop.tv_sec-start.tv_sec)+             // elapsed time
           (stop.tv_nsec-start.tv_nsec)/((double)BILLION;
    printf("syevd time: %lf sec.\n",accum); // print elapsed time
    printf("after syevd: info = %d\n", *Info);
    printf("eigenvalues:\n");                  // print first eigenvalues
    for(int i = 0 ; i < 3 ; i++){
        printf("W[%d] = %E\n", i+1, W[i]);
    }
    // free memory
    cudaFree(A);
    cudaFree(W);
    cudaFree(Info);
    cudaFree(work);
    cusolverDnDestroy(cusolverH);
    cudaDeviceReset();
    return 0;
}
//syevd time: 2.246703 sec.
//after syevd: info = 0
//eigenvalues:
//W[1] = -2.582270E+01
//W[2] = -2.566824E+01
//W[3] = -2.563596E+01

```

3.7.3 cusolverDnDsyeval - eigenvalues and eigenvectors for symmetric matrices in double precision

This function computes in double precision all eigenvalues and, optionally, eigenvectors of a real symmetric matrix A . The second parameter can take the values `CUSOLVER_EIG_MODE_VECTOR` or `CUSOLVER_EIG_MODE_NOVECTOR` and answers the question whether the eigenvectors are desired. The symmetric matrix A can be stored in lower (`CUBLAS_FILL_MODE_LOWER`) or upper (`CUBLAS_FILL_MODE_UPPER`) mode. If the eigenvectors are desired, then on exit A contains orthonormal eigenvectors. The eigenvalues are stored in an array W .

```

#include <time.h>
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include <cusolverDn.h>
#define BILLION 1000000000L;
int main(int argc, char*argv[])
{
    struct timespec start,stop;           // variables for timing
    double accum;                         // elapsed time variable
    cusolverDnHandle_t cusolverH;
    cusolverStatus_t cusolver_status = CUSOLVER_STATUS_SUCCESS;

```

```

    cudaError_t cudaStat = cudaSuccess;
    const int m = 2048;          // number of rows and columns of A
    const int lda = m;          // leading dimension of A
    double *A;                  // mxm matrix
    double *V;                  // mxm matrix of eigenvectors
    double *W;                  // m-vector of eigenvalues
// prepare memory on the host
A = (double*)malloc(lda*m*sizeof(double));
V = (double*)malloc(lda*m*sizeof(double));
W = (double*)malloc(m*sizeof(double));
// define random A
for(int i=0;i<lda*m;i++) A[i] = rand()/((double)RAND_MAX);
// declare arrays on the device
double *d_A;                  // mxm matrix A on the device
double *d_W;                  // m-vector of eigenvalues on the device
int *devInfo;                 // info on the device
double *d_work;               // workspace on the device
int lwork = 0;                // workspace size
int info_gpu = 0;             // info copied from device to host
// create cusolver handle
cusolver_status = cusolverDnCreate(&cusolverH);
// prepare memory on the device
cudaStat = cudaMalloc ((void**)&d_A, sizeof(double)*lda*m);
cudaStat = cudaMalloc ((void**)&d_W, sizeof(double)*m);
cudaStat = cudaMalloc ((void**)&devInfo, sizeof(int));
cudaStat = cudaMemcpy(d_A, A, sizeof(double)*lda*m,
                      cudaMemcpyHostToDevice); // copy A->d_A
// compute eigenvalues and eigenvectors
cusolverEigMode_t jobz = CUSOLVER_EIG_MODE_VECTOR;
// use lower left triangle of the matrix
cublasFillMode_t uplo = CUBLAS_FILL_MODE_LOWER;
// compute buffer size and prepare workspace
cusolver_status = cusolverDnDsyevd_bufferSize(cusolverH,
                                              jobz, uplo, m, d_A, lda, d_W, &lwork);
cudaStat = cudaMalloc((void**)&d_work, sizeof(double)*lwork);
clock_gettime(CLOCK_REALTIME, &start); // start timer
// compute the eigenvalues and eigenvectors for a symmetric,
// real mxm matrix (only the lower left triangle of A is used)

    cusolver_status = cusolverDnDsyevd(cusolverH, jobz, uplo, m,
                                       d_A, lda, d_W, d_work, lwork, devInfo);

    cudaStat = cudaDeviceSynchronize();
    clock_gettime(CLOCK_REALTIME, &stop); // stop timer
    accum=(stop.tv_sec-start.tv_sec)+ // elapsed time
           (stop.tv_nsec-start.tv_nsec)/((double)BILLION;
    printf("Dsyevd time: %lf sec.\n", accum); //print elapsed time
    cudaStat = cudaMemcpy(W, d_W, sizeof(double)*m,
                          cudaMemcpyDeviceToHost); // copy d_W->W
    cudaStat = cudaMemcpy(V, d_A, sizeof(double)*lda*m,
                          cudaMemcpyDeviceToHost); // copy d_A->V
    cudaStat = cudaMemcpy(&info_gpu, devInfo, sizeof(int),

```

```

        cudaMemcpyDeviceToHost); // copy devInfo->info_gpu
printf("after syevd: info_gpu = %d\n", info_gpu);
printf("eigenvalues:\n");           // print first eigenvalues
for(int i = 0 ; i < 3 ; i++){
    printf("W[%d] = %E\n", i+1, W[i]);
}
// free memory
cudaFree(d_A);
cudaFree(d_W);
cudaFree(devInfo);
cudaFree(d_work);
cusolverDnDestroy(cusolverH);
cudaDeviceReset();
return 0;
}
//Dsyevd time: 3.279903 sec.
//after syevd: info_gpu = 0
//eigenvalues:
//W[1] = -2.582273E+01
//W[2] = -2.566824E+01
//W[3] = -2.563596E+01

```

3.7.4 cusolverDnDsyevd - unified memory version

```

#include <time.h>
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include <cusolverDn.h>
#define BILLION 1000000000L;
int main(int argc, char*argv[])
{
    struct timespec start, stop;           // variables for timing
    double accum;                         // elapsed time variable
    cusolverDnHandle_t cusolverH;
    const int m = 2048;                   // number of rows and columns of A
    const int lda = m;                     // leading dimension of A
    double *A;                             // mxm matrix
    double *W;                             // m-vector of eigenvalues
    // prepare memory
    cudaMallocManaged(&A, lda*m*sizeof(double)); //unif.mem.for A
    cudaMallocManaged(&W, m*sizeof(double));      //unif.mem.for W
    // define random A
    for(int i=0; i<lda*m; i++) A[i] = rand()/(double)RAND_MAX;
    int *Info;                             // info
    double *work;                           // workspace
    int lwork = 0;                          // workspace size
    // create cusolver handle
    cusolverDnCreate(&cusolverH);
    cudaMallocManaged(&Info, sizeof(int)); //unified mem. for Info

```

```

// compute eigenvalues and eigenvectors
cusolverEigMode_t jobz = CUSOLVER_EIG_MODE_VECTOR;
// use lower left triangle of the matrix
cublasFillMode_t uplo = CUBLAS_FILL_MODE_LOWER;
// compute buffer size and prepare workspace
cusolverDnDsyevd_bufferSize(cusolverH,
                             jobz, uplo, m, A, lda, W, &lwork);
cudaMallocManaged(&work, lwork*sizeof(double)); // mem. for work
clock_gettime(CLOCK_REALTIME, &start);           // start timer
// compute the eigenvalues and eigenvectors for a symmetric,
// real mxm matrix (only the lower left triangle of A is used)

cusolverDnDsyevd(cusolverH, jobz, uplo, m, A, lda, W, work,
                  lwork, Info);

cudaDeviceSynchronize();
clock_gettime(CLOCK_REALTIME, &stop);           // stop timer
accum=(stop.tv_sec-start.tv_sec)+                // elapsed time
       (stop.tv_nsec-start.tv_nsec)/((double)BILLION;
printf("Dsyevd time: %lf sec.\n", accum); // print elapsed time
printf("after syevd: info = %d\n", *Info);
printf("eigenvalues:\n");                  // print first eigenvalues
for(int i = 0 ; i < 3 ; i++){
    printf("W[%d] = %E\n", i+1, W[i]);
}
// free memory
cudaFree(A);
cudaFree(W);
cudaFree(Info);
cudaFree(work);
cusolverDnDestroy(cusolverH);
cudaDeviceReset();
return 0;
}

//Dsyevd time: 3.395064 sec.
//after syevd: info = 0
//eigenvalues:
//W[1] = -2.582273E+01
//W[2] = -2.566824E+01
//W[3] = -2.563596E+01

```

Chapter 4

MAGMA by example

4.1 General remarks on Magma

MAGMA is an abbreviation for Matrix Algebra for GPU and Multicore Architectures (<http://icl.cs.utk.edu/magma/>). It is a collection of linear algebra routines for dense and sparse matrices. It is a successor of Lapack and ScaLapack, specially developed for heterogeneous CPU-GPU architectures. Magma is an open-source project developed by Innovative Computing Laboratory (ICL), University of Tennessee, Knoxville, USA.

The main ingredients of (dense) Magma are:

- LU, QR and Cholesky factorization.
- Linear solvers based on LU, QR and Cholesky decompositions.
- Eigenvalue and singular value problem solvers.
- Generalized Hermitian-definite eigenproblem solver.
- Mixed-precision iterative refinement solvers based on LU, QR and Cholesky factorizations.

Magma Sparse contains (among other things):

- Sparse linear solvers.
- Sparse eigenvalues.
- Sparse preconditioners.

There is also Magma Batched, which allows for parallel computations on a set of small matrices.

A more detailed information on procedures contained in Magma can be

found in *MAGMA Users's Guide*: <http://icl.cs.utk.edu/projectsfiles/magma/doxygen/>.

Let us notice that the source files in Magma `src` directory contain precise syntax descriptions of Magma functions, so we do not repeat that information in our text (the syntax is also easily available on the Internet).

Instead, we present a series of examples how to use the (dense part of) library.

All subprograms have four versions corresponding to four data types

- `s` - `float` – real single-precision
- `d` - `double` – real double-precision,
- `c` - `magmaFloatComplex` – complex single-precision,
- `z` - `magmaDoubleComplex` – complex double-precision.

To be precise, there exist also some mixed precision routines of the type `sc`, `dz`, `ds`, `zc`, but we have decided to omit the corresponding examples.

- We shall restrict our examples to the most popular real, single and double precision versions. The single precision versions are important because in users hands there are millions of inexpensive GPUs which have restricted double precision capabilities. Installing Magma on such devices can be a good starting point to more advanced studies. On the other hand in many applications the double precision is necessary, so we have decided to present our examples in both versions (in Magma BLAS case only in single precision). In most examples we measure the computations times, so one can compare the performance in single and double precision.
- Ideally we should check for errors on every function call. Unfortunately such an approach doubles the length of our sample codes (which are as short as possible by design). Since our set of Magma sample code (without error checking) is over two hundred pages long, we have decided to ignore the error checking and to focus on the explanations which cannot be found in the syntax description.
- To obtain more compact explanations in our examples we restrict the full generality of Magma to the special case where the leading dimension of matrices is equal to the number of rows and the stride between consecutive elements of vectors is equal to 1. Magma allows for more flexible approach giving the user the access to submatrices and subvectors. The corresponding generalizations can be found in syntax descriptions in source files.

4.1.1 Remarks on installation and compilation

Magma can be downloaded from <http://icl.cs.utk.edu/magma/software/index.html>. In the Magma directory obtained after extraction of the downloaded `magma-X.Y.Z.tar.gz` file there is `README` file which contains installation instructions. The user has to provide `make.inc` which specifies where CUDA, BLAS and LAPACK are installed in the system. Some sample `make.inc` files are contained in Magma directory. After proper modification of the `make.inc` file, running

```
$make
```

creates `libmagma.a` and `libmagma_sparse.a` in Magma `lib` subdirectory and testing drivers in `testing` directory.

The method of compilation of examples depends on the libraries specified in `make.inc`. In the present version of our text we used `Openblas` and Magma directory was a subdirectory of `$HOME` directory. We compiled examples in two steps:

```
g++ -O3 -fopenmp -std=c++11 -DHAVE_CUBLAS -I/usr/local/cuda/  
include -I../include -c -o 001isamax_v2u.o 001isamax_v2u.cpp
```

```
g++ -fopenmp -o 001isamax_v2u 001isamax_v2u.o -L../lib -lm  
-lmagma -L/usr/local/cuda/lib64 -L/usr/lib -lopenblas -lcublas  
-lcudart
```

Let us remark, that only two examples of the present chapter contain the `cudaDeviceSynchronize()` function. The function `magma_sync_wtime` used in the remaining examples contains the synchronization command and `cudaDeviceSynchronize()` is not necessary. Note however that, for example in subsection 4.2.4 (vectors swapping with unified memory), omitting `cudaDeviceSynchronize()` leads to wrong results (vectors are not swapped).

4.1.2 Remarks on hardware used in examples

In most examples we have measured the computations times. The times were obtained on the machine with Ubuntu 16.04, CUDA 8.0, magma-2.2.0 compiled with Openblas library

- Intel(R) Core(TM) i7-6700K CPU, 4.00GHz
- Nvidia(R) GeForce GTX 1080

4.2 Magma BLAS

We restrict ourselves to presentation of the following subset of Magma BLAS single precision functions.

Level 1 BLAS : magma_isamax, magma_sswap,

Level 2 BLAS : magma_sgemv, magma_ssylv,

Level 3 BLAS : magma_sgemm, magma_ssylv, magma_ssylvk, magma_ssylv2k, magma_strmm, magma_sgeadd.

4.2.1 magma_isamax - find element with maximal absolute value

This functions finds the smallest index of the element of an array with the maximum magnitude.

```
#include <stdlib.h>
#include <stdio.h>
#include "magma_v2.h"
int main( int argc, char** argv ){
    magma_init();                                // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create( dev,&queue);
    magma_int_t m = 1024;                        // length of a
    float *a;                                    // a - m-vector on the host
    float *d_a;                                  // d_a - m-vector a on the device
    // allocate array on the host
    magma_smalloc_cpu( &a , m );                // host memory for a
    // allocate array on the device
    magma_smalloc( &d_a, m );                  // device memory for a
    // a={sin(0),sin(1),...,sin(m-1)}
    for(int j=0;j<m;j++) a[j]=sin((float)j);
    // copy data from host to device
    magma_ssetvector(m, a, 1, d_a,1, queue);    // copy a -> d_a
    // find the smallest index of the element of d_a with maximum
    // absolute value

    int i = magma_isamax( m, d_a, 1, queue );

    printf("max |a[i]|: %f\n",fabs(a[i-1]));
    printf("fortran index: %d\n",i);
    magma_free_cpu(a);                          // free host memory
    magma_free(d_a);                            // free device memory
    magma_queue_destroy(queue);
    magma_finalize();
    return 0;
}
// max |a[i]|: 0.999990
// fortran index: 700
```



```

    magma_int_t err;
// allocate the vectors on the host
    err = magma_smalloc_cpu( &a , m );          // host mem. for a
    err = magma_smalloc_cpu( &b , m );          // host mem. for b
// allocate the vector on the device
    err = magma_smalloc( &d_a, m );             // device memory for a
    err = magma_smalloc( &d_b, m );             // device memory for b
// a={sin(0),sin(1),...,sin(m-1)}
    for(int j=0;j<m;j++) a[j]=sin((float)j);
// b={cos(0),cos(1),...,cos(m-1)}
    for(int j=0;j<m;j++) b[j]=cos((float)j);
    printf("a: ");
    for(int j=0;j<4;j++) printf("%6.4f",a[j]);printf("...\n");
    printf("b: ");
    for(int j=0;j<4;j++) printf("%6.4f",b[j]);printf("...\n");
// copy data from host to device
    magma_ssetvector( m, a, 1, d_a,1,queue);    // copy a -> d_a
    magma_ssetvector( m, b, 1, d_b,1,queue);    // copy b -> d_b
// swap the vectors

    magma_sswap( m, d_a, 1, d_b, 1, queue );

    magma_sgetvector( m, d_a, 1, a, 1,queue);    // copy d_a -> a
    magma_sgetvector( m, d_b, 1, b, 1,queue);    // copy d_b -> b
    printf("after magma_sswap:\n");
    printf("a: ");
    for(int j=0;j<4;j++) printf("%6.4f",a[j]);printf("...\n");
    printf("b: ");
    for(int j=0;j<4;j++) printf("%6.4f",b[j]);printf("...\n");
    free(a);                                     // free host memory
    free(b);                                     // free host memory
    magma_free(d_a);                             // free device memory
    magma_free(d_b);                             // free device memory
    magma_queue_destroy(queue);
    magma_finalize();
    return 0;
}
// a: 0.0000,0.8415,0.9093,0.1411,...
// b: 1.0000,0.5403,-0.4161,-0.9900,...
// after magma_sswap:
// a: 1.0000,0.5403,-0.4161,-0.9900,...
// b: 0.0000,0.8415,0.9093,0.1411,...

```

4.2.4 magma_sswap - unified memory version

```

#include <stdlib.h>
#include <stdio.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
int main( int argc, char** argv ){
    magma_init();                                // initialize Magma

```

```

magma_queue_t queue=NULL;
magma_int_t dev=0;
magma_queue_create(dev,&queue);
magma_int_t m = 1024;
float *a; // length of a,b
float *b; // a- m-vector
           // b- m-vector
cudaMallocManaged(&a,m*sizeof(float)); // unif.memory for a
cudaMallocManaged(&b,m*sizeof(float)); // unif.memory for b
           // a={sin(0),sin(1),...,sin(m-1)}
for(int j=0;j<m;j++) a[j]=sin((float)j);
           // b={cos(0),cos(1),...,cos(m-1)}
for(int j=0;j<m;j++) b[j]=cos((float)j);
printf("a: ");
for(int j=0;j<4;j++) printf("%6.4f",a[j]);printf("...\n");
printf("b: ");
for(int j=0;j<4;j++) printf("%6.4f",b[j]);printf("...\n");
// swap the vectors

magma_sswap( m, a, 1, b, 1, queue );

cudaDeviceSynchronize();
printf("after magma_sswap:\n");
printf("a: ");
for(int j=0;j<4;j++) printf("%6.4f",a[j]);printf("...\n");
printf("b: ");
for(int j=0;j<4;j++) printf("%6.4f",b[j]);printf("...\n");
magma_free(a); // free memory
magma_free(b); // free memory
magma_queue_destroy(queue);
magma_finalize();
return 0;
}
// a: 0.0000,0.8415,0.9093,0.1411,...
// b: 1.0000,0.5403,-0.4161,-0.9900,...
// after magma_sswap:
// a: 1.0000,0.5403,-0.4161,-0.9900,...
// b: 0.0000,0.8415,0.9093,0.1411,...

```

4.2.5 magma_sgemv - matrix-vector multiplication

This function performs matrix-vector multiplication

$$c = \alpha \, op(A)b + \beta c,$$

where A is a matrix, b, c are vectors, α, β are scalars and $op(A)$ can be equal to A (MagmaNoTrans case), A^T (transposition) in MagmaTrans case or A^H (conjugate transposition) in MagmaConjTrans case.

```

#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"

```

```

#include "magma_lapack.h"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    real_Double_t gpu_time;
    magma_int_t m = 4096; // number of rows of a
    magma_int_t n = 2048; // number of columns of a
    magma_int_t mn=m*n; // size of a
    float *a; // a- mxn matrix on the host
    float *b; // b- n-vector on the host
    float *c,*c2; // c,c2- m-vectors on the host
    float *d_a; // d_a- mxn matrix a on the device
    float *d_b; // d_b- n-vector b on the device
    float *d_c; // d_c - m-vector on the device
    float alpha = MAGMA_S_MAKE( 1.0, 0.0 ); // alpha=1
    float beta = MAGMA_S_MAKE( 1.0, 0.0 ); // beta=1
    magma_int_t ione = 1; // random uniform distr. in (0,1)
    magma_int_t ISEED[4] = { 0,1,2,3 }; // seed
    magma_int_t err;
    // allocate matrix and vectors on the host
    err = magma_smalloc_pinned( &a , m*n ); // host mem. for a
    err = magma_smalloc_pinned( &b , n ); // host mem. for b
    err = magma_smalloc_pinned( &c , m ); // host mem. for c
    err = magma_smalloc_pinned( &c2, m ); // host mem. for c2
    // allocate matrix and vectors on the device
    err = magma_smalloc( &d_a, m*n ); // device memory for a
    err = magma_smalloc( &d_b, n ); // device memory for b
    err = magma_smalloc( &d_c, m ); // device memory for c
    // generate random matrix a and vectors b,c
    lapackf77_slarnv(&ione,ISEED,&mn,a); // randomize a
    lapackf77_slarnv(&ione,ISEED,&n,b); // randomize b
    lapackf77_slarnv(&ione,ISEED,&m,c); // randomize c
    // copy data from host to device
    magma_ssetmatrix( m, n, a,m,d_a,m,queue); // copy a -> d_a
    magma_ssetvector( n, b, 1, d_b, 1,queue); // copy b -> d_b
    magma_ssetvector( m, c, 1, d_c, 1,queue); // copy c -> d_c
    // matrix-vector multiplication:
    // d_c = alpha*d_a*d_b + beta*d_c;
    // d_a- mxn matrix; b - n-vector; c - m-vector
    gpu_time = magma_sync_wtime(NULL);

    magma_sgemv( MagmaNoTrans,m,n,alpha,d_a,m,d_b,1,beta,d_c,1,
                queue);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("magma_sgemv time: %7.5f sec.\n",gpu_time);
    // copy data from device to host
    magma_sgetvector( m, d_c, 1, c2, 1,queue); // copy d_c -> c2
    printf("after magma_sgemv:\n");
    printf("c2: ");

```

```

    for(int j=0;j<4;j++) printf("%9.4f",c2[j]);
    printf("...\n");
    magma_free_pinned(a);           // free host memory
    magma_free_pinned(b);           // free host memory
    magma_free_pinned(c);           // free host memory
    magma_free_pinned(c2);          // free host memory
    magma_free(d_a);                // free device memory
    magma_free(d_b);                // free device memory
    magma_free(d_c);                // free device memory
    magma_queue_destroy(queue);
    magma_finalize();               // finalize Magma
    return 0;
}
//magma_sgemv time: 0.00016 sec.
//after magma_sgemv:
//c2:  507.9388, 498.1866, 503.1055, 508.1643,...

```

4.2.6 magma_sgemv - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv ){
    magma_init();                       // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    real_Double_t gpu_time;
    magma_int_t m = 4096;               // number of rows of a
    magma_int_t n = 2048;               // number of columns of a
    magma_int_t mn=m*n;                // size of a
    float *a;                          // a- mxn matrix
    float *b;                          // b- n-vector
    float *c;                          // c- m-vector
    float alpha = MAGMA_S_MAKE( 1.0, 0.0 ); // alpha=1
    float beta  = MAGMA_S_MAKE( 1.0, 0.0 ); // beta=1
    magma_int_t ione = 1;               // random uniform distr. in (0,1)
    magma_int_t ISEED[4] = { 0,1,2,3 }; // seed
    cudaMallocManaged(&a,m*n*sizeof(float)); // unif.mem.for a
    cudaMallocManaged(&b,n*sizeof(float));   // unif.mem.for b
    cudaMallocManaged(&c,m*sizeof(float));   // unif.mem.for c
    // generate random matrix a and vectors b,c
    lapackf77_slarnv(&ione,ISEED,&mn,a);      // randomize a
    lapackf77_slarnv(&ione,ISEED,&n,b);        // randomize b
    lapackf77_slarnv(&ione,ISEED,&m,c);        // randomize c
    // matrix-vector multiplication:
    // c = alpha*a*b + beta*c;
    // a- mxn matrix; b - n-vector; c - m-vector

```

```

    gpu_time = magma_sync_wtime(NULL);

    magma_sgemv(MagmaNoTrans,m,n,alpha,a,m,b,1,beta,c,1,queue);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("magma_sgemv time: %7.5f sec.\n",gpu_time);
    printf("after magma_sgemv:\n");
    printf("c: ");
    for(int j=0;j<4;j++) printf("%9.4f",c[j]);
    printf("...\n");
    magma_free(a); // free memory
    magma_free(b); // free memory
    magma_free(c); // free memory
    magma_queue_destroy(queue);
    magma_finalize(); // finalize Magma
    return 0;
}
//magma_sgemv time: 0.00504 sec.
//after magma_sgemv:
//c: 507.9388, 498.1866, 503.1055, 508.1643,...
```

4.2.7 magma_ssymv - symmetric matrix-vector multiplication

This function performs the symmetric matrix-vector multiplication.

$$c = \alpha Ab + \beta c,$$

where A is an $m \times m$ symmetric matrix, b, c are vectors and α, β are scalars. The matrix A can be stored in lower (MagmaLower) or upper (MagmaUpper) mode.

```

#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    real_Double_t gpu_time;
    magma_int_t m = 4096; // number of rows and columns of a
    magma_int_t mm=m*m; // size of a
    float *a; // a- mxm matrix on the host
    // lower triangular part of a contains the lower triangular
    // part of some matrix
    float *b; // b- m-vector on the host
    float *c,*c2; // c,c2- m-vectors on the host
    float *d_a; // d_a- mxm matrix a on the device
    float *d_b; // d_b- m-vector b on the device
    float *d_c; //d_c - m-vector on the device
```

```

float alpha = MAGMA_S_MAKE( 1.0, 0.0 );           // alpha=1
float beta  = MAGMA_S_MAKE( 1.0, 0.0 );           // beta=1
magma_int_t ione = 1;                             //random uniform distr. in (0,1)
magma_int_t ISEED[4] = { 0,1,2,3 };               // seed
magma_int_t err;
// allocate matrix and vectors on the host
err = magma_smalloc_pinned( &a , mm );           // host mem. for a
err = magma_smalloc_pinned( &b , m );             // host mem. for b
err = magma_smalloc_pinned( &c , m );             // host mem. for c
err = magma_smalloc_pinned( &c2, m );             // host mem. for c2
// allocate matrix and vectors on the device
err = magma_smalloc( &d_a, mm );                 // device memory for a
err = magma_smalloc( &d_b, m );                 // device memory for b
err = magma_smalloc( &d_c , m );                 // device memory for c
// generate random matrix a and vectors b,c; only the lower
// triangular part of a is to be referenced
lapackf77_slarnv(&ione,ISEED,&mm,a);             // randomize a
lapackf77_slarnv(&ione,ISEED,&m,b);             // randomize b
lapackf77_slarnv(&ione,ISEED,&m,c);             // randomize c
// copy data from host to device
magma_ssetmatrix( m, m, a, m, d_a,m,queue);      // copy a -> d_a
magma_ssetvector( m, b, 1, d_b,1,queue);         // copy b -> d_b
magma_ssetvector( m, c, 1, d_c,1,queue);         // copy c -> d_c
// symmetric matrix-vector multiplication:
// d_c = alpha*d_a*d_b + beta*d_c;
// d_a- mxm symmetric matrix; b - m-vector; c - m-vector
gpu_time = magma_sync_wtime(NULL);

magma_ssymv( MagmaLower,m,alpha,d_a,m,d_b,1,beta,d_c,1,queue);

gpu_time = magma_sync_wtime(NULL)-gpu_time;
printf("magma_ssymv time: %7.5f sec.\n",gpu_time);
// copy data from device to host
magma_sgetvector( m, d_c, 1, c2,1,queue);        // copy d_c ->c2
printf("after magma_ssymv:\n");
printf("c2: ");
for(int j=0;j<4;j++) printf("%10.4f,",c2[j]);
printf("...\n");
magma_free_pinned(a);                           // free host memory
magma_free_pinned(b);                           // free host memory
magma_free_pinned(c);                           // free host memory
magma_free_pinned(c2);                          // free host memory
magma_free(d_a);                                // free device memory
magma_free(d_b);                                // free device memory
magma_free(d_c);                                // free device memory
magma_queue_destroy(queue);
magma_finalize();                               // finalize Magma
return 0;
}
//magma_ssymv time: 0.00033 sec.
//after magma_ssymv:
//c2: 1003.9608, 1029.2787, 1008.7328, 1042.9585,...
```


4.2.8 magma_ssymv - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    real_Double_t gpu_time;
    magma_int_t m = 4096; // number of rows and columns of a
    magma_int_t mm=m*m; // size of a
    float *a; // a- mxm matrix
    // lower triangular part of a contains the lower triangular
    // part of some matrix
    float *b; // b- m-vector
    float *c; // c- m-vector
    float alpha = MAGMA_S_MAKE( 1.0, 0.0 ); // alpha=1
    float beta = MAGMA_S_MAKE( 1.0, 0.0 ); // beta=1
    magma_int_t ione = 1; //random uniform distr. in (0,1)
    magma_int_t ISEED[4] = { 0,1,2,3 }; // seed
    cudaMallocManaged(&a,mm*sizeof(float)); // unified mem. for a
    cudaMallocManaged(&b,m*sizeof(float)); // unified mem. for b
    cudaMallocManaged(&c,m*sizeof(float)); // unified mem. for c
    // generate random matrix a and vectors b,c; only the lower
    // triangular part of a is to be referenced
    lapackf77_slarnv(&ione,ISEED,&mm,a); // randomize a
    lapackf77_slarnv(&ione,ISEED,&m,b); // randomize b
    lapackf77_slarnv(&ione,ISEED,&m,c); // randomize c
    // symmetric matrix-vector multiplication:
    // c = alpha*a*b + beta*c;
    // a- mxm symmetric matrix; b - m-vector; c - m-vector
    gpu_time = magma_sync_wtime(NULL);

    magma_ssymv( MagmaLower,m,alpha,a,m,b,1,beta,c,1,queue);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("magma_ssymv time: %7.5f sec.\n",gpu_time);
    printf("after magma_ssymv:\n");
    printf("c: ");
    for(int j=0;j<4;j++) printf("%10.4f,",c[j]);
    printf("...\n");
    magma_free(a); // free memory
    magma_free(b); // free memory
    magma_free(c); // free memory
    magma_queue_destroy(queue);
    magma_finalize(); // finalize Magma
    return 0;
}

```

```
//magma_ssymv time: 0.01379 sec.
//after magma_ssymv:
//c: 1003.9608, 1029.2787, 1008.7328, 1042.9585,...
```

4.2.9 magma_sgemm - matrix-matrix multiplication

This function performs the matrix-matrix multiplication

$$C = \alpha op(A)op(B) + \beta C,$$

where A, B, C are matrices and α, β are scalars. The value of $op(A)$ can be equal to A in `MagmaNoTrans` case, A^T (transposition) in `MagmaTrans` case, or A^H (conjugate transposition) in `MagmaConjTrans` case and similarly for $op(B)$.

```
#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    real_Double_t gpu_time;
    magma_int_t m = 8192; // a - mxk matrix
    magma_int_t n = 4096; // b - kxn matrix
    magma_int_t k = 2048; // c - mxn matrix
    magma_int_t mk=m*k; // size of a
    magma_int_t kn=k*n; // size of b
    magma_int_t mn=m*n; // size of c
    float *a; // a- mxk matrix on the host
    float *b; // b- kxn matrix on the host
    float *c; // c- mxn matrix on the host
    float *d_a; // d_a- mxk matrix a on the device
    float *d_b; // d_b- kxn matrix b on the device
    float *d_c; // d_c- mxn matrix c on the device
    float alpha = MAGMA_S_MAKE( 1.0, 0.0 ); // alpha=1
    float beta = MAGMA_S_MAKE( 1.0, 0.0 ); // beta=1
    magma_int_t ione = 1; //random uniform distr. in (0,1)
    magma_int_t ISEED[4] = { 0,1,2,3 }; // seed
    magma_int_t err;
    // allocate matrices on the host
    err = magma_smalloc_pinned( &a , mk ); // host mem. for a
    err = magma_smalloc_pinned( &b , kn ); // host mem. for b
    err = magma_smalloc_pinned( &c , mn ); // host mem. for c
    // allocate matrices and on the device
    err = magma_smalloc( &d_a, mk ); // device memory for a
    err = magma_smalloc( &d_b, kn ); // device memory for b
    err = magma_smalloc( &d_c, mn ); // device memory for c
```

```

// generate random matrices a, b, c;
lapackf77_slarnv(&ione, ISEED, &mk, a);           // randomize a
lapackf77_slarnv(&ione, ISEED, &kn, b);           // randomize b
lapackf77_slarnv(&ione, ISEED, &mn, c);           // randomize c
// copy data from host to device
magma_ssetmatrix( m, k, a, m, d_a, m, queue);     // copy a -> d_a
magma_ssetmatrix( k, n, b, k, d_b, k, queue);     // copy b -> d_b
magma_ssetmatrix( m, n, c, m, d_c, m, queue);     // copy c -> d_c
// matrix-matrix multiplication: d_c = al*d_a*d_b + bet*d_c
// d_a -mxk matrix, d_b -kxn matrix, d_c -mxn matrix;
// al,bet - scalars
gpu_time = magma_sync_wtime(NULL);

magma_sgemm(MagmaNoTrans, MagmaNoTrans, m, n, k, alpha, d_a, m,
            d_b, k, beta, d_c, m, queue);

gpu_time = magma_sync_wtime(NULL) - gpu_time;
printf("magma_sgemm time: %7.5f sec.\n", gpu_time);
// copy data from device to host
magma_sgetmatrix( m, n, d_c, m, c, m, queue); // copy d_c -> c
printf("after magma_sgemm:\n");
printf("c:\n");
for(int i=0; i<4; i++){
for(int j=0; j<4; j++) printf("%10.4f, ", c[i*m+j]);
printf("...\n");
}
printf(".....\n");
magma_free_pinned(a);           // free host memory
magma_free_pinned(b);           // free host memory
magma_free_pinned(c);           // free host memory
magma_free(d_a);                // free device memory
magma_free(d_b);                // free device memory
magma_free(d_c);                // free device memory
magma_queue_destroy(queue);     // destroy queue
magma_finalize();               // finalize Magma
return 0;
}
//magma_sgemm time: 0.01936 sec.
//after magma_sgemm:
//c:
// 498.3723, 521.3933, 507.0844, 515.5119,...
// 504.1406, 517.1718, 509.3519, 511.3415,...
// 511.1694, 530.6165, 517.5001, 524.9462,...
// 505.5946, 522.4631, 511.7729, 516.2770,...
// .....

```

4.2.10 magma_sgemm - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"

```

```

#include "magma_lapack.h"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    real_Double_t gpu_time;
    magma_int_t m = 8192; // a - mxk matrix
    magma_int_t n = 4096; // b - kxn matrix
    magma_int_t k = 2048; // c - mxn matrix
    magma_int_t mk=m*k; // size of a
    magma_int_t kn=k*n; // size of b
    magma_int_t mn=m*n; // size of c
    float *a; // a- mxk matrix
    float *b; // b- kxn matrix
    float *c; // c- mxn matrix
    float alpha = MAGMA_S_MAKE( 1.0, 0.0 ); // alpha=1
    float beta = MAGMA_S_MAKE( 1.0, 0.0 ); // beta=1
    magma_int_t ione = 1; //random uniform distr. in (0,1)
    magma_int_t ISEED[4] = { 0,1,2,3 }; // seed
    cudaMallocManaged(&a,mk*sizeof(float)); // unified mem. for a
    cudaMallocManaged(&b,kn*sizeof(float)); // unified mem. for b
    cudaMallocManaged(&c,mn*sizeof(float)); // unified mem. for c
    // generate random matrices a, b, c;
    lapackf77_slarnv(&ione,ISEED,&mk,a); // randomize a
    lapackf77_slarnv(&ione,ISEED,&kn,b); // randomize b
    lapackf77_slarnv(&ione,ISEED,&mn,c); // randomize c
    // matrix-matrix multiplication: c = al*a*b + bet*c
    // a -mxk matrix, b -kxn matrix, c -mxn matrix;
    // al,bet - scalars
    gpu_time = magma_sync_wtime(NULL);

    magma_sgemm(MagmaNoTrans,MagmaNoTrans,m,n,k,alpha,a,m,b,k,
                beta,c,m,queue);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("magma_sgemm time: %7.5f sec.\n",gpu_time);
    printf("after magma_sgemm:\n");
    printf("c:\n");
    for(int i=0;i<4;i++){
        for(int j=0;j<4;j++) printf("%10.4f,",c[i*m+j]);
        printf("...\n");}
    printf(".....\n");
    magma_free(a); // free memory
    magma_free(b); // free memory
    magma_free(c); // free memory
    magma_queue_destroy(queue); // destroy queue
    magma_finalize(); // finalize Magma
    return 0;
}
//magma_sgemm time: 0.05634 sec.
//after magma_sgemm:

```

```
//c:
// 498.3723, 521.3933, 507.0844, 515.5119,...
// 504.1406, 517.1718, 509.3519, 511.3415,...
// 511.1694, 530.6165, 517.5001, 524.9462,...
// 505.5946, 522.4631, 511.7729, 516.2770,...
//.....
```

4.2.11 magma_ssymm - symmetric matrix-matrix multiplication

This function performs the left or right symmetric matrix-matrix multiplications

$$C = \alpha AB + \beta C \quad \text{in MagmaLeft case,}$$

$$C = \alpha BA + \beta C \quad \text{in MagmaRight case.}$$

The symmetric matrix A has dimension $m \times m$ in the first case and $n \times n$ in the second one. The general matrices B, C have dimensions $m \times n$ and α, β are scalars. The matrix A can be stored in lower (MagmaLower) or upper (MagmaUpper) mode.

```
#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    real_Double_t gpu_time;
    magma_int_t info;
    magma_int_t m = 8192; // a - mxm matrix
    magma_int_t n = 4096; // b,c - mxn matrices
    magma_int_t mm=m*m; // size of a
    magma_int_t mn=m*n; // size of b,c
    float *a; // a- mxm matrix on the host
    float *b; // b- mxn matrix on the host
    float *c; // c- mxn matrix on the host
    float *d_a; // d_a- mxm matrix a on the device
    float *d_b; // d_b- mxn matrix b on the device
    float *d_c; // d_c- mxn matrix c on the device
    float alpha = MAGMA_S_MAKE( 1.0, 0.0 ); // alpha=1
    float beta = MAGMA_S_MAKE( 1.0, 0.0 ); // beta=1
    magma_int_t ione = 1; //random uniform distr. in (0,1)
    magma_int_t ISEED[4] = { 0,1,2,3 }; // seed
    magma_int_t err;
    // allocate matrices on the host
    err = magma_smalloc_pinned( &a , mm ); // host memory for a
    err = magma_smalloc_pinned( &b , mn ); // host memory for b
    err = magma_smalloc_pinned( &c , mn ); // host memory for c
```

```

// allocate matrices on the device
err = magma_smalloc( &d_a, mm ); // device memory for a
err = magma_smalloc( &d_b, mn ); // device memory for b
err = magma_smalloc( &d_c, mn ); // device memory for c
// generate random matrices a, b, c;
lapackf77_slarnv(&ione, ISEED, &mm, a); // randomize a
// lower triangular part of a is the lower triangular part
// of some matrix, the strictly upper triangular
// part of a is not referenced
lapackf77_slarnv(&ione, ISEED, &mn, b); // randomize b
lapackf77_slarnv(&ione, ISEED, &mn, c); // randomize c
// copy data from host to device
magma_ssetmatrix( m, m, a, m, d_a, m, queue ); // copy a -> d_a
magma_ssetmatrix( m, n, b, m, d_b, m, queue ); // copy b -> d_b
magma_ssetmatrix( m, n, c, m, d_c, m, queue ); // copy c -> d_c
// matrix-matrix multiplication: d_c = al*d_a*d_b + bet*d_c
// d_a -mxm symmetric matrix, d_b, d_c -mxn matrices;
// al, bet - scalars
gpu_time = magma_sync_wtime(NULL);

magma_ssymm( MagmaLeft, MagmaLower, m, n, alpha, d_a, m, d_b, m, beta,
             d_c, m, queue);

gpu_time = magma_sync_wtime(NULL) - gpu_time;
printf("magma_ssymm time: %7.5f sec.\n", gpu_time);
// copy data from device to host
magma_sgetmatrix( m, n, d_c, m, c, m, queue ); // copy d_c -> c
printf("after magma_ssymm:\n");
printf("c:\n");
for(int i=0; i<4; i++){
for(int j=0; j<4; j++) printf("%10.4f, ", c[i*m+j]);
printf("...\n");
}
printf(".....\n");
magma_free_pinned(a); // free host memory
magma_free_pinned(b); // free host memory
magma_free_pinned(c); // free host memory
magma_free(d_a); // free device memory
magma_free(d_b); // free device memory
magma_free(d_c); // free device memory
magma_queue_destroy(queue); // destroy queue
magma_finalize(); // finalize Magma
return 0;
}

//magma_ssymm time: 0.20599 sec.
//after magma_ssymm:
//c:
// 2021.3811, 2045.4391, 2048.6990, 2019.2104,...
// 2037.0023, 2050.8364, 2047.5414, 2031.6825,...
// 2053.6797, 2084.0034, 2077.5015, 2068.3196,...
// 2023.3375, 2045.9795, 2051.4314, 2013.8230,...
// .....

```

4.2.12 magma_ssymm - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    real_Double_t gpu_time;
    magma_int_t m = 8192; // a - mxm matrix
    magma_int_t n = 4096; // b,c - mxn matrices
    magma_int_t mm=m*m; // size of a
    magma_int_t mn=m*n; // size of b,c
    float *a; // a- mxm matrix
    float *b; // b- mxn matrix
    float *c; // c- mxn matrix
    float alpha = MAGMA_S_MAKE( 1.0, 0.0 ); // alpha=1
    float beta = MAGMA_S_MAKE( 1.0, 0.0 ); // beta=1
    magma_int_t ione = 1; //random uniform distr. in (0,1)
    magma_int_t ISEED[4] = { 0,1,2,3 }; // seed
    cudaMallocManaged(&a,mm*sizeof(float)); // unified mem.for a
    cudaMallocManaged(&b,mn*sizeof(float)); // unified mem.for b
    cudaMallocManaged(&c,mn*sizeof(float)); // unified mem.for c
    // generate random matrices a, b, c;
    lapackf77_slarnv(&ione,ISEED,&mm,a); // randomize a
    // lower triangular part of a is the lower triangular part
    // of some matrix, the strictly upper triangular
    // part of a is not referenced
    lapackf77_slarnv(&ione,ISEED,&mn,b); // randomize b
    lapackf77_slarnv(&ione,ISEED,&mn,c); // randomize c
    // matrix-matrix multiplication: c = al*a*b + bet*c
    // a -mxm symmetric matrix, b, c -mxn matrices;
    // al,bet - scalars
    gpu_time = magma_sync_wtime(NULL);

    magma_ssymm( MagmaLeft,MagmaLower,m,n,alpha,a,m,b,m,beta,c,m,
                queue);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("magma_ssymm time: %7.5f sec.\n",gpu_time);
    printf("after magma_ssymm:\n");
    printf("c:\n");
    for(int i=0;i<4;i++){
        for(int j=0;j<4;j++) printf("%10.4f",c[i*m+j]);
        printf("...\n");
    }
    printf(".....\n");
    magma_free(a); // free memory
    magma_free(b); // free memory
    magma_free(c); // free memory
}

```

```

    magma_queue_destroy(queue);           // destroy queue
    magma_finalize();                     // finalize Magma
    return 0;
}
//magma_ssymm time: 0.27643 sec.
//after magma_ssymm:
//c:
// 2021.3811, 2045.4391, 2048.6990, 2019.2104,...
// 2037.0023, 2050.8364, 2047.5414, 2031.6825,...
// 2053.6797, 2084.0034, 2077.5015, 2068.3196,...
// 2023.3375, 2045.9795, 2051.4314, 2013.8230,...
//.....

```

4.2.13 magma_ssyrrk - symmetric rank-k update

This function performs the symmetric rank-k update

$$C = \alpha op(A)op(A)^T + \beta C,$$

where $op(A)$ is an $m \times k$ matrix, C is a symmetric $m \times m$ matrix stored in lower (MagmaLower) or upper (MagmaUpper) mode and α, β are scalars. The value of $op(A)$ can be equal to A in MagmaNoTrans case or A^T (transposition) in MagmaTrans case.

```

#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv ){
    magma_init();                               // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    real_Double_t gpu_time;
    magma_int_t info;
    magma_int_t m = 8192;                       // a - mxk matrix
    magma_int_t k = 4096;                       // c - mxm matrix
    magma_int_t mm=m*m;                         // size of c
    magma_int_t mk=m*k;                         // size of a
    float *a;                                  // a- mxk matrix on the host
    float *c;                                  // c- mxm matrix on the host
    float *d_a;                                // d_a- mxk matrix a on the device
    float *d_c;                                // d_c- mxm matrix c on the device
    float alpha = 1.0;                         // alpha=1
    float beta = 1.0;                          // beta=1
    magma_int_t ione = 1;                      //random uniform distr. in (0,1)
    magma_int_t ISEED[4] = { 0,1,2,3 };       // seed
    magma_int_t err;
    // allocate matrices on the host
    err = magma_smallocc_pinned( &a , mk );    // host memory for a

```



```

    err = magma_smalloc_pinned( &c , mm ); // host memory for c
// allocate matrices on the device
    err = magma_smalloc( &d_a, mk ); // device memory for a
    err = magma_smalloc( &d_c, mm ); // device memory for c
// generate random matrices a, c;
    lapackf77_slarnv(&ione, ISEED, &mk, a); // randomize a
    lapackf77_slarnv(&ione, ISEED, &mm, c); // randomize c
// upper triangular part of c is the upper triangular part
// of some matrix, the strictly lower triangular
// part of c is not referenced
// copy data from host to device
    magma_ssetmatrix( m, k, a, m, d_a, m, queue); // copy a -> d_a
    magma_ssetmatrix( m, m, c, m, d_c, m, queue); // copy c -> d_c
// symmetric rank-k update: d_c=alpha*d_a*d_a^T+beta*d_c
// d_c -mxm symmetric matrix, d_a -mxk matrix;
// alpha,beta - scalars
    gpu_time = magma_sync_wtime(NULL);

    magma_ssyrrk( MagmaUpper,MagmaNoTrans,m,k,alpha,d_a,m,beta,
                  d_c,m,queue);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("magma_ssyrrk time: %7.5f sec.\n",gpu_time);
// copy data from device to host
    magma_sgetmatrix( m, m, d_c, m,c,m,queue); // copy d_c -> c
    printf("after magma_ssyrrk:\n");
    printf("c:\n");
    for(int i=0;i<4;i++){
    for(int j=0;j<4;j++) if(i>=j) printf("%10.4f",c[i*m+j]);
    printf("...\n");}
    printf(".....\n");
    magma_free_pinned(a); // free host memory
    magma_free_pinned(c); // free host memory
    magma_free(d_a); // free device memory
    magma_free(d_c); // free device memory
    magma_queue_destroy(queue); // destroy queue
    magma_finalize(); // finalize Magma
    return 0;
}
//magma_ssyrrk time: 0.03725 sec.
//after magma_ssyrrk:
//c:
// 1358.9562,...
// 1027.0094, 1382.1946,...
// 1011.2416, 1022.4153, 1351.7262,...
// 1021.8580, 1037.6437, 1025.0333, 1376.4917,...
//.....

```

4.2.14 magma_ssyk - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    real_Double_t gpu_time;
    magma_int_t m = 8192; // a - mxk matrix
    magma_int_t k = 4096; // c - mxm matrix
    magma_int_t mm=m*m; // size of c
    magma_int_t mk=m*k; // size of a
    float *a; // a- mxk matrix
    float *c; // c- mxm matrix
    float alpha = 1.0; // alpha=1
    float beta = 1.0; // beta=1
    magma_int_t ione = 1; //random uniform distr. in (0,1)
    magma_int_t ISEED[4] = { 0,1,2,3 }; // seed
    cudaMallocManaged(&a,mk*sizeof(float)); // unified mem.for a
    cudaMallocManaged(&c,mm*sizeof(float)); // unified mem.for c
    // generate random matrices a, c;
    lapackf77_slarnv(&ione,ISEED,&mk,a); // randomize a
    lapackf77_slarnv(&ione,ISEED,&mm,c); // randomize c
    // upper triangular part of c is the upper triangular part
    // of some matrix, the strictly lower triangular
    // part of c is not referenced
    // symmetric rank-k update: c=alpha*a*a^T+beta*c
    // c -mxm symmetric matrix, a -mxk matrix;
    // alpha,beta - scalars
    gpu_time = magma_sync_wtime(NULL);

    magma_ssyk( MagmaUpper,MagmaNoTrans,m,k,alpha,a,m,beta,
               c,m,queue);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("magma_ssyk time: %.5f sec.\n",gpu_time);
    printf("after magma_ssyk:\n");
    printf("c:\n");
    for(int i=0;i<4;i++){
        for(int j=0;j<4;j++) if(i>=j) printf("%10.4f",c[i*m+j]);
        printf("...\n");
    }
    printf(".....\n");
    magma_free(a); // free memory
    magma_free(c); // free memory
    magma_queue_destroy(queue); // destroy queue
    magma_finalize(); // finalize Magma
    return 0;

```

```

}
//magma_ssyrrk time: 0.09162 sec.
//after magma_ssyrrk:
//c:
// 1358.9562,...
// 1027.0094, 1382.1946,...
// 1011.2416, 1022.4153, 1351.7262,...
// 1021.8580, 1037.6437, 1025.0333, 1376.4917,...
//.....

```

4.2.15 magma_ssyrr2k - symmetric rank-2k update

This function performs the symmetric rank-2k update

$$C = \alpha(op(A)op(B)^T + op(B)op(A)^T) + \beta C,$$

where $op(A)$, $op(B)$ are $m \times k$ matrices, C is a symmetric $m \times m$ matrix stored in lower (MagmaLower) or upper (MagmaUpper) mode and α, β are scalars. The value of $op(A)$ can be equal to A in MagmaNoTrans case or A^T (transposition) in MagmaTrans case and similarly for $op(B)$.

```

#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    real_Double_t gpu_time;
    magma_int_t info;
    magma_int_t m = 8192; // a,b - mxk matrices
    magma_int_t k = 4096; // c - mxm matrix
    magma_int_t mm=m*m; // size of c
    magma_int_t mk=m*k; // size of a
    float *a; // a- mxk matrix on the host
    float *b; // b- mxk matrix on the host
    float *c; // c- mxm matrix on the host
    float *d_a; // d_a- mxk matrix a on the device
    float *d_b; // d_b- mxk matrix a on the device
    float *d_c; // d_c- mxm matrix c on the device
    float alpha = 1.0; // alpha=1
    float beta = 1.0; // beta=1
    magma_int_t ione = 1; //random uniform distr. in (0,1)
    magma_int_t ISEED[4] = { 0,1,2,3 }; // seed
    magma_int_t err;
    // allocate matrices on the host
    err = magma_smallocc_pinned( &a , mk ); // host memory for a
    err = magma_smallocc_pinned( &b , mk ); // host memory for b

```

```

    err = magma_smalloc_pinned( &c , mm ); // host memory for c
// allocate matrices on the device
    err = magma_smalloc( &d_a, mk ); // device memory for a
    err = magma_smalloc( &d_b, mk ); // device memory for b
    err = magma_smalloc( &d_c, mm ); // device memory for c
// generate random matrices a,b,c;
    lapackf77_slarnv(&ione, ISEED, &mk, a); // randomize a
    lapackf77_slarnv(&ione, ISEED, &mk, b); // randomize b
    lapackf77_slarnv(&ione, ISEED, &mm, c); // randomize c
// upper triangular part of c is the upper triangular part
// of some matrix, the strictly lower triangular
// part of c is not referenced
// copy data from host to device
    magma_ssetmatrix( m, k, a, m, d_a, m, queue); // copy a -> d_a
    magma_ssetmatrix( m, k, a, m, d_b, m, queue); // copy b -> d_b
    magma_ssetmatrix( m, m, c, m, d_c, m, queue); // copy c -> d_c
// symmetric rank-2k update:
// d_c=alpha*d_a*d_b^T+ alpha*d_b*d_a^T+beta*d_c
// d_c -mxm symmetric matrix, d_a,d_b -mxk matrices;
// alpha,beta - scalars
    gpu_time = magma_sync_wtime(NULL);

    magma_ssyr2k( MagmaUpper,MagmaNoTrans,m,k,alpha,d_a,m,d_b,m,
                 beta,d_c,m,queue);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("magma_ssyr2k time: %7.5f sec.\n",gpu_time);
// copy data from device to host
    magma_sgetmatrix( m, m, d_c, m, c, m, queue); // copy d_c -> c
    printf("after magma_ssyr2k:\n");
    printf("c:\n");
    for(int i=0;i<4;i++){
    for(int j=0;j<4;j++) if(i>=j) printf("%10.4f",c[i*m+j]);
    printf("...\n");}
    printf(".....\n");
    magma_free_pinned(a); // free host memory
    magma_free_pinned(c); // free host memory
    magma_free(d_a); // free device memory
    magma_free(d_c); // free device memory
    magma_queue_destroy(queue); // destroy queue
    magma_finalize(); // finalize Magma
    return 0;
}
//magma_ssyr2k time: 0.07446 sec.
//after magma_ssyr2k:
//c:
// 2718.7930,...
// 2054.1855, 2763.3325,...
// 2022.0312, 2043.4248, 2702.5745,...
// 2043.3660, 2075.6743, 2048.9951, 2753.3296,...
//.....

```

4.2.16 magma_ssy2k - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    real_Double_t gpu_time;
    magma_int_t m = 8192; // a,b - mxk matrices
    magma_int_t k = 4096; // c - mxm matrix
    magma_int_t mm=m*m; // size of c
    magma_int_t mk=m*k; // size of a
    float *a; // a- mxk matrix
    float *b; // b- mxk matrix
    float *c; // c- mxm matrix
    float alpha = 1.0; // alpha=1
    float beta = 1.0; // beta=1
    magma_int_t ione = 1; //random uniform distr. in (0,1)
    magma_int_t ISEED[4] = { 0,1,2,3 }; // seed
    cudaMallocManaged(&a,mk*sizeof(float)); // unified mem.for a
    cudaMallocManaged(&b,mk*sizeof(float)); // unified mem.for b
    cudaMallocManaged(&c,mm*sizeof(float)); // unified mem.for c
    // generate random matrices a,b,c;
    lapackf77_slarnv(&ione,ISEED,&mk,a); // randomize a
    lapackf77_slarnv(&ione,ISEED,&mk,b); // randomize b
    lapackf77_slarnv(&ione,ISEED,&mm,c); // randomize c
    // upper triangular part of c is the upper triangular part
    // of some matrix, the strictly lower triangular
    // part of c is not referenced

    // symmetric rank-2k update:
    // c=alpha*a*b^T+ alpha*b*a^T+beta*c
    // c -mxm symmetric matrix, a,b -mxk matrices;
    // alpha,beta - scalars
    gpu_time = magma_sync_wtime(NULL);

    magma_ssy2k( MagmaUpper,MagmaNoTrans,m,k,alpha,a,m,b,m,
                beta,c,m,queue);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("magma_ssy2k time: %7.5f sec.\n",gpu_time);
    printf("after magma_ssy2k:\n");
    printf("c:\n");
    for(int i=0;i<4;i++){
    for(int j=0;j<4;j++) if(i>=j) printf("%10.4f",c[i*m+j]);
    printf("...\n");}
    printf(".....\n");

```

```

    magma_free(a);                // free memory
    magma_free(b);                // free memory
    magma_free(c);                // free memory
    magma_queue_destroy(queue);   // destroy queue
    magma_finalize();             // finalize Magma
    return 0;
}
//magma_ssyrr2k time: 0.13833 sec.
//after magma_ssyrr2k:
//c:
// 2047.3660,...
// 2044.8237, 2041.2444,...
// 2041.6855, 2038.5908, 2023.9705,...
// 2050.2649, 2057.5630, 2046.7908, 2059.3657,...
//.....

```

4.2.17 magma_strmm - triangular matrix-matrix multiplication

This function performs the left or right triangular matrix-matrix multiplications

$$\begin{aligned}
 C &= \alpha \operatorname{op}(A) B && \text{in MagmaLeft case,} \\
 C &= \alpha B \operatorname{op}(A) && \text{in MagmaRight case,}
 \end{aligned}$$

where A is a triangular matrix, C, B are $m \times n$ matrices and α is a scalar. The value of $\operatorname{op}(A)$ can be equal to A in `MagmaNoTrans` case, A^T (transposition) in `MagmaTrans` case or A^H (conjugate transposition) in `MagmaConjTrans` case. A has dimension $m \times m$ in the first case and $n \times n$ in the second case. A can be stored in lower (`MagmaLower`) or upper (`MagmaUpper`) mode. If the diagonal of the matrix A has non-unit elements, then the parameter `MagmaNonUnit` should be used (in the opposite case - `MagmaUnit`).

```

#include <stdio.h>
#include <stdlib.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv ){
    magma_init();                // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    real_Double_t gpu_time;
    magma_int_t info;
    magma_int_t m = 8192;        // a - mxm matrix
    magma_int_t n = 4096;        // c - mxn matrix
    magma_int_t mm=m*m;          // size of a
    magma_int_t mn=m*n;          // size of c
    float *a;                    // a- mxm matrix on the host

```

```

float *c;                                // c- mxn matrix on the host
float *d_a;                              // d_a- mxm matrix a on the device
float *d_c;                              // d_c- mxn matrix c on the device
float alpha = 1.0;                        // alpha=1
magma_int_t ione = 1;
magma_int_t ISEED[4] = { 0,1,2,3 };      // seed
magma_int_t err;
// allocate matrices on the host
err = magma_smalloc_pinned( &a , mm );   // host memory for a
err = magma_smalloc_pinned( &c , mn );   // host memory for c
// allocate matrices on the device
err = magma_smalloc( &d_a, mm );        // device memory for a
err = magma_smalloc( &d_c, mn );        // device memory for c
// generate random matrices a, c;
lapackf77_slarnv(&ione, ISEED, &mm, a); // randomize a
lapackf77_slarnv(&ione, ISEED, &mn, c); // randomize c
// upper triangular part of a is the upper triangular part
// of some matrix, the strictly lower
// triangular part of a is not referenced
// copy data from host to device
magma_ssetmatrix( m, m, a, m, d_a, m, queue); // copy a -> d_a
magma_ssetmatrix( m, n, c, m, d_c, m, queue); // copy c -> d_c
// triangular matrix-matrix multiplication
// d_c=alpha*d_a*d_c
// d_c -mxn matrix, d_a -mxm triangular matrix;
// alpha - scalar
gpu_time = magma_sync_wtime(NULL);

magma_strmm(MagmaLeft, MagmaUpper, MagmaNoTrans, MagmaNonUnit,
            m, n, alpha, d_a, m, d_c, m, queue);

gpu_time = magma_sync_wtime(NULL) - gpu_time;
printf("magma_strmm time: %7.5f sec.\n", gpu_time);
// copy data from device to host
magma_sgetmatrix( m, n, d_c, m, c, m, queue); // copy d_c -> c
printf("after magma_strmm:\n");
printf("c:\n");
for(int i=0; i<4; i++){
for(int j=0; j<4; j++) printf("%10.4f", c[i*m+j]);
printf("...\n");}
printf(".....\n");
magma_free_pinned(a); // free host memory
magma_free_pinned(c); // free host memory
magma_free(d_a);      // free device memory
magma_free(d_c);      // free device memory
magma_queue_destroy(queue); // destroy queue
magma_finalize();      // finalize Magma
return 0;
}
//magma_strmm time: 0.04829 sec.
//after magma_strmm:
//c:

```

```
// 2051.0024, 2038.8608, 2033.2482, 2042.2589,...
// 2040.4783, 2027.2789, 2025.2496, 2041.6721,...
// 2077.4158, 2052.2390, 2050.5039, 2074.0791,...
// 2028.7070, 2034.3572, 2003.8625, 2031.4501,...
//.....
```

4.2.18 magma_strmm - unified memory version

```
#include <stdio.h>
#include <stdlib.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    real_Double_t gpu_time;
    magma_int_t m = 8192; // a - mxm matrix
    magma_int_t n = 4096; // c - mxn matrix
    magma_int_t mm=m*m; // size of a
    magma_int_t mn=m*n; // size of c
    float *a; // a- mxm matrix
    float *c; // c- mxn matrix
    float alpha = 1.0; // alpha=1
    magma_int_t ione = 1;
    magma_int_t ISEED[4] = { 0,1,2,3 }; // seed
    cudaMallocManaged(&a,mm*sizeof(float)); // unified mem.for a
    cudaMallocManaged(&c,mn*sizeof(float)); // unified mem.for c
    // generate random matrices a, c;
    lapackf77_slarnv(&ione,ISEED,&mm,a); // randomize a
    lapackf77_slarnv(&ione,ISEED,&mn,c); // randomize c
    // upper triangular part of a is the upper triangular part
    // of some matrix, the strictly lower
    // triangular part of a is not referenced

    // triangular matrix-matrix multiplication c=alpha*a*c
    // c -mxn matrix, a -mxm triangular matrix; alpha - scalar
    gpu_time = magma_sync_wtime(NULL);

    magma_strmm(MagmaLeft,MagmaUpper,MagmaNoTrans,MagmaNonUnit,
                m,n,alpha,a,m,c,m,queue);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("magma_strmm time: %7.5f sec.\n",gpu_time);
    printf("after magma_strmm:\n");
    printf("c:\n");
    for(int i=0;i<4;i++){
```



```

    for(int j=0;j<4;j++) if(i>=j) printf("%10.4f",c[i*m+j]);
    printf("...\n");}
    printf(".....\n");
    magma_free(a); // free memory
    magma_free(c); // free memory
    magma_queue_destroy(queue); // destroy queue
    magma_finalize(); // finalize Magma
    return 0;
}
//magma_strmm time: 0.12141 sec.
//after magma_strmm:
//c:
// 2051.0024, 2038.8608, 2033.2482, 2042.2589,...
// 2040.4783, 2027.2789, 2025.2496, 2041.6721,...
// 2077.4158, 2052.2390, 2050.5039, 2074.0791,...
// 2028.7070, 2034.3572, 2003.8625, 2031.4501,...
//.....

```

4.2.19 magmablas_sgeadd - matrix-matrix addition

This function performs the addition of matrices

$$C = \alpha A + C,$$

where A, C are $m \times n$ matrices and α is a scalar.

```

#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    real_Double_t gpu_time;
    magma_int_t m = 8192; // a - mxn matrix
    magma_int_t n = 4096; // c - mxn matrix
    magma_int_t mn=m*n; // size of c
    float *a; // a- mxn matrix on the host
    float *c; // c- mxn matrix on the host
    float *d_a; // d_a- mxn matrix a on the device
    float *d_c; // d_c- mxn matrix c on the device
    float alpha = 2.0; // alpha=2
    magma_int_t ione = 1;
    magma_int_t ISEED[4] = { 0,1,2,3 }; // seed
    magma_int_t err;
    // allocate matrices the host
    err = magma_smalloc_pinned( &a , mn ); // host memory for a
    err = magma_smalloc_pinned( &c , mn ); // host memory for c
    // allocate matrices on the device

```

```

    err = magma_smalloc( &d_a,  mn );          // device memory for a
    err = magma_smalloc( &d_c,  mn );          // device memory for c
// generate random matrices a, c;
    lapackf77_slarnv(&ione, ISEED, &mn, a);      // randomize a
    lapackf77_slarnv(&ione, ISEED, &mn, c);      // randomize c
    printf("a:\n");
    for(int i=0; i<4; i++){
    for(int j=0; j<4; j++)    printf("%10.4f", a[i*m+j]);
    printf("...\n");}
    printf(".....\n");
    printf("c:\n");
    for(int i=0; i<4; i++){
    for(int j=0; j<4; j++)    printf("%10.4f", c[i*m+j]);
    printf("...\n");}
    printf(".....\n");
// copy data from host to device
    magma_ssetmatrix( m, n, a, m, d_a, m, queue); // copy a -> d_a
    magma_ssetmatrix( m, n, c, m, d_c, m, queue); // copy c -> d_c
// d_c=alpha*d_a+d_c
// d_a, d_c -mxn matrices;
// alpha - scalar
    gpu_time = magma_sync_wtime(NULL);

    magmablas_sgeadd(m,n,alpha,d_a,m,d_c,m,queue);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("magmablas_sgeadd time: %7.5f sec.\n",gpu_time);
// copy data from device to host
    magma_sgetmatrix( m, n, d_c, m, c, m, queue); // copy d_c -> c
    printf("after magmablas_sgeadd:\n");
    printf("c:\n");
    for(int i=0; i<4; i++){
    for(int j=0; j<4; j++)    printf("%10.4f", c[i*m+j]);
    printf("...\n");}
    printf(".....\n");
    magma_free_pinned(a);          // free host memory
    magma_free_pinned(c);          // free host memory
    magma_free(d_a);               // free device memory
    magma_free(d_c);               // free device memory
    magma_queue_destroy(queue);    // destroy queue
    magma_finalize();              // finalize Magma
    return 0;
}
//a:
//    0.1319,    0.2338,    0.3216,    0.7105,...
//    0.6137,    0.0571,    0.4461,    0.8876,...
//    0.5486,    0.9655,    0.8833,    0.8968,...
//    0.5615,    0.0839,    0.2581,    0.8629,...
//.....
//c:

```

```
//      0.0443,      0.4490,      0.8054,      0.1554,...
//      0.1356,      0.5692,      0.6642,      0.2544,...
//      0.6798,      0.7744,      0.8358,      0.1854,...
//      0.3021,      0.1897,      0.9450,      0.0734,...
//.....
//magmablas_sgeadd time: 0.00174 sec.
//after magmablas_sgeadd:
//c:
//      0.3080,      0.9166,      1.4487,      1.5765,...
//      1.3630,      0.6835,      1.5565,      2.0297,...
//      1.7771,      2.7055,      2.6023,      1.9789,...
//      1.4252,      0.3575,      1.4612,      1.7992,...
//.....
```

4.2.20 magmablas_sgeadd - unified memory version

```
#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    real_Double_t gpu_time;
    magma_int_t m = 8192; // a - mxn matrix
    magma_int_t n = 4096; // c - mxn matrix
    magma_int_t mn=m*n; // size of c
    float *a; // a- mxn matrix
    float *c; // c- mxn matrix
    float alpha = 2.0; // alpha=2
    magma_int_t ione = 1;
    magma_int_t ISEED[4] = { 0,1,2,3 }; // seed
    cudaMallocManaged(&a,mn*sizeof(float)); // unified mem.for a
    cudaMallocManaged(&c,mn*sizeof(float)); // unified mem.for c
    // generate random matrices a, c;
    lapackf77_slarnv(&ione,ISEED,&mn,a); // randomize a
    lapackf77_slarnv(&ione,ISEED,&mn,c); // randomize c
    printf("a:\n");
    for(int i=0;i<4;i++){
        for(int j=0;j<4;j++) printf("%10.4f",a[i*m+j]);
        printf("...\n");}
    printf(".....\n");
    printf("c:\n");
    for(int i=0;i<4;i++){
        for(int j=0;j<4;j++) printf("%10.4f",c[i*m+j]);
        printf("...\n");}
    printf(".....\n");
```

```

// c=alpha*a+c;   a, c -mxn matrices;   alpha - scalar
gpu_time = magma_sync_wtime(NULL);

magma_blas_sgeadd(m,n,alpha,a,m,c,m,queue);

gpu_time = magma_sync_wtime(NULL)-gpu_time;
printf("magma_blas_sgeadd time: %7.5f sec.\n",gpu_time);
printf("after magma_blas_sgeadd:\n");
printf("c:\n");
for(int i=0;i<4;i++){
for(int j=0;j<4;j++) printf("%10.4f",c[i*m+j]);
printf("...\n");}
printf(".....\n");
magma_free(a); // free memory
magma_free(c); // free memory
magma_queue_destroy(queue); // destroy queue
magma_finalize(); // finalize Magma
return 0;
}
//a:
//   0.1319,    0.2338,    0.3216,    0.7105,...
//   0.6137,    0.0571,    0.4461,    0.8876,...
//   0.5486,    0.9655,    0.8833,    0.8968,...
//   0.5615,    0.0839,    0.2581,    0.8629,...
//.....
//c:
//   0.0443,    0.4490,    0.8054,    0.1554,...
//   0.1356,    0.5692,    0.6642,    0.2544,...
//   0.6798,    0.7744,    0.8358,    0.1854,...
//   0.3021,    0.1897,    0.9450,    0.0734,...
//.....
//magma_blas_sgeadd time: 0.03860 sec.
//after magma_blas_sgeadd:
//c:
//   0.3080,    0.9166,    1.4487,    1.5765,...
//   1.3630,    0.6835,    1.5565,    2.0297,...
//   1.7771,    2.7055,    2.6023,    1.9789,...
//   1.4252,    0.3575,    1.4612,    1.7992,...
//.....

```

4.3 LU decomposition and solving general linear systems

4.3.1 magma_sgesv - solve a general linear system in single precision, CPU interface

This function solves in single precision a general real, linear system

$$A X = B,$$

where A is an $m \times m$ matrix and X, B are $m \times n$ matrices. A, B are defined on the host. In the solution, the LU decomposition of A with partial pivoting and row interchanges is used. See [magma-X.Y.Z/src/sgesv.cpp](#) for more details.

```
#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    real_Double_t gpu_time;
    magma_int_t *piv, info;
    magma_int_t m = 8192; // a - mxm matrix
    magma_int_t n = 100; // c - mxn matrix
    magma_int_t mm=m*m; // size of a
    magma_int_t mn=m*n; // size of c
    float *a; // a- mxm matrix on the host
    float *b; // b- mxn matrix on the host
    float *c; // c- mxn matrix on the host
    magma_int_t ione = 1; //random uniform distr. in (0,1)
    magma_int_t ISEED[4] = { 0,0,0,1 }; // seed
    magma_int_t err;
    const float alpha = 1.0; // alpha=1
    const float beta = 0.0; // beta=0
    // allocate matrices on the host
    err = magma_smalloc_pinned( &a , mm ); // host memory for a
    err = magma_smalloc_pinned( &b , mn ); // host memory for b
    err = magma_smalloc_pinned( &c , mn ); // host memory for c
    piv=(magma_int_t*)malloc(m*sizeof(magma_int_t));
    // generate random matrices a, b;
    lapackf77_slarnv(&ione, ISEED, &mm, a); // randomize a
    lapackf77_slarnv(&ione, ISEED, &mn, b); // randomize b
    printf("upper left corner of the expected solution:\n");
    magma_sprint( 4, 4, b, m );
    // right hand side c=a*b
    blasf77_sgemm("N", "N", &m, &n, &n, &alpha, a, &m, b, &m, &beta, c, &m);
    // solve the linear system a*x=c
    // c -mxn matrix, a -mxm matrix;
    // c is overwritten by the solution
    gpu_time = magma_sync_wtime(NULL);

    magma_sgesv(m,n,a,m,piv,c,m,&info);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("magma_sgesv time: %7.5f sec.\n", gpu_time); // time
    printf("upper left corner of the solution:\n");
    magma_sprint( 4, 4, c, m ); // part of the solution
    magma_free_pinned(a); // free host memory
    magma_free_pinned(b); // free host memory
    magma_free_pinned(c); // free host memory
```

```

    free(piv);
    magma_finalize();
    return 0;
}
//upper left corner of the expected solution:
//[
//  0.3924    0.5546    0.6481    0.5479
//  0.9790    0.7204    0.4220    0.4588
//  0.5246    0.0813    0.8202    0.6163
//  0.6624    0.8634    0.8748    0.0717
//];
//magma_sgesv time: 0.61733 sec.
//upper left corner of the solution:
//[
//  0.3927    0.5548    0.6484    0.5483
//  0.9788    0.7204    0.4217    0.4586
//  0.5242    0.0815    0.8199    0.6161
//  0.6626    0.8631    0.8749    0.0717
//];

```

4.3.2 magma_sgesv - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv ){
    magma_init();
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    real_Double_t gpu_time;
    magma_int_t *piv, info;
    magma_int_t m = 8192;
    magma_int_t n = 100;
    magma_int_t mm=m*m;
    magma_int_t mn=m*n;
    float *a;
    float *b;
    float *c;
    magma_int_t ione = 1;
    magma_int_t ISEED[4] = { 0,0,0,1 };
    const float alpha = 1.0;
    const float beta = 0.0;
    cudaMallocManaged(&a,mm*sizeof(float)); // unified mem.for a
    cudaMallocManaged(&b,mn*sizeof(float)); // unified mem.for b
    cudaMallocManaged(&c,mn*sizeof(float)); // unified mem.for c
    cudaMallocManaged(&piv,m*sizeof(int)); // unified mem.for piv
    // generate random matrices a, b;

```

```

    lapackf77_slarnv(&ione, ISEED, &mm, a);           // randomize a
    lapackf77_slarnv(&ione, ISEED, &mn, b);           // randomize b
    printf("expected solution:\n");
    magma_sprint( 4, 4, b, m );
    // right hand side c=a*b
    blasf77_sgemm("N", "N", &m, &n, &n, &alpha, a, &m, b, &m, &beta, c, &m);
    // solve the linear system a*x=c
    // c -mxn matrix, a -mxm matrix;
    // c is overwritten by the solution
    gpu_time = magma_sync_wtime(NULL);

    magma_sgesv(m, n, a, m, piv, c, m, &info);

    gpu_time = magma_sync_wtime(NULL) - gpu_time;
    printf("magma_sgesv time: %7.5f sec.\n", gpu_time); // time
    printf("solution:\n");
    magma_sprint( 4, 4, c, m );           // part of the solution
    magma_free(a);                         // free memory
    magma_free(b);                         // free memory
    magma_free(c);                         // free memory
    magma_free(piv);                       // free memory
    magma_finalize();                      // finalize Magma
    return 0;
}
//expected solution:
//[
//  0.3924  0.5546  0.6481  0.5479
//  0.9790  0.7204  0.4220  0.4588
//  0.5246  0.0813  0.8202  0.6163
//  0.6624  0.8634  0.8748  0.0717
//];
//magma_sgesv time: 0.42720 sec.
//solution:
//[
//  0.3927  0.5548  0.6484  0.5483
//  0.9788  0.7204  0.4217  0.4586
//  0.5242  0.0815  0.8199  0.6161
//  0.6626  0.8631  0.8749  0.0717
//];

```

4.3.3 magma_dgesv - solve a general linear system in double precision, CPU interface

This function solves in double precision a general real, linear system

$$A X = B,$$

where A is an $m \times m$ matrix and X, B are $m \times n$ matrices. A, B are defined on the host. In the solution, the LU decomposition of A with partial pivoting and row interchanges is used. See [magma-X.Y.Z/src/dgesv.cpp](#) for more details.

```

#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    real_Double_t gpu_time;
    magma_int_t *piv, info;
    magma_int_t m = 8192; // a - mxm matrix
    magma_int_t n = 100; // c - mxn matrix
    magma_int_t mm=m*m; // size of a
    magma_int_t mn=m*n; // size of c
    double *a; // a- mxm matrix on the host
    double *b; // b- mxn matrix on the host
    double *c; // c- mxn matrix on the host
    magma_int_t ione = 1; //random uniform distr. in (0,1)
    magma_int_t ISEED[4] = { 0,0,0,1 }; // seed
    magma_int_t err;
    const double alpha = 1.0; // alpha=1
    const double beta = 0.0; // beta=0
    // allocate matrices on the host
    err = magma_dmalloc_pinned( &a , mm ); // host memory for a
    err = magma_dmalloc_pinned( &b , mn ); // host memory for b
    err = magma_dmalloc_pinned( &c , mn ); // host memory for c
    piv=(magma_int_t*)malloc(m*sizeof(magma_int_t));
    // generate random matrices a, b;
    lapackf77_dlarnv(&ione,ISEED,&mm,a); // randomize a
    lapackf77_dlarnv(&ione,ISEED,&mn,b); // randomize b
    printf("expected solution:\n");
    magma_dprint( 4, 4, b, m );
    // right hand side c=a*b
    blasf77_dgemm("N","N",&m,&n,&n,&alpha,a,&m,b,&m,&beta,c,&m);
    // solve the linear system a*x=c
    // c -mxn matrix, a -mxm matrix;
    // c is overwritten by the solution
    gpu_time = magma_sync_wtime(NULL);

    magma_dgesv(m,n,a,m,piv,c,m,&info);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("magma_dgesv time: %7.5f sec.\n",gpu_time); // time
    printf("solution:\n");
    magma_dprint( 4, 4, c, m ); // part of the solution
    magma_free_pinned(a); // free host memory
    magma_free_pinned(b); // free host memory
    magma_free_pinned(c); // free host memory
    free(piv); // free host memory
    magma_finalize(); // finalize Magma
    return 0;
}
//expected solution:

```



```
//[
//    0.5440    0.5225    0.8499    0.4012
//    0.4878    0.9321    0.2277    0.7495
//    0.0124    0.7743    0.5884    0.3296
//    0.2166    0.6253    0.8843    0.3685
//];
//magma_dgesv time: 1.81342 sec.
//solution:
//[
//    0.5440    0.5225    0.8499    0.4012
//    0.4878    0.9321    0.2277    0.7495
//    0.0124    0.7743    0.5884    0.3296
//    0.2166    0.6253    0.8843    0.3685
//];
```

4.3.4 magma_dgesv - unified memory version

```
#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    real_Double_t gpu_time;
    magma_int_t *piv, info;
    magma_int_t m = 8192; // a - mxm matrix
    magma_int_t n = 100; // c - mxn matrix
    magma_int_t mm=m*m; // size of a
    magma_int_t mn=m*n; // size of c
    double *a; // a- mxm matrix
    double *b; // b- mxn matrix
    double *c; // c- mxn matrix
    magma_int_t ione = 1; //random uniform distr. in (0,1)
    magma_int_t ISEED[4] = { 0,0,0,1 }; // seed
    const double alpha = 1.0; // alpha=1
    const double beta = 0.0; // beta=0
    cudaMallocManaged(&a,mm*sizeof(double)); // unified mem.for a
    cudaMallocManaged(&b,mn*sizeof(double)); // unified mem.for b
    cudaMallocManaged(&c,mn*sizeof(double)); // unified mem.for c
    cudaMallocManaged(&piv,m*sizeof(int)); // unified mem.for piv
    // generate random matrices a, b;
    lapackf77_dlarnv(&ione,ISEED,&mm,a); // randomize a
    lapackf77_dlarnv(&ione,ISEED,&mn,b); // randomize b
    printf("expected solution:\n");
    magma_dprint( 4, 4, b, m );
    // right hand side c=a*b
```

```

    blasf77_dgemm("N","N",&m,&n,&n,&alpha,a,&m,b,&m,&beta,c,&m);
// solve the linear system a*x=c
// c -mxn matrix, a -mxm matrix;
// c is overwritten by the solution
    gpu_time = magma_sync_wtime(NULL);

    magma_dgesv(m,n,a,m,piv,c,m,&info);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("magma_dgesv time: %7.5f sec.\n",gpu_time);    // time
    printf("solution:\n");
    magma_dprint( 4, 4, c, m );                        // part of the solution
    magma_free(a);                                     // free memory
    magma_free(b);                                     // free memory
    magma_free(c);                                     // free memory
    magma_free(piv);                                   // free memory
    magma_finalize();                                  // finalize Magma
    return 0;
}
//expected solution:
//[
//    0.5440    0.5225    0.8499    0.4012
//    0.4878    0.9321    0.2277    0.7495
//    0.0124    0.7743    0.5884    0.3296
//    0.2166    0.6253    0.8843    0.3685
//];
//magma_dgesv time: 1.69905 sec.
//solution:
//[
//    0.5440    0.5225    0.8499    0.4012
//    0.4878    0.9321    0.2277    0.7495
//    0.0124    0.7743    0.5884    0.3296
//    0.2166    0.6253    0.8843    0.3685
//];

```

4.3.5 magma_sgesv_gpu - solve a general linear system in single precision, GPU interface

This function solves in single precision a general real, linear system

$$A X = B,$$

where A is an $m \times m$ matrix and X, B are $m \times n$ matrices. A, B are defined on the device. In the solution, the LU decomposition of A with partial pivoting and row interchanges is used. See [magma-X.Y.Z/src/sgesv_gpu.cpp](#) for more details.

```

#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"

```

```

int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    real_Double_t gpu_time;
    magma_int_t *piv, info; // piv - array of indices of inter-
    magma_int_t m = 8192; // changed rows; a,d_a - mxm matrices
    magma_int_t n = 100; // b,c,d_c - mxn matrices
    magma_int_t mm=m*m; // size of a,d_a
    magma_int_t mn=m*n; // size of b,c,d_c
    float *a; // a- mxm matrix on the host
    float *b; // b- mxn matrix on the host
    float *c; // c- mxn matrix on the host
    float *d_a; // d_a- mxm matrix a on the device
    float *d_c; // d_c- mxn matrix c on the device
    magma_int_t ione = 1;
    magma_int_t ISEED[4] = { 0,0,0,1 }; // seed
    magma_int_t err;
    const float alpha = 1.0; // alpha=1
    const float beta = 0.0; // beta=0
    // allocate matrices
    err = magma_smalloc_cpu( &a , mm ); // host memory for a
    err = magma_smalloc_cpu( &b , mn ); // host memory for b
    err = magma_smalloc_cpu( &c , mn ); // host memory for c
    err = magma_smalloc( &d_a , mm ); // device memory for a
    err = magma_smalloc( &d_c , mn ); // device memory for c
    piv=(magma_int_t*)malloc(m*sizeof(magma_int_t)); // host mem.
    // generate matrices // for piv
    lapackf77_slarnv(&ione,ISEED,&mm,a); // randomize a
    lapackf77_slarnv(&ione,ISEED,&mn,b); // randomize b
    printf("expected solution:\n");
    magma_sprint( 4, 4, b, m ); // part of the expected solution
    // right hand side c=a*b
    blasf77_sgemm("N","N",&m,&n,&m,&alpha,a,&m,b,&m,&beta,c,&m);
    magma_ssetmatrix( m, m, a, m, d_a,m,queue); // copy a -> d_a
    magma_ssetmatrix( m, n, c, m, d_c,m,queue); // copy c -> d_c
    // MAGMA
    // solve the linear system d_a*x=d_c, d_a -mxm matrix,
    // d_c -mxn matrix, d_c is overwritten by the solution;
    // LU decomposition with partial pivoting and row
    // interchanges is used, row i is interchanged with row piv(i)
    gpu_time = magma_sync_wtime(NULL);

    magma_sgesv_gpu(m,n,d_a,m,piv,d_c,m,&info);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("magma_sgesv_gpu time: %7.5f sec.\n",gpu_time);
    magma_sgetmatrix( m, n, d_c, m, c,m,queue);
    printf("solution:\n");
    magma_sprint( 4, 4, c, m ); // part of Magma solution
    free(a); // free host memory

```

```

    free(b); // free host memory
    free(c); // free host memory
    free(piv); // free host memory
    magma_free(d_a); // free device memory
    magma_free(d_c); // free device memory
    magma_queue_destroy(queue); // destroy queue
    magma_finalize(); // finalize Magma
    return 0;
}
//expected solution:
//[
//  0.3924  0.5546  0.6481  0.5479
//  0.9790  0.7204  0.4220  0.4588
//  0.5246  0.0813  0.8202  0.6163
//  0.6624  0.8634  0.8748  0.0717
//];
//magma_sgesv_gpu time: 0.29100 sec.
//solution:
//[
//  0.3546  0.5629  0.6696  0.4666
//  1.0140  0.7044  0.4187  0.4630
//  0.5813  0.0568  0.8220  0.5983
//  0.6398  0.8704  0.8650  0.1062
//];

```

4.3.6 magma_sgesv_gpu - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    real_Double_t gpu_time;
    magma_int_t *piv, info; // piv - array of indices of inter-
    magma_int_t m = 8192; // changed rows; a - mxm matrix
    magma_int_t n = 100; // b,c - mxn matrices
    magma_int_t mm=m*m; // size of a
    magma_int_t mn=m*n; // size of b,c
    float *a; // a- mxm matrix
    float *b; // b- mxn matrix
    float *c; // c- mxn matrix
    magma_int_t ione = 1;
    magma_int_t ISEED[4] = { 0,0,0,1 }; // seed
    const float alpha = 1.0; // alpha=1
    const float beta = 0.0; // beta=0
    // allocate matrices

```

```

    cudaMallocManaged(&a,mm*sizeof(float)); // unified mem.for a
    cudaMallocManaged(&b,mn*sizeof(float)); // unified mem.for b
    cudaMallocManaged(&c,mn*sizeof(float)); // unified mem.for c
    cudaMallocManaged(&piv,m*sizeof(int)); // unified mem.for piv
// generate matrices
    lapackf77_slarnv(&ione,ISEED,&mm,a); // randomize a
    lapackf77_slarnv(&ione,ISEED,&mn,b); // randomize b
    printf("expected solution:\n");
    magma_sprint( 4, 4, b, m ); // part of the expected solution
// right hand side c=a*b
    blasf77_sgemm("N","N",&m,&n,&m,&alpha,a,&m,b,&m,&beta,c,&m);
// solve the linear system a*x=c, a -mxm matrix,
// c -mxn matrix, c is overwritten by the solution;
// LU decomposition with partial pivoting and row
// interchanges is used, row i is interchanged with row piv(i)
    gpu_time = magma_sync_wtime(NULL);

    magma_sgesv_gpu(m,n,a,m,piv,c,m,&info);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("magma_sgesv_gpu time: %7.5f sec.\n",gpu_time);
    printf("solution:\n");
    magma_sprint( 4, 4, c, m ); // part of Magma solution
    magma_free(piv); // free memory
    magma_free(a); // free memory
    magma_free(b); // free memory
    magma_free(c); // free memory
    magma_queue_destroy(queue); // destroy queue
    magma_finalize(); // finalize Magma
    return 0;
}
//expected solution:
//[
// 0.3924 0.5546 0.6481 0.5479
// 0.9790 0.7204 0.4220 0.4588
// 0.5246 0.0813 0.8202 0.6163
// 0.6624 0.8634 0.8748 0.0717
//];
//magma_sgesv_gpu time: 0.31976 sec.
//solution:
//[
// 0.3546 0.5629 0.6696 0.4666
// 1.0140 0.7044 0.4187 0.4630
// 0.5813 0.0568 0.8220 0.5983
// 0.6398 0.8704 0.8650 0.1062
//];

```

4.3.7 magma_dgesv_gpu - solve a general linear system in double precision, GPU interface

This function solves in double precision a general, f real linear system

$$A X = B,$$

where A is an $m \times m$ matrix and X, B are $m \times n$ matrices. A, B are defined on the device. In the solution, the LU decomposition of A with partial pivoting and row interchanges is used. See [magma-X.Y.Z/src/dgesv_gpu.cpp](#) for more details.

```
#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    real_Double_t gpu_time;
    magma_int_t *piv, info; // piv - array of indices of inter-
    magma_int_t m = 8192; // changed rows; a,d_a - mxm matrices
    magma_int_t n = 100; // b,c,d_c - mxn matrices
    magma_int_t mm=m*m; // size of a,d_a
    magma_int_t mn=m*n; // size of b,c,d_c
    double *a; // a- mxm matrix on the host
    double *b; // b- mxn matrix on the host
    double *c; // c- mxn matrix on the host
    double *d_a; // d_a- mxm matrix a on the device
    double *d_c; // d_c- mxn matrix c on the device
    magma_int_t ione = 1;
    magma_int_t ISEED[4] = { 0,0,0,1 }; // seed
    magma_int_t err;
    const double alpha = 1.0; // alpha=1
    const double beta = 0.0; // beta=0
    // allocate matrices
    err = magma_dmalloc_cpu( &a , mm ); // host memory for a
    err = magma_dmalloc_cpu( &b , mn ); // host memory for b
    err = magma_dmalloc_cpu( &c , mn ); // host memory for c
    err = magma_dmalloc( &d_a , mm ); // device memory for a
    err = magma_dmalloc( &d_c , mn ); // device memory for c
    piv=(magma_int_t*)malloc(m*sizeof(magma_int_t)); // host mem.
    // generate matrices // for piv
    lapackf77_dlarnv(&ione,ISEED,&mm,a); // randomize a
    lapackf77_dlarnv(&ione,ISEED,&mn,b); // randomize b
    printf("expected solution:\n");
    magma_dprint( 4, 4, b, m ); // part of the expected solution
    // right hand side c=a*b
    blasf77_dgemm("N","N",&m,&n,&m,&alpha,a,&m,b,&m,&beta,c,&m);
```

```

    magma_dsetmatrix( m, m, a, m, d_a,m,queue); // copy a -> d_a
    magma_dsetmatrix( m, n, c, m, d_c,m,queue); // copy c -> d_c
// MAGMA
// solve the linear system d_a*x=d_c, d_a -mxm matrix,
// d_c -mxn matrix, d_c is overwritten by the solution;
// LU decomposition with partial pivoting and row
// interchanges is used, row i is interchanged with row piv(i)
    gpu_time = magma_sync_wtime(NULL);

    magma_dgesv_gpu(m,n,d_a,m,piv,d_c,m,&info);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("magma_dgesv_gpu time: %7.5f sec.\n",gpu_time);
    magma_dgetmatrix( m, n, d_c, m, c,m,queue);
    printf("solution:\n");
    magma_dprint( 4, 4, c, m );           // part of Magma solution
    free(a);                             // free host memory
    free(b);                             // free host memory
    free(c);                             // free host memory
    free(piv);                           // free host memory
    magma_free(d_a);                     // free device memory
    magma_free(d_c);                     // free device memory
    magma_queue_destroy(queue);          // destroy queue
    magma_finalize();                     // finalize Magma
    return 0;
}
//expected solution:
//[
//    0.5440    0.5225    0.8499    0.4012
//    0.4878    0.9321    0.2277    0.7495
//    0.0124    0.7743    0.5884    0.3296
//    0.2166    0.6253    0.8843    0.3685
//];
//magma_dgesv_gpu time: 1.47404 sec.
//solution:
//[
//    0.5440    0.5225    0.8499    0.4012
//    0.4878    0.9321    0.2277    0.7495
//    0.0124    0.7743    0.5884    0.3296
//    0.2166    0.6253    0.8843    0.3685
//];

```

4.3.8 magma_dgesv_gpu - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv ){

```

```

magma_init(); // initialize Magma
magma_queue_t queue=NULL;
magma_int_t dev=0;
magma_queue_create(dev,&queue);
real_Double_t gpu_time;
magma_int_t *piv, info; // piv - array of indices of inter-
magma_int_t m = 8192; // changed rows; a,d_a - mxm matrices
magma_int_t n = 100; // b,c,d_c - mxn matrices
magma_int_t mm=m*m; // size of a
magma_int_t mn=m*n; // size of b,c
double *a; // a- mxm matrix
double *b; // b- mxn matrix
double *c; // c- mxn matrix
magma_int_t ione = 1;
magma_int_t ISEED[4] = { 0,0,0,1 }; // seed
const double alpha = 1.0; // alpha=1
const double beta = 0.0; // beta=0
// allocate matrices
cudaMallocManaged(&a,mm*sizeof(double)); // unified mem.for a
cudaMallocManaged(&b,mn*sizeof(double)); // unified mem.for b
cudaMallocManaged(&c,mn*sizeof(double)); // unified mem.for c
cudaMallocManaged(&piv,m*sizeof(int)); // unified mem.for piv
// generate matrices
lapackf77_dlarnv(&ione,ISEED,&mm,a); // randomize a
lapackf77_dlarnv(&ione,ISEED,&mn,b); // randomize b
printf("expected solution:\n");
magma_dprint( 4, 4, b, m ); // part of the expected solution
// right hand side c=a*b
blasf77_dgemm("N","N",&m,&n,&m,&alpha,a,&m,b,&m,&beta,c,&m);
// solve the linear system a*x=c, a -mxm matrix,
// c -mxn matrix, c is overwritten by the solution;
// LU decomposition with partial pivoting and row
// interchanges is used, row i is interchanged with row piv(i)
gpu_time = magma_sync_wtime(NULL);

magma_dgesv_gpu(m,n,a,m,piv,c,m,&info);

gpu_time = magma_sync_wtime(NULL)-gpu_time;
printf("magma_dgesv_gpu time: %7.5f sec.\n",gpu_time);
printf("solution:\n");
magma_dprint( 4, 4, c, m ); // part of Magma solution
magma_free(piv); // free memory
magma_free(a); // free memory
magma_free(b); // free memory
magma_free(c); // free memory
magma_queue_destroy(queue); // destroy queue
magma_finalize(); // finalize Magma
return 0;
}

//expected solution:

```



```

//[
//  0.5440    0.5225    0.8499    0.4012
//  0.4878    0.9321    0.2277    0.7495
//  0.0124    0.7743    0.5884    0.3296
//  0.2166    0.6253    0.8843    0.3685
//];
//magma_dgesv_gpu time: 1.55957 sec.
//solution:
//[
//  0.5440    0.5225    0.8499    0.4012
//  0.4878    0.9321    0.2277    0.7495
//  0.0124    0.7743    0.5884    0.3296
//  0.2166    0.6253    0.8843    0.3685
//];

```

4.3.9 magma_sgetrf, lapackf77_sgetrs - LU factorization and solving factorized systems in single precision, CPU interface

The first function using single precision computes an LU factorization of a general $m \times n$ matrix A using partial pivoting with row interchanges:

$$A = P L U,$$

where P is a permutation matrix, L is lower triangular with unit diagonal, and U is upper diagonal. The matrix A to be factored is defined on the host. On exit A contains the factors L, U . The information on the interchanged rows is contained in piv . See [magma-X.Y.Z/src/sgetrf.cpp](#) for more details.

Using the obtained factorization one can replace the problem of solving a general linear system by solving two triangular systems with matrices L and U respectively. The Lapack function `sgetrs` uses the LU factorization to solve a general linear system (it is faster to use Lapack `sgetrs` than to copy A to the device).

```

#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    real_Double_t gpu_time;
    magma_int_t *piv, info; // piv - array of indices of inter-
    magma_int_t m = 8192, n=8192; // changed rows; a - mxn matrix
    magma_int_t nrhs = 100; // b - nxnrhs, c - mxnrhs matrices
    magma_int_t mn=m*n; // size of a
    magma_int_t nnrhs=n*nrhs; // size of b
    magma_int_t mnrhs=m*nrhs; // size of c
    float *a; // a- mxn matrix on the host
    float *b; // b- nxnrhs matrix on the host

```

```

float *c;                                // c- mxnrhs matrix on the host
magma_int_t ione = 1;
magma_int_t ISEED[4] = {0,0,0,1};        // seed
magma_int_t err;
const float alpha = 1.0;                  // alpha=1
const float beta = 0.0;                   // beta=0
// allocate matrices on the host
err = magma_smalloc_pinned(&a , mn );    // host memory for a
err = magma_smalloc_pinned(&b, nnrhs );  // host memory for b
err = magma_smalloc_pinned(&c, mnrhs );  // host memory for c
piv=(magma_int_t*)malloc(m*sizeof(magma_int_t)); // host mem.
// generate matrices                                // for piv
lapackf77_slarnv(&ione, ISEED, &mn, a);    // randomize a
lapackf77_slaset(MagmaFullStr, &n, &nnrhs, &alpha, &alpha, b, &n);
// b - nxnrhs matrix of ones
printf("upper left corner of the expected solution:\n");
magma_sprint( 4, 4, b, m ); // part of the expected solution
blasf77_sgemm("N", "N", &m, &nnrhs, &n, &alpha, a, &m, b, &m, &beta, c,
              &m);           // right hand side c=a*b
// MAGMA
// solve the linear system a*x=c, a -mxn matrix, c -mxnrhs ma-
// trix, c is overwritten by the solution; LU decomposition
// with partial pivoting and row interchanges is used,
// row i is interchanged with row piv(i)
gpu_time = magma_sync_wtime(NULL);

magma_sgetrf( m, n, a, m, piv, &info);
lapackf77_sgetrs("N", &m, &nnrhs, a, &m, piv, c, &m, &info);

gpu_time = magma_sync_wtime(NULL)-gpu_time;
printf("magma_sgetrf + lapackf77_sgetrs time: %7.5f sec.\n",
      gpu_time); // Magma/Lapack time
printf("upper left corner of the solution:\n");
magma_sprint( 4, 4, c, m ); // part of the Magma/Lapack sol.
magma_free_pinned(a);        // free host memory
magma_free_pinned(b);        // free host memory
magma_free_pinned(c);        // free host memory
free(piv);                   // free host memory
magma_finalize();            // finalize Magma
return 0;
}

//upper left corner of the expected solution:
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];
//magma_sgetrf + lapackf77_sgetrs time: 0.77011 sec.
//upper left corner of the solution:

```

```

//[
//  0.9682    0.9682    0.9682    0.9682
//  1.0134    1.0134    1.0134    1.0134
//  1.0147    1.0147    1.0147    1.0147
//  1.0034    1.0034    1.0034    1.0034
//];

```

4.3.10 magma_sgetrf, lapackf77_sgetrs - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    real_Double_t gpu_time;
    magma_int_t *piv, info; // piv - array of indices of inter-
    magma_int_t m = 8192, n=8192; // changed rows; a - mxn matrix
    magma_int_t nrhs = 100; // b - nxnrhs, c - mxnrhs matrices
    magma_int_t mn=m*n; // size of a
    magma_int_t nnrhs=n*nrhs; // size of b
    magma_int_t mnrhs=m*nrhs; // size of c
    float *a; // a- mxn matrix
    float *b; // b- nxnrhs matrix
    float *c; // c- mxnrhs matrix
    magma_int_t ione = 1;
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    const float alpha = 1.0; // alpha=1
    const float beta = 0.0; // beta=0
    cudaMallocManaged(&a, mn*sizeof(float)); // unified mem.for a
    cudaMallocManaged(&b, nnrhs*sizeof(float)); //unif. mem.for b
    cudaMallocManaged(&c, mnrhs*sizeof(float)); //unif. mem.for c
    cudaMallocManaged(&piv, m*sizeof(int)); // unified mem.for piv
    // generate matrices
    lapackf77_slarnv(&ione, ISEED, &mn, a); // randomize a
    lapackf77_slaset(MagmaFullStr, &n, &nrhs, &alpha, &alpha, b, &n);
    // b - nxnrhs matrix of ones
    printf("upper left corner of the expected solution:\n");
    magma_sprint( 4, 4, b, m ); // part of the expected solution
    // right hand side c=a*b
    blasf77_sgemm("N", "N", &m, &nrhs, &n, &alpha, a, &m, b, &m, &beta,
        c, &m); // right hand side c=a*b
    // solve the linear system a*x=c, a -mxn matrix, c -mxnrhs ma-
    // trix, c is overwritten by the solution; LU decomposition
    // with partial pivoting and row interchanges is used,
    // row i is interchanged with row piv(i)
    gpu_time = magma_sync_wtime(NULL);

```

```

magma_sgetrf( m, n, a, m, piv, &info);
lapackf77_sgetrs("N",&m,&nrhs,a,&m,piv,c,&m,&info);

gpu_time = magma_sync_wtime(NULL)-gpu_time;
printf("magma_sgetrf + lapackf77_sgetrs time: %7.5f sec.\n",
      gpu_time); // Magma/Lapack time
printf("upper left corner of the Magma solution:\n");
magma_sprint( 4, 4, c, m );//part of the Magma/Lap. solution
magma_free(a);           // free memory
magma_free(b);           // free memory
magma_free(c);           // free memory
magma_free(piv);         // free memory
magma_finalize();        // finalize Magma
return 0;
}
//upper left corner of the expected solution:
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];
//magma_sgetrf + lapackf77_sgetrs time: 0.80166 sec.
//upper left corner of the Magma solution:
//[
//  0.9682  0.9682  0.9682  0.9682
//  1.0134  1.0134  1.0134  1.0134
//  1.0147  1.0147  1.0147  1.0147
//  1.0034  1.0034  1.0034  1.0034
//];

```

4.3.11 magma_dgetrf, lapackf77_dgetrs - LU factorization and solving factorized systems in double precision, CPU interface

The first function using double precision computes an LU factorization of a general $m \times n$ matrix A using partial pivoting with row interchanges:

$$A = P L U,$$

where P is a permutation matrix, L is lower triangular with unit diagonal, and U is upper diagonal. The matrix A to be factored is defined on the host. On exit A contains the factors L, U . The information on the interchanged rows is contained in piv . See [magma-X.Y.Z/src/sgetrf.cpp](#) for more details.

Using the obtained factorization one can replace the problem of solving a general linear system by solving two triangular systems with matrices L and U respectively. The Lapack function `dgetrs` uses the LU factorization to solve a general linear system (it is faster to use Lapack `dgetrs` than to copy A to the device).

```

#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    real_Double_t gpu_time;
    magma_int_t *piv, info; // piv - array of indices of inter-
    magma_int_t m = 8192,n=8192; // changed rows; a - mxn matrix
    magma_int_t nrhs = 100; // b - nxnrhs, c - mxnrhs matrices
    magma_int_t mn=m*n; // size of a
    magma_int_t nnrhs=n*nrhs; // size of b
    magma_int_t mnrhs=m*nrhs; // size of c
    double *a; // a- mxn matrix on the host
    double *b; // b- nxnrhs matrix on the host
    double *c; // c- mxnrhs matrix on the host
    magma_int_t ione = 1;
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    magma_int_t err;
    const double alpha = 1.0; // alpha=1
    const double beta = 0.0; // beta=0
    // allocate matrices on the host
    err = magma_dmalloc_pinned(&a,mn); // host memory for a
    err = magma_dmalloc_pinned(&b,nnrhs); // host memory for b
    err = magma_dmalloc_pinned(&c,mnrhs); // host memory for c
    piv=(magma_int_t*)malloc(m*sizeof(magma_int_t)); // host mem.
    // generate matrices // for piv
    lapackf77_dlarnv(&ione,ISEED,&mn,a); // randomize a
    lapackf77_dlaset(MagmaFullStr,&n,&nrhs,&alpha,&alpha,b,&n);
    // b - nxnrhs matrix of ones
    printf("upper left corner of the expected solution:\n");
    magma_dprint( 4, 4, b, m ); // part of the expected solution
    // right hand side c=a*b
    blasf77_dgemm("N","N",&m,&nrhs,&n,&alpha,a,&m,b,&m,&beta,
    c,&m); // right hand side c=a*b
    // MAGMA
    // solve the linear system a*x=c, a -mxn matrix, c -mxnrhs ma-
    // trix, c is overwritten by the solution; LU decomposition
    // with partial pivoting and row interchanges is used,
    // row i is interchanged with row piv(i)
    gpu_time = magma_sync_wtime(NULL);

    magma_dgetrf( m, n, a, m, piv, &info);
    lapackf77_dgetrs("N",&m,&nrhs,a,&m,piv,c,&m,&info);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("magma_dgetrf + lapackf77_dgetrs time: %7.5f sec.\n",
    gpu_time); // Magma/Lapack time
    printf("upper left corner of the Magma solution:\n");
    magma_dprint( 4, 4, c, m );//part of the Magma/Lap. solution
    magma_free_pinned(a); // free host memory
    magma_free_pinned(b); // free host memory

```

```

    magma_free_pinned(c);           // free host memory
    free(piv);                      // free host memory
    magma_finalize();               // finalize Magma
    return 0;
}
//upper left corner of the expected solution:
//[
//  1.0000    1.0000    1.0000    1.0000
//  1.0000    1.0000    1.0000    1.0000
//  1.0000    1.0000    1.0000    1.0000
//  1.0000    1.0000    1.0000    1.0000
//];
//magma_dgetrf + lapackf77_dgetrs time: 1.89429 sec.
//upper left corner of the Magma solution:
//[
//  1.0000    1.0000    1.0000    1.0000
//  1.0000    1.0000    1.0000    1.0000
//  1.0000    1.0000    1.0000    1.0000
//  1.0000    1.0000    1.0000    1.0000
//];

```

4.3.12 magma_dgetrf, lapackf77_dgetrs - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv ){
    magma_init();                      // initialize Magma
    real_Double_t  gpu_time;
    magma_int_t *piv, info; // piv - array of indices of inter-
    magma_int_t m = 8192,n=8192; // changed rows; a - mxn matrix
    magma_int_t nrhs = 100; // b - nxnrhs, c - mxnrhs matrices
    magma_int_t mn=m*n;                // size of a
    magma_int_t nnrhs=n*nrhs;          // size of b
    magma_int_t mnrhs=m*nrhs;          // size of c
    double *a;                         // a- mxn matrix
    double *b;                         // b- nxnrhs matrix
    double *c;                         // c- mxnrhs matrix
    magma_int_t ione = 1;
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    const double alpha = 1.0;          // alpha=1
    const double beta = 0.0;           // beta=0
    cudaMallocManaged(&a,mn*sizeof(double)); // unified mem.for a
    cudaMallocManaged(&b,nnrhs*sizeof(double)); // unif.mem.for b
    cudaMallocManaged(&c,mnrhs*sizeof(double)); // unif.mem.for c
    cudaMallocManaged(&piv,m*sizeof(int)); // unified mem.for piv
    // generate matrices
    lapackf77_dlarnv(&ione,ISEED,&mn,a); // randomize a

```

```

    lapackf77_dlaset(MagmaFullStr,&n,&nrhs,&alpha,&alpha,b,&n);
                                // b - nxnrhs matrix of ones
    printf("upper left corner of the expected solution:\n");
    magma_dprint( 4, 4, b, m ); // part of the expected solution
// right hand side c=a*b
    blasf77_dgemm("N","N",&m,&nrhs,&n,&alpha,a,&m,b,&m,&beta,
                                c,&m); // right hand side c=a*b
// solve the linear system a*x=c, a -mxn matrix, c -mxnrhs ma-
// trix, c is overwritten by the solution; LU decomposition
// with partial pivoting and row interchanges is used,
// row i is interchanged with row piv(i)
    gpu_time = magma_sync_wtime(NULL);

    magma_dgetrf( m, n, a, m, piv, &info);
    lapackf77_dgetrs("N",&m,&nrhs,a,&m,piv,c,&m,&info);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("magma_dgetrf + lapackf77_dgetrs time: %7.5f sec.\n",
                                gpu_time); // Magma/Lapack time
    printf("upper left corner of the Magma solution:\n");
    magma_dprint( 4, 4, c, m );//part of the Magma/Lap. solution
    magma_free(a); // free memory
    magma_free(b); // free memory
    magma_free(c); // free memory
    magma_free(piv); // free memory
    magma_finalize(); // finalize Magma
    return 0;
}
//upper left corner of the expected solution:
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];
//magma_dgetrf + lapackf77_dgetrs time: 2.03707 sec.
//upper left corner of the Magma solution:
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];

```

4.3.13 magma_sgetrf_gpu, magma_sgetrs_gpu - LU factorization and solving factorized systems in single precision, GPU interface

The function `magma_sgetrf_gpu` computes in single precision an LU factorization of a general $m \times n$ matrix A using partial pivoting with row

interchanges:

$$A = P L U,$$

where P is a permutation matrix, L is lower triangular with unit diagonal, and U is upper diagonal. The matrix A to be factored and the factors L, U are defined on the device. The information on the interchanged rows is contained in *piv*. See [magma-X.Y.Z/src/sgetrf_gpu.cpp](#) for more details. Using the obtained factorization one can replace the problem of solving a general linear system by solving two triangular systems with matrices L and U respectively. The function `magma_sgetrs_gpu` uses the L, U factors defined on the device by `magma_sgetrf_gpu` to solve in single precision a general linear system

$$A X = B.$$

The right hand side B is a matrix defined on the device. On exit it is replaced by the solution. See [magma-X.Y.Z/src/sgetrs_gpu.cpp](#) for more details.

```
#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    real_Double_t gpu_time;
    magma_int_t *piv, info; // piv - array of indices of inter-
    magma_int_t m = 8192,n=8192; // changed rows; a - mxn matrix
    magma_int_t nrhs = 100; // b - nxnrhs, c - mxnrhs matrices
    magma_int_t mn=m*n; // size of a
    magma_int_t nnrhs=n*nrhs; // size of b
    magma_int_t mnrhs=m*nrhs; // size of c
    float *a; // a- mxn matrix on the host
    float *b; // b- nxnrhs matrix on the host
    float *c; // c- mxnrhs matrix on the host
    float *d_a; // d_a- mxn matrix a on the device
    float *d_c; // d_c- mxnrhs matrix c on the device
    magma_int_t ione = 1;
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    magma_int_t err;
    const float alpha = 1.0; // alpha=1
    const float beta = 0.0; // beta=0
    // allocate matrices on the host
    err = magma_smalloc_cpu( &a , mn ); // host memory for a
    err = magma_smalloc_cpu( &b , nnrhs ); // host memory for b
    err = magma_smalloc_cpu( &c , mnrhs ); // host memory for c
    err = magma_smalloc( &d_a, mn ); // device memory for a
    err = magma_smalloc( &d_c, mnrhs ); // device memory for c
```



```

    piv=(magma_int_t*)malloc(m*sizeof(magma_int_t)); // host mem.
// generate matrices // for piv
lapackf77_slarnv(&ione, ISEED, &mn, a); // randomize a
lapackf77_slaset(MagmaFullStr, &n, &nrhs, &alpha, &alpha, b, &n);
// b - nxnrhs matrix of ones

printf("upper left corner of the expected solution:\n");
magma_sprint( 4, 4, b, n ); // part of the expected solution
// right hand side c=a*b
blasf77_sgemm("N", "N", &m, &nrhs, &n, &alpha, a, &m, b, &m, &beta, c, &m);
magma_ssetmatrix( m, n, a, m, d_a, m, queue); // copy a -> d_a
magma_ssetmatrix( m, nrhs, c, m, d_c, m, queue); // copy c -> d_c
// MAGMA
// solve the linear system d_a*x=d_c, d_a -mxn matrix,
// d_c -mxnrhs matrix, d_c is overwritten by the solution;
// LU decomposition with partial pivoting and row interchanges
// is used, row i is interchanged with row piv(i)
gpu_time = magma_sync_wtime(NULL);

magma_sgetrf_gpu( m, n, d_a, m, piv, &info);
magma_sgetrs_gpu(MagmaNoTrans, m, nrhs, d_a, m, piv, d_c, m, &info);

gpu_time = magma_sync_wtime(NULL)-gpu_time;
printf("magma_sgetrf_gpu+magma_sgetrs_gpu time: %7.5f sec.\n",
      gpu_time); // Magma time
magma_sgetmatrix( m, nrhs, d_c, m, c, m, queue);
printf("upper left corner of the Magma solution:\n");
magma_sprint( 4, 4, c, m ); // part of the Magma solution
free(a); // free host memory
free(b); // free host memory
free(c); // free host memory
free(piv); // free host memory
magma_free(d_a); // free device memory
magma_free(d_c); // free device memory
magma_queue_destroy(queue); // destroy queue
magma_finalize(); // finalize Magma
return 0;
}
//upper left corner of the expected solution:
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];
//magma_sgetrf_gpu+magma_sgetrs_gpu time: 0.28847 sec.
//upper left corner of the Magma solution:
//[
//  1.0431  1.0431  1.0431  1.0431
//  1.0446  1.0446  1.0446  1.0446
//  1.1094  1.1094  1.1094  1.1094
//  0.9207  0.9207  0.9207  0.9207
//];

```

4.3.14 magma_sgetrf_gpu, magma_sgetrs_gpu - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    real_Double_t gpu_time;
    magma_int_t *piv, info; // piv - array of indices of inter-
    magma_int_t m = 8192,n=8192; // changed rows; a - mxn matrix
    magma_int_t nrhs = 100; // b - nxnrhs, c - mxnrhs matrices
    magma_int_t mn=m*n; // size of a
    magma_int_t nnrhs=n*nrhs; // size of b
    magma_int_t mnrhs=m*nrhs; // size of c
    float *a; // a- mxn matrix
    float *b; // b- nxnrhs matrix
    float *c; // c- mxnrhs matrix
    magma_int_t ione = 1;
    magma_int_t ISEED[4] = { 0,0,0,1 }; // seed
    const float alpha = 1.0; // alpha=1
    const float beta = 0.0; // beta=0
    cudaMallocManaged(&a,mn*sizeof(float)); // unified mem.for a
    cudaMallocManaged(&b,nnrhs*sizeof(float)); // unif. mem.for b
    cudaMallocManaged(&c,mnrhs*sizeof(float)); // unif. mem.for c
    cudaMallocManaged(&piv,m*sizeof(int)); // unified mem.for piv
    // generate matrices a, b;
    lapackf77_slarnv(&ione,ISEED,&mn,a); // randomize a
    lapackf77_slaset(MagmaFullStr,&n,&nrhs,&alpha,&alpha,b,&n);
    // b - nxnrhs matrix of ones
    printf("upper left corner of the expected solution:\n");
    magma_sprint( 4, 4, b, m ); // part of the expected solution
    // right hand side c=a*b
    blasf77_sgemm("N","N",&m,&nrhs,&n,&alpha,a,&m,b,&m,&beta,c,&m);
    // solve the linear system a*x=c, a -mxn matrix,
    // c -mxnrhs matrix, c is overwritten by the solution;
    // LU decomposition with partial pivoting and row interchanges
    // is used, row i is interchanged with row piv(i)
    gpu_time = magma_sync_wtime(NULL);

    magma_sgetrf_gpu( m, n, a, m, piv, &info);
    magma_sgetrs_gpu(MagmaNoTrans,m,nrhs,a,m,piv,c,m,&info);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("magma_sgetrf_gpu+magma_sgetrs_gpu time: %7.5f sec.\n",
        gpu_time); // Magma time

```

```

printf("upper left corner of the solution:\n");
magma_sprint( 4, 4, c, m );      // part of the Magma solution
magma_free(piv);                  // free memory
magma_free(a);                    // free memory
magma_free(b);                    // free memory
magma_free(c);                    // free memory
magma_queue_destroy(queue);       // destroy queue
magma_finalize();                 // finalize Magma
return 0;
}
//upper left corner of the expected solution:
//[
//  1.0000    1.0000    1.0000    1.0000
//  1.0000    1.0000    1.0000    1.0000
//  1.0000    1.0000    1.0000    1.0000
//  1.0000    1.0000    1.0000    1.0000
//];
//magma_sgetrf_gpu+magma_sgetrs_gpu time: 0.31721 sec.
//upper left corner of the solution:
//[
//  1.0431    1.0431    1.0431    1.0431
//  1.0446    1.0446    1.0446    1.0446
//  1.1094    1.1094    1.1094    1.1094
//  0.9207    0.9207    0.9207    0.9207
//];

```

4.3.15 magma_dgetrf_gpu, magma_dgetrs_gpu - LU factorization and solving factorized systems in double precision , GPU interface

The function `magma_dgetrf_gpu` computes in double precision an LU factorization of a general $m \times n$ matrix A using partial pivoting with row interchanges:

$$A = P L U,$$

where P is a permutation matrix, L is lower triangular with unit diagonal, and U is upper diagonal. The matrix A to be factored and the factors L, U are defined on the device. The information on the interchanged rows is contained in `piv`. See [magma-X.Y.Z/src/dgetrf_gpu.cpp](#) for more details. Using the obtained factorization one can replace the problem of solving a general linear system by solving two triangular systems with matrices L and U respectively. The function `magma_dgetrs_gpu` uses the L, U factors defined on the device by `magma_dgetrf_gpu` to solve in double precision a general linear system

$$A X = B.$$

The right hand side B is a matrix defined on the device. On exit it is replaced by the solution. See [magma-X.Y.Z/src/dgetrs_gpu.cpp](#) for more details.

```

#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    real_Double_t gpu_time;
    magma_int_t *piv, info; // piv - array of indices of inter-
    magma_int_t m = 8192,n=8192; // changed rows; a - mxn matrix
    magma_int_t nrhs = 100; // b - nxnrhs, c - mxnrhs matrices
    magma_int_t mn=m*n; // size of a
    magma_int_t nnrhs=n*nrhs; // size of b
    magma_int_t mnrhs=m*nrhs; // size of c
    double *a; // a- mxn matrix on the host
    double *b; // b- nxnrhs matrix on the host
    double *c; // c- mxnrhs matrix on the host
    double *d_a; // d_a- mxn matrix a on the device
    double *d_c; // d_c- mxnrhs matrix c on the device
    magma_int_t ione = 1;
    magma_int_t ISEED[4] = { 0,0,0,1 }; // seed
    magma_int_t err;
    const double alpha = 1.0; // alpha=1
    const double beta = 0.0; // beta=0
    // allocate matrices on the host
    err = magma_dmalloc_cpu( &a , mn ); // host memory for a
    err = magma_dmalloc_cpu( &b , nnrhs ); // host memory for b
    err = magma_dmalloc_cpu( &c , mnrhs ); // host memory for c
    err = magma_dmalloc( &d_a , mn ); // device memory for a
    err = magma_dmalloc( &d_c , mnrhs ); // device memory for c
    piv=(magma_int_t*)malloc(m*sizeof(magma_int_t));
    // generate matrices a, b;
    lapackf77_dlarnv(&ione,ISEED,&mn,a); // randomize a
    lapackf77_dlaset(MagmaFullStr,&n,&nrhs,&alpha,&alpha,b,&n);
    // b - nxnrhs matrix of ones
    printf("upper left corner of the expected solution:\n");
    magma_dprint( 4, 4, b, m ); // part of the expected solution
    // right hand side c=a*b
    blasf77_dgemm("N","N",&m,&nrhs,&n,&alpha,a,&m,b,&m,&beta,c,&m);
    magma_dsetmatrix( m, n, a, m, d_a,m,queue); // copy a -> d_a
    magma_dsetmatrix( m, nrhs, c,m,d_c,m,queue); // copy c -> d_c
    // MAGMA
    // solve the linear system d_a*x=d_c, d_a -mxn matrix,
    // d_c -mxnrhs matrix, d_c is overwritten by the solution;
    // LU decomposition with partial pivoting and row interchanges
    // is used, row i is interchanged with row piv(i)
    gpu_time = magma_sync_wtime(NULL);

    magma_dgetrf_gpu( m, n, d_a, m, piv, &info);
    magma_dgetrs_gpu(MagmaNoTrans,m,nrhs,d_a,m,piv,d_c,m,&info);

```

```

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("magma_dgetrf_gpu+magma_dgetrs_gpu time: %7.5f sec.\n",
           gpu_time); // Magma time
    magma_dgetmatrix( m, nrhs, d_c, m, c, m,queue);
    printf("upper left corner of the solution:\n");
    magma_dprint( 4, 4, c, m ); // part of the Magma solution
    free(a); // free host memory
    free(b); // free host memory
    free(c); // free host memory
    free(piv); // free host memory
    magma_free(d_a); // free device memory
    magma_free(d_c); // free device memory
    magma_queue_destroy(queue); // destroy queue
    magma_finalize(); // finalize Magma
    return 0;
}
//upper left corner of the expected solution:
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];
//magma_dgetrf_gpu+magma_dgetrs_gpu time: 1.47816 sec.
//upper left corner of the solution:
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];

```

4.3.16 magma_dgetrf_gpu, magma_dgetrs_gpu - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    real_Double_t gpu_time;
    magma_int_t *piv, info; // piv - array of indices of inter-
    magma_int_t m = 8192,n=8192; // changed rows; a - mxn matrix
    magma_int_t nrhs = 100; // b - nxnrhs, c - mxnrhs matrices

```

```

magma_int_t mn=m*n; // size of a
magma_int_t nnrhs=n*nrhs; // size of b
magma_int_t mnrhs=m*nrhs; // size of c
double *a; // a- mxn matrix
double *b; // b- nxnrhs matrix
double *c; // c- mxnrhs matrix
magma_int_t ione = 1;
magma_int_t ISEED[4] = { 0,0,0,1 }; // seed
const double alpha = 1.0; // alpha=1
const double beta = 0.0; // beta=0
cudaMallocManaged(&a,mn*sizeof(double)); // unified mem.for a
cudaMallocManaged(&b,nnrhs*sizeof(double)); // unif.mem.for b
cudaMallocManaged(&c,mnrhs*sizeof(double)); // unif.mem.for c
cudaMallocManaged(&piv,m*sizeof(int)); // unified mem.for piv
// generate matrices a, b;
lapackf77_dlarv(&ione,ISEED,&mn,a); // randomize a
lapackf77_dlaset(MagmaFullStr,&n,&nrhs,&alpha,&alpha,b,&n);
// b - nxnrhs matrix of ones
printf("upper left corner of the expected solution:\n");
magma_dprint( 4, 4, b, m ); // part of the expected solution
// right hand side c=a*b
blasf77_dgemm("N","N",&m,&nrhs,&n,&alpha,a,&m,b,&m,&beta,c,&m);
// solve the linear system a*x=c, a -mxn matrix,
// c -mxnrhs matrix, c is overwritten by the solution;
// LU decomposition with partial pivoting and row interchanges
// is used, row i is interchanged with row piv(i)
gpu_time = magma_sync_wtime(NULL);

magma_dgetrf_gpu( m, n, a, m, piv, &info);
magma_dgetrs_gpu(MagmaNoTrans,m,nrhs,a,m,piv,c,m,&info);

gpu_time = magma_sync_wtime(NULL)-gpu_time;
printf("magma_dgetrf_gpu+magma_dgetrs_gpu time: %7.5f sec.\n",
      gpu_time); // Magma time
printf("upper left corner of the solution:\n");
magma_dprint( 4, 4, c, m ); // part of the Magma solution
magma_free(piv); // free memory
magma_free(a); // free memory
magma_free(b); // free memory
magma_free(c); // free memory
magma_queue_destroy(queue); // destroy queue
magma_finalize(); // finalize Magma
return 0;
}
//upper left corner of the expected solution:
//[
// 1.0000 1.0000 1.0000 1.0000
// 1.0000 1.0000 1.0000 1.0000
// 1.0000 1.0000 1.0000 1.0000
// 1.0000 1.0000 1.0000 1.0000
//];
//magma_dgetrf_gpu+magma_dgetrs_gpu time: 1.51280 sec.

```

```
//upper left corner of the solution:
//[
//  1.0000    1.0000    1.0000    1.0000
//  1.0000    1.0000    1.0000    1.0000
//  1.0000    1.0000    1.0000    1.0000
//  1.0000    1.0000    1.0000    1.0000
//];
```

4.3.17 magma_sgetrf_mgpu - LU factorization in single precision on multiple GPUs

The function `magma_sgetrf_mgpu` computes in single precision an LU factorization of a general $m \times n$ matrix A using partial pivoting with row interchanges:

$$A = P L U,$$

where P is a permutation matrix, L is lower triangular with unit diagonal, and U is upper diagonal. The blocks of matrix A to be factored and the blocks of factors L, U are distributed on `num_gpus` devices. The information on the interchanged rows is contained in `ipiv`. See [magma-X.Y.Z/src/sgetrf_mgpu.cpp](#) for more details.

Using the obtained factorization one can replace the problem of solving a general linear system by solving two triangular systems with matrices L and U respectively. The Lapack function `lapackf77_sgetrs` uses the L, U factors copied from `num_gpus` devices to solve in single precision a general linear system

$$A X = B.$$

The right hand side B is a matrix defined on the host. On exit it is replaced by the solution.

```
#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b) (((a)<(b))?(a):(b))
int main( int argc, char** argv)
{
    magma_init(); // initialize Magma
    int num_gpus = 1;
    magma_setdevice(0);
    magma_queue_t queues[num_gpus];
    for( int dev = 0; dev < num_gpus; ++dev ) {
        magma_queue_create( dev, &queues[dev] );
    }
    magma_int_t err;
    real_Double_t cpu_time, gpu_time;
    magma_int_t m = 8192, n = 8192; // a,r - mxn matrices
    magma_int_t nrhs = 100; // b - nxnrhs, c - mxnrhs matrices
    magma_int_t *ipiv; // array of indices of interchanged rows
```

```

magma_int_t n2=m*n; // size of a,r
magma_int_t nnrhs=n*nrhs; // size of b
magma_int_t mnrhs=m*nrhs; // size of c
float *a, *r; // a,r - mxn matrices on the host
float *b, *c; // b - nxnrhs, c - mxnrhs matrices on the host
magmaFloat_ptr d_la[num_gpus];
float alpha=1.0, beta=0.0; // alpha=1,beta=0
magma_int_t n_local;
magma_int_t ione = 1, info;
magma_int_t i, min_mn=min(m,n), nb;
magma_int_t ldn_local; // m*ldn_local - size of the part of a
magma_int_t ISEED[4] = {0,0,0,1}; // on i-th device
nb =magma_get_sgetrf_nb(m,n); //optim.block size for sgetrf
// allocate memory on cpu
ipiv=(magma_int_t*)malloc(min_mn*sizeof(magma_int_t));
// host memory for ipiv
err = magma_smalloc_cpu(&a,n2); // host memory for a
err = magma_smalloc_pinned(&r,n2); // host memory for r
err = magma_smalloc_pinned(&b,nnrhs); // host memory for b
err = magma_smalloc_pinned(&c,mnrhs); // host memory for c
// allocate device memory on num_gpus devices
for(i=0; i<num_gpus; i++){
    n_local = ((n/nb)/num_gpus)*nb;
    if (i < (n/nb)%num_gpus)
        n_local += nb;
    else if (i == (n/nb)%num_gpus)
        n_local += n%nb;
    ldn_local = ((n_local+31)/32)*32;
    magma_setdevice(i);
    err = magma_smalloc(&d_la[i],m*ldn_local); //device memory
} // on i-th device
magma_setdevice(0);
// generate matrices
lapackf77_slarnv( &ione, ISEED, &n2, a ); // randomize a
lapackf77_slaset(MagmaFullStr,&n,&nrhs,&alpha,&alpha,b,&n);
// b - nxnrhs matrix of ones
lapackf77_slacpy( MagmaFullStr,&m,&n,a,&m,r,&m); //a->r
printf("upper left corner of the expected solution:\n");
magma_sprint(4,4,b,m); // part of the expected solution
blasf77_sgemm("N","N",&m,&nrhs,&n,&alpha,a,&m,b,&m,
    &beta,c,&m); // right hand side c=a*b
// LAPACK version of LU decomposition
cpu_time=magma_wtime();
lapackf77_sgetrf(&m, &n, a, &m, ipiv, &info);
cpu_time=magma_wtime()-cpu_time;
printf("lapackf77_sgetrf time: %7.5f sec.\n",cpu_time);
// copy the corresponding parts of the matrix r to num_gpus
magma_ssetmatrix_1D_col_bcyclic( num_gpus, m, n, nb, r, m,
    d_la, m, queues );
// MAGMA
// LU decomposition on num_gpus devices with partial pivoting
// and row interchanges, row i is interchanged with row ipiv(i)

```



```

    gpu_time = magma_sync_wtime(NULL);

    magma_sgetrf_mgpu( num_gpus, m, n, d_la, m, ipiv, &info);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("magma_sgetrf_mgpu time: %7.5f sec.\n",gpu_time);
    // copy the decomposition from num_gpus devices to r on the
    // host
    magma_sgetmatrix_1D_col_bccyclic( num_gpus, m, n, nb, d_la,
                                     m, r, m, queues );

    magma_setdevice(0);
    // solve on the host the system r*x=c; x overwrites c,
    // using the LU decomposition obtained on num_gpus devices
    lapackf77_sgetrs("N",&m,&nrhs,r,&m,ipiv,c,&m,&info);
    // print part of the solution from sgetrf_mgpu and sgetrs
    printf("upper left corner of the solution \n\
    from sgetrf_mgpu+sgetrs:\n");    // part of the solution from
    magma_sprint( 4, 4, c, m);    // magma_sgetrf_mgpu + sgetrs
    free(ipiv);    // free host memory
    free(a);    // free host memory
    magma_free_pinned(r);    // free host memory
    magma_free_pinned(b);    // free host memory
    magma_free_pinned(c);    // free host memory
    for(i=0; i<num_gpus; i++){
        magma_free(d_la[i] );    // free device memory
    }
    for( int dev = 0; dev < num_gpus; ++dev ) {
        magma_queue_destroy( queues[dev] );
    }
    magma_finalize();    // finalize Magma
}
//upper left corner of the expected solution:
//[
//  1.0000    1.0000    1.0000    1.0000
//  1.0000    1.0000    1.0000    1.0000
//  1.0000    1.0000    1.0000    1.0000
//  1.0000    1.0000    1.0000    1.0000
//];
//lapackf77_sgetrf time: 1.39675 sec.
//magma_sgetrf_mgpu time: 0.28165 sec.
//upper left corner of the solution
// from sgetrf_mgpu+sgetrs:
//[
//  0.9682    0.9682    0.9682    0.9682
//  1.0134    1.0134    1.0134    1.0134
//  1.0147    1.0147    1.0147    1.0147
//  1.0034    1.0034    1.0034    1.0034
//];

```

4.3.18 magma_dgetrf_mgpu - LU factorization in double precision on multiple GPUs

The function `magma_dgetrf_mgpu` computes in double precision an LU factorization of a general $m \times n$ matrix A using partial pivoting with row interchanges:

$$A = P L U,$$

where P is a permutation matrix, L is lower triangular with unit diagonal, and U is upper diagonal. The blocks of matrix A to be factored and the blocks of factors L, U are distributed on `num_gpus` devices. The information on the interchanged rows is contained in `ipiv`. See [magma-X.Y.Z/src/dgetrf_mgpu.cpp](#) for more details.

Using the obtained factorization one can replace the problem of solving a general linear system by solving two triangular systems with matrices L and U respectively. The Lapack function `lapackf77_dgetrs` uses the L, U factors copied from `num_gpus` devices to solve in double precision a general linear system

$$A X = B.$$

The right hand side B is a matrix defined on the host. On exit it is replaced by the solution.

```
#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b) (((a)<(b))?(a):(b))
int main( int argc, char** argv)
{
    magma_init(); // initialize Magma
    int num_gpus = 1;
    magma_setdevice(0);
    magma_queue_t queues[num_gpus];
    for( int dev = 0; dev < num_gpus; ++dev ) {
        magma_queue_create( dev, &queues[dev] );
    }
    magma_int_t err;
    real_Double_t cpu_time, gpu_time;
    magma_int_t m = 8192, n = 8192; // a,r - mxn matrices
    magma_int_t nrhs = 100; // b - nxnrhs, c - mxnrhs matrices
    magma_int_t *ipiv; // array of indices of interchanged rows
    magma_int_t n2=m*n; // size of a,r
    magma_int_t nnrhs=n*nrhs; // size of b
    magma_int_t mnrhs=m*nrhs; // size of c
    double *a, *r; // a,r - mxn matrices on the host
    double *b, *c; // b - nxnrhs, c - mxnrhs matrices on the host
    magmaDouble_ptr d_la[num_gpus];
    double alpha=1.0, beta=0.0; // alpha=1, beta=0
```

```

magma_int_t  n_local;
magma_int_t  ione = 1, info;
magma_int_t  i, min_mn=min(m,n), nb;
magma_int_t  ldn_local; // mxldn_local - size of the part of a
magma_int_t  ISEED[4] = {0,0,0,1}; // on i-th device
nb =magma_get_dgetrf_nb(m,n); //optim.block size for dgetrf
// allocate memory on cpu
ipiv=(magma_int_t*)malloc(min_mn*sizeof(magma_int_t));
// host memory for ipiv
err = magma_dmalloc_cpu(&a,n2); // host memory for a
err = magma_dmalloc_pinned(&r,n2); // host memory for r
err = magma_dmalloc_pinned(&b,nnrhs); // host memory for b
err = magma_dmalloc_pinned(&c,mnrhs); // host memory for c
// allocate device memory on num_gpus devices
for(i=0; i<num_gpus; i++){
    n_local = ((n/nb)/num_gpus)*nb;
    if (i < (n/nb)%num_gpus)
        n_local += nb;
    else if (i == (n/nb)%num_gpus)
        n_local += n%nb;
    ldn_local = ((n_local+31)/32)*32;
    magma_setdevice(i);
    err = magma_dmalloc(&d_la[i],m*ldn_local); //device memory
} // on i-th device
magma_setdevice(0);
// generate matrices
lapackf77_dlarnv( &ione, ISEED, &n2, a ); // randomize a
lapackf77_dlaset(MagmaFullStr,&n,&nnrhs,&alpha,&alpha,b,&n);
// b - nxnrhs matrix of ones
lapackf77_dlacpy( MagmaFullStr,&m,&n,a,&m,r,&m); //a->r
printf("upper left corner of the expected solution:\n");
magma_dprint(4,4,b,m); // part of the expected solution
blasf77_dgemm("N","N",&m,&nnrhs,&n,&alpha,a,&m,b,&m,
    &beta,c,&m); // right hand side c=a*b
// LAPACK version of LU decomposition
cpu_time=magma_wtime();
lapackf77_dgetrf(&m, &n, a, &m, ipiv, &info);
cpu_time=magma_wtime()-cpu_time;
printf("lapackf77_dgetrf time: %7.5f sec.\n",cpu_time);
// copy the corresponding parts of the matrix r to num_gpus
magma_dsetmatrix_1D_col_bcyclic( num_gpus, m, n, nb, r, m,
    d_la, m, queues );
// MAGMA
// LU decomposition on num_gpus devices with partial pivoting
// and row interchanges, row i is interchanged with row ipiv(i)
gpu_time = magma_sync_wtime(NULL);

magma_dgetrf_mgpu( num_gpus, m, n, d_la, m, ipiv, &info);

gpu_time = magma_sync_wtime(NULL)-gpu_time;
printf("magma_dgetrf_mgpu time: %7.5f sec.\n",gpu_time);
// copy the decomposition from num_gpus devices to r on the

```

```

// host
magma_dgetmatrix_1D_col_bccyclic( num_gpus, m, n, nb, d_la,
                                  m, r, m, queues );

magma_setdevice(0);
// solve on the host the system r*x=c; x overwrites c,
// using the LU decomposition obtained on num_gpus devices
lapackf77_dgetrs("N",&m,&nrhs,r,&m,ipiv,c,&m,&info);
// print part of the solution from dgetrf_mgpu and dgetrs
printf("upper left corner of the solution \n\
from dgetrf_mgpu+dgetrs:\n"); // part of the solution from
magma_dprint( 4, 4, c, m); // magma_dgetrf_mgpu + dgetrs
free(ipiv); // free host memory
free(a); // free host memory
magma_free_pinned(r); // free host memory
magma_free_pinned(b); // free host memory
magma_free_pinned(c); // free host memory
for(i=0; i<num_gpus; i++){
    magma_free(d_la[i] ); // free device memory
}
for( int dev = 0; dev < num_gpus; ++dev ) {
    magma_queue_destroy( queues[dev] );
}
magma_finalize(); // finalize Magma
}
//upper left corner of the expected solution:
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];
//lapackf77_dgetrf time: 2.82328 sec.
//magma_dgetrf_mgpu time: 1.41692 sec.
//upper left corner of the solution
// from dgetrf_mgpu+dgetrs:
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];

```

4.3.19 magma_sgetri_gpu - inverse matrix in single precision, GPU interface

This function computes in single precision the inverse A^{-1} of an $m \times m$ matrix A :

$$A A^{-1} = A^{-1} A = I.$$

It uses the LU decomposition with partial pivoting and row interchanges computed by `magma_sgetrf_gpu`. The information on pivots is contained in an array `piv`. The function uses also a workspace array `dwork` of size `ldwork`. The matrix A is defined on the device and on exit it is replaced by its inverse. See [magma-X.Y.Z/src/sgetri_gpu.cpp](#) for more details.

```
#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    double gpu_time;
    float *dwork; // dwork - workspace
    magma_int_t ldwork; // size of dwork
    magma_int_t *piv, info; // piv - array of indices of inter-
    magma_int_t m = 8192; // changed rows; a - mxm matrix
    magma_int_t mm=m*m; // size of a, r, c
    float *a; // a- mxm matrix on the host
    float *d_a; // d_a- mxm matrix a on the device
    float *d_r; // d_r- mxm matrix r on the device
    float *d_c; // d_c- mxm matrix c on the device
    magma_int_t ione = 1;
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    magma_int_t err;
    const float alpha = 1.0; // alpha=1
    const float beta = 0.0; // beta=0
    ldwork = m * magma_get_sgetri_nb( m ); // optimal block size
    // allocate matrices
    err = magma_smallocc_cpu( &a , mm ); // host memory for a
    err = magma_smallocc( &d_a, mm ); // device memory for a
    err = magma_smallocc( &d_r, mm ); // device memory for r
    err = magma_smallocc( &d_c, mm ); // device memory for c
    err = magma_smallocc( &dwork, ldwork); // dev. mem. for ldwork
    piv=(magma_int_t*)malloc(m*sizeof(magma_int_t)); // host mem.
    // generate random matrix a // for piv
    lapackf77_slarnv(&ione,ISEED,&mm,a); // randomize a
    magma_ssetmatrix( m, m, a,m, d_a,m,queue); // copy a -> d_a
    magmablas_slacpy(MagmaFull,m,m,d_a,m,d_r,m,queue); //d_a->d_r
    // find the inverse matrix: d_a*X=I using the LU factorization
    // with partial pivoting and row interchanges computed by
    // magma_sgetrf_gpu; row i is interchanged with row piv(i);
    // d_a -mxm matrix; d_a is overwritten by the inverse
    gpu_time = magma_sync_wtime(NULL);

    magma_sgetrf_gpu( m, m, d_a, m, piv, &info);
    magma_sgetri_gpu(m,d_a,m,piv,dwork,ldwork,&info);
```

```

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    magma_sgemm(MagmaNoTrans,MagmaNoTrans,m,m,m,alpha,d_a,m,
               d_r,m,beta,d_c,m,queue); // multiply a-1*a
    printf("magma_sgetrf_gpu + magma_sgetri_gpu time: %7.5f sec.\n",gpu_time);
    magma_sgetmatrix( m, m, d_c, m, a, m, queue); // copy d_c->a
    printf("upper left corner of a-1*a:\n");
    magma_sprint( 4, 4, a, m ); // part of a-1*a
    free(a); // free host memory
    free(piv); // free host memory
    magma_free(d_a); // free device memory
    magma_free(d_r); // free device memory
    magma_free(d_c); // free device memory
    magma_queue_destroy(queue); // destroy queue
    magma_finalize(); // finalize Magma
    return 0;
}
//magma_sgetrf_gpu + magma_sgetri_gpu time: 0.58294 sec.
//upper left corner of a-1*a:
//[
//  1.0000  0.0000  0.0000 -0.0000
// -0.0000  1.0000 -0.0000 -0.0000
//  0.0000 -0.0000  1.0000 -0.0000
// -0.0000  0.0000 -0.0000  1.0000
//];

```

4.3.20 magma_sgetri_gpu - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    double gpu_time;
    float *dwork; // dwork - workspace
    magma_int_t ldwork; // size of dwork
    magma_int_t *piv, info; // piv - array of indices of inter-
    magma_int_t m = 8192; // changed rows; a - mxm matrix
    magma_int_t mm=m*m; // size of a, r, c
    float *a; // a- mxm matrix
    float *r; // r- mxm matrix
    float *c; // c- mxm matrix
    magma_int_t ione = 1;
    magma_int_t ISEED[4] = { 0,0,0,1 }; // seed
    const float alpha = 1.0; // alpha=1
}

```

```

    const float beta = 0.0; // beta=0
    ldwork = m * magma_get_sgetri_nb( m ); // optimal block size
// allocate matrices
    cudaMallocManaged(&a,mm*sizeof(float)); // unified mem.for a
    cudaMallocManaged(&r,mm*sizeof(float)); // unified mem.for b
    cudaMallocManaged(&c,mm*sizeof(float)); // unified mem.for c
    cudaMallocManaged(&dwork,ldwork*sizeof(float)); //m.for dwork
    cudaMallocManaged(&piv,m*sizeof(int)); // unified mem.for piv
// generate random matrix a
    lapackf77_slarnv(&ione,ISEED,&mm,a); // randomize a
    magmablas_slacpy(MagmaFull,m,m,a,m,r,m,queue); //a->r
// find the inverse matrix: a*X=I using the LU factorization
// with partial pivoting and row interchanges computed by
// magma_sgetrf_gpu; row i is interchanged with row piv(i);
// a -mxm matrix; a is overwritten by the inverse
    gpu_time = magma_sync_wtime(NULL);

    magma_sgetrf_gpu( m, m, a, m, piv, &info);
    magma_sgetri_gpu(m,a,m,piv,dwork,ldwork,&info);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    magma_sgemm(MagmaNoTrans,MagmaNoTrans,m,m,m,alpha,a,m,
               r,m,beta,c,m,queue); // multiply a^-1*a
    printf("magma_sgetrf_gpu + magma_sgetri_gpu time: %7.5f sec.\n",gpu_time);
    magma_sgetmatrix( m, m, c, m, a, m, queue); // copy c->a
    printf("upper left corner of a^-1*a:\n");
    magma_sprint( 4, 4, a, m ); // part of a^-1*a
    magma_free(piv); // free memory
    magma_free(a); // free memory
    magma_free(r); // free memory
    magma_free(c); // free memory
    magma_queue_destroy(queue); // destroy queue
    magma_finalize(); // finalize Magma
    return 0;
}
//magma_sgetrf_gpu + magma_sgetri_gpu time: 0.53595 sec.
//upper left corner of a^-1*a:
//[
//  1.0000  0.0000  0.0000 -0.0000
// -0.0000  1.0000 -0.0000 -0.0000
//  0.0000 -0.0000  1.0000 -0.0000
// -0.0000  0.0000 -0.0000  1.0000
//];

```

4.3.21 magma_dgetri_gpu - inverse matrix in double precision, GPU interface

This function computes in double precision the inverse A^{-1} of an $m \times m$ matrix A :

$$A A^{-1} = A^{-1} A = I.$$

It uses the LU decomposition with partial pivoting and row interchanges computed by `magma_dgetrf_gpu`. The information on pivots is contained in an array `piv`. The function uses also a workspace array `dwork` of size `ldwork`. The matrix A is defined on the device and on exit it is replaced by its inverse. See [magma-X.Y.Z/src/dgetri_gpu.cpp](#) for more details.

```
#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    double gpu_time, *dwork; // dwork - workspace
    magma_int_t ldwork; // size of dwork
    magma_int_t *piv, info; // piv - array of indices of inter-
    magma_int_t m = 8192; // changed rows; a - mxm matrix
    magma_int_t mm=m*m; // size of a, r, c
    double *a; // a- mxm matrix on the host
    double *d_a; // d_a- mxm matrix a on the device
    double *d_r; // d_r- mxm matrix r on the device
    double *d_c; // d_c- mxm matrix c on the device
    magma_int_t ione = 1;
    magma_int_t ISEED[4] = { 0,0,0,1 }; // seed
    magma_int_t err;
    const double alpha = 1.0; // alpha=1
    const double beta = 0.0; // beta=0
    ldwork = m * magma_get_dgetri_nb( m ); // optimal block size
    // allocate matrices
    err = magma_dmalloc_cpu( &a , mm ); // host memory for a
    err = magma_dmalloc( &d_a, mm ); // device memory for a
    err = magma_dmalloc( &d_r, mm ); // device memory for r
    err = magma_dmalloc( &d_c, mm ); // device memory for c
    err = magma_dmalloc( &dwork, ldwork ); // dev. mem. for ldwork
    piv=(magma_int_t*)malloc(m*sizeof(magma_int_t)); // host mem.
    // generate random matrix a // for piv
    lapackf77_dlarnv(&ione,ISEED,&mm,a); // randomize a
    magma_dsetmatrix( m, m, a,m, d_a,m,queue); // copy a -> d_a
    magmablas_dlacpy(MagmaFull,m,m,d_a,m,d_r,m,queue); //d_a->d_r
    // find the inverse matrix: d_a*X=I using the LU factorization
    // with partial pivoting and row interchanges computed by
    // magma_dgetrf_gpu; row i is interchanged with row piv(i);
    // d_a -mxm matrix; d_a is overwritten by the inverse
    gpu_time = magma_sync_wtime(NULL);

    magma_dgetrf_gpu( m, m, d_a, m, piv, &info);
    magma_dgetri_gpu(m,d_a,m,piv,dwork,ldwork,&info);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
```



```

magma_dgemm(MagmaNoTrans, MagmaNoTrans, m, m, m, alpha, d_a, m,
            d_r, m, beta, d_c, m, queue); // multiply  $a^{-1} * a$ 
printf("magma_dgetrf_gpu + magma_dgetri_gpu time: %7.5f sec.\n", gpu_time);
magma_dgetmatrix( m, m, d_c, m, a, m, queue); // copy  $d_c \rightarrow a$ 
printf("upper left corner of  $a^{-1} * a$ :\n");
magma_dprint( 4, 4, a, m ); // part of  $a^{-1} * a$ 
free(a); // free host memory
free(piv); // free host memory
magma_free(d_a); // free device memory
magma_free(d_r); // free device memory
magma_free(d_c); // free device memory
magma_queue_destroy(queue); // destroy queue
magma_finalize(); // finalize Magma
return 0;
}
//magma_dgetrf_gpu + magma_dgetri_gpu time: 4.79694 sec.
//upper left corner of  $a^{-1} * a$ :
//[
//  1.0000  -0.0000  -0.0000  0.0000
//  0.0000   1.0000  -0.0000 -0.0000
// -0.0000  0.0000   1.0000  0.0000
// -0.0000  0.0000  -0.0000  1.0000
//];

```

4.3.22 magma_dgetri_gpu - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev, &queue);
    double gpu_time, *dwork; // dwork - workspace
    magma_int_t ldwork; // size of dwork
    magma_int_t *piv, info; // piv - array of indices of inter-
    magma_int_t m = 8192; // changed rows; a - mxm matrix
    magma_int_t mm=m*m; // size of a, r, c
    double *a; // a- mxm matrix
    double *r; // r- mxm matrix
    double *c; // c- mxm matrix
    magma_int_t ione = 1;
    magma_int_t ISEED[4] = { 0,0,0,1 }; // seed
    const double alpha = 1.0; // alpha=1
    const double beta = 0.0; // beta=0
    ldwork = m * magma_get_dgetri_nb( m ); // optimal block size

```

```

// allocate matrices
cudaMallocManaged(&a,mm*sizeof(double));// unified mem.for a
cudaMallocManaged(&r,mm*sizeof(double));// unified mem.for r
cudaMallocManaged(&c,mm*sizeof(double));// unified mem.for c
cudaMallocManaged(&dwork,ldwork*sizeof(double));//mem. dwork
cudaMallocManaged(&piv,m*sizeof(int));// unified mem.for piv
// generate random matrix a
lapackf77_dlarnv(&ione,ISEED,&mm,a);           // randomize a
magmablas_dlacpy(MagmaFull,m,m,a,m,r,m,queue); //a->r
// find the inverse matrix: a*X=I using the LU factorization
// with partial pivoting and row interchanges computed by
// magma_dgetrf_gpu; row i is interchanged with row piv(i);
// a -mxm matrix; a is overwritten by the inverse
gpu_time = magma_sync_wtime(NULL);

magma_dgetrf_gpu( m, m, a, m, piv, &info);
magma_dgetri_gpu(m,a,m,piv,dwork,ldwork,&info);

gpu_time = magma_sync_wtime(NULL)-gpu_time;
magma_dgemm(MagmaNoTrans,MagmaNoTrans,m,m,m,alpha,a,m,
            r,m,beta,c,m,queue); // multiply a^-1*a
printf("magma_dgetrf_gpu + magma_dgetri_gpu time: %7.5f sec.\n",gpu_time);

magma_dgetmatrix( m, m, c, m, a, m, queue); // copy c->a
printf("upper left corner of a^-1*a:\n");
magma_dprint( 4, 4, a, m ); // part of a^-1*a
magma_free(piv); // free memory
magma_free(a); // free memory
magma_free(r); // free memory
magma_free(c); // free memory
magma_queue_destroy(queue); // destroy queue
magma_finalize(); // finalize Magma
return 0;
}
//magma_dgetrf_gpu + magma_dgetri_gpu time: 4.77694 sec.
//upper left corner of a^-1*a:
//[
//  1.0000  -0.0000  -0.0000  0.0000
//  0.0000   1.0000  -0.0000 -0.0000
// -0.0000  0.0000   1.0000  0.0000
// -0.0000  0.0000 -0.0000   1.0000
//];

```

4.4 Cholesky decomposition and solving systems with positive definite matrices

4.4.1 magma_sposv - solve a system with a positive definite matrix in single precision, CPU interface

This function computes in single precision the solution of a real linear system

$$A X = B,$$

where A is an $m \times m$ symmetric positive definite matrix and B, X are general $m \times n$ matrices. The Cholesky factorization

$$A = \begin{cases} U^T U & \text{in MagmaUpper case,} \\ L L^T & \text{in MagmaLower case} \end{cases}$$

is used, where U is an upper triangular matrix and L is a lower triangular matrix. The matrices A, B and the solution X are defined on the host. See [magma-X.Y.Z/src/sposv.cpp](#) for more details.

```
#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#include "magma_sutil.cpp"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    double gpu_time;
    magma_int_t info;
    magma_int_t m = 8192; // a - mxm matrix
    magma_int_t n = 100; // b,c - mxn matrices
    magma_int_t mm=m*m; // size of a
    magma_int_t mn=m*n; // size of b,c
    float *a; // a- mxm matrix on the host
    float *b; // b- mxn matrix on the host
    float *c; // c- mxn matrix on the host
    magma_int_t ione = 1;
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    magma_int_t err;
    const float alpha = 1.0; // alpha=1
    const float beta = 0.0; // beta=0
    // allocate matrices on the host
    err = magma_smalloc_cpu( &a , mm ); // host memory for a
    err = magma_smalloc_cpu( &b , mn ); // host memory for b
    err = magma_smalloc_cpu( &c , mn ); // host memory for c
    // generate matrices
    lapackf77_slarnv(&ione, ISEED, &mm, a); // randomize a
    // b - mxn matrix of ones
    lapackf77_slaset(MagmaFullStr, &m, &n, &alpha, &alpha, b, &m);
```

```

// symmetrize a and increase diagonal
magma_smake_hpd( m, a, m );
printf("upper left corner of the expected solution:\n");
magma_sprint( 4, 4, b, m ); // part of the expected solution
// right hand side c=a*b
blasf77_sgemm("N","N",&m,&n,&m,&alpha,a,&m,b,&m,&beta,c,&m);
// solve the linear system a*x=c
// c -mxn matrix, a -mxm symmetric, positive def. matrix;
// c is overwritten by the solution,
// use the Cholesky factorization a=L*L^T
gpu_time = magma_sync_wtime(NULL);

magma_sposv(MagmaLower,m,n,a,m,c,m,&info);

gpu_time = magma_sync_wtime(NULL)-gpu_time;
printf("magma_sposv time: %7.5f sec.\n",gpu_time); // Magma
printf("upper left corner of the Magma solution:\n"); //time
magma_sprint( 4, 4, c, m ); // part of Magma solution
free(a); // free host memory
free(b); // free host memory
free(c); // free host memory
magma_finalize(); // finalize Magma
return 0;
}
//upper left corner of the expected solution:
//[
// 1.0000 1.0000 1.0000 1.0000
// 1.0000 1.0000 1.0000 1.0000
// 1.0000 1.0000 1.0000 1.0000
// 1.0000 1.0000 1.0000 1.0000
//];
//magma_sposv time: 0.41469 sec.
//upper left corner of the Magma solution:
//[
// 1.0000 1.0000 1.0000 1.0000
// 1.0000 1.0000 1.0000 1.0000
// 1.0000 1.0000 1.0000 1.0000
// 1.0000 1.0000 1.0000 1.0000
//];

```

4.4.2 magma_sposv - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#include "magma_dutil.cpp"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma

```

```

double    gpu_time;
magma_int_t info;
magma_int_t m = 8192;
magma_int_t n = 100;
magma_int_t mm=m*m;
magma_int_t mn=m*n;
float *a;
float *b;
float *c;
magma_int_t ione = 1;
magma_int_t ISEED[4] = { 0,0,0,1 };
const float alpha = 1.0;
const float beta = 0.0;
// allocate matrices
cudaMallocManaged(&a,mm*sizeof(float)); // unif.memory for a
cudaMallocManaged(&b,mn*sizeof(float)); // unif.memory for b
cudaMallocManaged(&c,mn*sizeof(float)); // unif.memory for c
// generate random matrices a, b;
lapackf77_slarnv(&ione,ISEED,&mm,a); // randomize a
lapackf77_slaset(MagmaFullStr,&m,&n,&n,&alpha,&alpha,b,&m);
// symmetrize a and increase diagonal
magma_smake_hpd( m, a, m );
printf("upper left corner of the expected solution:\n");
magma_sprint( 4, 4, b, m ); // part of the expected solution
// right hand side c=a*b
blasf77_sgemm("N","N",&m,&n,&n,&alpha,a,&m,b,&m,&beta,c,&m);
// solve the linear system a*x=c
// c -mxn matrix, a -mxm symmetric, positive def. matrix;
// c is overwritten by the solution,
// use the Cholesky factorization a=L*L^T
gpu_time = magma_sync_wtime(NULL);

magma_sposv(MagmaLower,m,n,a,m,c,m,&info);

gpu_time = magma_sync_wtime(NULL)-gpu_time;
printf("magma_sposv time: %7.5f sec.\n",gpu_time); // Magma
printf("upper left corner of the Magma solution:\n"); //time
magma_sprint( 4, 4, c, m ); // part of Magma solution
magma_free(a); // free memory
magma_free(b); // free memory
magma_free(c); // free memory
magma_finalize(); // finalize Magma
return 0;
}
//upper left corner of the expected solution:
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];
//magma_sposv time: 0.44253 sec.

```

```
//upper left corner of the Magma solution:
//[
//  1.0000    1.0000    1.0000    1.0000
//  1.0000    1.0000    1.0000    1.0000
//  1.0000    1.0000    1.0000    1.0000
//  1.0000    1.0000    1.0000    1.0000
//];
```

4.4.3 magma_dposv - solve a system with a positive definite matrix in double precision, CPU interface

This function computes in double precision the solution of a real linear system

$$A X = B,$$

where A is an $m \times m$ symmetric positive definite matrix and B, X are general $m \times n$ matrices. The Cholesky factorization

$$A = \begin{cases} U^T U & \text{in MagmaUpper case,} \\ L L^T & \text{in MagmaLower case} \end{cases}$$

is used, where U is an upper triangular matrix and L is a lower triangular matrix. The matrices A, B and the solution X are defined on the host. See [magma-X.Y.Z/src/dposv.cpp](#) for more details.

```
#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#include "magma_dutil.cpp"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    double gpu_time;
    magma_int_t info;
    magma_int_t m = 8192; // a - mxm matrix
    magma_int_t n = 100; // b,c - mxn matrices
    magma_int_t mm=m*m; // size of a
    magma_int_t mn=m*n; // size of b,c
    double *a; // a- mxm matrix on the host
    double *b; // b- mxn matrix on the host
    double *c; // c- mxn matrix on the host
    magma_int_t ione = 1;
    magma_int_t ISEED[4] = { 0,0,0,1 }; // seed
    magma_int_t err;
    const double alpha = 1.0; // alpha=1
    const double beta = 0.0; // beta=0
    // allocate matrices on the host
    err = magma_dmalloc_cpu( &a , mm ); // host memory for a
    err = magma_dmalloc_cpu( &b , mn ); // host memory for b
    err = magma_dmalloc_cpu( &c , mn ); // host memory for c
    // generate random matrices a, b;
```

```

    lapackf77_dlarnv(&ione, ISEED, &mm, a);           // randomize a
    lapackf77_dlaset(MagmaFullStr, &m, &n, &alpha, &alpha, b, &m);
// symmetrize a and increase diagonal
    magma_dmake_hpd( m, a, m );
    printf("upper left corner of the expected solution:\n");
    magma_dprint( 4, 4, b, m ); // part of the expected solution
// right hand side c=a*b
    blasf77_dgemm("N", "N", &m, &n, &n, &alpha, a, &m, b, &m, &beta, c, &m);
// solve the linear system a*x=c
// c -mxn matrix, a -mxm symmetric, positive def. matrix;
// c is overwritten by the solution,
// use the Cholesky factorization a=L*L^T
    gpu_time = magma_sync_wtime(NULL);

    magma_dposv(MagmaLower, m, n, a, m, c, m, &info);

    gpu_time = magma_sync_wtime(NULL) - gpu_time;
    printf("magma_dposv time: %7.5f sec.\n", gpu_time); // Magma
    printf("upper left corner of the Magma solution:\n"); //time
    magma_dprint( 4, 4, c, m );           // part of Magma solution
    free(a);                             // free host memory
    free(b);                             // free host memory
    free(c);                             // free host memory
    magma_finalize();                    // finalize Magma
    return 0;
}
//upper left corner of the expected solution:
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];
//magma_dposv time: 1.39989 sec.
//upper left corner of the Magma solution:
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];

```

4.4.4 magma_dposv - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#include "magma_dutil.cpp"
int main( int argc, char** argv ){

```

```

magma_init(); // initialize Magma
double gpu_time;
magma_int_t info;
magma_int_t m = 8192; // a - mxm matrix
magma_int_t n = 100; // b,c - mxn matrices
magma_int_t mm=m*m; // size of a
magma_int_t mn=m*n; // size of b,c
double *a; // a- mxm matrix
double *b; // b- mxn matrix
double *c; // c- mxn matrix
magma_int_t ione = 1;
magma_int_t ISEED[4] = { 0,0,0,1 }; // seed
const double alpha = 1.0; // alpha=1
const double beta = 0.0; // beta=0
// allocate matrices
cudaMallocManaged(&a,mm*sizeof(double)); // unif.memory for a
cudaMallocManaged(&b,mn*sizeof(double)); // unif.memory for b
cudaMallocManaged(&c,mn*sizeof(double)); // unif.memory for c
// generate random matrices a, b;
lapackf77_dlarnv(&ione,ISEED,&mm,&a); // randomize a
lapackf77_dlaset(MagmaFullStr,&m,&n,&n,&alpha,&alpha,&b,&m);
// symmetrize a and increase diagonal
magma_dmake_hpd( m, a, m );
printf("upper left corner of the expected solution:\n");
magma_dprint( 4, 4, b, m ); // part of the expected solution
// right hand side c=a*b
blasf77_dgemm("N","N",&m,&n,&n,&alpha,&a,&m,&b,&m,&beta,&c,&m);
// solve the linear system a*x=c
// c -mxn matrix, a -mxm symmetric, positive def. matrix;
// c is overwritten by the solution,
// use the Cholesky factorization a=L*L^T
gpu_time = magma_sync_wtime(NULL);

magma_dposv(MagmaLower,m,n,a,m,c,m,&info);

gpu_time = magma_sync_wtime(NULL)-gpu_time;
printf("magma_dposv time: %7.5f sec.\n",gpu_time); // Magma
printf("upper left corner of the Magma solution:\n"); //time
magma_dprint( 4, 4, c, m ); // part of Magma solution
magma_free(a); // free memory
magma_free(b); // free memory
magma_free(c); // free memory
magma_finalize(); // finalize Magma
return 0;
}
//upper left corner of the expected solution:
//[
// 1.0000 1.0000 1.0000 1.0000
// 1.0000 1.0000 1.0000 1.0000
// 1.0000 1.0000 1.0000 1.0000
// 1.0000 1.0000 1.0000 1.0000
//];

```



```
//magma_dposv time: 1.39497 sec.
//upper left corner of the Magma solution:
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];
```

4.4.5 magma_sposv_gpu - solve a system with a positive definite matrix in single precision, GPU interface

This function computes in single precision the solution of a real linear system

$$A X = B,$$

where A is an $m \times m$ symmetric positive definite matrix and B, X are general $m \times n$ matrices. The Cholesky factorization

$$A = \begin{cases} U^T U & \text{in MagmaUpper case,} \\ L L^T & \text{in MagmaLower case} \end{cases}$$

is used, where U is an upper triangular matrix and L is a lower triangular matrix. The matrices A, B and the solution X are defined on the device. See [magma-X.Y.Z/src/sposv_gpu.cpp](#) for more details.

```
#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#include "magma_sutil.cpp"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    double gpu_time;
    magma_int_t info;
    magma_int_t m = 8192; // a - mxm matrix
    magma_int_t n = 100; // b,c - mxn matrices
    magma_int_t mm=m*m; // size of a
    magma_int_t mn=m*n; // size of b,c
    float *a; // a- mxm matrix on the host
    float *b; // b- mxn matrix on the host
    float *c; // c- mxn matrix on the host
    float *d_a; // d_a- mxm matrix a on the device
    float *d_c; // d_c- mxn matrix c on the device
    magma_int_t ione = 1;
    magma_int_t ISEED[4] = { 0,0,0,1 }; // seed
    magma_int_t err;
```

```

    const float alpha = 1.0;           // alpha=1
    const float beta = 0.0;           // beta=0
    // allocate matrices on the host
    err = magma_smalloc_cpu( &a , mm ); // host memory for a
    err = magma_smalloc_cpu( &b , mn ); // host memory for b
    err = magma_smalloc_cpu( &c , mn ); // host memory for c
    err = magma_smalloc( &d_a, mm );    // device memory for a
    err = magma_smalloc( &d_c, mn );    // device memory for c
    // generate matrices
    lapackf77_slarnv(&ione, ISEED, &mm, a); // randomize a
    // b - mxn matrix of ones
    lapackf77_slaset(MagmaFullStr, &m, &n, &alpha, &alpha, b, &m);
    // symmetrize a and increase diagonal
    magma_smake_hpd( m, a, m );
    printf("upper left corner of the expected solution:\n");
    magma_sprint( 4, 4, b, m ); // part of the expected solution
    // right hand side c=a*b
    blasf77_sgemm("N", "N", &m, &n, &m, &alpha, a, &m, b, &m, &beta, c, &m);
    magma_ssetmatrix( m, m, a, m, d_a, m, queue); // copy a -> d_a
    magma_ssetmatrix( m, n, c, m, d_c, m, queue); // copy c -> d_c
    // solve the linear system d_a*x=d_c
    // d_c -mxn matrix, d_a -mxm symmetric, positive def. matrix;
    // d_c is overwritten by the solution
    // use the Cholesky factorization d_a=L*L^T
    gpu_time = magma_sync_wtime(NULL);

    magma_sposv_gpu(MagmaLower, m, n, d_a, m, d_c, m, &info);

    gpu_time = magma_sync_wtime(NULL) - gpu_time;
    printf("magma_sposv_gpu time: %7.5f sec.\n", gpu_time);
    magma_sgetmatrix( m, n, d_c, m, c, m, queue );
    printf("upper left corner of the Magma solution:\n");
    magma_sprint( 4, 4, c, m ); // part of Magma solution
    free(a); // free host memory
    free(b); // free host memory
    free(c); // free host memory
    magma_free(d_a); // free device memory
    magma_free(d_c); // free device memory
    magma_queue_destroy(queue); // destroy queue
    magma_finalize(); // finalize Magma
    return 0;
}
//upper left corner of the expected solution:
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];
//magma_sposv_gpu time: 0.05821 sec.
//upper left corner of the Magma solution:

```

```
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];
```

4.4.6 magma_sposv_gpu - unified memory version

```
#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#include "magma_dutil.cpp"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    double gpu_time;
    magma_int_t info;
    magma_int_t m = 8192; // a - mxm matrix
    magma_int_t n = 100; // b,c - mxn matrices
    magma_int_t mm=m*m; // size of a
    magma_int_t mn=m*n; // size of b,c
    float *a; // a- mxm matrix
    float *b; // b- mxn matrix
    float *c; // c- mxn matrix
    magma_int_t ione = 1;
    magma_int_t ISEED[4] = { 0,0,0,1 }; // seed
    const float alpha = 1.0; // alpha=1
    const float beta = 0.0; // beta=0
    // allocate matrices
    cudaMallocManaged(&a,mm*sizeof(float)); // unif.memory for a
    cudaMallocManaged(&b,mn*sizeof(float)); // unif.memory for b
    cudaMallocManaged(&c,mn*sizeof(float)); // unif.memory for c
    // generate matrices
    lapackf77_slarnv(&ione,ISEED,&mm,a); // randomize a
    // b - mxn matrix of ones
    lapackf77_slaset(MagmaFullStr,&m,&n,&alpha,&alpha,b,&m);
    // symmetrize a and increase diagonal
    magma_smake_hpd( m, a, m );
    printf("upper left corner of the expected solution:\n");
    magma_sprint( 4, 4, b, m ); // part of the expected solution
    // right hand side c=a*b
    blasf77_sgemm("N","N",&m,&n,&m,&alpha,a,&m,b,&m,&beta,c,&m);
    // solve the linear system a*x=c
    // c -mxn matrix, a -mxm symmetric, positive def. matrix;
    // c is overwritten by the solution
    // use the Cholesky factorization a=L*L^T
```

```

    gpu_time = magma_sync_wtime(NULL);

    magma_sposv_gpu(MagmaLower,m,n,a,m,c,m,&info);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("magma_sposv_gpu time: %7.5f sec.\n",gpu_time);
    printf("upper left corner of the solution:\n");
    magma_sprint( 4, 4, c, m );          // part of Magma solution
    magma_free(a);                       // free memory
    magma_free(b);                       // free memory
    magma_free(c);                       // free memory
    magma_queue_destroy(queue);          // destroy queue
    magma_finalize();                    // finalize Magma
    return 0;
}

//upper left corner of the expected solution:
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];
//magma_sposv_gpu time: 0.09663 sec.
//upper left corner of the solution:
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];

```

4.4.7 magma_dposv_gpu - solve a system with a positive definite matrix in double precision, GPU interface

This function computes in double precision the solution of a real linear system

$$A X = B,$$

where A is an $m \times m$ symmetric positive definite matrix and B, X are general $m \times n$ matrices. The Cholesky factorization

$$A = \begin{cases} U^T U & \text{in MagmaUpper case,} \\ L L^T & \text{in MagmaLower case} \end{cases}$$

is used, where U is an upper triangular matrix and L is a lower triangular matrix. The matrices A, B and the solution X are defined on the device. See [magma-X.Y.Z/src/dposv_gpu.cpp](#) for more details.

```

#include <stdio.h>
#include <cuda.h>

```

```

#include "magma_v2.h"
#include "magma_lapack.h"
#include "magma_dutil.cpp"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    double gpu_time;
    magma_int_t info;
    magma_int_t m = 8192; // a - mxm matrix
    magma_int_t n = 100; // b,c - mxn matrices
    magma_int_t mm=m*m; // size of a
    magma_int_t mn=m*n; // size of b,c
    double *a; // a- mxm matrix on the host
    double *b; // b- mxn matrix on the host
    double *c; // c- mxn matrix on the host
    double *d_a; // d_a- mxm matrix a on the device
    double *d_c; // d_c- mxn matrix c on the device
    magma_int_t ione = 1;
    magma_int_t ISEED[4] = { 0,0,0,1 }; // seed
    magma_int_t err;
    const double alpha = 1.0; // alpha=1
    const double beta = 0.0; // beta=0
    // allocate matrices on the host
    err = magma_dmalloc_cpu( &a , mm ); // host memory for a
    err = magma_dmalloc_cpu( &b , mn ); // host memory for b
    err = magma_dmalloc_cpu( &c , mn ); // host memory for c
    err = magma_dmalloc( &d_a , mm ); // device memory for a
    err = magma_dmalloc( &d_c , mn ); // device memory for c
    // generate matrices
    lapackf77_dlarnv(&ione,ISEED,&mm,a); // randomize a
    // b - mxn matrix of ones
    lapackf77_dlaset(MagmaFullStr,&m,&n,&alpha,&alpha,b,&m);
    // symmetrize a and increase diagonal
    magma_dmake_hpd( m , a , m );
    printf("upper left corner of the expected solution:\n");
    magma_dprint( 4, 4, b , m ); // part of the expected solution
    // right hand side c=a*b
    blasf77_dgemm("N","N",&m,&n,&m,&alpha,a,&m,b,&m,&beta,c,&m);
    magma_dsetmatrix( m , m , a,m , d_a,m,queue ); // copy a -> d_a
    magma_dsetmatrix( m , n , c,m , d_c,m,queue ); // copy c -> d_c
    // solve the linear system d_a*x=d_c
    // d_c -mxn matrix, d_a -mxm symmetric, positive def. matrix;
    // d_c is overwritten by the solution
    // use the Cholesky factorization d_a=L*L^T
    gpu_time = magma_sync_wtime(NULL);

    magma_dposv_gpu(MagmaLower,m,n,d_a,m,d_c,m,&info);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("magma_dposv_gpu time: %7.5f sec.\n",gpu_time);
}

```

```

    magma_dgetmatrix( m, n, d_c, m, c, m, queue );
    printf("upper left corner of the solution:\n");
    magma_dprint( 4, 4, c, m );          // part of Magma solution
    free(a);                             // free host memory
    free(b);                             // free host memory
    free(c);                             // free host memory
    magma_free(d_a);                     // free device memory
    magma_free(d_c);                     // free device memory
    magma_queue_destroy(queue);          // destroy queue
    magma_finalize();                    // finalize Magma
    return 0;
}
//upper left corner of the expected solution:
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];
//magma_dposv_gpu time: 0.93042 sec.
//upper left corner of the solution:
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];

```

4.4.8 magma_dposv_gpu - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#include "magma_dutil.cpp"
int main( int argc, char** argv ){
    magma_init();                      // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    double gpu_time;
    magma_int_t info;
    magma_int_t m = 8192;
    magma_int_t n = 100;
    magma_int_t mm=m*m;
    magma_int_t mn=m*n;
    double *a;
    double *b;
    double *c;
    magma_int_t ione = 1;
    // a - mxm matrix
    // b,c - mxn matrices
    // size of a
    // size of b,c
    // a- mxm matrix
    // b- mxn matrix
    // c- mxn matrix

```

```

    magma_int_t ISEED[4] = { 0,0,0,1 };           // seed
    const double alpha = 1.0;                     // alpha=1
    const double beta = 0.0;                      // beta=0
    // allocate matrices
    cudaMallocManaged(&a,mm*sizeof(double)); // unif.memory for a
    cudaMallocManaged(&b,mn*sizeof(double)); // unif.memory for b
    cudaMallocManaged(&c,mn*sizeof(double)); // unif.memory for c
    // generate matrices
    lapackf77_dlarnv(&ione,ISEED,&mm,a);          // randomize a
    // b - mxn matrix of ones
    lapackf77_dlaset(MagmaFullStr,&m,&n,&alpha,&alpha,b,&m);
    // symmetrize a and increase diagonal
    magma_dmake_hpd( m, a, m );
    printf("upper left corner of the expected solution:\n");
    magma_dprint( 4, 4, b, m ); // part of the expected solution
    // right hand side c=a*b
    blasf77_dgemm("N","N",&m,&n,&m,&alpha,a,&m,b,&m,&beta,c,&m);
    // solve the linear system a*x=c
    // c -mxn matrix, a -mxm symmetric, positive def. matrix;
    // c is overwritten by the solution
    // use the Cholesky factorization a=L*L^T
    gpu_time = magma_sync_wtime(NULL);

    magma_dposv_gpu(MagmaLower,m,n,a,m,c,m,&info);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("magma_dposv_gpu time: %7.5f sec.\n",gpu_time);
    printf("upper left corner of the solution:\n");
    magma_dprint( 4, 4, c, m );           // part of Magma solution
    magma_free(a);                        // free memory
    magma_free(b);                        // free memory
    magma_free(c);                        // free memory
    magma_queue_destroy(queue);           // destroy queue
    magma_finalize();                     // finalize Magma
    return 0;
}
//upper left corner of the expected solution:
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];
//magma_dposv_gpu time: 0.94532 sec.
//upper left corner of the solution:
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];

```

4.4.9 magma_spotrf, lapackf77_spotrs - Cholesky decomposition and solving a system with a positive definite matrix in single precision, CPU interface

The function `magma_spotrf` computes in single precision the Cholesky factorization for a symmetric, positive definite $m \times m$ matrix A :

$$A = \begin{cases} U^T U & \text{in MagmaUpper case,} \\ L L^T & \text{in MagmaLower case,} \end{cases}$$

where U is an upper triangular matrix and L is a lower triangular matrix. The matrix A and the factors are defined on the host. See [magma-X.Y.Z/src/spotrf.cpp](#) for more details. Using the obtained factorization the function `lapackf77_spotrs` computes on the host in single precision the solution of the linear system

$$A X = B,$$

where B, X are general $m \times n$ matrices defined on the host. The solution X overwrites B .

```
#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#include "magma_sutil.cpp"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    double gpu_time;
    magma_int_t info;
    magma_int_t m = 8192; // a - mxm matrix
    magma_int_t n = 100; // b,c - mxn matrices
    magma_int_t mm=m*m; // size of a
    magma_int_t mn=m*n; // size of b,c
    float *a; // a- mxm matrix on the host
    float *b; // b- mxn matrix on the host
    float *c; // c- mxn matrix on the host
    magma_int_t ione = 1;
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    magma_int_t err;
    const float alpha = 1.0; // alpha=1
    const float beta = 0.0; // beta=0
    // allocate matrices on the host
    err = magma_smalloc_cpu( &a , mm ); // host memory for a
    err = magma_smalloc_cpu( &b , mn ); // host memory for b
    err = magma_smalloc_cpu( &c , mn ); // host memory for c
    // generate matrices
    lapackf77_slarnv(&ione, ISEED, &mm, a); // randomize a
    // b - mxn matrix of ones
    lapackf77_slaset(MagmaFullStr, &m, &n, &alpha, &alpha, b, &m);
```



```

// symmetrize a and increase diagonal
magma_smake_hpd( m, a, m );
printf("upper left corner of of the expected solution:\n");
magma_sprint( 4, 4, b, m ); // part of the expected solution
// right hand side c=a*b
blasf77_sgemm("N","N",&m,&n,&m,&alpha,a,&m,b,&m,&beta,c,&m);
// compute the Cholesky factorization a=L*L^T for a real
// symmetric, positive definite mxm matrix a;
// using this factorization solve the linear system a*x=c
// for a general mxn matrix c, c is overwritten by the
// solution
gpu_time = magma_sync_wtime(NULL);

magma_spotrf(MagmaLower, m, a, m, &info);
lapackf77_spotrs("L",&m,&n,a,&m,c,&m,&info);

gpu_time = magma_sync_wtime(NULL)-gpu_time;
printf("magma_spotrf+spotrs time: %7.5f sec.\n",gpu_time);
printf("upper left corner of the Magma/Lapack solution:\n");
magma_sprint( 4, 4, c, m ); // part of the Magma/Lapack sol.
free(a); // free host memory
free(b); // free host memory
free(c); // free host memory
magma_finalize(); // finalize Magma
return 0;
}
//upper left corner of of the expected solution:
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];
//magma_spotrf+spotrs time: 0.49789 sec.
//upper left corner of the Magma/Lapack solution:
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];

```

4.4.10 magma_spotrf, lapackf77_spotrs - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#include "magma_dutil.cpp"

```

```

int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    double gpu_time;
    magma_int_t info;
    magma_int_t m = 8192; // a - mxm matrix
    magma_int_t n = 100; // b,c - mxn matrices
    magma_int_t mm=m*m; // size of a
    magma_int_t mn=m*n; // size of b,c
    float *a; // a- mxm matrix
    float *b; // b- mxn matrix
    float *c; // c- mxn matrix
    magma_int_t ione = 1;
    magma_int_t ISEED[4] = { 0,0,0,1 }; // seed
    // magma_int_t err;
    const float alpha = 1.0; // alpha=1
    const float beta = 0.0; // beta=0
    // allocate matrices
    cudaMallocManaged(&a,mm*sizeof(float)); // unif.memory for a
    cudaMallocManaged(&b,mn*sizeof(float)); // unif.memory for b
    cudaMallocManaged(&c,mn*sizeof(float)); // unif.memory for c
    // generate random matrices a, b;
    lapackf77_slarnv(&ione,ISEED,&mm,a); // randomize a
    // b - mxn matrix of ones
    lapackf77_slaset(MagmaFullStr,&m,&n,&alpha,&alpha,b,&m);
    // symmetrize a and increase diagonal
    magma_smake_hpd( m, a, m );
    printf("upper left corner of of the expected solution:\n");
    magma_sprint( 4, 4, b, m ); // part of the expected solution
    // right hand side c=a*b
    blasf77_sgemm("N","N",&m,&n,&m,&alpha,a,&m,b,&m,&beta,c,&m);
    // compute the Cholesky factorization a=L*L^T for a real
    // symmetric, positive definite mxm matrix a;
    // using this factorization solve the linear system a*x=c
    // for a general mxn matrix c, c is overwritten by the
    // solution
    gpu_time = magma_sync_wtime(NULL);

    magma_spotrf(MagmaLower, m, a, m, &info);
    lapackf77_spotrs("L",&m,&n,a,&m,c,&m,&info);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("magma_spotrf+spotrs time: %7.5f sec.\n",gpu_time);
    printf("upper left corner of the Magma/Lapack solution:\n");
    magma_sprint( 4, 4, c, m ); // part of the Magma/Lapack sol.
    magma_free(a); // free memory
    magma_free(b); // free memory
    magma_free(c); // free memory
    magma_finalize(); // finalize Magma
    return 0;
}
//upper left corner of of the expected solution:

```

```
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];
//magma_spotrf+spotrs time: 0.48314 sec.
//upper left corner of the Magma/Lapack solution:
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];
```

4.4.11 magma_dpotrf, lapackf77_dpotrs - Cholesky decomposition and solving a system with a positive definite matrix in double precision, CPU interface

The function `magma_dpotrf` computes in double precision the Cholesky factorization for a symmetric, positive definite $m \times m$ matrix A :

$$A = \begin{cases} U^T U & \text{in MagmaUpper case,} \\ L L^T & \text{in MagmaLower case,} \end{cases}$$

where U is an upper triangular matrix and L is a lower triangular matrix. The matrix A and the factors are defined on the host. See [magma-X.Y.Z/src/dpotrf.cpp](#) for more details. Using the obtained factorization the function `lapackf77_dpotrs` computes on the host in double precision the solution of the linear system

$$A X = B,$$

where B, X are general $m \times n$ matrices defined on the host. The solution X overwrites B .

```
#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#include "magma_dutil.cpp"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    double gpu_time;
    magma_int_t info;
    magma_int_t m = 8192; // a - mxm matrix
    magma_int_t n = 100; // b,c - mxn matrices
    magma_int_t mm=m*m; // size of a
    magma_int_t mn=m*n; // size of b,c
```

```

double *a;                // a- mxm matrix on the host
double *b;                // b- mxn matrix on the host
double *c;                // c- mxn matrix on the host
magma_int_t ione = 1;
magma_int_t ISEED[4] = { 0,0,0,1 };           // seed
magma_int_t err;
const double alpha = 1.0;                     // alpha=1
const double beta = 0.0;                      // beta=0
// allocate matrices on the host
err = magma_dmalloc_cpu( &a , mm );           // host memory for a
err = magma_dmalloc_cpu( &b , mn );           // host memory for b
err = magma_dmalloc_cpu( &c , mn );           // host memory for c
// generate random matrices a, b;
lapackf77_dlarnv(&ione, ISEED, &mm, a);       // randomize a
// b - mxn matrix of ones
lapackf77_dlaset(MagmaFullStr, &m, &n, &alpha, &alpha, b, &m);
// symmetrize a and increase diagonal
magma_dmake_hpd( m, a, m );
printf("upper left corner of of the expected solution:\n");
magma_dprint( 4, 4, b, m ); // part of the expected solution
// right hand side c=a*b
blasf77_dgemm("N", "N", &m, &n, &m, &alpha, a, &m, b, &m, &beta, c, &m);
// compute the Cholesky factorization a=L*L^T for a real
// symmetric, positive definite mxm matrix a;
// using this factorization solve the linear system a*x=c
// for a general mxn matrix c, c is overwritten by the
// solution
gpu_time = magma_sync_wtime(NULL);

magma_dpotrf(MagmaLower, m, a, m, &info);
lapackf77_dpotrs("L", &m, &n, a, &m, c, &m, &info);

gpu_time = magma_sync_wtime(NULL)-gpu_time;
printf("magma_dpotrf+dpotrs time: %7.5f sec.\n", gpu_time);
printf("upper left corner of the Magma/Lapack solution:\n");
magma_dprint( 4, 4, c, m ); // part of the Magma/Lapack sol.
free(a);                      // free host memory
free(b);                      // free host memory
free(c);                      // free host memory
magma_finalize();             // finalize Magma
return 0;
}
//upper left corner of of the expected solution:
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];
//magma_dpotrf+dpotrs time: 1.40168 sec.
//upper left corner of the Magma/Lapack solution:

```

```
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];
```

4.4.12 magma_dpotrf, lapackf77_dpotsr - unified memory version

```
#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#include "magma_dutil.cpp"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    double gpu_time;
    magma_int_t info;
    magma_int_t m = 8192; // a - mxm matrix
    magma_int_t n = 100; // b,c - mxn matrices
    magma_int_t mm=m*m; // size of a
    magma_int_t mn=m*n; // size of b,c
    double *a; // a- mxm matrix
    double *b; // b- mxn matrix
    double *c; // c- mxn matrix
    magma_int_t ione = 1;
    magma_int_t ISEED[4] = { 0,0,0,1 }; // seed
    // magma_int_t err;
    const double alpha = 1.0; // alpha=1
    const double beta = 0.0; // beta=0
    // allocate matrices
    cudaMallocManaged(&a,mm*sizeof(double)); // unif.memory for a
    cudaMallocManaged(&b,mn*sizeof(double)); // unif.memory for b
    cudaMallocManaged(&c,mn*sizeof(double)); // unif.memory for c
    // generate random matrices a, b;
    lapackf77_dlarnv(&ione,ISEED,&mm,a); // randomize a
    // b - mxn matrix of ones
    lapackf77_dlaset(MagmaFullStr,&m,&n,&alpha,&alpha,b,&m);
    // symmetrize a and increase diagonal
    magma_dmake_hpd( m, a, m );
    printf("upper left corner of of the expected solution:\n");
    magma_dprint( 4, 4, b, m ); // part of the expected solution
    // right hand side c=a*b
    blasf77_dgemm("N","N",&m,&n,&m,&alpha,a,&m,b,&m,&beta,c,&m);
    // compute the Cholesky factorization a=L*L^T for a real
    // symmetric, positive definite mxm matrix a;
    // using this factorization solve the linear system a*x=c
    // for a general mxn matrix c, c is overwritten by the
    // solution
    gpu_time = magma_sync_wtime(NULL);
```

```

magma_dpotrf(MagmaLower, m, a, m, &info);
lapackf77_dpotrs("L",&m,&n,a,&m,c,&m,&info);

gpu_time = magma_sync_wtime(NULL)-gpu_time;
printf("magma_dpotrf+dpotrs time: %7.5f sec.\n",gpu_time);
printf("upper left corner of the Magma/Lapack solution:\n");
magma_dprint( 4, 4, c, m ); // part of the Magma/Lapack sol.
magma_free(a);              // free memory
magma_free(b);              // free memory
magma_free(c);              // free memory
magma_finalize();           // finalize Magma
return 0;
}
//upper left corner of of the expected solution:
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];
//magma_dpotrf+dpotrs time: 1.30345 sec.
//upper left corner of the Magma/Lapack solution:
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];

```

4.4.13 magma_spotrf_gpu, magma_spotrs_gpu - Cholesky decomposition and solving a system with a positive definite matrix in single precision, GPU interface

The function `magma_spotrf_gpu` computes in single precision the Cholesky factorization for a symmetric, positive definite $m \times m$ matrix A :

$$A = \begin{cases} U^T U & \text{in MagmaUpper case,} \\ L L^T & \text{in MagmaLower case,} \end{cases}$$

where U is an upper triangular matrix and L is a lower triangular matrix. The matrix A and the factors are defined on the device. See [magma-X.Y.Z/src/spotrf_gpu.cpp](#) for more details. Using the obtained factorization the function `magma_spotrs_gpu` computes on the device in single precision the solution of the linear system

$$A X = B,$$

where B, X are general $m \times n$ matrices defined on the device. The solution X overwrites B .

```
#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#include "magma_sutil.cpp"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    double gpu_time;
    magma_int_t info;
    magma_int_t m = 8192; // a - mxm matrix
    magma_int_t n = 100; // b,c - mxn matrices
    magma_int_t mm=m*m; // size of a
    magma_int_t mn=m*n; // size of b,c
    float *a; // a- mxm matrix on the host
    float *b; // b- mxn matrix on the host
    float *c; // c- mxn matrix on the host
    float *d_a; // d_a- mxm matrix a on the device
    float *d_c; // d_c- mxn matrix c on the device
    magma_int_t ione = 1;
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    magma_int_t err;
    const float alpha = 1.0; // alpha=1
    const float beta = 0.0; // beta=0
    // allocate matrices on the host
    err = magma_smalloc_cpu( &a , mm ); // host memory for a
    err = magma_smalloc_cpu( &b , mn ); // host memory for b
    err = magma_smalloc_cpu( &c , mn ); // host memory for c
    err = magma_smalloc( &d_a , mm ); // device memory for a
    err = magma_smalloc( &d_c , mn ); // device memory for c
    // generate matrices
    lapackf77_slarnv(&ione,ISEED,&mm,a); // randomize a
    lapackf77_slaset(MagmaFullStr,&m,&n,&alpha,&alpha,b,&m);
    // b - mxn matrix of ones
    // symmetrize a and increase diagonal
    magma_smake_hpd( m, a, m );
    printf("upper left corner of of the expected solution:\n");
    magma_sprint( 4, 4, b, m ); // part of the expected solution
    // right hand side c=a*b
    blasf77_sgemm("N","N",&m,&n,&m,&alpha,a,&m,b,&m,&beta,c,&m);
    magma_ssetmatrix( m, m, a,m, d_a,m,queue); // copy a -> d_a
    magma_ssetmatrix( m, n, c,m, d_c,m,queue); // copy c -> d_c
    // compute the Cholesky factorization d_a=L*L^T for a real
    // symmetric, positive definite mxm matrix d_a;
    // using this factorization solve the linear system d_a*x=d_c
    // for a general mxn matrix d_c, d_c is overwritten by the
    // solution
    gpu_time = magma_sync_wtime(NULL);
```

```

magma_spotrf_gpu(MagmaLower, m, d_a, m, &info);
magma_spotrs_gpu(MagmaLower, m, n, d_a, m, d_c, m, &info);

gpu_time = magma_sync_wtime(NULL) - gpu_time;
printf("magma_spotrf_gpu+magma_spotrs_gpu time: %7.5f sec.\n",
      gpu_time);
magma_sgetmatrix( m, n, d_c, m, c, m, queue); // copy d_c -> c
printf("upper left corner of the Magma solution:\n");
magma_sprint( 4, 4, c, m ); // part of the Magma solution
free(a); // free host memory
free(b); // free host memory
free(c); // free host memory
magma_free(d_a); // free device memory
magma_free(d_c); // free device memory
magma_queue_destroy(queue); // destroy queue
magma_finalize(); // finalize Magma
return 0;
}
//upper left corner of of the expected solution:
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];
//magma_spotrf_gpu+magma_spotrs_gpu time: 0.05582 sec.
//upper left corner of the Magma solution:
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];

```

4.4.14 magma_spotrf_gpu, magma_spotrs_gpu - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#include "magma_dutil.cpp"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    double gpu_time;
    magma_int_t info;
    magma_int_t m = 8192; // a - mxm matrix

```



```

magma_int_t n = 100; // b,c - mxn matrices
magma_int_t mm=m*m; // size of a
magma_int_t mn=m*n; // size of b,c
float *a; // a- mxm matrix
float *b; // b- mxn matrix
float *c; // c- mxn matrix
magma_int_t ione = 1;
magma_int_t ISEED[4] = {0,0,0,1}; // seed
const float alpha = 1.0; // alpha=1
const float beta = 0.0; // beta=0
// allocate matrices
cudaMallocManaged(&a,mm*sizeof(float)); // unif.memory for a
cudaMallocManaged(&b,mn*sizeof(float)); // unif.memory for b
cudaMallocManaged(&c,mn*sizeof(float)); // unif.memory for c
// generate matrices
lapackf77_slarnv(&ione,ISEED,&mm,a); // randomize a
lapackf77_slaset(MagmaFullStr,&m,&n,&n,&alpha,&alpha,b,&m);
// b - mxn matrix of ones
// symmetrize a and increase diagonal
magma_smake_hpd( m, a, m );
printf("upper left corner of of the expected solution:\n");
magma_sprint( 4, 4, b, m ); // part of the expected solution
// right hand side c=a*b
blasf77_sgemm("N","N",&m,&n,&m,&alpha,a,&m,b,&m,&beta,c,&m);
// compute the Cholesky factorization a=L*L^T for a real
// symmetric, positive definite mxm matrix a;
// using this factorization solve the linear system a*x=c
// for a general mxn matrix c, c is overwritten by the
// solution
gpu_time = magma_sync_wtime(NULL);

magma_spotrf_gpu(MagmaLower, m, a, m, &info);
magma_spotrs_gpu(MagmaLower,m,n,a,m,c,m,&info);

gpu_time = magma_sync_wtime(NULL)-gpu_time;
printf("magma_spotrf_gpu+magma_spotrs_gpu time: %7.5f sec.\n",
      gpu_time);

printf("upper left corner of the solution:\n");
magma_sprint( 4, 4, c, m ); // part of the Magma solution
magma_free(a); // free memory
magma_free(b); // free memory
magma_free(c); // free memory
magma_queue_destroy(queue); // destroy queue
magma_finalize(); // finalize Magma
return 0;
}
//upper left corner of of the expected solution:
//[
// 1.0000 1.0000 1.0000 1.0000
// 1.0000 1.0000 1.0000 1.0000
// 1.0000 1.0000 1.0000 1.0000
// 1.0000 1.0000 1.0000 1.0000

```

```
//];
//magma_spotrf_gpu+magma_spotrs_gpu time: 0.09600 sec.
//upper left corner of the solution:
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];
```

4.4.15 magma_dpotrf_gpu, magma_dpotsr_gpu - Cholesky decomposition and solving a system with a positive definite matrix in double precision, GPU interface

The function `magma_dpotrf_gpu` computes in double precision the Cholesky factorization for a symmetric, positive definite $m \times m$ matrix A :

$$A = \begin{cases} U^T U & \text{in MagmaUpper case,} \\ L L^T & \text{in MagmaLower case,} \end{cases}$$

where U is an upper triangular matrix and L is a lower triangular matrix. The matrix A and the factors are defined on the device. See [magma-X.Y.Z/src/dpotrf_gpu.cpp](#) for more details. Using the obtained factorization the function `magma_dpotsr_gpu` computes on the device in double precision the solution of the linear system

$$A X = B,$$

where B, X are general $m \times n$ matrices defined on the device. The solution X overwrites B .

```
#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#include "magma_dutil.cpp"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    double gpu_time;
    magma_int_t info;
    magma_int_t m = 8192; // a - mxm matrix
    magma_int_t n = 100; // b,c - mxn matrices
    magma_int_t mm=m*m; // size of a
    magma_int_t mn=m*n; // size of b,c
    double *a; // a- mxm matrix on the host
    double *b; // b- mxn matrix on the host
```

```

double *c;                                // c- mxn matrix on the host
double *d_a;                              // d_a- mxm matrix a on the device
double *d_c;                              // d_c- mxn matrix c on the device
magma_int_t ione = 1;
magma_int_t ISEED[4] = {0,0,0,1};        // seed
magma_int_t err;
const double alpha = 1.0;                // alpha=1
const double beta = 0.0;                 // beta=0
// allocate matrices
err = magma_dmalloc_cpu( &a , mm );      // host memory for a
err = magma_dmalloc_cpu( &b , mn );      // host memory for b
err = magma_dmalloc_cpu( &c , mn );      // host memory for c
err = magma_dmalloc( &d_a, mm );         // device memory for a
err = magma_dmalloc( &d_c, mn );         // device memory for c
// generate matrices
lapackf77_dlarnv(&ione, ISEED, &mm, a);  // randomize a
lapackf77_dlaset(MagmaFullStr, &m, &n, &n, &alpha, &alpha, b, &m);
// b - mxn matrix of ones
// symmetrize a and increase diagonal
magma_dmake_hpd( m, a, m );
printf("upper left corner of of the expected solution:\n");
magma_dprint( 4, 4, b, m ); // part of the expected solution
// right hand side c=a*b
blasf77_dgemm("N", "N", &m, &n, &m, &alpha, a, &m, b, &m, &beta, c, &m);
magma_dsetmatrix( m, m, a, m, d_a, m, queue); // copy a -> d_a
magma_dsetmatrix( m, n, c, m, d_c, m, queue); // copy c -> d_c
// compute the Cholesky factorization d_a=L*L^T for a real
// symmetric, positive definite mxm matrix d_a;
// using this factorization solve the linear system d_a*x=d_c
// for a general mxn matrix d_c, d_c is overwritten by the
// solution
gpu_time = magma_sync_wtime(NULL);

magma_dpotrf_gpu(MagmaLower, m, d_a, m, &info);
magma_dpotrs_gpu(MagmaLower, m, n, d_a, m, d_c, m, &info);

gpu_time = magma_sync_wtime(NULL)-gpu_time;
printf("magma_dpotrf_gpu+magma_dpotrs_gpu time:%7.5f sec.\n",
      gpu_time);
magma_dgetmatrix( m, n, d_c, m, c, m, queue); // copy d_c -> c
printf("upper left corner of the solution:\n");
magma_dprint( 4, 4, c, m ); // part of the Magma solution
free(a); // free host memory
free(b); // free host memory
free(c); // free host memory
magma_free(d_a); // free device memory
magma_free(d_c); // free device memory
magma_queue_destroy(queue); // destroy queue
magma_finalize(); // finalize Magma
return 0;
}
//upper left corner of of the expected solution:

```

```
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];
//magma_dpotrf_gpu+magma_dpotrs_gpu time: 0.93016 sec.
//upper left corner of the solution:
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];
```

4.4.16 magma_dpotrf_gpu, magma_dpotrs_gpu - unified memory version

```
#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#include "magma_dutil.cpp"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    double gpu_time;
    magma_int_t info;
    magma_int_t m = 8192; // a - mxm matrix
    magma_int_t n = 100; // b,c - mxn matrices
    magma_int_t mm=m*m; // size of a
    magma_int_t mn=m*n; // size of b,c
    double *a; // a- mxm matrix
    double *b; // b- mxn matrix
    double *c; // c- mxn matrix
    magma_int_t ione = 1;
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    const double alpha = 1.0; // alpha=1
    const double beta = 0.0; // beta=0
    // allocate matrices
    cudaMallocManaged(&a,mm*sizeof(double)); // unif.memory for a
    cudaMallocManaged(&b,mn*sizeof(double)); // unif.memory for b
    cudaMallocManaged(&c,mn*sizeof(double)); // unif.memory for c
    // generate matrices
    lapackf77_dlarnv(&ione,ISEED,&mm,a); // randomize a
    lapackf77_dlaset(MagmaFullStr,&m,&n,&n,&alpha,&alpha,b,&m);
    // b - mxn matrix of ones
```

```

// symmetrize a and increase diagonal
magma_dmake_hpd( m, a, m );
printf("upper left corner of of the expected solution:\n");
magma_dprint( 4, 4, b, m ); // part of the expected solution
// right hand side c=a*b
blasf77_dgemm("N","N",&m,&n,&m,&alpha,a,&m,b,&m,&beta,c,&m);
// compute the Cholesky factorization a=L*L^T for a real
// symmetric, positive definite mxm matrix a;
// using this factorization solve the linear system a*x=c
// for a general mxn matrix c, c is overwritten by the
// solution
gpu_time = magma_sync_wtime(NULL);

magma_dpotrf_gpu(MagmaLower, m, a, m, &info);
magma_dpotrs_gpu(MagmaLower,m,n,a,m,c,m,&info);

gpu_time = magma_sync_wtime(NULL)-gpu_time;
printf("magma_dpotrf_gpu+magma_dpotrs_gpu time: %7.5f sec.\n",
      gpu_time);

printf("upper left corner of the solution:\n");
magma_dprint( 4, 4, c, m ); // part of the Magma solution

magma_free(a); // free memory
magma_free(b); // free memory
magma_free(c); // free memory
magma_queue_destroy(queue); // destroy queue
magma_finalize(); // finalize Magma
return 0;
}
//upper left corner of of the expected solution:
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];
//magma_dpotrf_gpu+magma_dpotrs_gpu time: 0.95875 sec.
//upper left corner of the solution:
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];

```

4.4.17 magma_spotrf_mgpu, lapackf77_spotrs - Cholesky decomposition on multiple GPUs and solving a system with a positive definite matrix in single precision

The function `magma_spotrf_mgpu` computes in single precision the Cholesky factorization for a symmetric, positive definite $m \times m$ matrix A :

$$A = \begin{cases} U^T U & \text{in MagmaUpper case,} \\ L L^T & \text{in MagmaLower case,} \end{cases}$$

where U is an upper triangular matrix and L is a lower triangular matrix. The matrix A and the factors are distributed to `num_gpus` devices. See [magma-X.Y.Z/src/spotrf_mgpu.cpp](#) for more details. Using the obtained factorization, after gathering the factors to some common matrix on the host, the function `lapackf77_spotrs` computes in single precision on the host the solution of the linear system

$$A X = B,$$

where B, X are general $m \times n$ matrices defined on the host. The solution X overwrites B .

```
#include <stdio.h>
#include <cublas.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#include "magma_sutil.cpp"
int main( int argc, char** argv) {
    magma_init(); // initialize Magma
    int num_gpus = 1;
    magma_setdevice(0);
    magma_queue_t queues[num_gpus];
    for( int dev = 0; dev < num_gpus; ++dev ) {
        magma_queue_create( dev, &queues[dev] );
    }
    double    cpu_time, gpu_time;
    magma_int_t err;
    magma_int_t m = 8192; // a,r - mxm matrices
    magma_int_t nrhs =100; // b,c - mxnrhs matrices
    magma_int_t mm=m*m; // size of a,r
    magma_int_t mnrhs=m*nrhs; // size of b,c
    float *a, *r; // a,r - mxn matrices on the host
    float *b, *c; // b,c - mxnrhs matrices on the host
    magmaFloat_ptr d_la[num_gpus];
    float alpha=1.0, beta=0.0;
    magma_int_t mb, nb;
    magma_int_t lda=m, ldda, n_local;
    magma_int_t i, info;
    magma_int_t ione = 1 ;
```

```

magma_int_t ISEED[4] = {0,0,0,1};
nb = magma_get_spotrf_nb(m); // optimal block size for spotrf
mb = nb;
n_local = nb*(1+m/(nb*num_gpus)) * mb*((m+mb-1)/mb);
ldda = n_local;
// allocate host memory for matrices
err = magma_smalloc_pinned(&a,mm); // host memory for a
err = magma_smalloc_pinned(&r,mm); // host memory for r
err = magma_smalloc_pinned(&b,mnrhs); // host memory for b
err = magma_smalloc_pinned(&c,mnrhs); // host memory for c
// allocate local matrix on the devices
for(i=0; i<num_gpus; i++){
    magma_setdevice(i);
    err = magma_smalloc(&d_la[i],ldda); //device memory
} // on i-th device
magma_setdevice(0);
lapackf77_slarnv( &ione, ISEED, &mm, a ); // randomize a
lapackf77_slaset(MagmaFullStr,&m,&nrhs,&alpha,&alpha,b,&m);
// b - mxnrhs matrix of ones
// Symmetrize a and increase diagonal
magma_smake_hpd( m, a, m );
// copy a -> r
lapackf77_slacpy( MagmaFullStr,&m,&m,a,&lda,r,&lda);
printf("upper left corner of the expected solution:\n");
magma_sprint(4,4,b,m); // expected solution
blasf77_sgemm("N","N",&m,&nrhs,&m,&alpha,a,&m,b,&m,&beta,
c,&m); // right hand c=a*b

// MAGMA
// distribute the matrix a to num_gpus devices
// going through each block-row
ldda = (1+m/(nb*num_gpus))*nb;
magma_ssetmatrix_1D_row_bcyclic( num_gpus, m, m, nb, r, lda,
d_la, ldda, queues );

magma_setdevice(0);
gpu_time = magma_sync_wtime(NULL);
// compute the Cholesky factorization a=L*L^T on num_gpus
// devices, blocks of a and blocks of factors are distributed
// to num_gpus devices

magma_spotrf_mgpu(num_gpus,MagmaLower,m,d_la,ldda,&info);

gpu_time = magma_sync_wtime(NULL)-gpu_time;
printf("magma_spotrf_mgpu time: %7.5f sec.\n", gpu_time);
// gather the resulting matrix from num_gpus devices to r
magma_sgetmatrix_1D_row_bcyclic( num_gpus, m, m, nb, d_la,
ldda, r, lda, queues );

magma_setdevice(0);
// use LAPACK to obtain the solution of a*x=c
lapackf77_spotrs("L",&m,&nrhs,r,&m,c,&m,&info);
printf("upper left corner of the solution \n\
from spotrf_mgpu+spotrs:\n");
magma_sprint( 4, 4, c, m); // Magma/Lapack solution

```

```

// LAPACK version of spotrf for time comparison
cpu_time=magma_wtime();
lapackf77_spotrf("L", &m, a, &lda, &info);
cpu_time=magma_wtime()-cpu_time;
printf("Lapack spotrf time: %7.5f sec.\n",cpu_time);
magma_free_pinned(a);           // free host memory
magma_free_pinned(r);           // free host memory
magma_free_pinned(b);           // free host memory
magma_free_pinned(c);           // free host memory
for(i=0; i<num_gpus; i++){
    magma_setdevice(i);
    magma_free(d_la[i] );       // free device memory
}
for( int dev = 0; dev < num_gpus; ++dev ) {
    magma_queue_destroy( queues[dev] );
}
magma_finalize();               // finalize Magma
}
//upper left corner of the expected solution:
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];
//magma_spotrf_mgpu time: 0.05060 sec.
//upper left corner of the solution
// from spotrf_mgpu+spotrs:
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];
//Lapack spotrf time: 0.70702 sec.

```

4.4.18 magma_dpotrf_mgpu, lapackf77_dpotrs - Cholesky decomposition and solving a system with a positive definite matrix in double precision on multiple GPUs

The function `magma_dpotrf_mgpu` computes in double precision the Cholesky factorization for a symmetric, positive definite $m \times m$ matrix A :

$$A = \begin{cases} U^T U & \text{in MagmaUpper case,} \\ L L^T & \text{in MagmaLower case,} \end{cases}$$

where U is an upper triangular matrix and L is a lower triangular matrix. The matrix A and the factors are distributed to `num_gpus` devices. See [magma-X.Y.Z/src/dpotrf_mgpu.cpp](http://magma-x.y.z/src/dpotrf_mgpu.cpp) for more details. Using the obtained factorization, after gathering the factors to some common matrix on the

host, the function `lapackf77_dpotrs` computes in double precision on the host the solution of the linear system

$$A X = B,$$

where B, X are general $m \times n$ matrices defined on the host. The solution X overwrites B .

```
#include <stdio.h>
#include <cublas.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#include "magma_dutil.cpp"
int main( int argc, char** argv) {
    magma_init(); // initialize Magma
    int num_gpus = 1;
    magma_setdevice(0);
    magma_queue_t queues[num_gpus];
    for( int dev = 0; dev < num_gpus; ++dev ) {
        magma_queue_create( dev, &queues[dev] );
    }
    double cpu_time, gpu_time;
    magma_int_t err;
    magma_int_t m = 8192; // a,r - m*m matrices
    magma_int_t nrhs = 100; // b,c - mxnrhs matrices
    magma_int_t mm=m*m; // size of a,r
    magma_int_t mnrhs=m*nrhs; // size of b,c
    double *a, *r; // a,r - mxn matrices on the host
    double *b, *c; // b,c - mxnrhs matrices on the host
    magmaDouble_ptr d_la[num_gpus];
    double alpha=1.0, beta=0.0;
    magma_int_t mb, nb;
    magma_int_t lda=m, ldda, n_local;
    magma_int_t i, info;
    magma_int_t ione = 1 ;
    magma_int_t ISEED[4] = {0,0,0,1};
    nb = magma_get_dpotrf_nb(m); // optimal block size for dpotrf
    mb = nb;
    n_local = nb*(1+m/(nb*num_gpus)) * mb*((m+mb-1)/mb);
    ldda = n_local;
    // allocate host memory for matrices
    err = magma_dmalloc_pinned(&a,mm); // host memory for a
    err = magma_dmalloc_pinned(&r,mm); // host memory for r
    err = magma_dmalloc_pinned(&b,mnrhs); // host memory for b
    err = magma_dmalloc_pinned(&c,mnrhs); // host memory for c
    // allocate local matrix on the devices
    for(i=0; i<num_gpus; i++){
        magma_setdevice(i);
        err = magma_dmalloc(&d_la[i],ldda); //device memory
    } // on i-th device
    magma_setdevice(0);
```

```

    lapackf77_dlarnv( &ione, ISEED, &mm, a );          // randomize a
    lapackf77_dlaset( MagmaFullStr, &m, &nrhs, &alpha, &alpha, b, &m );
                                     // b - mxnrhs matrix of ones
// Symmetrize a and increase diagonal
    magma_dmake_hpd( m, a, m );
// copy a -> r
    lapackf77_dlacpy( MagmaFullStr, &m, &m, a, &lda, r, &lda );
    printf("upper left corner of the expected solution:\n");
    magma_dprint(4, 4, b, m);                      // expected solution
    blasf77_dgemm("N", "N", &m, &nrhs, &m, &alpha, a, &m, b, &m, &beta,
                  c, &m); // right hand c=a*b

// MAGMA
// distribute the matrix a to num_gpus devices
// going through each block-row
    ldda = (1+m/(nb*num_gpus))*nb;
    magma_dsetmatrix_1D_row_bccyclic( num_gpus, m, m, nb, r, lda,
                                      d_la, ldda, queues );

    magma_setdevice(0);
    gpu_time = magma_sync_wtime(NULL);
// compute the Cholesky factorization a=L*L^T on num_gpus
// devices, blocks of a and blocks of factors are distributed
// to num_gpus devices

    magma_dpotrf_mgpu(num_gpus, MagmaLower, m, d_la, ldda, &info);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("magma_dpotrf_mgpu time: %7.5f sec.\n", gpu_time);
// gather the resulting matrix from num_gpus devices to r
    magma_dgetmatrix_1D_row_bccyclic( num_gpus, m, m, nb, d_la,
                                      ldda, r, lda, queues );

    magma_setdevice(0);
// use LAPACK to obtain the solution of a*x=c
    lapackf77_dpotrs("L", &m, &nrhs, r, &m, c, &m, &info);
    printf("upper left corner of the solution \n\
from dpotrf_mgpu+dpotrs:\n");
    magma_dprint( 4, 4, c, m);                      // Magma/Lapack solution
// LAPACK version of dpotrf for time comparison
    cpu_time=magma_wtime();
    lapackf77_dpotrf("L", &m, a, &lda, &info);
    cpu_time=magma_wtime()-cpu_time;
    printf("Lapack dpotrf time: %7.5f sec.\n", cpu_time);
    magma_free_pinned(a);                          // free host memory
    magma_free_pinned(r);                          // free host memory
    magma_free_pinned(b);                          // free host memory
    magma_free_pinned(c);                          // free host memory
    for(i=0; i<num_gpus; i++){
        magma_setdevice(i);
        magma_free(d_la[i] );                      // free device memory
    }
    for( int dev = 0; dev < num_gpus; ++dev ) {
        magma_queue_destroy( queues[dev] );
    }

```

```

    magma_finalize(); // finalize Magma
}
//upper left corner of the expected solution:
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];
//magma_dpotrf_mgpu time: 0.79751 sec.
//upper left corner of the solution
// from dpotrf_mgpu+dpotrs:
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];
//Lapack dpotrf time: 1.72130 sec.

```

4.4.19 magma_spotri - invert a symmetric positive definite matrix in single precision, CPU interface

This function computes in single precision the inverse A^{-1} of an $m \times m$ symmetric, positive definite matrix A :

$$A A^{-1} = A^{-1} A = I.$$

It uses the Cholesky decomposition:

$$A = \begin{cases} U^T U & \text{in MagmaUpper case,} \\ L L^T & \text{in MagmaLower case,} \end{cases}$$

computed by `magma_spotrf`. The matrix A is defined on the host and on exit it is replaced by its inverse. See [magma-X.Y.Z/src/spotri.cpp](#) for more details.

```

#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#include "magma_sutil.cpp"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    double gpu_time ;
    magma_int_t info;

```

```

magma_int_t m = 8192;                                // a - mxm matrix
magma_int_t mm=m*m;                                  // size of a, r, c
float *a;                                             // a- mxm matrix on the host
float *r;                                             // r- mxm matrix on the host
float *c;                                             // c- mxm matrix on the host
magma_int_t ione = 1;
magma_int_t ISEED[4] = {0,0,0,1};                  // seed
magma_int_t err;
const float alpha = 1.0;                            // alpha=1
const float beta = 0.0;                             // beta=0
// allocate matrices on the host
err = magma_smalloc_cpu( &a , mm );                // host memory for a
err = magma_smalloc_cpu( &r , mm );                // host memory for r
err = magma_smalloc_cpu( &c , mm );                // host memory for c
// generate random matrix a
lapackf77_slarnv(&ione, ISEED, &mm, a);            // randomize a
// symmetrize a and increase diagonal
magma_smake_hpd( m, a, m );
lapackf77_slacpy(MagmaFullStr, &m, &m, a, &m, r, &m); // a->r
// find the inverse matrix a-1: a*X=I for mxm symmetric,
// positive definite matrix a using the Cholesky decomposition
// obtained by magma_spotrf; a is overwritten by the inverse
gpu_time = magma_sync_wtime(NULL);

magma_spotrf( MagmaLower, m, a, m, &info);
magma_spotri( MagmaLower, m, a, m, &info);
//a overwritten by a-1
gpu_time = magma_sync_wtime(NULL)-gpu_time;
printf("magma_spotrf + magma_spotri time: %7.5f sec.\n", gpu_time);

// compute a-1*a
blasf77_ssymm("L", "L", &m, &m, &alpha, a, &m, r, &m, &beta, c, &m);
printf("upper left corner of a-1*a:\n");
magma_sprint( 4, 4, c, m );                        // part of a-1*a
free(a);                                           // free host memory
free(r);                                           // free host memory
free(c);                                           // free host memory
magma_finalize();                                  // finalize Magma
return 0;
}
//magma_spotrf + magma_spotri time: 0.58457 sec.
//upper left corner of a-1*a:

//[
//  1.0000  0.0000 -0.0000  0.0000
//  0.0000  1.0000  0.0000 -0.0000
// -0.0000  0.0000  1.0000  0.0000
//  0.0000 -0.0000  0.0000  1.0000
//];

```

4.4.20 magma_spotri - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#include "magma_dutil.cpp"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    double gpu_time ;
    magma_int_t info;
    magma_int_t m = 8192; // a - mxm matrix
    magma_int_t mm=m*m; // size of a, r, c
    float *a; // a- mxm matrix
    float *r; // r- mxm matrix
    float *c; // c- mxm matrix
    magma_int_t ione = 1;
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    const float alpha = 1.0; // alpha=1
    const float beta = 0.0; // beta=0
    // allocate matrices
    cudaMallocManaged(&a,mm*sizeof(float)); // unified mem.for a
    cudaMallocManaged(&r,mm*sizeof(float)); // unified mem.for r
    cudaMallocManaged(&c,mm*sizeof(float)); // unified mem.for c
    // generate random matrix a
    lapackf77_slarnv(&ione,ISEED,&mm,a); // randomize a
    // symmetrize a and increase diagonal
    magma_smake_hpd( m, a, m );
    lapackf77_slacpy(MagmaFullStr,&m,&m,a,&m,r,&m); // a->r
    // find the inverse matrix a^-1: a*X=I for mxm symmetric,
    // positive definite matrix a using the Cholesky decomposition
    // obtained by magma_spotrf; a is overwritten by the inverse
    gpu_time = magma_sync_wtime(NULL);

    magma_spotrf( MagmaLower, m, a, m, &info);
    magma_spotri( MagmaLower, m, a, m, &info);
    // a overwritten by a^-1
    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("magma_spotrf + magma_spotri time: %7.5f sec.\n",gpu_time);

    // compute a^-1*a
    blasf77_ssymm("L","L",&m,&m,&alpha,a,&m,r,&m,&beta,c,&m);
    printf("upper left corner of a^-1*a:\n");
    magma_sprint( 4, 4, c, m ); // part of a^-1*a
    magma_free(a); // free memory
    magma_free(r); // free memory
    magma_free(c); // free memory
    magma_finalize(); // finalize Magma
    return 0;
}
//magma_spotrf + magma_spotri time: 0.57705 sec.

```

```
//upper left corner of a^-1*a:
//[
//  1.0000    0.0000   -0.0000    0.0000
//  0.0000    1.0000    0.0000   -0.0000
//  -0.0000    0.0000    1.0000    0.0000
//  0.0000   -0.0000    0.0000    1.0000
//];
```

4.4.21 magma_dpotri - invert a positive definite matrix in double precision, CPU interface

This function computes in double precision the inverse A^{-1} of an $m \times m$ symmetric, positive definite matrix A :

$$A A^{-1} = A^{-1} A = I.$$

It uses the Cholesky decomposition:

$$A = \begin{cases} U^T U & \text{in MagmaUpper case,} \\ L L^T & \text{in MagmaLower case,} \end{cases}$$

computed by `magma_dpotrf`. The matrix A is defined on the host and on exit it is replaced by its inverse. See [magma-X.Y.Z/src/dpotri.cpp](#) for more details.

```
#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#include "magma_dutil.cpp"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    double gpu_time ;
    magma_int_t info;
    magma_int_t m = 8192; // a - mxm matrix
    magma_int_t mm=m*m; // size of a, r, c
    double *a; // a- mxm matrix on the host
    double *r; // r- mxm matrix on the host
    double *c; // c- mxm matrix on the host
    magma_int_t ione = 1;
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    magma_int_t err;
    const double alpha = 1.0; // alpha=1
    const double beta = 0.0; // beta=0
    // allocate matrices on the host
    err = magma_dmalloc_cpu( &a , mm ); // host memory for a
    err = magma_dmalloc_cpu( &r , mm ); // host memory for r
    err = magma_dmalloc_cpu( &c , mm ); // host memory for c
    // generate random matrix a
    lapackf77_dlarnv(&ione, ISEED, &mm, a); // randomize a
```

```

// symmetrize a and increase diagonal
magma_dmake_hpd( m, a, m );
lapackf77_dlacpy( MagmaFullStr, &m, &m, a, &m, r, &m ); // a->r
// find the inverse matrix a^-1: a*X=I for mxm symmetric,
// positive definite matrix a using the Cholesky decomposition
// obtained by magma_dpotrf; a is overwritten by the inverse
gpu_time = magma_sync_wtime(NULL);

magma_dpotrf( MagmaLower, m, a, m, &info);
magma_dpotri( MagmaLower, m, a, m, &info);
// a overwritten by a^-1
gpu_time = magma_sync_wtime(NULL)-gpu_time;
printf("magma_dpotrf + magma_dpotri time: %7.5f sec.\n", gpu_time);

// compute a^-1*a
blasf77_dsymm("L", "L", &m, &m, &alpha, a, &m, r, &m, &beta, c, &m);
printf("upper left corner of a^-1*a:\n");
magma_dprint( 4, 4, c, m ); // part of a^-1*a
free(a); // free host memory
free(r); // free host memory
free(c); // free host memory
magma_finalize(); // finalize Magma
return 0;
}
//magma_dpotrf + magma_dpotri time: 3.06706 sec.
//upper left corner of a^-1*a:
//[
//  1.0000  -0.0000  0.0000  0.0000
// -0.0000  1.0000 -0.0000 -0.0000
//  0.0000 -0.0000  1.0000 -0.0000
//  0.0000  0.0000 -0.0000  1.0000
//];

```

4.4.22 magma_dpotri - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#include "magma_dutil.cpp"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    double gpu_time ;
    magma_int_t info;
    magma_int_t m = 8192; // a - mxm matrix
    magma_int_t mm=m*m; // size of a, r, c
    double *a; // a- mxm matrix
    double *r; // r- mxm matrix
    double *c; // c- mxm matrix

```

```

    magma_int_t ione = 1;
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    const double alpha = 1.0; // alpha=1
    const double beta = 0.0; // beta=0
    // allocate matrices
    cudaMallocManaged(&a,mm*sizeof(double)); // unified mem.for a
    cudaMallocManaged(&r,mm*sizeof(double)); // unified mem.for r
    cudaMallocManaged(&c,mm*sizeof(double)); // unified mem.for c
    // generate random matrix a
    lapackf77_dlarnv(&ione,ISEED,&mm,a); // randomize a
    // symmetrize a and increase diagonal
    magma_dmake_hpd( m, a, m );
    lapackf77_dlacpy(MagmaFullStr,&m,&m,a,&m,r,&m); // a->r
    // find the inverse matrix a^-1: a*X=I for mxm symmetric,
    // positive definite matrix a using the Cholesky decomposition
    // obtained by magma_dpotrf; a is overwritten by the inverse
    gpu_time = magma_sync_wtime(NULL);

    magma_dpotrf( MagmaLower, m, a, m, &info);
    magma_dpotri( MagmaLower, m, a, m, &info);
    // a overwritten by a^-1
    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("magma_dpotrf + magma_dpotri time: %7.5f sec.\n",gpu_time);

    // compute a^-1*a
    blasf77_dsymm("L","L",&m,&m,&alpha,a,&m,r,&m,&beta,c,&m);
    printf("upper left corner of a^-1*a:\n");
    magma_dprint( 4, 4, c, m ); // part of a^-1*a
    magma_free(a); // free memory
    magma_free(r); // free memory
    magma_free(c); // free memory
    magma_finalize(); // finalize Magma
    return 0;
}
//magma_dpotrf + magma_dpotri time: 3.06806 sec.
//upper left corner of a^-1*a:
//[
//  1.0000  -0.0000  0.0000  0.0000
//  -0.0000  1.0000  -0.0000  -0.0000
//   0.0000  -0.0000  1.0000  -0.0000
//   0.0000  0.0000  -0.0000  1.0000
//];

```

4.4.23 magma_spotri_gpu - invert a positive definite matrix in single precision, GPU interface

This function computes in single precision the inverse A^{-1} of an $m \times m$ symmetric, positive definite matrix A :

$$A A^{-1} = A^{-1} A = I.$$

It uses the Cholesky decomposition:

$$A = \begin{cases} U^T U & \text{in MagmaUpper case,} \\ L L^T & \text{in MagmaLower case,} \end{cases}$$

computed by `magma_spotrf_gpu`. The matrix A is defined on the device and on exit it is replaced by its inverse. See `magma-X.Y.Z/src/spotri_gpu.cpp` for more details.

```
#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#include "magma_dutil.cpp"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    double gpu_time ;
    magma_int_t info;
    magma_int_t m = 8192; // a - mxm matrix
    magma_int_t mm=m*m; // size of a, r, c
    float *a; // a- mxm matrix on the host
    float *d_a; // d_a- mxm matrix a on the device
    float *d_r; // d_r- mxm matrix r on the device
    float *d_c; // d_c- mxm matrix c on the device
    magma_int_t ione = 1;
    magma_int_t ISEED[4] = { 0,0,0,1 }; // seed
    magma_int_t err;
    const float alpha = 1.0; // alpha=1
    const float beta = 0.0; // beta=0
    // allocate matrices on the host
    err = magma_smalloc_cpu( &a , mm ); // host memory for a
    err = magma_smalloc( &d_a, mm ); // device memory for a
    err = magma_smalloc( &d_r, mm ); // device memory for r
    err = magma_smalloc( &d_c, mm ); // device memory for c
    // generate random matrix a
    lapackf77_slarnv(&ione,ISEED,&mm,a); // randomize a
    // symmetrize a and increase diagonal
    magma_smake_hpd( m, a, m );
    magma_ssetmatrix( m, m, a,m,d_a,m,queue ); // copy a -> d_a
    magmablas_slacpy(MagmaFull,m,m,d_a,m,d_r,m,queue); //d_a->d_r
    // find the inverse matrix (d_a)^-1: d_a*X=I for mxm symmetric
    // positive definite matrix d_a using the Cholesky decompos.
    // obtained by magma_spotrf_gpu;
    // d_a is overwritten by the inverse
    gpu_time = magma_sync_wtime(NULL);

    magma_spotrf_gpu( MagmaLower, m, d_a, m, &info);
    magma_spotri_gpu( MagmaLower, m, d_a, m, &info);
```

```

//d_a overwritten by d_a^-1
gpu_time = magma_sync_wtime(NULL)-gpu_time;
// compute d_a^-1*d_a
magma_ssymm(MagmaLeft,MagmaLower,m,m,alpha,d_a,m,d_r,m,beta,
            d_c,m,queue);
printf("magma_spotrf_gpu + magma_spotri_gpu time: %7.5f sec.\n",gpu_time);
magma_sgetmatrix( m, m, d_c, m, a, m,queue); // copy d_c->a
printf("upper left corner of a^-1*a:\n");
magma_sprint( 4, 4, a, m ); // part of a^-1*a
free(a); // free host memory
magma_free(d_a); // free device memory
magma_free(d_r); // free device memory
magma_free(d_c); // free device memory
magma_queue_destroy(queue); // destroy queue
magma_finalize(); // finalize Magma
return 0;
}
//magma_spotrf_gpu + magma_spotri_gpu time: 0.16664 sec.
//upper left corner of a^-1*a:
//[
//  1.0000  0.0000  0.  0.
//  0.0000  1.0000  0.0000  0.0000
//  0.0000  0.0000  1.0000 -0.0000
//  0.0000 -0.0000 -0.0000  1.0000
//];

```

4.4.24 magma_spotri_gpu - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#include "magma_dutil.cpp"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    double gpu_time ;
    magma_int_t info;
    magma_int_t m = 8192; // a - mxm matrix
    magma_int_t mm=m*m; // size of a, r, c
    float *a; // a- mxm matrix
    float *r; // r- mxm matrix
    float *c; // c- mxm matrix
    magma_int_t ione = 1;
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    const float alpha = 1.0; // alpha=1
}

```

```

    const float beta = 0.0;                                // beta=0
// allocate matrices
    cudaMallocManaged(&a,mm*sizeof(float)); // unified mem.for a
    cudaMallocManaged(&r,mm*sizeof(float)); // unified mem.for r
    cudaMallocManaged(&c,mm*sizeof(float)); // unified mem.for c
// generate random matrix a
    lapackf77_slarnv(&ione, ISEED, &mm, a);                // randomize a
// symmetrize a and increase diagonal
    magma_smake_hpd( m, a, m );
    magma_blas_slacpy(MagmaFull,m,m,a,m,r,m,queue);        //a->r
// find the inverse matrix (a)^-1: a*X=I for mxm symmetric
// positive definite matrix a using the Cholesky factoriza-
// tion obtained by magma_spotrf_gpu;
// a is overwritten by the inverse
    gpu_time = magma_sync_wtime(NULL);

    magma_spotrf_gpu( MagmaLower, m, a, m, &info);
    magma_spotri_gpu( MagmaLower, m, a, m, &info);
                                                    //inv overwrites a

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    magma_ssymm(MagmaLeft,MagmaLower,m,m,alpha,a,m,r,m,beta,
                c,m,queue);                        // c=a^-1*a
    printf("magma_spotrf_gpu + magma_spotri_gpu time: %7.5f sec.\n",gpu_time);

    magma_sgetmatrix( m, m, c, m, a, m,queue);        // copy c->a
    printf("upper left corner of a^-1*a:\n");
    magma_sprint( 4, 4, a, m );                        // part of a^-1*a
    magma_free(a);                                     // free memory
    magma_free(r);                                     // free memory
    magma_free(c);                                     // free memory
    magma_queue_destroy(queue);                        // destroy queue
    magma_finalize();                                  // finalize Magma
    return 0;
}
//magma_spotrf_gpu + magma_spotri_gpu time: 0.15814 sec.
//upper left corner of a^-1*a:
//[
//  1.0000  0.0000  0.      0.
//  0.0000  1.0000  0.0000  0.0000
//  0.0000  0.0000  1.0000 -0.0000
//  0.0000 -0.0000 -0.0000  1.0000
//];

```

4.4.25 magma_dpotri_gpu - invert a positive definite matrix in double precision, GPU interface

This function computes in double precision the inverse A^{-1} of an $m \times m$ symmetric, positive definite matrix A :

$$A A^{-1} = A^{-1} A = I.$$

It uses the Cholesky decomposition:

$$A = \begin{cases} U^T U & \text{in MagmaUpper case,} \\ L L^T & \text{in MagmaLower case,} \end{cases}$$

computed by `magma_dpotrf_gpu`. The matrix A is defined on the device and on exit it is replaced by its inverse. See `magma-X.Y.Z/src/dpotri_gpu.cpp` for more details.

```
#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#include "magma_dutil.cpp"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create( dev,&queue);
    double gpu_time ;
    magma_int_t info;
    magma_int_t m = 8192; // a - mxm matrix
    magma_int_t mm=m*m; // size of a, r, c
    double *a; // a- mxm matrix on the host
    double *d_a; // d_a- mxm matrix a on the device
    double *d_r; // d_r- mxm matrix r on the device
    double *d_c; // d_c- mxm matrix c on the device
    magma_int_t ione = 1;
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    magma_int_t err;
    const double alpha = 1.0; // alpha=1
    const double beta = 0.0; // beta=0
    // allocate matrices on the host
    err = magma_dmalloc_cpu( &a , mm ); // host memory for a
    err = magma_dmalloc( &d_a, mm ); // device memory for a
    err = magma_dmalloc( &d_r, mm ); // device memory for r
    err = magma_dmalloc( &d_c, mm ); // device memory for c
    // generate random matrix a
    lapackf77_dlarnv(&ione,ISEED,&mm,a); // randomize a
    // symmetrize a and increase diagonal
    magma_dmake_hpd( m, a, m );
    magma_dsetmatrix( m, m, a, m,d_a,m,queue ); // copy a -> d_a

    magmablas_dlacpy(MagmaFull,m,m,d_a,m,d_r,m,queue); //d_a->d_r
    // find the inverse matrix (d_a)^-1: d_a*X=I for mxm symmetric
    // positive definite matrix d_a using the Cholesky factoriza-
    // tion obtained by magma_dpotrf_gpu;
    // d_a is overwritten by the inverse
    gpu_time = magma_sync_wtime(NULL);

    magma_dpotrf_gpu( MagmaLower, m, d_a, m, &info);
    magma_dpotri_gpu( MagmaLower, m, d_a, m, &info);
```

```

//inv overwrites d_a

gpu_time = magma_sync_wtime(NULL)-gpu_time;
magma_dsymm(MagmaLeft,MagmaLower,m,m,alpha,d_a,m,d_r,m,beta,
            d_c,m,queue); // d_c=d_a^-1*d_a
printf("magma_dpotrf_gpu + magma_dpotri_gpu time: %7.5f sec.\n",gpu_time);
magma_dgetmatrix( m, m, d_c, m, a, m,queue); // copy d_c->a
printf("upper left corner of a^-1*a:\n");
magma_dprint( 4, 4, a, m ); // part of a^-1*a
free(a); // free host memory
magma_free(d_a); // free device memory
magma_free(d_r); // free device memory
magma_free(d_c); // free device memory
magma_queue_destroy(queue); // destroy queue
magma_finalize(); // finalize Magma
return 0;
}
//magma_dpotrf_gpu + magma_dpotri_gpu time: 2.51915 sec.
//upper left corner of a^-1*a:
//[
//  1.0000 -0.0000 -0.0000  0.0000
//  0.0000  1.0000 -0.0000 -0.0000
//  0.0000 -0.0000  1.0000 -0.0000
//  0.0000  0.0000 -0.0000  1.0000
//];

```

4.4.26 magma_dpotri_gpu - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#include "magma_dutil.cpp"
int main( int argc, char** argv ){
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    double gpu_time ;
    magma_int_t info;
    magma_int_t m = 8192; // a - mxm matrix
    magma_int_t mm=m*m; // size of a, r, c
    double *a; // a- mxm matrix
    double *r; // r- mxm matrix
    double *c; // c- mxm matrix
    magma_int_t ione = 1;
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    const double alpha = 1.0; // alpha=1
}

```

```

    const double beta = 0.0;                                // beta=0
// allocate matrices
    cudaMallocManaged(&a,mm*sizeof(double)); // unified mem.for a
    cudaMallocManaged(&r,mm*sizeof(double)); // unified mem.for r
    cudaMallocManaged(&c,mm*sizeof(double)); // unified mem.for c
// generate random matrix a
    lapackf77_dlarnv(&ione, ISEED, &mm, a);                // randomize a
// symmetrize a and increase diagonal
    magma_dmake_hpd( m, a, m );
    magma_dblas_dlacpy(MagmaFull,m,m,a,m,r,m,queue);        //a->r
// find the inverse matrix (a)^-1: a*X=I for mxm symmetric
// positive definite matrix a using the Cholesky factoriza-
// tion obtained by magma_dpotrf_gpu;
// a is overwritten by the inverse
    gpu_time = magma_sync_wtime(NULL);

    magma_dpotrf_gpu( MagmaLower, m, a, m, &info);
    magma_dpotri_gpu( MagmaLower, m, a, m, &info);
                                //inv overwrites a
    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    magma_dsymm(MagmaLeft,MagmaLower,m,m,alpha,a,m,r,m,beta,
                c,m,queue);      // c=a^-1*a
    printf("magma_dpotrf_gpu + magma_dpotri_gpu time: %7.5f sec.\n",gpu_time);

    magma_dgetmatrix( m, m, c, m, a, m,queue);              // copy c->a
    printf("upper left corner of a^-1*a:\n");
    magma_dprint( 4, 4, a, m );                             // part of a^-1*a
    magma_free(a);                                           // free memory
    magma_free(r);                                           // free memory
    magma_free(c);                                           // free memory
    magma_queue_destroy(queue);                             // destroy queue
    magma_finalize();                                       // finalize Magma
    return 0;
}
//magma_dpotrf_gpu + magma_dpotri_gpu time: 2.53001 sec.
//upper left corner of a^-1*a:
//[
//  1.0000  -0.0000  -0.0000   0.0000
//  0.0000   1.0000  -0.0000  -0.0000
//  0.0000  -0.0000   1.0000  -0.0000
//  0.0000   0.0000  -0.0000   1.0000
//];

```

4.5 QR decomposition and the least squares solution of general systems

4.5.1 magma_sgels_gpu - the least squares solution of a linear system using QR decomposition in single precision, GPU interface

This function solves in single precision the least squares problem

$$\min_X \|A X - B\|,$$

where A is an $m \times n$ matrix, $m \geq n$ and B is an $m \times nrhs$ matrix, both defined on the device. In the solution the QR factorization of A is used. The solution X overwrites B . See [magma-X.Y.Z/src/sgels_gpu.cpp](https://github.com/oneapi-src/magma/blob/master/src/sgels_gpu.cpp) for more details.

```
#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b) (((a)<(b))? (a):(b))
#define max(a,b) (((a)<(b))? (b):(a))
int main( int argc, char** argv)
{
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    double gpu_time, cpu_time;
    magma_int_t m = 8192, n = 8192; // a - mxn matrix
    magma_int_t nrhs = 100; // b - nxnrhs, c - mxnrhs matrix
    float *a; // a - mxn matrix on the host
    float *b, *c; // b - nxnrhs, c - mxnrhs matrix on the host
    float *d_a, *d_c; // d_a - mxn matrix, d_c - mxnrhs matrix
                        // on the device

    magma_int_t mn = m*n; // size of a
    magma_int_t nnrhs=n*nrhs; // size of b
    magma_int_t mnrhs=m*nrhs; // size of c
    magma_int_t ldda, lddb; // leading dim. of d_a and d_c
    float *tau, *hwork, tmp[1]; // used in workspace preparation
    magma_int_t lworkgpu, lhwork; // workspace sizes
    magma_int_t info, min_mn, nb, l1, l2;
    magma_int_t ione = 1;
    const float alpha = 1.0; // alpha=1
    const float beta = 0.0; // beta=0
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    ldda = ((m+31)/32)*32; // ldda=m if 32 divides m
    lddb = ldda;
    min_mn = min(m, n);
```

```

nb = magma_get_sgeqrf_nb(m,n); //optim.block size for sgeqrf
lworkgpu = (m-n + nb)*(nrhs+2*nb);
magma_smalloc_cpu(&tau,min_mn); // host memory for tau
magma_smalloc_cpu(&a,mn); // host memory for a
magma_smalloc_cpu(&b,nrhs); // host memory for b
magma_smalloc_cpu(&c,mnrhs); // host memory for c
magma_smalloc(&d_a,ldda*n); // device memory for d_a
magma_smalloc(&d_c,lddb*nrhs); // device memory for d_c
// Get size for workspace
lhwork = -1;
lapackf77_sgeqrf(&m, &n, a, &m, tau, tmp, &lhwork, &info);
l1 = (magma_int_t)MAGMA_S_REAL( tmp[0] );
lhwork = -1;
lapackf77_sormqr( MagmaLeftStr, MagmaTransStr,
                  &m, &nrhs, &min_mn, a, &m, tau,
                  c, &m, tmp, &lhwork, &info);
l2 = (magma_int_t)MAGMA_S_REAL( tmp[0] );
lhwork = max( max( l1, l2 ), lworkgpu );
magma_smalloc_cpu(&hwork,lhwork); // host memory for worksp.
lapackf77_slarnv( &ione, ISEED, &mn, a ); // randomize a
lapackf77_slaset(MagmaFullStr,&n,&nrhs,&alpha,&alpha,b,&m);
// b - mxnrhs matrix of ones
blasf77_sgemm("N","N",&m,&nrhs,&n,&alpha,a,&m,b,&m,&beta,
              c,&m); // right hand side c=a*b
// so the exact solution is the matrix of ones
// MAGMA
magma_ssetmatrix( m, n, a,m,d_a,ldda,queue); // copy a -> d_a
magma_ssetmatrix( m, nrhs, c,m,d_c,lddb,queue); // c -> d_c
gpu_time = magma_sync_wtime(NULL);
// solve the least squares problem min ||d_a*x-d_c||
// using the QR decomposition,
// the solution overwrites d_c

magma_sgels_gpu(MagmaNoTrans, m, n, nrhs, d_a, ldda, d_c, lddb,
                hwork, lworkgpu, &info);

gpu_time = magma_sync_wtime(NULL)-gpu_time;
printf("MAGMA time: %7.3f sec.\n",gpu_time); // Magma time
// Get the solution in b
magma_sgetmatrix( n, nrhs, d_c, lddb, b,n,queue); // d_c -> b
printf("upper left corner of of the magma_sgels sol.:\n");
magma_sprint( 4, 4, b, n ); // part of the Magma QR solution
// LAPACK version of sgels
cpu_time=magma_wtime();
lapackf77_sgels( MagmaNoTransStr, &m, &n, &nrhs,
                a, &m, c, &m, hwork, &lhwork, &info);
cpu_time=magma_wtime()-cpu_time;
printf("LAPACK time: %7.3f sec.\n",cpu_time); // Lapack time
printf("upper left corner of the lapackf77_sgels sol.:\n");
magma_sprint( 4, 4, c, m ); // part of the Lapack QR solution
free(tau); // free host memory
free(a); // free host memory

```



```

    free(b);                                // free host memory
    free(c);                                // free host memory
    free(hwork);                            // free host memory
    magma_free(d_a);                        // free device memory
    magma_free(d_c);                        // free device memory
    magma_queue_destroy(queue);              // destroy queue
    magma_finalize( );                      // finalize Magma
    return EXIT_SUCCESS;
}
//MAGMA time: 0.358 sec.
//upper left corner of of the magma_sgels solution:
//[
//  0.9811  0.9811  0.9811  0.9811
//  1.0186  1.0186  1.0186  1.0186
//  1.0216  1.0216  1.0216  1.0216
//  0.9952  0.9952  0.9952  0.9952
//];
//LAPACK time: 11.352 sec.
//upper left corner of the lapackf77_sgels solution:
//[
//  0.9963  0.9963  0.9963  0.9963
//  0.9969  0.9969  0.9969  0.9969
//  0.9925  0.9925  0.9925  0.9925
//  1.0070  1.0070  1.0070  1.0070
//];

```

4.5.2 magma_sgels_gpu - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b) (((a)<(b))? (a):(b))
#define max(a,b) (((a)<(b))? (b):(a))
int main( int argc, char** argv)
{
    magma_init();                          // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    double gpu_time, cpu_time;
    magma_int_t m = 8192, n = 8192;        // a - mxn matrix
    magma_int_t nrhs = 100;                // b - nxnrhs, c - mxnrhs matrix
    float *a;                             // a - mxn matrix
    float *b, *c;                          // b - nxnrhs, c - mxnrhs matrix
    float *a1, *c1;                        // a1 - mxn matrix, c1 - mxnrhs matrix
    magma_int_t mn = m*n;                  // size of a
    magma_int_t nnrhs=n*nrhs;              // size of b
    magma_int_t mnrhs=m*nrhs;              // size of c

```

```

magma_int_t  ldda, lddb;                // leading dim of a and c
float *tau, *hwork, tmp[1]; // used in workspace preparation
magma_int_t lworkgpu, lhwork;          // workspace sizes
magma_int_t  info, min_mn, nb, l1, l2;
magma_int_t  ione = 1;
const float alpha = 1.0;                // alpha=1
const float beta = 0.0;                 // beta=0
magma_int_t ISEED[4] = {0,0,0,1};      // seed
ldda = ((m+31)/32)*32;                  // ldda=m if 32 divides m
lddb = ldda;
min_mn = min(m, n);
nb = magma_get_sgeqrf_nb(m,n); //optim. block size for sgeqrf
lworkgpu = (m-n + nb)*(nrhs+2*nb);
// prepare unified memory
cudaMallocManaged(&tau,min_mn*sizeof(float)); //u.mem.for tau
cudaMallocManaged(&a,mn*sizeof(float)); // unified mem.for a
cudaMallocManaged(&b,nrhs*sizeof(float)); //unif. mem.for b
cudaMallocManaged(&c,mnrhs*sizeof(float)); //unif. mem.for c
cudaMallocManaged(&a1,mn*sizeof(float)); //unified mem.for a1
cudaMallocManaged(&c1,mnrhs*sizeof(float)); //unif.mem for c1
// Get size for workspace
lhwork = -1;
lapackf77_sgeqrf(&m, &n, a, &m, tau, tmp, &lhwork, &info);
l1 = (magma_int_t)MAGMA_D_REAL( tmp[0] );
lhwork = -1;
lapackf77_sormqr( MagmaLeftStr, MagmaTransStr,
                  &m, &nrhs, &min_mn, a, &m, tau,
                  c, &m, tmp, &lhwork, &info);
l2 = (magma_int_t)MAGMA_D_REAL( tmp[0] );
lhwork = max( max( l1, l2 ), lworkgpu );
cudaMallocManaged(&hwork,lhwork*sizeof(float)); //mem.f.hwork
lapackf77_slarnv( &ione, ISEED, &mn, a ); // randomize a
lapackf77_slaset(MagmaFullStr,&n,&nrhs,&alpha,&alpha,b,&m);
// b - mxnrhs matrix of ones
blasf77_sgemm("N","N",&m,&nrhs,&n,&alpha,a,&m,b,&m,&beta,
              c,&m); // right hand side c=a*b
// so the exact solution is the matrix of ones
// MAGMA
magma_ssetmatrix( m, n, a,m,a1,ldda,queue); // copy a -> a1
magma_ssetmatrix( m, nrhs, c,m,c1,lddb,queue); // c -> c1
gpu_time = magma_sync_wtime(NULL);
// solve the least squares problem min ||a1*x-c1||
// using the QR decomposition,
// the solution overwrites c

magma_sgels_gpu(MagmaNoTrans, m, n, nrhs, a1, ldda, c1, lddb,
                hwork, lworkgpu, &info);

gpu_time = magma_sync_wtime(NULL)-gpu_time;
printf("MAGMA time: %7.3f sec.\n",gpu_time); // Magma time
printf("upper left corner of of the magma_sgels sol.:\n");
magma_sprint( 4, 4, c1, n ); // part of the Magma QR solution

```

```

// LAPACK version of sgels
cpu_time=magma_wtime();
lapackf77_sgels( MagmaNoTransStr, &m, &n, &nrhs,
                a, &m, c, &m, hwork, &lhwork, &info);
cpu_time=magma_wtime()-cpu_time;
printf("LAPACK time: %7.3f sec.\n",cpu_time); // Lapack time
printf("upper left corner of the lapackf77_sgels sol.:\n");
magma_sprint( 4, 4, c, m );// part of the Lapack QR solution
magma_free(tau);           // free memory
magma_free(a);             // free memory
magma_free(b);             // free memory
magma_free(c);             // free memory
magma_free(hwork);         // free memory
magma_free(a1);            // free memory
magma_free(c1);            // free memory
magma_queue_destroy(queue); // destroy queue
magma_finalize( );         // finalize Magma
return EXIT_SUCCESS;
}

//MAGMA time: 0.358 sec.
//upper left corner of of the magma_sgels sol.:
//[
// 0.9899 0.9899 0.9899 0.9899
// 1.0087 1.0087 1.0087 1.0087
// 1.0115 1.0115 1.0115 1.0115
// 0.9993 0.9993 0.9993 0.9993
//];
//LAPACK time: 12.776 sec.
//upper left corner of the lapackf77_sgels sol.:
//[
// 0.9960 0.9960 0.9960 0.9960
// 0.9966 0.9966 0.9966 0.9966
// 0.9952 0.9952 0.9952 0.9952
// 1.0066 1.0066 1.0066 1.0066
//];

```

4.5.3 magma_dgels_gpu - the least squares solution of a linear system using QR decomposition in double precision, GPU interface

This function solves in double precision the least squares problem

$$\min_X \|A X - B\|,$$

where A is an $m \times n$ matrix, $m \geq n$ and B is an $m \times nrhs$ matrix, both defined on the device. In the solution the QR factorization of A is used. The solution X overwrites B . See [magma-X.Y.Z/src/dgels_gpu.cpp](#) for more details.

```

#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b) (((a)<(b))? (a):(b))
#define max(a,b) (((a)<(b))? (b):(a))
int main( int argc, char** argv)
{
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    double gpu_time, cpu_time;
    magma_int_t m = 8192, n = 8192; // a - mxn matrix
    magma_int_t nrhs = 100; // b - nxnrhs, c - mxnrhs matrix
    double *a; // a - mxn matrix on the host
    double *b, *c; // b - nxnrhs, c - mxnrhs matrix on the host
    double *d_a, *d_c; // d_a - mxn matrix, d_c - mxnrhs matrix
                        // on the device
    magma_int_t mn = m*n; // size of a
    magma_int_t nnrhs=n*nrhs; // size of b
    magma_int_t mnrhs=m*nrhs; // size of c
    magma_int_t ldda, lddb; // leading dim of d_a and d_c
    double *tau, *hwork, tmp[1]; // used in workspace preparation
    magma_int_t lworkgpu, lhwork; // workspace sizes
    magma_int_t info, min_mn, nb, l1, l2;
    magma_int_t ione = 1;
    const double alpha = 1.0; // alpha=1
    const double beta = 0.0; // beta=0
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    ldda = ((m+31)/32)*32; // ldda=m if 32 divides m
    lddb = ldda;
    min_mn = min(m, n);
    nb = magma_get_dgeqrf_nb(m,n); //optim. block size for dgeqrf
    lworkgpu = (m-n + nb)*(nrhs+2*nb);
    magma_dmalloc_cpu(&tau,min_mn); // host memory for tau
    magma_dmalloc_cpu(&a,mn); // host memory for a
    magma_dmalloc_cpu(&b,nnrhs); // host memory for b
    magma_dmalloc_cpu(&c,mnrhs); // host memory for c
    magma_dmalloc(&d_a,ldda*n); // device memory for d_a
    magma_dmalloc(&d_c,lddb*nrhs); // device memory for d_c
    // Get size for workspace
    lhwork = -1;
    lapackf77_dgeqrf(&m, &n, a, &m, tau, tmp, &lhwork, &info);
    l1 = (magma_int_t)MAGMA_D_REAL( tmp[0] );
    lhwork = -1;
    lapackf77_dormqr( MagmaLeftStr, MagmaTransStr,
                     &m, &nrhs, &min_mn, a, &m, tau,
                     c, &m, tmp, &lhwork, &info);
    l2 = (magma_int_t)MAGMA_D_REAL( tmp[0] );
    lhwork = max( max( l1, l2 ), lworkgpu );
    magma_dmalloc_cpu(&hwork,lhwork); // host memory for worksp.

```

```

    lapackf77_dlarnv( &ione, ISEED, &mn, a ); // randomize a
    lapackf77_dlaset(MagmaFullStr,&n,&nrhs,&alpha,&alpha,b,&m);
                                // b - mxnrhs matrix of ones
    blasf77_dgemm("N","N",&m,&nrhs,&n,&alpha,a,&m,b,&m,&beta,
                  c,&m); // right hand side c=a*b
// so the exact solution is the matrix of ones
// MAGMA
    magma_dsetmatrix( m, n, a,m,d_a,ldda,queue); // copy a -> d_a
    magma_dsetmatrix( m, nrhs, c,m,d_c,lddb,queue); // c -> d_c
    gpu_time = magma_sync_wtime(NULL);
// solve the least squares problem min ||d_a*x-d_c||
// using the QR decomposition,
// the solution overwrites d_c

    magma_dgels_gpu(MagmaNoTrans, m, n, nrhs, d_a, ldda, d_c, lddb,
                    hwork, lworkgpu, &info);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("MAGMA time: %7.3f sec.\n",gpu_time); // Magma time
// Get the solution in b
    magma_dgetmatrix( n, nrhs, d_c, lddb,b,n,queue); // d_c -> b
    printf("upper left corner of of the magma_dgels sol.:\n");
    magma_dprint( 4, 4, b, n ); // part of the Magma QR solution
// LAPACK version of dgels
    cpu_time=magma_wtime();
    lapackf77_dgels( MagmaNoTransStr, &m, &n, &nrhs,
                    a, &m, c, &m, hwork, &lwork, &info);
    cpu_time=magma_wtime()-cpu_time;
    printf("LAPACK time: %7.3f sec.\n",cpu_time); // Lapack time
    printf("upper left corner of the lapackf77_dgels sol.:\n");
    magma_dprint( 4, 4, c, m ); // part of the Lapack QR solution
    free(tau); // free host memory
    free(a); // free host memory
    free(b); // free host memory
    free(c); // free host memory
    free(hwork); // free host memory
    magma_free(d_a); // free device memory
    magma_free(d_c); // free device memory
    magma_queue_destroy(queue); // destroy queue
    magma_finalize( ); // finalize Magma
    return EXIT_SUCCESS;
}
//MAGMA time: 3.157 sec.
//upper left corner of of the magma_dgels solution:
//[
// 1.0000 1.0000 1.0000 1.0000
// 1.0000 1.0000 1.0000 1.0000
// 1.0000 1.0000 1.0000 1.0000
// 1.0000 1.0000 1.0000 1.0000
//];
//LAPACK time: 18.927 sec.
//upper left corner of the lapackf77_dgels solution:

```

```
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];
```

4.5.4 magma_dgels_gpu - unified memory version

```
#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b) (((a)<(b))? (a):(b))
#define max(a,b) (((a)<(b))? (b):(a))
int main( int argc, char** argv)
{
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    double gpu_time, cpu_time;
    magma_int_t m = 8192, n = 8192; // a - mxn matrix
    magma_int_t nrhs = 100; // b - nxnrhs, c - mxnrhs matrix
    double *a; // a - mxn matrix
    double *b, *c; // b - nxnrhs, c - mxnrhs matrix
    double *a1, *c1; // a1 - mxn matrix, c1 - mxnrhs matrix
    magma_int_t mn = m*n; // size of a
    magma_int_t nnrhs=n*nrhs; // size of b
    magma_int_t mnrhs=m*nrhs; // size of c
    magma_int_t ldda, lddb; // leading dim of a and c
    double *tau, *hwork, tmp[1]; // used in workspace preparation
    magma_int_t lworkgpu, lhwork; // workspace sizes
    magma_int_t info, min_mn, nb, l1, l2;
    magma_int_t ione = 1;
    const double alpha = 1.0; // alpha=1
    const double beta = 0.0; // beta=0
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    ldda = ((m+31)/32)*32; // ldda=m if 32 divides m
    lddb = ldda;
    min_mn = min(m, n);
    nb = magma_get_dgeqrf_nb(m,n); //optim. block size for dgeqrf
    lworkgpu = (m-n + nb)*(nrhs+2*nb);
    // prepare unified memory
    cudaMallocManaged(&tau,min_mn*sizeof(double)); //mem.for tau
    cudaMallocManaged(&a,mn*sizeof(double)); // unified mem.for a
    cudaMallocManaged(&b,nnrhs*sizeof(double)); // unif.mem.for b
    cudaMallocManaged(&c,mnrhs*sizeof(double)); // unif.mem.for c
    cudaMallocManaged(&a1,mn*sizeof(double)); // unif. mem.for a1
```

```

    cudaMallocManaged(&c1,mnrhs*sizeof(double)); //uni.mem.for c1
// Get size for workspace
    lhwork = -1;
    lapackf77_dgeqrf(&m, &n, a, &m, tau, tmp, &lhwork, &info);
    l1 = (magma_int_t)MAGMA_D_REAL( tmp[0] );
    lhwork = -1;
    lapackf77_dormqr( MagmaLeftStr, MagmaTransStr,
                     &m, &nrhs, &min_mn, a, &m, tau,
                     c, &m, tmp, &lhwork, &info);
    l2 = (magma_int_t)MAGMA_D_REAL( tmp[0] );
    lhwork = max( max( l1, l2 ), lworkgpu );
    cudaMallocManaged(&hwork,lhwork*sizeof(double)); //mem.-hwork
    lapackf77_dlarnv( &ione, ISEED, &mn, a ); // randomize a
    lapackf77_dlaset(MagmaFullStr,&n,&nrhs,&alpha,&alpha,b,&m);
                                // b - mxnrhs matrix of ones
    blasf77_dgemm("N","N",&m,&nrhs,&n,&alpha,a,&m,b,&m,&beta,
                  c,&m); // right hand side c=a*b
// so the exact solution is the matrix of ones
// MAGMA
    magma_dsetmatrix( m, n, a,m,a1,ldda,queue); // copy a -> a1
    magma_dsetmatrix( m, nrhs, c,m,c1,lddb,queue); // c -> c1
    gpu_time = magma_sync_wtime(NULL);
// solve the least squares problem min ||a1*x-c1||
// using the QR decomposition,
// the solution overwrites c

    magma_dgels_gpu(MagmaNoTrans, m, n, nrhs, a1, ldda, c1, lddb,
                    hwork, lworkgpu, &info);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("MAGMA time: %7.3f sec.\n",gpu_time); // Magma time
    printf("upper left corner of of the magma_dgels sol.:\n");
    magma_dprint( 4, 4, c1, n );// part of the Magma QR solution
// LAPACK version of dgels
    cpu_time=magma_wtime();
    lapackf77_dgels( MagmaNoTransStr, &m, &n, &nrhs,
                    a, &m, c, &m, hwork, &lhwork, &info);
    cpu_time=magma_wtime()-cpu_time;
    printf("LAPACK time: %7.3f sec.\n",cpu_time); // Lapack time
    printf("upper left corner of the lapackf77_dgels sol.:\n");
    magma_dprint( 4, 4, c, m );// part of the Lapack QR solution
    magma_free(tau); // free memory
    magma_free(a); // free memory
    magma_free(b); // free memory
    magma_free(c); // free memory
    magma_free(hwork); // free memory
    magma_free(a1); // free memory
    magma_free(c1); // free memory
    magma_queue_destroy(queue); // destroy queue
    magma_finalize( ); // finalize Magma
    return EXIT_SUCCESS;
}

```

```

//MAGMA time: 3.168 sec.
//upper left corner of of the magma_dgels sol.:
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];
//LAPACK time: 19.405 sec.
//upper left corner of the lapackf77_dgels sol.:
//[
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//  1.0000  1.0000  1.0000  1.0000
//];

```

4.5.5 magma_sgels3_gpu - the least squares solution of a linear system using QR decomposition in single precision, GPU interface

This function solves in single precision the least squares problem

$$\min_X \|A X - B\|,$$

where A is an $m \times n$ matrix, $m \geq n$ and B is an $m \times nrhs$ matrix, both defined on the device. In the solution the QR factorization of A is used. The solution X overwrites B . See [magma-X.Y.Z/src/sgels3_gpu.cpp](#) for more details.

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime_api.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b) (((a)<(b))?(a):(b))
#define max(a,b) (((a)<(b))?(b):(a))
int main( int argc, char** argv)
{
    magma_init();
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    double gpu_perf, cpu_perf;
    float matnorm, work[1];
    float c_one = MAGMA_S_ONE;
    float c_neg_one = MAGMA_S_NEG_ONE;
    magma_int_t m = 8192, n = 8192, n2;
    magma_int_t nrhs = 4;
    float *a, *a2;
    // initialize Magma
    // a, a2 - mxn matrices
    // on the host

```



```

float *b, *x, *r, *tau, *hwork, tmp[1]; // b, x, r - mxnrhs
// matrices on the host
float *d_a, *d_b; // d_a - mxn matrix, d_b - mxnrhs matrix
// on the device

magma_int_t lda, ldb, ldda, lddb, lworkgpu, lhwork;
magma_int_t i, info, min_mn, nb, l1, l2;
magma_int_t *piv,ione = 1;
magma_int_t ISEED[4] = {0,0,0,1};
ldda = ((m+31)/32)*32;
lddb = ldda;
n2 = m * n;
min_mn = min(m, n);
nb = magma_get_sgeqrf_nb(m,n);
lda = ldb = m;
lworkgpu = (m-n + nb)*(nrhs+2*nb);
magma_smalloc_cpu(&tau,min_mn); // host memory for tau
magma_smalloc_cpu(&a,lda*n); // host memory for a
magma_smalloc_cpu(&a2,lda*n); // host memory for a2
magma_smalloc_cpu(&b,ldb*nrhs); // host memory for b
magma_smalloc_cpu(&x,ldb*nrhs); // host memory for x
magma_smalloc_cpu(&r,ldb*nrhs); // host memory for r
magma_smalloc(&d_a,ldda*n); // device memory for d_a
magma_smalloc(&d_b,lddb*nrhs); // device memory for d_b
piv=(magma_int_t*)malloc(n*sizeof(magma_int_t)); // host mem.
// Get size for host workspace // for piv
lhwork = -1;
lapackf77_sgeqrf(&m, &n, a, &m, tau, tmp, &lhwork, &info);
l1 = (magma_int_t)MAGMA_S_REAL( tmp[0] );
lhwork = -1;
lapackf77_sormqr( MagmaLeftStr, MagmaTransStr,
                  &m, &nrhs, &min_mn, a, &lda, tau,
                  x, &ldb, tmp, &lhwork, &info);
l2 = (magma_int_t)MAGMA_S_REAL( tmp[0] );
lhwork = max( max( l1, l2 ), lworkgpu );
magma_smalloc_cpu(&hwork,lhwork); // host memory for hwork
// randomize the matrices a, b
lapackf77_slarnv( &ione, ISEED, &n2, a );
n2 = m*nrhs;
lapackf77_slarnv( &ione, ISEED, &n2, b );
// make copies of a and b: a-> a2, b -> r
lapackf77_slacpy(MagmaFullStr,&m,&nrhs,b,&ldb,r,&ldb);
lapackf77_slacpy(MagmaFullStr,&m,&m,a,&lda,a2,&lda);
// MAGMA
magma_ssetmatrix( m,n,a,lda,d_a,ldda,queue); // copy a -> d_a
magma_ssetmatrix( m,nrhs,b,ldb,d_b,lddb,queue); // b -> d_b
gpu_perf = magma_sync_wtime(NULL);
// solve the least squares problem min ||d_a*x-d_b||
// using the QR decomposition,
// the solution overwrites d_b

magma_sgels3_gpu( MagmaNoTrans, m, n, nrhs, d_a, ldda, d_b,
                  lddb, hwork, lworkgpu, &info);

```

```

    gpu_perf = magma_sync_wtime(NULL)-gpu_perf;
    printf("MAGMA time: %7.3f sec.\n",gpu_perf);    // Magma time
// Get the solution in x
    magma_sgetmatrix( n, nrhs, d_b,lddb,x,ldb,queue); // d_b -> x
    printf("upper left corner of of the Magma solution:\n");
    magma_sprint( 4, 4, x, m );    // small part of Magma solution
// LAPACK version of sgels
    cpu_perf=magma_wtime();
    lapackf77_sgels( MagmaNoTransStr, &m, &n, &nrhs,
                    a, &lda, b, &ldb, hwork, &lhwork, &info);
    cpu_perf=magma_wtime()-cpu_perf;
    printf("LAPACK time: %7.3f sec.\n",cpu_perf); // Lapack time
    printf("upper left corner of of the Lapack solution:\n");
    magma_sprint( 4, 4, b, m ); // small part of Lapack solution
    magma_sgesv(n,nrhs,a2,n,piv,r,n,&info);
    printf("upper left corner of of the Lapack sgesv solution\n"
           "for comparison:\n");
    magma_sprint( 4, 4, r, m ); // small part of Lapack solution
                                // using LU decomposition
// Free memory
    free(tau);                  // free host memory
    free(a);                    // free host memory
    free(b);                    // free host memory
    free(x);                    // free host memory
    free(r);                    // free host memory
    free(hwork);                // free host memory
    magma_free(d_a);            // free device memory
    magma_free(d_b);            // free device memory
    magma_queue_destroy(queue); // destroy queue
    magma_finalize( );          // finalize Magma
    return EXIT_SUCCESS;
}
//MAGMA time:    0.361 sec.
//upper left corner of of the Magma solution:
//[
//    1.4699 -10.2185  -5.7395  -5.9746
//    -2.2209   2.5485  -0.9190   3.5244
//    -2.9939   1.9242  -2.9884   4.5815
//    1.8249   3.5963   3.4809  -1.2047
//];
//LAPACK time:   11.019 sec.
//upper left corner of of the Lapack solution:
//[
//    1.4686 -10.2189  -5.7409  -5.9751
//    -2.2217   2.5500  -0.9187   3.5261
//    -2.9954   1.9264  -2.9890   4.5840
//    1.8256   3.5976   3.4825  -1.2042
//];
//upper left corner of of the Lapack sgesv solution
//for comparison:

```

```
//[
//  1.4768 -10.2685 -5.7679 -5.9999
//  -2.2205  2.5265 -0.9392  3.5168
//  -2.9938  1.8580 -3.0401  4.5591
//  1.8242  3.6547  3.5300 -1.1885
//];
```

4.5.6 magma_sgels3_gpu - unified memory version

```
#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b) (((a)<(b))? (a):(b))
#define max(a,b) (((a)<(b))? (b):(a))
int main( int argc, char** argv)
{
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    double gpu_perf, cpu_perf;
    magma_int_t m = 8192, n = 8192, n2;
    magma_int_t nrhs = 4;
    float *a, *a2; // a, a2 -mxn matrices
    // a used in Lapack sgels, a2 -copy of a used in Magma sgesv
    float *b, *x, *r, *tau, *hwork, tmp[1]; // b, x, r - mxnrhs
    // matr.: b used in Lapack sgels, r copy used in Magma sgesv
    float *a1, *b1; // a1 - mxn matrix, b1 - mxnrhs matrix
    // copies of a, b used in Magma sgels
    magma_int_t lda, ldb, ldda, lddb, lworkgpu, lhwork;
    magma_int_t info, min_mn, nb, l1, l2;
    magma_int_t *piv,ione = 1;
    magma_int_t ISEED[4] = {0,0,0,1};
    ldda = ((m+31)/32)*32;
    lddb = ldda;
    n2 = m * n;
    min_mn = min(m, n);
    nb = magma_get_sgeqrf_nb(m,n);
    lda = ldb = m;
    lworkgpu = (m-n + nb)*(nrhs+2*nb);
    // prepare unified memory
    cudaMallocManaged(&tau,min_mn*sizeof(float)); //mem.for tau
    cudaMallocManaged(&a,lda*n*sizeof(float)); // unif.mem.for a
    cudaMallocManaged(&a2,lda*n*sizeof(float)); //unif.mem.for a2
    cudaMallocManaged(&b,ldb*nrhs*sizeof(float)); //uni.mem.for b
    cudaMallocManaged(&x,ldb*nrhs*sizeof(float)); //uni.mem.for x
    cudaMallocManaged(&r,ldb*nrhs*sizeof(float)); //uni.mem.for r
    cudaMallocManaged(&a1,ldda*n*sizeof(float)); //uni.mem.for a1
```

```

    cudaMallocManaged(&b1, lddb*nrhs*sizeof(float)); //mem.for b1
    cudaMallocManaged(&piv, n*sizeof(magma_int_t)); //mem.for piv
// Get size for workspace
    lwork = -1;
    lapackf77_sgeqrf(&m, &n, a, &m, tau, tmp, &lwork, &info);
    l1 = (magma_int_t)MAGMA_D_REAL( tmp[0] );
    lwork = -1;
    lapackf77_sormqr( MagmaLeftStr, MagmaTransStr,
                     &m, &nrhs, &min_mn, a, &lda, tau,
                     x, &ldb, tmp, &lwork, &info);
    l2 = (magma_int_t)MAGMA_S_REAL( tmp[0] );
    lwork = max( max( l1, l2 ), lworkgpu );
// magma_sgels3 needs this workspace
    cudaMallocManaged(&hwork, lwork*sizeof(float)); //mem.f.hwork
// randomize the matrices a, b
    lapackf77_slarnv( &ione, ISEED, &n2, a ); // random a
    n2 = m*nrhs; // size of b, x, r
    lapackf77_slarnv( &ione, ISEED, &n2, b ); // random b
// make copies of a and b: a-> a2, b -> r (they are overwrit.)
    lapackf77_slacpy(MagmaFullStr, &m, &nrhs, b, &ldb, r, &ldb);
    lapackf77_slacpy(MagmaFullStr, &m, &m, a, &lda, a2, &lda);
// copies of a, b for MAGMA
    magma_ssetmatrix(m, n, a, lda, a1, ldda, queue); // copy a -> a1
    magma_ssetmatrix( m, nrhs, b, ldb, b1, lddb, queue); // b -> b1
    gpu_perf = magma_sync_wtime(NULL);
// solve the least squares problem min ||a1*x-b1||
// using the QR decomposition,
// the solution overwrites b1
// MAGMA version

    magma_sgels3_gpu( MagmaNoTrans, m, n, nrhs, a1, ldda, b1,
                     lddb, hwork, lworkgpu, &info);

    gpu_perf = magma_sync_wtime(NULL)-gpu_perf;
    printf("MAGMA time: %7.3f sec.\n", gpu_perf); // Magma time
    printf("upper left corner of of the Magma solution:\n");
    magma_sprint( 4, 4, b1, m ); // small part of Magma solution
// LAPACK version of sgels
    cpu_perf=magma_wtime();
    lapackf77_sgels( MagmaNoTransStr, &m, &n, &nrhs,
                     a, &lda, b, &ldb, hwork, &lwork, &info);
    cpu_perf=magma_wtime()-cpu_perf;
    printf("LAPACK time: %7.3f sec.\n", cpu_perf); // Lapack time
    printf("upper left corner of of the Lapack solution:\n");
    magma_sprint( 4, 4, b, m ); // small part of Lapack solution
// MAGMA sgesv for comparison
    magma_sgesv(n, nrhs, a2, n, piv, r, n, &info);
    printf("upper left corner of of the Lapack sgesv solution\n"
           "for comparison:\n");
    magma_sprint( 4, 4, r, m ); // small part of dgesv solution
                                // using LU decomposition
// Free unified memory

```

```

    magma_free(tau);           // free memory
    magma_free(a);             // free memory
    magma_free(a2);            // free memory
    magma_free(b);             // free memory
    magma_free(x);             // free memory
    magma_free(r);             // free memory
    magma_free(hwork);         // free memory
    magma_free(a1);            // free memory
    magma_free(b1);            // free memory
    magma_queue_destroy(queue); // destroy queue
    magma_finalize( );         // finalize Magma
    return EXIT_SUCCESS;
}

//MAGMA time: 0.357 sec.
//upper left corner of of the Magma solution:
//[
//  1.4699 -10.2185 -5.7395 -5.9746
//  -2.2209  2.5485 -0.9190  3.5244
//  -2.9939  1.9242 -2.9884  4.5815
//  1.8249  3.5963  3.4809 -1.2047
//];
//LAPACK time: 12.676 sec.
//upper left corner of of the Lapack solution:
//[
//  1.4686 -10.2189 -5.7409 -5.9751
//  -2.2217  2.5500 -0.9187  3.5261
//  -2.9954  1.9264 -2.9890  4.5840
//  1.8256  3.5976  3.4825 -1.2042
//];
//upper left corner of of the Lapack sgesv solution
//for comparison:
//[
//  1.4768 -10.2685 -5.7679 -5.9999
//  -2.2205  2.5265 -0.9392  3.5168
//  -2.9938  1.8580 -3.0401  4.5591
//  1.8242  3.6547  3.5300 -1.1885
//];

```

4.5.7 magma_dgels3_gpu - the least squares solution of a linear system using QR decomposition in double precision, GPU interface

This function solves in double precision the least squares problem

$$\min_X \|A X - B\|,$$

where A is an $m \times n$ matrix, $m \geq n$ and B is an $m \times nrhs$ matrix, both defined on the device. In the solution the QR factorization of A is used. The solution X overwrites B . See [magma-X.Y.Z/src/dgels3_gpu.cpp](#) for more details.

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime_api.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b) (((a)<(b))? (a):(b))
#define max(a,b) (((a)<(b))? (b):(a))
int main( int argc, char** argv)
{
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    double gpu_perf, cpu_perf;
    double matnorm, work[1];
    double c_one = MAGMA_D_ONE;
    double c_neg_one = MAGMA_D_NEG_ONE;
    magma_int_t m = 8192, n = 8192, n2;
    magma_int_t nrhs = 4;
    double *a, *a2; // a, a2 - mxn matrices
                    // on the host
    double *b, *x, *r, *tau, *hwork, tmp[1]; // b, x, r - mxnrhs
                    // matrices on the host
    double *d_a, *d_b; // d_a - mxn matrix, d_b - mxnrhs matrix
                    // on the device

    magma_int_t lda, ldb, ldda, lddb, lworkgpu, lhwork;
    magma_int_t i, info, min_mn, nb, l1, l2;
    magma_int_t *piv, ione = 1;
    magma_int_t ISEED[4] = {0,0,0,1};
    ldda = ((m+31)/32)*32;
    lddb = ldda;
    n2 = m * n;
    min_mn = min(m, n);
    nb = magma_get_dgeqrf_nb(m,n);
    lda = ldb = m;
    lworkgpu = (m-n + nb)*(nrhs+2*nb);
    magma_dmalloc_cpu(&tau,min_mn); // host memory for tau
    magma_dmalloc_cpu(&a,lda*n); // host memory for a
    magma_dmalloc_cpu(&a2,lda*n); // host memory for a2
    magma_dmalloc_cpu(&b,ldb*nrhs); // host memory for b
    magma_dmalloc_cpu(&x,ldb*nrhs); // host memory for x
    magma_dmalloc_cpu(&r,ldb*nrhs); // host memory for r
    magma_dmalloc(&d_a,ldda*n); // device memory for d_a
    magma_dmalloc(&d_b,lddb*nrhs); // device memory for d_b
    piv=(magma_int_t*)malloc(n*sizeof(magma_int_t)); // host mem.
    // Get size for host workspace // for piv
    lhwork = -1;
    lapackf77_dgeqrf(&m, &n, a, &m, tau, tmp, &lhwork, &info);
    l1 = (magma_int_t)MAGMA_D_REAL( tmp[0] );
    lhwork = -1;
    lapackf77_dormqr( MagmaLeftStr, MagmaTransStr,
                    &m, &nrhs, &min_mn, a, &lda, tau,

```

```

        x, &lhb, tmp, &lhwork, &info);
l2 = (magma_int_t)MAGMA_S_REAL( tmp[0] );
lhwork = max( max( l1, l2 ), lworkgpu );
magma_dmalloc_cpu(&hwork,lhwork); // host memory for hwork
// randomize the matrices a, b
lapackf77_dlarnv( &ione, ISEED, &n2, a );
n2 = m*nrhs; // size of b, x, r
lapackf77_dlarnv( &ione, ISEED, &n2, b );
// make copies of a and b: a-> a2, b -> r
lapackf77_dlacpy(MagmaFullStr,&m,&nrhs,b,&lhb,r,&lhb);
lapackf77_dlacpy(MagmaFullStr,&m,&m,a,&lhb,a2,&lhb);
// MAGMA
magma_dsetmatrix(m,n,a,lhb,d_a,lhb,queue); // copy a -> d_a
magma_dsetmatrix( m,nrhs,b,lhb,d_b,lhb,queue); // b -> d_b
gpu_perf = magma_sync_wtime(NULL);
// solve the least squares problem min ||d_a*x-d_b||
// using the QR decomposition,
// the solution overwrites d_b

magma_dgels3_gpu( MagmaNoTrans, m, n, nrhs, d_a, lhb, d_b,
                  lhb, hwork, lworkgpu, &info);

gpu_perf = magma_sync_wtime(NULL)-gpu_perf;
printf("MAGMA time: %7.3f sec.\n",gpu_perf); // Magma time
// copy the solution to x
magma_dgetmatrix(n, nrhs, d_b, lhb,x,lhb,queue); // d_b -> x
printf("upper left corner of of the Magma solution:\n");
magma_dprint( 4, 4, x, m ); // small part of Magma solution
// LAPACK version of dgels
cpu_perf=magma_wtime();
lapackf77_dgels( MagmaNoTransStr, &m, &n, &nrhs,
                a, &lhb, b, &lhb, hwork, &lhwork, &info);
cpu_perf=magma_wtime()-cpu_perf;
printf("LAPACK time: %7.3f sec.\n",cpu_perf); // Lapack time
printf("upper left corner of of the Lapack solution:\n");
magma_dprint( 4, 4, b, m ); // small part of Lapack solution
magma_dgesv(n,nrhs,a2,n,piv,r,n,&info);
printf("upper left corner of of the Lapack dgesv solution\n"
       "for comparison:\n");
magma_dprint( 4, 4, r, m ); // small part of Lapack solution
                             // using LU decomposition
// Free memory
free(tau); // free host memory
free(a); // free host memory
free(b); // free host memory
free(x); // free host memory
free(r); // free host memory
free(hwork); // free host memory
magma_free(d_a); // free device memory
magma_free(d_b); // free device memory
magma_queue_destroy(queue); // destroy queue
magma_finalize( ); // finalize Magma

```

```

    return EXIT_SUCCESS;
}
//MAGMA time: 3.032 sec.
//upper left corner of of the Magma solution:
//[
// -2.9416  0.1624  0.2631 -2.0923
// -0.0242  0.5965 -0.4656 -0.3765
//  0.6595  0.5525  0.5783 -0.1609
// -0.5521 -1.2515  0.0901 -0.2223
//];
//LAPACK time: 18.957 sec.
//upper left corner of of the Lapack solution:
//[
// -2.9416  0.1624  0.2631 -2.0923
// -0.0242  0.5965 -0.4656 -0.3765
//  0.6595  0.5525  0.5783 -0.1609
// -0.5521 -1.2515  0.0901 -0.2223
//];
//upper left corner of of the Lapack dgesv solution
//for comparison:
//[
// -2.9416  0.1624  0.2631 -2.0923
// -0.0242  0.5965 -0.4656 -0.3765
//  0.6595  0.5525  0.5783 -0.1609
// -0.5521 -1.2515  0.0901 -0.2223
//];

```

4.5.8 magma_dgels3_gpu - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b) (((a)<(b))?(a):(b))
#define max(a,b) (((a)<(b))?(b):(a))
int main( int argc, char** argv)
{
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    double gpu_perf, cpu_perf;
    magma_int_t m = 8192, n = 8192, n2;
    magma_int_t nrhs = 4;
    double *a, *a2; // a, a2 -mxn matrices
    // a used in Lapack dgels, a2 -copy of a used in Magma dgesv
    double *b, *x, *r, *tau, *hwork, tmp[1]; // b, x, r - mxnrhs
    // matr.: b used in Lapack dgels, r copy used in Magma dgesv
    double *a1, *b1; // a1 - mxn matrix, b1 - mxnrhs matrix

```



```

// copies of a, b used in Magma dgels
magma_int_t lda, ldb, ldda, lddb, lworkgpu, lhwork;
magma_int_t info, min_mn, nb, l1, l2;
magma_int_t *piv,ione = 1;
magma_int_t ISEED[4] = {0,0,0,1};
ldda = ((m+31)/32)*32;
lddb = ldda;
n2 = m * n;
min_mn = min(m, n);
nb = magma_get_dgeqrf_nb(m,n);
lda = ldb = m;
lworkgpu = (m-n + nb)*(nrhs+2*nb);
// prepare unified memory
cudaMallocManaged(&tau,min_mn*sizeof(double)); //mem.for tau
cudaMallocManaged(&a,lda*n*sizeof(double)); // unif.mem.for a
cudaMallocManaged(&a2,lda*n*sizeof(double)); //un.mem.for a2
cudaMallocManaged(&b,ldb*nrhs*sizeof(double)); //u.mem.for b
cudaMallocManaged(&x,ldb*nrhs*sizeof(double)); //u.mem.for x
cudaMallocManaged(&r,ldb*nrhs*sizeof(double)); //u.mem.for r
cudaMallocManaged(&a1,ldda*n*sizeof(double)); //un.mem.for a1
cudaMallocManaged(&b1,lddb*nrhs*sizeof(double)); //mem.for b1
cudaMallocManaged(&piv,n*sizeof(magma_int_t)); //mem.for piv
// Get size for workspace
lhwork = -1;
lapackf77_dgeqrf(&m, &n, a, &m, tau, tmp, &lhwork, &info);
l1 = (magma_int_t)MAGMA_D_REAL( tmp[0] );
lhwork = -1;
lapackf77_dormqr( MagmaLeftStr, MagmaTransStr,
                 &m, &nrhs, &min_mn, a, &lda, tau,
                 x, &ldb, tmp, &lhwork, &info);
l2 = (magma_int_t)MAGMA_S_REAL( tmp[0] );
lhwork = max( max( l1, l2 ), lworkgpu );
// magma_dgels3_gpu needs this workspace
cudaMallocManaged(&hwork,lhwork*sizeof(double)); //mem.-hwork
// randomize the matrices a, b
lapackf77_dlarnv( &ione, ISEED, &n2, a ); // random a
n2 = m*nrhs; // size of b, x, r
lapackf77_dlarnv( &ione, ISEED, &n2, b ); // random b
// make copies of a and b: a-> a2, b -> r (they are overwrit.)
lapackf77_dlacpy(MagmaFullStr,&m,&nrhs,b,&ldb,r,&ldb);
lapackf77_dlacpy(MagmaFullStr,&m,&m,a,&lda,a2,&lda);
// copies of a,b for MAGMA
magma_dsetmatrix(m,n,a,lda,a1,ldda,queue); // copy a -> a1
magma_dsetmatrix( m,nrhs,b,ldb,b1,lddb,queue); // b -> b1
gpu_perf = magma_sync_wtime(NULL);
// solve the least squares problem min ||a1*x-b1||
// using the QR decomposition,
// the solution overwrites b1
// MAGMA version

magma_dgels3_gpu( MagmaNoTrans, m, n, nrhs, a1, ldda, b1,
                 lddb, hwork, lworkgpu, &info);

```

```

    gpu_perf = magma_sync_wtime(NULL)-gpu_perf;
    printf("MAGMA time: %7.3f sec.\n",gpu_perf);    // Magma time
    printf("upper left corner of of the Magma  solution:\n");
    magma_dprint( 4, 4, b1, m ); // small part of Magma solution
// LAPACK version of dgels
    cpu_perf=magma_wtime();
    lapackf77_dgels( MagmaNoTransStr, &m, &n, &nrhs,
                    a, &lda, b, &ldb, hwork, &lhwork, &info);
    cpu_perf=magma_wtime()-cpu_perf;
    printf("LAPACK time: %7.3f sec.\n",cpu_perf); // Lapack time
    printf("upper left corner of of the Lapack solution:\n");
    magma_dprint( 4, 4, b, m ); // small part of Lapack solution
// MAGMA dgesv for comparison
    magma_dgesv(n,nrhs,a2,n,piv,r,n,&info);
    printf("upper left corner of of the Lapack dgesv solution\n"
           "for comparison:\n");
    magma_dprint( 4, 4, r, m ); // small part of dgesv solution
                                // using LU decomposition
// Free unified memory
    magma_free(tau);           // free memory
    magma_free(a);             // free memory
    magma_free(a2);            // free memory
    magma_free(b);             // free memory
    magma_free(x);             // free memory
    magma_free(r);             // free memory
    magma_free(hwork);         // free memory
    magma_free(a1);            // free memory
    magma_free(b1);            // free memory
    magma_queue_destroy(queue); // destroy queue
    magma_finalize( );         // finalize Magma
    return EXIT_SUCCESS;
}
//MAGMA time:    3.047 sec.
//upper left corner of of the Magma  solution:
//[
//  -2.9416    0.1624    0.2631   -2.0923
//  -0.0242    0.5965   -0.4656   -0.3765
//   0.6595    0.5525    0.5783   -0.1609
//  -0.5521   -1.2515    0.0901   -0.2223
//];
//LAPACK time:   21.545 sec.
//upper left corner of of the Lapack solution:
//[
//  -2.9416    0.1624    0.2631   -2.0923
//  -0.0242    0.5965   -0.4656   -0.3765
//   0.6595    0.5525    0.5783   -0.1609
//  -0.5521   -1.2515    0.0901   -0.2223
//];
//upper left corner of of the Lapack dgesv solution
//for comparison:

```

```
//[
//  -2.9416   0.1624   0.2631  -2.0923
//  -0.0242   0.5965  -0.4656  -0.3765
//   0.6595   0.5525   0.5783  -0.1609
//  -0.5521  -1.2515   0.0901  -0.2223
//];
```

4.5.9 magma_sgeqrf - QR decomposition in single precision, CPU interface

This function computes in single precision the QR factorization:

$$A = Q R,$$

where A is an $m \times n$ general matrix defined on the host, R is upper triangular (upper trapezoidal in general case) and Q is orthogonal. On exit the upper triangle (trapezoid) of A contains R . The orthogonal matrix Q is represented as a product of elementary reflectors $H(1) \dots H(\min(m, n))$, where $H(k) = I - \tau_k v_k v_k^T$. The real scalars τ_k are stored in an array τ and the nonzero components of vectors v_k are stored on exit in parts of columns of A corresponding to the lower triangular (trapezoidal) part of A : $v_k(1 : k - 1) = 0$, $v_k(k) = 1$ and $v_k(k + 1 : m)$ is stored in $A(k + 1 : m, k)$. See [magma-X.Y.Z/src/sgeqrf.cpp](#) for more details.

```
#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b) (((a)<(b))? (a):(b))
#define max(a,b) (((a)<(b))? (b):(a))
int main( int argc, char** argv)
{
    magma_init(); // initialize Magma
    double gpu_time, cpu_time;
    magma_int_t m = 4096, n = 4096, n2=m*n;
    float *a, *r; // a, r - mxn matrices on the host
    float *tau; // scalars defining the elementary reflectors
    float *hwork, tmp[1]; // hwork - workspace; tmp -used in
    magma_int_t info, min_mn, nb; // workspace query
    magma_int_t ione = 1, lhwork; // lhwork - workspace size
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    min_mn = min(m, n);
    float mzone= MAGMA_S_NEG_ONE;
    float matnorm, work[1]; // used in difference computations
    magma_smallocc_cpu(&tau, min_mn); // host memory for tau
    magma_smallocc_pinned(&a, n2); // host memory for a
    magma_smallocc_pinned(&r, n2); // host memory for r
    // Get size for workspace
    nb = magma_get_sgeqrf_nb(m, n); //optim.block size for sgetrf
    lhwork = -1;
```

```

    lapackf77_sgeqrf(&m,&n,a,&m,tau,tmp,&lhwork,&info);
    lhwork = (magma_int_t)MAGMA_S_REAL( tmp[0] );
    lhwork = max(lhwork,max(n*nb,2*nb*nb));
    magma_smalloc_cpu(&hwork,lhwork);    // host memory for hwork
// Randomize the matrix
    lapackf77_slarnv( &ione, ISEED, &n2, a );    // randomize a
    lapackf77_slacpy(MagmaFullStr,&m,&n,a,&m,r,&m);    // a->r
// MAGMA
    gpu_time = magma_sync_wtime(NULL);
// compute a QR factorization of a real mxn matrix a
// a=Q*R, Q - orthogonal, R - upper triangular

    magma_sgeqrf( m, n, a, m, tau, hwork, lhwork, &info);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("MAGMA time: %7.3f sec.\n",gpu_time);    // print Magma
// LAPACK                                                                    time
    cpu_time=magma_wtime();
    lapackf77_sgeqrf(&m,&n,r,&m,tau,hwork,&lhwork,&info);
    cpu_time=magma_wtime()-cpu_time;
    printf("LAPACK time: %7.3f sec.\n",cpu_time);    //print Lapack
// difference                                                                    time
    matnorm = lapackf77_slange("f", &m, &n, a, &m, work);
    blasf77_saxpy(&n2, &mzone, a, &ione, r, &ione);
    printf("difference:  %e\n",    // ||a-r||_F/||a||_F
    lapackf77_slange("f", &m, &n, r, &m, work) / matnorm);
// Free memory
    free(tau);    // free host memory
    free(hwork);    // free host memory
    magma_free_pinned(a);    // free host memory
    magma_free_pinned(r);    // free host memory
    magma_finalize( );    // finalize Magma
    return EXIT_SUCCESS;
}
//MAGMA time:    0.310 sec.
//LAPACK time:    1.397 sec.
//difference:    1.860795e-06

```

4.5.10 magma_sgeqrf - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b) (((a)<(b))? (a):(b))
#define max(a,b) (((a)<(b))? (b):(a))
int main( int argc, char** argv)
{
    magma_init();    // initialize Magma

```

```

double gpu_time, cpu_time;
magma_int_t m = 3072, n = 4096, n2=m*n;
float *a, *r; // a,r - mxn matrices
float *tau; // scalars defining the elementary reflectors
float *hwork, tmp[1]; // hwork - workspace; tmp -used in
magma_int_t info, min_mn,nb; // workspace query
magma_int_t ione = 1,lhwork; // lhwork - workspace size
magma_int_t ISEED[4] = {0,0,0,1}; // seed
min_mn = min(m, n);
float mzone= MAGMA_S_NEG_ONE;
float matnorm, work[1]; // used in difference computations
cudaMallocManaged(&tau,min_mn*sizeof(float)); //mem.for tau
cudaMallocManaged(&a,n2*sizeof(float)); // unified.mem.for b
cudaMallocManaged(&r,n2*sizeof(float)); // unified.mem.for r
// Get size for workspace
nb = magma_get_sgeqrf_nb(m,n); //optim.block size for sgetrf
lhwork = -1;
lapackf77_sgeqrf(&m,&n,a,&m,tau,tmp,&lhwork,&info);
lhwork = (magma_int_t)MAGMA_D_REAL( tmp[0] );
lhwork = max(lhwork,max(n*nb,2*nb*nb));
cudaMallocManaged(&hwork,lhwork*sizeof(float)); //mem.f.hwork
// Randomize the matrix
lapackf77_slarnv( &ione, ISEED, &n2, a ); // randomize a
lapackf77_slacpy(MagmaFullStr,&m,&n,a,&m,r,&m); // a->r
// MAGMA
gpu_time = magma_sync_wtime(NULL);
// compute a QR factorization of a real mxn matrix a
// a=Q*R

magma_sgeqrf( m, n, a, m, tau, hwork, lhwork, &info);

gpu_time = magma_sync_wtime(NULL)-gpu_time;
printf("MAGMA time: %7.3f sec.\n",gpu_time); // print Magma
// LAPACK time
cpu_time=magma_wtime();
lapackf77_sgeqrf(&m,&n,r,&m,tau,hwork,&lhwork,&info);
cpu_time=magma_wtime()-cpu_time;
printf("LAPACK time: %7.3f sec.\n",cpu_time); //print Lapack
// difference time
matnorm = lapackf77_slange("f", &m, &n, a, &m, work);
blasf77_saxpy(&n2, &mzone, a, &ione, r, &ione);
printf("difference: %e\n", // ||a-r||_F/||a||_F
lapackf77_slange("f", &m, &n, r, &m, work) / matnorm);
// Free memory
magma_free(tau); // free memory
magma_free(hwork); // free memory
magma_free(a); // free memory
magma_free(r); // free memory
magma_finalize( ); // finalize Magma
return EXIT_SUCCESS;
}
//MAGMA time: 0.321 sec.

```

```
//LAPACK time:    0.775 sec.
//difference:    2.625919e-06
```

4.5.11 magma_dgeqrf - QR decomposition in double precision, CPU interface

This function computes in double precision the QR factorization:

$$A = Q R,$$

where A is an $m \times n$ general matrix defined on the host, R is upper triangular (upper trapezoidal in general case) and Q is orthogonal. On exit the upper triangle (trapezoid) of A contains R . The orthogonal matrix Q is represented as a product of elementary reflectors $H(1) \dots H(\min(m, n))$, where $H(k) = I - \tau_k v_k v_k^T$. The real scalars τ_k are stored in an array τ and the nonzero components of vectors v_k are stored on exit in parts of columns of A corresponding to the lower triangular (trapezoidal) part of A : $v_k(1 : k - 1) = 0$, $v_k(k) = 1$ and $v_k(k + 1 : m)$ is stored in $A(k + 1 : m, k)$. See [magma-X.Y.Z/src/dgeqrf.cpp](#) for more details.

```
#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b) (((a)<(b))?(a):(b))
#define max(a,b) (((a)<(b))?(b):(a))
int main( int argc, char** argv)
{
    magma_init(); // initialize Magma
    double gpu_time, cpu_time;
    magma_int_t m = 3072, n = 4096, n2=m*n; // a,r - mxn matrices
    double *a, *r; // a, r - mxn matrices on the host
    double *tau; // scalars defining the elementary reflectors
    double *hwork, tmp[1]; // hwork - workspace; tmp -used in
    magma_int_t i, info, min_mn, nb; // workspace query
    magma_int_t ione = 1, lhwork; // lhwork - workspace size
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    min_mn = min(m, n);
    double mzone= MAGMA_S_NEG_ONE;
    double matnorm, work[1]; // used in difference computations
    magma_dmalloc_cpu(&tau, min_mn); // host memory for tau
    magma_dmalloc_pinned(&a, n2); // host memory for a
    magma_dmalloc_pinned(&r, n2); // host memory for r
    // Get size for workspace
    nb = magma_get_dgeqrf_nb(m, n); //optim.block size for dgetrf
    lhwork = -1;
    lapackf77_dgeqrf(&m, &n, a, &m, tau, tmp, &lhwork, &info);
    lhwork = (magma_int_t)MAGMA_D_REAL( tmp[0] );
    lhwork = max(lhwork, max(n*nb, 2*nb*nb));
    magma_dmalloc_cpu(&hwork, lhwork); // host memory for hwork
```

```

// Randomize the matrix
lapackf77_dlarnv( &ione, ISEED, &n2, a );          // randomize a
lapackf77_dlacpy(MagmaFullStr,&m,&n,a,&m,r,&m);      // a->r
// MAGMA
gpu_time = magma_sync_wtime(NULL);
// compute a QR factorization of a real mxn matrix a
// a=Q*R

magma_dgeqrf( m, n, a, m, tau, hwork, lhwork, &info);

gpu_time = magma_sync_wtime(NULL)-gpu_time;
printf("MAGMA time: %7.3f sec.\n",gpu_time); // print Magma
// LAPACK                                     time
cpu_time=magma_wtime();
lapackf77_dgeqrf(&m,&n,r,&m,tau,hwork,&lhwork,&info);
cpu_time=magma_wtime()-cpu_time;
printf("LAPACK time: %7.3f sec.\n",cpu_time); //print Lapack
// difference                                     time
matnorm = lapackf77_dlange("f", &m, &n, a, &m, work);
blasf77_daxpy(&n2, &mzone, a, &ione, r, &ione);
printf("difference: %e\n",          // ||a-r||_F/||a||_F
lapackf77_dlange("f", &m, &n, r, &m, work) / matnorm);
// Free memory
free(tau);          // free host memory
free(hwork);        // free host memory
magma_free_pinned(a); // free host memory
magma_free_pinned(r); // free host memory
magma_finalize( );   // finalize Magma
return EXIT_SUCCESS;
}
//MAGMA time:    0.561 sec.
//LAPACK time:   1.521 sec.
//difference:    4.705079e-15

```

4.5.12 magma_dgeqrf - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b) (((a)<(b))?(a):(b))
#define max(a,b) (((a)<(b))?(b):(a))
int main( int argc, char** argv)
{
    magma_init();          // initialize Magma
    double gpu_time, cpu_time;
    magma_int_t m = 3072, n = 4096, n2=m*n; // a,r - mxn matrices
    double *a, *r;         // on the host
    double *tau;           // scalars defining the elementary reflectors

```

```

double *hwork, tmp[1];    // hwork - workspace; tmp -used in
magma_int_t  info, min_mn,nb;           // workspace query
magma_int_t  ione = 1,lhwork;           // lhwork - workspace size
magma_int_t  ISEED[4] = {0,0,0,1};      // seed
min_mn = min(m, n);
double mzone= MAGMA_S_NEG_ONE;
double matnorm, work[1]; // used in difference computations
cudaMallocManaged(&tau,min_mn*sizeof(double)); //u.mem.for tau
cudaMallocManaged(&a,n2*sizeof(double)); // unified mem.for a
cudaMallocManaged(&r,n2*sizeof(double)); // unified mem.for r
// Get size for workspace
nb = magma_get_dgeqrf_nb(m,n); // optim.block size for dgetrf
lhwork = -1;
lapackf77_dgeqrf(&m,&n,a,&m,tau,tmp,&lhwork,&info);
lhwork = (magma_int_t)MAGMA_D_REAL( tmp[0] );
lhwork = max(lhwork,max(n*nb,2*nb*nb));
cudaMallocManaged(&hwork,lhwork*sizeof(double)); //mem.f.hwork
// Randomize the matrix
lapackf77_dlarnv( &ione, ISEED, &n2, a ); // randomize a
lapackf77_dlacpy(MagmaFullStr,&m,&n,a,&m,r,&m); // a->r
// MAGMA
gpu_time = magma_sync_wtime(NULL);
// compute a QR factorization of a real mxn matrix a
// a=Q*R

magma_dgeqrf( m, n, a, m, tau, hwork, lhwork, &info);

gpu_time = magma_sync_wtime(NULL)-gpu_time;
printf("MAGMA time: %7.3f sec.\n",gpu_time); // print Magma
// LAPACK                                     time
cpu_time=magma_wtime();
lapackf77_dgeqrf(&m,&n,r,&m,tau,hwork,&lhwork,&info);
cpu_time=magma_wtime()-cpu_time;
printf("LAPACK time: %7.3f sec.\n",cpu_time); //print Lapack
// difference                                     time
matnorm = lapackf77_dlange("f", &m, &n, a, &m, work);
blasf77_daxpy(&n2, &mzone, a, &ione, r, &ione);
printf("difference:  %e\n",           // ||a-r||_F||a||_F
lapackf77_dlange("f", &m, &n, r, &m, work) / matnorm);
// Free memory
magma_free(tau); // free memory
magma_free(hwork); // free memory
magma_free(a); // free memory
magma_free(r); // free memory
magma_finalize( ); // finalize Magma
return EXIT_SUCCESS;
}
//MAGMA time:  0.614 sec.
//LAPACK time:  1.521 sec.
//difference:  4.705079e-15

```


4.5.13 magma_sgeqrf_gpu - QR decomposition in single precision, GPU interface

This function computes in single precision the QR factorization:

$$A = Q R,$$

where A is an $m \times n$ general matrix defined on the device, R is upper triangular (upper trapezoidal in general case) and Q is orthogonal. On exit the upper triangle (trapezoid) of A contains R . The orthogonal matrix Q is represented as a product of elementary reflectors $H(1) \dots H(\min(m, n))$, where $H(k) = I - \tau_k v_k v_k^T$. The real scalars τ_k are stored in an array τ and the nonzero components of vectors v_k are stored on exit in parts of columns of A corresponding to the lower triangular (trapezoidal) part of A : $v_k(1 : k - 1) = 0, v_k(k) = 1$ and $v_k(k + 1 : m)$ is stored in $A(k + 1 : m, k)$. See [magma-X.Y.Z/src/sgeqrf_gpu.cpp](#) for more details.

```
#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b) (((a)<(b))?(a):(b))
int main( int argc, char** argv)
{
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    double gpu_time, cpu_time;
    magma_int_t m = 8192, n = 8192, n2=m*n, ldda;
    float *a, *r; // a, r - mxn matrices on the host
    float *d_a; // d_a mxn matrix on the device
    float *tau; // scalars defining the elementary reflectors
    float *hwork, tmp[1]; // hwork - workspace; tmp -used
    magma_int_t info, min_mn; // in worksp.size comp.
    magma_int_t ione = 1, lhwork; // lhwork - workspace size
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    ldda = ((m+31)/32)*32; // ldda = m if 32 divides m
    min_mn = min(m, n);
    float mzone= MAGMA_S_NEG_ONE;
    float matnorm, work[1]; // used in difference computations
    magma_smallocc_cpu(&tau,min_mn); // host memory for tau
    magma_smallocc_pinned(&a,n2); // host memory for a
    magma_smallocc_pinned(&r,n2); // host memory for r
    magma_smallocc(&d_a,ldda*n); // device memory for d_a
    // Get size for workspace
    lhwork = -1;
    lapackf77_sgeqrf(&m,&n,a,&m,tau,tmp,&lhwork,&info);
    lhwork = (magma_int_t)MAGMA_S_REAL( tmp[0] );
    magma_smallocc_cpu(&hwork,lhwork); // host memory for hwork
```

```

// Lapack needs this array
// Randomize the matrix
lapackf77_slarnv( &ione, ISEED, &n2, a ); // randomize a
// MAGMA
magma_ssetmatrix( m, n, a,m,d_a,ldda,queue); // copy a -> d_a
gpu_time = magma_sync_wtime(NULL);
// compute a QR factorization of a real mxn matrix d_a
// d_a=Q*R, Q - orthogonal, R - upper triangular

magma_sgeqrf2_gpu( m, n, d_a, ldda, tau, &info);

gpu_time = magma_sync_wtime(NULL)-gpu_time;
printf("MAGMA time: %7.3f sec.\n",gpu_time); // Magma time
// LAPACK
cpu_time=magma_wtime();
lapackf77_sgeqrf(&m,&n,a,&m,tau,hwork,&lhwork,&info);
cpu_time=magma_wtime()-cpu_time;
printf("LAPACK time: %7.3f sec.\n",cpu_time); // Lapack time
// difference
magma_sgetmatrix( m, n,d_a,ldda,r,m,queue); // copy d_a -> r
matnorm = lapackf77_slange("f", &m, &n, a, &m, work);
blasf77_saxpy(&n2, &mzone, a, &ione, r, &ione);
printf("difference: %e\n", // ||a-r||_F/||a||_F
lapackf77_slange("f", &m, &n, r, &m, work) / matnorm);
// Free memory
free(tau); // free host memory
free(hwork); // free host memory
magma_free_pinned(a); // free host memory
magma_free_pinned(r); // free host memory
magma_free(d_a); // free device memory
magma_queue_destroy(queue); // destroy queue
magma_finalize( ); // finalize Magma
return EXIT_SUCCESS;
}
//MAGMA time: 0.335 sec.
//LAPACK time: 10.037 sec.
//difference: 2.670853e-06

```

4.5.14 magma_sgeqrf_gpu - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b) (((a)<(b))?(a):(b))
int main( int argc, char** argv)

```

```

{
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    double gpu_time, cpu_time;
    magma_int_t m = 8192, n = 8192, n2=m*n, ldda;
    float *a, *r; // a, r - mxn matrices
    float *a1; // a1 mxn matrix, used by Magma sgeqrf2_gpu
    float *tau; // scalars defining the elementary reflectors
    float *hwork, tmp[1]; // hwork - workspace; tmp -used in
    magma_int_t info, min_mn; // comp. workspace size
    magma_int_t ione = 1, lhwork; // lhwork - workspace size
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    ldda = ((m+31)/32)*32; // ldda = m if 32 divides m
    min_mn = min(m, n);
    float mzone= MAGMA_D_NEG_ONE;
    float matnorm, work[1]; // used in difference computations
    // prepare unified memory
    cudaMallocManaged(&tau,min_mn*sizeof(float)); //u.mem.for tau
    cudaMallocManaged(&a,n2*sizeof(float)); //unified mem.for a
    cudaMallocManaged(&r,n2*sizeof(float)); //unified mem.for r
    cudaMallocManaged(&a1,ldda*n*sizeof(float)); //uni.mem.for a1
    // Get size for workspace
    lhwork = -1;
    lapackf77_sgeqrf(&m,&n,a,&m,tau,tmp,&lhwork,&info);
    lhwork = (magma_int_t)MAGMA_D_REAL( tmp[0] );
    cudaMallocManaged(&hwork,lhwork*sizeof(float)); //mem.f.hwork
    // Randomize the matrix
    lapackf77_slarnv( &ione, ISEED, &n2, a ); // randomize a
    // MAGMA
    magma_ssetmatrix( m, n, a,m,a1,ldda,queue); // copy a -> a1
    gpu_time = magma_sync_wtime(NULL);
    // compute a QR factorization of a real mxn matrix a1
    // a1=Q*R, Q - orthogonal, R - upper triangular

    magma_sgeqrf2_gpu( m, n, a1, ldda, tau, &info);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("MAGMA time: %7.3f sec.\n",gpu_time); // Magma time
    // LAPACK
    cpu_time=magma_wtime();
    lapackf77_sgeqrf(&m,&n,a,&m,tau,hwork,&lhwork,&info);
    cpu_time=magma_wtime()-cpu_time;
    printf("LAPACK time: %7.3f sec.\n",cpu_time); // Lapack time
    // difference
    magma_sgetmatrix( m, n, a1,ldda,r,m,queue); // copy a1 -> r
    matnorm = lapackf77_slange("f", &m, &n, a, &m, work);
    blasf77_saxpy(&n2, &mzone, a, &ione, r, &ione);
    printf("difference: %e\n", // ||a-r||_F/||a||_F
    lapackf77_slange("f", &m, &n, r, &m, work) / matnorm);
    // Free memory

```

```

    magma_free(tau);           // free memory
    magma_free(hwork);        // free memory
    magma_free(a);            // free memory
    magma_free(r);            // free memory
    magma_free(a1);           // free memory
    magma_queue_destroy(queue); // destroy queue
    magma_finalize( );         // finalize Magma
    return EXIT_SUCCESS;
}
//MAGMA time:    0.341 sec.
//LAPACK time:   11.828 sec.
//difference:    2.670853e-06

```

4.5.15 magma_dgeqrf_gpu - QR decomposition in double precision, GPU interface

This function computes in double precision the QR factorization:

$$A = Q R,$$

where A is an $m \times n$ general matrix defined on the device, R is upper triangular (upper trapezoidal in general case) and Q is orthogonal. On exit the upper triangle (trapezoid) of A contains R . The orthogonal matrix Q is represented as a product of elementary reflectors $H(1) \dots H(\min(m, n))$, where $H(k) = I - \tau_k v_k v_k^T$. The real scalars τ_k are stored in an array τ and the nonzero components of vectors v_k are stored on exit in parts of columns of A corresponding to the lower triangular (trapezoidal) part of A : $v_k(1 : k - 1) = 0$, $v_k(k) = 1$ and $v_k(k + 1 : m)$ is stored in $A(k + 1 : m, k)$. See [magma-X.Y.Z/src/dgeqrf_gpu.cpp](#) for more details.

```

#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b) (((a)<(b))?(a):(b))
int main( int argc, char** argv)
{
    magma_init();           // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    double gpu_time, cpu_time;
    magma_int_t m = 8192, n = 8192, n2=m*n, ldda;
    double *a, *r;          // a, r - mxn matrices on the host
    double *d_a;            // d_a mxn matrix on the device
    double *tau;            // scalars defining the elementary reflectors
    double *hwork, tmp[1];  // hwork - workspace; tmp -used in
    magma_int_t info, min_mn; // workspace query
    magma_int_t ione = 1, lhwork; // lhwork - workspace size
    magma_int_t ISEED[4] = {0,0,0,1}; // seed

```

```

    ldda = ((m+31)/32)*32;           // ldda = m if 32 divides m
    min_mn = min(m, n);
    double mzone= MAGMA_D_NEG_ONE;
    double matnorm, work[1]; // used in difference computations
    magma_dmalloc_cpu(&tau,min_mn); // host memory for tau
    magma_dmalloc_pinned(&a,n2);    // host memory for a
    magma_dmalloc_pinned(&r,n2);    // host memory for r
    magma_dmalloc(&d_a,ldda*n);     // device memory for d_a
// Get size for workspace
    lhwork = -1;
    lapackf77_dgeqrf(&m,&n,a,&m,tau,tmp,&lhwork,&info);
    lhwork = (magma_int_t)MAGMA_D_REAL( tmp[0] );
    magma_dmalloc_cpu(&hwork,lhwork); // host memory for hwork
                                     // Lapack version needs this array
// Randomize the matrix
    lapackf77_dlarnv( &ione, ISEED, &n2, a ); // randomize a
// MAGMA
    magma_dsetmatrix( m, n, a,m,d_a,ldda,queue); // copy a -> d_a
    gpu_time = magma_sync_wtime(NULL);
// compute a QR factorization of a real mxn matrix d_a
// d_a=Q*R, Q - orthogonal, R - upper triangular

    magma_dgeqrf2_gpu( m, n, d_a, ldda, tau, &info);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("MAGMA time: %7.3f sec.\n",gpu_time); // Magma time
// LAPACK
    cpu_time=magma_wtime();
    lapackf77_dgeqrf(&m,&n,a,&m,tau,hwork,&lhwork,&info);
    cpu_time=magma_wtime()-cpu_time;
    printf("LAPACK time: %7.3f sec.\n",cpu_time); // Lapack time
// difference
    magma_dgetmatrix( m, n, d_a,ldda,r,m,queue); // copy d_a -> r
    matnorm = lapackf77_dlange("f", &m, &n, a, &m, work);
    blasf77_daxpy(&n2, &mzone, a, &ione, r, &ione);
    printf("difference: %e\n", // ||a-r||_F/||a||_F
    lapackf77_dlange("f", &m, &n, r, &m, work) / matnorm);
// Free memory
    free(tau); // free host memory
    free(hwork); // free host memory
    magma_free_pinned(a); // free host memory
    magma_free_pinned(r); // free host memory
    magma_free(d_a); // free device memory
    magma_queue_destroy(queue); // destroy queue
    magma_finalize( ); // finalize Magma
    return EXIT_SUCCESS;
}
//MAGMA time: 2.955 sec.
//LAPACK time: 16.932 sec.
//difference: 4.933266e-15

```

4.5.16 magma_dgeqrf_gpu - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b) (((a)<(b))?(a):(b))
int main( int argc, char** argv)
{
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    double gpu_time, cpu_time;
    magma_int_t m = 8192, n = 8192, n2=m*n, ldda;
    double *a, *r; // a, r - mxn matrices
    double *a1; // a1 mxn matrix, used by Magma dgeqrf2_gpu
    double *tau; // scalars defining the elementary reflectors
    double *hwork, tmp[1]; // hwork - workspace; tmp -used in
    magma_int_t info, min_mn; // comp. workspace size
    magma_int_t ione = 1, lhwork; // lhwork - workspace size
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    ldda = ((m+31)/32)*32; // ldda = m if 32 divides m
    min_mn = min(m, n);
    double mzone= MAGMA_D_NEG_ONE;
    double matnorm, work[1]; // used in difference computations
    // prepare unified memory
    cudaMallocManaged(&tau,min_mn*sizeof(double)); //mem.for tau
    cudaMallocManaged(&a,n2*sizeof(double)); //unified.mem.for a
    cudaMallocManaged(&r,n2*sizeof(double)); //unified.mem.for r
    cudaMallocManaged(&a1,ldda*n*sizeof(double)); //un.mem.for a1
    // Get size for workspace
    lhwork = -1;
    lapackf77_dgeqrf(&m,&n,a,&m,tau,tmp,&lhwork,&info);
    lhwork = (magma_int_t)MAGMA_D_REAL( tmp[0] );
    cudaMallocManaged(&hwork,lhwork*sizeof(double)); //mem-hwork
    // Randomize the matrix
    lapackf77_dlarnv( &ione, ISEED, &n2, a ); // randomize a
    // MAGMA
    magma_dsetmatrix( m, n, a,m,a1,ldda,queue); // copy a -> a1
    gpu_time = magma_sync_wtime(NULL);
    // compute a QR factorization of a real mxn matrix a1
    // a1=Q*R, Q - orthogonal, R - upper triangular

    magma_dgeqrf2_gpu( m, n, a1, ldda, tau, &info);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("MAGMA time: %7.3f sec.\n",gpu_time); // Magma time
    // LAPACK
    cpu_time=magma_wtime();
    lapackf77_dgeqrf(&m,&n,a,&m,tau,hwork,&lhwork,&info);

```

```

    cpu_time=magma_wtime()-cpu_time;
    printf("LAPACK time: %7.3f sec.\n",cpu_time); // Lapack time
// difference
    magma_dgetmatrix( m, n, a1,ldda,r,m,queue); // copy a1 -> r
    matnorm = lapackf77_dlange("f", &m, &n, a, &m, work);
    blasf77_daxpy(&n2, &mzone, a, &ione, r, &ione);
    printf("difference: %e\n", // ||a-r||_F/||a||_F
    lapackf77_dlange("f", &m, &n, r, &m, work) / matnorm);
// Free memory
    magma_free(tau); // free memory
    magma_free(hwork); // free memory
    magma_free(a); // free memory
    magma_free(r); // free memory
    magma_free(a1); // free memory
    magma_queue_destroy(queue); // destroy queue
    magma_finalize( ); // finalize Magma
    return EXIT_SUCCESS;
}
//MAGMA time: 3.014 sec.
//LAPACK time: 17.932 sec.
//difference: 4.933266e-15

```

4.5.17 magma_sgeqrf_mgpu - QR decomposition in single precision on multiple GPUs

This function computes in single precision the QR factorization:

$$A = Q R,$$

where A is an $m \times n$ general matrix, R is upper triangular (upper trapezoidal in general case) and Q is orthogonal. The matrix A and the factors are distributed on `num_gpus` devices. On exit the upper triangle (trapezoid) of A contains R . The orthogonal matrix Q is represented as a product of elementary reflectors $H(1) \dots H(\min(m, n))$, where $H(k) = I - \tau_k v_k v_k^T$. The real scalars τ_k are stored in an array τ and the nonzero components of vectors v_k are stored on exit in parts of columns of A corresponding to the lower triangular (trapezoidal) part of A : $v_k(1 : k-1) = 0$, $v_k(k) = 1$ and $v_k(k+1 : m)$ is stored in $A(k+1 : m, k)$. See [magma-X.Y.Z/src/sgeqrf_mgpu.cpp](#) for more details.

```

#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b) (((a)<(b))?(a):(b))
#define max(a,b) (((a)<(b))?(b):(a))
int main( int argc, char** argv)
{
    magma_init(); // initialize Magma

```

```

int num_gpus = 1; // for num_gpus GPUs
magma_setdevice(0);
magma_queue_t queues[num_gpus];
for( int dev = 0; dev < num_gpus; ++dev ) {
    magma_queue_create( dev, &queues[dev] );
}
double cpu_time, gpu_time;
magma_int_t m = 8192, n = m, n2=m*n;
float *a, *r; // a, r - mxn matrices on the host
magmaFloat_ptr d_la[num_gpus]; // pointers to memory
// on num_gpus devices
float *tau; // scalars defining the elementary reflectors
float *h_work, tmp[1]; // hwork - workspace; tmp -used in
// workspace query
magma_int_t n_local[4]; // sizes of local parts of matrix
magma_int_t i, info, min_mn= min(m, n);

magma_int_t ione = 1, lhwork; // lhwork - workspace size
float c_neg_one = MAGMA_D_NEG_ONE;
float matnorm, work[1]; // used in difference computations
magma_int_t ISEED[4] = {0,0,0,1}; // seed
magma_int_t ldda = ((m+31)/32)*32; //ldda = m if 32 divides m
magma_int_t nb = magma_get_sgeqrf_nb(m,n); //optim.block size
printf("Number of GPUs to be used = %d\n", (int) num_gpus);
// Allocate host memory for matrices
magma_smallocc_cpu(&tau,min_mn); // host memory for tau
magma_smallocc_pinned(&a,n2); // host memory for a
magma_smallocc_pinned(&r,n2); // host memory for r
for(i=0; i<num_gpus; i++){
    n_local[i] = ((n/nb)/num_gpus)*nb;
    if (i < (n/nb)%num_gpus)
        n_local[i] += nb;
    else if (i == (n/nb)%num_gpus)
        n_local[i] += n%nb;
    magma_setdevice(i);
    magma_smallocc(&d_la[i],ldda*n_local[i]); //device memory
// on num_gpus GPUs
    printf("device %2d n_local=%4d\n",(int)i,(int)n_local[i]);
}
magma_setdevice(0);
// Get size for host workspace
lhwork = -1;
lapackf77_sgeqrf(&m, &n, a, &m, tau, tmp, &lhwork, &info);
lhwork = (magma_int_t)MAGMA_S_REAL( tmp[0] );
magma_smallocc_cpu(&h_work,lhwork); // host memory for h_work
//Lapack sgeqrf needs this array
// Randomize the matrix a and copy a -> r
lapackf77_slarnv(&ione,ISEED,&n2,a);
lapackf77_slacpy(MagmaFullStr,&m,&n,a,&m,r,&m); // a->r
// LAPACK
cpu_time=magma_wtime();
// QR decomposition on the host

```



```

lapackf77_sgeqrf(&m,&n,a,&m,tau,h_work,&lhwork,&info);
cpu_time=magma_wtime()-cpu_time;
printf("Lapack dgeqrf time: %7.5f sec.\n",cpu_time);
// MAGMA // print Lapack time
magma_ssetmatrix_1D_col_bcyclic(num_gpus, m, n,nb, r,m,d_la,
                                ldda, queues); // distribute r -> num_gpus devices
gpu_time = magma_sync_wtime(NULL);
// QR decomposition on num_gpus devices

magma_sgeqrf2_mgpu( num_gpus, m, n, d_la, ldda, tau, &info);

gpu_time = magma_sync_wtime(NULL)-gpu_time;
printf("Magma dgeqrf_mgpu time: %7.5f sec.\n",gpu_time);
// print Magma time
magma_sgetmatrix_1D_col_bcyclic(num_gpus,m, n,nb, d_la,ldda,
                                r, m, queues); // gather num_gpus devices -> r
// difference
matnorm = lapackf77_slange("f", &m, &n, a, &m, work);
blasf77_saxpy(&n2, &c_neg_one, a, &i_one, r, &i_one);
printf("difference: %e\n",
        lapackf77_slange("f",&m,&n,r,&m,work)/matnorm);
free(tau); // free host memory
free(h_work); // free host memory
magma_free_pinned(a); // free host memory
magma_free_pinned(r); // free host memory
for(i=0; i<num_gpus; i++){
    magma_setdevice(i);
    magma_free(d_la[i] ); // free device memory
}
for( int dev = 0; dev < num_gpus; ++dev ) {
    magma_queue_destroy( queues[dev] );
}
magma_finalize( ); // finalize Magma
return EXIT_SUCCESS;
}
//Number of GPUs to be used = 1
//device 0 n_local=8192
//Lapack dgeqrf time: 10.11191 sec.
//Magma dgeqrf_mgpu time: 0.33583 sec.
//difference: 2.670853e-06

```

4.5.18 magma_dgeqrf_mgpu - QR decomposition in double precision on multiple GPUs

This function computes in double precision the QR factorization:

$$A = Q R,$$

where A is an $m \times n$ general matrix, R is upper triangular (upper trapezoidal in general case) and Q is orthogonal. The matrix A and the factors are

distributed on `num_gpus` devices. On exit the upper triangle (trapezoid) of A contains R . The orthogonal matrix Q is represented as a product of elementary reflectors $H(1) \dots H(\min(m, n))$, where $H(k) = I - \tau_k v_k v_k^T$. The real scalars τ_k are stored in an array τ and the nonzero components of vectors v_k are stored on exit in parts of columns of A corresponding to the lower triangular (trapezoidal) part of A : $v_k(1 : k-1) = 0$, $v_k(k) = 1$ and $v_k(k+1 : m)$ is stored in $A(k+1 : m, k)$. See [magma-X.Y.Z/src/dgeqrf_mgpu.cpp](#) for more details.

```
#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b) (((a)<(b))?(a):(b))
#define max(a,b) (((a)<(b))?(b):(a))
int main( int argc, char** argv)
{
    magma_init(); // initialize Magma
    int num_gpus = 1; // for num_gpus GPUs
    magma_setdevice(0);
    magma_queue_t queues[num_gpus];
    for( int dev = 0; dev < num_gpus; ++dev ) {
        magma_queue_create( dev, &queues[dev] );
    }
    double cpu_time, gpu_time;
    magma_int_t m = 8192, n = m, n2=m*n;
    double *a, *r; // a, r - mxn matrices on the host
    magmaDouble_ptr d_la[num_gpus]; // pointers to memory
    // on num_gpus devices
    double *tau; // scalars defining the elementary reflectors
    double *h_work, tmp[1]; // hwork - workspace; tmp -used in
    // workspace query
    magma_int_t n_local[4]; // sizes of local parts of matrix
    magma_int_t i, info, min_mn= min(m, n);

    magma_int_t ione = 1, lhwork; // lhwork - workspace size
    double c_neg_one = MAGMA_D_NEG_ONE;
    double matnorm, work[1]; // used in difference computations
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    magma_int_t ldda = ((m+31)/32)*32; //ldda = m if 32 divides m
    magma_int_t nb = magma_get_dgeqrf_nb(m,n); //optim.block size
    printf("Number of GPUs to be used = %d\n", (int) num_gpus);
    // Allocate host memory for matrices
    magma_dmalloc_cpu(&tau,min_mn); // host memory for tau
    magma_dmalloc_pinned(&a,n2); // host memory for a
    magma_dmalloc_pinned(&r,n2); // host memory for r
    for(i=0; i<num_gpus; i++){
        n_local[i] = ((n/nb)/num_gpus)*nb;
        if (i < (n/nb)%num_gpus)
            n_local[i] += nb;
        else if (i == (n/nb)%num_gpus)
```

```

        n_local[i] += n%nb;
        magma_setdevice(i);
        magma_dmalloc(&d_la[i], ldda*n_local[i]); //device memory
                                                // on num_gpus GPUs
        printf("device %2d n_local=%4d\n", (int)i, (int)n_local[i]);
    }
    magma_setdevice(0);
// Get size for host workspace
    lhwork = -1;
    lapackf77_dgeqrf(&m, &n, a, &m, tau, tmp, &lhwork, &info);
    lhwork = (magma_int_t)MAGMA_D_REAL( tmp[0] );
    magma_dmalloc_cpu(&h_work, lhwork); // host memory for h_work
                                        //Lapack sgeqrf needs this array
// Randomize the matrix a and copy a -> r
    lapackf77_dlarnv(&ione, ISEED, &n2, a);
    lapackf77_dlacpy(MagmaFullStr, &m, &n, a, &m, r, &m); // a->r
// LAPACK
    cpu_time=magma_wtime();
// QR decomposition on the host
    lapackf77_dgeqrf(&m, &n, a, &m, tau, h_work, &lhwork, &info);
    cpu_time=magma_wtime()-cpu_time;
    printf("Lapack dgeqrf time: %7.5f sec.\n", cpu_time);
// MAGMA // print Lapack time
    magma_dsetmatrix_1D_col_bcyclic(num_gpus, m, n, nb, r, m, d_la,
        ldda, queues); // distribute r -> num_gpus devices
    gpu_time = magma_sync_wtime(NULL);
// QR decomposition on num_gpus devices

    magma_dgeqrf2_mgpu( num_gpus, m, n, d_la, ldda, tau, &info);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("Magma dgeqrf_mgpu time: %7.5f sec.\n", gpu_time);
                                                // print Magma time
    magma_dgetmatrix_1D_col_bcyclic(num_gpus, m, n, nb, d_la, ldda,
        r, m, queues); // gather num_gpus devices -> r
// difference
    matnorm = lapackf77_dlange("f", &m, &n, a, &m, work);
    blasf77_daxpy(&n2, &c_neg_one, a, &ione, r, &ione);
    printf("difference: %e\n",
        lapackf77_dlange("f", &m, &n, r, &m, work)/matnorm);
    free(tau); // free host memory
    free(h_work); // free host memory
    magma_free_pinned(a); // free host memory
    magma_free_pinned(r); // free host memory
    for(i=0; i<num_gpus; i++){
        magma_setdevice(i);
        magma_free(d_la[i] ); // free device memory
    }
    for( int dev = 0; dev < num_gpus; ++dev ) {
        magma_queue_destroy( queues[dev] );
    }
    magma_finalize( ); // finalize Magma

```

```

    return EXIT_SUCCESS;
}
//Number of GPUs to be used = 1
//device 0 n_local=8192
//Lapack dgeqrf time: 16.91422 sec.
//Magma dgeqrf_mgpu time: 2.99641 sec.
//difference: 4.933266e-15

```

4.5.19 magma_sgelqf - LQ decomposition in single precision, CPU interface

This function computes in single precision the LQ factorization:

$$A = L Q,$$

where A is an $m \times n$ general matrix defined on the host, L is lower triangular (lower trapezoidal in general case) and Q is orthogonal. On exit the lower triangle (trapezoid) of A contains L . The orthogonal matrix Q is represented as a product of elementary reflectors $H(1) \dots H(\min(m, n))$, where $H(k) = I - \tau_k v_k v_k^T$. The real scalars τ_k are stored in an array τ and the nonzero components of vectors v_k are stored on exit in parts of rows of A corresponding to the upper triangular (trapezoidal) part of A : $v_k(1 : k-1) = 0, v_k(k) = 1$ and $v_k(k+1 : n)$ is stored in $A(k, k+1 : n)$. See [magma-X.Y.Z/src/sgelqf.cpp](#) for more details.

```

#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b) (((a)<(b))? (a):(b))
#define max(a,b) (((a)<(b))? (b):(a))
int main( int argc, char** argv){
    magma_init(); magma_init(); // initialize Magma
    double gpu_time, cpu_time;
    magma_int_t m = 4096, n = 4096, n2=m*n;
    float *a, *r; // a, r - mxn matrices on the host
    float *tau; // scalars defining the elementary reflectors
    float *h_work, tmp[1]; // h_work - workspace; tmp -used in
                                // workspace query

    magma_int_t info, min_mn, nb;
    magma_int_t lwork = 1, lwork; // lwork - workspace size
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    float matnorm, work[1]; // used in difference computations
    float mzone= MAGMA_S_NEG_ONE;
    min_mn = min(m, n);
    nb = magma_get_sgeqrf_nb(m,n); //optim.block size for sgeqrf
    magma_smallocc_cpu(&tau,min_mn); // host memory for tau
    magma_smallocc_pinned(&a,n2); // host memory for a

```

```

    magma_smalloc_pinned(&r,n2);                // host memory for r
// Get size for host workspace
    lwork = -1;
    lapackf77_sgelqf(&m, &n, a, &m, tau, tmp, &lwork, &info);
    lwork = (magma_int_t)MAGMA_S_REAL( tmp[0] );
    lwork = max( lwork, m*nb );
    magma_smalloc_pinned(&h_work,lwork); //host memory for h_work
// Randomize the matrix a and copy a -> r
    lapackf77_slarnv( &ione, ISEED, &n2, a );
    lapackf77_slacpy(MagmaFullStr,&m,&n,a,&m,r,&m );
// MAGMA
    gpu_time = magma_sync_wtime(NULL);
// LQ factorization for a real matrix r=L*Q using Magma
// L - lower triangular, Q - orthogonal

    magma_sgelqf(m,n,r,m,tau,h_work,lwork,&info);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("MAGMA time: %7.3f sec.\n",gpu_time); // print Magma
// LAPACK                                     time
    cpu_time=magma_wtime();
// LQ factorization for a real matrix a=L*Q on the host
    lapackf77_sgelqf(&m,&n,a,&m,tau,h_work,&lwork,&info);
    cpu_time=magma_wtime()-cpu_time;
    printf("LAPACK time: %7.3f sec.\n",cpu_time); // print Lapack
// difference                                     time
    matnorm = lapackf77_slange("f", &m, &n, a, &m, work);
    blasf77_saxpy(&n2, &mzone, a, &ione, r, &ione);
    printf("difference: %e\n", // ||a-r||_F/||a||_F
    lapackf77_slange("f", &m, &n, r, &m, work) / matnorm);
// Free emory
    free(tau); // free host memory
    magma_free_pinned(a); // free host memory
    magma_free_pinned(r); // free host memory
    magma_free_pinned(h_work); // free host memory
    magma_finalize( ); // finalize Magma
    return EXIT_SUCCESS;
}
//MAGMA time: 0.318 sec.
//LAPACK time: 2.394 sec.
//difference: 1.846170e-06

```

4.5.20 magma_sgelqf - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b) (((a)<(b))?(a):(b))

```

```

#define max(a,b)  (((a)<(b))?(b):(a))
int main( int argc, char** argv){
    magma_init();  magma_init();           // initialize Magma
    double  gpu_time, cpu_time;
    magma_int_t m = 4096, n = 4096, n2=m*n;
    float *a, *r;                          // a, r - mxn matrices
    float *tau;    // scalars defining the elementary reflectors
    float *h_work, tmp[1];    // h_work - workspace; tmp -used in
                                // workspace query

    magma_int_t  info, min_mn, nb;
    magma_int_t  ione = 1, lwork;           // lwork - workspace size
    magma_int_t  ISEED[4] = {0,0,0,1};      // seed
    float matnorm, work[1];    // used in difference computations
    float mzone= MAGMA_D_NEG_ONE;
    min_mn = min(m, n);
    nb = magma_get_sgeqrf_nb(m,n); //optim. block size for sgeqrf
    // prepare unified memory
    cudaMallocManaged(&tau,min_mn*sizeof(float)); //u.mem.for tau
    cudaMallocManaged(&a,n2*sizeof(float)); //unif. memory for a
    cudaMallocManaged(&r,n2*sizeof(float)); //unif. memory for r
    // Get size for workspace
    lwork = -1;
    lapackf77_sgelqf(&m, &n, a, &m, tau, tmp, &lwork, &info);
    lwork = (magma_int_t)MAGMA_D_REAL( tmp[0] );
    lwork = max( lwork, m*nb );
    cudaMallocManaged(&h_work,lwork*sizeof(float)); //mem.h_work
    // Randomize the matrix a and copy a -> r
    lapackf77_slarnv( &ione, ISEED, &n2, a );
    lapackf77_slacpy(MagmaFullStr,&m,&n,a,&m,r,&m );
    // MAGMA
    gpu_time = magma_sync_wtime(NULL);
    // LQ factorization for a real matrix r=L*Q using Magma
    // L - lower triangular, Q - orthogonal

    magma_sgelqf(m,n,r,m,tau,h_work,lwork,&info);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("MAGMA time: %7.3f sec.\n",gpu_time); // print Magma
    // LAPACK                                     time
    cpu_time=magma_wtime();
    // LQ factorization for a real matrix a=L*Q using Lapack
    lapackf77_sgelqf(&m,&n,a,&m,tau,h_work,&lwork,&info);
    cpu_time=magma_wtime()-cpu_time;
    printf("LAPACK time: %7.3f sec.\n",cpu_time); // print Lapack
    // difference                                     time
    matnorm = lapackf77_slange("f", &m, &n, a, &m, work);
    blasf77_saxpy(&n2, &mzone, a, &ione, r, &ione);
    printf("difference: %e\n",          // ||a-r||_F/||a||_F
    lapackf77_slange("f", &m, &n, r, &m, work) / matnorm);
    // Free emory
    magma_free(tau);           // free memory
    magma_free(a);            // free memory

```

```

    magma_free(r);                                // free memory
    magma_free(h_work);                            // free memory
    magma_finalize( );                             // finalize Magma
    magma_finalize( );                             // finalize Magma
    return EXIT_SUCCESS;
}
//MAGMA time:    0.342 sec.
//LAPACK time:   1.818 sec.
//difference: 1.846170e-06

```

4.5.21 magma_dgelqf - LQ decomposition in double precision, CPU interface

This function computes in double precision the LQ factorization:

$$A = L Q,$$

where A is an $m \times n$ general matrix defined on the host, L is lower triangular (lower trapezoidal in general case) and Q is orthogonal. On exit the lower triangle (trapezoid) of A contains L . The orthogonal matrix Q is represented as a product of elementary reflectors $H(1) \dots H(\min(m, n))$, where $H(k) = I - \tau_k v_k v_k^T$. The real scalars τ_k are stored in an array τ and the nonzero components of vectors v_k are stored on exit in parts of rows of A corresponding to the upper triangular (trapezoidal) part of A : $v_k(1 : k-1) = 0$, $v_k(k) = 1$ and $v_k(k+1 : n)$ is stored in $A(k, k+1 : n)$. See [magma-X.Y.Z/src/dgelqf.cpp](#) for more details.

```

#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b) (((a)<(b))?(a):(b))
#define max(a,b) (((a)<(b))?(b):(a))
int main( int argc, char** argv){
    magma_init(); magma_init();                    // initialize Magma
    double gpu_time, cpu_time;
    magma_int_t m = 4096, n = 4096, n2=m*n;
    double *a, *r;                                // a, r - mxn matrices on the host
    double *tau;                                    // scalars defining the elementary reflectors
    double *h_work, tmp[1];                        // h_work - workspace; tmp -used in
                                                    // workspace query

    magma_int_t info, min_mn, nb;
    magma_int_t ione = 1, lwork;                  // lwork - workspace size
    magma_int_t ISEED[4] = {0,0,0,1};             // seed
    double matnorm, work[1];                       // used in difference computations
    double mzone= MAGMA_D_NEG_ONE;
    min_mn = min(m, n);
    nb = magma_get_dgeqrf_nb(m,n); //optim. block size for dgeqrf
    magma_dmalloc_cpu(&tau,min_mn);               // host memory for tau
    magma_dmalloc_pinned(&a,n2);                  // host memory for a

```

```

    magma_dmalloc_pinned(&r,n2);                // host memory for r
// Get size for host workspace
    lwork = -1;
    lapackf77_dgelqf(&m, &n, a, &m, tau, tmp, &lwork, &info);
    lwork = (magma_int_t)MAGMA_D_REAL( tmp[0] );
    lwork = max( lwork, m*nb );
    magma_dmalloc_pinned(&h_work,lwork); //host memory for h_work
// Randomize the matrix a and copy a -> r
    lapackf77_dlarnv( &ione, ISEED, &n2, a );
    lapackf77_dlacpy( MagmaFullStr, &m, &n, a, &m, r, &m );
// MAGMA
    gpu_time = magma_sync_wtime(NULL);
// LQ factorization for a real matrix r=L*Q using Magma
// L - lower triangular, Q - orthogonal

    magma_dgelqf(m,n,r,m,tau,h_work,lwork,&info);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("MAGMA time: %7.3f sec.\n",gpu_time); // print Magma
// LAPACK                                     time
    cpu_time=magma_wtime();
// LQ factorization for a real matrix a=L*Q on the host
    lapackf77_dgelqf(&m,&n,a,&m,tau,h_work,&lwork,&info);
    cpu_time=magma_wtime()-cpu_time;
    printf("LAPACK time: %7.3f sec.\n",cpu_time); // print Lapack
// difference                                     time
    matnorm = lapackf77_dlange("f", &m, &n, a, &m, work);
    blasf77_daxpy(&n2, &mzone, a, &ione, r, &ione);
    printf("difference: %e\n", // ||a-r||_F/||a||_F
    lapackf77_dlange("f", &m, &n, r, &m, work) / matnorm);
// Free emory
    free(tau); // free host memory
    magma_free_pinned(a); // free host memory
    magma_free_pinned(r); // free host memory
    magma_free_pinned(h_work); // free host memory
    magma_finalize( ); // finalize Magma
    magma_finalize( ); // finalize Magma
    return EXIT_SUCCESS;
}
//MAGMA time: 0.715 sec.
//LAPACK time: 3.240 sec.
//difference: 3.434041e-15

```

4.5.22 magma_dgelqf - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"

```



```

#define min(a,b)  (((a)<(b))?(a):(b))
#define max(a,b)  (((a)<(b))?(b):(a))
int main( int argc, char** argv){
    magma_init(); magma_init(); // initialize Magma
    double gpu_time, cpu_time;
    magma_int_t m = 4096, n = 4096, n2=m*n;
    double *a, *r; // a, r - mxn matrices
    double *tau; // scalars defining the elementary reflectors
    double *h_work, tmp[1]; // h_work - workspace; tmp -used in
                                // workspace query

    magma_int_t info, min_mn, nb;
    magma_int_t ione = 1, lwork; // lwork - workspace size
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    double matnorm, work[1]; // used in difference computations
    double mzone= MAGMA_D_NEG_ONE;
    min_mn = min(m, n);
    nb = magma_get_dgeqrf_nb(m,n); //optim. block size for dgeqrf
    // prepare unified memory
    cudaMallocManaged(&tau,min_mn*sizeof(double)); //mem.for tau
    cudaMallocManaged(&a,n2*sizeof(double)); //unif. memory for a
    cudaMallocManaged(&r,n2*sizeof(double)); //unif. memory for r
    // Get size for workspace
    lwork = -1;
    lapackf77_dgelqf(&m, &n, a, &m, tau, tmp, &lwork, &info);
    lwork = (magma_int_t)MAGMA_D_REAL( tmp[0] );
    lwork = max( lwork, m*nb );
    cudaMallocManaged(&h_work,lwork*sizeof(double)); //mem.h_work
    // Randomize the matrix a and copy a -> r
    lapackf77_dlarnv( &ione, ISEED, &n2, a );
    lapackf77_dlacpy(MagmaFullStr,&m,&n,a,&m,r,&m );
    // MAGMA
    gpu_time = magma_sync_wtime(NULL);
    // LQ factorization for a real matrix r=L*Q using Magma
    // L - lower triangular, Q - orthogonal

    magma_dgelqf(m,n,r,m,tau,h_work,lwork,&info);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("MAGMA time: %7.3f sec.\n",gpu_time); // print Magma
    // LAPACK time
    cpu_time=magma_wtime();
    // LQ factorization for a real matrix a=L*Q using Lapack
    lapackf77_dgelqf(&m,&n,a,&m,tau,h_work,&lwork,&info);
    cpu_time=magma_wtime()-cpu_time;
    printf("LAPACK time: %7.3f sec.\n",cpu_time); // print Lapack
    // difference time
    matnorm = lapackf77_dlange("f", &m, &n, a, &m, work);
    blasf77_daxpy(&n2, &mzone, a, &ione, r, &ione);
    printf("difference: %e\n", // ||a-r||_F/||a||_F
    lapackf77_dlange("f", &m, &n, r, &m, work) / matnorm);
    // Free emory
    magma_free(tau); // free memory

```

```

    magma_free(a);                // free memory
    magma_free(r);                // free memory
    magma_free(h_work);          // free memory
    magma_finalize( );           // finalize Magma
    magma_finalize( );           // finalize Magma
    return EXIT_SUCCESS;
}
//MAGMA time:    0.728 sec.
//LAPACK time:   2.827 sec.
//difference: 3.434041e-15

```

4.5.23 magma_sgelqf_gpu - LQ decomposition in single precision, GPU interface

This function computes in single precision the LQ factorization:

$$A = L Q,$$

where A is an $m \times n$ general matrix defined on the device, L is lower triangular (lower trapezoidal in general case) and Q is orthogonal. On exit the lower triangle (trapezoid) of A contains L . The orthogonal matrix Q is represented as a product of elementary reflectors $H(1) \dots H(\min(m, n))$, where $H(k) = I - \tau_k v_k v_k^T$. The real scalars τ_k are stored in an array τ and the nonzero components of vectors v_k are stored on exit in parts of rows of A corresponding to the upper triangular (trapezoidal) part of A : $v_k(1 : k-1) = 0$, $v_k(k) = 1$ and $v_k(k+1 : n)$ is stored in $A(k, k+1 : n)$. See [magma-X.Y.Z/src/sgelqf_gpu.cpp](#) for more details.

```

#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b) (((a)<(b))?(a):(b))
#define max(a,b) (((a)<(b))?(b):(a))
int main( int argc, char** argv){
    magma_init(); magma_init();                // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    double gpu_time, cpu_time;
    magma_int_t m = 4096, n = 4096, n2=m*n;
    float *a, *r;                // a, r - mxn matrices on the host
    float *h_work, tmp[1];       // h_work - workspace; tmp -used in
                                // workspace query
    float *tau;                  // scalars defining the elementary reflectors
    float *d_a;                  // d_a - mxn matrix on the device
    magma_int_t info, min_mn, nb;
    magma_int_t lwork = 1, lwork; // lwork - workspace size
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    float matnorm, work[1];       // used in difference computations

```

```

float mzone= MAGMA_S_NEG_ONE;
min_mn = min(m, n);
nb = magma_get_sgeqrf_nb(m,n); //optim.block size for sgeqrf
magma_smallocc_cpu(&tau,min_mn); // host memory for tau
magma_smallocc_pinned(&a,n2); // host memory for a
magma_smallocc_pinned(&r,n2); // host memory for r
magma_smallocc(&d_a,m*n); // device memory for d_a
// Get size for host workspace
lwork = -1;
lapackf77_sgelqf(&m, &n, a, &m, tau, tmp, &lwork, &info);
lwork = (magma_int_t)MAGMA_S_REAL( tmp[0] );
lwork = max( lwork, m*nb );
magma_smallocc_pinned(&h_work,lwork); //host memory for h_work
// Randomize the matrix a and copy a -> r
lapackf77_slarnv( &ione, ISEED, &n2, a );
lapackf77_slacpy(MagmaFullStr,&m,&n,a,&m,r,&m );
// MAGMA
magma_ssetmatrix( m, n, r, m, d_a,m,queue); // copy r -> d_a
gpu_time = magma_sync_wtime(NULL);
// LQ factorization for a real matrix d_a=L*Q on the device
// L - lower triangular, Q - orthogonal

magma_sgelqf_gpu(m,n,d_a,m,tau,h_work,lwork,&info);

gpu_time = magma_sync_wtime(NULL)-gpu_time;
printf("MAGMA time: %7.3f sec.\n",gpu_time); // print Magma
// LAPACK time
cpu_time=magma_wtime();
// LQ factorization for a real matrix a=L*Q on the host
lapackf77_sgelqf(&m,&n,a,&m,tau,h_work,&lwork,&info);
cpu_time=magma_wtime()-cpu_time;
printf("LAPACK time: %7.3f sec.\n",cpu_time); // print Lapack
// difference time
magma_sgetmatrix( m, n, d_a, m, r, m,queue);
matnorm = lapackf77_slange("f", &m, &n, a, &m, work);
blasf77_saxpy(&n2, &mzone, a, &ione, r, &ione);
printf("difference: %e\n", // ||a-r||_F/||a||_F
lapackf77_slange("f", &m, &n, r, &m, work) / matnorm);
// Free emory
free(tau); // free host memory
magma_free_pinned(a); // free host memory
magma_free_pinned(r); // free host memory
magma_free_pinned(h_work); // free host memory
magma_free(d_a); // free device memory
magma_queue_destroy(queue); // destroy queue
magma_finalize( ); // finalize Magma
return EXIT_SUCCESS;
}
//MAGMA time: 0.067 sec.
//LAPACK time: 2.364 sec.
//difference: 1.846170e-06

```

4.5.24 magma_sgelqf_gpu - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b) (((a)<(b))?(a):(b))
#define max(a,b) (((a)<(b))?(b):(a))
int main( int argc, char** argv){
    magma_init(); magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    double gpu_time, cpu_time;
    magma_int_t m = 4096, n = 4096, n2=m*n;
    float *a, *r; // a, r - mxn matrices
    float *h_work, tmp[1]; // h_work - workspace; tmp -used in
                                // workspace query
    float *tau; // scalars defining the elementary reflectors
    float *a1; // a1 - mxn matrix used by Magma sgelqf_gpu
    magma_int_t info, min_mn, nb;
    magma_int_t ione = 1, lwork; // lwork - workspace size
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    float matnorm, work[1]; // used in difference computations
    float mzone= MAGMA_S_NEG_ONE;
    min_mn = min(m, n);
    nb = magma_get_sgeqrf_nb(m,n); //optim.block size for sgeqrf
    cudaMallocManaged(&tau,min_mn*sizeof(float)); //u.mem.for tau
    cudaMallocManaged(&a,n2*sizeof(float)); //unif. memory for a
    cudaMallocManaged(&r,n2*sizeof(float)); //unif. memory for r
    cudaMallocManaged(&a1,n2*sizeof(float)); //unif. mem. for a1
    // Get size for workspace
    lwork = -1;
    lapackf77_sgelqf(&m, &n, a, &m, tau, tmp, &lwork, &info);
    lwork = (magma_int_t)MAGMA_S_REAL( tmp[0] );
    lwork = max( lwork, m*nb );
    cudaMallocManaged(&h_work,lwork*sizeof(float)); //mem.h_work
    // Randomize the matrix a and copy a -> r
    lapackf77_slarnv( &ione, ISEED, &n2, a );
    lapackf77_slacpy(MagmaFullStr,&m,&n,a,&m,r,&m ); // a->r
    // MAGMA
    magma_ssetmatrix( m, n, r, m, a1,m,queue); // copy r->a1
    gpu_time = magma_sync_wtime(NULL);
    // LQ factorization for a real matrix a1=L*Q using Magma
    // L - lower triangular, Q - orthogonal

    magma_sgelqf_gpu(m,n,a1,m,tau,h_work,lwork,&info);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("MAGMA time: %7.3f sec.\n",gpu_time); // print Magma
    // LAPACK time

```

```

    cpu_time=magma_wtime();
// LQ factorization for a real matrix a=L*Q using Lapack
    lapackf77_sgelqf(&m,&n,a,&m,tau,h_work,&lwork,&info);
    cpu_time=magma_wtime()-cpu_time;
    printf("LAPACK time: %7.3f sec.\n",cpu_time);// print Lapack
// difference
    matnorm = lapackf77_slange("f", &m, &n, a, &m, work);// norm
    blasf77_saxpy(&n2, &mzone, a, &ione, a1, &ione); // a - a1
    printf("difference: %e\n", // ||a-a1||_F||a1||_F
    lapackf77_slange("f", &m, &n, a1, &m, work) / matnorm);
// Free emory
    magma_free(tau); // free memory
    magma_free(a); // free memory
    magma_free(r); // free memory
    magma_free(h_work); // free memory
    magma_free(a1); // free memory
    magma_queue_destroy(queue); // destroy queue
    magma_finalize( ); // finalize Magma
    return EXIT_SUCCESS;
}

//MAGMA time: 0.068 sec.
//LAPACK time: 1.810 sec.
//difference: 1.846170e-06

```

4.5.25 magma_dgelqf_gpu - LQ decomposition in double precision, GPU interface

This function computes in double precision the LQ factorization:

$$A = L Q,$$

where A is an $m \times n$ general matrix defined on the device, L is lower triangular (lower trapezoidal in general case) and Q is orthogonal. On exit the lower triangle (trapezoid) of A contains L . The orthogonal matrix Q is represented as a product of elementary reflectors $H(1) \dots H(\min(m, n))$, where $H(k) = I - \tau_k v_k v_k^T$. The real scalars τ_k are stored in an array τ and the nonzero components of vectors v_k are stored on exit in parts of rows of A corresponding to the upper triangular (trapezoidal) part of A : $v_k(1 : k-1) = 0$, $v_k(k) = 1$ and $v_k(k+1 : n)$ is stored in $A(k, k+1 : n)$. See [magma-X.Y.Z/src/dgelqf_gpu.cpp](#) for more details.

```

#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b) (((a)<(b))?(a):(b))
#define max(a,b) (((a)<(b))?(b):(a))
int main( int argc, char** argv){

```

```

magma_init(); magma_init(); // initialize Magma
magma_queue_t queue=NULL;
magma_int_t dev=0;
magma_queue_create(dev,&queue);
double gpu_time, cpu_time;
magma_int_t m = 4096, n = 4096, n2=m*n;
double *a, *r; // a, r - mxn matrices on the host
double *h_work, tmp[1]; // h_work - workspace; tmp -used in
// workspace query
double *tau; // scalars defining the elementary reflectors
double *d_a; // d_a - mxn matrix on the device
magma_int_t info, min_mn, nb;
magma_int_t ione = 1, lwork; // lwork - workspace size
magma_int_t ISEED[4] = {0,0,0,1}; // seed
double matnorm, work[1]; // used in difference computations
double mzone= MAGMA_D_NEG_ONE;
min_mn = min(m, n);
nb = magma_get_dgeqrf_nb(m,n); //optim.block size for dgeqrf
magma_dmalloc_cpu(&tau,min_mn); // host memory for tau
magma_dmalloc_pinned(&a,n2); // host memory for a
magma_dmalloc_pinned(&r,n2); // host memory for r
magma_dmalloc(&d_a,m*n); // device memory for d_a
// Get size for host workspace
lwork = -1;
lapackf77_dgelqf(&m, &n, a, &m, tau, tmp, &lwork, &info);
lwork = (magma_int_t)MAGMA_D_REAL( tmp[0] );
lwork = max( lwork, m*nb );
magma_dmalloc_pinned(&h_work,lwork); // host mem.for h_work
// Randomize the matrix a and copy a -> r
lapackf77_dlarv( &ione, ISEED, &n2, a );
lapackf77_dlacpy(MagmaFullStr,&m,&n,a,&m,r,&m );
// MAGMA
magma_dsetmatrix( m, n, r, m, d_a,m,queue); // copy r -> d_a
gpu_time = magma_sync_wtime(NULL);
// LQ factorization for a real matrix d_a=L*Q on the device
// L - lower triangular, Q - orthogonal

magma_dgelqf_gpu(m,n,d_a,m,tau,h_work,lwork,&info);

gpu_time = magma_sync_wtime(NULL)-gpu_time;
printf("MAGMA time: %7.3f sec.\n",gpu_time); // print Magma
// LAPACK time
cpu_time=magma_wtime();
// LQ factorization for a real matrix a=L*Q on the host
lapackf77_dgelqf(&m,&n,a,&m,tau,h_work,&lwork,&info);
cpu_time=magma_wtime()-cpu_time;
printf("LAPACK time: %7.3f sec.\n",cpu_time); // print Lapack
// difference time
magma_dgetmatrix( m, n, d_a, m, r, m, queue);
matnorm = lapackf77_dlange("f", &m, &n, a, &m, work);
blasf77_daxpy(&n2, &mzone, a, &ione, r, &ione);
printf("difference: %e\n", // ||a-r||_F/||a||_F

```

```

    lapackf77_dlange("f", &m, &n, r, &m, work) / matnorm);
// Free emory
free(tau); // free host memory
magma_free_pinned(a); // free host memory
magma_free_pinned(r); // free host memory
magma_free_pinned(h_work); // free host memory
magma_free(d_a); // free device memory
magma_queue_destroy(queue); // destroy queue
magma_finalize( ); // finalize Magma
return EXIT_SUCCESS;
}
//MAGMA time: 0.466 sec.
//LAPACK time: 3.197 sec.
//difference: 3.434041e-15

```

4.5.26 magma_dgelqf_gpu - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b) (((a)<(b))?(a):(b))
#define max(a,b) (((a)<(b))?(b):(a))
int main( int argc, char** argv){
    magma_init(); magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    double gpu_time, cpu_time;
    magma_int_t m = 4096, n = 4096, n2=m*n;
    double *a, *r; // a, r - mxn matrices
    double *h_work, tmp[1]; // h_work - workspace; tmp -used in
    // workspace query

    double *tau; // scalars defining the elementary reflectors
    double *a1; // a1 - mxn matrix used by Magma dgelqf_gpu
    magma_int_t info, min_mn, nb;
    magma_int_t ione = 1, lwork; // lwork - workspace size
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    double matnorm, work[1]; // used in difference computations
    double mzone= MAGMA_D_NEG_ONE;
    min_mn = min(m, n);
    nb = magma_get_dgeqrf_nb(m,n); //optim.block size for dgeqrf
    cudaMallocManaged(&tau,min_mn*sizeof(double)); //mem.for tau
    cudaMallocManaged(&a,n2*sizeof(double)); //unif. memory for a
    cudaMallocManaged(&r,n2*sizeof(double)); //unif. memory for r
    cudaMallocManaged(&a1,n2*sizeof(double)); //un. memory for a1
    // Get size for workspace
    lwork = -1;
    lapackf77_dgelqf(&m, &n, a, &m, tau, tmp, &lwork, &info);

```

```

    lwork = (magma_int_t)MAGMA_D_REAL( tmp[0] );
    lwork = max( lwork, m*nb );
    cudaMallocManaged(&h_work, lwork*sizeof(double)); //unif.mem
// Randomize the matrix a and copy a -> r // for h_work
    lapackf77_dlarnv( &ione, ISEED, &n2, a );
    lapackf77_dlacpy( MagmaFullStr, &m, &n, a, &m, r, &m ); // a->r
// MAGMA
    magma_dsetmatrix( m, n, r, m, a1, m, queue); // copy r->a1
    gpu_time = magma_sync_wtime(NULL);
// LQ factorization for a real matrix a1=L*Q using Magma
// L - lower triangular, Q - orthogonal

    magma_dgelqf_gpu(m,n,a1,m,tau,h_work,lwork,&info);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("MAGMA time: %7.3f sec.\n",gpu_time); // print Magma
// LAPACK time
    cpu_time=magma_wtime();
// LQ factorization for a real matrix a=L*Q using Lapack
    lapackf77_dgelqf(&m,&n,a,&m,tau,h_work,&lwork,&info);
    cpu_time=magma_wtime()-cpu_time;
    printf("LAPACK time: %7.3f sec.\n",cpu_time); // print Lapack
// difference time
    matnorm = lapackf77_dlange("f", &m, &n, a, &m, work); // norm
    blasf77_daxpy(&n2, &mzone, a, &ione, a1, &ione); // a - a1
    printf("difference: %e\n", // ||a-a1||_F/||a||_F
    lapackf77_dlange("f", &m, &n, a1, &m, work) / matnorm);
// Free emory
    magma_free(tau); // free memory
    magma_free(a); // free memory
    magma_free(r); // free memory
    magma_free(h_work); // free memory
    magma_free(a1); // free memory
    magma_queue_destroy(queue); // destroy queue
    magma_finalize( ); // finalize Magma
    return EXIT_SUCCESS;
}
//MAGMA time: 0.472 sec.
//LAPACK time: 2.832 sec.
//difference: 3.434041e-15

```

4.5.27 magma_sgeqp3 - QR decomposition with column pivoting in single precision, CPU interface

This function computes in single precision a QR factorization with column pivoting:

$$A P = Q R,$$

where A is an $m \times n$ matrix defined on the host, R is upper triangular (trapezoidal), Q is orthogonal and P is a permutation matrix. On exit the upper triangle (trapezoid) of A contains R . The orthogonal matrix Q

is represented as a product of elementary reflectors $H(1) \dots H(\min(m, n))$, where $H(k) = I - \tau_k v_k v_k^T$. The real scalars τ_k are stored in an array τ and the nonzero components of the vectors v_k are stored on exit in parts of columns of A corresponding to its upper triangular (trapezoidal) part: $v_k(1 : k - 1) = 0, v_k(k) = 1$ and $v_k(k + 1 : m)$ is stored in $A(k + 1 : m, k)$. The information on columns pivoting is contained in $jpvt$. On exit if $jpvt(j) = k$, then j -th column of AP was the k -th column of A . See [magma-X.Y.Z/src/sgeqp3.cpp](#) for more details.

```
#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b) (((a)<(b))?(a):(b))
#define max(a,b) (((a)<(b))?(b):(a))
int main( int argc, char** argv)
{
    magma_init(); // initialize Magma
    double gpu_time, cpu_time;
    magma_int_t m = 4096, n = 4096, n2=m*n;
    float *a, *r; // a, r - mxn matrices on the host
    float *h_work; // workspace
    float *tau; // scalars defining the elementary reflectors
    magma_int_t *jpvt; // pivoting information
    magma_int_t j, info, min_mn=min(m, n), nb;
    magma_int_t ione = 1, lwork; // lwork - workspace size
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    nb = magma_get_sgeqp3_nb(m,n); // optimal block size
    jpvt=(magma_int_t*)malloc(n*sizeof(magma_int_t)); //host mem.
    // for jpvt
    magma_smalloc_cpu(&tau,min_mn); // host memory for tau
    magma_smalloc_pinned(&a,n2); // host memory for a
    magma_smalloc_pinned(&r,n2); // host memory for r
    lwork = 2*n + ( n+1 )*nb;
    lwork = max(lwork, m * n + n);
    magma_smalloc_cpu(&h_work,lwork); // host memory for h_work
    // Randomize the matrix a and copy a -> r
    lapackf77_slarnv(&ione,ISEED,&n2,a);
    lapackf77_slacpy(MagmaFullStr,&m,&n,a,&m,r,&m); // a->r
    // LAPACK
    for (j = 0; j < n; j++)
        jpvt[j] = 0;
    cpu_time=magma_wtime();
    // QR decomposition with column pivoting, Lapack version
    lapackf77_sgeqp3(&m,&n,r,&m,jpvt,tau,h_work,&lwork,&info);
    cpu_time=magma_wtime()-cpu_time;
    printf("LAPACK time: %7.3f sec.\n",cpu_time); // Lapack time
    // MAGMA
    lapackf77_slacpy(MagmaFullStr,&m,&n,a,&m,r,&m); // a->r
    for (j = 0; j < n; j++)
        jpvt[j] = 0 ;
}
```

```

    gpu_time = magma_sync_wtime(NULL);
    // QR decomposition with column pivoting, Magma version

    magma_sgeqp3(m,n,r,m,jpvt,tau,h_work,lwork,&info);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("MAGMA time: %7.3f sec.\n",gpu_time);    // Magma time
    // Free memory
    free(jpvt);                                // free host memory
    free(tau);                                  // free host memory
    magma_free_pinned(a);                      // free host memory
    magma_free_pinned(r);                      // free host memory
    free( h_work );                            // free host memory
    magma_finalize( );                          // finalize Magma
    return EXIT_SUCCESS;
}
//LAPACK time:    6.402 sec.
//MAGMA time:    1.568 sec.

```

4.5.28 magma_sgeqp3 - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b) (((a)<(b))?(a):(b))
#define max(a,b) (((a)<(b))?(b):(a))
int main( int argc, char** argv)
{
    magma_init();                                // initialize Magma
    double gpu_time, cpu_time;
    magma_int_t m = 4096, n = 4096, n2=m*n;
    float *a, *r;                                // a, r - mxn matrices
    float *h_work;                                // workspace
    float *tau;    // scalars defining the elementary reflectors
    magma_int_t *jpvt;    // pivoting information
    magma_int_t j, info, min_mn=min(m, n), nb;
    magma_int_t ione = 1, lwork;    // lwork - workspace size
    magma_int_t ISEED[4] = {0,0,0,1};    // seed
    nb = magma_get_sgeqp3_nb(m,n);    // optimal block size
    cudaMallocManaged(&jpvt,n*sizeof(magma_int_t)); // m.for jpvt
    cudaMallocManaged(&tau,min_mn*sizeof(float)); // u.mem.for tau
    cudaMallocManaged(&a,n2*sizeof(float)); //unif. memory for a
    cudaMallocManaged(&r,n2*sizeof(float)); //unif. memory for r
    lwork = 2*n + ( n+1 )*nb;
    lwork = max(lwork, m * n + n);

    cudaMallocManaged(&h_work,lwork*sizeof(float)); //mem.h_work
    // Randomize the matrix a and copy a -> r

```

```

    lapackf77_slarnv(&ione, ISEED, &n2, a);
    lapackf77_slacpy(MagmaFullStr, &m, &n, a, &m, r, &m);           // a->r
// LAPACK
    for (j = 0; j < n; j++)
        jpvt[j] = 0;
    cpu_time=magma_wtime();
// QR decomposition with column pivoting, Lapack version
    lapackf77_sgeqp3(&m, &n, r, &m, jpvt, tau, h_work, &lwork, &info);
    cpu_time=magma_wtime()-cpu_time;
    printf("LAPACK time: %7.3f sec.\n", cpu_time); // Lapack time
// MAGMA
    lapackf77_slacpy(MagmaFullStr, &m, &n, a, &m, r, &m);           // a->r
    for (j = 0; j < n; j++)
        jpvt[j] = 0;
    gpu_time = magma_sync_wtime(NULL);
// QR decomposition with column pivoting, Magma version

    magma_sgeqp3(m, n, r, m, jpvt, tau, h_work, lwork, &info);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("MAGMA time: %7.3f sec.\n", gpu_time); // Magma time
// Free memory
    magma_free(jpvt);           // free memory
    magma_free(tau);           // free memory
    magma_free(a);             // free memory
    magma_free(r);             // free memory
    magma_free(h_work);       // free memory
    magma_finalize( );         // finalize Magma
    return EXIT_SUCCESS;
}

```

4.5.29 magma_dgeqp3 - QR decomposition with column pivoting in double precision, CPU interface

This function computes in double precision a QR factorization with column pivoting:

$$A P = Q R,$$

where A is an $m \times n$ matrix defined on the host, R is upper triangular (trapezoidal), Q is orthogonal and P is a permutation matrix. On exit the upper triangle (trapezoid) of A contains R . The orthogonal matrix Q is represented as a product of elementary reflectors $H(1) \dots H(\min(m, n))$, where $H(k) = I - \tau_k v_k v_k^T$. The real scalars τ_k are stored in an array τ and the nonzero components of the vectors v_k are stored on exit in parts of columns of A corresponding to its upper triangular (trapezoidal) part: $v_k(1 : k-1) = 0$, $v_k(k) = 1$ and $v_k(k+1 : m)$ is stored in $A(k+1 : m, k)$. The information on columns pivoting is contained in $jpvt$. On exit if $jpvt(j) = k$, then j -th column of AP was the k -th column of A . See [magma-X.Y.Z/src/dgeqp3.cpp](#) for more details.

```

#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b) (((a)<(b))?(a):(b))
#define max(a,b) (((a)<(b))?(b):(a))
int main( int argc, char** argv)
{
    magma_init(); // initialize Magma
    double gpu_time, cpu_time;
    magma_int_t m = 4096, n = 4096, n2=m*n;
    double *a, *r; // a, r - mxn matrices on the host
    double *h_work; // workspace
    double *tau; // scalars defining the elementary reflectors
    magma_int_t *jpvt; // pivoting information
    magma_int_t j, info, min_mn=min(m, n), nb;
    magma_int_t ione = 1, lwork; // lwork - workspace size
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    nb = magma_get_dgeqp3_nb(m,n); // optimal block size
    jpvt=(magma_int_t*)malloc(n*sizeof(magma_int_t)); //host mem.
    // for jpvt
    magma_dmalloc_cpu(&tau,min_mn); // host memory for tau
    magma_dmalloc_pinned(&a,n2); // host memory for a
    magma_dmalloc_pinned(&r,n2); // host memory for r
    lwork = 2*n + ( n+1 )*nb;
    lwork = max(lwork, m * n + n);
    magma_dmalloc_cpu(&h_work,lwork); // host memory for h_work
    // Randomize the matrix a and copy a -> r
    lapackf77_dlarnv(&ione,ISEED,&n2,a);
    lapackf77_dlacpy(MagmaFullStr,&m,&n,a,&m,r,&m); // a->r
    // LAPACK
    for (j = 0; j < n; j++)
        jpvt[j] = 0;
    cpu_time=magma_wtime();
    // QR decomposition with column pivoting, Lapack version
    lapackf77_dgeqp3(&m,&n,r,&m,jpvt,tau,h_work,&lwork,&info);
    cpu_time=magma_wtime()-cpu_time;
    printf("LAPACK time: %7.3f sec.\n",cpu_time); // Lapack time
    // MAGMA
    lapackf77_dlacpy(MagmaFullStr,&m,&n,a,&m,r,&m); // a->r
    for (j = 0; j < n; j++)
        jpvt[j] = 0 ;
    gpu_time = magma_sync_wtime(NULL);
    // QR decomposition with column pivoting, Magma version

    magma_dgeqp3(m,n,r,m,jpvt,tau,h_work,lwork,&info);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("MAGMA time: %7.3f sec.\n",gpu_time); // Magma time
    // Free memory
    free(jpvt); // free host memory
    free(tau); // free host memory
}

```

```

    magma_free_pinned(a);                // free host memory
    magma_free_pinned(r);                // free host memory
    free( h_work );                      // free host memory
    magma_finalize( );                   // finalize Magma
    return EXIT_SUCCESS;
}
//LAPACK time:  14.135 sec.
//MAGMA time:   2.173 sec.

```

4.5.30 magma_dgeqp3 - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b)  (((a)<(b))?(a):(b))
#define max(a,b)  (((a)<(b))?(b):(a))
int main( int argc, char** argv)
{
    magma_init();                        // initialize Magma
    double gpu_time, cpu_time;
    magma_int_t m = 4096, n = 4096, n2=m*n;
    double *a, *r;                      // a, r - mxn matrices
    double *h_work;                      // workspace
    double *tau;                         // scalars defining the elementary reflectors
    magma_int_t *jpvt;                   // pivoting information
    magma_int_t j, info, min_mn=min(m, n), nb;
    magma_int_t ione = 1, lwork;         // lwork - workspace size
    magma_int_t ISEED[4] = {0,0,0,1};   // seed
    nb = magma_get_dgeqp3_nb(m,n);       // optimal block size
    cudaMallocManaged(&jpvt,n*sizeof(magma_int_t)); // m.for jpvt
    cudaMallocManaged(&tau,min_mn*sizeof(double)); //mem.for tau
    cudaMallocManaged(&a,n2*sizeof(double)); //unif.memory for a
    cudaMallocManaged(&r,n2*sizeof(double)); //unif.memory for r
    lwork = 2*n + ( n+1 )*nb;
    lwork = max(lwork, m * n + n);

    cudaMallocManaged(&h_work,lwork*sizeof(double)); //mem.h_work
    // Randomize the matrix a and copy a -> r
    lapackf77_dlarnv(&ione,ISEED,&n2,a);
    lapackf77_dlacpy(MagmaFullStr,&m,&n,a,&m,r,&m); // a->r
    // LAPACK
    for (j = 0; j < n; j++)
        jpvt[j] = 0;
    cpu_time=magma_wtime();
    // QR decomposition with column pivoting, Lapack version
    lapackf77_dgeqp3(&m,&n,r,&m,jpvt,tau,h_work,&lwork,&info);
    cpu_time=magma_wtime()-cpu_time;
    printf("LAPACK time: %7.3f sec.\n",cpu_time); // Lapack time
}

```



```

float *VL, *VR;           // VL,VR - nxn matrices of left and
                           // right eigenvectors
float *wr1, *wr2;         // wr1,wr2 - real parts of eigenvalues
float *wi1, *wi2, error;   // wi1,wi2 - imaginary parts of
                           // eigenvalues

magma_int_t ione = 1, i, j, info, nb;
float mione = -1.0f, *h_work; // h_work - workspace
magma_int_t incr = 1, inci = 1, lwork; // lwork -worksp. size
nb = magma_get_sgehrd_nb(n); // optimal block size for sgehrd
float work[1]; // used in difference computations
lwork = n*(2+nb);
lwork = max(lwork, n*(5+2*n));
magma_smallocc_cpu(&wr1, n); // host memory for real
magma_smallocc_cpu(&wr2, n); // and imaginary parts
magma_smallocc_cpu(&wi1, n); // of eigenvalues
magma_smallocc_cpu(&wi2, n);
magma_smallocc_cpu(&a, n2); // host memory for a
magma_smallocc_cpu(&r, n2); // host memory for r
magma_smallocc_cpu(&VL, n2); // host memory for left
magma_smallocc_cpu(&VR, n2); // and right eigenvectors
magma_smallocc_cpu(&h_work, lwork); // host memory for h_work
// define a, r // [1 0 0 0 0 ...]
for(i=0; i<n; i++){ // [0 2 0 0 0 ...]
    a[i*n+i]=(float)(i+1); // a = [0 0 3 0 0 ...]
    r[i*n+i]=(float)(i+1); // [0 0 0 4 0 ...]
} // [0 0 0 0 5 ...]
printf("upper left corner of a:\n"); // .....
magma_sprint(5,5,a,n); // print a
// compute the eigenvalues and the right eigenvectors
// for a general, real nxn matrix,
// Magma version, left eigenvectors not computed,
// right eigenvectors are computed

magma_sgeev(MagmaNoVec, MagmaVec, n, r, n, wr1, wi1, VL, n, VR, n,
            h_work, lwork, &info);

printf("first 5 eigenvalues of a:\n");
for(j=0; j<5; j++){
    printf("%f+%f*I\n", wr1[j], wi1[j]); // print eigenvalues
    printf("left upper corner of right eigenvectors matrix:\n");
    magma_sprint(5,5,VR,n); // print right eigenvectors
// Lapack version // in columns
lapackf77_sgeev("N", "V", &n, a, &n, wr2, wi2, VL, &n, VR, &n,
                h_work, &lwork, &info);
// difference in real parts of eigenvalues
blasf77_saxpy( &n, &mione, wr1, &incr, wr2, &incr);
error = lapackf77_slange( "M", &n, &ione, wr2, &n, work );
printf("difference in real parts: %e\n", error);
// difference in imaginary parts of eigenvalues
blasf77_saxpy( &n, &mione, wi1, &inci, wi2, &inci);
error = lapackf77_slange( "M", &n, &ione, wi2, &n, work );
printf("difference in imaginary parts: %e\n", error);

```

```

    free(wr1);                // free host memory
    free(wr2);                // free host memory
    free(wi1);                // free host memory
    free(wi2);                // free host memory
    free(a);                  // free host memory
    free(r);                  // free host memory
    free(VL);                 // free host memory
    free(VR);                 // free host memory
    free(h_work);             // free host memory
    magma_finalize();         // finalize Magma
    return EXIT_SUCCESS;
}
//upper left corner of a:
//[
//  1.0000  0.      0.      0.      0.
//  0.      2.0000  0.      0.      0.
//  0.      0.      3.0000  0.      0.
//  0.      0.      0.      4.0000  0.
//  0.      0.      0.      0.      5.0000
//];
//first 5 eigenvalues of a:
//1.000000+0.000000*I
//2.000000+0.000000*I
//3.000000+0.000000*I
//4.000000+0.000000*I
//5.000000+0.000000*I
//left upper corner of right eigenvectors matrix:
//[
//  1.0000  0.      0.      0.      0.
//  0.      1.0000  0.      0.      0.
//  0.      0.      1.0000  0.      0.
//  0.      0.      0.      1.0000  0.
//  0.      0.      0.      0.      1.0000
//];
//difference in real parts: 0.000000e+00
//difference in imaginary parts: 0.000000e+00

```

4.6.2 magma_sgeev - unified memory version, small matrix

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define max(a,b) (((a)<(b))? (b):(a))
int main( int argc, char** argv) {
    magma_init();                // initialize Magma
    magma_int_t n=1024, n2=n*n;
    float *a, *r;                // a, r - nxn matrices
    float *VL, *VR;              // VL,VR - nxn matrices of left and

```



```

// right eigenvectors
float *wr1, *wr2;      // wr1,wr2 - real parts of eigenvalues
float *wi1, *wi2;      // wi1,wi2 - imaginary parts of
magma_int_t ione = 1, i, j, info, nb;      // eigenvalues
float mione = -1.0, error, *h_work;      // h_work - workspace
magma_int_t incr = 1, inci = 1, lwork; // lwork -worksp. size
nb = magma_get_sgehrd_nb(n); // optimal block size for sgehrd
float work[1];          // used in difference computations
lwork = n*(2+nb);
lwork = max(lwork, n*(5+2*n));
cudaMallocManaged(&wr1, n*sizeof(float)); //unified memory for
cudaMallocManaged(&wr2, n*sizeof(float)); //real parts of eig
cudaMallocManaged(&wi1, n*sizeof(float)); //unified memory for
cudaMallocManaged(&wi2, n*sizeof(float)); //imag.parts of eig
cudaMallocManaged(&a, n2*sizeof(float)); //unif. memory for a
cudaMallocManaged(&r, n2*sizeof(float)); //unif. memory for r
cudaMallocManaged(&VL, n2*sizeof(float)); //u.mem.for left and
cudaMallocManaged(&VR, n2*sizeof(float)); //right eigenvect.
cudaMallocManaged(&h_work, lwork*sizeof(float));
// define a, r
for(i=0; i<n; i++){
    a[i*n+i]=(float)(i+1);
    r[i*n+i]=(float)(i+1);
}
printf("upper left corner of a:\n");
magma_sprint(5,5,a,n);
// compute the eigenvalues and the right eigenvectors
// for a general, real nxn matrix,
// Magma version, left eigenvectors not computed,
// right eigenvectors are computed

magma_sgeev(MagmaNoVec, MagmaVec, n, r, n, wr1, wi1, VL, n, VR, n,
            h_work, lwork, &info);

printf("first 5 eigenvalues of a:\n");
for(j=0; j<5; j++){
    printf("%f+%f*I\n", wr1[j], wi1[j]); // print eigenvalues
printf("left upper corner of right eigenvectors matrix:\n");
magma_sprint(5,5,VR,n); // print right eigenvectors
// Lapack version // in columns
lapackf77_sgeev("N", "V", &n, a, &n, wr2, wi2, VL, &n, VR, &n,
                h_work, &lwork, &info);
// difference in real parts of eigenvalues
blasf77_saxpy( &n, &mione, wr1, &incr, wr2, &incr);
error = lapackf77_slange( "M", &n, &ione, wr2, &n, work );
printf("difference in real parts: %e\n", error);
// difference in imaginary parts of eigenvalues
blasf77_saxpy( &n, &mione, wi1, &inci, wi2, &inci);
error = lapackf77_slange( "M", &n, &ione, wi2, &n, work );
printf("difference in imaginary parts: %e\n", error);
magma_free(wr1); // free memory
magma_free(wr2); // free memory

```

```

    magma_free(wi1);           // free memory
    magma_free(wi2);           // free memory
    magma_free(a);             // free memory
    magma_free(r);             // free memory
    magma_free(VL);            // free memory
    magma_free(VR);            // free memory
    magma_free(h_work);        // free memory
    magma_finalize();          // finalize Magma
    return EXIT_SUCCESS;
}
//upper left corner of a:
//[
//  1.0000  0.      0.      0.      0.
//  0.      2.0000  0.      0.      0.
//  0.      0.      3.0000  0.      0.
//  0.      0.      0.      4.0000  0.
//  0.      0.      0.      0.      5.0000
//];
//first 5 eigenvalues of a:
//1.000000+0.000000*I
//2.000000+0.000000*I
//3.000000+0.000000*I
//4.000000+0.000000*I
//5.000000+0.000000*I
//left upper corner of right eigenvectors matrix:
//[
//  1.0000  0.      0.      0.      0.
//  0.      1.0000  0.      0.      0.
//  0.      0.      1.0000  0.      0.
//  0.      0.      0.      1.0000  0.
//  0.      0.      0.      0.      1.0000
//];
//difference in real parts: 0.000000e+00
//difference in imaginary parts: 0.000000e+00

```

4.6.3 magma_dgeev - compute the eigenvalues and optionally eigenvectors of a general real matrix in double precision, CPU interface, small matrix

This function computes in double precision the eigenvalues and, optionally, the left and/or right eigenvectors for an $n \times n$ matrix A defined on the host. The first parameter can take the values `MagmaNoVec` or `MagmaVec` and answers the question whether the left eigenvectors are to be computed. Similarly the second parameter answers the question whether the right eigenvectors are to be computed. The computed eigenvectors are normalized to have Euclidean norm equal to one. If computed, the left eigenvectors are stored in columns of an array `VL` and the right eigenvectors in columns of `VR`. The real and imaginary parts of eigenvalues are stored in arrays `wr`, `wi` respectively. See [magma-X.Y.Z/src/dgeev.cpp](#) for more details.

```
#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define max(a,b) (((a)<(b))?(b):(a))
int main( int argc, char** argv) {
    magma_init(); // initialize Magma
    magma_int_t n=1024, n2=n*n;
    double *a, *r; // a, r - nxn matrices on the host
    double *VL, *VR; // VL,VR - nxn matrices of left and
                    // right eigenvectors
    double *wr1, *wr2; // wr1,wr2 - real parts of eigenvalues
    double *wi1, *wi2; // wi1,wi2 - imaginary parts of
    magma_int_t ione = 1, i, j, info, nb; // eigenvalues
    double mione = -1.0, error, *h_work; // h_work - workspace
    magma_int_t incr = 1, inci = 1, lwork; // lwork -worksp. size
    nb = magma_get_dgehrd_nb(n); // optimal block size for dgehrd
    double work[1]; // used in difference computations
    lwork = n*(2+nb);
    lwork = max(lwork, n*(5+2*n));
    magma_dmalloc_cpu(&wr1, n); // host memory for real
    magma_dmalloc_cpu(&wr2, n); // and imaginary parts
    magma_dmalloc_cpu(&wi1, n); // of eigenvalues
    magma_dmalloc_cpu(&wi2, n);
    magma_dmalloc_cpu(&a, n2); // host memory for a
    magma_dmalloc_cpu(&r, n2); // host memory for r
    magma_dmalloc_cpu(&VL, n2); // host memory for left
    magma_dmalloc_cpu(&VR, n2); // and right eigenvectors
    magma_dmalloc_cpu(&h_work, lwork); // host memory for h_work
    // define a, r // [1 0 0 0 0 ...]
    for(i=0; i<n; i++){ // [0 2 0 0 0 ...]
        a[i*n+i]=(double)(i+1); // a = [0 0 3 0 0 ...]
        r[i*n+i]=(double)(i+1); // [0 0 0 4 0 ...]
    } // [0 0 0 0 5 ...]
    printf("upper left corner of a:\n"); // .....
    magma_dprint(5,5,a,n); // print a
    // compute the eigenvalues and the right eigenvectors
    // for a general, real nxn matrix,
    // Magma version, left eigenvectors not computed,
    // right eigenvectors are computed

    magma_dgeev(MagmaNoVec, MagmaVec, n, r, n, wr1, wi1, VL, n, VR, n,
                h_work, lwork, &info);

    printf("first 5 eigenvalues of a:\n");
    for(j=0; j<5; j++)
        printf("%f+%f*I\n", wr1[j], wi1[j]); // print eigenvalues
    printf("left upper corner of right eigenvectors matrix:\n");
    magma_dprint(5,5,VR,n); // print right eigenvectors
    // Lapack version // in columns
    lapackf77_dgeev("N", "V", &n, a, &n, wr2, wi2, VL, &n, VR, &n,
                    h_work, &lwork, &info);
```

```

// difference in real parts of eigenvalues
blasf77_daxpy( &n, &mione, wr1, &incr, wr2, &incr);
error = lapackf77_dlange( "M", &n, &ione, wr2, &n, work );
printf("difference in real parts: %e\n",error);
// difference in imaginary parts of eigenvalues
blasf77_daxpy( &n, &mione, wi1, &inci, wi2, &inci);
error = lapackf77_dlange( "M", &n, &ione, wi2, &n, work );
printf("difference in imaginary parts: %e\n",error);
free(wr1); // free host memory
free(wr2); // free host memory
free(wi1); // free host memory
free(wi2); // free host memory
free(a); // free host memory
free(r); // free host memory
free(VL); // free host memory
free(VR); // free host memory
free(h_work); // free host memory
magma_finalize(); // finalize Magma
return EXIT_SUCCESS;
}
//upper left corner of a:
//[
//  1.0000  0.      0.      0.      0.
//  0.      2.0000  0.      0.      0.
//  0.      0.      3.0000  0.      0.
//  0.      0.      0.      4.0000  0.
//  0.      0.      0.      0.      5.0000
//];
//first 5 eigenvalues of a:
//1.000000+0.000000*I
//2.000000+0.000000*I
//3.000000+0.000000*I
//4.000000+0.000000*I
//5.000000+0.000000*I
//left upper corner of right eigenvectors matrix:
//[
//  1.0000  0.      0.      0.      0.
//  0.      1.0000  0.      0.      0.
//  0.      0.      1.0000  0.      0.
//  0.      0.      0.      1.0000  0.
//  0.      0.      0.      0.      1.0000
//];
//difference in real parts: 0.000000e+00
//difference in imaginary parts: 0.000000e+00

```

4.6.4 magma_dgeev - unified memory version, small matrix

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>

```

```

#include "magma_v2.h"
#include "magma_lapack.h"
#define max(a,b) (((a)<(b))? (b):(a))
int main( int argc, char** argv) {
    magma_init(); // initialize Magma
    magma_int_t n=1024, n2=n*n;
    double *a, *r; // a, r - nxn matrices
    double *VL, *VR; // VL,VR - nxn matrices of left and
                    // right eigenvectors
    double *wr1, *wr2; // wr1,wr2 - real parts of eigenvalues
    double *wi1, *wi2; // wi1,wi2 - imaginary parts of
    magma_int_t ione = 1, i, j, info, nb; // eigenvalues
    double mione = -1.0, error, *h_work; // h_work - workspace
    magma_int_t incr = 1, inci = 1, lwork; // lwork -worksp. size
    nb = magma_get_dgehrd_nb(n); // optimal block size for dgehrd
    double work[1]; // used in difference computations
    lwork = n*(2+nb);
    lwork = max(lwork, n*(5+2*n));
    cudaMallocManaged(&wr1, n*sizeof(double)); //unified mem. for
    cudaMallocManaged(&wr2, n*sizeof(double)); //real parts of eig
    cudaMallocManaged(&wi1, n*sizeof(double)); //unified mem. for
    cudaMallocManaged(&wi2, n*sizeof(double)); //imag. parts of eig
    cudaMallocManaged(&a, n2*sizeof(double)); //unif. memory for a
    cudaMallocManaged(&r, n2*sizeof(double)); //unif. memory for r
    cudaMallocManaged(&VL, n2*sizeof(double)); //mem. for left and
    cudaMallocManaged(&VR, n2*sizeof(double)); //right eigenvect.
    cudaMallocManaged(&h_work, lwork*sizeof(double));
    // define a, r // [1 0 0 0 0 ...]
    for(i=0; i<n; i++){ // [0 2 0 0 0 ...]
        a[i*n+i]=(double)(i+1); // a = [0 0 3 0 0 ...]
        r[i*n+i]=(double)(i+1); // [0 0 0 4 0 ...]
    } // [0 0 0 0 5 ...]
    printf("upper left corner of a:\n"); // [.....]
    magma_dprint(5,5,a,n); // print a
    // compute the eigenvalues and the right eigenvectors
    // for a general, real nxn matrix,
    // Magma version, left eigenvectors not computed,
    // right eigenvectors are computed

    magma_dgeev(MagmaNoVec, MagmaVec, n, r, n, wr1, wi1, VL, n, VR, n,
                h_work, lwork, &info);

    printf("first 5 eigenvalues of a:\n");
    for(j=0; j<5; j++){
        printf("%f+%f*I\n", wr1[j], wi1[j]); // print eigenvalues
    }
    printf("left upper corner of right eigenvectors matrix:\n");
    magma_dprint(5,5,VR,n); // print right eigenvectors
    // Lapack version // in columns
    lapackf77_dgeev("N", "V", &n, a, &n, wr2, wi2, VL, &n, VR, &n,
                    h_work, &lwork, &info);
    // difference in real parts of eigenvalues
    blasf77_daxpy( &n, &mione, wr1, &incr, wr2, &incr);

```

```

    error = lapackf77_dlange( "M", &n, &ione, wr2, &n, work );
    printf("difference in real parts: %e\n",error);
// difference in imaginary parts of eigenvalues
    blasf77_daxpy( &n, &mione, wi1, &inci, wi2, &inci);
    error = lapackf77_dlange( "M", &n, &ione, wi2, &n, work );
    printf("difference in imaginary parts: %e\n",error);
    magma_free(wr1); // free memory
    magma_free(wr2); // free memory
    magma_free(wi1); // free memory
    magma_free(wi2); // free memory
    magma_free(a); // free memory
    magma_free(r); // free memory
    magma_free(VL); // free memory
    magma_free(VR); // free memory
    magma_free(h_work); // free memory
    magma_finalize(); // finalize Magma
    return EXIT_SUCCESS;
}
//upper left corner of a:
//[
//  1.0000  0.      0.      0.      0.
//  0.      2.0000  0.      0.      0.
//  0.      0.      3.0000  0.      0.
//  0.      0.      0.      4.0000  0.
//  0.      0.      0.      0.      5.0000
//];
//first 5 eigenvalues of a:
//1.000000+0.000000*I
//2.000000+0.000000*I
//3.000000+0.000000*I
//4.000000+0.000000*I
//5.000000+0.000000*I
//left upper corner of right eigenvectors matrix:
//[
//  1.0000  0.      0.      0.      0.
//  0.      1.0000  0.      0.      0.
//  0.      0.      1.0000  0.      0.
//  0.      0.      0.      1.0000  0.
//  0.      0.      0.      0.      1.0000
//];
//difference in real parts: 0.000000e+00
//difference in imaginary parts: 0.000000e+00

```

4.6.5 magma_sgeev - compute the eigenvalues and optionally eigenvectors of a general real matrix in single precision, CPU interface, big matrix

This function computes in single precision the eigenvalues and, optionally, the left and/or right eigenvectors for an $n \times n$ matrix A defined on the host. The first parameter can take the values `MagmaNoVec` or `MagmaVec` and answers the question whether the left eigenvectors are to be computed. Sim-

ilarly the second parameter answers the question whether the right eigenvectors are to be computed. The computed eigenvectors are normalized to have Euclidean norm equal to one. If computed, the left eigenvectors are stored in columns of an array VL and the right eigenvectors in columns of VR. The real and imaginary parts of eigenvalues are stored in arrays wr, wi respectively. See [magma-X.Y.Z/src/sgeev.cpp](#) for more details.

```
#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define max(a,b) (((a)<(b))?(b):(a))
int main( int argc, char** argv) {
    magma_init(); // initialize Magma
    magma_int_t n=8192, n2=n*n;
    float *a, *r; // a, r - nxn matrices on the host
    float *VL, *VR; // VL,VR - nxn matrices of left and
                    // right eigenvectors
    float *wr1, *wr2; // wr1,wr2 - real parts of eigenvalues
    float *wi1, *wi2; // wi1,wi2 - imaginary parts of eigenvals
    float gpu_time, cpu_time, *h_work; // h_work - workspace
    magma_int_t ione=1, info, nb, lwork; // lwork - worksp. size
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    nb = magma_get_sgehrd_nb(n); // optimal block size for sgehrd
    lwork = n*(2+nb);
    lwork = max(lwork, n*(5+2*n));
    magma_smallocc_cpu(&wr1, n); // host memory for real
    magma_smallocc_cpu(&wr2, n); // and imaginary parts
    magma_smallocc_cpu(&wi1, n); // of eigenvalues
    magma_smallocc_cpu(&wi2, n);
    magma_smallocc_cpu(&a, n2); // host memory for a
    magma_smallocc_pinned(&r, n2); // host memory for r
    magma_smallocc_pinned(&VL, n2); // host memory for left
    magma_smallocc_pinned(&VR, n2); // and right eigenvectors
    magma_smallocc_pinned(&h_work, lwork); // host memory for h_work
    // Randomize the matrix a and copy a -> r
    lapackf77_slarnv(&ione, ISEED, &n2, a);
    lapackf77_slacpy(MagmaFullStr, &n, &n, a, &n, r, &n);
    // MAGMA
    gpu_time = magma_sync_wtime(NULL);
    // compute the eigenvalues of a general, real nxn matrix,
    // Magma version, left and right eigenvectors not computed

    magma_sgeev(MagmaNoVec, MagmaNoVec, n, r, n, wr1, wi1, VL, n, VR, n,
                h_work, lwork, &info);

    gpu_time = magma_sync_wtime(NULL) - gpu_time;
    printf("sgeev gpu time: %7.5f sec.\n", gpu_time); // Magma
    // LAPACK // time
    cpu_time = magma_wtime();
    // compute the eigenvalues of a general, real nxn matrix a,
```

```

// Lapack version
lapackf77_sgeev("N", "N", &n, a, &n,
               wr2, wi2, VL, &n, VR, &n, h_work, &lwork, &info);
cpu_time=magma_wtime()-cpu_time;
printf("sgeev cpu time: %7.5f sec.\n",cpu_time); // Lapack
free(wr1); // time
free(wr2); // free host memory
free(wi1); // free host memory
free(wi2); // free host memory
free(a); // free host memory
magma_free_pinned(r); // free host memory
magma_free_pinned(VL); // free host memory
magma_free_pinned(VR); // free host memory
magma_free_pinned(h_work); // free host memory
magma_finalize( ); // finalize Magma
return EXIT_SUCCESS;
}
//sgeev gpu time: 46.21376 sec.
//sgeev cpu time: 157.79790 sec.

```

4.6.6 magma_sgeev - unified memory version, big matrix

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define max(a,b) (((a)<(b))?(b):(a))
int main( int argc, char** argv) {
    magma_init(); // initialize Magma
    magma_int_t n=8192, n2=n*n;
    float *a, *r; // a, r - nxn matrices
    float *VL, *VR; // VL,VR - nxn matrices of left and
                    // right eigenvectors
    float *wr1, *wr2; // wr1,wr2 - real parts of eigenvalues
    float *wi1, *wi2; // wi1,wi2 - imaginary parts of eigenvalues
    double gpu_time, cpu_time;
    float *h_work; // h_work - workspace
    magma_int_t ione=1,info,nb,lwork; // lwork - worksp. size
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    nb = magma_get_sgehrd_nb(n); // optimal block size for sgehrd
    lwork = n*(2+nb);
    lwork = max(lwork, n*(5+2*n));
    cudaMallocManaged(&wr1,n*sizeof(float)); //unified memory for
    cudaMallocManaged(&wr2,n*sizeof(float)); //real parts of eig
    cudaMallocManaged(&wi1,n*sizeof(float)); //unified memory for
    cudaMallocManaged(&wi2,n*sizeof(float)); //imag.parts of eig
    cudaMallocManaged(&a,n2*sizeof(float)); //unif. memory for a
    cudaMallocManaged(&r,n2*sizeof(float)); //unif. memory for r
    cudaMallocManaged(&VL,n2*sizeof(float)); //u.mem.for left and

```



```

    cudaMallocManaged(&VR,n2*sizeof(float)); //right eigenvect.
    cudaMallocManaged(&h_work,lwork*sizeof(float)); //m.f. h_work
// Randomize the matrix a and copy a -> r
    lapackf77_slarnv(&ione,ISEED,&n2,a);
    lapackf77_slacpy(MagmaFullStr,&n,&n,a,&n,r,&n);
// MAGMA
    gpu_time = magma_sync_wtime(NULL);
// compute the eigenvalues of a general, real nxn matrix,
// Magma version, left and right eigenvectors not computed

    magma_sgeev(MagmaNoVec,MagmaNoVec,n,r,n,wr1,wi1,VL,n,VR,n,
                h_work,lwork,&info);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("sgeev gpu time: %7.5f sec.\n",gpu_time); // Magma
// LAPACK // time
    cpu_time=magma_wtime();
// compute the eigenvalues of a general, real nxn matrix a,
// Lapack version
    lapackf77_sgeev("N", "N", &n, a, &n,
                    wr2, wi2, VL, &n, VR, &n, h_work, &lwork, &info);
    cpu_time=magma_wtime()-cpu_time;
    printf("sgeev cpu time: %7.5f sec.\n",cpu_time); // Lapack
    magma_free(wr1); // time
    magma_free(wr2); // free memory
    magma_free(wi1); // free memory
    magma_free(wi2); // free memory
    magma_free(a); // free memory
    magma_free(r); // free memory
    magma_free(VL); // free memory
    magma_free(VR); // free memory
    magma_free(h_work); // free memory
    magma_finalize( ); // finalize Magma
    return EXIT_SUCCESS;
}
//sgeev gpu time: 40.60117 sec.
//sgeev cpu time: 108.51452 sec.

```

4.6.7 magma_dgeev - compute the eigenvalues and optionally eigenvectors of a general real matrix in double precision, CPU interface, big matrix

This function computes in double precision the eigenvalues and, optionally, the left and/or right eigenvectors for an $n \times n$ matrix A defined on the host. The first parameter can take the values `MagmaNoVec` or `MagmaVec` and answers the question whether the left eigenvectors are to be computed. Similarly the second parameter answers the question whether the right eigenvectors are to be computed. The computed eigenvectors are normalized to have Euclidean norm equal to one. If computed, the left eigenvectors are stored in columns of an array `VL` and the right eigenvectors in columns of

VR. The real and imaginary parts of eigenvalues are stored in arrays `wr`, `wi` respectively. See [magma-X.Y.Z/src/dgeev.cpp](#) for more details.

```
#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define max(a,b) (((a)<(b))?(b):(a))
int main( int argc, char** argv) {
    magma_init(); // initialize Magma
    magma_int_t n=8192, n2=n*n;
    double *a, *r; // a, r - nxn matrices on the host
    double *VL, *VR; // VL,VR - nxn matrices of left and
    // right eigenvectors
    double *wr1, *wr2; // wr1,wr2 - real parts of eigenvalues
    double *wi1, *wi2; // wi1,wi2 - imaginary parts of eigenvalues
    double gpu_time, cpu_time, *h_work; // h_work - workspace
    magma_int_t ione=1, info, nb, lwork; // lwork - workspace size
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    nb = magma_get_dgehrd_nb(n); // optimal block size for dgehrd
    lwork = n*(2+nb);
    lwork = max(lwork, n*(5+2*n));
    magma_dmalloc_cpu(&wr1, n); // host memory for real
    magma_dmalloc_cpu(&wr2, n); // and imaginary parts
    magma_dmalloc_cpu(&wi1, n); // of eigenvalues
    magma_dmalloc_cpu(&wi2, n);
    magma_dmalloc_cpu(&a, n2); // host memory for a
    magma_dmalloc_pinned(&r, n2); // host memory for r
    magma_dmalloc_pinned(&VL, n2); // host memory for left
    magma_dmalloc_pinned(&VR, n2); // and right eigenvectors
    magma_dmalloc_pinned(&h_work, lwork); // host memory for h_work
    // Randomize the matrix a and copy a -> r
    lapackf77_dlarv(&ione, ISEED, &n2, a);
    lapackf77_dlacpy(MagmaFullStr, &n, &n, a, &n, r, &n);
    // MAGMA
    gpu_time = magma_sync_wtime(NULL);
    // compute the eigenvalues of a general, real nxn matrix,
    // Magma version, left and right eigenvectors not computed

    magma_dgeev(MagmaNoVec, MagmaNoVec, n, r, n, wr1, wi1, VL, n, VR, n,
                h_work, lwork, &info);

    gpu_time = magma_sync_wtime(NULL) - gpu_time;
    printf("dgeev gpu time: %7.5f sec.\n", gpu_time); // Magma
    // LAPACK // time
    cpu_time = magma_wtime();
    // compute the eigenvalues of a general, real nxn matrix a,
    // Lapack version
    lapackf77_dgeev("N", "N", &n, a, &n,
                    wr2, wi2, VL, &n, VR, &n, h_work, &lwork, &info);
    cpu_time = magma_wtime() - cpu_time;
    printf("dgeev cpu time: %7.5f sec.\n", cpu_time); // Lapack
```

```

    free(wr1);
    free(wr2);
    free(wi1);
    free(wi2);
    free(a);
    magma_free_pinned(r);
    magma_free_pinned(VL);
    magma_free_pinned(VR);
    magma_free_pinned(h_work);
    magma_finalize( );
    return EXIT_SUCCESS;
}
//dgeev gpu time: 95.42350 sec.
//dgeev cpu time: 211.23290 sec.

```

4.6.8 magma_dgeev - unified memory version, big matrix

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define max(a,b) (((a)<(b))? (b):(a))
int main( int argc, char** argv) {
    magma_init();
    magma_int_t n=8192, n2=n*n;
    double *a, *r;
    double *VL, *VR;
    double *wr1, *wr2;
    double *wi1, *wi2;
    double gpu_time, cpu_time, *h_work;
    magma_int_t ione=1, info, nb, lwork;
    magma_int_t ISEED[4] = {0,0,0,1};
    nb = magma_get_dgehrd_nb(n);
    lwork = n*(2+nb);
    lwork = max(lwork, n*(5+2*n));
    cudaMallocManaged(&wr1, n*sizeof(double));
    cudaMallocManaged(&wr2, n*sizeof(double));
    cudaMallocManaged(&wi1, n*sizeof(double));
    cudaMallocManaged(&wi2, n*sizeof(double));
    cudaMallocManaged(&a, n2*sizeof(double));
    cudaMallocManaged(&r, n2*sizeof(double));
    cudaMallocManaged(&VL, n2*sizeof(double));
    cudaMallocManaged(&VR, n2*sizeof(double));
    cudaMallocManaged(&h_work, lwork*sizeof(double));
    // Randomize the matrix a and copy a -> r
    lapackf77_dlarnv(&ione, ISEED, &n2, a);
    lapackf77_dlacpy(MagmaFullStr, &n, &n, a, &n, r, &n);
    // MAGMA

```

```

    gpu_time = magma_sync_wtime(NULL);
    // compute the eigenvalues of a general, real nxn matrix,
    // Magma version, left and right eigenvectors not computed

    magma_dgeev(MagmaNoVec,MagmaNoVec,n,r,n,wr1,wi1,VL,n,VR,n,
                h_work,lwork,&info);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("dgeev gpu time: %7.5f sec.\n",gpu_time);    // Magma
    // LAPACK                                           // time
    cpu_time=magma_wtime();
    // compute the eigenvalues of a general, real nxn matrix a,
    // Lapack version
    lapackf77_dgeev("N", "N", &n, a, &n,
                    wr2, wi2, VL, &n, VR, &n, h_work, &lwork, &info);
    cpu_time=magma_wtime()-cpu_time;
    printf("dgeev cpu time: %7.5f sec.\n",cpu_time);    // Lapack
    magma_free(wr1);                                   // time
    magma_free(wr2);                                   // free memory
    magma_free(wi1);                                   // free memory
    magma_free(wi2);                                   // free memory
    magma_free(a);                                     // free memory
    magma_free(r);                                     // free memory
    magma_free(VL);                                    // free memory
    magma_free(VR);                                    // free memory
    magma_free(h_work);                                // free memory
    magma_finalize( );                                  // finalize Magma
    return EXIT_SUCCESS;
}
//dgeev gpu time: 62.50911 sec.
//dgeev cpu time: 185.40615 sec.

```

4.6.9 magma_sgehrd - reduce a general matrix to the upper Hessenberg form in single precision, CPU interface

This function using the single precision reduces a general real $n \times n$ matrix A defined on the host to upper Hessenberg form:

$$Q^T A Q = H,$$

where Q is an orthogonal matrix and H has zero elements below the first subdiagonal. The orthogonal matrix Q is represented as a product of elementary reflectors $H(i\ell o) \dots H(ihi)$, where $H(k) = I - \tau_k v_k v_k^T$. The real scalars τ_k are stored in an array τ and the information on vectors v_k is stored on exit in the lower triangular part of A below the first subdiagonal: $v_k(1 : k) = 0, v_k(k + 1) = 1$ and $v_k(ihi + 1 : n) = 0$; $v_k(k + 2 : ihi)$ is stored in $A(k + 2 : ihi, k)$. The function uses also an array dT defined on the device, storing blocks of triangular matrices used in the reduction process. See [magma-X.Y.Z/src/sgehrd.cpp](#) for more details.

```

#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv)
{
    magma_init(); // initialize Magma
    double gpu_time, cpu_time;
    magma_int_t n=2048, n2=n*n;
    float *a, *r, *r1; // a,r,r1 - nxn matrices on the host
    float *tau; // scalars defining the elementary reflectors
    float *h_work; // workspace
    magma_int_t info;
    magma_int_t ione = 1, nb, lwork; // lwork - workspace size
    float *dT; // store nb*nb blocks of triangular matrices used
    magma_int_t ilo=ione, ihi=n; // in reduction
    float mone= MAGMA_S_NEG_ONE;
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    float work[1]; // used in difference computations
    nb = magma_get_sgehrd_nb(n); // optimal block size for sgehrd
    lwork = n*nb;
    magma_smallocc_cpu(&a,n2); // host memory for a
    magma_smallocc_cpu(&tau,n); // host memory for tau
    magma_smallocc_pinned(&r,n2); // host memory for r
    magma_smallocc_pinned(&r1,n2); // host memory for r1
    magma_smallocc_pinned(&h_work,lwork); // host memory for h_work
    magma_smallocc(&dT,nb*n); // device memory for dT
    // Randomize the matrix a and copy a -> r, a -> r1
    lapackf77_slarnv( &ione, ISEED, &n2, a );
    lapackf77_slacpy(MagmaFullStr,&n,&n,a,&n,r,&n);
    lapackf77_slacpy(MagmaFullStr,&n,&n,a,&n,r1,&n);
    // MAGMA
    gpu_time = magma_sync_wtime(NULL);
    // reduce the matrix r to upper Hessenberg form by an
    // orthogonal transformation, Magma version

    magma_sgehrd(n,ilo,ihi,r,n,tau,h_work,lwork,dT,&info);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("Magma time: %7.3f sec.\n",gpu_time); // Magma time
    {
        int i, j;
        for(j=0; j<n-1; j++)
            for(i=j+2; i<n; i++)
                r[i+j*n] = MAGMA_S_ZERO;
    }
    printf("upper left corner of the Hessenberg form:\n");
    magma_sprint(5,5,r,n); // print the Hessenberg form
    // LAPACK
    cpu_time=magma_wtime();
    // reduce the matrix r1 to upper Hessenberg form by an
    // orthogonal transformation, Lapack version

```

```

lapackf77_sgehrd(&n,&ione,&n,r1,&n,tau,h_work,&lwork,&info);
cpu_time=magma_wtime()-cpu_time;
printf("Lapack time: %7.3f sec.\n",cpu_time);
{
    int i, j;
    for(j=0; j<n-1; j++)
        for(i=j+2; i<n; i++)
            r1[i+j*n] = MAGMA_S_ZERO;
}
// difference
blasf77_saxpy(&n2,&mone,r,&ione,r1,&ione);
printf("max difference: %e\n",
        lapackf77_slange("M", &n, &n, r1, &n, work));
free(a); // free host memory
free(tau); // free host memory
magma_free_pinned(h_work); // free host memory
magma_free_pinned(r); // free host memory
magma_free_pinned(r1); // free host memory
magma_free(dT); // free device memory
magma_finalize( ); // finalize Magma
return EXIT_SUCCESS;
}
//Magma time: 0.365 sec.
//upper left corner of the Hessenberg form:
//[
// 0.1206 -19.4276 -11.6704 0.5872 -0.0777
// -26.2667 765.4211 444.0294 -6.4941 0.5035
// 0. 444.5269 258.5998 -4.0942 0.2565
// 0. 0. -15.2374 0.3507 0.0222
// 0. 0. 0. -13.0577 -0.1760
//];
//Lapack time: 0.916 sec.
//max difference: 1.018047e-03

```

4.6.10 magma_sgehrd - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv)
{
    magma_init(); // initialize Magma
    double gpu_time, cpu_time;
    magma_int_t n=2048, n2=n*n;
    float *a, *r, *r1; // a,r,r1 - nxn matrices
    float *tau; // scalars defining the elementary reflectors
    float *h_work; // workspace
    magma_int_t info;

```

```

magma_int_t ione = 1, nb, lwork; // lwork - workspace size
float *dT; // store nb*nb blocks of triangular matrices used
magma_int_t ilo=ione, ihi=n; // in reduction
float mone= MAGMA_S_NEG_ONE;
magma_int_t ISEED[4] = {0,0,0,1}; // seed
float work[1]; // used in difference computations
nb = magma_get_sgehrd_nb(n); // optimal block size for sgehrd
lwork = n*nb;
cudaMallocManaged(&tau,n*sizeof(float)); //unif. mem.for tau
cudaMallocManaged(&a,n2*sizeof(float)); //unif. memory for a
cudaMallocManaged(&r,n2*sizeof(float)); //unif. memory for r
cudaMallocManaged(&r1,n2*sizeof(float)); //unif. memory for r1
cudaMallocManaged(&h_work,lwork*sizeof(float)); //m.f.h_work
cudaMallocManaged(&dT,nb*n*sizeof(float)); //unif. mem.for dT
// Randomize the matrix a and copy a -> r, a -> r1
lapackf77_slarnv( &ione, ISEED, &n2, a );
lapackf77_slacpy(MagmaFullStr,&n,&n,a,&n,r,&n);
lapackf77_slacpy(MagmaFullStr,&n,&n,a,&n,r1,&n);
// MAGMA
gpu_time = magma_sync_wtime(NULL);
// reduce the matrix r to upper Hessenberg form by an
// orthogonal transformation, Magma version

magma_sgehrd(n,ilo,ihl,r,n,tau,h_work,lwork,dT,&info);

gpu_time = magma_sync_wtime(NULL)-gpu_time;
printf("Magma time: %7.3f sec.\n",gpu_time); // Magma time
{
    int i, j;
    for(j=0; j<n-1; j++)
        for(i=j+2; i<n; i++)
            r[i+j*n] = MAGMA_S_ZERO;
}
printf("upper left corner of the Hessenberg form:\n");
magma_sprint(5,5,r,n); // print the Hessenberg form
// LAPACK
cpu_time=magma_wtime();
// reduce the matrix r1 to upper Hessenberg form by an
// orthogonal transformation, Lapack version
lapackf77_sgehrd(&n,&ione,&n,r1,&n,tau,h_work,&lwork,&info);
cpu_time=magma_wtime()-cpu_time;
printf("Lapack time: %7.3f sec.\n",cpu_time);
{
    int i, j;
    for(j=0; j<n-1; j++)
        for(i=j+2; i<n; i++)
            r1[i+j*n] = MAGMA_S_ZERO;
}
// difference
blasf77_saxpy(&n2,&mone,r,&ione,r1,&ione);
printf("max difference: %e\n",
        lapackf77_slange("M", &n, &n, r1, &n, work));

```

```

    magma_free(a);                // free memory
    magma_free(tau);              // free memory
    magma_free(h_work);          // free memory
    magma_free(r);               // free memory
    magma_free(r1);              // free memory
    magma_free(dT);              // free memory
    magma_finalize( );           // finalize Magma
    return EXIT_SUCCESS;
}
//Magma time:    0.403 sec.
//upper left corner of the Hessenberg form:
//[
//   0.1206  -19.4276  -11.6704   0.5872   -0.0777
//  -26.2667  765.4211  444.0294  -6.4941   0.5035
//    0.      444.5269  258.5998  -4.0942   0.2565
//    0.      0.      -15.2374   0.3507   0.0222
//    0.      0.      0.      -13.0577  -0.1760
//];
//Lapack time:    0.644 sec.
//max difference: 1.018047e-03

```

4.6.11 magma_dgehrd - reduce a general matrix to the upper Hessenberg form in double precision, CPU interface

This function using the double precision reduces a general real $n \times n$ matrix A defined on the host to upper Hessenberg form:

$$Q^T A Q = H,$$

where Q is an orthogonal matrix and H has zero elements below the first subdiagonal. The orthogonal matrix Q is represented as a product of elementary reflectors $H(ilo) \dots H(ihi)$, where $H(k) = I - \tau_k v_k v_k^T$. The real scalars τ_k are stored in an array τ and the information on vectors v_k is stored on exit in the lower triangular part of A below the first subdiagonal: $v_k(1 : k) = 0, v_k(k + 1) = 1$ and $v_k(ihi + 1 : n) = 0$; $v_k(k + 2 : ihi)$ is stored in $A(k + 2 : ihi, k)$. The function uses also an array dT defined on the device, storing blocks of triangular matrices used in the reduction process. See [magma-X.Y.Z/src/dgehrd.cpp](#) for more details.

```

#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv)
{
    magma_init();                // initialize Magma
    double  gpu_time, cpu_time;
    magma_int_t n=2048, n2=n*n;
    double *a, *r, *r1;         // a,r,r1 - nxn matrices on the host

```



```

double *tau;    // scalars defining the elementary reflectors
double *h_work;                               // workspace
magma_int_t info;
magma_int_t ione = 1, nb, lwork;    // lwork - workspace size
double *dT; // store nb*nb blocks of triangular matrices used
magma_int_t ilo=ione, ihi=n;        // in reduction
double mone= MAGMA_D_NEG_ONE;
magma_int_t ISEED[4] = {0,0,0,1};    // seed
double work[1];    // used in difference computations
nb = magma_get_dgehrd_nb(n); // optimal block size for dgehrd
lwork = n*nb;
magma_dmalloc_cpu(&a,n2);    // host memory for a
magma_dmalloc_cpu(&tau,n);    // host memory for tau
magma_dmalloc_pinned(&r,n2);    // host memory for r
magma_dmalloc_pinned(&r1,n2);    // host memory for r1
magma_dmalloc_pinned(&h_work,lwork); // host memory for h_work
magma_dmalloc(&dT,nb*n);    // device memory for dT
// Randomize the matrix a and copy a -> r, a -> r1
lapackf77_dlarnv( &ione, ISEED, &n2, a );
lapackf77_dlacpy(MagmaFullStr,&n,&n,a,&n,r,&n);
lapackf77_dlacpy(MagmaFullStr,&n,&n,a,&n,r1,&n);
// MAGMA
gpu_time = magma_sync_wtime(NULL);
// reduce the matrix r to upper Hessenberg form by an
// orthogonal transformation, Magma version

magma_dgehrd(n,ilo,ihl,r,n,tau,h_work,lwork,dT,&info);

gpu_time = magma_sync_wtime(NULL)-gpu_time;
printf("Magma time: %7.3f sec.\n",gpu_time);    // Magma time
{
    int i, j;
    for(j=0; j<n-1; j++)
        for(i=j+2; i<n; i++)
            r[i+j*n] = MAGMA_D_ZERO;
}
printf("upper left corner of the Hessenberg form:\n");
magma_dprint(5,5,r,n);    // print the Hessenberg form
// LAPACK
cpu_time=magma_wtime();
// reduce the matrix r1 to upper Hessenberg form by an
// orthogonal transformation, Lapack version
lapackf77_dgehrd(&n,&ione,&n,r1,&n,tau,h_work,&lwork,&info);
cpu_time=magma_wtime()-cpu_time;
printf("Lapack time: %7.3f sec.\n",cpu_time);
{
    int i, j;
    for(j=0; j<n-1; j++)
        for(i=j+2; i<n; i++)
            r1[i+j*n] = MAGMA_D_ZERO;
}
// difference

```

```

blasf77_daxpy(&n2,&mone,r,&ione,r1,&ione);
printf("max difference: %e\n",
        lapackf77_dlange("M", &n, &n, r1, &n, work));
free(a); // free host memory
free(tau); // free host memory
magma_free_pinned(h_work); // free host memory
magma_free_pinned(r); // free host memory
magma_free_pinned(r1); // free host memory
magma_free(dT); // free device memory
magma_finalize( ); // finalize Magma
return EXIT_SUCCESS;
}
//Magma time: 0.525 sec.
//upper left corner of the Hessenberg form:
//[
// 0.1206 -19.4276 -11.6704 0.5872 -0.0777
// -26.2667 765.4211 444.0295 -6.4941 0.5035
// 0. 444.5269 258.5999 -4.0943 0.2565
// 0. 0. -15.2374 0.3507 0.0222
// 0. 0. 0. -13.0577 -0.1760
//];
//Lapack time: 2.067 sec.
//max difference: 1.444213e-12

```

4.6.12 magma_dgehrd - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv)
{
    magma_init(); // initialize Magma
    double gpu_time, cpu_time;
    magma_int_t n=2048, n2=n*n;
    double *a, *r, *r1; // a,r,r1 - nxn matrices
    double *tau; // scalars defining the elementary reflectors
    double *h_work; // workspace
    magma_int_t info;
    magma_int_t ione = 1, nb, lwork; // lwork - workspace size
    double *dT; // store nb*nb blocks of triangular matrices used
    magma_int_t ilo=ione, ihi=n; // in reduction
    double mone= MAGMA_D_NEG_ONE;
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    double work[1]; // used in difference computations
    nb = magma_get_dgehrd_nb(n); // optimal block size for dgehrd
    lwork = n*nb;
    cudaMallocManaged(&tau,n*sizeof(double)); //unif. mem.for tau
    cudaMallocManaged(&a,n2*sizeof(double)); //unif. memory for a

```

```

    cudaMallocManaged(&r,n2*sizeof(double)); //unif. memory for r
    cudaMallocManaged(&r1,n2*sizeof(double)); //unif. mem. for r1
    cudaMallocManaged(&h_work,lwork*sizeof(double)); //m.f.h_work
    cudaMallocManaged(&dT,nb*n*sizeof(double)); //unif.mem.for dT
// Randomize the matrix a and copy a -> r, a -> r1
    lapackf77_dlarv( &ione, ISEED, &n2, a );
    lapackf77_dlacpy(MagmaFullStr,&n,&n,a,&n,r,&n);
    lapackf77_dlacpy(MagmaFullStr,&n,&n,a,&n,r1,&n);
// MAGMA
    gpu_time = magma_sync_wtime(NULL);
// reduce the matrix r to upper Hessenberg form by an
// orthogonal transformation, Magma version

    magma_dgehrd(n,ilo,ihi,r,n,tau,h_work,lwork,dT,&info);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("Magma time: %7.3f sec.\n",gpu_time);    // Magma time
{
    int i, j;
    for(j=0; j<n-1; j++)
        for(i=j+2; i<n; i++)
            r[i+j*n] = MAGMA_D_ZERO;
}
    printf("upper left corner of the Hessenberg form:\n");
    magma_dprint(5,5,r,n);    // print the Hessenberg form
// LAPACK
    cpu_time=magma_wtime();
// reduce the matrix r1 to upper Hessenberg form by an
// orthogonal transformation, Lapack version
    lapackf77_dgehrd(&n,&ione,&n,r1,&n,tau,h_work,&lwork,&info);
    cpu_time=magma_wtime()-cpu_time;
    printf("Lapack time: %7.3f sec.\n",cpu_time);
{
    int i, j;
    for(j=0; j<n-1; j++)
        for(i=j+2; i<n; i++)
            r1[i+j*n] = MAGMA_D_ZERO;
}
// difference
    blasf77_daxpy(&n2,&mone,r,&ione,r1,&ione);
    printf("max difference: %e\n",
        lapackf77_dlange("M", &n, &n, r1, &n, work));
    magma_free(a);    // free memory
    magma_free(tau);    // free memory
    magma_free(h_work);    // free memory
    magma_free(r);    // free memory
    magma_free(r1);    // free memory
    magma_free(dT);    // free memory
    magma_finalize( );    // finalize Magma
    return EXIT_SUCCESS;
}
//Magma time:    0.572 sec.

```

```
//upper left corner of the Hessenberg form:
//[
//  0.1206 -19.4276 -11.6704  0.5872 -0.0777
// -26.2667 765.4211 444.0295 -6.4941  0.5035
//  0.      444.5269 258.5999 -4.0943  0.2565
//  0.      0.      -15.2374  0.3507  0.0222
//  0.      0.      0.      -13.0577 -0.1760
//];
//Lapack time:  1.753 sec.
//max difference: 1.444213e-12
```

4.7 Eigenvalues and eigenvectors for symmetric matrices

4.7.1 magma_ssyevd - compute the eigenvalues and optionally eigenvectors of a symmetric real matrix in single precision, CPU interface, small matrix

This function computes in single precision all eigenvalues and, optionally, eigenvectors of a real symmetric matrix A defined on the host. The first parameter can take the values `MagmaVec` or `MagmaNoVec` and answers the question whether the eigenvectors are desired. If the eigenvectors are desired, it uses a divide and conquer algorithm. The symmetric matrix A can be stored in lower (`MagmaLower`) or upper (`MagmaUpper`) mode. If the eigenvectors are desired, then on exit A contains orthonormal eigenvectors. The eigenvalues are stored in an array `w`. See [magma-X.Y.Z/src/ssyevd.cpp](#) for more details.

```
#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv) {
    magma_init(); // initialize Magma
    magma_int_t n=1024, n2=n*n;
    float *a, *r; // a, r - nxn matrices on the host
    float *h_work; // workspace
    magma_int_t lwork; // h_work size
    magma_int_t *iwork; // workspace
    magma_int_t liwork; // iwork size
    float *w1, *w2; // w1,w2 - vectors of eigenvalues
    float error, work[1]; // used in difference computations
    magma_int_t ione = 1, i, j, info;
    float mione = -1.0f;
    magma_int_t incr = 1;
    magma_smallocc_cpu(&w1,n); // host memory for real
    magma_smallocc_cpu(&w2,n); // eigenvalues
```

```

    magma_smallocc_cpu(&a,n2);                // host memory for a
    magma_smallocc_cpu(&r,n2);                // host memory for r
// Query for workspace sizes
    float aux_work[1];
    magma_int_t aux_iwork[1];
    magma_ssyevd(MagmaVec,MagmaLower,n,r,n,w1,aux_work,-1,
                aux_iwork,-1,&info );

    lwork  = (magma_int_t) aux_work[0];
    liwork = aux_iwork[0];
    iwork=(magma_int_t*)malloc(liwork*sizeof(magma_int_t));
    magma_smallocc_cpu(&h_work,lwork);        // memory for workspace
// define a, r                                //      [1 0 0 0 0 ...]
    for(i=0;i<n;i++){                          //      [0 2 0 0 0 ...]
        a[i*n+i]=(float)(i+1);                  // a = [0 0 3 0 0 ...]
        r[i*n+i]=(float)(i+1);                  //      [0 0 0 4 0 ...]
    }                                              //      [0 0 0 0 5 ...]
    printf("upper left corner of a:\n");        // .....
    magma_sprint(5,5,a,n);                      // print part of a
// compute the eigenvalues and eigenvectors for a symmetric,
// real nxn matrix; Magma version

    magma_ssyevd(MagmaVec,MagmaLower,n,r,n,w1,h_work,lwork,iwork,
                liwork,&info);

    printf("first 5 eigenvalues of a:\n");
    for(j=0;j<5;j++){
        printf("%f\n",w1[j]);                    // print first eigenvalues
    }
    printf("left upper corner of the matrix of eigenvectors:\n");
    magma_sprint(5,5,r,n); // part of the matrix of eigenvectors
// Lapack version
    lapackf77_ssyevd("V","L",&n,a,&n,w2,h_work,&lwork,iwork,
                    &liwork,&info);

// difference in eigenvalues
    blasf77_saxpy( &n, &mione, w1, &incr, w2, &incr);
    error = lapackf77_slange( "M", &n, &ione, w2, &n, work );
    printf("difference in eigenvalues: %e\n",error);
    free(w1);                                     // free host memory
    free(w2);                                     // free host memory
    free(a);                                     // free host memory
    free(r);                                     // free host memory
    free(h_work);                                // free host memory
    magma_finalize();                             // finalize Magma
    return EXIT_SUCCESS;
}
//upper left corner of a:
//[
//  1.0000   0.      0.      0.      0.
//  0.      2.0000   0.      0.      0.
//  0.      0.      3.0000   0.      0.
//  0.      0.      0.      4.0000   0.
//  0.      0.      0.      0.      5.0000
//];

```

```

//first 5 eigenvalues of a:
//1.000000
//2.000000
//3.000000
//4.000000
//5.000000
//left upper corner of the matrix of eigenvectors:
//[
//  1.0000  0.      0.      0.      0.
//  0.      1.0000  0.      0.      0.
//  0.      0.      1.0000  0.      0.
//  0.      0.      0.      1.0000  0.
//  0.      0.      0.      0.      1.0000
//];
//difference in eigenvalues: 0.000000e+00

```

4.7.2 magma_ssyevd - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv) {
    magma_init(); // initialize Magma
    magma_int_t n=1024, n2=n*n;
    float *a, *r; // a, r - nxn matrices
    float *h_work; // workspace
    magma_int_t lwork; // h_work size
    magma_int_t *iwork; // workspace
    magma_int_t liwork; // iwork size
    float *w1, *w2; // w1,w2 - vectors of eigenvalues
    float error, work[1]; // used in difference computations
    magma_int_t ione = 1, i, j, info;
    float mione = -1.0;
    magma_int_t incr = 1;
    cudaMallocManaged(&w1, n*sizeof(float)); //unified memory for
    cudaMallocManaged(&w2, n*sizeof(float)); //eigenvalues
    cudaMallocManaged(&a, n2*sizeof(float)); //unif. memory for a
    cudaMallocManaged(&r, n2*sizeof(float)); //unif. memory for r
    // Query for workspace sizes
    float aux_work[1];
    magma_int_t aux_iwork[1];
    magma_ssyevd(MagmaVec, MagmaLower, n, r, n, w1, aux_work, -1,
                aux_iwork, -1, &info );
    lwork = (magma_int_t) aux_work[0];
    liwork = aux_iwork[0]; // unified memory for workspace:
    cudaMallocManaged(&iwork, liwork*sizeof(magma_int_t));
    cudaMallocManaged(&h_work, lwork*sizeof(float));

```

```

// define a, r
for(i=0; i<n; i++){
    a[i*n+i]=(float)(i+1);
    r[i*n+i]=(float)(i+1);
}
printf("upper left corner of a:\n");
magma_sprint(5,5,a,n);
// compute the eigenvalues and eigenvectors for a symmetric,
// real nxn matrix; Magma version

magma_ssyeval(MagmaVec,MagmaLower,n,r,n,w1,h_work,lwork,iwork,
              liwork,&info);

printf("first 5 eigenvalues of a:\n");
for(j=0; j<5; j++){
    printf("%f\n",w1[j]);
    printf("left upper corner of the matrix of eigenvectors:\n");
    magma_sprint(5,5,r,n); // part of the matrix of eigenvectors
// Lapack version
    lapackf77_ssyeval("V","L",&n,a,&n,w2,h_work,&lwork,iwork,
                      &liwork,&info);
// difference in eigenvalues
    blasf77_saxpy( &n, &mione, w1, &incr, w2, &incr);
    error = lapackf77_slange( "M", &n, &iione, w2, &n, work );
    printf("difference in eigenvalues: %e\n",error);
    magma_free(w1);
    magma_free(w2);
    magma_free(a);
    magma_free(r);
    magma_free(h_work);
    magma_finalize();
    return EXIT_SUCCESS;
}

//upper left corner of a:
//[
//  1.0000  0.      0.      0.      0.
//  0.      2.0000  0.      0.      0.
//  0.      0.      3.0000  0.      0.
//  0.      0.      0.      4.0000  0.
//  0.      0.      0.      0.      5.0000
//];

//first 5 eigenvalues of a:
//1.000000
//2.000000
//3.000000
//4.000000
//5.000000

//left upper corner of the matrix of eigenvectors:

```

```
//[
//  1.0000  0.      0.      0.      0.
//  0.      1.0000  0.      0.      0.
//  0.      0.      1.0000  0.      0.
//  0.      0.      0.      1.0000  0.
//  0.      0.      0.      0.      1.0000
//];
//difference in eigenvalues: 0.000000e+00
```

4.7.3 magma_dsyevd - compute the eigenvalues and optionally eigenvectors of a symmetric real matrix in double precision, CPU interface, small matrix

This function computes in double precision all eigenvalues and, optionally, eigenvectors of a real symmetric matrix A defined on the host. The first parameter can take the values `MagmaVec` or `MagmaNoVec` and answers the question whether the eigenvectors are desired. If the eigenvectors are desired, it uses a divide and conquer algorithm. The symmetric matrix A can be stored in lower (`MagmaLower`) or upper (`MagmaUpper`) mode. If the eigenvectors are desired, then on exit A contains orthonormal eigenvectors. The eigenvalues are stored in an array `w`. See [magma-X.Y.Z/src/dsyevd.cpp](#) for more details.

```
#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv) {
    magma_init(); // initialize Magma
    magma_int_t n=1024, n2=n*n;
    double *a, *r; // a, r - nxn matrices on the host
    double *h_work; // workspace
    magma_int_t lwork; // h_work size
    magma_int_t *iwork; // workspace
    magma_int_t liwork; // iwork size
    double *w1, *w2; // w1,w2 - vectors of eigenvalues
    double error, work[1]; // used in difference computations
    magma_int_t ione = 1, i, j, info;
    double mione = -1.0;
    magma_int_t incr = 1;
    magma_dmalloc_cpu(&w1,n); // host memory for real
    magma_dmalloc_cpu(&w2,n); // eigenvalues
    magma_dmalloc_cpu(&a,n2); // host memory for a
    magma_dmalloc_cpu(&r,n2); // host memory for r
    // Query for workspace sizes
    double aux_work[1];
    magma_int_t aux_iwork[1];
    magma_dsyevd(MagmaVec, MagmaLower, n, r, n, w1, aux_work, -1,
                aux_iwork, -1, &info );
    lwork = (magma_int_t) aux_work[0];
```



```

    liwork = aux_iwork[0];
    iwork=(magma_int_t*)malloc(liwork*sizeof(magma_int_t));
    magma_dmalloc_cpu(&h_work,lwork);    // memory for workspace
// define a, r                                //      [1 0 0 0 0 ...]
for(i=0;i<n;i++){                            //      [0 2 0 0 0 ...]
    a[i*n+i]=(double)(i+1);                // a = [0 0 3 0 0 ...]
    r[i*n+i]=(double)(i+1);                //      [0 0 0 4 0 ...]
}                                            //      [0 0 0 0 5 ...]
printf("upper left corner of a:\n");    //      .....
magma_dprint(5,5,a,n);                    // print part of a
// compute the eigenvalues and eigenvectors for a symmetric,
// real nxn matrix; Magma version

    magma_dsyevd(MagmaVec,MagmaLower,n,r,n,w1,h_work,lwork,iwork,
                liwork,&info);

    printf("first 5 eigenvalues of a:\n");
    for(j=0;j<5;j++)
        printf("%f\n",w1[j]);                // print first eigenvalues
    printf("left upper corner of the matrix of eigenvectors:\n");
    magma_dprint(5,5,r,n); // part of the matrix of eigenvectors
// Lapack version
    lapackf77_dsyevd("V","L",&n,a,&n,w2,h_work,&lwork,iwork,
                    &liwork,&info);

// difference in eigenvalues
    blasf77_daxpy( &n, &mione, w1, &incr, w2, &incr);
    error = lapackf77_dlange( "M", &n, &ione, w2, &n, work );
    printf("difference in eigenvalues: %e\n",error);
    free(w1);                                // free host memory
    free(w2);                                // free host memory
    free(a);                                // free host memory
    free(r);                                // free host memory
    free(h_work);                            // free host memory
    magma_finalize();                        // finalize Magma
    return EXIT_SUCCESS;
}

//upper left corner of a:
//[
//  1.0000   0.      0.      0.      0.
//  0.      2.0000   0.      0.      0.
//  0.      0.      3.0000   0.      0.
//  0.      0.      0.      4.0000   0.
//  0.      0.      0.      0.      5.0000
//];

//first 5 eigenvalues of a:
//1.000000
//2.000000
//3.000000
//4.000000
//5.000000

```

```
//left upper corner of the matrix of eigenvectors:
//[
//  1.0000  0.      0.      0.      0.
//  0.      1.0000  0.      0.      0.
//  0.      0.      1.0000  0.      0.
//  0.      0.      0.      1.0000  0.
//  0.      0.      0.      0.      1.0000
//];
//difference in eigenvalues: 0.000000e+00
```

4.7.4 magma_dsyevd - unified memory version

```
#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv) {
    magma_init(); // initialize Magma
    magma_int_t n=1024, n2=n*n;
    double *a, *r; // a, r - nxn matrices
    double *h_work; // workspace
    magma_int_t lwork; // h_work size
    magma_int_t *iwork; // workspace
    magma_int_t liwork; // iwork size
    double *w1, *w2; // w1,w2 - vectors of eigenvalues
    double error, work[1]; // used in difference computations
    magma_int_t ione = 1, i, j, info;
    double mione = -1.0;
    magma_int_t incr = 1;
    cudaMallocManaged(&w1, n*sizeof(double)); //unified memory for
    cudaMallocManaged(&w2, n*sizeof(double)); //eigenvalues
    cudaMallocManaged(&a, n2*sizeof(double)); //unif. memory for a
    cudaMallocManaged(&r, n2*sizeof(double)); //unif. memory for r
    // Query for workspace sizes
    double aux_work[1];
    magma_int_t aux_iwork[1];
    magma_dsyevd(MagmaVec, MagmaLower, n, r, n, w1, aux_work, -1,
                aux_iwork, -1, &info );

    lwork = (magma_int_t) aux_work[0];
    liwork = aux_iwork[0]; // unified memory for workspace:
    cudaMallocManaged(&iwork, liwork*sizeof(magma_int_t));
    cudaMallocManaged(&h_work, lwork*sizeof(double));
    // define a, r // [1 0 0 0 0 ...]
    for(i=0; i<n; i++){ // [0 2 0 0 0 ...]
        a[i*n+i]=(double)(i+1); // a = [0 0 3 0 0 ...]
        r[i*n+i]=(double)(i+1); // [0 0 0 4 0 ...]
    } // [0 0 0 0 5 ...]
    printf("upper left corner of a:\n"); // .....
    magma_dprint(5,5,a,n); // print part of a
```

```

// compute the eigenvalues and eigenvectors for a symmetric,
// real nxn matrix; Magma version

magma_dsyevd(MagmaVec,MagmaLower,n,r,n,w1,h_work,lwork,iwork,
             liwork,&info);

printf("first 5 eigenvalues of a:\n");
for(j=0;j<5;j++)
    printf("%f\n",w1[j]);           // print first eigenvalues
printf("left upper corner of the matrix of eigenvectors:\n");
magma_dprint(5,5,r,n); // part of the matrix of eigenvectors
// Lapack version
lapackf77_dsyevd("V","L",&n,a,&n,w2,h_work,&lwork,iwork,
                 &liwork,&info);

// difference in eigenvalues
blasf77_daxpy( &n, &mione, w1, &incr, w2, &incr);
error = lapackf77_dlange( "M", &n, &iione, w2, &n, work );
printf("difference in eigenvalues: %e\n",error);
magma_free(w1);           // free memory
magma_free(w2);           // free memory
magma_free(a);            // free memory
magma_free(r);            // free memory
magma_free(h_work);       // free memory
magma_finalize();         // finalize Magma
return EXIT_SUCCESS;
}

//upper left corner of a:
//[
//  1.0000  0.      0.      0.      0.
//  0.      2.0000  0.      0.      0.
//  0.      0.      3.0000  0.      0.
//  0.      0.      0.      4.0000  0.
//  0.      0.      0.      0.      5.0000
//];
//first 5 eigenvalues of a:
//1.000000
//2.000000
//3.000000
//4.000000
//5.000000
//left upper corner of the matrix of eigenvectors:
//[
//  1.0000  0.      0.      0.      0.
//  0.      1.0000  0.      0.      0.
//  0.      0.      1.0000  0.      0.
//  0.      0.      0.      1.0000  0.
//  0.      0.      0.      0.      1.0000
//];
//difference in eigenvalues: 0.000000e+00

```

4.7.5 magma_ssyevd - compute the eigenvalues and optionally eigenvectors of a symmetric real matrix in single precision, CPU interface, big matrix

This function computes in single precision all eigenvalues and, optionally, eigenvectors of a real symmetric matrix A defined on the host. The first parameter can take the values `MagmaVec` or `MagmaNoVec` and answers the question whether the eigenvectors are desired. If the eigenvectors are desired, it uses a divide and conquer algorithm. The symmetric matrix A can be stored in lower (`MagmaLower`) or upper (`MagmaUpper`) mode. If the eigenvectors are desired, then on exit A contains orthonormal eigenvectors. The eigenvalues are stored in an array `w`. See [magma-X.Y.Z/src/ssyevd.cpp](#) for more details.

```
#include <stdio.h>
#include <cuda.h>
#include "magma.h"
#include "magma_lapack.h"
int main( int argc, char** argv) {
    magma_init(); // initialize Magma
    double gpu_time, cpu_time;
    magma_int_t n=8192, n2=n*n;
    float *a, *r; // a, r - nxn matrices on the host
    float *h_work; // workspace
    magma_int_t lwork; // h_work size
    magma_int_t *iwork; // workspace
    magma_int_t liwork; // iwork size
    float *w1, *w2; // w1,w2 - vectors of eigenvalues
    float error, work[1]; // used in difference computations
    magma_int_t ione = 1, info;
    float mione = -1.0f;
    magma_int_t incr = 1;
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    magma_smallocc_cpu(&w1,n); // host memory for real
    magma_smallocc_cpu(&w2,n); // eigenvalues
    magma_smallocc_cpu(&a,n2); // host memory for a
    magma_smallocc_cpu(&r,n2); // host memory for r
    // Query for workspace sizes
    float aux_work[1];
    magma_int_t aux_iwork[1];
    magma_ssyevd(MagmaVec, MagmaLower, n, r, n, w1, aux_work, -1,
                aux_iwork, -1, &info );

    lwork = (magma_int_t) aux_work[0];
    liwork = aux_iwork[0];
    iwork=(magma_int_t*) malloc(liwork*sizeof(magma_int_t));
    magma_smallocc_cpu(&h_work,lwork); // memory for workspace
    // Randomize the matrix a and copy a -> r
    lapackf77_slarnv(&ione, ISEED, &n2, a);
    lapackf77_slacpy(MagmaFullStr, &n, &n, a, &n, r, &n);
    gpu_time = magma_sync_wtime(NULL);
```

```

// compute the eigenvalues and eigenvectors for a symmetric,
// real nxn matrix; Magma version

magma_ssyevd(MagmaVec,MagmaLower,n,r,n,w1,h_work,lwork,iwork,
             liwork,&info);

gpu_time = magma_sync_wtime(NULL)-gpu_time;
printf("ssyevd gpu time: %7.5f sec.\n",gpu_time); // Magma
// Lapack version // time
cpu_time=magma_wtime();
lapackf77_ssyevd("V","L",&n,&a,&n,w2,h_work,&lwork,iwork,
                &liwork,&info);

cpu_time=magma_wtime()-cpu_time;
printf("ssyevd cpu time: %7.5f sec.\n",cpu_time); // Lapack
// difference in eigenvalues // time
blasf77_saxpy( &n, &mione, w1, &incr, w2, &incr);
error = lapackf77_slange( "M", &n, &ione, w2, &n, work );
printf("difference in eigenvalues: %e\n",error);
free(w1); // free host memory
free(w2); // free host memory
free(a); // free host memory
free(r); // free host memory
free(h_work); // free host memory
magma_finalize(); // finalize Magma
return EXIT_SUCCESS;
}
//ssyevd gpu time: 5.58410 sec.
//ssyevd cpu time: 49.01886 sec.
//difference in eigenvalues: 9.765625e-04

```

4.7.6 magma_ssyevd - unified memory version, big matrix

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv) {
    magma_init(); // initialize Magma
    double gpu_time, cpu_time;
    magma_int_t n=8192, n2=n*n;
    float *a, *r; // a, r - nxn matrices
    float *h_work; // workspace
    magma_int_t lwork; // h_work size
    magma_int_t *iwork; // workspace
    magma_int_t liwork; // iwork size
    float *w1, *w2; // w1,w2 - vectors of eigenvalues
    float error, work[1]; // used in difference computations
    magma_int_t ione = 1, info;
    float mione = -1.0f;

```

```

magma_int_t incr = 1;
magma_int_t ISEED[4] = {0,0,0,1}; // seed
cudaMallocManaged(&w1,n*sizeof(float)); //unified memory for
cudaMallocManaged(&w2,n*sizeof(float)); //eigenvalues
cudaMallocManaged(&a,n2*sizeof(float)); //unif. memory for a
cudaMallocManaged(&r,n2*sizeof(float)); //unif. memory for r
// Query for workspace sizes
float aux_work[1];
magma_int_t aux_iwork[1];
magma_ssyevd(MagmaVec,MagmaLower,n,r,n,w1,aux_work,-1,
             aux_iwork,-1,&info );

lwork = (magma_int_t) aux_work[0];
liwork = aux_iwork[0]; //unified memory for workspace:
cudaMallocManaged(&iwork,liwork*sizeof(magma_int_t));
cudaMallocManaged(&h_work,lwork*sizeof(float));
// Randomize the matrix a and copy a -> r
lapackf77_slarnv(&ione,ISEED,&n2,a);
lapackf77_slacpy(MagmaFullStr,&n,&n,a,&n,r,&n);
gpu_time = magma_sync_wtime(NULL);
// compute the eigenvalues and eigenvectors for a symmetric,
// real nxn matrix; Magma version

magma_ssyevd(MagmaVec,MagmaLower,n,r,n,w1,h_work,lwork,iwork,
             liwork,&info);

gpu_time = magma_sync_wtime(NULL)-gpu_time;
printf("ssyevd gpu time: %7.5f sec.\n",gpu_time); // Magma
// Lapack version // time
cpu_time=magma_wtime();
lapackf77_ssyevd("V","L",&n,a,&n,w2,h_work,&lwork,iwork,
                &liwork,&info);

cpu_time=magma_wtime()-cpu_time;
printf("ssyevd cpu time: %7.5f sec.\n",cpu_time); // Lapack
// difference in eigenvalues // time
blasf77_saxpy( &n, &mione, w1, &incr, w2, &incr);
error = lapackf77_slange( "M", &n, &ione, w2, &n, work );
printf("difference in eigenvalues: %e\n",error);
magma_free(w1); // free memory
magma_free(w2); // free memory
magma_free(a); // free memory
magma_free(r); // free memory
magma_free(h_work); // free memory
magma_finalize(); // finalize Magma
return EXIT_SUCCESS;
}
//ssyevd gpu time: 5.77196 sec.
//ssyevd cpu time: 51.33320 sec.
//difference in eigenvalues: 9.765625e-04

```

4.7.7 magma_dsyevd - compute the eigenvalues and optionally eigenvectors of a symmetric real matrix in double precision, CPU interface, big matrix

This function computes in double precision all eigenvalues and, optionally, eigenvectors of a real symmetric matrix A defined on the host. The first parameter can take the values `MagmaVec` or `MagmaNoVec` and answers the question whether the eigenvectors are desired. If the eigenvectors are desired, it uses a divide and conquer algorithm. The symmetric matrix A can be stored in lower (`MagmaLower`) or upper (`MagmaUpper`) mode. If the eigenvectors are desired, then on exit A contains orthonormal eigenvectors. The eigenvalues are stored in an array `w`. See [magma-X.Y.Z/src/dsyevd.cpp](#) for more details.

```
#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv) {
    magma_init(); // initialize Magma
    double gpu_time, cpu_time;
    magma_int_t n=8192, n2=n*n;
    double *a, *r; // a, r - nxn matrices on the host
    double *h_work; // workspace
    magma_int_t lwork; // h_work size
    magma_int_t *iwork; // workspace
    magma_int_t liwork; // iwork size
    double *w1, *w2; // w1,w2 - vectors of eigenvalues
    double error, work[1]; // used in difference computations
    magma_int_t ione = 1, info;
    double mione = -1.0;
    magma_int_t incr = 1;
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    magma_dmalloc_cpu(&w1,n); // host memory for real
    magma_dmalloc_cpu(&w2,n); // eigenvalues
    magma_dmalloc_cpu(&a,n2); // host memory for a
    magma_dmalloc_cpu(&r,n2); // host memory for r
    // Query for workspace sizes
    double aux_work[1];
    magma_int_t aux_iwork[1];
    magma_dsyevd(MagmaVec, MagmaLower, n, r, n, w1, aux_work, -1,
                 aux_iwork, -1, &info );

    lwork = (magma_int_t) aux_work[0];
    liwork = aux_iwork[0];
    iwork=(magma_int_t*) malloc(liwork*sizeof(magma_int_t));
    magma_dmalloc_cpu(&h_work,lwork); // memory for workspace
    // Randomize the matrix a and copy a -> r
    lapackf77_dlarnv(&ione, ISEED, &n2, a);
    lapackf77_dlacpy(MagmaFullStr, &n, &n, a, &n, r, &n);
    gpu_time = magma_sync_wtime(NULL);
```

```

// compute the eigenvalues and eigenvectors for a symmetric,
// real nxn matrix; Magma version

magma_dsyevd(MagmaVec,MagmaLower,n,r,n,w1,h_work,lwork,iwork,
             liwork,&info);

gpu_time = magma_sync_wtime(NULL)-gpu_time;
printf("dsyevd gpu time: %7.5f sec.\n",gpu_time); // Magma
// Lapack version // time
cpu_time=magma_wtime();
lapackf77_dsyevd("V","L",&n,&a,&n,w2,h_work,&lwork,iwork,
                &liwork,&info);

cpu_time=magma_wtime()-cpu_time;
printf("dsyevd cpu time: %7.5f sec.\n",cpu_time); // Lapack
// difference in eigenvalues // time
blasf77_daxpy( &n, &mione, w1, &incr, w2, &incr);
error = lapackf77_dlange( "M", &n, &ione, w2, &n, work );
printf("difference in eigenvalues: %e\n",error);
free(w1); // free host memory
free(w2); // free host memory
free(a); // free host memory
free(r); // free host memory
free(h_work); // free host memory
magma_finalize(); // finalize Magma
return EXIT_SUCCESS;
}
//dsyevd gpu time: 17.29120 sec.
//dsyevd cpu time: 91.15194 sec.
//difference in eigenvalues: 1.364242e-11

```

4.7.8 magma_dsyevd - unified memory version, big matrix

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv) {
    magma_init(); // initialize Magma
    double gpu_time, cpu_time;
    magma_int_t n=8192, n2=n*n;
    double *a, *r; // a, r - nxn matrices
    double *h_work; // workspace
    magma_int_t lwork; // h_work size
    magma_int_t *iwork; // workspace
    magma_int_t liwork; // iwork size
    double *w1, *w2; // w1,w2 - vectors of eigenvalues
    double error, work[1]; // used in difference computations
    magma_int_t ione = 1, info;
    double mione = -1.0;

```



```

magma_int_t incr = 1;
magma_int_t ISEED[4] = {0,0,0,1}; // seed
cudaMallocManaged(&w1,n*sizeof(double)); //unified memory for
cudaMallocManaged(&w2,n*sizeof(double)); //eigenvalues
cudaMallocManaged(&a,n2*sizeof(double)); //unif. memory for a
cudaMallocManaged(&r,n2*sizeof(double)); //unif. memory for r
// Query for workspace sizes
double aux_work[1];
magma_int_t aux_iwork[1];
magma_dsyevd(MagmaVec,MagmaLower,n,r,n,w1,aux_work,-1,
             aux_iwork,-1,&info );

lwork = (magma_int_t) aux_work[0];
liwork = aux_iwork[0]; // unified memory for workspace:
cudaMallocManaged(&iwork,liwork*sizeof(magma_int_t));
cudaMallocManaged(&h_work,lwork*sizeof(double));
// Randomize the matrix a and copy a -> r
lapackf77_dlarnv(&ione,ISEED,&n2,a);
lapackf77_dlacpy(MagmaFullStr,&n,&n,a,&n,r,&n);
gpu_time = magma_sync_wtime(NULL);
// compute the eigenvalues and eigenvectors for a symmetric,
// real nxn matrix; Magma version

magma_dsyevd(MagmaVec,MagmaLower,n,r,n,w1,h_work,lwork,iwork,
             liwork,&info);

gpu_time = magma_sync_wtime(NULL)-gpu_time;
printf("dsyevd gpu time: %7.5f sec.\n",gpu_time); // Magma
// Lapack version // time
cpu_time=magma_wtime();
lapackf77_dsyevd("V","L",&n,a,&n,w2,h_work,&lwork,iwork,
                &liwork,&info);

cpu_time=magma_wtime()-cpu_time;
printf("dsyevd cpu time: %7.5f sec.\n",cpu_time); // Lapack
// difference in eigenvalues // time
blasf77_daxpy( &n, &mione, w1, &incr, w2, &incr);
error = lapackf77_dlange( "M", &n, &ione, w2, &n, work );
printf("difference in eigenvalues: %e\n",error);
magma_free(w1); // free memory
magma_free(w2); // free memory
magma_free(a); // free memory
magma_free(r); // free memory
magma_free(h_work); // free memory
magma_finalize(); // finalize Magma
return EXIT_SUCCESS;
}
//dsyevd gpu time: 17.29073 sec.
//dsyevd cpu time: 96.53918 sec.
//difference in eigenvalues: 1.364242e-11

```

4.7.9 magma_ssyevd_gpu - compute the eigenvalues and optionally eigenvectors of a symmetric real matrix in single precision, GPU interface, small matrix

This function computes in single precision all eigenvalues and, optionally, eigenvectors of a real symmetric matrix A defined on the device. The first parameter can take the values `MagmaVec` or `MagmaNoVec` and answers the question whether the eigenvectors are desired. If the eigenvectors are desired, it uses a divide and conquer algorithm. The symmetric matrix A can be stored in lower (`MagmaLower`) or upper (`MagmaUpper`) mode. If the eigenvectors are desired, then on exit A contains orthonormal eigenvectors. The eigenvalues are stored in an array `w`. See [magma-X.Y.Z/src/ssyevd_gpu.cpp](#) for more details.

```
#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv) {
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    magma_int_t n=1024, n2=n*n;
    float *a, *r; // a, r - nxn matrices on the host
    float *d_r; // nxn matrix on the device
    float *h_work; // workspace
    magma_int_t lwork; // h_work size
    magma_int_t *iwork; // workspace
    magma_int_t liwork; // iwork size
    float *w1, *w2; // w1,w2 - vectors of eigenvalues
    float error, work[1]; // used in difference computations
    magma_int_t ione = 1, i, j, info;
    float mione = -1.0f;
    magma_int_t incr = 1;
    magma_smallocc_cpu(&w1,n); // host memory for real
    magma_smallocc_cpu(&w2,n); // eigenvalues
    magma_smallocc_cpu(&a,n2); // host memory for a
    magma_smallocc_cpu(&r,n2); // host memory for r
    magma_smallocc(&d_r,n2); // device memory for d_r
    // Query for workspace sizes
    float aux_work[1];
    magma_int_t aux_iwork[1];
    magma_ssyevd_gpu(MagmaVec,MagmaLower,n,d_r,n,w1,r,n,aux_work,
        -1,aux_iwork,-1,&info );
    lwork = (magma_int_t) aux_work[0];
    liwork = aux_iwork[0];
    iwork=(magma_int_t*) malloc(liwork*sizeof(magma_int_t));
    magma_smallocc_cpu(&h_work,lwork); // memory for workspace
```

```

// define a, r
for(i=0;i<n;i++){
    a[i*n+i]=(float)(i+1);
    r[i*n+i]=(float)(i+1);
}
printf("upper left corner of a:\n");
magma_sprint(5,5,a,n);
magma_ssetmatrix( n, n, a, n, d_r, n,queue); // copy a -> d_r
// compute the eigenvalues and eigenvectors for a symmetric,
// real nxn matrix; Magma version

magma_ssyevd_gpu(MagmaVec,MagmaLower,n,d_r,n,w1,r,n,h_work,
                lwork,iwork,liwork,&info);

printf("first 5 eigenvalues of a:\n");
for(j=0;j<5;j++){
    printf("%f\n",w1[j]);
    printf("left upper corner of the matrix of eigenvectors:\n");
    magma_sgetmatrix( n, n, d_r, n, r, n,queue); // copy d_r -> r
    magma_sprint(5,5,r,n); // part of the matrix of eigenvectors
// Lapack version
    lapackf77_ssyevd("V","L",&n,a,&n,w2,h_work,&lwork,iwork,
                    &liwork,&info);

// difference in eigenvalues
    blasf77_saxpy( &n, &mione, w1, &incr, w2, &incr);
    error = lapackf77_slange( "M", &n, &ione, w2, &n, work );
    printf("difference in eigenvalues: %e\n",error);
    free(w1);
    free(w2);
    free(a);
    free(r);
    free(h_work);
    magma_free(d_r);
    magma_queue_destroy(queue);
    magma_finalize();
    return EXIT_SUCCESS;
}

//upper left corner of a:
//[
//  1.0000  0.      0.      0.      0.
//  0.      2.0000  0.      0.      0.
//  0.      0.      3.0000  0.      0.
//  0.      0.      0.      4.0000  0.
//  0.      0.      0.      0.      5.0000
//];
//first 5 eigenvalues of a:
//1.000000
//2.000000
//3.000000
//4.000000
//5.000000
//left upper corner of the matrix of eigenvectors:

```

```
//[
//  1.0000  0.      0.      0.      0.
//  0.      1.0000  0.      0.      0.
//  0.      0.      1.0000  0.      0.
//  0.      0.      0.      1.0000  0.
//  0.      0.      0.      0.      1.0000
//];
//difference in eigenvalues: 0.000000e+00
```

4.7.10 magma_ssyevd_gpu - unified memory version, small matrix

```
#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv) {
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    magma_int_t n=1024, n2=n*n;
    float *a, *r; // a, r - nxn matrices
    float *a1; // nxn matrix, copy of a used in magma_ssyevd_gpu
    float *h_work; // workspace
    magma_int_t lwork; // h_work size
    magma_int_t *iwork; // workspace
    magma_int_t liwork; // iwork size
    float *w1, *w2; // w1,w2 - vectors of eigenvalues
    float error, work[1]; // used in difference computations
    magma_int_t ione = 1, i, j, info;
    float mione = -1.0;
    magma_int_t incr = 1;
    cudaMallocManaged(&w1,n*sizeof(float)); //unified memory for
    cudaMallocManaged(&w2,n*sizeof(float)); //eigenvalues
    cudaMallocManaged(&a,n2*sizeof(float)); //unif. memory for a
    cudaMallocManaged(&r,n2*sizeof(float)); //unif. memory for r
    cudaMallocManaged(&a1,n2*sizeof(float)); //unif.memory for a1
    // Query for workspace sizes
    float aux_work[1];
    magma_int_t aux_iwork[1];
    magma_ssyevd_gpu(MagmaVec, MagmaLower, n, a1, n, w1, r, n, aux_work,
                    -1, aux_iwork, -1, &info );

    lwork = (magma_int_t) aux_work[0];
    liwork = aux_iwork[0]; // unified memory for workspace:
    cudaMallocManaged(&iwork,liwork*sizeof(magma_int_t));
    cudaMallocManaged(&h_work,lwork*sizeof(float));
    // define a, r // [1 0 0 0 0 ...]
    for(i=0;i<n;i++){ // [0 2 0 0 0 ...]
        a[i*n+i]=(float)(i+1); // a = [0 0 3 0 0 ...]
        r[i*n+i]=(float)(i+1); // [0 0 0 4 0 ...]
```

```

}
printf("upper left corner of a:\n"); // [0 0 0 0 5 ...]
magma_sprint(5,5,a,n); // .....
magma_ssetmatrix( n, n, a, n, a1, n, queue); // copy a -> a1
// compute the eigenvalues and eigenvectors for a symmetric,
// real nxn matrix; Magma version

magma_ssyevevd_gpu(MagmaVec,MagmaLower,n,a1,n,w1,r,n,h_work,
                  lwork,iwork,liwork,&info);

printf("first 5 eigenvalues of a:\n");
for(j=0;j<5;j++)
    printf("%f\n",w1[j]); // print first eigenvalues
printf("left upper corner of the matrix of eigenvectors:\n");
magma_sprint(5,5,a1,n); // part of the matrix of eigenvectors
// Lapack version
lapackf77_ssyevevd("V","L",&n,a,&n,w2,h_work,&lwork,iwork,
                  &liwork,&info);
// difference in eigenvalues
blasf77_saxpy( &n, &mione, w1, &incr, w2, &incr);
error = lapackf77_slange( "M", &n, &ione, w2, &n, work );
printf("difference in eigenvalues: %e\n",error);
magma_free(w1); // free memory
magma_free(w2); // free memory
magma_free(a); // free memory
magma_free(r); // free memory
magma_free(h_work); // free memory
magma_free(a1); // free memory
magma_queue_destroy(queue); // destroy queue
magma_finalize(); // finalize Magma
return EXIT_SUCCESS;
}

//upper left corner of a:
//[
//  1.0000  0.      0.      0.      0.
//  0.      2.0000  0.      0.      0.
//  0.      0.      3.0000  0.      0.
//  0.      0.      0.      4.0000  0.
//  0.      0.      0.      0.      5.0000
//];

//first 5 eigenvalues of a:
//1.000000
//2.000000
//3.000000
//4.000000
//5.000000

//left upper corner of the matrix of eigenvectors:

```

```
//[
//  1.0000  0.      0.      0.      0.
//  0.      1.0000  0.      0.      0.
//  0.      0.      1.0000  0.      0.
//  0.      0.      0.      1.0000  0.
//  0.      0.      0.      0.      1.0000
//];
//difference in eigenvalues: 0.000000e+00
```

4.7.11 magma_dsyevd_gpu - compute the eigenvalues and optionally eigenvectors of a symmetric real matrix in double precision, GPU interface, small matrix

This function computes in double precision all eigenvalues and, optionally, eigenvectors of a real symmetric matrix A defined on the device. The first parameter can take the values `MagmaVec` or `MagmaNoVec` and answers the question whether the eigenvectors are desired. If the eigenvectors are desired, it uses a divide and conquer algorithm. The symmetric matrix A can be stored in lower (`MagmaLower`) or upper (`MagmaUpper`) mode. If the eigenvectors are desired, then on exit A contains orthonormal eigenvectors. The eigenvalues are stored in an array `w`. See [magma-X.Y.Z/src/dsyevd_gpu.cpp](#) for more details.

```
#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv) {
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    magma_int_t n=1024, n2=n*n;
    double *a, *r; // a, r - nxn matrices on the host
    double *d_r; // nxn matrix on the device
    double *h_work; // workspace
    magma_int_t lwork; // h_work size
    magma_int_t *iwork; // workspace
    magma_int_t liwork; // iwork size
    double *w1, *w2; // w1,w2 - vectors of eigenvalues
    double error, work[1]; // used in difference computations
    magma_int_t ione = 1, i, j, info;
    double mione = -1.0;
    magma_int_t incr = 1;
    magma_dmalloc_cpu(&w1,n); // host memory for real
    magma_dmalloc_cpu(&w2,n); // eigenvalues
    magma_dmalloc_cpu(&a,n2); // host memory for a
    magma_dmalloc_cpu(&r,n2); // host memory for r
    magma_dmalloc(&d_r,n2); // device memory for d_r
    // Query for workspace sizes
    double aux_work[1];
```

```

magma_int_t aux_iwork[1];
magma_dsyevd_gpu(MagmaVec, MagmaLower, n, d_r, n, w1, r, n, aux_work,
                -1, aux_iwork, -1, &info );           // workspace query
lwork  = (magma_int_t) aux_work[0];
liwork = aux_iwork[0];
iwork=(magma_int_t*)malloc(liwork*sizeof(magma_int_t));
magma_dmalloc_cpu(&h_work, lwork);           // memory for workspace
// define a, r                                //      [1 0 0 0 0 ...]
for(i=0; i<n; i++){                          //      [0 2 0 0 0 ...]
    a[i*n+i]=(double)(i+1);                  // a = [0 0 3 0 0 ...]
    r[i*n+i]=(double)(i+1);                  //      [0 0 0 4 0 ...]
}                                              //      [0 0 0 0 5 ...]
printf("upper left corner of a:\n");          // .....
magma_dprint(5,5,a,n);                      // print part of a
magma_dsetmatrix( n, n, a, n, d_r, n, queue); // copy a -> d_r
// compute the eigenvalues and eigenvectors for a symmetric,
// real nxn matrix; Magma version

magma_dsyevd_gpu(MagmaVec, MagmaLower, n, d_r, n, w1, r, n, h_work,
                lwork, iwork, liwork, &info);

printf("first 5 eigenvalues of a:\n");
for(j=0; j<5; j++){
    printf("%f\n", w1[j]);                    // print first eigenvalues
printf("left upper corner of the matrix of eigenvectors:\n");
magma_dgetmatrix( n, n, d_r, n, r, n, queue); // copy d_r -> r
magma_dprint(5,5,r,n); // part of the matrix of eigenvectors
// Lapack version
lapackf77_dsyevd("V", "L", &n, a, &n, w2, h_work, &lwork, iwork,
                &liwork, &info);

// difference in eigenvalues
blasf77_daxpy( &n, &mione, w1, &incr, w2, &incr);
error = lapackf77_dlange( "M", &n, &ione, w2, &n, work );
printf("difference in eigenvalues: %e\n", error);
free(w1);                                     // free host memory
free(w2);                                     // free host memory
free(a);                                     // free host memory
free(r);                                     // free host memory
free(h_work);                                // free host memory
magma_free(d_r);                             // free device memory
magma_queue_destroy(queue);                  // destroy queue
magma_finalize();                            // finalize Magma
return EXIT_SUCCESS;
}
//upper left corner of a:
//[
//  1.0000   0.      0.      0.      0.
//  0.      2.0000   0.      0.      0.
//  0.      0.      3.0000   0.      0.
//  0.      0.      0.      4.0000   0.
//  0.      0.      0.      0.      5.0000
//];

```

```

//first 5 eigenvalues of a:
//1.000000
//2.000000
//3.000000
//4.000000
//5.000000

//left upper corner of the matrix of eigenvectors:
//[
//  1.0000  0.      0.      0.      0.
//  0.      1.0000  0.      0.      0.
//  0.      0.      1.0000  0.      0.
//  0.      0.      0.      1.0000  0.
//  0.      0.      0.      0.      1.0000
//];
//difference in eigenvalues: 0.000000e+00

```

4.7.12 magma_dsyevd_gpu - unified memory version, small matrix

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv) {
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    magma_int_t n=1024, n2=n*n;
    double *a, *r; // a, r - nxn matrices
    double *a1; // nxn matrix, copy of a used in magma_dsyevd_gpu
    double *h_work; // workspace
    magma_int_t lwork; // h_work size
    magma_int_t *iwork; // workspace
    magma_int_t liwork; // iwork size
    double *w1, *w2; // w1,w2 - vectors of eigenvalues
    double error, work[1]; // used in difference computations
    magma_int_t ione = 1, i, j, info;
    double mione = -1.0;
    magma_int_t incr = 1;
    cudaMallocManaged(&w1,n*sizeof(double)); //unified memory for
    cudaMallocManaged(&w2,n*sizeof(double)); //eigenvalues
    cudaMallocManaged(&a,n2*sizeof(double)); //unif. memory for a
    cudaMallocManaged(&r,n2*sizeof(double)); //unif. memory for r
    cudaMallocManaged(&a1,n2*sizeof(double)); //uni.memory for a1
    // Query for workspace sizes
    double aux_work[1];
    magma_int_t aux_iwork[1];

```



```

magma_dsyevd_gpu(MagmaVec,MagmaLower,n,a1,n,w1,r,n,aux_work,
                -1, aux_iwork,-1,&info );

lwork  = (magma_int_t) aux_work[0];
liwork = aux_iwork[0];           // unified memory for workspace:
cudaMallocManaged(&iwork,liwork*sizeof(magma_int_t));
cudaMallocManaged(&h_work,lwork*sizeof(double));
// define a, r                                //      [1 0 0 0 0 ...]
for(i=0;i<n;i++){                             //      [0 2 0 0 0 ...]
    a[i*n+i]=(double)(i+1);                     // a = [0 0 3 0 0 ...]
    r[i*n+i]=(double)(i+1);                     //      [0 0 0 4 0 ...]
}                                                  //      [0 0 0 0 5 ...]
printf("upper left corner of a:\n"); //      .....
magma_dprint(5,5,a,n);                          // print part of a
magma_dsetmatrix( n, n, a, n, a1, n,queue); // copy a -> a1
// compute the eigenvalues and eigenvectors for a symmetric,
// real nxn matrix; Magma version

magma_dsyevd_gpu(MagmaVec,MagmaLower,n,a1,n,w1,r,n,h_work,
                lwork,iwork,liwork,&info);

printf("first 5 eigenvalues of a:\n");
for(j=0;j<5;j++){
    printf("%f\n",w1[j]);                // print first eigenvalues
printf("left upper corner of the matrix of eigenvectors:\n");
magma_dprint(5,5,a1,n); // part of the matrix of eigenvectors
// Lapack version
lapackf77_dsyevd("V","L",&n,a,&n,w2,h_work,&lwork,iwork,
                &liwork,&info);

// difference in eigenvalues
blasf77_daxpy( &n, &mione, w1, &incr, w2, &incr);
error = lapackf77_dlange( "M", &n, &ione, w2, &n, work );
printf("difference in eigenvalues: %e\n",error);
magma_free(w1);                                // free memory
magma_free(w2);                                // free memory
magma_free(a);                                // free memory
magma_free(r);                                // free memory
magma_free(h_work);                           // free memory
magma_free(a1);                               // free memory
magma_queue_destroy(queue);                   // destroy queue
magma_finalize();                             // finalize Magma
return EXIT_SUCCESS;
}

//upper left corner of a:
//[
//  1.0000  0.      0.      0.      0.
//  0.      2.0000  0.      0.      0.
//  0.      0.      3.0000  0.      0.
//  0.      0.      0.      4.0000  0.
//  0.      0.      0.      0.      5.0000
//];
//first 5 eigenvalues of a:

```

```

//1.000000
//2.000000
//3.000000
//4.000000
//5.000000

//left upper corner of the matrix of eigenvectors:
//[
//  1.0000  0.      0.      0.      0.
//  0.      1.0000  0.      0.      0.
//  0.      0.      1.0000  0.      0.
//  0.      0.      0.      1.0000  0.
//  0.      0.      0.      0.      1.0000
//];
//difference in eigenvalues: 0.000000e+00

```

4.7.13 magma_ssyevegpu - compute the eigenvalues and optionally eigenvectors of a symmetric real matrix in single precision, GPU interface, big matrix

This function computes in single precision all eigenvalues and, optionally, eigenvectors of a real symmetric matrix A defined on the device. The first parameter can take the values `MagmaVec` or `MagmaNoVec` and answers the question whether the eigenvectors are desired. If the eigenvectors are desired, it uses a divide and conquer algorithm. The symmetric matrix A can be stored in lower (`MagmaLower`) or upper (`MagmaUpper`) mode. If the eigenvectors are desired, then on exit A contains orthonormal eigenvectors. The eigenvalues are stored in an array `w`. See [magma-X.Y.Z/src/ssyevegpu.cpp](#) for more details.

```

#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv) {
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    double gpu_time, cpu_time;
    magma_int_t n=8192, n2=n*n;
    float *a, *r; // a, r - nxn matrices on the host
    float *d_r; // nxn matrix on the device
    float *h_work; // workspace
    magma_int_t lwork; // h_work size
    magma_int_t *iwork; // workspace
    magma_int_t liwork; // iwork size
    float *w1, *w2; // w1,w2 - vectors of eigenvalues
    float error, work[1]; // used in difference computations
    magma_int_t ione = 1, info;

```

```

float mione = -1.0f;
magma_int_t incr = 1;
magma_int_t ISEED[4] = {0,0,0,1};
magma_smallocc_cpu(&w1,n); // host memory for real
magma_smallocc_cpu(&w2,n); // eigenvalues
magma_smallocc_cpu(&a,n2); // host memory for a
magma_smallocc_cpu(&r,n2); // host memory for r
magma_smallocc(&d_r,n2); // device memory for d_r
// Query for workspace sizes
float aux_work[1];
magma_int_t aux_iwork[1];
magma_ssyevd_gpu(MagmaVec,MagmaLower,n,d_r,n,w1,r,n,aux_work,
                -1,aux_iwork,-1,&info );

lwork = (magma_int_t) aux_work[0];
liwork = aux_iwork[0];
iwork=(magma_int_t*)malloc(liwork*sizeof(magma_int_t));
magma_smallocc_cpu(&h_work,lwork); // memory for workspace
// Randomize the matrix a and copy a -> r
lapackf77_slarnv(&mione,ISEED,&n2,a);
lapackf77_slacpy(MagmaFullStr,&n,&n,a,&n,r,&n);
magma_ssetmatrix( n, n, a, n, d_r,n,queue); // copy a -> d_r
// compute the eigenvalues and eigenvectors for a symmetric,
// real nxn matrix; Magma version
gpu_time = magma_sync_wtime(NULL);

magma_ssyevd_gpu(MagmaVec,MagmaLower,n,d_r,n,w1,r,n,h_work,
                lwork,iwork,liwork,&info);

gpu_time = magma_sync_wtime(NULL)-gpu_time;
printf("ssyevd gpu time: %7.5f sec.\n",gpu_time); // Magma
// Lapack version // time
cpu_time=magma_wtime();
lapackf77_ssyevd("V","L",&n,a,&n,w2,h_work,&lwork,iwork,
                &liwork,&info);

cpu_time=magma_wtime()-cpu_time;
printf("ssyevd cpu time: %7.5f sec.\n",cpu_time); // Lapack
// difference in eigenvalues // time
blasf77_saxpy( &n, &mione, w1, &incr, w2, &incr);
error = lapackf77_slange( "M", &n, &mione, w2, &n, work );
printf("difference in eigenvalues: %e\n",error);
free(w1); // free host memory
free(w2); // free host memory
free(a); // free host memory
free(r); // free host memory
free(h_work); // free host memory
magma_free(d_r); // free device memory
magma_queue_destroy(queue); // destroy queue
magma_finalize(); // finalize Magma
return EXIT_SUCCESS;
}
//ssyevd gpu time: 5.11538 sec.
//ssyevd cpu time: 49.32742 sec.

```

```
//difference in eigenvalues: 9.765625e-04
```

4.7.14 magma_ssyevd_gpu - unified memory version, big matrix

```
#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv) {
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    double gpu_time, cpu_time;
    magma_int_t n=8192, n2=n*n;
    float *a, *r; // a, r - nxn matrices
    float *a1; // nxn matrix, copy of a used in magma_ssyevd_gpu
    float *h_work; // workspace
    magma_int_t lwork; // h_work size
    magma_int_t *iwork; // workspace
    magma_int_t liwork; // iwork size
    float *w1, *w2; // w1,w2 - vectors of eigenvalues
    float error, work[1]; // used in difference computations
    magma_int_t ione = 1, info;
    float mione = -1.0;
    magma_int_t incr = 1;
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    cudaMallocManaged(&w1,n*sizeof(float)); //unified memory for
    cudaMallocManaged(&w2,n*sizeof(float)); //eigenvalues
    cudaMallocManaged(&a,n2*sizeof(float)); //unif. memory for a
    cudaMallocManaged(&r,n2*sizeof(float)); //unif. memory for r
    cudaMallocManaged(&a1,n2*sizeof(float)); //unif.memory for a1
    // Query for workspace sizes
    float aux_work[1];
    magma_int_t aux_iwork[1];
    magma_ssyevd_gpu(MagmaVec, MagmaLower, n, a1, n, w1, r, n, aux_work,
                    -1, aux_iwork, -1, &info );

    lwork = (magma_int_t) aux_work[0];
    liwork = aux_iwork[0]; // unified memory for workspace:
    cudaMallocManaged(&iwork,liwork*sizeof(magma_int_t));
    cudaMallocManaged(&h_work,lwork*sizeof(float));
    // Randomize the matrix a and copy a -> r
    lapackf77_slarnv(&ione, ISEED, &n2, a);
    lapackf77_slacpy(MagmaFullStr, &n, &n, a, &n, r, &n);
    magma_ssetmatrix( n, n, a, n, a1, n, queue); // copy a -> a1
    // compute the eigenvalues and eigenvectors for a symmetric,
    // real nxn matrix; Magma version
    gpu_time = magma_sync_wtime(NULL);
```

```

    magma_ssyevd_gpu(MagmaVec,MagmaLower,n,a1,n,w1,r,n,h_work,
                    lwork,iwork,liwork,&info);

    gpu_time = magma_sync_wtime(NULL)-gpu_time;
    printf("ssyevd_gpu gpu time: %7.5f sec.\n",gpu_time); //Magma
// Lapack version                                     time
    cpu_time=magma_wtime();
    lapackf77_ssyevd("V","L",&n,a,&n,w2,h_work,&lwork,iwork,
                    &liwork,&info);

    cpu_time=magma_wtime()-cpu_time;
    printf("ssyevd cpu time: %7.5f sec.\n",cpu_time); // Lapack
// difference in eigenvalues                          // time
    blasf77_saxpy( &n, &mione, w1, &incr, w2, &incr);
    error = lapackf77_slange( "M", &n, &iione, w2, &n, work );
    printf("difference in eigenvalues: %e\n",error);
    magma_free(w1); // free memory
    magma_free(w2); // free memory
    magma_free(a); // free memory
    magma_free(r); // free memory
    magma_free(h_work); // free memory
    magma_free(a1); // free memory
    magma_queue_destroy(queue); // destroy queue
    magma_finalize(); // finalize Magma
    return EXIT_SUCCESS;
}
//ssyevd_gpu gpu time: 5.29559 sec.
//ssyevd cpu time: 51.07547 sec.
//difference in eigenvalues: 9.765625e-04

```

4.7.15 magma_dsyevd_gpu - compute the eigenvalues and optionally eigenvectors of a symmetric real matrix in double precision, GPU interface, big matrix

This function computes in double precision all eigenvalues and, optionally, eigenvectors of a real symmetric matrix A defined on the device. The first parameter can take the values `MagmaVec` or `MagmaNoVec` and answers the question whether the eigenvectors are desired. If the eigenvectors are desired, it uses a divide and conquer algorithm. The symmetric matrix A can be stored in lower (`MagmaLower`) or upper (`MagmaUpper`) mode. If the eigenvectors are desired, then on exit A contains orthonormal eigenvectors. The eigenvalues are stored in an array `w`. See [magma-X.Y.Z/src/dsyevd_gpu.cpp](#) for more details.

```

#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv) {
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;

```

```

magma_int_t dev=0;
magma_queue_create(dev,&queue);
double gpu_time, cpu_time;
magma_int_t n=8192, n2=n*n;
double *a, *r; // a, r - nxn matrices on the host
double *d_r; // nxn matrix on the device
double *h_work; // workspace
magma_int_t lwork; // h_work size
magma_int_t *iwork; // workspace
magma_int_t liwork; // iwork size
double *w1, *w2; // w1,w2 - vectors of eigenvalues
double error, work[1]; // used in difference computations
magma_int_t ione = 1, info;
double mione = -1.0;
magma_int_t incr = 1;
magma_int_t ISEED[4] = {0,0,0,1}; // seed
magma_dmalloc_cpu(&w1,n); // host memory for real
magma_dmalloc_cpu(&w2,n); // eigenvalues
magma_dmalloc_cpu(&a,n2); // host memory for a
magma_dmalloc_cpu(&r,n2); // host memory for r
magma_dmalloc(&d_r,n2); // device memory for d_r
// Query for workspace sizes
double aux_work[1];
magma_int_t aux_iwork[1];
magma_dsyevd_gpu(MagmaVec,MagmaLower,n,d_r,n,w1,r,n,aux_work,
-1,aux_iwork,-1,&info );

lwork = (magma_int_t) aux_work[0];
liwork = aux_iwork[0];
iwork=(magma_int_t*) malloc(liwork*sizeof(magma_int_t));
magma_dmalloc_cpu(&h_work,lwork); // memory for workspace
// Randomize the matrix a and copy a -> r
lapackf77_dlarnv(&ione,ISEED,&n2,a);
lapackf77_dlacpy(MagmaFullStr,&n,&n,a,&n,r,&n);
magma_dsetmatrix( n, n, a, n, d_r,n,queue); // copy a -> d_r
// compute the eigenvalues and eigenvectors for a symmetric,
// real nxn matrix; Magma version
gpu_time = magma_sync_wtime(NULL);

magma_dsyevd_gpu(MagmaVec,MagmaLower,n,d_r,n,w1,r,n,h_work,
lwork,iwork,liwork,&info);

gpu_time = magma_sync_wtime(NULL)-gpu_time;
printf("dsyevd_gpu gpu time: %7.5f sec.\n",gpu_time); //Magma
// Lapack version time
cpu_time=magma_wtime();
lapackf77_dsyevd("V","L",&n,a,&n,w2,h_work,&lwork,iwork,
&liwork,&info);

cpu_time=magma_wtime()-cpu_time;
printf("dsyevd cpu time: %7.5f sec.\n",cpu_time); // Lapack
// difference in eigenvalues // time
blasf77_daxpy( &n, &mione, w1, &incr, w2, &incr);
error = lapackf77_dlange( "M", &n, &ione, w2, &n, work );

```

```

printf("difference in eigenvalues: %e\n",error);
free(w1); // free host memory
free(w2); // free host memory
free(a); // free host memory
free(r); // free host memory
free(h_work); // free host memory
magma_free(d_r); // free device memory
magma_queue_destroy(queue); // destroy queue
magma_finalize(); // finalize Magma
return EXIT_SUCCESS;
}
//dsyevd_gpu gpu time: 16.50546 sec.
//dsyevd cpu time: 91.54085 sec.
//difference in eigenvalues: 1.364242e-11

```

4.7.16 magma_dsyevd_gpu - unified memory version, big matrix

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
int main( int argc, char** argv) {
    magma_init(); // initialize Magma
    magma_queue_t queue=NULL;
    magma_int_t dev=0;
    magma_queue_create(dev,&queue);
    double gpu_time, cpu_time;
    magma_int_t n=8192, n2=n*n;
    double *a, *r; // a, r - nxn matrices
    double *a1; // nxn matrix, copy of a used in magma_dsyevd_gpu
    double *h_work; // workspace
    magma_int_t lwork; // h_work size
    magma_int_t *iwork; // workspace
    magma_int_t liwork; // iwork size
    double *w1, *w2; // w1,w2 - vectors of eigenvalues
    double error, work[1]; // used in difference computations
    magma_int_t ione = 1, info;
    double mione = -1.0;
    magma_int_t incr = 1;
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    cudaMallocManaged(&w1,n*sizeof(double)); //unified memory for
    cudaMallocManaged(&w2,n*sizeof(double)); //eigenvalues
    cudaMallocManaged(&a,n2*sizeof(double)); //unif. memory for a
    cudaMallocManaged(&r,n2*sizeof(double)); //unif. memory for r
    cudaMallocManaged(&a1,n2*sizeof(double)); //uni.memory for a1
    // Query for workspace sizes
    double aux_work[1];
    magma_int_t aux_iwork[1];
    magma_dsyevd_gpu(MagmaVec,MagmaLower,n,a1,n,w1,r,n,aux_work,

```

```

                                -1,aux_iwork,-1,&info );
lwork  = (magma_int_t) aux_work[0];
liwork = aux_iwork[0];          // unified memory for workspace:
cudaMallocManaged(&iwork,liwork*sizeof(magma_int_t));
cudaMallocManaged(&h_work,lwork*sizeof(double));
// Randomize the matrix a and copy a -> r
lapackf77_dlarnv(&ione,ISEED,&n2,a);
lapackf77_dlacpy(MagmaFullStr,&n,&n,a,&n,r,&n);
magma_dsetmatrix( n, n, a, n, a1,n,queue); // copy a -> a1
// compute the eigenvalues and eigenvectors for a symmetric,
// real nxn matrix; Magma version
gpu_time = magma_sync_wtime(NULL);

magma_dsyevd_gpu(MagmaVec,MagmaLower,n,a1,n,w1,r,n,h_work,
                lwork,iwork,liwork,&info);

gpu_time = magma_sync_wtime(NULL)-gpu_time;
printf("dsyevd_gpu gpu time: %7.5f sec.\n",gpu_time); //Magma
// Lapack version                                     time
cpu_time=magma_wtime();
lapackf77_dsyevd("V","L",&n,a,&n,w2,h_work,&lwork,iwork,
                &liwork,&info);

cpu_time=magma_wtime()-cpu_time;
printf("dsyevd cpu time: %7.5f sec.\n",cpu_time); // Lapack
// difference in eigenvalues                          // time
blasf77_daxpy( &n, &mione, w1, &incr, w2, &incr);
error = lapackf77_dlange( "M", &n, &ione, w2, &n, work );
printf("difference in eigenvalues: %e\n",error);
magma_free(w1);                                     // free memory
magma_free(w2);                                     // free memory
magma_free(a);                                       // free memory
magma_free(r);                                       // free memory
magma_free(h_work);                                  // free memory
magma_free(a1);                                      // free memory
magma_queue_destroy(queue);                          // destroy queue
magma_finalize();                                    // finalize Magma
return EXIT_SUCCESS;
}
//dsyevd_gpu gpu time: 16.55437 sec.
//dsyevd cpu time: 95.21645 sec.
//difference in eigenvalues: 1.364242e-11

```

4.8 Singular value decomposition

4.8.1 magma_sgesvd - compute the singular value decomposition of a general real matrix in single precision, CPU interface

This function computes in single precision the singular value decomposition of an $m \times n$ matrix defined on the host:

$$A = u \sigma v^T,$$

where σ is an $m \times n$ matrix which is zero except for its $\min(m, n)$ diagonal elements (singular values), u is an $m \times m$ orthogonal matrix and v is an $n \times n$ orthogonal matrix. The first $\min(m, n)$ columns of u and v are the left and right singular vectors of A . The first argument can take the following values:

MagmaAllVec - all m columns of u are returned in an array u ;
MagmaSomeVec - the first $\min(m, n)$ columns of u (the left singular vectors) are returned in the array u ;
MagmaOverwriteVec - the first $\min(m, n)$ columns of u are overwritten on the array A ;
MagmaNoVec - no left singular vectors are computed.

Similarly the second argument can take the following values:

MagmaAllVec - all n rows of v^T are returned in an array vt ;
MagmaSomeVec - the first $\min(m, n)$ rows of v^T (the right singular vectors) are returned in the array vt ;
MagmaOverwriteVec - the first $\min(m, n)$ rows of v^T are overwritten on the array A ;
MagmaNoVec - no right singular vectors are computed.

The singular values are stored in an array s .

See [magma-X.Y.Z/src/sgesvd.cpp](#) for more details.

```
#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b) (((a)<(b))?(a):(b))
int main( int argc, char** argv)
{
    magma_init(); // initialize Magma
    real_Double_t gpu_time, cpu_time;
    // Matrix size
    magma_int_t m=8192, n=8192, n2=m*n, min_mn=min(m,n);
    float *a, *r; // a,r - mxn matrices
    float *u, *vt; // u - mxm matrix, vt - nxn matrix on the host
    float *s1, *s2; // vectors of singular values
    magma_int_t info;
    magma_int_t ione = 1;
    float work[1], error = 1.; // used in difference computations
    float mone = -1.0, *h_work; // h_work - workspace
    magma_int_t lwork; // workspace size
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    // Allocate host memory
    magma_smallocc_cpu(&a, n2); // host memory for a
    magma_smallocc_cpu(&vt, n*n); // host memory for vt
    magma_smallocc_cpu(&u, m*m); // host memory for u
    magma_smallocc_cpu(&s1, min_mn); // host memory for s1
```

```

    magma_smallocc_cpu(&s2,min_mn);          // host memory for s2
    magma_smallocc_pinned(&r,n2);           // host memory for r
    magma_int_t nb = magma_get_sgesvd_nb(m,n); //optim.block size
    lwork=min_mn*min_mn+2*min_mn+2*min_mn*nb;
    magma_smallocc_pinned(&h_work,lwork); // host mem. for h_work
// Randomize the matrix a
    lapackf77_slarnv(&ione,ISEED,&n2,a);
    lapackf77_slacpy(MagmaFullStr,&m,&n,a,&m,r,&m); //a->r
// MAGMA
    gpu_time = magma_wtime();
// compute the singular value decomposition of r (copy of a)
// and optionally the left and right singular vectors:
// r = u*sigma*vt; the diagonal elements of sigma (s1 array)
// are the singular values of a in descending order
// the first min(m,n) columns of u contain the left sing. vec.
// the first min(m,n) columns of vt contain the right sing.vec.

    magma_sgesvd(MagmaNoVec,MagmaNoVec,m,n,r,m,s1,u,m,vt,n,h_work,
                lwork,&info );

    gpu_time = magma_wtime() - gpu_time;
    printf("sgesvd gpu time:  %7.5f\n", gpu_time); // Magma time
// LAPACK
    cpu_time = magma_wtime();
    lapackf77_sgesvd("N","N",&m,&n,a,&m,s2,u,&m,vt,&n,h_work,
                    &lwork,&info);

    cpu_time = magma_wtime() - cpu_time;
    printf("sgesvd cpu time:  %7.5f\n", cpu_time); // Lapack time
// difference
    error=lapackf77_slange("f",&min_mn,&ione,s1,&min_mn,work);
    blasf77_saxpy(&min_mn,&mone,s1,&ione,s2,&ione);
    error=lapackf77_slange("f",&min_mn,&ione,s2,&min_mn,work);
//error;
    printf("difference:  %e\n", error );// difference in singul.
// values

// Free memory
    free(a);          // free host memory
    free(vt);         // free host memory
    free(s1);         // free host memory
    free(s2);         // free host memory
    free(u);          // free host memory
    magma_free_pinned(h_work); // free host memory
    magma_free_pinned(r);    // free host memory
    magma_finalize( );      // finalize Magma
    return EXIT_SUCCESS;
}
//sgesvd gpu time:  15.00651
//sgesvd cpu time:  115.81860
//difference:  5.943540e-07

```

4.8.2 magma_sgesvd - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b) (((a)<(b))?(a):(b))
int main( int argc, char** argv)
{
    magma_init(); // initialize Magma
    real_Double_t gpu_time, cpu_time;
    // Matrix size
    magma_int_t m=8192, n=8192, n2=m*n, min_mn=min(m,n);
    float *a, *r; // a,r - mxn matrices
    float *u, *vt; // u - mxm matrix, vt - nxn matrix
    float *s1, *s2; // vectors of singular values
    magma_int_t info;
    magma_int_t ione = 1;
    float work[1], error = 1.; //used in difference computations
    float mone = -1.0, *h_work; // h_work - workspace
    magma_int_t lwork; // workspace size
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    cudaMallocManaged(&a,n2*sizeof(float)); //unif. memory for a
    cudaMallocManaged(&vt,n*n*sizeof(float)); //uni.memory for vt
    cudaMallocManaged(&u,m*m*sizeof(float)); //unif. memory for u
    cudaMallocManaged(&s1,min_mn*sizeof(float)); //uni.mem.for s1
    cudaMallocManaged(&s2,min_mn*sizeof(float)); //uni.mem.for s2
    cudaMallocManaged(&r,n2*sizeof(float)); //unif. memory for r
    magma_int_t nb = magma_get_sgesvd_nb(m,n); //optim.block size
    lwork=min_mn*min_mn+2*min_mn+2*min_mn*nb;
    cudaMallocManaged(&h_work,lwork*sizeof(float)); //m.f.h_work
    // Randomize the matrix a
    lapackf77_slarnv(&ione, ISEED, &n2, a);
    lapackf77_slacpy(MagmaFullStr, &m, &n, a, &m, r, &m); //a->r
    // MAGMA
    gpu_time = magma_wtime();
    // compute the singular value decomposition of r (copy of a)
    // and optionally the left and right singular vectors:
    // r = u*sigma*vt; the diagonal elements of sigma (s1 array)
    // are the singular values of a in descending order
    // the first min(m,n) columns of u contain the left sing. vec.
    // the first min(m,n) columns of vt contain the right sing.vec.

    magma_sgesvd(MagmaNoVec, MagmaNoVec, m, n, r, m, s1, u, m, vt, n, h_work,
                lwork, &info );

    gpu_time = magma_wtime() - gpu_time;
    printf("sgesvd gpu time: %7.5f\n", gpu_time); // Magma time
    // LAPACK
    cpu_time = magma_wtime();
    lapackf77_sgesvd("N", "N", &m, &n, a, &m, s2, u, &m, vt, &n, h_work,

```

```

                                                                    &lwork,&info);
    cpu_time = magma_wtime() - cpu_time;
    printf("sgesvd cpu time:  %7.5f\n", cpu_time); // Lapack time
// difference
    error=lapackf77_slange("f",&min_mn,&ione,s1,&min_mn,work);
    blasf77_saxpy(&min_mn,&mone,s1,&ione,s2,&ione);
    error=lapackf77_slange("f",&min_mn,&ione,s2,&min_mn,work);
                                                                    // error;
    printf("difference:  %e\n", error ); // difference in singul.
                                                                    // values
// Free memory
    magma_free(a); // free memory
    magma_free(vt); // free memory
    magma_free(s1); // free memory
    magma_free(s2); // free memory
    magma_free(u); // free memory
    magma_free(h_work); // free memory
    magma_free(r); // free memory
    magma_finalize( ); // finalize Magma
    return EXIT_SUCCESS;
}
//sgesvd gpu time:  16.51667
//sgesvd cpu time:  115.20410
//difference:  2.810940e-03

```

4.8.3 magma_dgesvd - compute the singular value decomposition of a general real matrix in double precision, CPU interface

This function computes in double precision the singular value decomposition of an $m \times n$ matrix defined on the host:

$$A = u \sigma v^T,$$

where σ is an $m \times n$ matrix which is zero except for its $\min(m, n)$ diagonal elements (singular values), u is an $m \times m$ orthogonal matrix and v is an $n \times n$ orthogonal matrix. The first $\min(m, n)$ columns of u and v are the left and right singular vectors of A . The first argument can take the following values:

MagmaAllVec - all m columns of u are returned in an array u ;
MagmaSomeVec - the first $\min(m, n)$ columns of u (the left singular vectors) are returned in the array u ;
MagmaOverwriteVec - the first $\min(m, n)$ columns of u are overwritten on the array A ;
MagmaNoVec - no left singular vectors are computed.

Similarly the second argument can take the following values:

MagmaAllVec - all n rows of v^T are returned in an array vt ;
MagmaSomeVec - the first $\min(m, n)$ rows of v^T (the right singular vectors)

are returned in the array *vt*;

MagmaOverwriteVec - the first $\min(m, n)$ rows of v^T are overwritten on the array *A*;

MagmaNoVec - no right singular vectors are computed.

The singular values are stored in an array *s*.

See [magma-X.Y.Z/src/dgesvd.cpp](#) for more details.

```
#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b) (((a)<(b))?(a):(b))
int main( int argc, char** argv)
{
    magma_init(); // initialize Magma
    real_Double_t    gpu_time, cpu_time;
// Matrix size
    magma_int_t m=8192, n=8192, n2=m*n, min_mn=min(m,n);
    double *a, *r; // a,r - mxn matrices
    double *u, *vt; // u - mxm matrix, vt - nxn matrix on the host
    double *s1, *s2; // vectors of singular values
    magma_int_t info;
    magma_int_t ione = 1;
    double work[1], error = 1.; //used in difference computations
    double mone = -1.0, *h_work; // h_work - workspace
    magma_int_t lwork; // workspace size
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
// Allocate host memory
    magma_dmalloc_cpu(&a,n2); // host memory for a
    magma_dmalloc_cpu(&vt,n*n); // host memory for vt
    magma_dmalloc_cpu(&u,m*m); // host memory for u
    magma_dmalloc_cpu(&s1,min_mn); // host memory for s1
    magma_dmalloc_cpu(&s2,min_mn); // host memory for s2
    magma_dmalloc_pinned(&r,n2); // host memory for r
    magma_int_t nb = magma_get_dgesvd_nb(m,n); //optim.block size
    lwork=min_mn*min_mn+2*min_mn+2*min_mn*nb;
    magma_dmalloc_pinned(&h_work,lwork); // host mem. for h_work
// Randomize the matrix a
    lapackf77_dlarnv(&ione,ISEED,&n2,a);
    lapackf77_dlacpy(MagmaFullStr,&m,&n,a,&m,r,&m); //a->r
// MAGMA
    gpu_time = magma_wtime();
// compute the singular value decomposition of r (copy of a)
// and optionally the left and right singular vectors:
// r = u*sigma*vt; the diagonal elements of sigma (s1 array)
// are the singular values of a in descending order
// the first min(m,n) columns of u contain the left sing. vec.
// the first min(m,n) columns of vt contain the right sing.vec.
```

```

    magma_dgesvd(MagmaNoVec,MagmaNoVec,m,n,r,m,s1,u,m,vt,n,h_work,
                lwork,&info );

    gpu_time = magma_wtime() - gpu_time;
    printf("dgesvd gpu time:  %7.5f\n", gpu_time); // Magma time
// LAPACK
    cpu_time = magma_wtime();
    lapackf77_dgesvd("N","N",&m,&n,&a,&m,&s2,&u,&m,&vt,&n,&h_work,
                    &lwork,&info);

    cpu_time = magma_wtime() - cpu_time;
    printf("dgesvd cpu time:  %7.5f\n", cpu_time); // Lapack time
// difference
    error=lapackf77_dlange("f",&min_mn,&ione,s1,&min_mn,work);
    blasf77_daxpy(&min_mn,&mone,s1,&ione,s2,&ione);
    error=lapackf77_dlange("f",&min_mn,&ione,s2,&min_mn,work);
                                // error;
    printf("difference:  %e\n", error );// difference in singul.
                                // values

// Free memory
    free(a);                    // free host memory
    free(vt);                   // free host memory
    free(s1);                   // free host memory
    free(s2);                   // free host memory
    free(u);                    // free host memory
    magma_free_pinned(h_work);  // free host memory
    magma_free_pinned(r);       // free host memory
    magma_finalize( );          // finalize Magma
    return EXIT_SUCCESS;
}
//dgesvd gpu time:  23.05454
//dgesvd cpu time:  228.58973
//difference:  1.526458e-15

```

4.8.4 magma_dgesvd - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b)  (((a)<(b))?(a):(b))
int main( int argc, char** argv)
{
    magma_init();                // initialize Magma
    real_Double_t    gpu_time, cpu_time;
// Matrix size
    magma_int_t m=8192, n=8192, n2=m*n, min_mn=min(m,n);
    double *a, *r;              // a,r - mxn matrices
    double *u, *vt;              // u - mxm matrix, vt - nxn matrix
    double *s1, *s2;             // vectors of singular values

```

```

magma_int_t info;
magma_int_t ione = 1;
double work[1], error = 1.; //used in difference computations
double mone = -1.0, *h_work; // h_work - workspace
magma_int_t lwork; // workspace size
magma_int_t ISEED[4] = {0,0,0,1}; // seed
cudaMallocManaged(&a,n2*sizeof(double)); //unif. memory for a
cudaMallocManaged(&vt,n*n*sizeof(double)); //unif .mem for vt
cudaMallocManaged(&u,m*m*sizeof(double)); //unif.memory for u
cudaMallocManaged(&s1,min_mn*sizeof(double)); //un.mem.for s1
cudaMallocManaged(&s2,min_mn*sizeof(double)); //un.mem.for s2
cudaMallocManaged(&r,n2*sizeof(double)); //unif. memory for r
magma_int_t nb = magma_get_dgesvd_nb(m,n); //optim.block size
lwork=min_mn*min_mn+2*min_mn+2*min_mn*nb;
cudaMallocManaged(&h_work,lwork*sizeof(double)); //m.f.h_work
// Randomize the matrix a
lapackf77_dlarnv(&ione,ISEED,&n2,a);
lapackf77_dlacpy(MagmaFullStr,&m,&n,a,&m,r,&m); //a->r
// MAGMA
gpu_time = magma_wtime();
// compute the singular value decomposition of r (copy of a)
// and optionally the left and right singular vectors:
// r = u*sigma*vt; the diagonal elements of sigma (s1 array)
// are the singular values of a in descending order
// the first min(m,n) columns of u contain the left sing. vec.
// the first min(m,n) columns of vt contain the right sing.vec.

magma_dgesvd(MagmaNoVec,MagmaNoVec,m,n,r,m,s1,u,m,vt,n,h_work,
            lwork,&info );

gpu_time = magma_wtime() - gpu_time;
printf("dgesvd gpu time:  %7.5f\n", gpu_time); // Magma time
// LAPACK
cpu_time = magma_wtime();
lapackf77_dgesvd("N","N",&m,&n,a,&m,s2,u,&m,vt,&n,h_work,
                &lwork,&info);

cpu_time = magma_wtime() - cpu_time;
printf("dgesvd cpu time:  %7.5f\n", cpu_time); // Lapack time
// difference
error=lapackf77_dlange("f",&min_mn,&ione,s1,&min_mn,work);
blasf77_daxpy(&min_mn,&mone,s1,&ione,s2,&ione);
error=lapackf77_dlange("f",&min_mn,&ione,s2,&min_mn,work);
// error;
printf("difference:  %e\n", error ); // difference in singul.
// values

// Free memory
magma_free(a); // free memory
magma_free(vt); // free memory
magma_free(s1); // free memory
magma_free(s2); // free memory
magma_free(u); // free memory
magma_free(h_work); // free memory

```

```

    magma_free(r);                                // free memory
    magma_finalize( );                             // finalize Magma
    return EXIT_SUCCESS;
}
//dgesvd gpu time: 25.20418
//dgesvd cpu time: 231.04632
//difference: 7.219722e-12

```

4.8.5 magma_sgebrd - reduce a real matrix to bidiagonal form by orthogonal transformations in single precision, CPU interface

This function reduces in single precision an $m \times n$ matrix A defined on the host to upper or lower bidiagonal form by orthogonal transformations:

$$Q^T A P = B,$$

where P, Q are orthogonal and B is bidiagonal. If $m \geq n$, B is upper bidiagonal; if $m < n$, B is lower bidiagonal. The obtained diagonal and the super/subdiagonal are written to `diag` and `offdiag` arrays respectively. If $m \geq n$, the elements below the diagonal, with the array `tauq` represent the orthogonal matrix Q as a product of elementary reflectors $H_k = I - \text{tauq}_k \cdot v_k \cdot v_k^T$, and the elements above the first superdiagonal with the array `taup` represent the orthogonal matrix P as a product of elementary reflectors $G_k = I - \text{taup}_k \cdot u_k \cdot u_k^T$. See [magma-X.Y.Z/src/sgebrd.cpp](#) for more details.

```

#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b) (((a)<(b))?(a):(b))
int main( int argc, char** argv){
    magma_init();                                // initialize Magma
    double gpu_time, cpu_time;
    magma_int_t m = 4096, n = 4096, n2=m*n;
    float *a, *r;                                // a,r - mxn matrices on the host
    float *h_work;                                // workspace
    magma_int_t lhwork;                           // size of h_work
    float *taup, *tauq; // arrays describ. elementary reflectors
    float *diag, *offdiag; // bidiagonal form in two arrays
    magma_int_t info, minmn=min(m,n), nb;
    magma_int_t ione = 1;
    magma_int_t ISEED[4] = {0,0,0,1};             // seed
    nb = magma_get_sgebrd_nb(m,n);                 // optimal block size
    magma_smalloc_cpu(&a,m*n);                     // host memory for a
    magma_smalloc_cpu(&tauq,minmn);                 // host memory for tauq
    magma_smalloc_cpu(&taup,minmn);                 // host memory for taup
    magma_smalloc_cpu(&diag,minmn);                 // host memory for diag

```



```

    magma_smalloc_cpu(&offdiag,minmn-1); // host mem. for offdiag
    magma_smalloc_pinned(&r,m*n);          // host memory for r
    lhwork = (m + n)*nb;
    magma_smalloc_pinned(&h_work,lhwork); // host mem. for h_work
// Randomize the matrix a
    lapackf77_slarnv( &ione, ISEED, &n2, a );
    lapackf77_slacpy(MagmaFullStr,&m,&n,a,&m,r,&m); // a->r
// MAGMA
    gpu_time = magma_wtime();
// reduce the matrix r to upper bidiagonal form by orthogonal
// transformations: q^T*r*p, the obtained diagonal and the
// superdiagonal are written to diag and offdiag arrays resp.;
// the elements below the diagonal, represent the orthogonal
// matrix q as a product of elementary reflectors described
// by tauq; elements above the first superdiagonal represent
// the orthogonal matrix p as a product of elementary reflect-
// ors described by taup;

    magma_sgebrd(m,n,r,m,diag,offdiag,tauq,taup,h_work,lhwork,
                &info);

    gpu_time = magma_wtime() - gpu_time;
    printf("sgebrd gpu time: %7.5f sec.\n",gpu_time);
// LAPACK
    cpu_time = magma_wtime();
    lapackf77_sgebrd(&m,&n,a,&m,diag,offdiag,tauq,taup,h_work,
                    &lhwork,&info);

    cpu_time = magma_wtime() - cpu_time;
    printf("sgebrd cpu time: %7.5f sec.\n",cpu_time);
// free memory
    free(a); // free host memory
    free(tauq); // free host memory
    free(taup); // free host memory
    free(diag); // free host memory
    magma_free_pinned(r); // free host memory
    magma_free_pinned(h_work); // free host memory
    magma_finalize( ); // finalize Magma
    return EXIT_SUCCESS;
}
//sgebrd gpu time: 2.28088 sec.
//sgebrd cpu time: 13.83244 sec.

```

4.8.6 magma_sgebrd - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b) (((a)<(b))?(a):(b))

```

```

int main( int argc, char** argv){
    magma_init(); // initialize Magma
    double gpu_time, cpu_time;
    magma_int_t m = 4096, n = 4096, n2=m*n;
    float *a, *r; // a,r - mxn matrices
    float *h_work; // workspace
    magma_int_t lhwork; // size of h_work
    float *taup, *tauq; // arrays describ. elementary reflectors
    float *diag, *offdiag; // bidiagonal form in two arrays
    magma_int_t info, minmn=min(m,n), nb;
    magma_int_t ione = 1;
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    nb = magma_get_sgebrd_nb(m,n); // optimal block size
    cudaMallocManaged(&a,n2*sizeof(float)); //unif. memory for a
    cudaMallocManaged(&r,n2*sizeof(float)); //unif. memory for r
    cudaMallocManaged(&tauq,minmn*sizeof(float)); //mem.for tauq
    cudaMallocManaged(&taup,minmn*sizeof(float)); //mem.for taup
    cudaMallocManaged(&diag,minmn*sizeof(float)); //mem.for diag
    cudaMallocManaged(&offdiag,(minmn-1)*sizeof(float)); //unif.
    lhwork = (m + n)*nb; //memory for offdiag
    cudaMallocManaged(&h_work,lhwork*sizeof(float)); //unif.mem.
    // Randomize the matrix a // for workspace
    lapackf77_slarnv( &ione, ISEED, &n2, a );
    lapackf77_slacpy(MagmaFullStr,&m,&n,a,&m,r,&m); // a->r
    // MAGMA
    gpu_time = magma_wtime();
    // reduce the matrix r to upper bidiagonal form by orthogonal
    // transformations: q^T*r*p, the obtained diagonal and the
    // superdiagonal are written to diag and offdiag arrays resp.;
    // the elements below the diagonal, represent the orthogonal
    // matrix q as a product of elementary reflectors described
    // by tauq; elements above the first superdiagonal represent
    // the orthogonal matrix p as a product of elementary reflect-
    // ors described by taup;

    magma_sgebrd(m,n,r,m,diag,offdiag,tauq,taup,h_work,lhwork,
                &info);

    gpu_time = magma_wtime() - gpu_time;
    printf("sgebrd gpu time: %7.5f sec.\n",gpu_time);
    // LAPACK
    cpu_time = magma_wtime();
    lapackf77_sgebrd(&m,&n,a,&m,diag,offdiag,tauq,taup,h_work,
                    &lhwork,&info);

    cpu_time = magma_wtime() - cpu_time;
    printf("sgebrd cpu time: %7.5f sec.\n",cpu_time);
    // free memory
    magma_free(a); // free memory
    magma_free(tauq); // free memory
    magma_free(taup); // free memory
    magma_free(diag); // free memory
    magma_free(r); // free memory

```

```

    magma_free(h_work); // free memory
    magma_finalize( ); // finalize Magma
    return EXIT_SUCCESS;
}
//sgebrd gpu time: 2.45375 sec.
//sgebrd cpu time: 12.81832 sec.

```

4.8.7 magma_dgebrd - reduce a real matrix to bidiagonal form by orthogonal transformations in double precision, CPU interface

This function reduces in double precision an $m \times n$ matrix A defined on the host to upper or lower bidiagonal form by orthogonal transformations:

$$Q^T A P = B,$$

where P, Q are orthogonal and B is bidiagonal. If $m \geq n$, B is upper bidiagonal; if $m < n$, B is lower bidiagonal. The obtained diagonal and the super/subdiagonal are written to `diag` and `offdiag` arrays respectively. If $m \geq n$, the elements below the diagonal, with the array `tauq` represent the orthogonal matrix Q as a product of elementary reflectors $H_k = I - \tau q_k \cdot v_k \cdot v_k^T$, and the elements above the first superdiagonal with the array `taup` represent the orthogonal matrix P as a product of elementary reflectors $G_k = I - \tau p_k \cdot u_k \cdot u_k^T$. See [magma-X.Y.Z/src/dgebrd.cpp](#) for more details.

```

#include <stdio.h>
#include <cuda.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b) (((a)<(b))?(a):(b))
int main( int argc, char** argv){
    magma_init(); // initialize Magma
    double gpu_time, cpu_time;
    magma_int_t m = 4096, n = 4096, n2=m*n;
    double *a, *r; // a,r - mxn matrices on the host
    double *h_work; // workspace
    magma_int_t lhwork; // size of h_work
    double *taup, *tauq; // arrays describ. elementary reflectors
    double *diag, *offdiag; // bidiagonal form in two arrays
    magma_int_t info, minmn=min(m,n), nb;
    magma_int_t ione = 1;
    magma_int_t ISEED[4] = {0,0,0,1}; // seed
    nb = magma_get_dgebrd_nb(m,n); // optimal block size
    magma_dmalloc_cpu(&a,m*n); // host memory for a
    magma_dmalloc_cpu(&tauq,minmn); // host memory for tauq
    magma_dmalloc_cpu(&taup,minmn); // host memory for taup
    magma_dmalloc_cpu(&diag,minmn); // host memory for diag
    magma_dmalloc_cpu(&offdiag,minmn-1); // host mem. for offdiag

```

```

    magma_dmalloc_pinned(&r,m*n);                // host memory for r
    lhwork = (m + n)*nb;
    magma_dmalloc_pinned(&h_work,lhwork); // host mem. for h_work
// Randomize the matrix a
    lapackf77_dlarnv( &ione, ISEED, &n2, a );
    lapackf77_dlacpy(MagmaFullStr,&m,&n,a,&m,r,&m); // a->r
// MAGMA
    gpu_time = magma_wtime();
// reduce the matrix r to upper bidiagonal form by orthogonal
// transformations: q^T*r*p, the obtained diagonal and the
// superdiagonal are written to diag and offdiag arrays resp.;
// the elements below the diagonal, represent the orthogonal
// matrix q as a product of elementary reflectors described
// by tauq; elements above the first superdiagonal represent
// the orthogonal matrix p as a product of elementary reflect-
// ors described by taup;

    magma_dgebrd(m,n,r,m,diag,offdiag,tauq,taup,h_work,lhwork,
                &info);

    gpu_time = magma_wtime() - gpu_time;
    printf("dgebrd gpu time: %7.5f sec.\n",gpu_time);
// LAPACK
    cpu_time = magma_wtime();
    lapackf77_dgebrd(&m,&n,a,&m,diag,offdiag,tauq,taup,h_work,
                    &lhwork,&info);

    cpu_time = magma_wtime() - cpu_time;
    printf("dgebrd cpu time: %7.5f sec.\n",cpu_time);
// free memory
    free(a);                // free host memory
    free(tauq);              // free host memory
    free(taup);              // free host memory
    free(diag);              // free host memory
    magma_free_pinned(r);    // free host memory
    magma_free_pinned(h_work); // free host memory
    magma_finalize( );        // finalize Magma
    return EXIT_SUCCESS;
}
//dgebrd gpu time: 3.54390 sec.
//dgebrd cpu time: 29.55658 sec.

```

4.8.8 magma_dgebrd - unified memory version

```

#include <stdio.h>
#include <cuda.h>
#include <cuda_runtime.h>
#include "magma_v2.h"
#include "magma_lapack.h"
#define min(a,b) (((a)<(b))? (a):(b))
int main( int argc, char** argv){

```

```

magma_init(); // initialize Magma
double gpu_time, cpu_time;
magma_int_t m = 4096, n = 4096, n2=m*n;
double *a, *r; // a,r - mxn matrices
double *h_work; // workspace
magma_int_t lhwork; // size of h_work
double *taup, *tauq; // arrays describ. elementary reflectors
double *diag, *offdiag; // bidiagonal form in two arrays
magma_int_t info, minmn=min(m,n), nb;
magma_int_t ione = 1;
magma_int_t ISEED[4] = {0,0,0,1}; // seed
nb = magma_get_dgebrd_nb(m,n); // optimal block size
cudaMallocManaged(&a,n2*sizeof(double)); //unif. memory for a
cudaMallocManaged(&r,n2*sizeof(double)); //unif. memory for r
cudaMallocManaged(&tauq,minmn*sizeof(double)); //mem.for tauq
cudaMallocManaged(&taup,minmn*sizeof(double)); //mem.for taup
cudaMallocManaged(&diag,minmn*sizeof(double)); //mem.for diag
cudaMallocManaged(&offdiag,(minmn-1)*sizeof(double)); //unif.
lhwork = (m + n)*nb; //memory for offdiag
cudaMallocManaged(&h_work,lhwork*sizeof(double)); //workspace
// Randomize the matrix a
lapackf77_dlarnv(&ione, ISEED, &n2, a );
lapackf77_dlacpy(MagmaFullStr,&m,&n,a,&m,r,&m); // a->r
// MAGMA
gpu_time = magma_wtime();
// reduce the matrix r to upper bidiagonal form by orthogonal
// transformations: q^T*r*p, the obtained diagonal and the
// superdiagonal are written to diag and offdiag arrays resp.;
// the elements below the diagonal, represent the orthogonal
// matrix q as a product of elementary reflectors described
// by tauq; elements above the first superdiagonal represent
// the orthogonal matrix p as a product of elementary reflect-
// ors described by taup;

magma_dgebrd(m,n,r,m,diag,offdiag,tauq,taup,h_work,lhwork,
             &info);

gpu_time = magma_wtime() - gpu_time;
printf("dgebrd gpu time: %7.5f sec.\n",gpu_time);
// LAPACK
cpu_time = magma_wtime();
lapackf77_dgebrd(&m,&n,a,&m,diag,offdiag,tauq,taup,h_work,
                &lhwork,&info);

cpu_time = magma_wtime() - cpu_time;
printf("dgebrd cpu time: %7.5f sec.\n",cpu_time);
// free memory
magma_free(a); // free memory
magma_free(tauq); // free memory
magma_free(taup); // free memory
magma_free(diag); // free memory
magma_free(r); // free memory
magma_free(h_work); // free memory

```

```
    magma_finalize( );                                // finalize Magma
    return EXIT_SUCCESS;
}
//dgebrd gpu time: 3.77582 sec.
//dgebrd cpu time: 28.54319 sec.
```

Bibliography

- [CUBLAS] *CUBLAS LIBRARY User Guide*, Nvidia, June 2017
http://docs.nvidia.com/cuda/pdf/CUBLAS_Library.pdf
- [CUSOLVER] *CUSOLVER LIBRARY*, Nvidia , June 2017
http://docs.nvidia.com/cuda/pdf/CUSOLVER_Library.pdf
- [MAGMA] *MAGMA Users' Guide* , Univ. of Tennessee, Knoxville, Univ. of California, Berkeley, Univ. of Colorado, Denver, November 2016
<http://icl.cs.utk.edu/projectsfiles/magma/doxygen/>
- [ARRAYFIRE] Chrzęszczuk A., *Matrix Computations on the GPU with ArrayFire-Python and ArrayFire-C/C++*. Version 2017, July 2017
https://www.researchgate.net/publication/319135914_Matrix_Computations_on_GPU_with_ArrayFire-Python_and_ArrayFire-CCVersion_2017
- [FUND] Watkins D. S., *Fundamentals of Matrix Computations*, 2nd ed., John Willey & Sons, New York 2002
- [MATR] Golub G. H, van Loan C. F., *Matrix Computations*, 3rd ed. Johns Hopkins Univ. Press, Baltimore 1996
- [LAPACK] Anderson E., Bai Z., Bischof C., Blackford S., Demmel J., Dongarra J., et al *LAPACK Users' Guide*, 3rd ed., August 1999
<http://www.netlib.org/lapack/lug/>